

HBASE CON2012

Supporting HBase: How to Stabilize, Diagnose and Repair

Jeff Bean, Jonathan Hsieh, Kathleen Ting
{jwfbean,jon,kathleen}@cloudera.com

5/22/12

Who Are We?

- Jeff Bean
 - Designated Support Engineer, Cloudera
 - Education Program Lead, Cloudera
- Kathleen Ting
 - Support Manager, Cloudera
 - ZooKeeper Subject Matter Expert
- Jonathan Hsieh
 - Software Engineer, Cloudera
 - Apache HBase Committer and PMC member



Outline

- Preventative HBase Medicine:
 - Tips for a healthy HBase
- The HBase Triage:
 - Fixes for acute HBase pains
- The HBase Surgery:
 - Repairing a Corrupted HBase



Outline

- Preventative HBase Medicine:
 - Tips for a healthy HBase
- The HBase Triage:
 - Fixes for acute HBase pains
- The HBase Surgery:
 - Repairing a Corrupted HBase




“Monitor your system, exercise your workload, and eat your vegetables.”

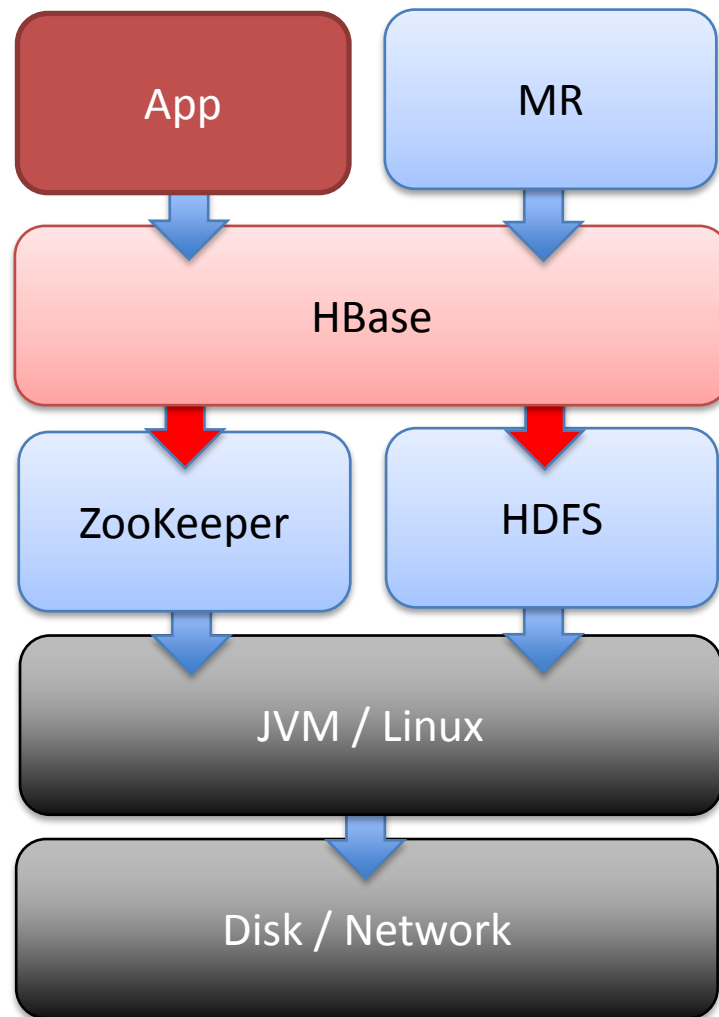


HAVING SAID THAT,
LET'S BEGIN THERAPY.
READY ?





HBase Cross-Section



Doctor's Advice: "A ounce of prevention worth a pound of cure."




- Understand your workload and test for it
- Size your cluster properly (see Cluster Sizer)
- Monitor, alert, and manage your cluster with Ganglia, Nagios, and/or Cloudera Manager
 - Don't be Dr. House!

A Case Study



Symptom: Long Running MapReduce job with blacklisted TaskTrackers


TaskTracker	No. of Failures
NodeX	4
NodeY	3
NodeQ	7
NodeB	10
NodeP	8
NodeV	6



Symptom: Node B Task Logs


```
$ find . | xargs grep "giving up"
```

```
./attempt_201107261334_0221_m_000962_1/syslog:2011-08-02  
11:09:34,248 INFO org.apache.hadoop.ipc.HbaseRPC: Server at  
NodeA:60020 could not be reached after 1 tries, giving up.  
./attempt_201107261334_0221_m_000962_1/syslog:2011-08-02  
11:09:37,328 INFO org.apache.hadoop.ipc.HbaseRPC: Server at  
NodeA:60020 could not be reached after 1 tries, giving up.  
./attempt_201107261334_0221_m_000962_1/syslog:2011-08-02  
11:09:40,465 INFO org.apache.hadoop.ipc.HbaseRPC: Server at  
NodeA:60020 could not be reached after 1 tries, giving up.
```




Symptom: RegionServer logs of Node A:

2011-08-02 11:04:20,324 FATAL
org.apache.hadoop.hbase.regionserver.HRegionServer
: ABORTING region server
serverName=**NodeA**,60020,1312228900706,
load=(requests=10847, regions=342, usedHeap=8193,
maxHeap=15350): regions
erver:60020-0x4316487a73e1626 regionserver:60020-
0x4316487a73e1626 received expired from **ZooKeeper**,
aborting
org.apache.zookeeper.KeeperException\$SessionExpiredEx
ception: KeeperErrorCode = Session expired





Cascading failure! Some other node says ouch...



```
2011-08-01 12:55:39,356 INFO org.apache.hadoop.hbase.regionserver.ShutdownHook:
Shutdown hook starting; hbase.shutdown.hook=true; fsShutdownHook=Thread[Thread-
15,5,main]
2011-08-01 12:55:39,629 INFO org.apache.hadoop.hbase.regionserver.HRegionServer:
STOPPED: Shutdown hook
2011-08-01 12:55:39,629 INFO org.apache.hadoop.hbase.regionserver.ShutdownHook:
Starting fs shutdown hook thread.
2011-08-01 12:55:39,695 ERROR org.apache.hadoop.hdfs.DFSCClient: Exception closing file
/hbase/.logsNodeA,60020,1311651881177NodeA%3A60020.1311656326143 :
java.io.IOException: Error Recovery for block blk_1102151039331207284_16350929
failed because recovery from primary datanode NodeA:50010 failed 6 times. Pipeline
was NodeA:50010. Aborting...
java.io.IOException: Error Recovery for block blk_1102151039331207284_16350929 failed
because recovery from primary datanode NodeA:50010 failed 6 times. Pipeline was
NodeA:50010. Aborting...
at
org.apache.hadoop.hdfs.DFSCClient$DFSOutputStream.processDatanodeError(DFSCli
ent.java:2841)
at
org.apache.hadoop.hdfs.DFSCClient$DFSOutputStream.access$1600(DFSCClient.java:2
305)
at
org.apache.hadoop.hdfs.DFSCClient$DFSOutputStream$DataStreamer.run(DFSCClient.j
ava:2477)
2011-08-01 12:55:39,842 INFO org.apache.hadoop.hbase.regionserver.ShutdownHook:
Shutdown hook finished.
```



Symptom: Ganglia Memory Graph on Node A...

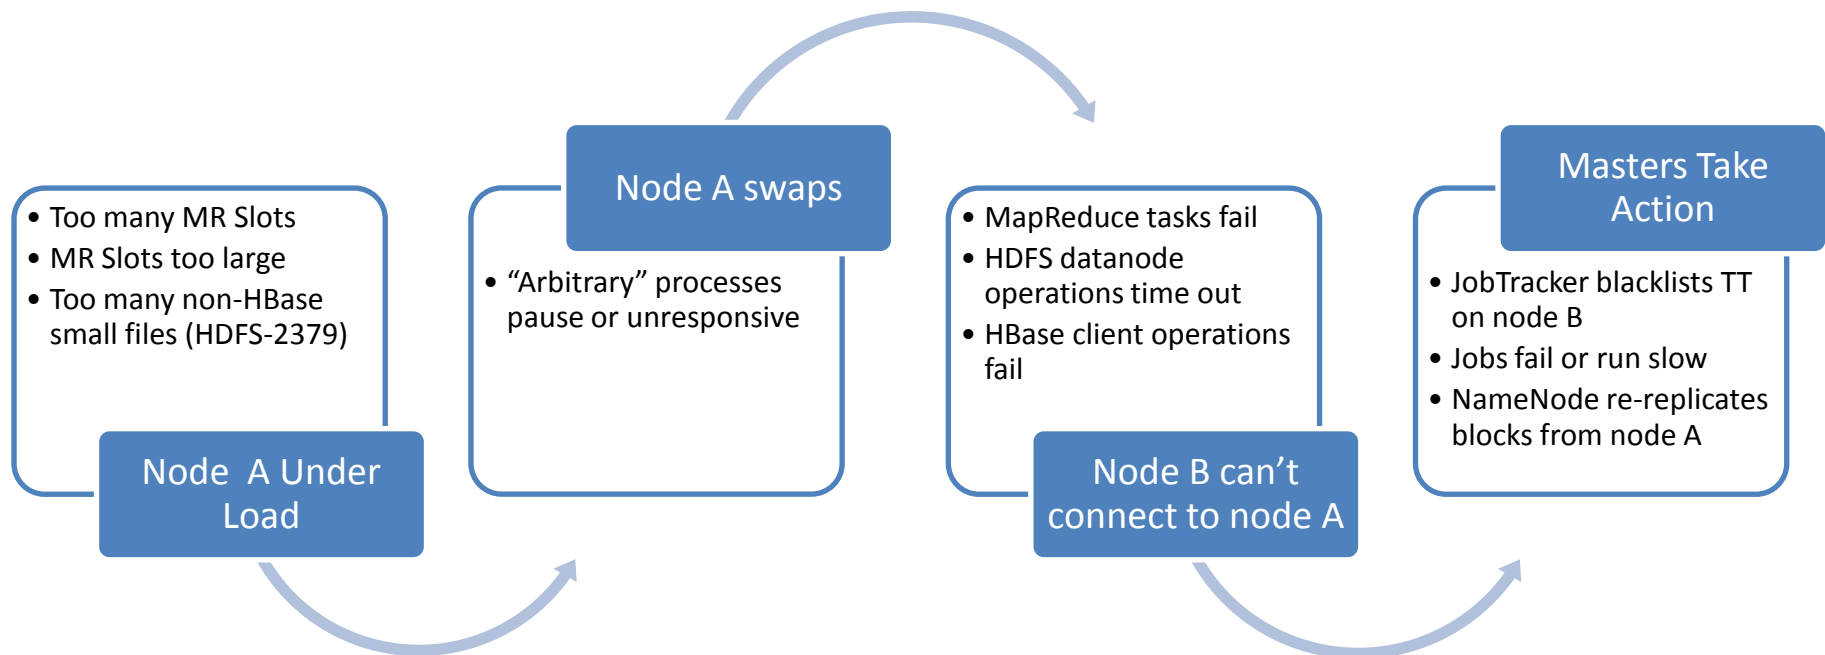


Symptom: Ganglia swap_free on Node A...

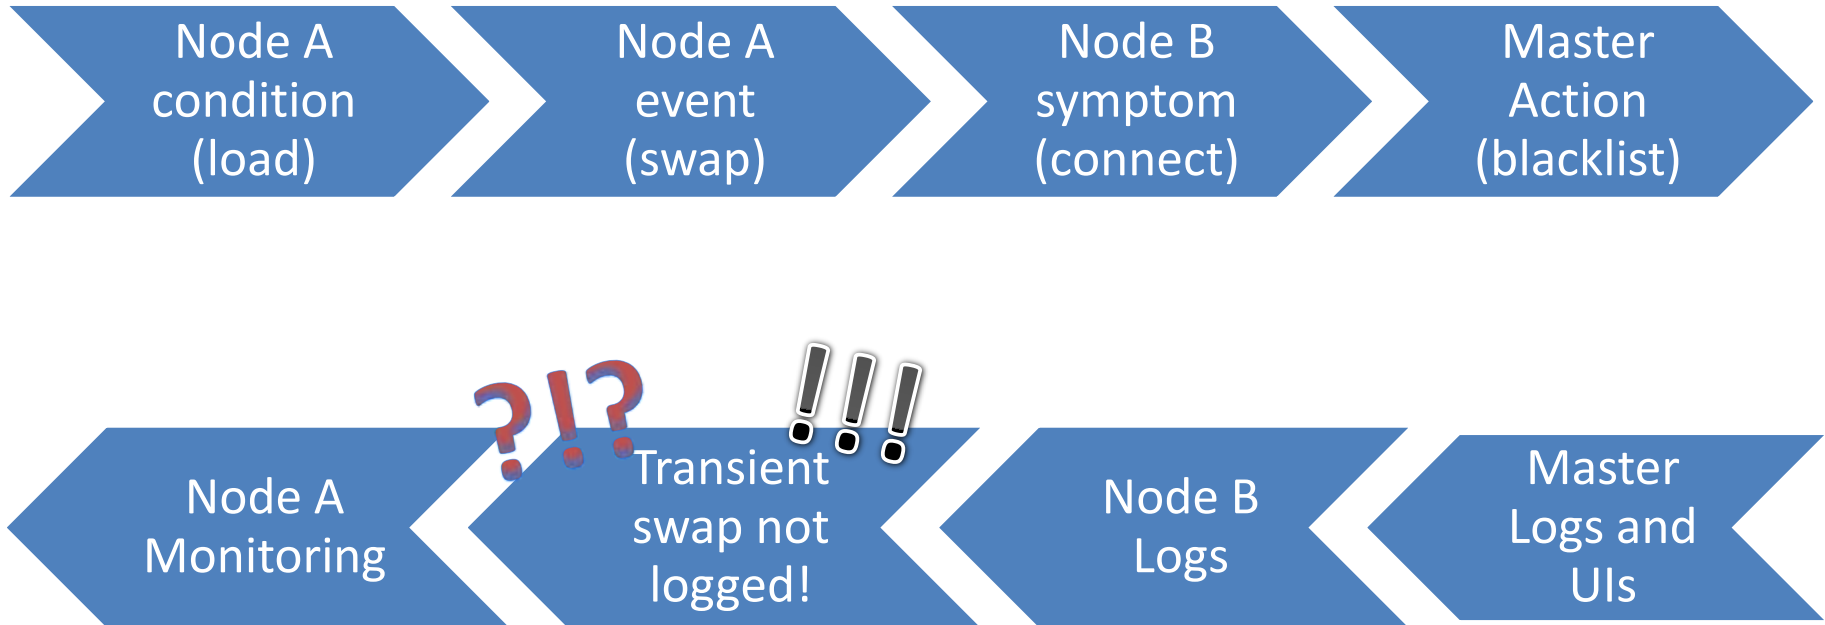


A Case study: Radiant Pain

“I was having back pains, and it turned out to be my heart!”



Event Trail and Evidence Trail



DOs and DON'Ts for keeping HBase Healthy

DOs

- Monitor and Alert
- Optimize network
- Know your logs

DON'Ts

- Swap
- Oversubscribe MR
- Share the network



Outline


- Preventative HBase Medicine:
 - Tips for a healthy HBase
- The HBase Triage:
 - Fixes for acute HBase pains
- The HBase Surgery:
 - Repairing a Corrupted HBase



“Cloudera 911 here, how can we help?”



HBase Support Tickets



Understanding the logs helps us diagnose issues

- Related events logged by different processes in different places
- Log messages point at each other
 - HDFS accesses by RS logged by NN and DN
 - HBase accesses by MR logged by JT, RS, NN, ZK
 - ZK logs indicate HBase health

The HBase Triage: Fixes for acute HBase pains

- Severe Pain
- Complete Unconsciousness

The HBase Triage: Fixes for acute HBase pains

- Severe Pain
- Complete Unconsciousness

Connection Reset


WARN - Session <id> for server <server id>, unexpected error, closing socket connection and attempting reconnect
java.io.IOException: Connection reset by peer

What causes this?

- Running out of ZK connections

How can it be resolved?

- Manually close connections
- Fixed in HBASE-5466 and HBASE-4773



Running out of DN Threads & File Descriptors

INFO hdfs.DFSCClient: Could not obtain block <blk id> from any node: java.io.IOException: No live nodes contain current block.


ERROR java.io.IOException: Too many open files

What causes this?

- HBase likes to keep data files open

How can it be resolved?

- Increase dfs.datanode.max.xcievers to 4096
- Increase /etc/security/limits.conf
 - hbase - nofile 32768




“Long Garbage Collecting Pause”


WARN org.apache.hadoop.hbase.util.Sleeper:
We slept 19118ms instead of 1000ms, this is
likely due to a long garbage collecting pause and
it's usually bad

How can it be resolved?

- zoo.cfg: maxSessionTimeout=180000
hbase-site.xml: zookeeper.session.timeout=180000
- Oversubscribed if MR & HBase are co-located



Heap Allocation Per Node



The HBase Triage: Fixes for acute HBase pains

- Severe Pain
- Complete Unconsciousness

ZK can't start & HBase hangs


INFO org.apache.hadoop.hdfs.DFSCClient: Could not complete file <name> retrying...

What causes this?

- High dfs.replication.min causes HBase hang - can't close file until created all replicas

How can it be resolved?

- Remove dfs.replication.min
- Temp increase dfs.balance.bandwidthPerSec
- Fixed in HDFS-2936



Unable to Load Database

FATAL


org.apache.zookeeper.server.quorum.QuorumPeer: Unable to load database on disk

What causes this?

- ZK data directories filled up

How can it be resolved?

- Wipe out /var/zookeeper/version-2
- Run zkCleanup.sh script via cron



Downed HBase Master and RegionServers

WARN

org.apache.zookeeper.server.quorum.Learner:

Exception when following the leader


java.net.SocketTimeoutException: Read timed out

What causes this?

- Session Timeout + Session Expiration = NW Prob

How can it be resolved?

- Monitor network (e.g. ifconfig)
- Run ≥ 3 ZK servers (majority rules)



The HBase Triage: Fixes for acute HBase pains

- Severe Pain
- Complete Unconsciousness

Outline


- Preventative HBase Medicine:
 - Tips for a healthy HBase
- The HBase Triage:
 - Fixes for acute HBase pains
- The HBase Surgery:
 - Repairing a Corrupted HBase



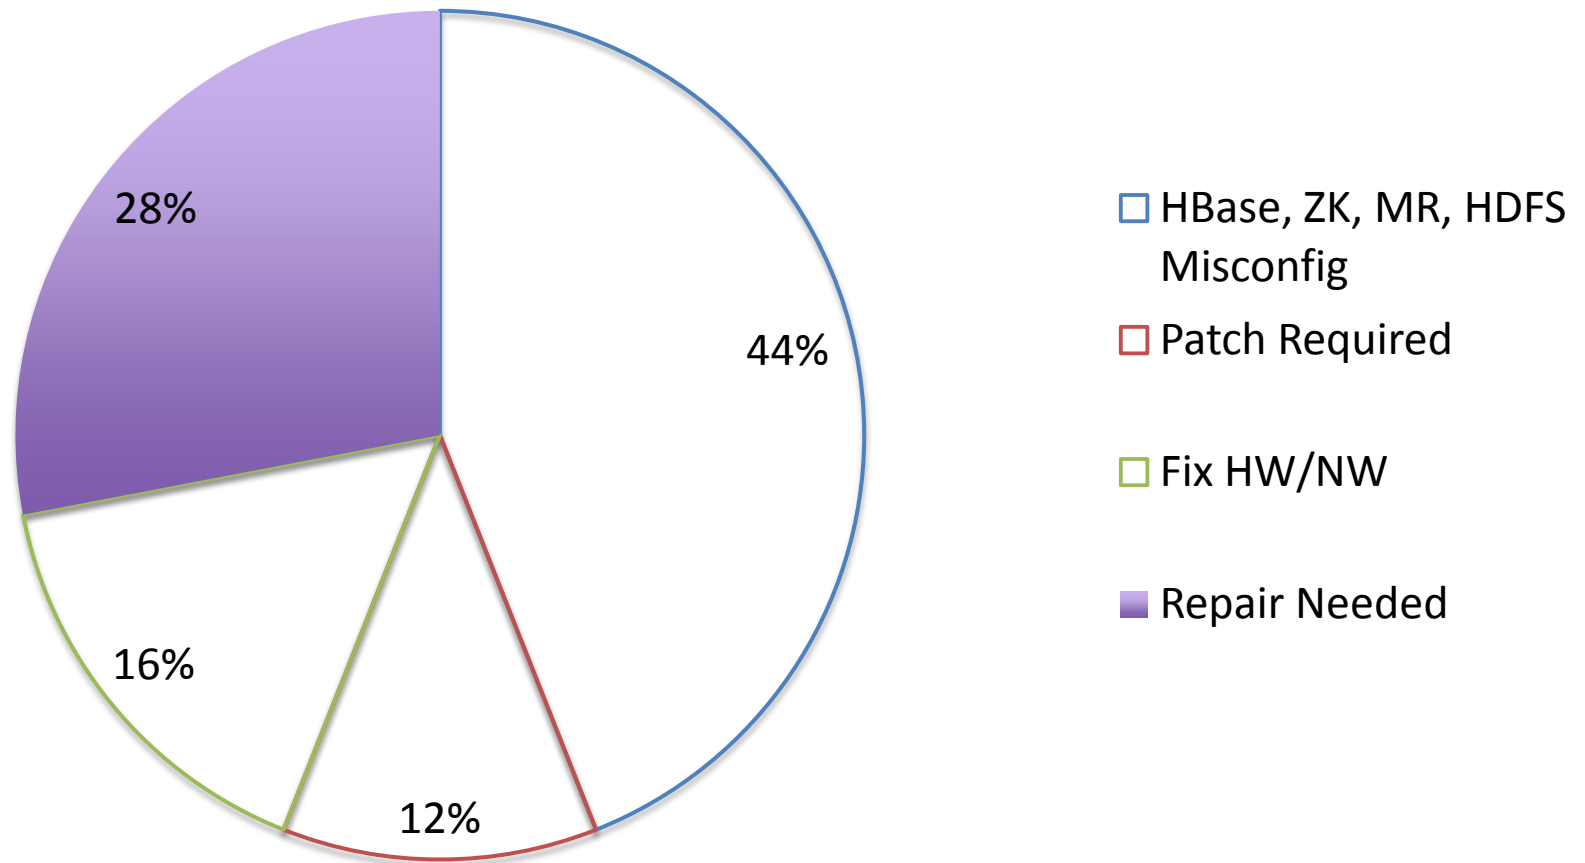
“To the operating room, please”



HBase Support Tickets



HBase Support Tickets



Detecting internal problems with hbck

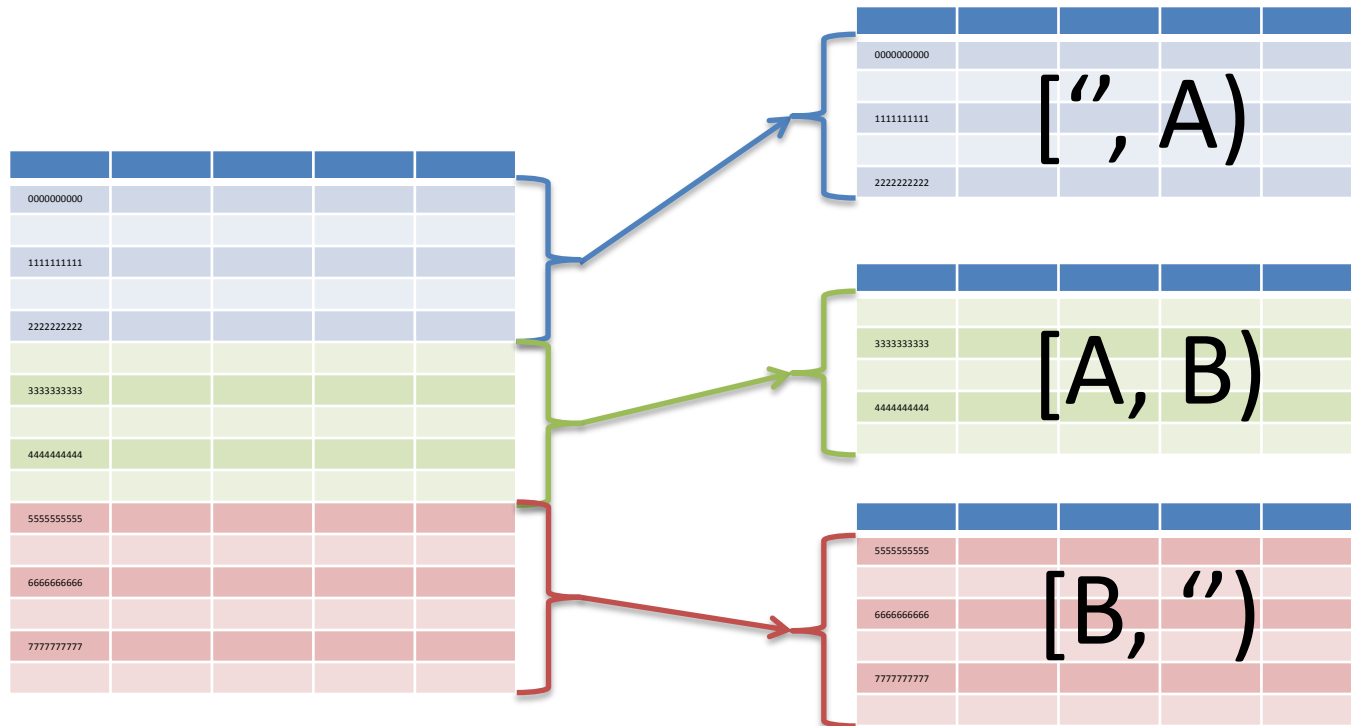
- HBase since 0.90 has included a tool for scanning an HBase instance's internals to find corruptions.

`hbase hbck`

`hbase hbck -details`



Tables are sharded into regions



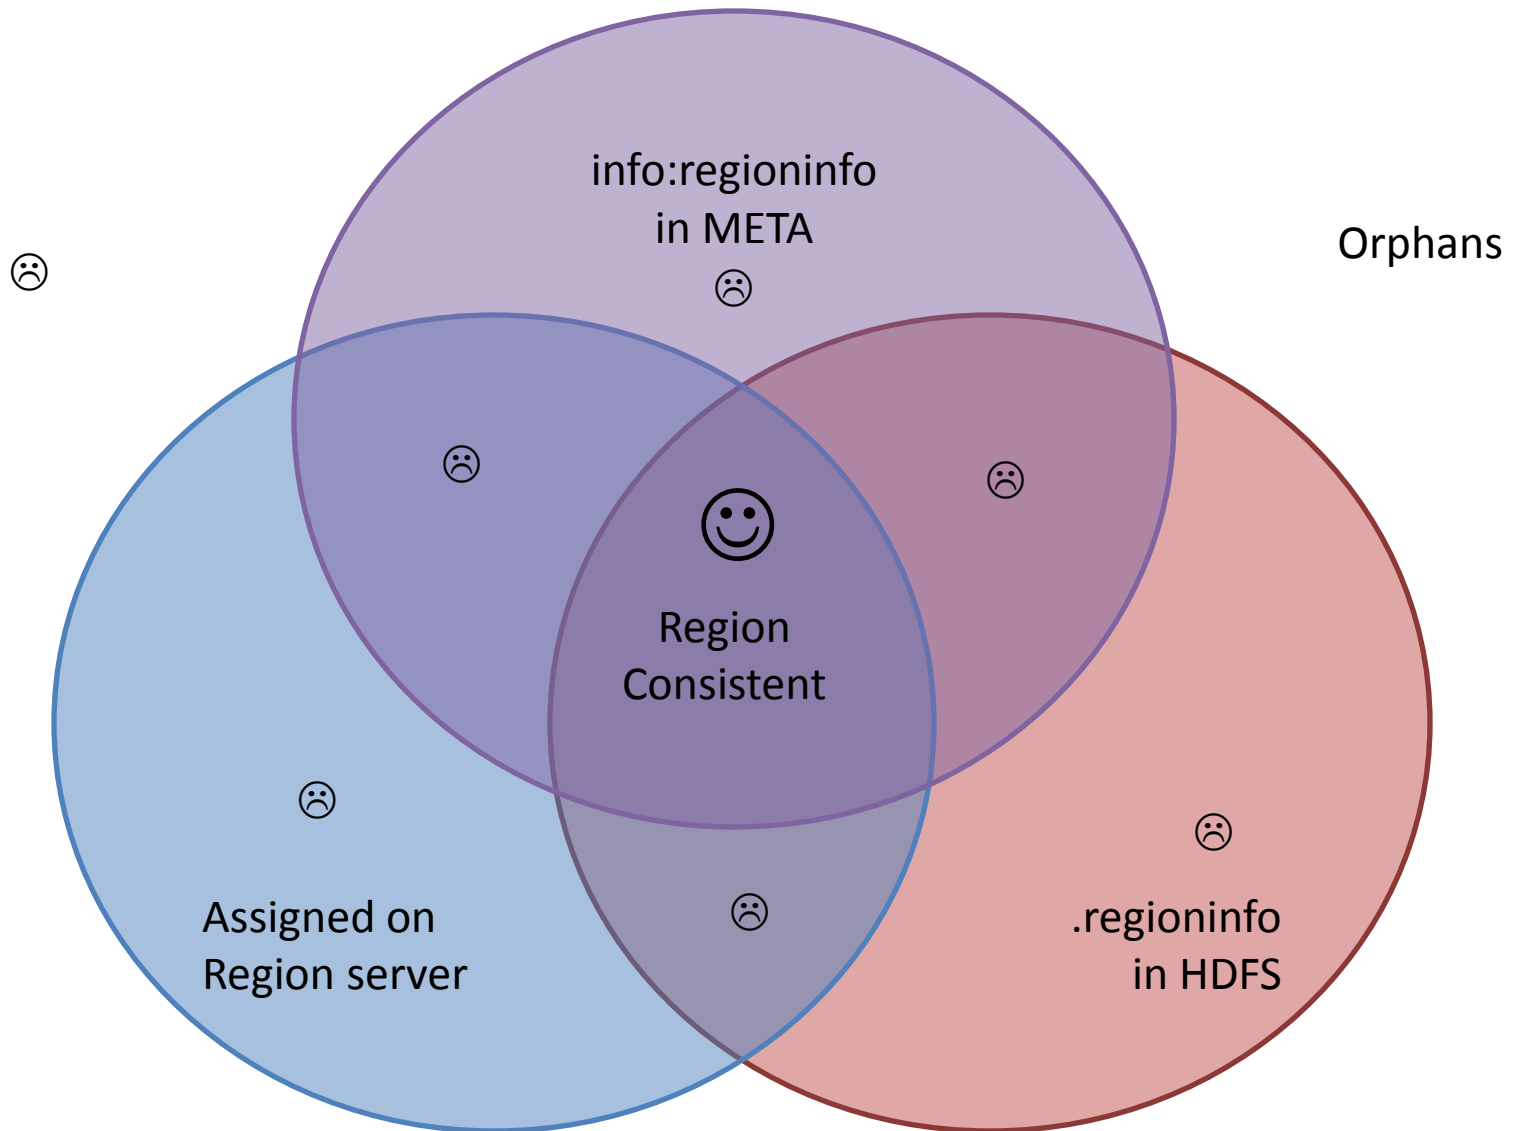
Invariants: Maintain **table integrity** and **region consistency** !

Table Integrity Invariants

[' ,A)
[A,B)
[B, C)
[C, D)
[D, E)
[E, F)
[F, G)
[G, ' ')


- Every key shall get assigned to a single region.
- Table Regions shall:
 - Cover the entire range of possible keys,
 - from the absolute start (“)
 - to the absolute end (unfortunately, also “).

Region Consistency Invariants

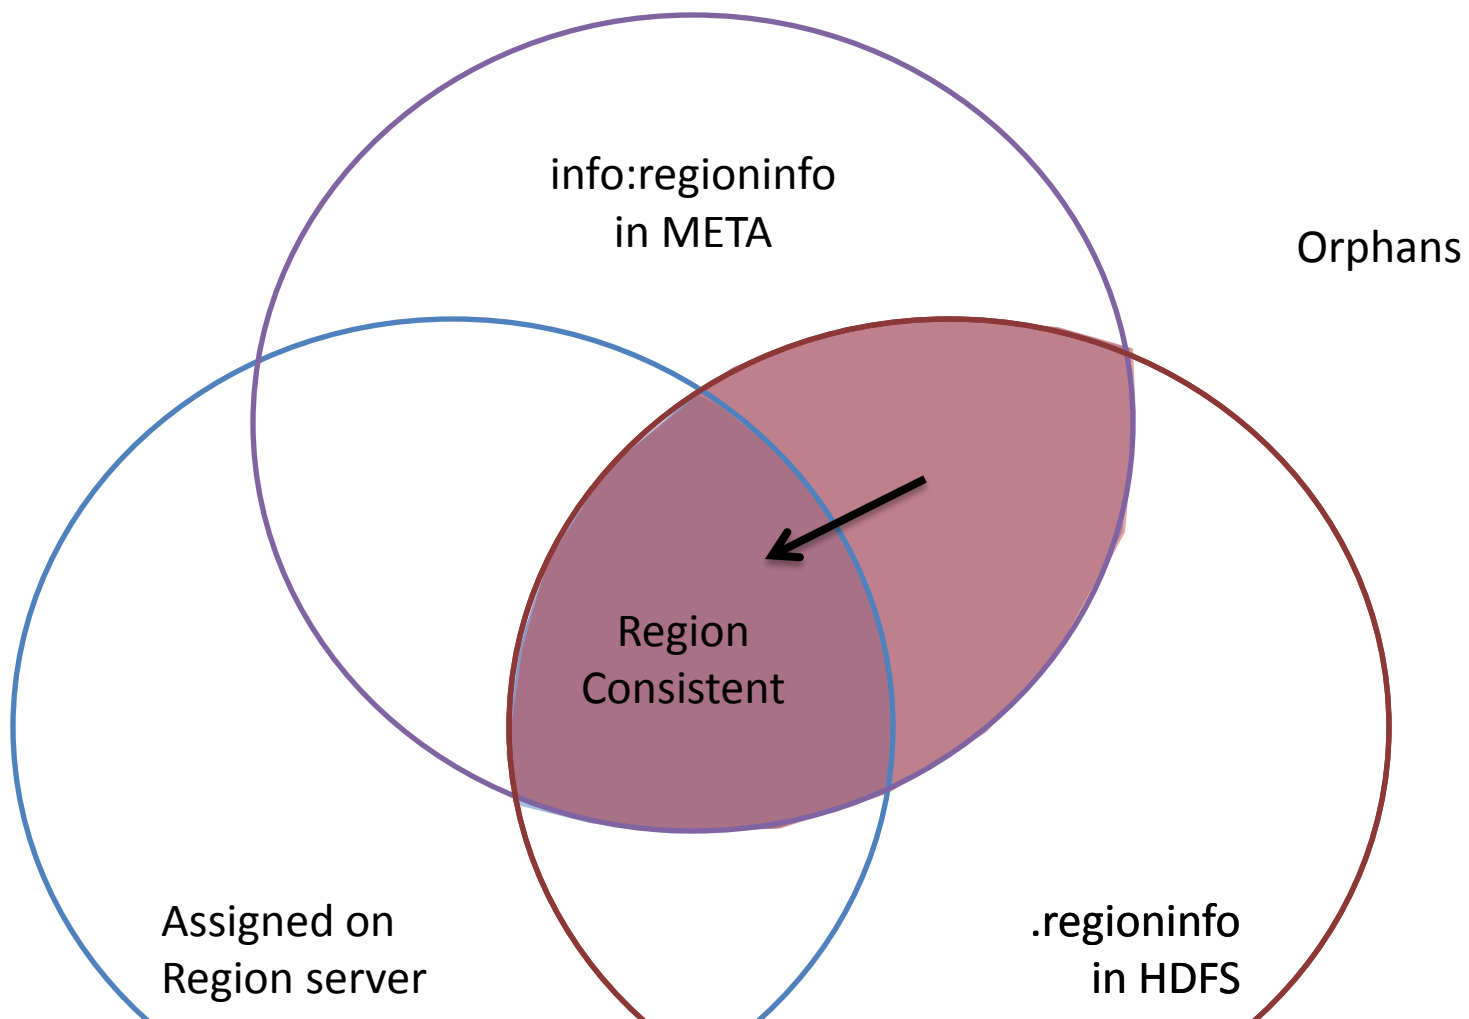


Repairing internal problems with hbck

- Newer and upcoming versions of HBase include an hbck that can **fix** internal problem as well as detect.
 - 0.90.7
 - 0.92.2
 - 0.94.0
 - CDH3u4+
 - CDH4b2+



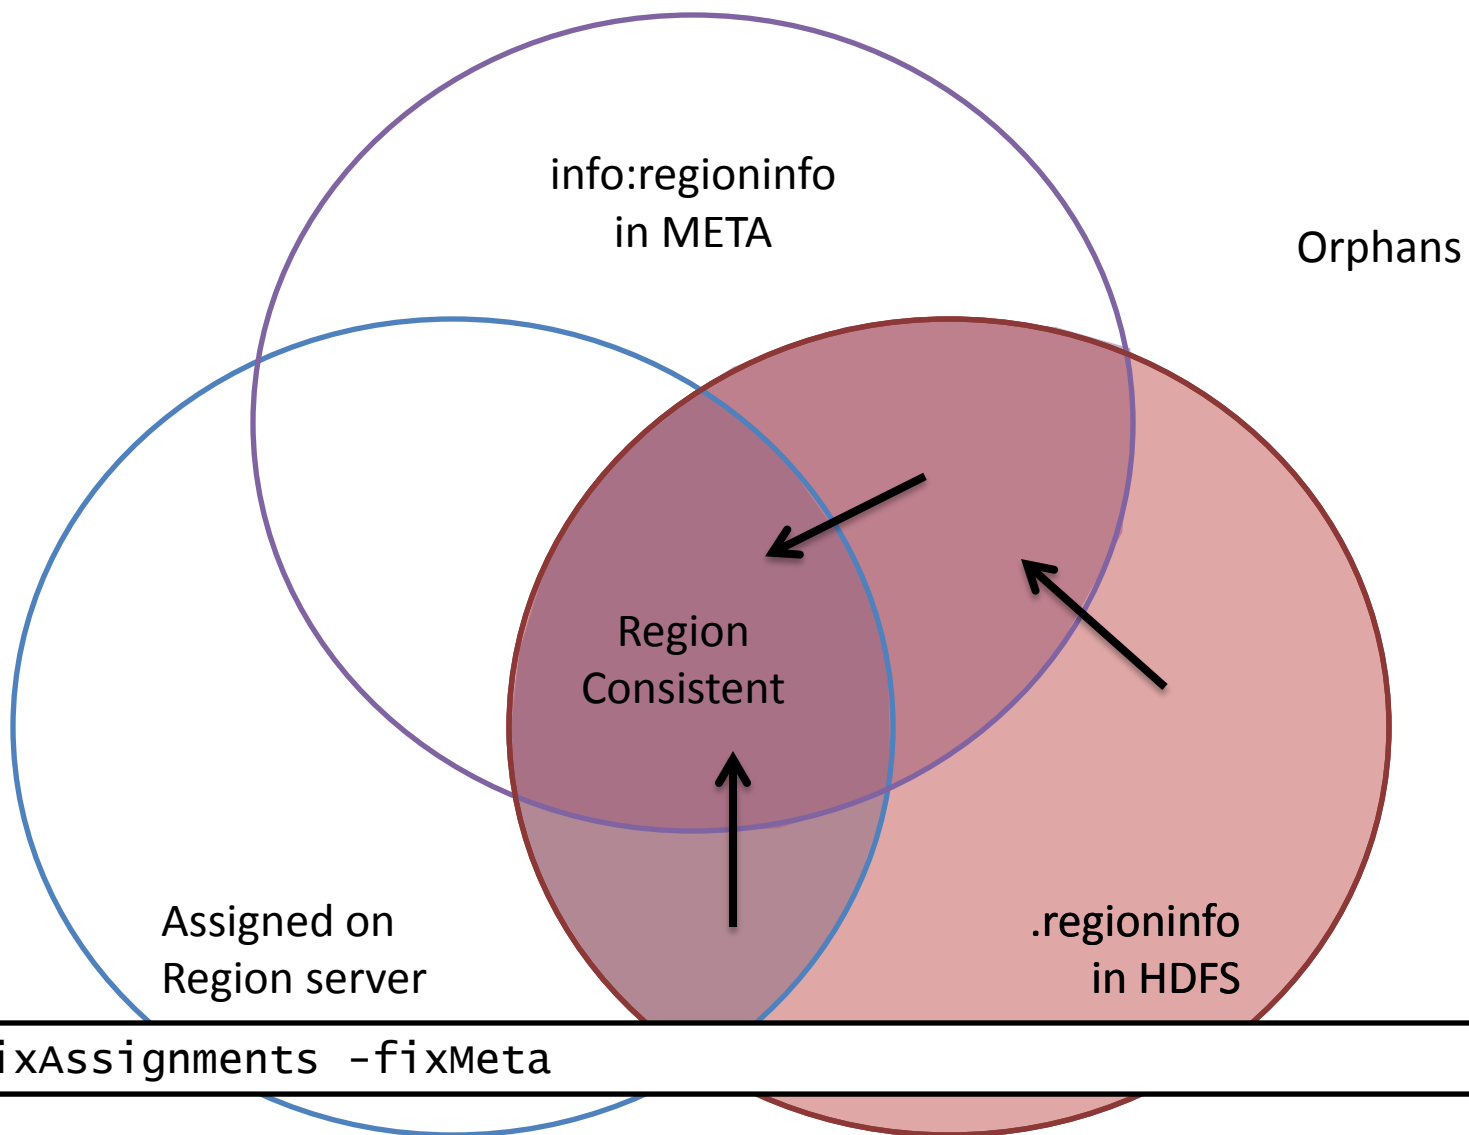
Bad region assignment



```
hbck -fix (0.90.x)
```

```
hbck -fixAssignments (0.90.7+, 0.92.2+, 0.94+)
```

Region not in META



Regioninfo not in HDFS

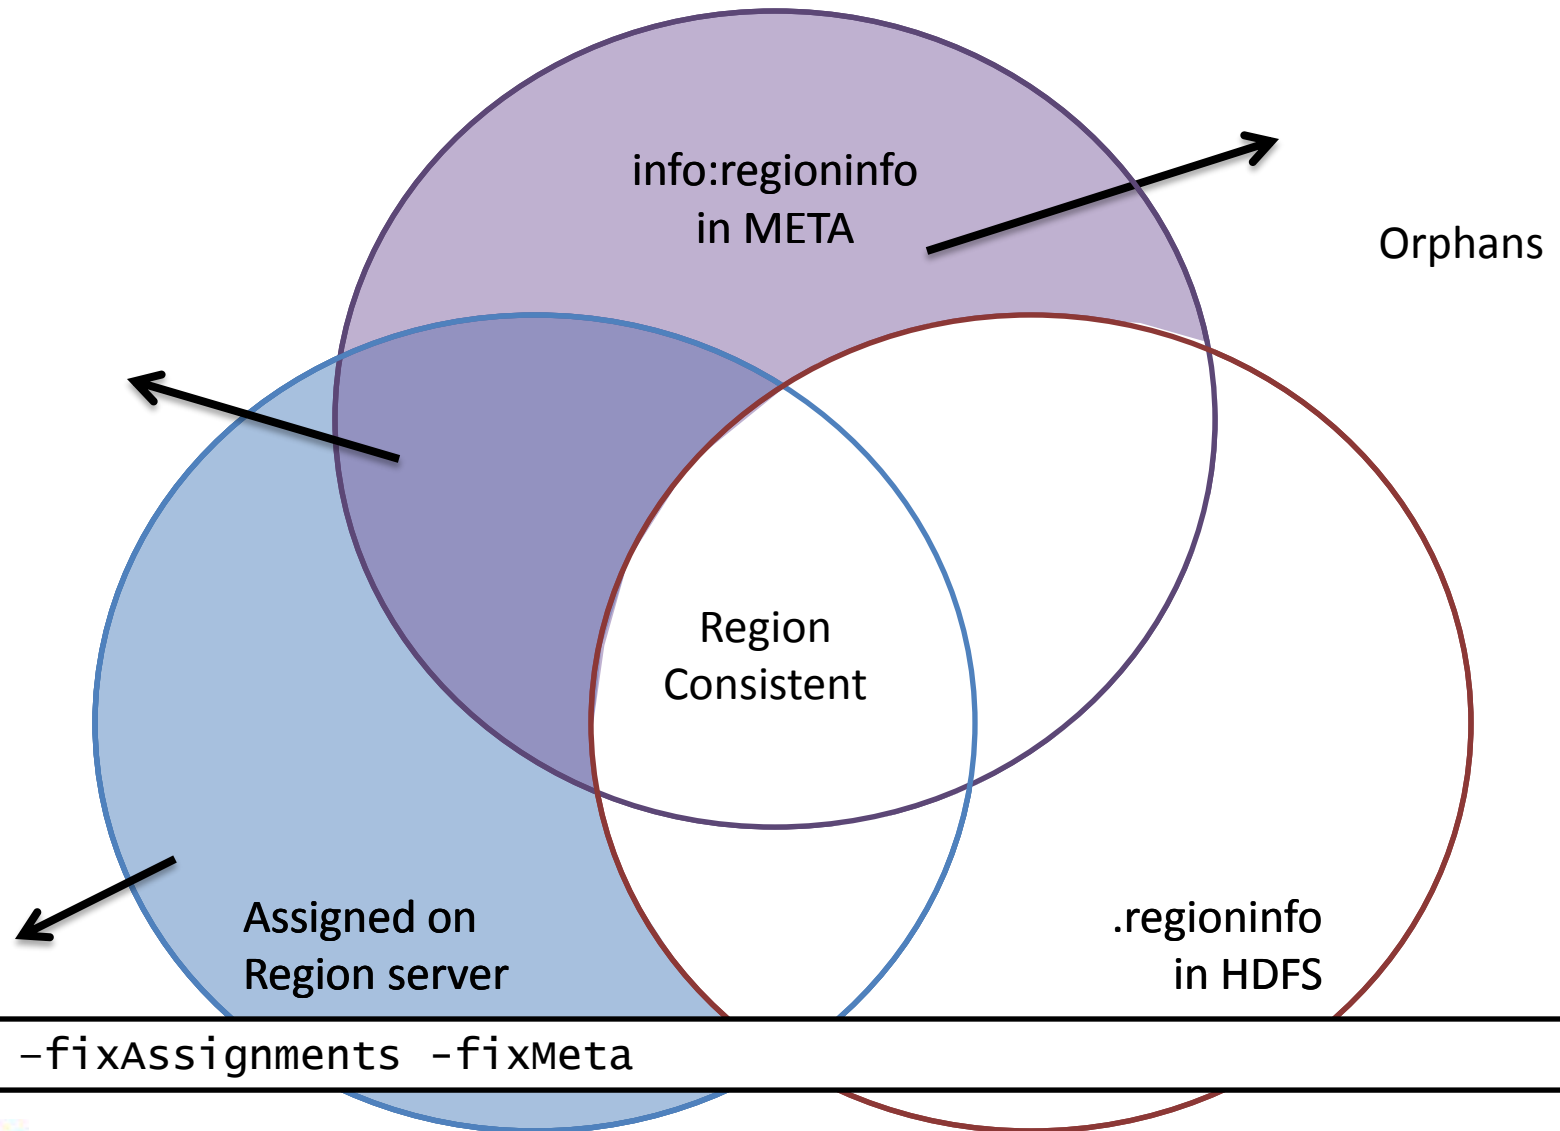



Table Regions must not have holes

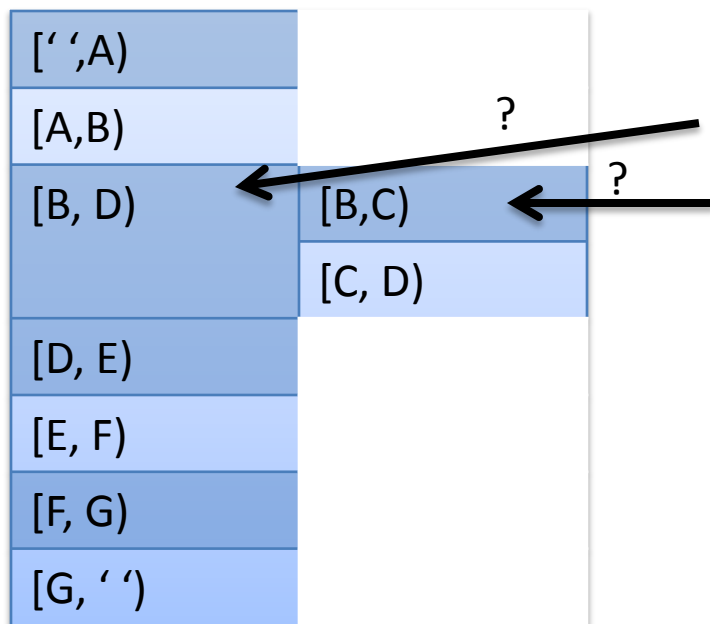
[' ,A)
[A,B)
[B, C)
[D, E)
[E, F)
[F, G)
[G, ')



- Where to I put row key “CRUD”?
- Where is region [C,D)?
- Repair:
 - Find the orphan and adopt it.
 - Fabricate a new region to fill the hole

```
# NOTE! HBase should be idle (no get/put/split/compacts)
hbck -fixHdfsHoles -fixHdfsOrphans -fixAssignments -fixMeta
```

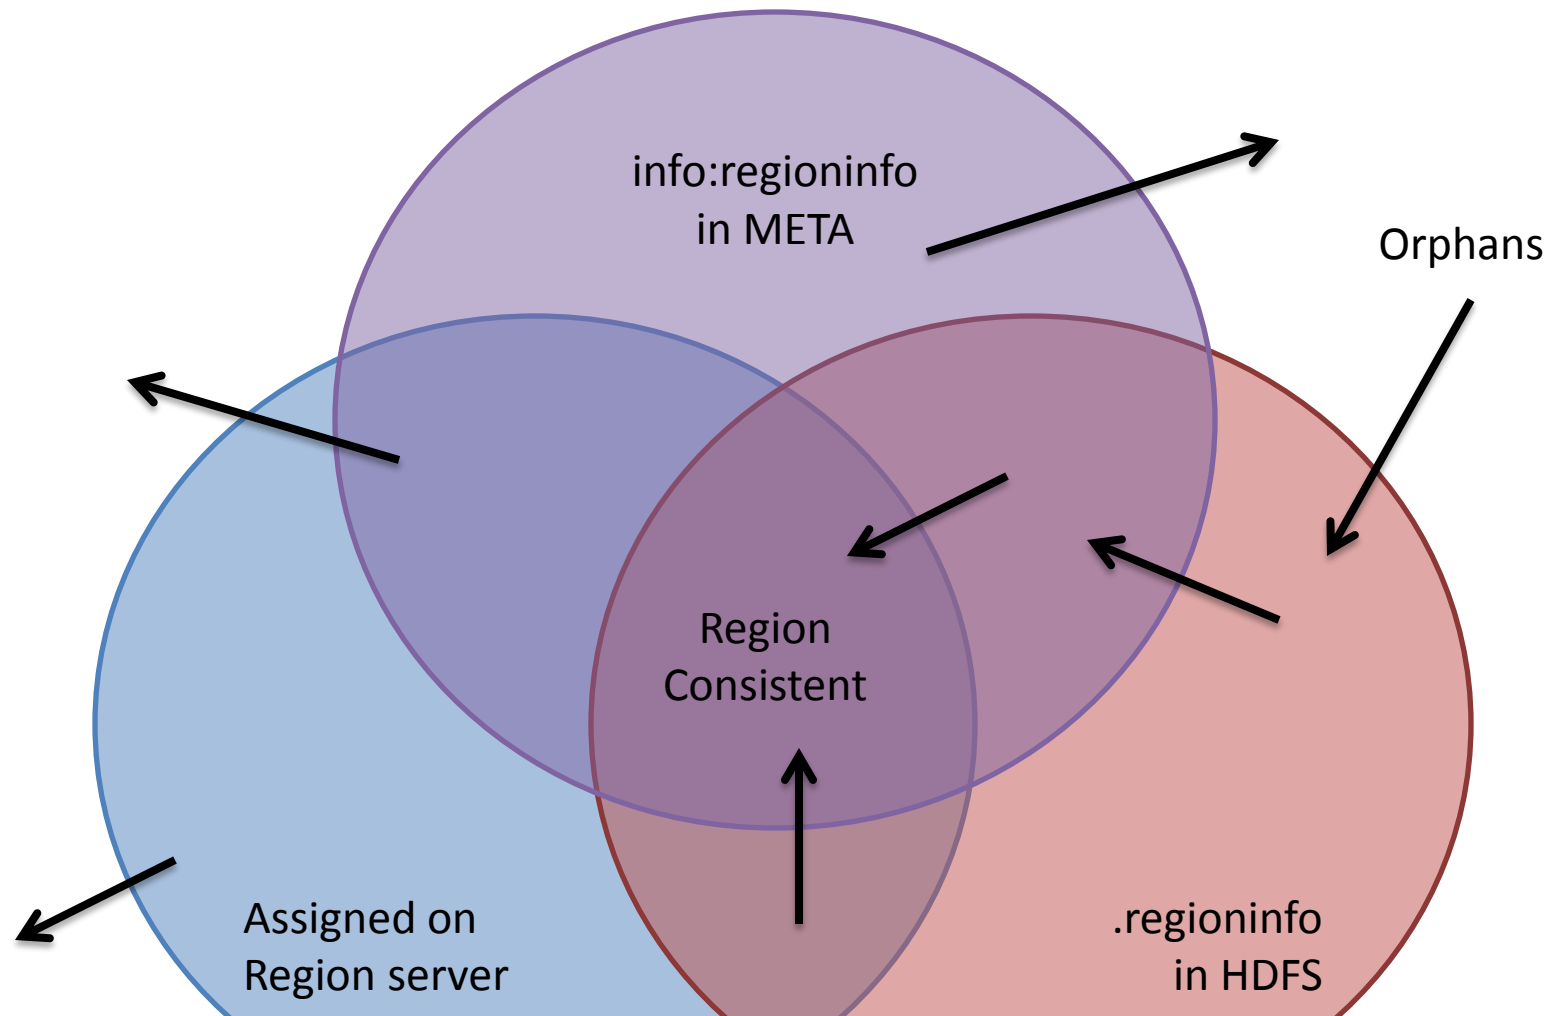
Table Regions must not overlap



- Hm.. Which region should “BAD” go?
- Is it [B, D) or is it [B,C)?
- Likely due to a bad split.
- Repair:
 - Merge regions or,
 - Sideline and bulk load

```
# NOTE! HBase should be idle (no get/put/split/compacts)
hbck -fixHdfsOverlaps -fixAssignments -fixMeta
```


Consistency problem summary



```
hbck -fixAssignments -fixMeta -fixHdfsHoles -fixHdfsOrphans -  
fixHdfsOverlaps
```

Investigating further

- HFile – examine contents of HFiles
- Hlog – examine contents of HLog file
- OfflineMetaRepair – Rebuild meta table from file system.
- Also, some scripts for manual repairs:
<https://github.com/jmhsieh/hbase-repair-scripts>

Outline

- Preventative HBase Medicine:
 - Tips for a healthy HBase
- The HBase Triage:
 - Fixes for acute HBase pains
- The HBase Surgery:
 - Repairing a Corrupted HBase



Questions?