




Pop-up Loft  
LONDON

# Deep Dive on Amazon RDS (Relational Database Service)


Richard Ainley, AWS Solutions Architect

18<sup>th</sup> September 2017



# Agenda

- Quick intro to RDS
- Security
- Metrics and Monitoring
- High Availability
- Scaling
- Backups and Snapshots
- Migrating



# Amazon Relational Database Service (RDS)

Launch



No infrastructure  
management



Cost-effective



Application  
compatibility



Instant provisioning



Scale up/down



Pop-up Loft  
LONDON


# Trade-offs with a managed service

## Fully managed host and OS

- No access to the database host operating system
- Limited ability to modify configuration that is managed on the host operating system
- No functions that rely on configuration from the host OS

## Fully managed storage

- Max storage limits
  - Microsoft SQL Server - 16 TB
  - MySQL, MariaDB, PostgreSQL, Oracle - 6 TB
  - Aurora - 64 TB



# Amazon RDS engines

## Commercial

ORACLE®



## Open source



PostgreSQL



MariaDB

## Cloud native



**Amazon** Aurora





Pop-up Loft  
LONDON



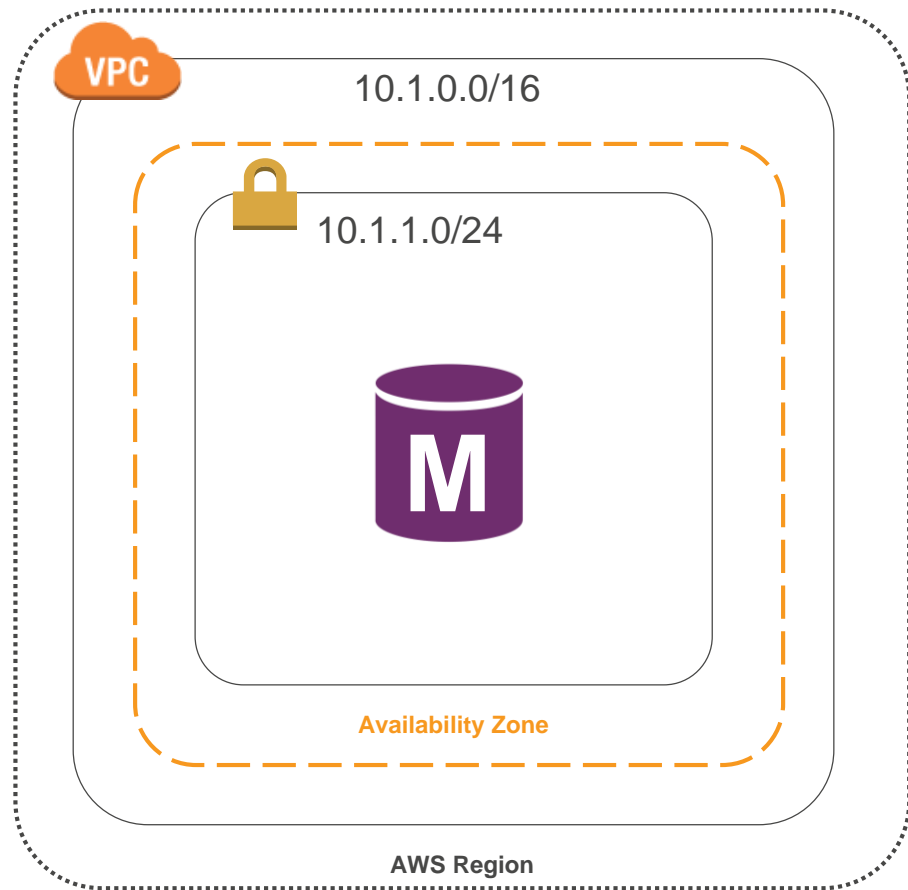
Pop-up Loft  
LONDON

# Security



# Amazon Virtual Private Cloud (Amazon VPC)


Securely control network configuration



# Amazon Virtual Private Cloud (Amazon VPC)

Securely control network configuration

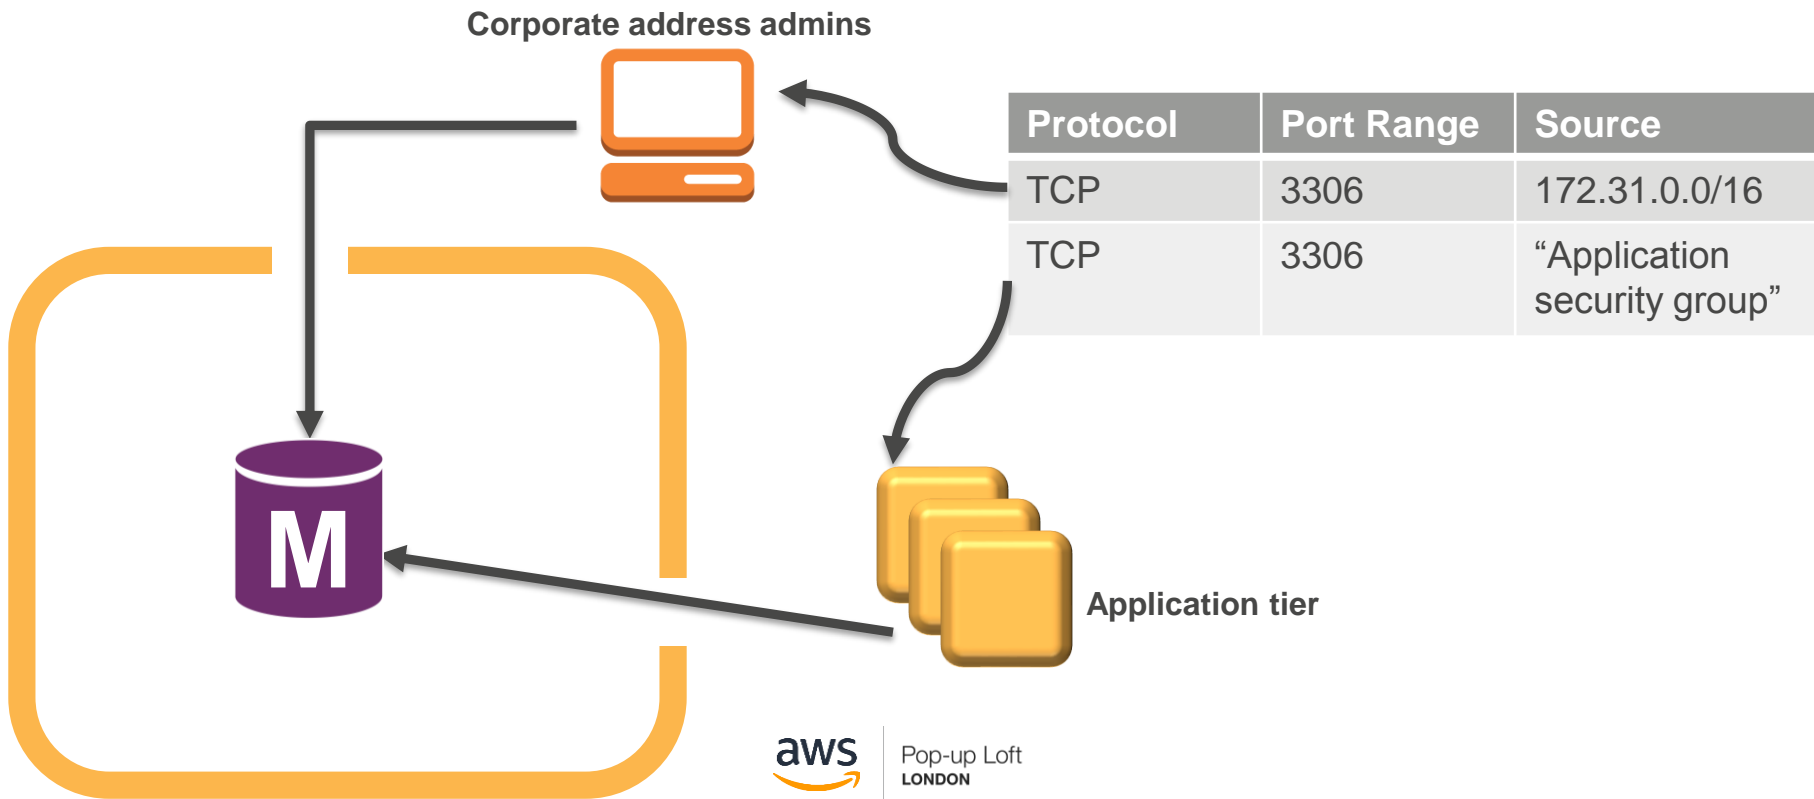
## Manage connectivity





# Security groups


## Database IP firewall protection




# AWS IAM governed access

You can use AWS Identity and Access Management (IAM) to:

- Control who can perform actions on RDS
- Authenticate to your RDS MySQL / Aurora DB
  - MySQL 5.6.34 / 5.7.16 or higher
  - Aurora 1.10 or higher
  - Not available for db.t1.micro / db.m1.small



# Compliance



Singapore MTCS



27001/9001  
27017/27018



Pop-up Loft  
LONDON

# Compliance

## Aurora

SOC 1, 2, 3  
ISO 20001/9001  
ISO 27107/27018  
PCI  
HIPAA BAA

## MySQL

SOC 1, 2, 3  
ISO 20001/9001  
ISO 27107/27018  
PCI  
FedRamp  
HIPAA BAA  
UK Gov. Programs  
Singapore MTCS

## Oracle

SOC 1, 2, 3  
ISO 20001/9001  
ISO 27107/27018  
PCI  
FedRamp  
HIPAA BAA  
UK Gov. Programs  
Singapore MTCS

## PostgreSQL

SOC 1, 2, 3  
ISO 20001/9001  
ISO 27107/27018  
PCI  
UK Gov. Programs  
Singapore MTCS  
HIPAA BAA

## MariaDB

SOC 1, 2, 3  
ISO 20001/9001  
ISO 27107/27018  
PCI  
HIPAA BAA


## SQL Server

SOC 1, 2, 3  
ISO 20001/9001  
ISO 27107/27018  
PCI  
UK Gov. Programs  
Singapore MTCS  
HIPAA BAA

# SSL

## Database traffic encryption

**Available for all six engines**





# **At Rest Encryption**

# Transparent Data Encryption



Pop-up Loft  
LONDON

# At Rest Encryption for all RDS Engines


## AWS Key Management Service (KMS)

Two-tiered key hierarchy using envelope encryption:

- Unique data key encrypts customer data
- AWS KMS master keys encrypt data keys
- Available for **ALL** RDS engines

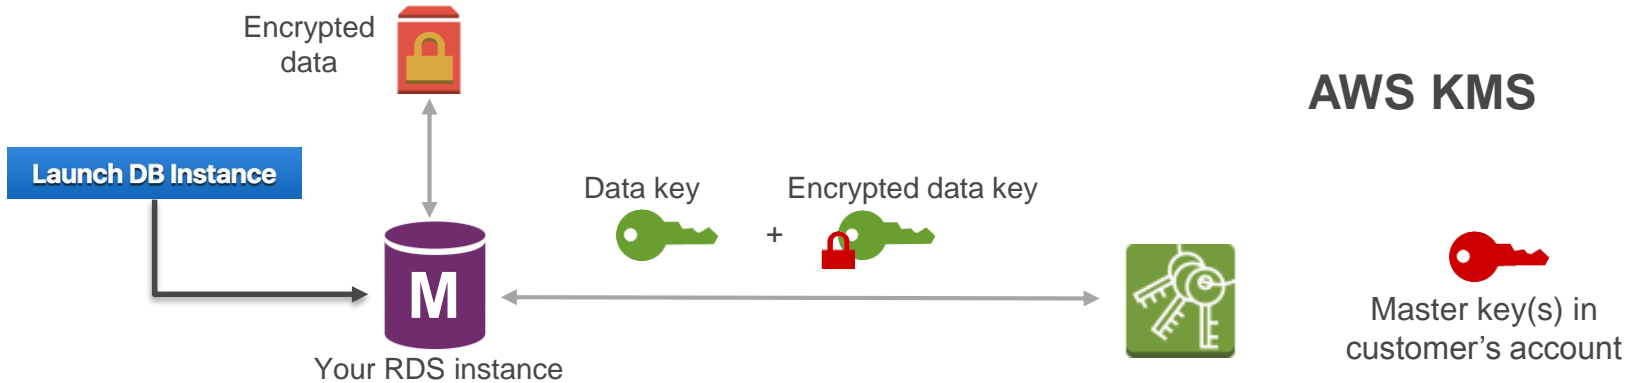
Benefits:

- Limits risk of compromised data key
- Better performance for encrypting large data
- Easier to manage small number of master keys than millions of data keys
- Centralized access and audit of key activity





# How keys are used to protect your data



1. Launch your RDS instance
2. RDS instance requests encryption key to use to encrypt data, passes reference to master key in account
3. Client request authenticated based on permissions set on both the user and the key
4. A unique data encryption key is created and encrypted under the KMS master key
5. Plaintext and encrypted data key returned to RDS
6. Plaintext data key stored in memory and used to encrypt/decrypt RDS data

# Enabling encryption

(default) aws/rds  
✓ Enter a key ARN

## Console

<b>Enable Encryption</b>	Yes
<b>Master Key</b>	(default) aws/rds
<b>Description</b>	Default master key that protects my RDS database volumes when no other key is defined
<b>Account</b>	This account ( )
<b>KMS Key ID</b>	( )

## AWS Command Line Interface (AWS CLI)


```
aws rds create-db-instance --region us-west-2 --db-instance-identifier sg-cli-test \  
--allocated-storage 20 --storage-encrypted \  
--db-instance-class db.m4.large --engine mysql \  
--master-username myawsuser --master-user-password myawsuser
```

```
aws rds create-db-instance --region us-west-2 --db-instance-identifier sg-cli-test1 \  
allocated-storage 20 --storage-encrypted --kms-key-id xxxxxxxxxxxxxxxxxxxx \  
instance-class db.m4.large --engine mysql \  
--master-username myawsuser --master-user-  
password myawsuser
```

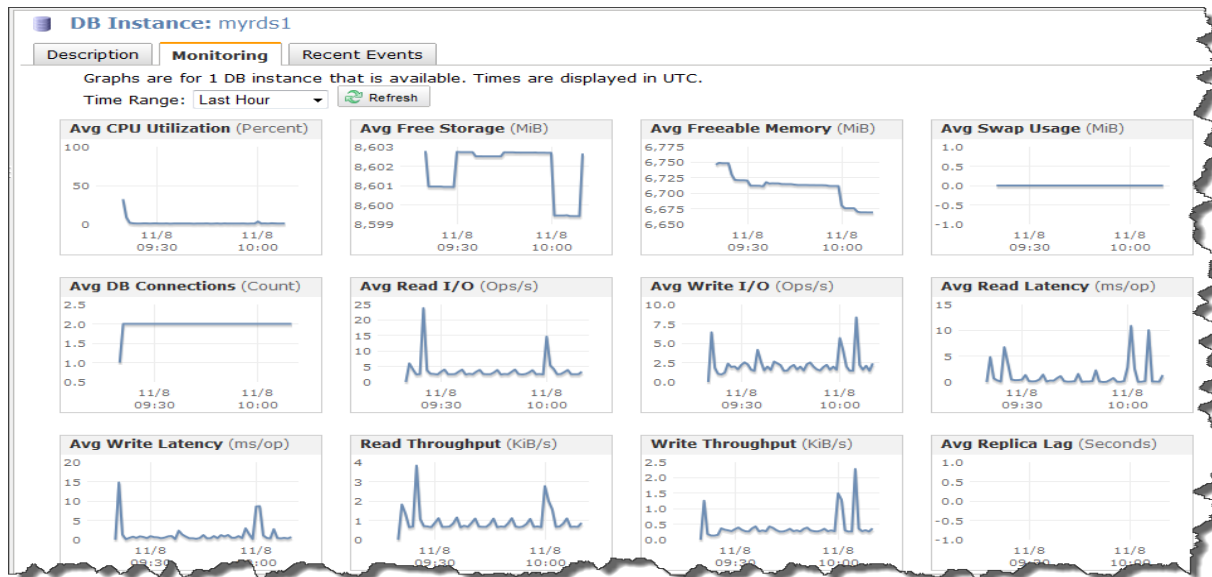


Pop-up Loft  
LONDON

# Metrics and Monitoring



# Standard monitoring



## Amazon CloudWatch metrics for Amazon RDS

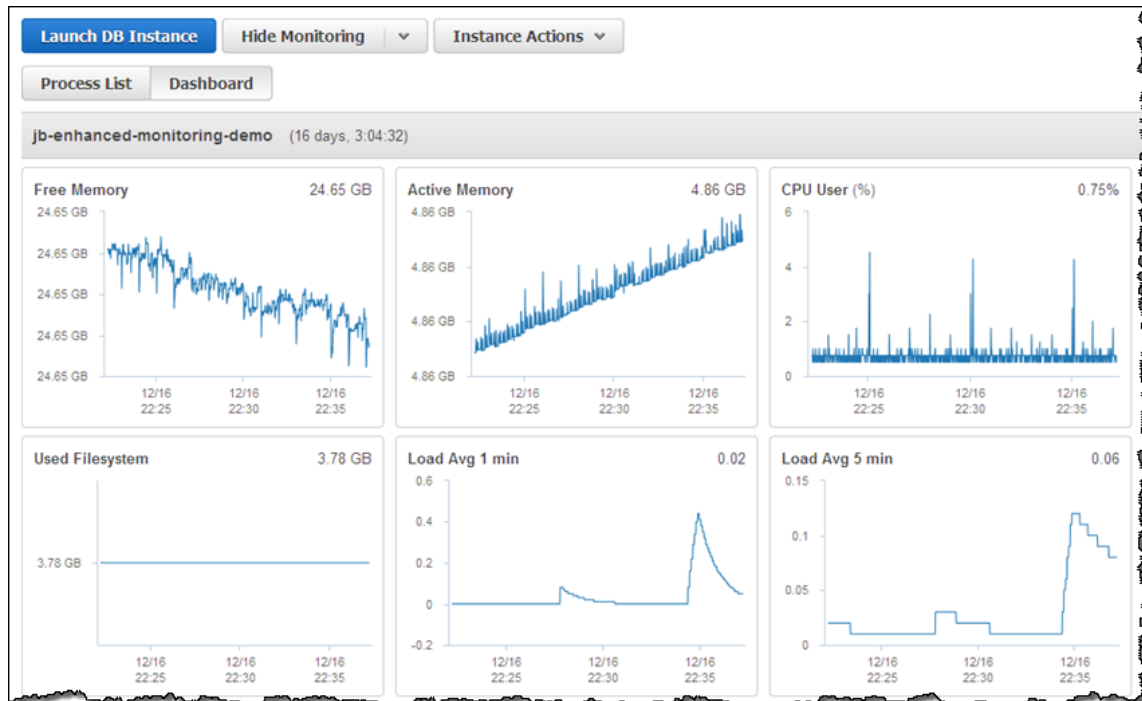
- CPU utilization
- Storage
- Memory
- Swap usage
- DB connections
- I/O (read and write)
- Latency (read and write)
- Throughput (read and write)
- Replica lag
- Many more

## Amazon CloudWatch Alarms

- Similar to on-premises custom monitoring tools

# Enhanced Monitoring

Access to over 50 new CPU, memory, file system, and disk I/O metrics as low as 1 second intervals



### Monitoring

Enable Enhanced Monitoring

Yes

Monitoring Role

Default

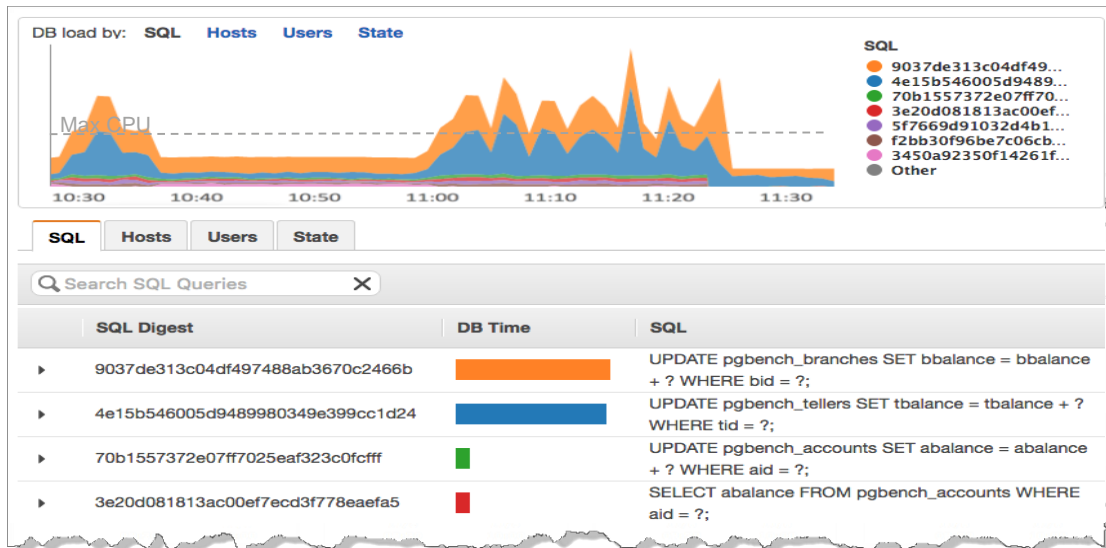
Granularity

1

second(s)

☒ I authorize RDS to create the IAM role `rds-monitoring-role`.

# Simplify monitoring with AWS Management Console



## Amazon Performance Insights for RDS



- Database Load : Identifies database bottlenecks
  - Easy
  - Powerful
- Identifies source of bottlenecks
  - Top SQL
- Adjustable Time frame
  - Hour, day, week and longer

AWS re:Invent 2016 DAT206: [https://youtu.be/ztmtJJTC8\\_Y?t=39m53s](https://youtu.be/ztmtJJTC8_Y?t=39m53s)

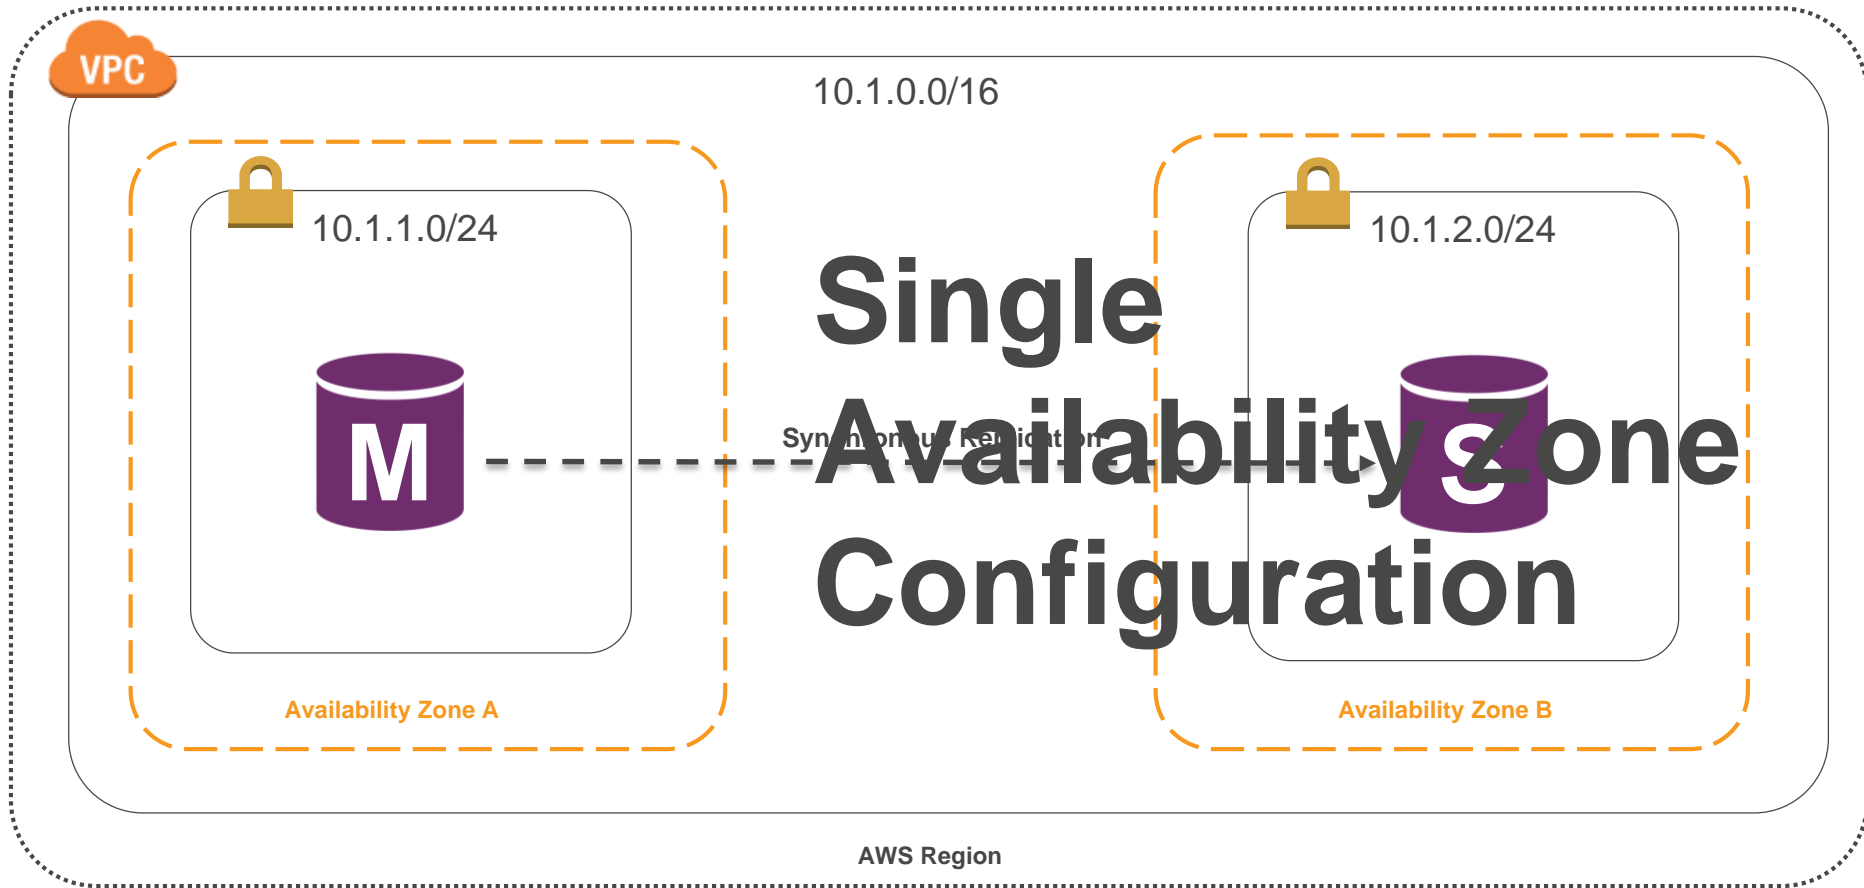


Pop-up Loft  
LONDON

# High Availability




# HA Multi Availability Zone Configuration





# High availability - Multi-AZ - DNS

mydatabase.us-west-2.rds.amazonaws.com:3306



# Read Replicas

Bring data close to your customer's applications in different regions

Relieve pressure on your master node for supporting reads and writes

Promote a Read Replica to a master for faster recovery in the event of disaster




Within or cross-region

- MySQL
- MariaDB
- PostgreSQL
- Aurora

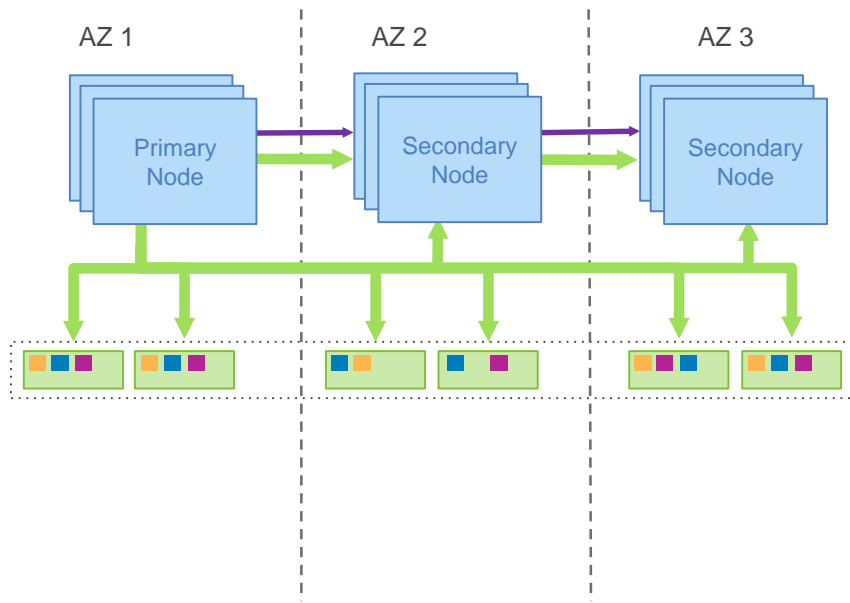
# High availability - Amazon Aurora storage

- Storage volume automatically grows up to 64 TB
- Quorum system for read/write; latency tolerant
- Peer-to-peer gossip replication to fill in holes
- Continuous backup to Amazon S3 (built for 11 9s durability)
- Continuous monitoring of nodes and disks for repair
- 10 GB segments as unit of repair or hotspot rebalance
- Quorum membership changes do not stall writes



# High availability - Amazon Aurora nodes



- Aurora cluster contains primary node and up to 15 secondary nodes
- Failing database nodes are automatically detected and replaced
- Failing database processes are automatically detected and recycled
- Secondary nodes automatically promoted on persistent outage, no single point of failure
- Customer application can scale out read traffic across secondary nodes






Pop-up Loft  
LONDON

# Scaling



# Why Scale?

- Handle higher load or lower usage
- Naturally grow over time
- Control costs



# What can I scale?

Database Instance



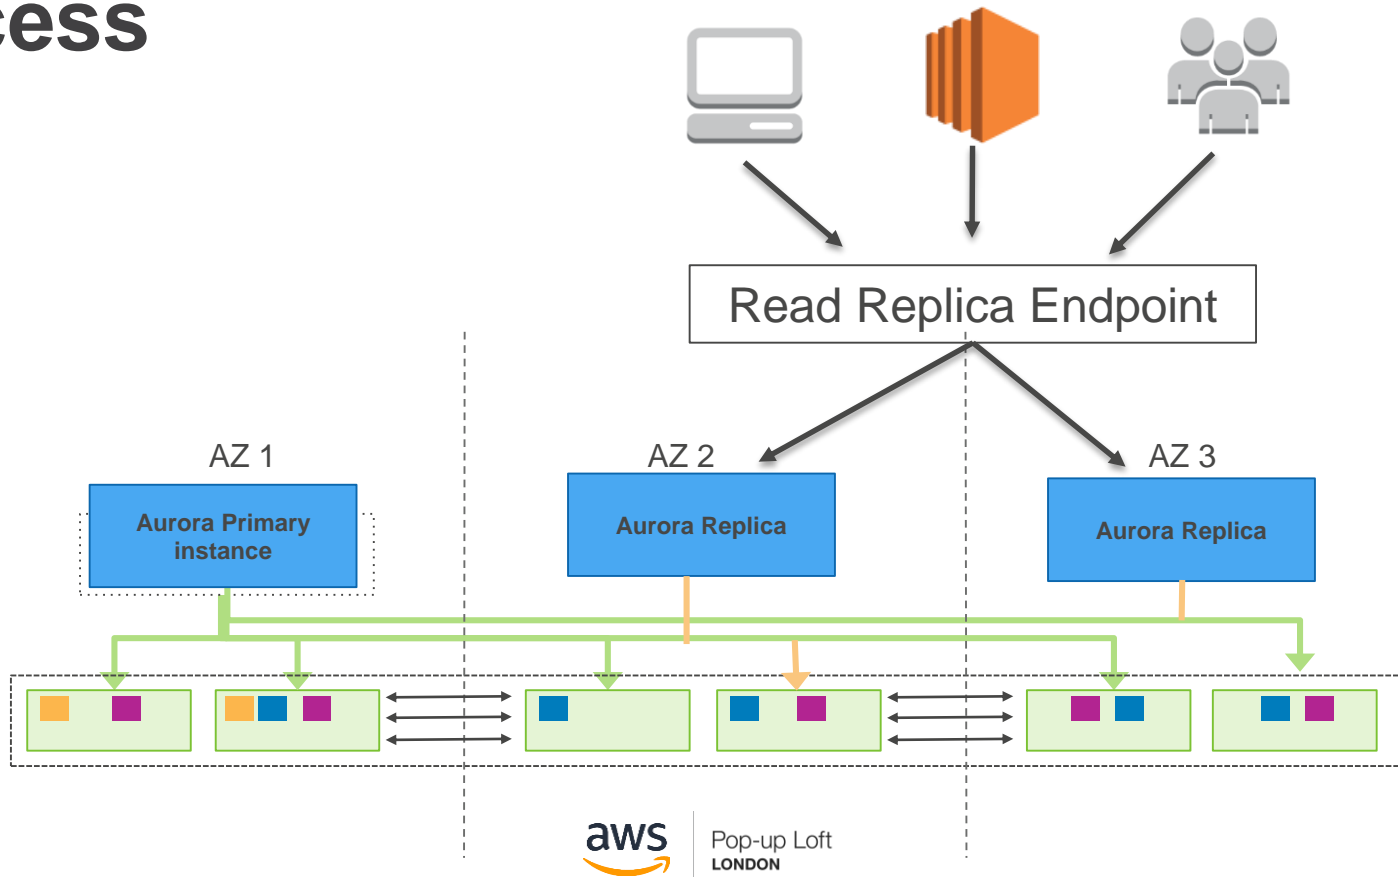
Read Replicas



Storage



# Amazon Aurora - Balanced Read Replica Access





# Scaling your instance up/down

AWS Management Console

**Instance Actions** ▾

- See Details
- Create Read Replica
- Promote Read Replica
- Take Snapshot
- Restore to Point in Time
- Migrate Latest Snapshot
- Modify**
- Reboot
- Delete

## Modify DB Instance: mysql-test

### Instance Specifications

DB Engine Version	MySQL 5.6.27 (default) ▾
DB Instance Class	db.m4.xlarge — 4 vCPU, 16 GiB RAM ▾
Multi-AZ Deployment	Yes ▾
Storage Type	General Purpose (SSD) ▾
Allocated Storage*	1600 GB

 **Apply Immediately**  

# Scaling - single Availability Zone

mydatabase.us-west-2.rds.amazonaws.com:3306

## Alarms and Recent Events

TIME (UTC-7)	EVENT
Oct 30 8:06 AM	Finished applying modification to DB instance class
Oct 30 8:06 AM	DB instance restarted
Oct 30 8:03 AM	DB instance shutdown
Oct 30 8:00 AM	Applying modification to database instance class

AWS Region

# Scaling - Multi-AZ

## Alarms and Recent Events

TIME (UTC-7)	EVENT
Oct 30 8:11 AM	Finished applying modification to DB instance class
Oct 30 8:03 AM	Multi-AZ instance failover completed
Oct 30 8:02 AM	DB instance restarted
Oct 30 8:02 AM	Multi-AZ instance failover started
Oct 30 7:52 AM	Applying modification to database instance class

onaws.com:3306



# Scaling - automation

## AWS CLI

```
aws rds modify-db-instance --db-instance-identifier sg-cli-test --db-instance-class db.m4.large --apply-immediately
```

## Scheduled CLI—cron

```
#Scale down at 8:00 PM on Friday  
0 20 * * 5 /home/ec2-user/scripts/scale_down_rds.sh  
  
#Scale up at 4:00 AM on Monday  
0 4 * * 1 /home/ec2-user/scripts/scale_up_rds.sh
```

# Scaling - automation

## Scheduled - AWS Lambda

No server but still runs on a schedule!

```
import boto3

client=boto3.client('rds')

def lambda_handler(event, context):
    response=client.modify_db_instance(DBInstanceIdentifier='sg-cli-test',
                                       DBInstanceClass='db.m4.xlarge',
                                       ApplyImmediately=True)

    print response
```

# Scaling - automation

## Metrics-based scaling

- Amazon CloudWatch and AWS Lambda!



# Scaling - automation

```
import boto3
import json

client=boto3.client('rds')


def lambda_handler(event, context):

    message = event['Records'][0]['Sns']['Message']
    parsed_message=json.loads(message)
    db_instance=parsed_message['Trigger']['Dimensions'][0]['value']
    print 'DB Instance: ' + db_instance

    response=client.modify_db_instance(DBInstanceIdentifier=db_instance,
                                       DBInstanceClass='db.m4.large',
                                       ApplyImmediately=True)

    print response
```

# Switch Off Dev Test Instances





MySQL  
MariaDB  
PostgreSQL  
Oracle  
SQL Server





Pop-up Loft  
LONDON

# Backups and snapshots



# Automated Backups

MySQL, PostgreSQL, MariaDB, Oracle, SQL Server

- Scheduled daily volume backup of entire instance
- Archive database change logs
- 35-day retention
- Taken from standby when running multi-AZ

## Aurora

- Automatic, continuous, incremental backups
- No impact on database performance
- 35-day retention



# How do automated RDS backups work?

Every day during your backup window, the RDS service creates a storage volume snapshot of your database

→ If database is Multi-AZ, the snapshot is taken from the standby

Every five minutes, RDS backs up the transaction logs of your database

## Availability and Durability

**DB Instance Status** available

**Multi AZ** Yes

**Automated Backups** Enabled (7 Days)

**Latest Restore Time** October 12, 2016 at 4:50:00 PM UTC-7



Pop-up Loft  
LONDON

# Restoring

- Creates an entire new database instance
- You define all the instance configuration, just like creating a new instance

## Restore DB Instance

You are creating a new DB Instance from a source DB Instance at a specified time. This new DB Instance will have the default DB Security Group and DB Parameter Groups.

**Use Latest Restorable Time** ☒ October 30, 2016 at 7:49:24 AM UTC-7


**Use Custom Restore Time** ☐ MMMM d, y 00 : 00 : 00 UTC-7



Pop-up Loft  
LONDON

# Snapshots



- Full copies of your RDS database
- Independent of scheduled backups
- Used to create a new RDS instance
- Taken from the standby when running multi-AZ



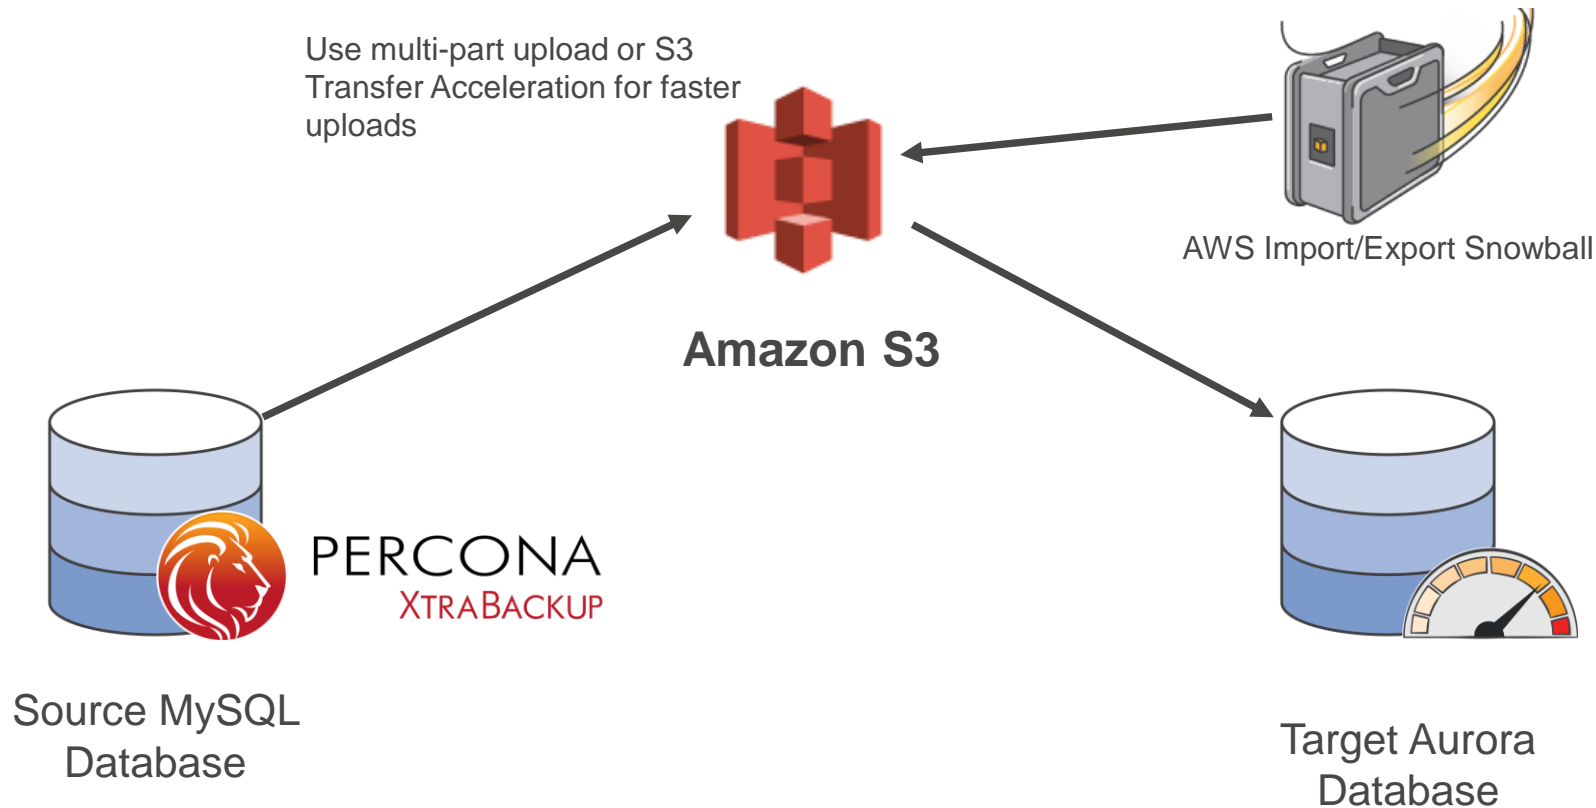


Pop-up Loft  
LONDON

# Migrating onto RDS



# MySQL Backup to Aurora via S3



# MySQL Backup to Aurora via S3

Restore Aurora DB Cluster from S3



### Specify Source Backup Details

#### Source Database Specifications

Source Engine

mysql

Source Engine Version

5.6

#### S3 Backup Location

Select S3 Bucket\*

mysql-demo-backups

S3 Bucket Prefix (Optional)

#### IAM Role

IAM Role\*

db-backup-file-access

Create a New Role

\* Required

Cancel

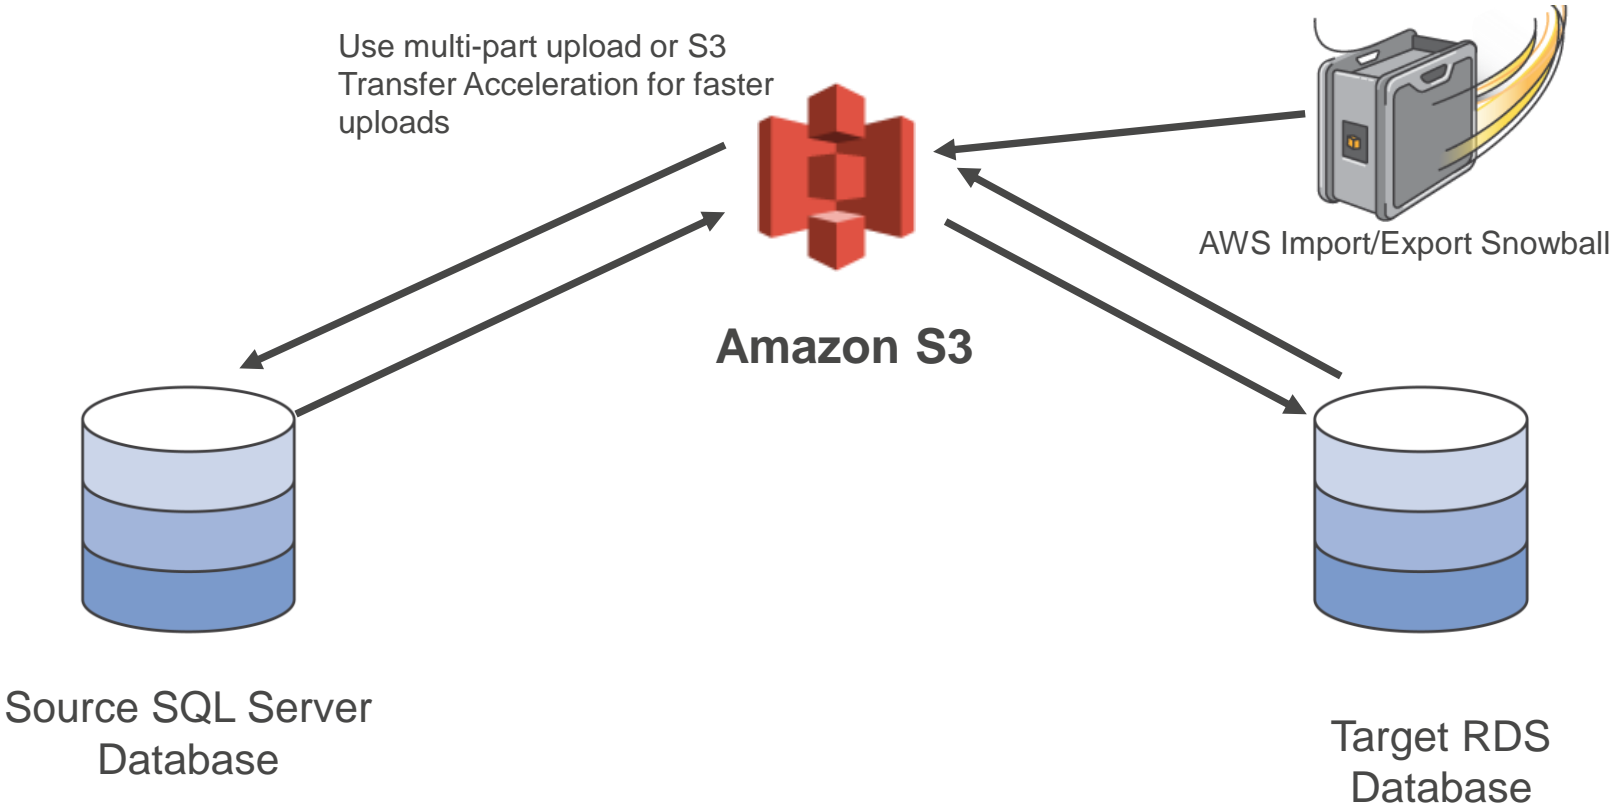
Next Step



Pop-up Loft  
LONDON



# SQL Server Backup to RDS SQL Server via S3



# SQL Server Backup to RDS SQL Server via S3

## Importing to RDS

### No Encryption

```
exec msdb.dbo.rds_restore_database  
    @restore_db_name='database_name',  
    @s3_arn_to_restore_from='arn:aws:s3:::bucket_name/file_name_and_extension';
```

### Encryption

```
exec msdb.dbo.rds_restore_database  
    @restore_db_name='database_name',  
    @s3_arn_to_restore_from='arn:aws:s3:::bucket_name/file_name_and_extension',  
    @kms_master_key_arn='arn:aws:kms:region:account-id:key/key-id';
```



## AWS Database Migration Service




ORACLE

**Amazon** Aurora



- ✓ Move data to the same or different database engine
- ✓ Keep your apps running during the migration
- ✓ Start your first migration in 10 minutes or less
- ✓ Replicate within, to, or from Amazon EC2 or RDS



## AWS Schema Conversion Tool

- ✓ Migrate from Oracle and SQL Server
- ✓ Move your tables, views, stored procedures, and data manipulation language (DML) to MySQL, MariaDB, and Aurora
- ✓ Highlight where manual edits are needed

# SCT supported OLTP conversions

Source Database	Target Database on Amazon RDS
Microsoft SQL Server (version 2008 and later)	Amazon Aurora (MySQL or PostgreSQL), Microsoft SQL Server, MySQL, PostgreSQL
MySQL (version 5.5 and later)	Amazon Aurora (PostgreSQL), MySQL, PostgreSQL
Oracle (version 10.2 and later)	Amazon Aurora (MySQL or PostgreSQL), MySQL, Oracle, PostgreSQL
PostgreSQL (version 9.1 and later)	Amazon Aurora (MySQL), MySQL, PostgreSQL



Pop-up Loft  
**LONDON**

# Thank you!

