



Masterclass

Amazon Elastic MapReduce



Ian Massingham – Technical Evangelist



ianmas@amazon.com



@IanMmmm



Masterclass

1

A technical deep dive that goes beyond the basics

2

Intended to educate you on how to get the best from AWS services

3

Show you how things work and how to get things done

Amazon Elastic MapReduce



Provides a managed Hadoop framework
Quickly & cost-effectively process vast amounts of data
Makes it easy, fast & cost-effective for you to process data
Run other popular distributed frameworks such as Spark

Low Cost

Easy to Use

Elastic



Amazon EMR

Flexible

Reliable

Secure



Amazon EMR: Example Use Cases

Clickstream Analysis

Amazon EMR can be used to analyze click stream data in order to segment users and understand user preferences. Advertisers can also analyze click streams and advertising impression logs to deliver more effective ads.

Genomics

Amazon EMR can be used to process vast amounts of genomic data and other large scientific data sets quickly and efficiently. Researchers can access genomic data hosted for free on AWS.

Log Processing

Amazon EMR can be used to process logs generated by web and mobile applications. Amazon EMR helps customers turn petabytes of un-structured or semi-structured data into useful insights about their applications or users.



Agenda

- Hadoop Fundamentals
- Core Features of Amazon EMR
- How to Get Started with Amazon EMR
- Supported Hadoop Tools
- Additional EMR Features
- Third Party Tools
- Resources where you can learn more

HADOOP FUNDAMENTALS

Very large
clickstream
logging data
(e.g TBs)

Lots of actions by
John Smith



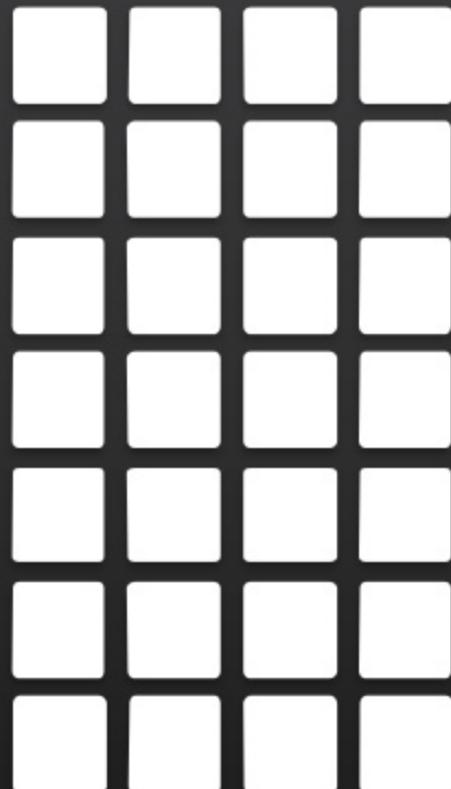
Very large
clickstream
logging data
(e.g TBs)

Lots of actions by
John Smith

Very large
clickstream
logging data
(e.g TBs)



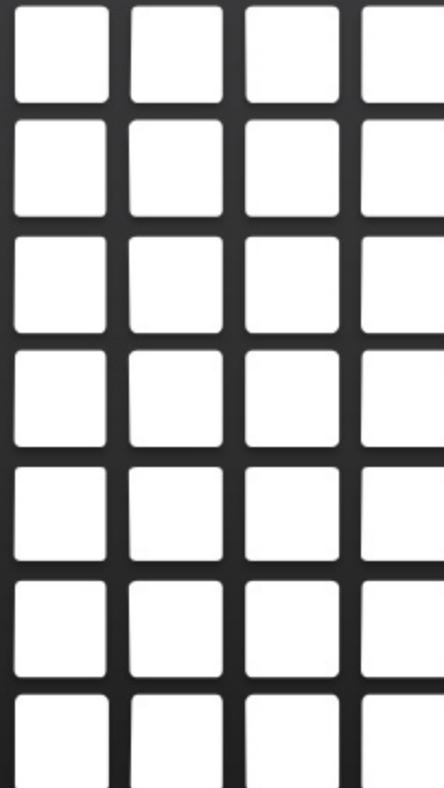
Split the log
into many
small pieces



Very large
clickstream
logging data
(e.g TBs)

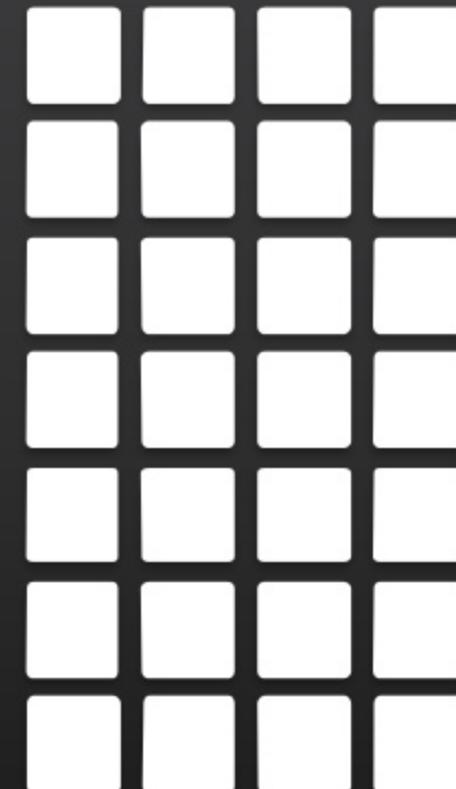
Lots of actions by
John Smith

Split the log
into many
small pieces



Very large clickstream logging data (e.g TBs)

Lots of actions by John Smith



Process in an EMR cluster

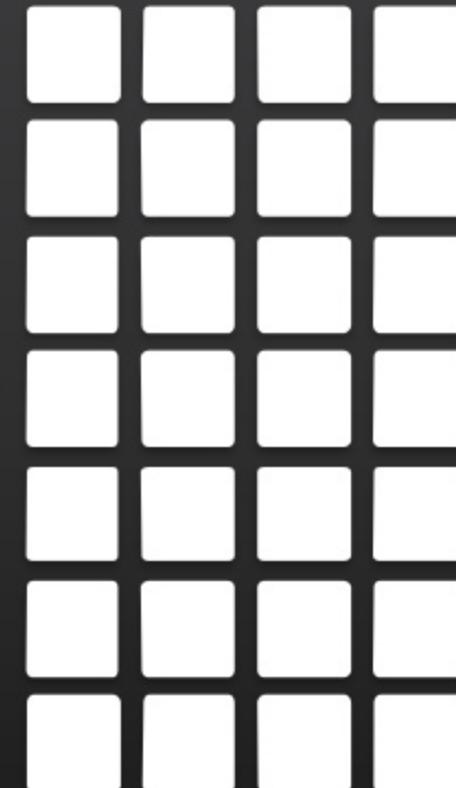
Split the log into many small pieces

Aggregate the results from all the nodes

Very large clickstream logging data (e.g TBs)

Split the log into many small pieces

Lots of actions by John Smith



Process in an EMR cluster

Aggregate the results from all the nodes

What John Smith did

Very large
clickstream
logging data
(e.g TBs)



Insight in a fraction of the time

What John
Smith did

CORE FEATURES OF AMAZON EMR

ELASTIC



Elastic

Provision as much capacity as you need
Add or remove capacity at any time

Deploy Multiple Clusters



Resize a Running Cluster



LOW COST



Low Cost

Low Hourly Pricing

Amazon EC2 Spot Integration



Amazon EC2 Reserved Instance Integration

Elasticity



Amazon S3 Integration



Low Cost

Accenture Hadoop Study:

Amazon EMR ‘offers better price-performance’

The image shows a white cloud icon connected by a black cable to a large blue right-pointing arrow. Below the arrow, the text "High performance. Delivered." is written. At the bottom left, there is a small red link that says "Where to deploy your Hadoop clusters? Executive Summary". The Accenture logo is at the bottom right.

High performance. Delivered.

Accenture Technology Labs
Where to deploy your Hadoop clusters?
Executive Summary

accenture

consulting | technology | outsourcing

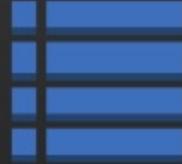
FLEXIBLE DATA STORES



Amazon
S3



Hadoop Distributed
File System



Amazon
DynamoDB



Amazon
EMR



Amazon
Redshift



Amazon
Glacier



Amazon
Relational
Database Service

Amazon S3 + Amazon EMR



- Allows you to decouple storage and computing resources
- Use Amazon S3 features such as server-side encryption
- When you launch your cluster, EMR streams data from S3
- Multiple clusters can process the same data concurrently

Hadoop Distributed
File System (HDFS)



Amazon
DynamoDB



AWS
Data Pipeline

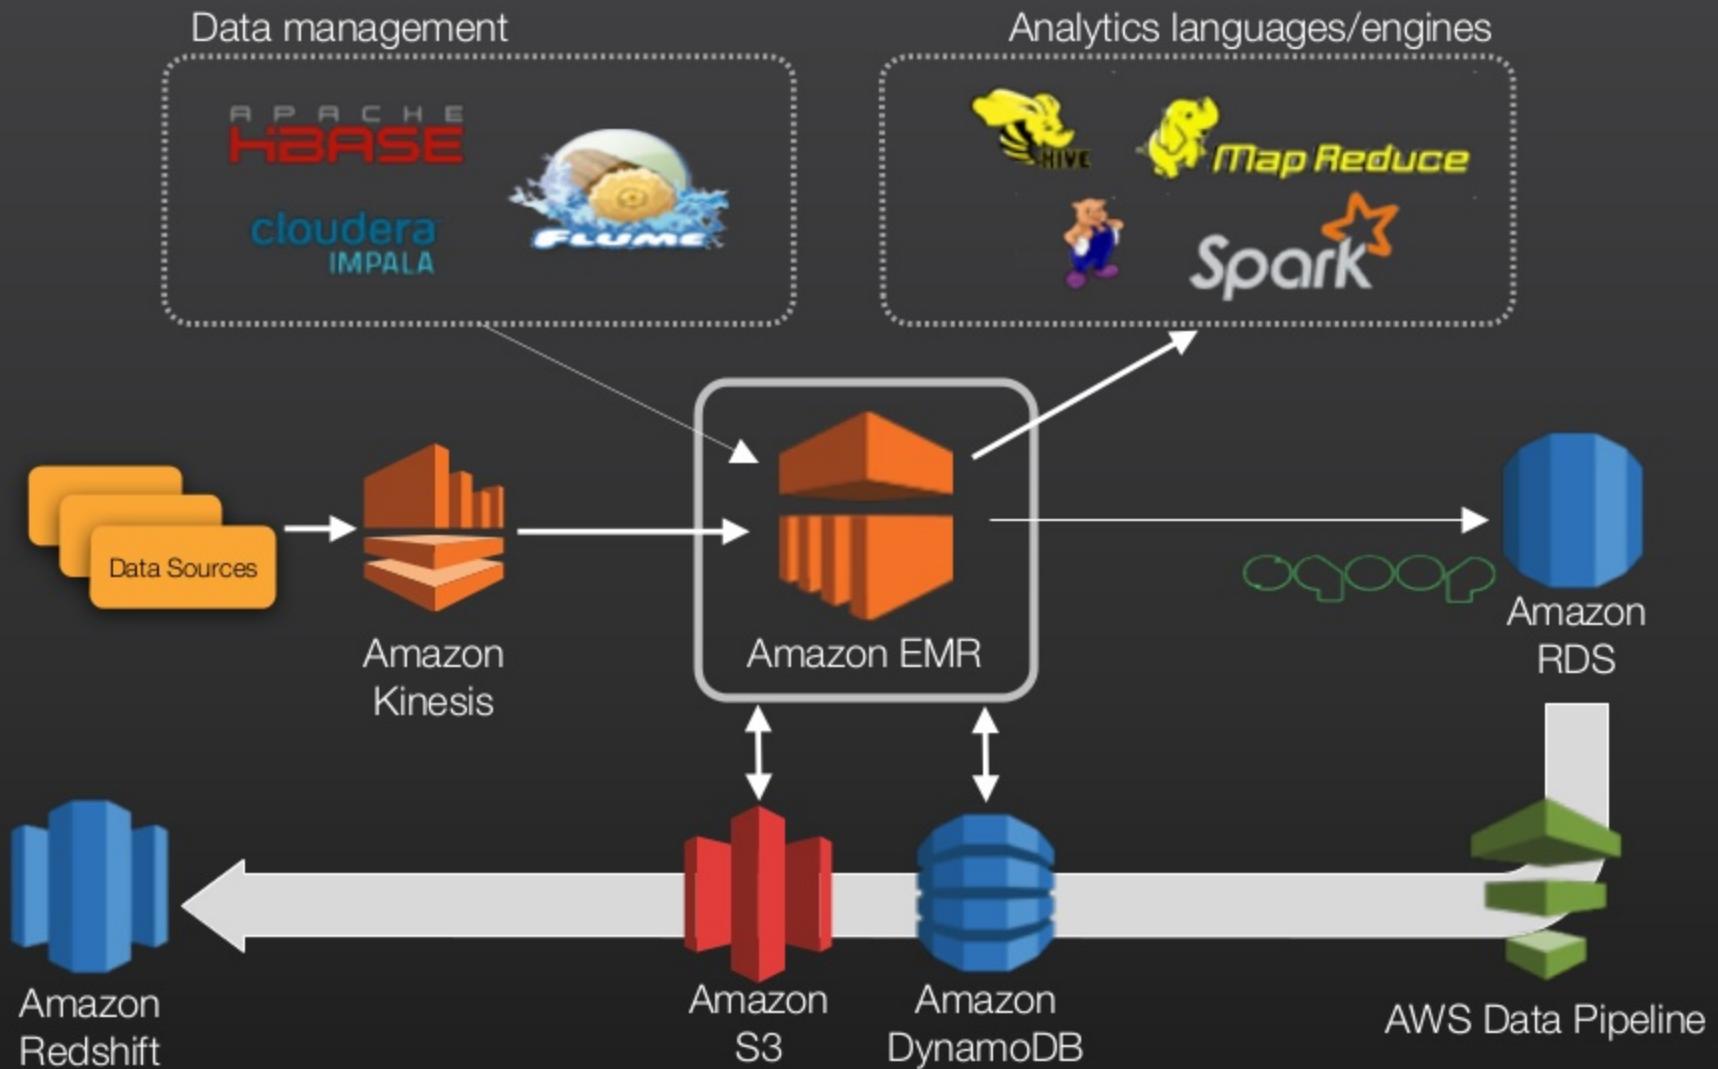


Amazon
RDS



Amazon
Redshift





GETTING STARTED WITH AMAZON ELASTIC MAPREDUCE

Develop your data processing application



The screenshot shows a list of articles under the 'Articles & Tutorials' section. The articles include:

- Run MapReduce on Amazon Hadoop MapReduce**: An introductory article about running MapReduce on Amazon Hadoop MapReduce.
- Commodity Advertising using Apache Hive and Amazon EMR**: An article on using Apache Hive to run commodity advertising on Amazon EMR.
- Using Hadoop with Amazon Elastic MapReduce**: An article on using Hadoop with Amazon EMR.
- Analyze Big Data with Apache Hive, Windows PowerShell, and Amazon EMR**: An article on using Apache Hive, Windows PowerShell, and Amazon EMR to analyze big data.
- Using Streaming MapReduce with EMR**: An article on using Streaming MapReduce with Amazon EMR.
- Amazon EMR with the MapReduce Distribution for MapReduce**: An article on using the MapReduce distribution with Amazon EMR.
- Building Machine Learning using Amazon EMR**: An article on building machine learning applications using Amazon EMR.
- Running MapReduce with the Java API for EMR**: An article on using the Java API for Amazon EMR to run MapReduce jobs.
- Identifying Similar Items using Apache Beeswarm and Apache Hive on Hadoop**: An article on using Apache Beeswarm and Apache Hive on Hadoop to identify similar items.
- Word Count Example**: An article on writing a word count example.
- CloudFront**: An article on CloudFront.

Develop your data processing application



Upload your application and data to Amazon S3

Develop your data processing application



Upload your application and data to Amazon S3



Develop your data processing application



Upload your application and data to Amazon S3



Develop your data processing application



Upload your application and data to Amazon S3



Develop your data processing application



Upload your application and data to Amazon S3



Configure and launch your cluster

Configure and launch your cluster

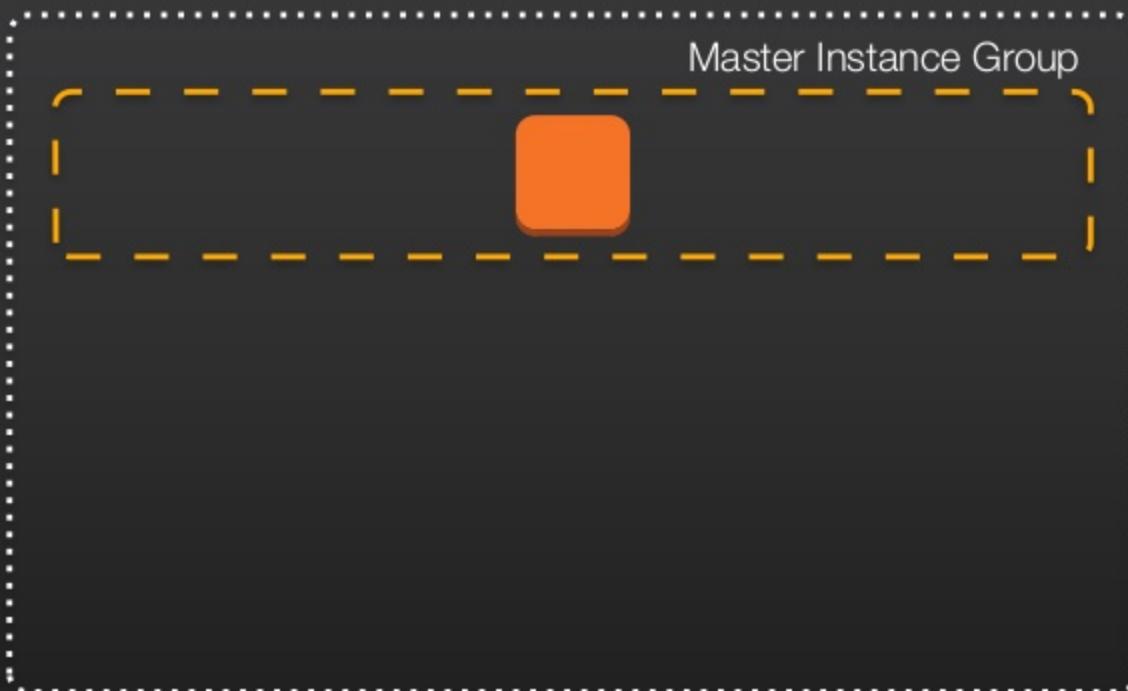
Amazon EMR Cluster

Start an EMR cluster
using console, CLI tools
or an AWS SDK

Configure and launch your cluster

Master instance group
created that controls the
cluster

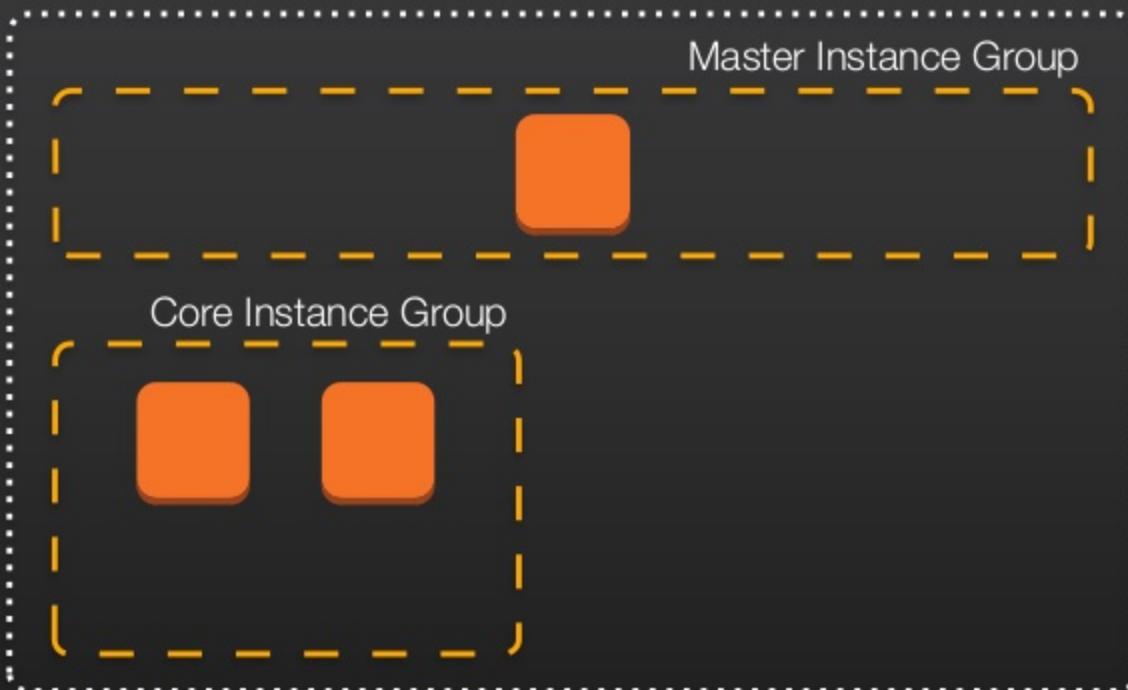
Amazon EMR Cluster



Configure and launch your cluster

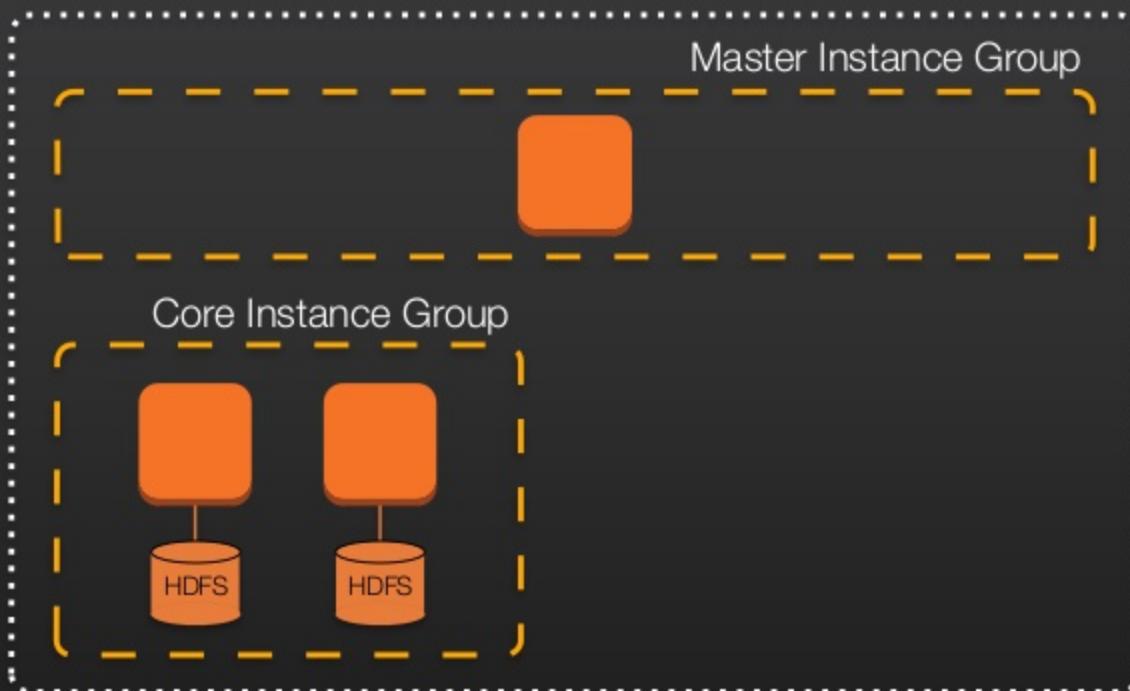
Core instance group
created for life of cluster

Amazon EMR Cluster



Configure and launch your cluster

Amazon EMR Cluster



Core instance group
created for life of cluster

Core instances run
DataNode and
TaskTracker daemons

Configure and launch your cluster

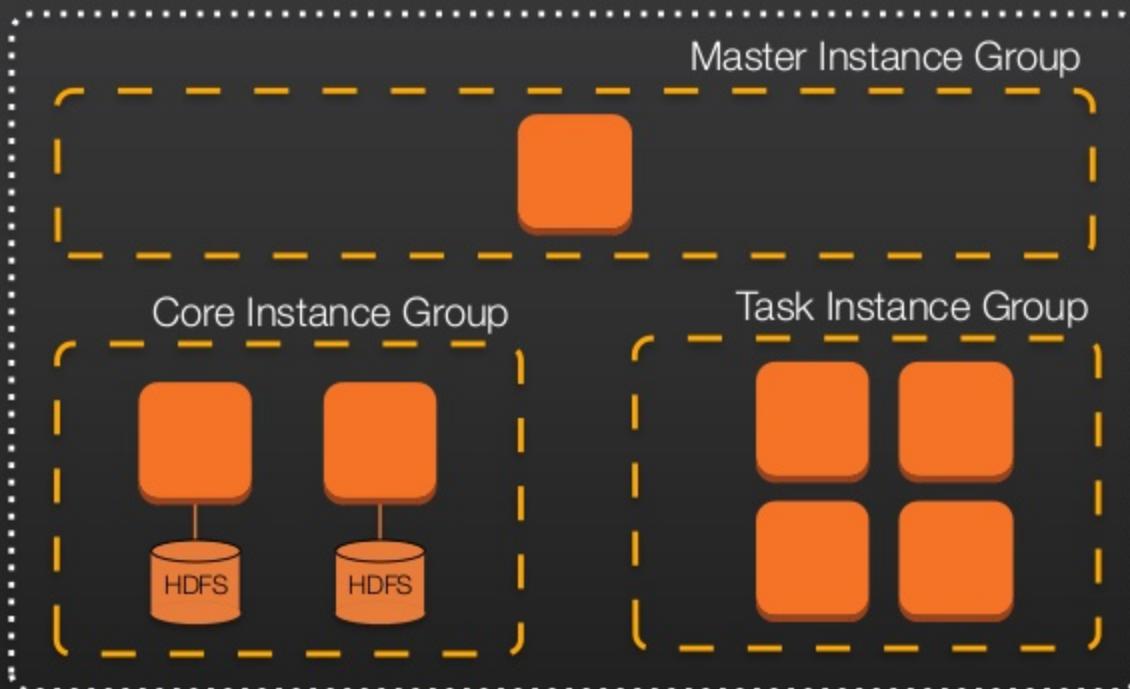
Amazon EMR Cluster

Master Instance Group

Core Instance Group

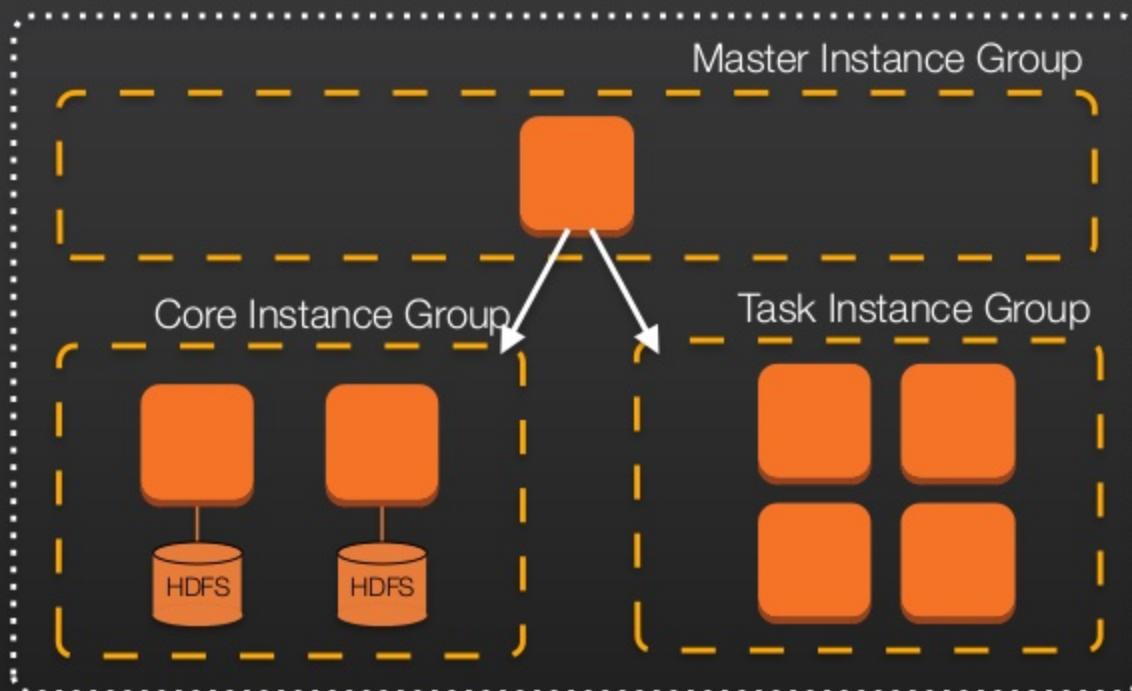
Task Instance Group

Optional task instances
can be added or
subtracted



Configure and launch your cluster

Amazon EMR Cluster



Master node
coordinates distribution
of work and manages
cluster state

Develop your data processing application



Upload your application and data to Amazon S3



Configure and launch your cluster



Optionally, monitor the cluster

Develop your data processing application



Upload your application and data to Amazon S3



Configure and launch your cluster



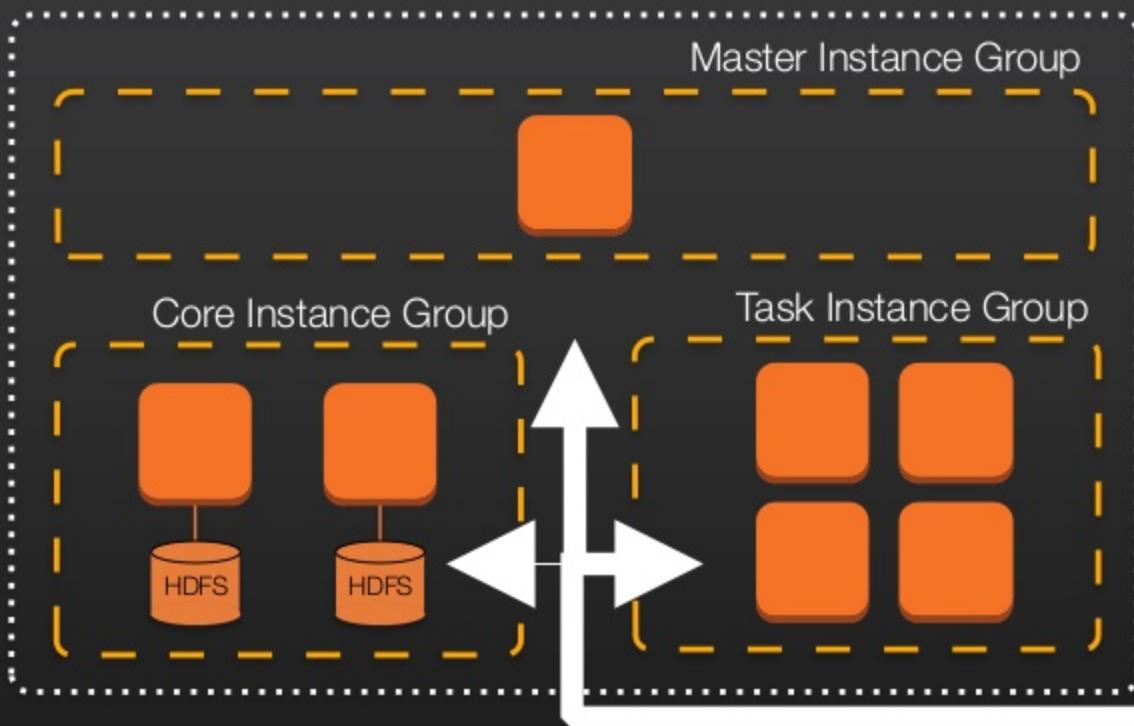
Optionally, monitor the cluster



Retrieve the output

Retrieve the output

Amazon EMR Cluster



S3 can be used as underlying ‘file system’ for input/output data



DEMO:

**GETTING STARTED WITH
AMAZON EMR USING A SAMPLE
HADOOP STREAMING APPLICATION**

Hadoop Streaming

Utility that comes with the Hadoop distribution

Allows you to create and run Map/Reduce jobs with any executable or script as the mapper and/or the reducer

Reads the input from standard input and the reducer outputs data through standard output

By default, each line of input/output represents a record with tab separated key/value

Job Flow for Sample Application

Steps

A step is a unit of work you submit to the cluster. A step might contain one or more Hadoop jobs, or contain instructions to install or configure an application. You can submit up to 256 steps to a cluster. [Learn more](#)

| Name | Action on failure | JAR location | Arguments |
|-------------------|-------------------|-----------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Word count | Terminate cluster | /home/hadoop/contrib/streaming/hadoop-streaming.jar | -files s3://eu-west-1.elasticmapreduce/samples/wordcount/wordSplitter.py -mapper wordSplitter.py -reducer aggregate -input s3://eu-west-1.elasticmapreduce/samples/wordcount/input -output s3://ianmas-aws-emr/intermediate/ |
| Streaming program | Terminate cluster | /home/hadoop/contrib/streaming/hadoop-streaming.jar | -mapper /bin/cat -reducer org.apache.hadoop.mapred.lib.IdentityReducer -input s3://ianmas-aws-emr/intermediate/ -output s3://ianmas-aws-emr/output -jobconf mapred.reduce.tasks=1 |

Job Flow: Step 1

JAR location: /home/hadoop/contrib/streaming/hadoop-streaming.jar

Arguments:

```
-files s3://eu-west-1.elasticmapreduce/samples/wordcount/wordSplitter.py  
-mapper wordSplitter.py  
-reducer aggregate  
-input s3://eu-west-1.elasticmapreduce/samples/wordcount/input  
-output s3://ianmas-aws-emr/intermediate/
```

Step 1: mapper: wordSplitter.py

```
#!/usr/bin/python
import sys
import re

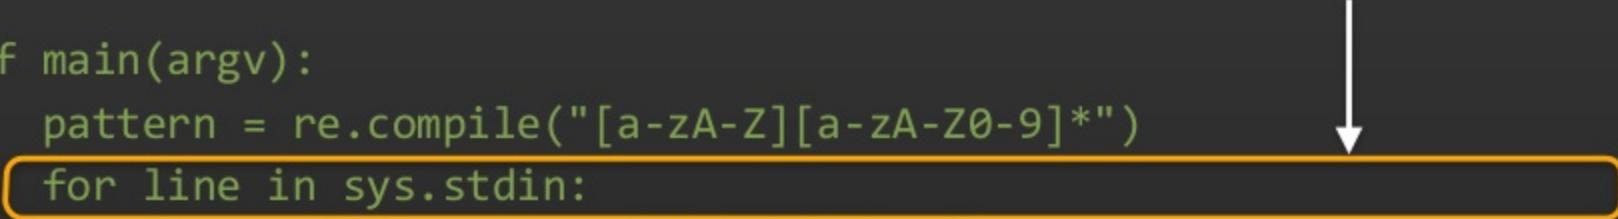
def main(argv):
    pattern = re.compile("[a-zA-Z][a-zA-Z0-9]*")
    for line in sys.stdin:
        for word in pattern.findall(line):
            print "LongValueSum:" + word.lower() + "\t" + "1"

if __name__ == "__main__":
    main(sys.argv)
```

Step 1: mapper: wordSplitter.py

```
#!/usr/bin/python
import sys
import re
def main(argv):
    pattern = re.compile("[a-zA-Z][a-zA-Z0-9]*")
    for line in sys.stdin:
        for word in pattern.findall(line):
            print "LongValueSum:" + word.lower() + "\t" + "1"
if __name__ == "__main__":
    main(sys.argv)
```

Read words from StdIn line by line



Step 1: mapper: wordSplitter.py

```
#!/usr/bin/python
import sys
import re
Output to StdOut tab delimited records
in the format "LongValueSum:abacus\t1"
def main(argv):
    pattern = re.compile("[a-zA-Z][a-zA-Z0-9]*")
    for line in sys.stdin:
        for word in pattern.findall(line):
            print "LongValueSum:" + word.lower() + "\t" + "1"
if __name__ == "__main__":
    main(sys.argv)
```



Step 1: reducer: aggregate

Sorts inputs and adds up totals:

“Abacus 1”

“Abacus 1”

“Abacus 1”

becomes

“Abacus 3”

Step 1: input/output

The input is all the objects in the S3 bucket/prefix:

```
s3://eu-west-1.elasticmapreduce/samples/wordcount/input
```

Output is written to the following S3 bucket/prefix to be used as input for the next step in the job flow:

```
s3://ianmas-aws-emr/intermediate/
```

One output object is created for each reducer (generally one per core)

Job Flow: Step 2

JAR location: /home/hadoop/contrib/streaming/hadoop-streaming.jar

Arguments:

Accept anything and return as text

-mapper /bin/cat

-reducer org.apache.hadoop.mapred.lib.IdentityReducer

-input s3://ianmas-aws-emr/intermediate/

-output s3://ianmas-aws-emr/output

-jobconf mapred.reduce.tasks=1

Job Flow: Step 2

JAR location: /home/hadoop/contrib/streaming/hadoop-streaming.jar

Arguments:

```
-mapper /bin/cat  
-reducer org.apache.hadoop.mapred.lib.IdentityReducer  
-input s3://ianmas-aws-emr/intermediate/  
-output s3://ianmas-aws-emr/output  
-jobconf mapred.reduce.tasks=1
```

Sort
↓

Job Flow: Step 2

JAR location: /home/hadoop/contrib/streaming/hadoop-streaming.jar

Arguments:

```
-mapper /bin/cat  
-reducer org.apache.hadoop.mapred.lib.IdentityReducer  
-input s3://ianmas-aws-emr/intermediate/  
-output s3://ianmas-aws-emr/output  
-jobconf mapred.reduce.tasks=1
```

Take previous output as input



Job Flow: Step 2

JAR location: /home/hadoop/contrib/streaming/hadoop-streaming.jar

Arguments:

Output location

```
-mapper /bin/cat  
-reducer org.apache.hadoop.mapred.lib.IdentityReducer  
-input s3://ianmas-aws-emr/intermediate/  
-output s3://ianmas-aws-emr/output  
-jobconf mapred.reduce.tasks=1
```



Job Flow: Step 2

JAR location: /home/hadoop/contrib/streaming/hadoop-streaming.jar

Arguments:

Use a single reduce task
to get a single output object

```
-mapper /bin/cat  
-reducer org.apache.hadoop.mapred.lib.IdentityReducer  
-input s3://ianmas-aws-emr/intermediate/  
-output s3://ianmas-aws-emr/output  
-jobconf mapred.reduce.tasks=1
```



SUPPORTED HADOOP TOOLS



Supported Hadoop Tools

Hive



An open source data warehouse & analytics package that runs on top of Hadoop. Operated by Hive QL, a SQL-based language which allows users to structure, summarize, and query data.

Pig



An open source analytics package that runs on top of Hadoop. Pig is operated by Pig Latin, a SQL-like language which allows users to structure, summarize, and query data. Allows processing of complex and unstructured data sources such as text documents and log files.

HBase



Provides you an efficient way of storing large quantities of sparse data using column-based storage. HBase provides fast lookup of data because data is stored in-memory instead of on disk. Optimized for sequential write operations, and it is highly efficient for batch inserts, updates, and deletes.



Supported Hadoop Tools

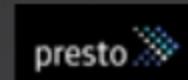
Impala



A tool in the Hadoop ecosystem for interactive, ad hoc querying using SQL syntax. It uses a massively parallel processing (MPP) engine similar to that found in a traditional RDBMS.

This lends Impala to interactive, low-latency analytics. You can connect to BI tools through ODBC and JDBC drivers.

Presto



An open source distributed SQL query engine for running interactive analytic queries against data sources of all sizes ranging from gigabytes to petabytes.

Hue



An open source user interface for Hadoop that makes it easier to run and develop Hive queries, manage files in HDFS, run and develop Pig scripts, and manage tables.

DEMO: APACHE HUE ON EMR



AWS Official Blog

New – Apache Spark on Amazon EMR

by Jeff Barr | on 16 JUN 2013 | in [Amazon EMR](#), [Permalink](#)

My colleague Jon Fritz wrote the guest post below to introduce a powerful new feature for Amazon EMR.

— [Jeff](#)

I'm happy to announce that Amazon EMR now supports [Apache Spark](#). Amazon EMR is a web service that makes it easy for you to process and analyze vast amounts of data using applications in the Hadoop ecosystem, including Hive, Pig, HBase, Presto, Impala, and others. We're delighted to officially add Spark to this list. Although many customers have previously been installing Spark using custom scripts, you can now launch an Amazon EMR cluster with Spark directly from the Amazon EMR Console, CLI, or API.

[Apache Spark: Beyond Hadoop MapReduce](#)

We have seen great customer successes using Hadoop MapReduce for large scale data processing, batch reporting, ad hoc analysis on unstructured data, and machine learning. Apache Spark, a newer distributed processing framework in the Hadoop ecosystem, is also proving to be an enticing engine by increasing job performance and development velocity for certain workloads.

By using a directed acyclic graph (DAG) execution engine, Spark can create a more efficient query plan for data transformations. Also, Spark uses in-memory, fault-tolerant resilient distributed datasets (RDDs), keeping intermediates, inputs, and outputs in memory instead of on-disk. These two elements of functionality can result in better performance for certain workloads when compared to Hadoop MapReduce, which will force jobs into a sequential map-reduce framework and incurs an I/O cost from writing intermediates out to disk. Spark's performance enhancements are particularly applicable for iterative workloads, which are common in machine learning and low-latency querying use cases.

Additionally, Spark natively supports Scala, Python, and Java APIs, and it includes libraries for SQL, popular machine learning algorithms, graph processing, and stream processing. With many tightly integrated development options, it can be easier to create and maintain applications for Spark than to work with the various abstractions wrapped around the Hadoop MapReduce API.

[Introducing Spark on Amazon EMR](#)

Today, we are introducing support for [Apache Spark](#) in Amazon EMR. You can quickly and easily create scalable, managed Spark clusters on a variety of [Amazon Elastic Compute Cloud \(EC2\)](#) instance types from the Amazon EMR console, [AWS Command Line Interface](#), or [Amazon EMR API](#). Apache EMR now supports the most widely adopted (ES5) block storage device available to deliver a no-maintenance, high-performance, distributed system that scales linearly to whatever your data volume needs. Apache EMR now makes it easy to build distributed systems for the world's largest data processing needs.



Create a Cluster with Spark

```
$ aws emr create-cluster --name "Spark cluster" \
--ami-version 3.8 --applications Name=Spark \
--ec2-attributes KeyName=myKey --instance-type m3.xlarge \
--instance-count 3 --use-default-roles
```

```
$ ssh -i myKey hadoop@masternode
```

invoke the spark shell with

```
$ spark-shell
```

or

```
$ pyspark
```

Working with the Spark Shell

Counting the occurrences of a string a text file stored in Amazon S3 with spark

```
$ pyspark
>>> sc
<pyspark.context.SparkContext object at 0x7fe7e659fa50>
>>> textfile = sc.textFile("s3://elasticmapreduce/samples/hive-ads/tables/impressions/
dt=2009-04-13-08-05/ec2-0-51-75-39.amazon.com-2009-04-13-08-05.log")
>>> linesWithCartoonNetwork = textfile.filter(lambda line: "cartoonnetwork.com" in
line).count()
15/06/04 17:12:22 INFO lzo.GPLNativeCodeLoader: Loaded native gpl library from the
embedded binaries
<snip>
<Spark program continues>
>>> linesWithCartoonNetwork
9
```

ADDITIONAL EMR FEATURES

CONTROL NETWORK ACCESS TO YOUR EMR CLUSTER

Using SSH local port forwarding

```
ssh -i EMRKeyPair.pem -N \
-L 8160:ec2-52-16-143-78.eu-west-1.compute.amazonaws.com:8888 \
hadoop@ec2-52-16-143-78.eu-west-1.compute.amazonaws.com
```

MANAGE USERS, PERMISSIONS AND ENCRYPTION

File System Configuration

The EMR File System (EMRFS) and the Hadoop Distributed File System (HDFS) are both installed on your EMR cluster. HDFS stores data on an EMR cluster, while EMRFS allows EMR clusters to store data on S3. You can enable [S3 server-side encryption](#) or [S3 client-side encryption](#) and [consistent view](#) for EMRFS below, or use a bootstrap action to configure additional settings for EMRFS.

| | | |
|----------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| EMRFS S3 Encryption | <input type="button" value="None"/> <input checked="" type="button" value="S3 server-side encryption"/> <input type="button" value="S3 client-side encryption with AWS Key Management Service (KMS)"/> <input type="button" value="S3 client-side encryption with custom encryption materials provider"/> |  Choose encryption method for objects written to or read from S3 using EMRFS. Please note that this will not encrypt files written to HDFS. Learn more |
| Consistent view | <input type="checkbox"/> Monitors list and read-after-write (for new puts) consistency for files in S3. Learn more | |

INSTALL ADDITIONAL SOFTWARE WITH BOOTSTRAP ACTIONS

Bootstrap Actions

i Bootstrap actions are scripts that are executed during setup before Hadoop starts on every cluster node. You can use them to install additional software and customize your applications. [Learn more](#)

| Bootstrap action type | Name | S3 location | Optional arguments |
|-----------------------|--------------------------------------------|----------------------------------------------------------|--------------------|
| Add bootstrap action | <input type="text" value="Custom action"/> | <input type="button" value="Select a bootstrap action"/> | |

EFFICIENTLY COPY DATA TO EMR FROM AMAZON S3

Run on a cluster master node:

```
$ hadoop jar /home/hadoop/lib/emr-s3distcp-1.0.jar -  
Dmapreduce.job.reduces=30 --src s3://s3bucketname/ --dest hdfs://  
$HADOOP_NAMENODE_HOST:$HADOOP_NAMENODE_PORT/data/ --outputCodec 'none'
```

SCHEDULE RECURRING WORKFLOWS

AWS Data Pipeline

AWS Data Pipeline is a web service that helps you reliably process and move data between different AWS compute and storage services, as well as on-premise data sources, at specified intervals. With AWS Data Pipeline, you can regularly access your data where it's stored, transform and process it at scale, and efficiently transfer the results to AWS services such as Amazon S3, Amazon RDS, Amazon DynamoDB, and Amazon Elastic MapReduce (EMR).

AWS Data Pipeline helps you easily create complex data processing workloads that are fault tolerant, repeatable, and highly available. You don't have to worry about ensuring resource availability, managing inter-task dependencies, retrying transient failures or timeouts in individual tasks, or creating a failure notification system. AWS Data Pipeline also allows you to move and process data that was previously locked up in on-premise data silos.



MONITOR YOUR CLUSTER

DEBUG YOUR APPLICATIONS

Log files generated by EMR Clusters include:

- Step logs
- Hadoop logs
- Bootstrap action logs
- Instance state logs

USE THE MAPR DISTRIBUTION

Amazon EMR with the MapR Distribution for Hadoop

Amazon Elastic MapReduce (Amazon EMR) makes it easy to provision and manage Hadoop in the AWS Cloud. Hadoop is available in multiple distributions and Amazon EMR gives you the option of using the Amazon Distribution or the [MapR Distribution](#) for Hadoop.

MapR delivers on the promise of Hadoop with a proven, enterprise-grade platform that supports a broad set of mission-critical and real-time production uses. MapR brings unprecedented dependability, ease-of-use and world-record speed to Hadoop, NoSQL, database and streaming applications in one unified Big Data platform. MapR is used across financial services, retail, media, healthcare, manufacturing, telecommunications and government organizations as well as by leading Fortune 100 and Web 2.0 companies. Investors include Lightspeed Venture Partners, Mayfield Fund, NEA, and Redpoint Ventures. Connect with MapR on [Facebook](#), [Linkedin](#), and [Twitter](#).



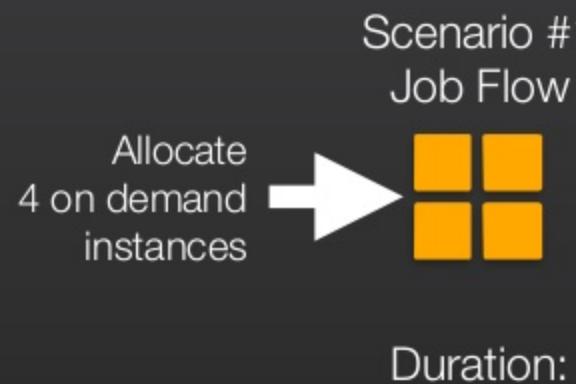
gubdseel vdi es llrw se snoflesasnglo jnemmerwbg bns snotfesdnimcomocafel
beedatglb dquibl strotewvl seisneqmo 0.5 dev bns 00t enutho
jennanD seelnevl jnolqebR bns AEN, bnuE belyeM, seneh9 etunneV
.jelllwT bns ,ulbaeknU ,hooobceE no RqslM trlw

TUNE YOUR CLUSTER FOR COST & PERFORMANCE

Supported EC2 instance types

- General Purpose
- Compute Optimized
- Memory Optimized
- Storage Optimized - D2 instance family 
- D2 instances are available in four sizes with 6TB, 12TB, 24TB, and 48TB storage options.
- GPU Instances

TUNE YOUR CLUSTER FOR COST & PERFORMANCE



4 instances *14 hrs * \$0.50
Total = \$28



4 instances *7 hrs * \$0.50 = \$14 +
5 instances * 7 hrs * \$0.25 = \$8.75
Total = \$22.75

THIRD PARTY TOOLS

MicroStrategy

MAPR

 **Datameer**
Move Data Simple™

 **ATTUNITY**

BI/Visualization

Hadoop Distribution

Graphical IDE

Data Transfer

MORTAR

 **SAP Business Objects**

 **bounday**

 **JASPERSOFT**
THE INTELLIGENCE INSIDE

Integration and Analytics

Business Intelligence

Monitoring

BI/Visualization

 **talend***
open data solutions

 **splunk**

 **Compuware**

 **able**

Graphical IDE

Data Exploration

Performance Tuning

BI/Visualization

Graphical IDE

Data Exploration

Performance Tuning

BI/Visualization

**RESOURCES YOU CAN USE
TO LEARN MORE**

aws.amazon.com/emr

Getting Started with Amazon EMR Tutorial guide:

docs.aws.amazon.com/ElasticMapReduce/latest/DeveloperGuide/emr-get-started.html

Customer Case Studies for Big Data Use-Cases

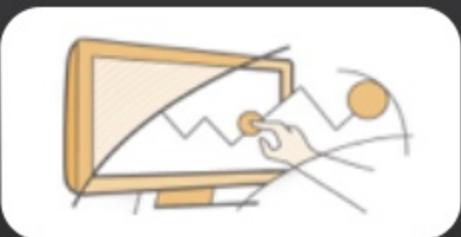
aws.amazon.com/solutions/case-studies/big-data/

Amazon EMR Documentation:

aws.amazon.com/documentation/emr/

AWS Training & Certification

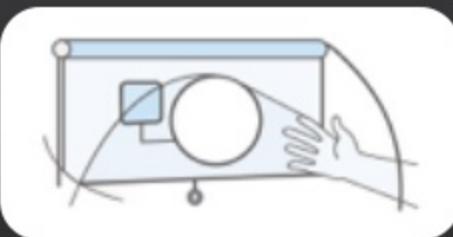
Self-Paced Labs



Try products, gain new skills, and get hands-on practice working with AWS technologies

[aws.amazon.com/training/
self-paced-labs](https://aws.amazon.com/training/self-paced-labs)

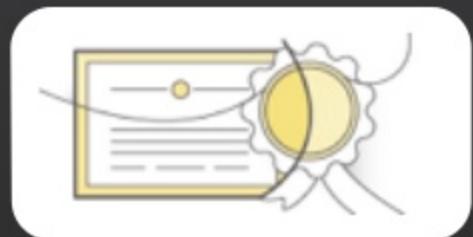
Training



Build technical expertise to design and operate scalable, efficient applications on AWS

aws.amazon.com/training

Certification



Validate your proven skills and expertise with the AWS platform

aws.amazon.com/certification

Follow us for more
events & webinars



Ian Massingham – Technical Evangelist

 @ianMmmm



@AWS_UK1 for local AWS events & news



@AWScloud for Global AWS News & Announcements