

# Deep Dive on Amazon Relational Database Service


Prahlad Rao, Solutions Architect

Sept 13, 2016



# What to expect


- Amazon RDS overview (super quick)
- Security
- Metrics and monitoring
- High availability
- Scaling on RDS
- Backups and snapshots
- Migrating to RDS
- Q&A!



# Amazon Relational Database Service (Amazon RDS)



Launch



No infrastructure  
management




Cost-effective



Application  
compatibility



Instant provisioning



Scale up/down

# Amazon RDS engines

## Commercial

ORACLE®



## Open source



PostgreSQL



MariaDB

## Amazon Aurora

Amazon  
Aurora

# Selected Amazon RDS customers



vodafone

intuit.



SEGA®

*Kempinski*  
HOTELIERS SINCE 1897

OUTBACK  
STEAKHOUSE®

ROVIO



FCBARCELONA  
*més que un club*

OVERSEAS VOTE  
FOUNDATION



Newsweek  
& THE DAILY BEAST

*P*  
The  
Washington Post  
Company

Trinity Mirror plc

coursera



ZUMBA®  
FITNESS

# Selected Amazon Aurora customers




# Trade-offs with a managed service

## Fully managed host and OS

- No access to the database host operating system
- Limited ability to modify configuration that is managed on the host operating system
- No functions that rely on configuration from the host OS

## Fully managed storage

- Max storage limits
  - Microsoft SQL Server—4 TB
  - MySQL, MariaDB, PostgreSQL, Oracle—6 TB
  - Aurora—64 TB
- Growing your database is a process





# Security

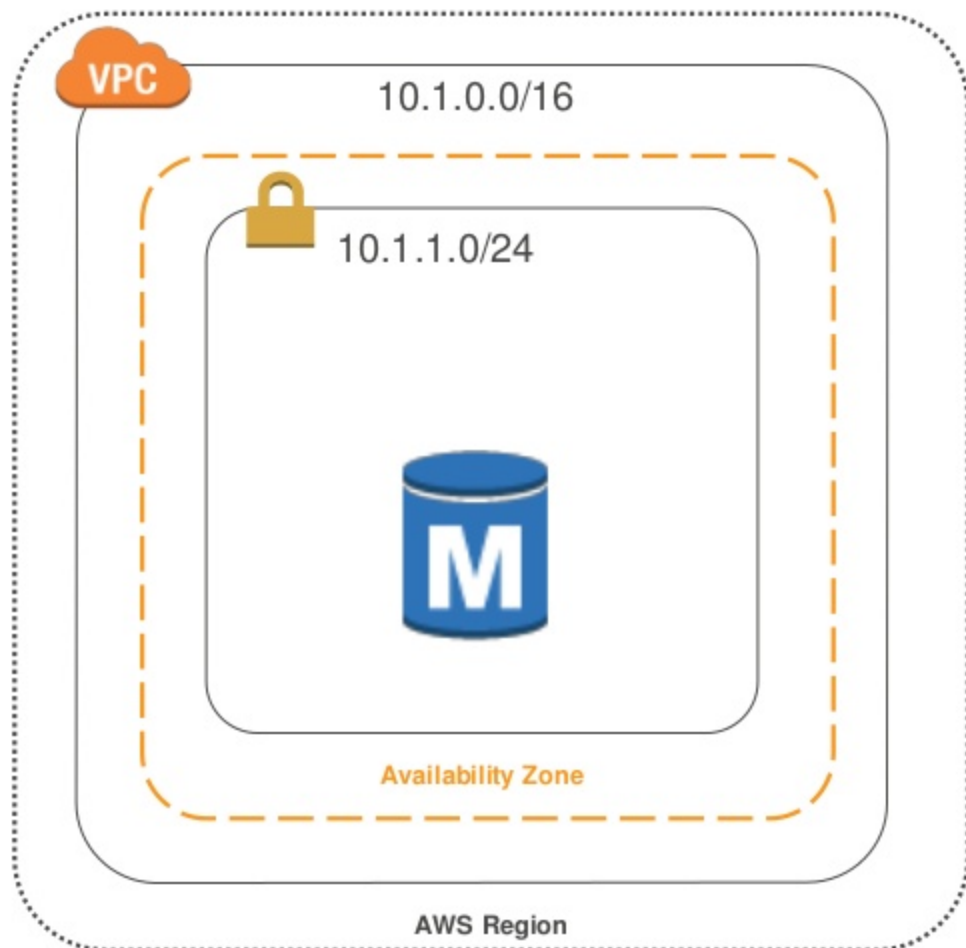




# Amazon Virtual Private Cloud (Amazon VPC)

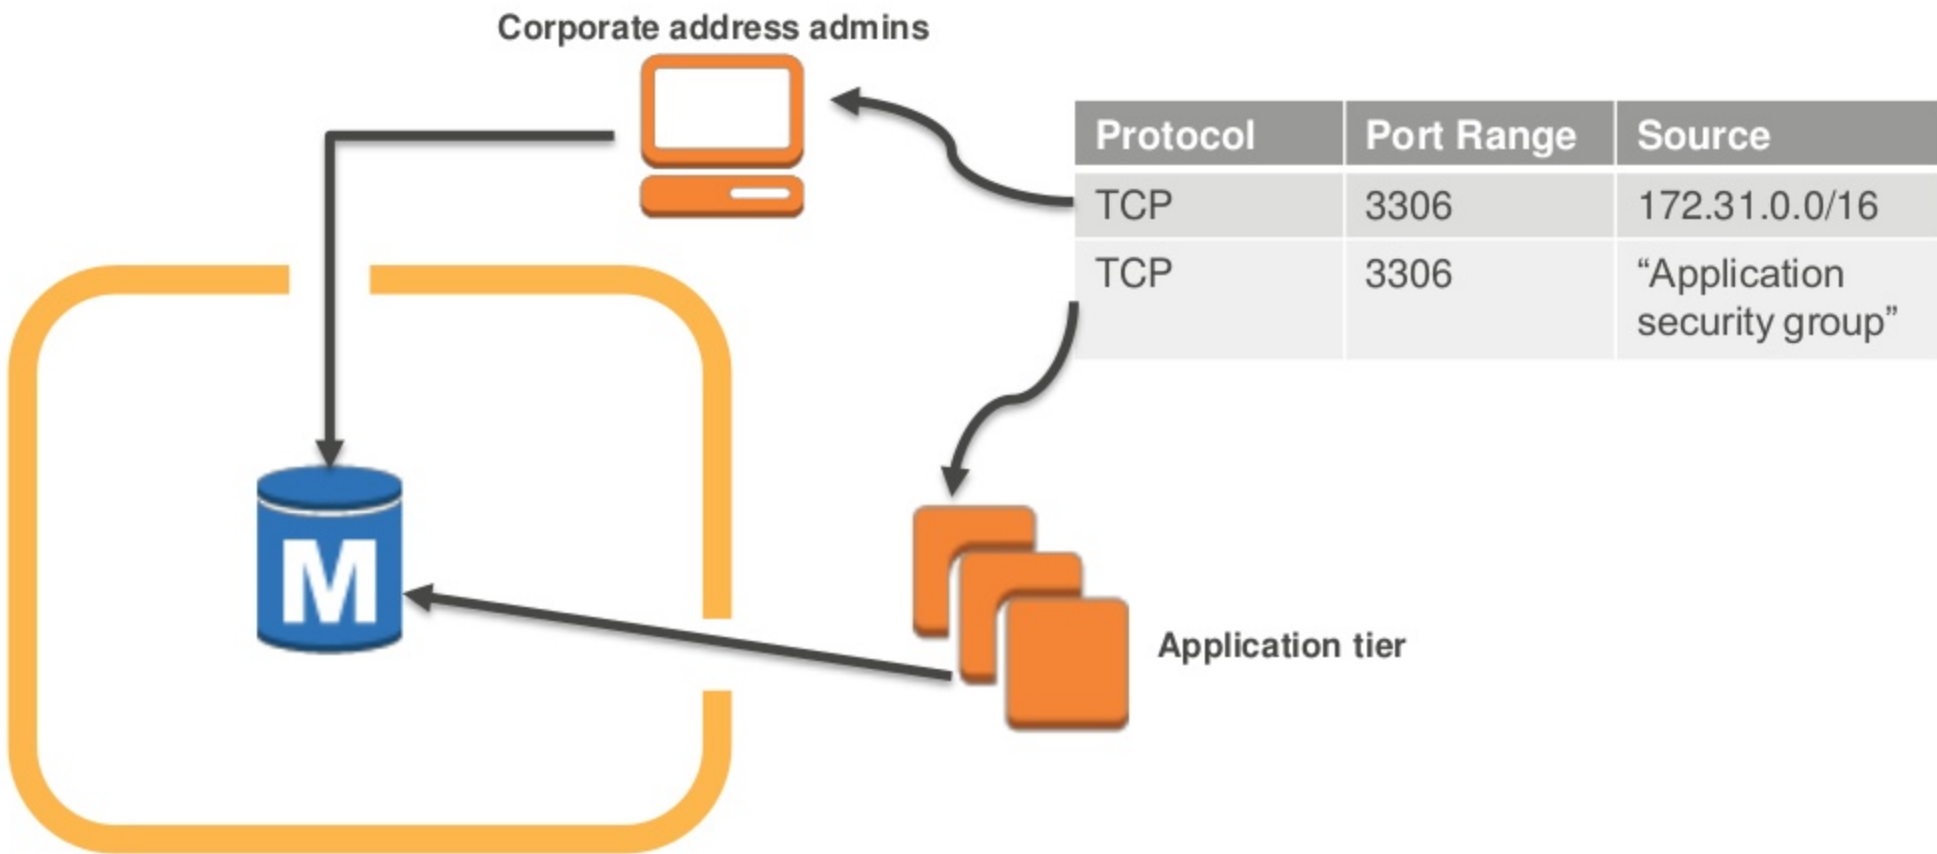
Securely control network configuration

## Manage connectivity




# Security groups

## Database IP firewall protection



# Compliance



Singapore MTCS



27001/9001  
27017/27018

# Compliance

## MySQL and Oracle

- SOC 1, 2, and 3
- ISO 27001/9001
- ISO 27017/27018
- PCI DSS
- FedRAMP
- HIPAA BAA
- UK government programs
- Singapore MTCS


## SQL Server and PostgreSQL

- SOC 1, 2, and 3
- ISO 27001/9001
- ISO 27017/27018
- PCI DSS
- UK government programs
- Singapore MTCS

# SSL


Database traffic encryption

**Available for all six engines**



# At-rest encryption

- DB instance storage
- Automated backups
- Read Replicas
- Snapshots



- **Available for all six engines**
- **No additional cost**
- **Support compliance requirements**


# AWS KMS—RDS standard encryption

Two-tiered key hierarchy using envelope encryption:

- Unique data key encrypts customer data
- AWS KMS master keys encrypt data keys


Benefits:

- Limits risk of compromised data key
- Better performance for encrypting large data
- Easier to manage small number of master keys than millions of data keys
- Centralized access and audit of key activity





# How keys are used to protect your data



1. RDS instance requests encryption key to use to encrypt data, passes reference to master key in account
2. Client request authenticated based on permissions set on both the user and the key
3. A unique data encryption key is created and encrypted under the KMS master key
4. Plaintext and encrypted data key returned to the client
5. Plaintext data key used to encrypt data and then deleted when practical
6. Encrypted data key is stored; it's sent back to KMS when needed for data decryption

# Enabling encryption

AWS Command Line Interface (AWS CLI)

```
aws rds create-db-instance --region us-west-2 --db-instance-identifier sg-cli-test \  
--allocated-storage 20 --storage-encrypted \  
--db-instance-class db.m4.large --engine mysql \  
--master-username myawsuser --master-user-password myawsuser
```

```
aws rds create-db-instance --region us-west-2 --db-instance-identifier sg-cli-test1 \  
--allocated-storage 20 --storage-encrypted --kms-key-id xxxxxxxxxxxxxxxxxxxx \  
--db-instance-class db.m4.large --engine mysql \  
--master-username myawsuser --master-user-password myawsuser
```


# Amazon RDS + AWS KMS useful hints

- You can only encrypt on new database creation
- Encryption cannot be removed
- Master and Read Replica must be encrypted
- Unencrypted snapshots cannot be restored to encrypted DB
  - Aurora will allow this
  - You can create encrypted copies of your unencrypted snapshots
- Cannot restore MySQL to Aurora or Aurora to MySQL
- Cannot copy snapshots or replicate DB across regions


# IAM governed access

You can use AWS Identity and Access Management (IAM) to control who can perform actions on RDS


## Controlled with database grants



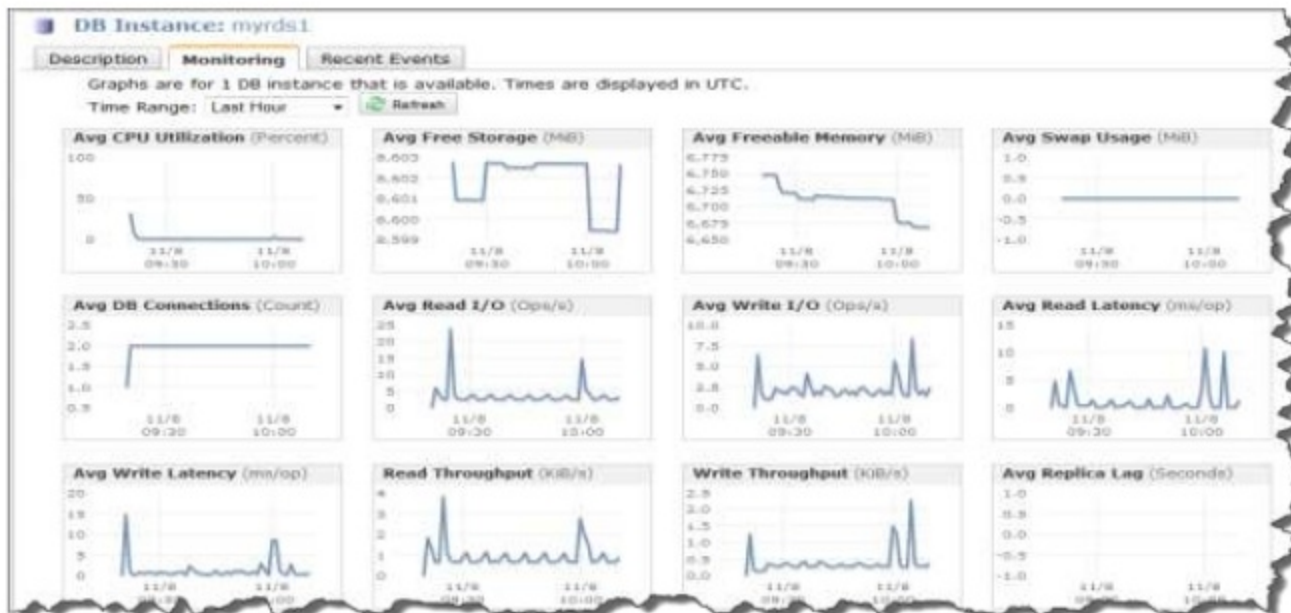
## Controlled with IAM



# Metrics and monitoring



# Standard monitoring



## Amazon CloudWatch metrics for Amazon RDS

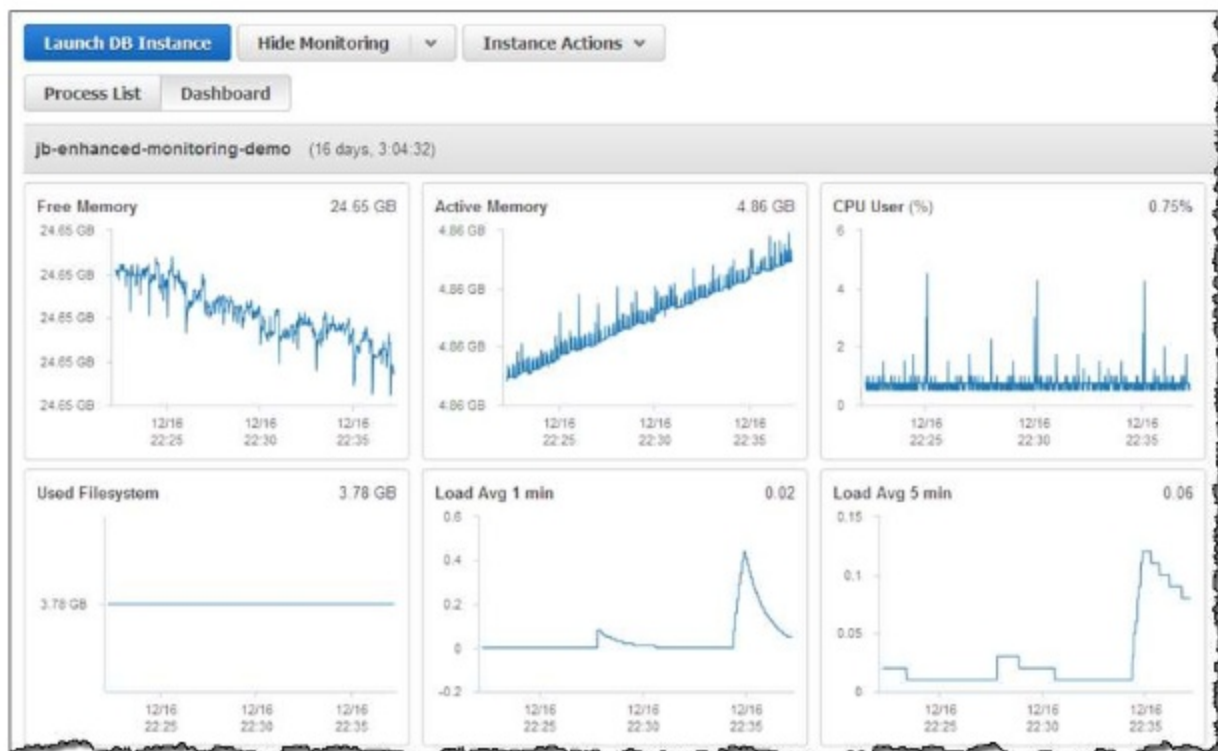
- CPU utilization
- Storage
- Memory
- Swap usage
- DB connections
- I/O (read and write)
- Latency (read and write)
- Throughput (read and write)
- Replica lag
- Many more

## Amazon CloudWatch Alarms

- Similar to on-premises custom monitoring tools

# Enhanced Monitoring

Access to over 50 new CPU, memory, file system, and disk I/O metrics as low as 1 second intervals



**Monitoring**

Enable Enhanced Monitoring

Monitoring Role

Granularity  second(s)

☐ I authorize RDS to create the IAM role rds-monitoring-role.



# Event notifications

- Uses Amazon Simple Notification Service (Amazon SNS) to notify users when an event occurs
- 17 different event categories (availability, backup, configuration change, and so on)

**Create Event Subscription**

Name  ⓘ

Send notifications to  ⓘ [create topic](#)

Source Type  ⓘ

Enabled ☒ Yes ☐ No

**Event Categories**

☐ Select All  
☒ Select specific


availability
backup
configuration change
creation
deletion
failover
failure
low storage
maintenance
notification
recovery
restoration

**DB Instances**


☐ Select All  
☒ Select specific

djr-mysqlexampledb
djr-mysqlexampledb-restore
djr-mysqlexampledb-rr
djr-mysqlexampledb4
djr-rr-v2
djr-rr-v3
djr-sqltest
myvpcdbinstance
sg-dp-target
sg-oracle11204
sg-postgresql1
sg-rest-snap
sg-sqlsvr-ec2

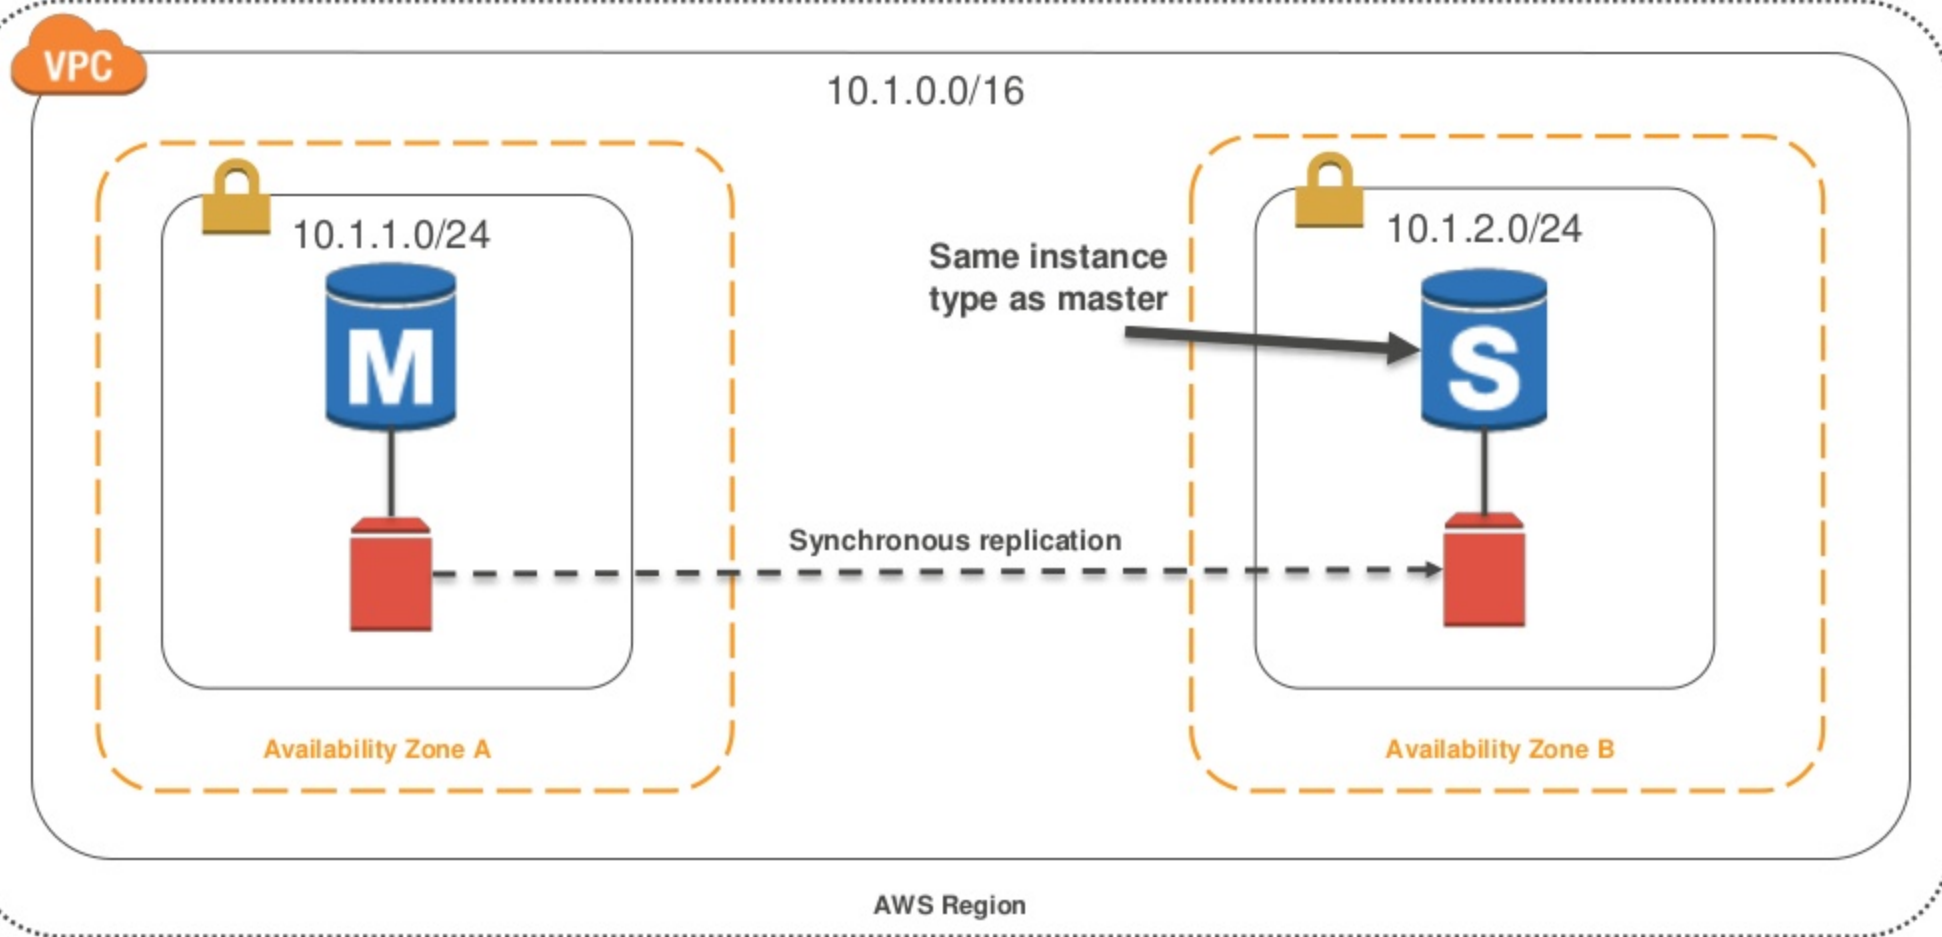
**High availability**




# Minimal deployment—single AZ



# High availability—Multi-AZ




# High availability—Multi-AZ to DNS



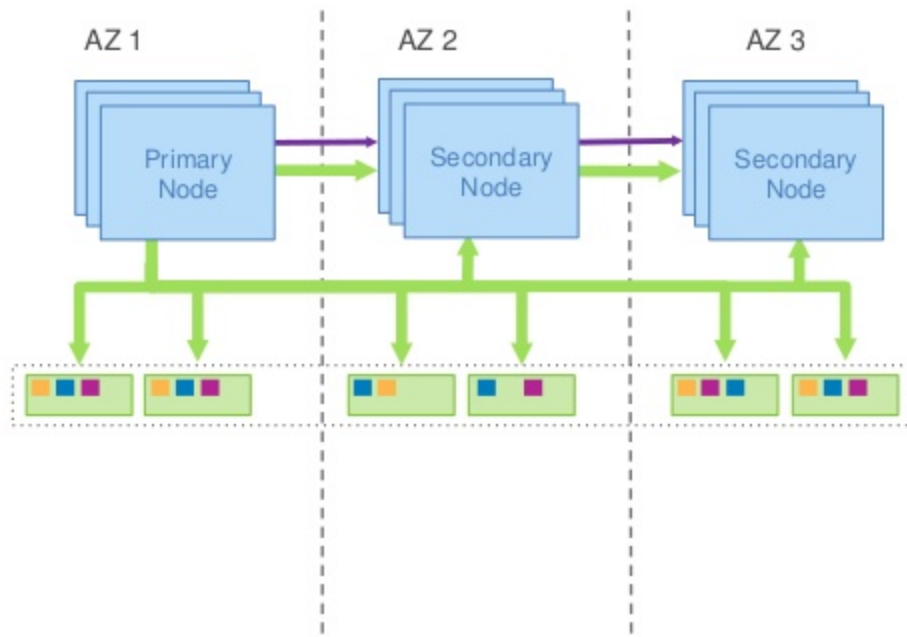
# High availability—Amazon Aurora storage

- Storage volume automatically grows up to 64 TB
- Quorum system for read/write; latency tolerant
- Peer-to-peer gossip replication to fill in holes
- Continuous backup to Amazon S3 (built for 11 9s durability)
- Continuous monitoring of nodes and disks for repair
- 10 GB segments as unit of repair or hotspot rebalance
- Quorum membership changes do not stall writes




# High availability—Aurora nodes

- Aurora cluster contains primary node and up to 15 secondary nodes
- Failing database nodes are automatically detected and replaced
- Failing database processes are automatically detected and recycled
- Secondary nodes automatically promoted on persistent outage, no single point of failure
- Customer application can scale out read traffic across secondary nodes





# Scaling on RDS



# Read Replicas

Bring data close to your customer's applications in different regions

Relieve pressure on your master node for supporting reads and writes

Promote a Read Replica to a master for faster recovery in the event of disaster



# Read Replicas

Within a region

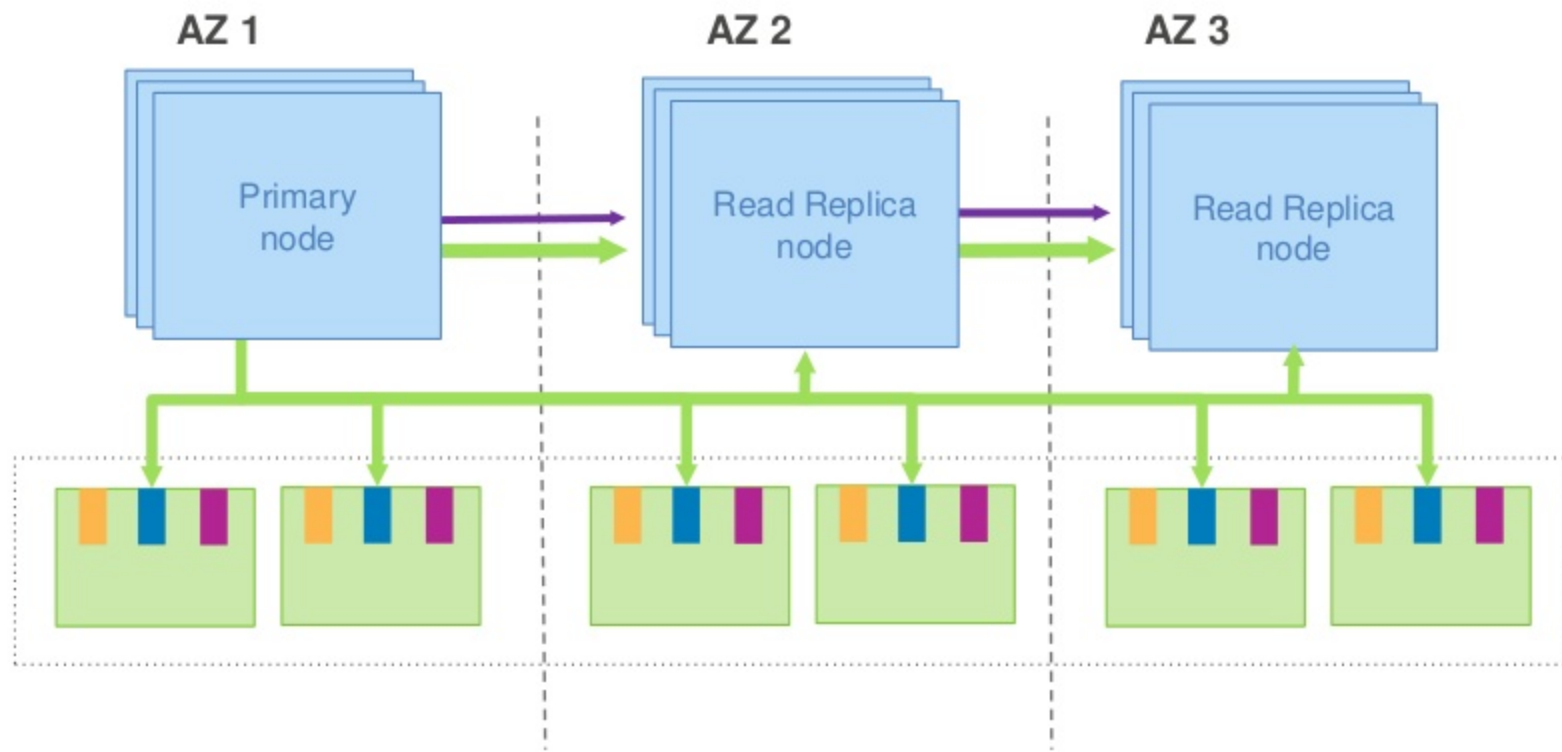
- MySQL
- MariaDB
- PostgreSQL
- Aurora

Cross-region

- MySQL
- MariaDB
- PostgreSQL
- Aurora



# Read Replicas for Amazon Aurora



# Read Replicas—Oracle and SQL Server

## Options


- Oracle GoldenGate
- Third-party replication products
- Snapshots

**ORACLE®**



# Scaling up—or down

- Handle higher load or lower usage
- Control costs



# Scaling up—or down

AWS Management Console

Instance Actions ▾

- See Details
- Create Read Replica
- Promote Read Replica
- Take Snapshot
- Restore to Point in Time
- Migrate Latest Snapshot
- Modify**
- Reboot
- Delete

## Modify DB Instance: sg-cli-test

### Instance Specifications

DB Engine Version	MySQL 5.6.27 (default)	⌵
DB Instance Class	db.m4.large — 2 vCPU, 8 GiB RAM	⌵
Multi-AZ Deployment	No	⌵
Storage Type	General Purpose (SSD)	⌵
Allocated Storage*	600	GB



Apply Immediately





# Scaling—single AZ

With single AZ deployment, the master takes an outage

## Alarms and Recent Events

TIME (UTC-7)	EVENT
Mar 26 7:01 AM	DB instance restarted
Mar 26 7:00 AM	Finished applying modification to DB instance class
Mar 26 6:53 AM	Applying modification to database instance class

# Scaling—Multi-AZ

With Multi-AZ, the standby gets upgraded first

## Alarms and Recent Events

TIME (UTC-7)	EVENT
Mar 26 6:34 AM	Finished applying modification to DB instance class
Mar 26 6:28 AM	Multi-AZ instance failover completed
Mar 26 6:28 AM	DB instance restarted
Mar 26 6:28 AM	Multi-AZ instance failover started
Mar 26 6:20 AM	Applying modification to database instance class

dbinstancenam

n:3006



# Scaling—automation

## AWS CLI

```
aws rds modify-db-instance --db-instance-identifier sg-cli-test --db-instance-class db.m4.large --apply-immediately
```

## Scheduled CLI—cron

```
#Scale down at 8:00 PM on Friday  
0 20 * * 5 /home/ec2-user/scripts/scale_down_rds.sh  
  
#Scale up at 4:00 AM on Monday  
0 4 * * 1 /home/ec2-user/scripts/scale_up_rds.sh
```

# Scaling—automation

## Scheduled—AWS Lambda

No server but still runs on a schedule!

```
import boto3

client=boto3.client('rds')

def lambda_handler(event, context):
    response=client.modify_db_instance(DBInstanceIdentifier='sg-cli-test',
                                       DBInstanceClass='db.m4.xlarge',
                                       ApplyImmediately=True)

    print response
```

# Scaling—automation

## Metrics-based scaling

- Amazon CloudWatch and AWS Lambda!



# Scaling—automation

```
import boto3
import json

client=boto3.client('rds')


def lambda_handler(event, context):

    message = event['Records'][0]['Sns']['Message']
    parsed_message=json.loads(message)
    db_instance=parsed_message['Trigger']['Dimensions'][0]['value']
    print 'DB Instance: ' + db_instance

    response=client.modify_db_instance(DBInstanceIdentifier=db_instance,
                                       DBInstanceClass='db.m4.large',
                                       ApplyImmediately=True)

    print response
```

# Backups and snapshots





# Backups

## MySQL, PostgreSQL, MariaDB, Oracle, SQL Server

- Scheduled daily backup of entire instance
- Archive database change logs
- 35 day retention for backups
- Multiple copies in each AZ where you have instances for a deployment

## Aurora

- Automatic, continuous, incremental backups
- Point-in-time restore
- No impact on database performance
- 35 day retention





# Restoring

- Restoring creates an entire new database instance
- You define all the instance configuration just like a new instance

## Restore DB Instance

You are creating a new DB Instance from a source DB Instance at a specified time. This new DB Instance will have the default DB Security Group and DB Parameter Groups.


This feature is currently supported for InnoDB storage engine only. If you are using MyISAM, refer to details [here](#).

**Use Latest Restorable Time** ☒ March 8, 2016 at 12:10:00 PM UTC-8

**Use Custom Restore Time** ☐   :  :  UTC-8

# Snapshots


- Full copies of your Amazon RDS database that are different from your scheduled backups
- Backed by Amazon S3
- Used to create a new RDS instance
- Remain encrypted if using encryption




# Snapshots

## Use cases

- Resolve production issues
- Nonproduction environments
- Point-in-time restore
- Final copy before terminating a database
- Disaster recovery
- Cross-region copy
- Copy between accounts



**Migrating onto RDS**





## AWS Database Migration Service




ORACLE

Amazon Aurora



- ✓ Move data to the same or different database engine
- ✓ Keep your apps running during the migration
- ✓ Start your first migration in 10 minutes or less
- ✓ Replicate within, to, or from Amazon EC2 or RDS



## AWS Schema Conversion Tool

- ✓ Migrate from Oracle and SQL Server
- ✓ Move your tables, views, stored procedures, and data manipulation language (DML) to MySQL, MariaDB, and Aurora
- ✓ Highlight where manual edits are needed

# Thank you!