




# Apache Kudu\*: Fast Analytics on Fast Data

Todd Lipcon (Kudu team lead) – [todd@cloudera.com](mailto:todd@cloudera.com)  
[@tlipcon](https://twitter.com/tlipcon)

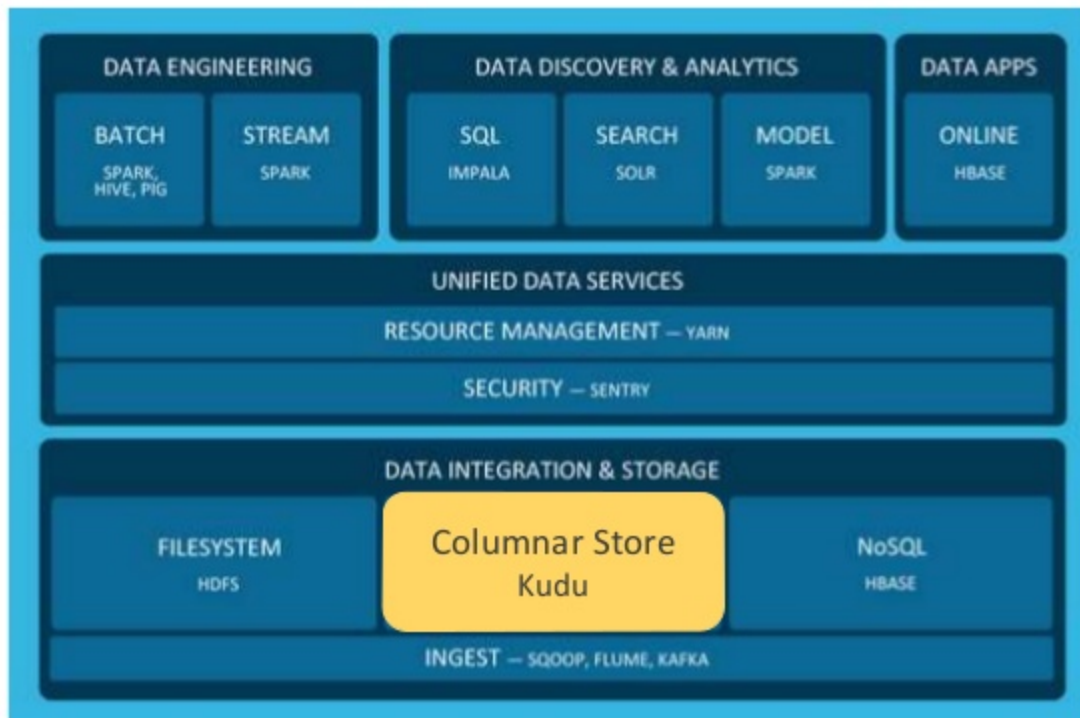
**Tweet about this talk:** [@apachekudu](https://twitter.com/apachekudu) or [#kudu](https://twitter.com/kudu)

\* Incubating at the Apache Software Foundation



# Apache Kudu

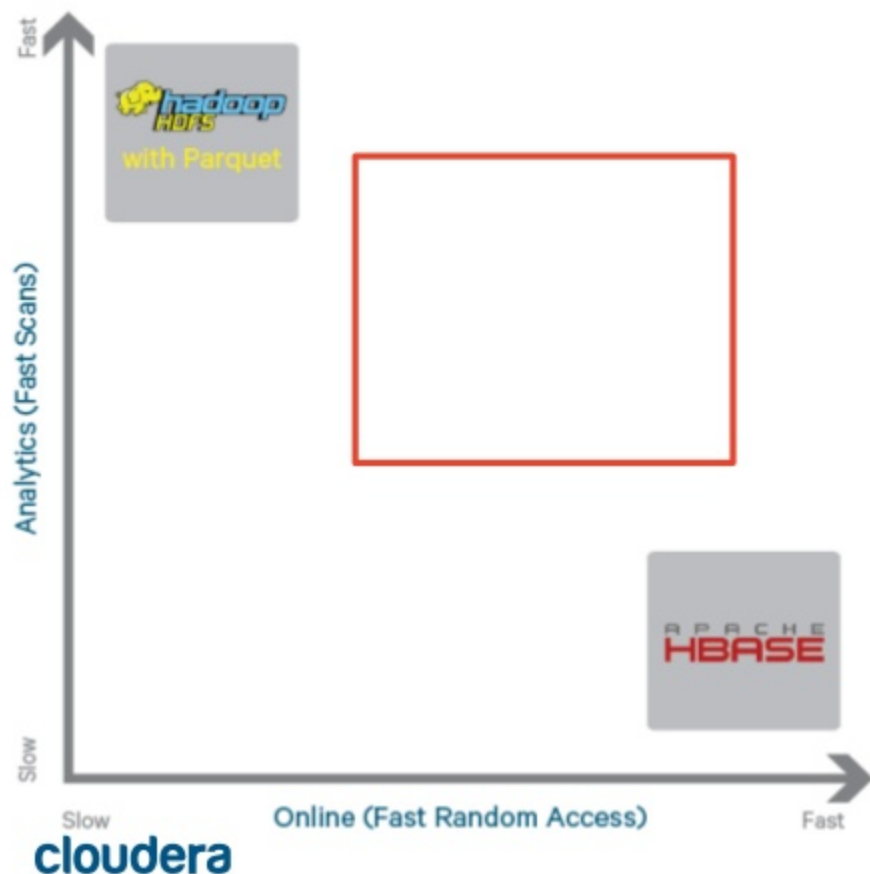
Storage for Fast Analytics on Fast Data



- New updatable column store for Hadoop
- Apache-licensed open source
- Beta now available

# Why Kudu?

# Current Storage Landscape in Hadoop Ecosystem



**HDFS** (GFS) excels at:

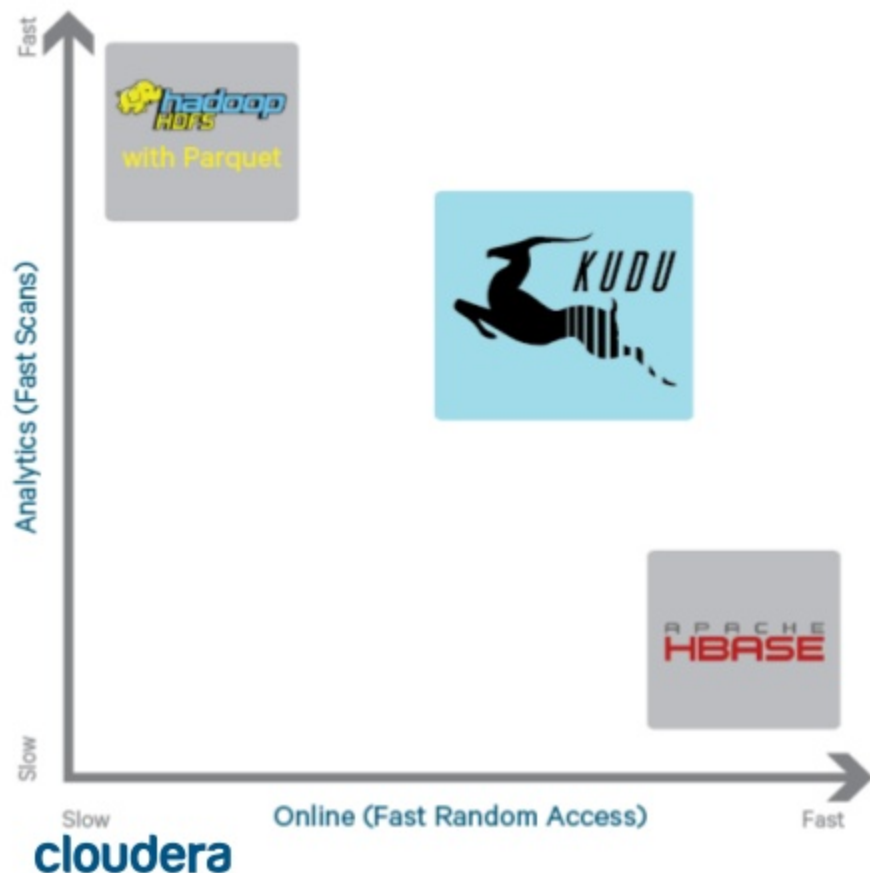
- Batch ingest only (eg hourly)
- Efficiently scanning large amounts of data (analytics)

**HBase** (BigTable) excels at:

- Efficiently finding and writing individual rows
- Making data mutable

Gaps exist when these properties are needed *simultaneously*

# Kudu Design Goals



- **High throughput** for big scans  
*Goal:* Within 2x of Parquet
- **Low-latency** for short accesses  
*Goal:* 1ms read/write on SSD
- **Database-like** semantics (initially single-row ACID)
- **Relational data model**
  - SQL queries are easy
  - “NoSQL” style scan/insert/update (Java/C++ client)

# Changing Hardware landscape

- **Spinning disk -> solid state storage**
  - **NAND flash:** Up to 450k read 250k write iops, about 2GB/sec read and 1.5GB/sec write throughput, at a price of less than \$3/GB and dropping
  - **3D XPoint memory** (1000x faster than NAND, cheaper than RAM)
- **RAM** is cheaper and more abundant:
  - 64->128->256GB over last few years
- **Takeaway:** The **next bottleneck is CPU**, and current storage systems weren't designed with CPU efficiency in mind.

# What's Kudu?

# Scalable and Fast Tabular Storage

- **Scalable**
  - Tested up to 275 nodes (~3PB cluster)
  - Designed to scale to **1000s of nodes, tens of PBs**
- **Fast**
  - **Millions** of read/write operations per second across cluster
  - **Multiple GB/second** read throughput per node
- **Tabular**
  - **SQL-like** schema: **finite number** of **typed** columns (unlike HBase/Cassandra)
  - **Fast ALTER TABLE**
  - **“NoSQL”** APIs: Java/C++/Python **or SQL** (Impala/Spark/etc)



# Use cases and architectures

# Kudu Use Cases

**Kudu is best for use cases requiring a simultaneous combination of sequential and random reads and writes**

- **Time Series**

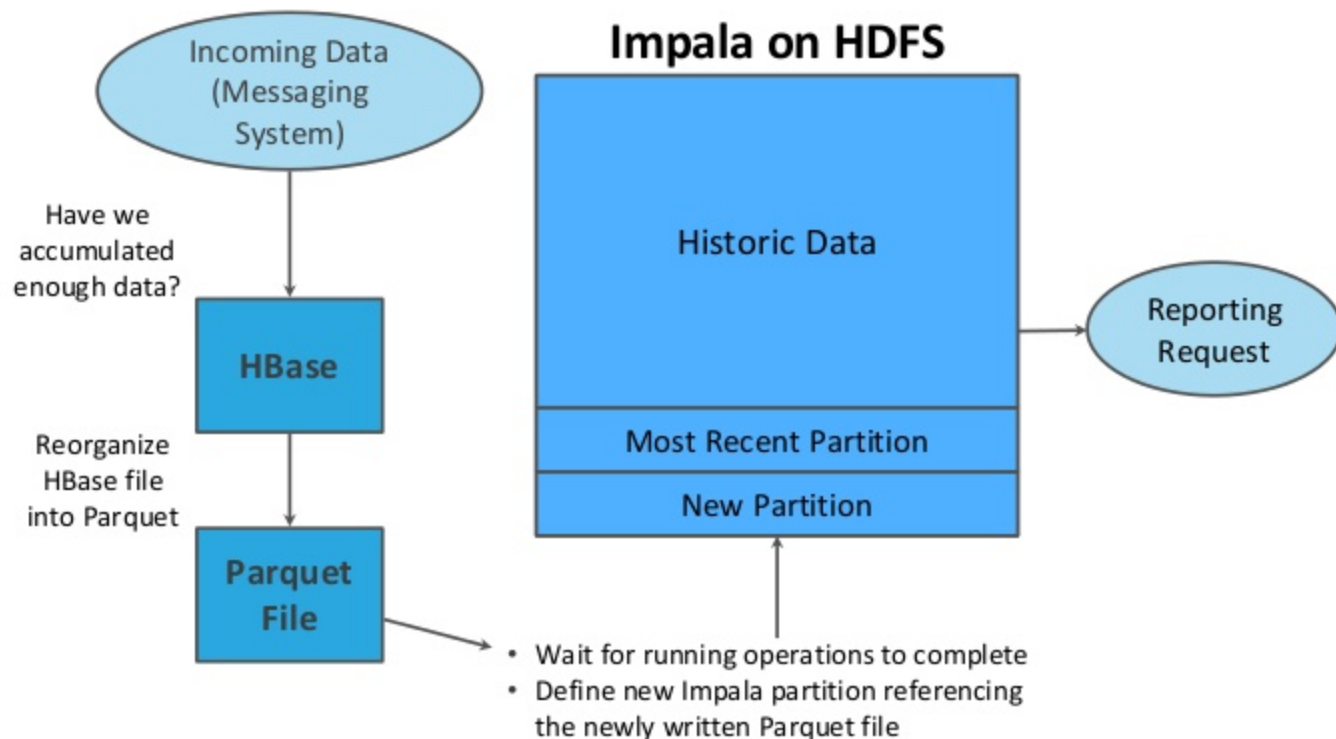
- Examples: Stream market data; fraud detection & prevention; network monitoring
- Workload: Insert, updates, scans, lookups

- **Online Reporting**

- Examples: ODS
- Workload: Inserts, updates, scans, lookups

# Real-Time Analytics in Hadoop Today


Fraud Detection in the Real World = Storage Complexity



## Considerations:

- How do I handle failure during this process?
- How often do I reorganize data streaming in into a format appropriate for reporting?
- When reporting, how do I see data that has not yet been reorganized?
- How do I ensure that important jobs aren't interrupted by maintenance?

# Real-Time Analytics in Hadoop with Kudu



## Improvements:

- One system to operate
- No cron jobs or background processes
- Handle late arrivals or data corrections with ease
- New data available immediately for analytics or operations

# Xiaomi Use Case

- World's 4<sup>th</sup> largest smart-phone maker (most popular in China)
- Gather important RPC tracing events from mobile app and backend service.
- Service monitoring & troubleshooting tool.



## High write throughput

- >5 Billion records/day and growing



## Query latest data and quick response

- Identify and resolve issues quickly



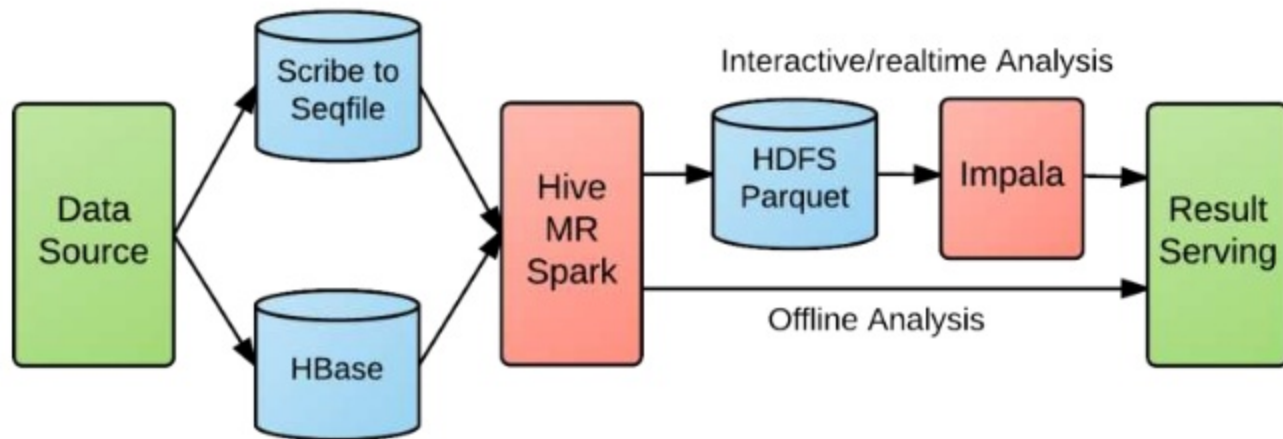
## Can search for individual records

- Easy for troubleshooting



# Xiaomi Big Data Analytics Pipeline

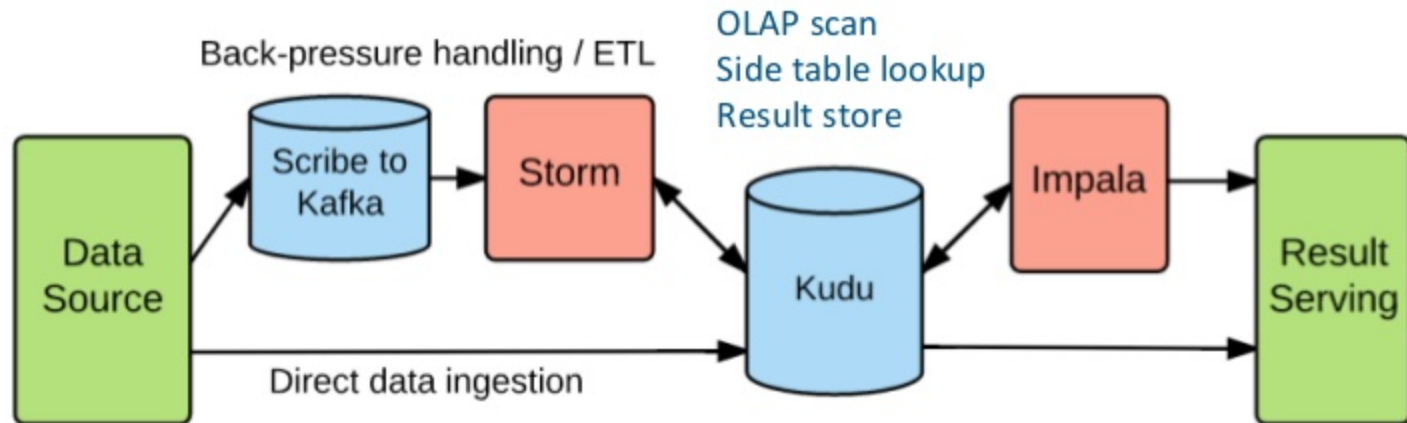
Before Kudu



- **Long pipeline**  
high latency(1 hour ~ 1 day), data conversion pains
- **No ordering**  
Log arrival(storage) order not exactly logical order  
e.g. read 2-3 days of log for data in 1 day

# Xiaomi Big Data Analysis Pipeline

Simplified With Kudu



- ETL Pipeline(0~10s latency)  
Apps that need to prevent backpressure or require ETL
- Direct Pipeline(no latency)  
Apps that don't require ETL and no backpressure issues

# How it Works

Replication and fault tolerance




# Tables, Tablets, and Tablet Servers

- Table is **horizontally partitioned into *tablets***
  - *Range* or *hash* partitioning
  - `PRIMARY KEY (host, metric, timestamp) DISTRIBUTE BY HASH(timestamp) INTO 100 BUCKETS`
    - `bucketNumber = hashCode(row['timestamp']) % 100`
- Each tablet has N **replicas** (3 or 5), with **Raft consensus**
  - Automatic **fault tolerance**
  - MTTR: ~5 seconds
- **Tablet servers** host tablets on local disk drives

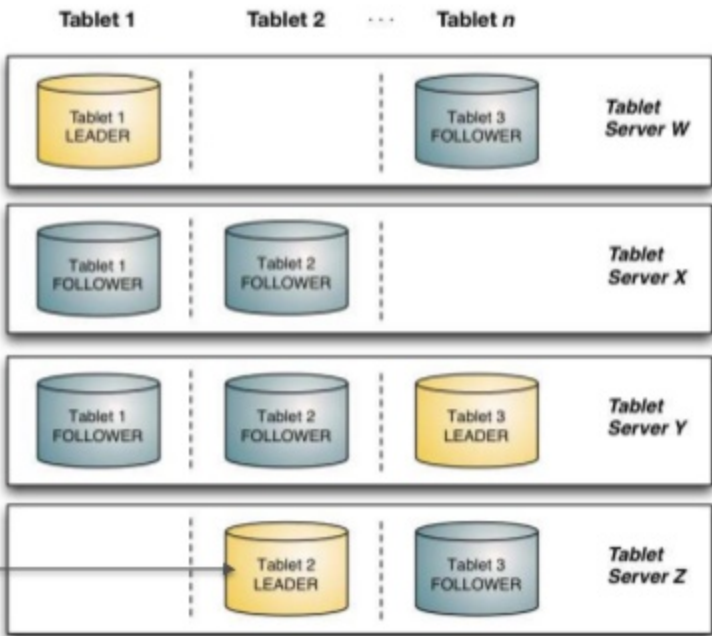

# Metadata and the Master

- **Replicated master**
  - Acts as a tablet directory
  - Acts as a catalog (which tables exist, etc)
  - Acts as a load balancer (tracks TS liveness, re-replicates under-replicated tablets)
- **Not a bottleneck**
  - super fast in-memory lookups



Hey Master! Where is the row for 'tlipcon' in table "T"?

It's part of tablet 2, which is on servers {Z,Y,X}.  
BTW, here's info on other tablets you might care about: T1, T2, T3, ...



UPDATE tlipcon  
SET col=foo

# Raft Consensus



# How it Works

Columnar storage

# Columnar Storage

Twitter Firehose Table			
tweet_id	user_name	created_at	text
INT64	STRING	TIMESTAMP	STRING
23059873	newsycbot	1442865158	Visual Explanation of the Raft Consensus Algorithm <a href="http://bit.ly/1DOUac0">http://bit.ly/1DOUac0</a> (cmts <a href="http://bit.ly/1HKmjfc">http://bit.ly/1HKmjfc</a> )
22309487	RideImpala	1442828307	Introducing the Ibis project: for the Python experience at Hadoop Scale
23059861	fastly	1442865156	Missed July's SF @papers_we_love? You can now watch @el_bhs talk about @google's globally-distributed database: <a href="http://fastly.us/1eVz8MM">http://fastly.us/1eVz8MM</a>
23010982	llvmorg	1442865155	LLVM 3.7 is out! Get it while it's HOT! <a href="http://llvm.org/releases/download.html#3.7.0">http://llvm.org/releases/download.html#3.7.0</a>

Tweet\_id

{23059873,  
22309487,  
23059861,  
23010982}

User\_name

{newsycbot,  
RideImpala,  
fastly,  
llvmorg}

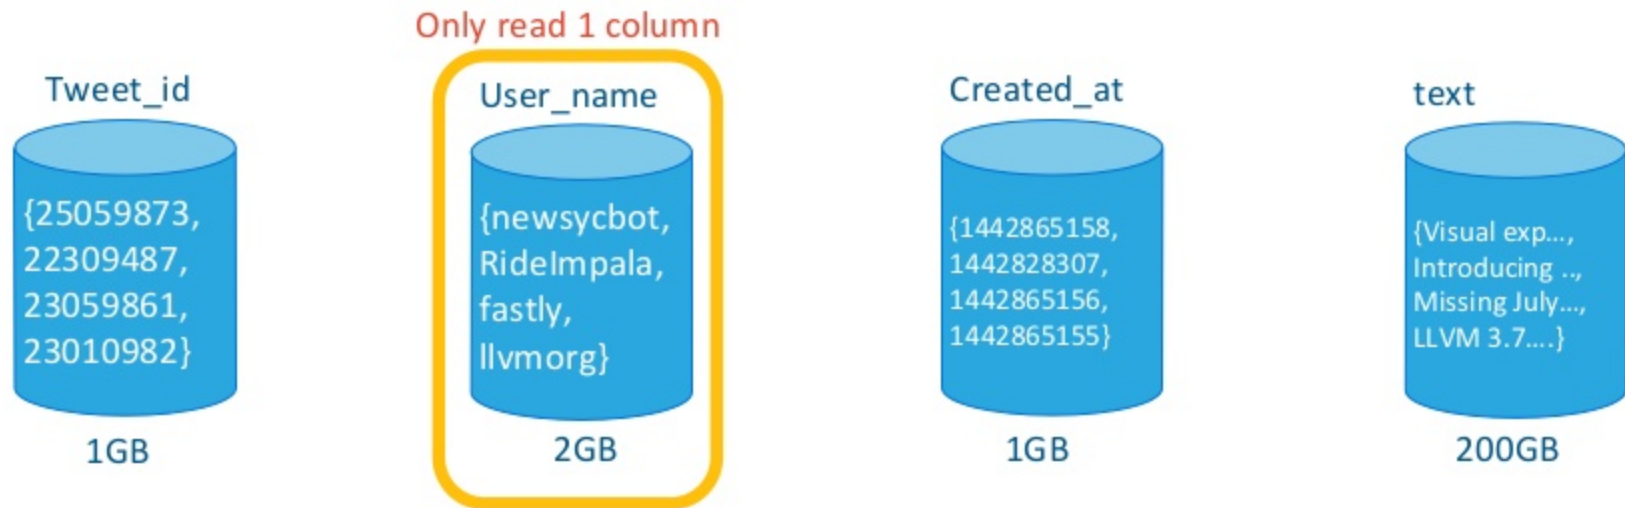
Created\_at

{1442865158,  
1442828307,  
1442865156,  
1442865155}

text

{Visual exp...,  
Introducing ...,  
Missing July...,  
LLVM 3.7....}

# Columnar Storage




```
SELECT COUNT(*) FROM tweets WHERE user_name = 'newsycbot';
```



# Columnar Compression

Created\_at



Created_at	Diff(created_at)
1442865158	n/a
1442828307	-36851
1442865156	36849
1442865155	-1
64 bits each	17 bits each

- **Many columns can compress to a few bits per row!**
- Especially:
  - Timestamps
  - Time series values
  - Low-cardinality strings
- **Massive space savings and throughput increase!**



# Handling Inserts and Updates

- Inserts go to an in-memory row store (MemRowSet)
  - Durable due to write-ahead logging
  - Later flush to columnar format on disk
- Updates go to in-memory “delta store”
  - Later flush to “delta files” on disk
  - Eventually “compact” into the previously-written columnar data files
- Details elided here due to time constraints
  - available in other slide decks online, or come to office hours to learn more!

# Integrations

## Spark DataSource Integration (WIP)

```
sqlContext.load("org.kududb.spark",  
    Map("kudu.table" -> "foo",  
        "kudu.master" -> "master.example.com"))  
    .registerTempTable("mytable")  
df = sqlContext.sql(  
    "select col_a, col_b from mytable " +  
    "where col_c = 123")
```

Available in Kudu 0.7.0, but still being improved

# Impala Integration

- `CREATE TABLE ... DISTRIBUTE BY HASH(col1) INTO 16 BUCKETS AS SELECT ... FROM ...`
- `INSERT/UPDATE/DELETE`
- Optimizations like predicate pushdown, scan parallelism, more on the way
- Not an Impala user? Community working on other integrations (Hive, Drill, Presto, Phoenix)

## MapReduce Integration


- Multi-framework cluster (MR + HDFS + Kudu on the same disks)
- **KuduTableInputFormat / KuduTableOutputFormat**
  - Support for pushing predicates, column projections, etc

# Performance

# TPC-H (Analytics benchmark)

- 75 server cluster
  - 12 (spinning) disk each, enough RAM to fit dataset
  - TPC-H Scale Factor 100 (100GB)
- Example query:
  - ```
SELECT n_name, sum(l_extendedprice * (1 - l_discount)) as revenue FROM customer,
orders, lineitem, supplier, nation, region WHERE c_custkey = o_custkey AND
l_orderkey = o_orderkey AND l_suppkey = s_suppkey AND c_nationkey = s_nationkey
AND s_nationkey = n_nationkey AND n_regionkey = r_regionkey AND r_name = 'ASIA'
AND o_orderdate >= date '1994-01-01' AND o_orderdate < '1995-01-01' GROUP BY
n_name ORDER BY revenue desc;
```

## TPC-H SF 100 @75 nodes




- Kudu outperforms Parquet by 31% (geometric mean) for RAM-resident data




# Versus other NoSQL Storage

- **Phoenix: SQL layer on HBase**
- 10 node cluster (9 worker, 1 master)
- TPC-H LINEITEM table only (6B rows)



# What about NoSQL-style Random Access? (YCSB)

- **YCSB** 0.5.0-snapshot
- 10 node cluster  
(9 worker, 1 master)
- 100M row data set
- 10M operations each workload



# Getting started

# Project Status

- Open source beta released in September
- Latest release 0.7.1 hot off the presses
  - Usable for many applications
  - Have not experienced unrecoverable data loss, reasonably stable (almost no crashes reported). Users testing up to 200 nodes so far.
  - Still requires some expert assistance, and you'll probably find some bugs
- Part of the **Apache Software Foundation** Incubator
  - Community-driven open source process

# Apache Kudu Community

**cloudera**




# Getting Started As a User

- <http://getkudu.io>
- [user@kudu.incubator.apache.org](mailto:user@kudu.incubator.apache.org)
- <http://getkudu-slack.herokuapp.com/>
  
- Quickstart VM
  - Easiest way to get started
  - Impala and Kudu in an easy-to-install VM
- CSD and Parcels
  - For installation on a Cloudera Manager-managed cluster

# Getting Started As a Developer

- <http://github.com/apache/incubator-kudu>
- Code reviews: <http://gerrit.cloudera.org>
- Public JIRA: <http://issues.apache.org/jira/browse/KUDU>
  - Includes bugs going back to 2013. Come see our dirty laundry!
- Mailing list: [dev@kudu.incubator.apache.org](mailto:dev@kudu.incubator.apache.org)
  
- Apache 2.0 license open source
- Contributions are welcome and encouraged!



<http://getkudu.io/>  
[@getkudu](#)

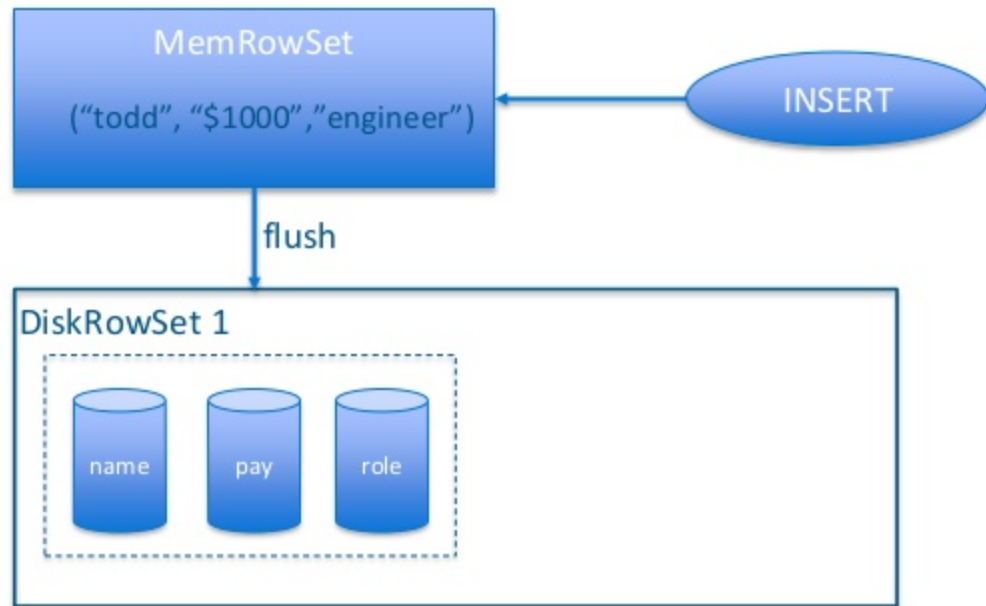


# Backup slides

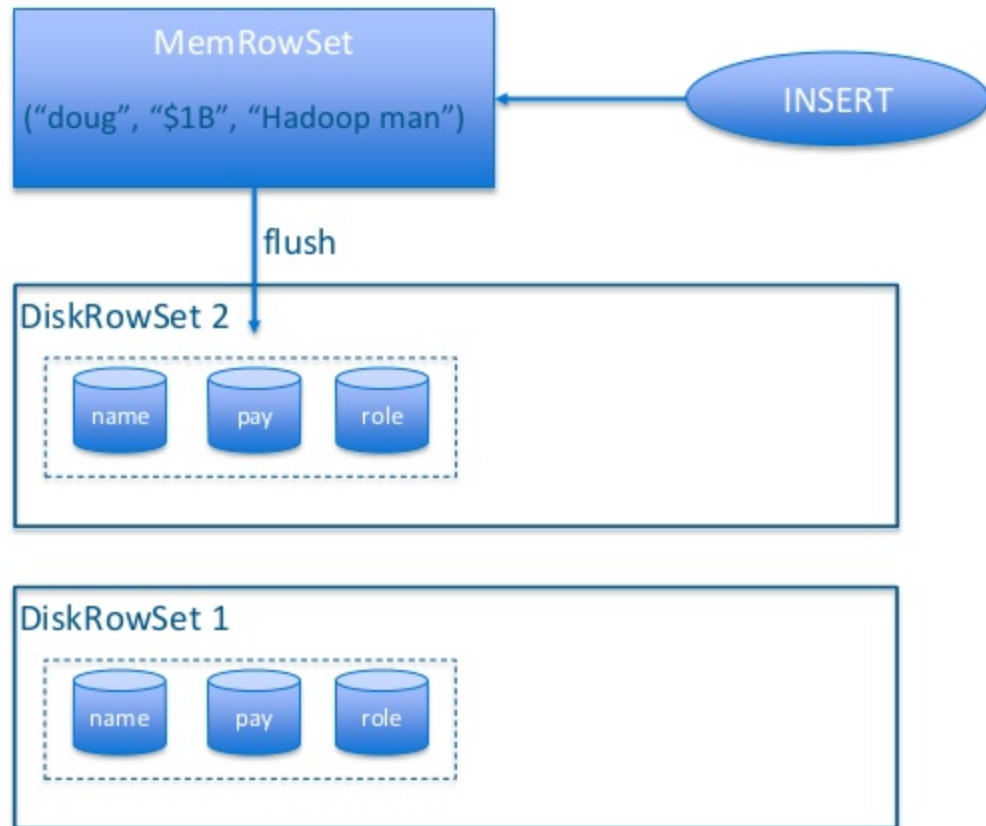
# How it works

Write and read paths

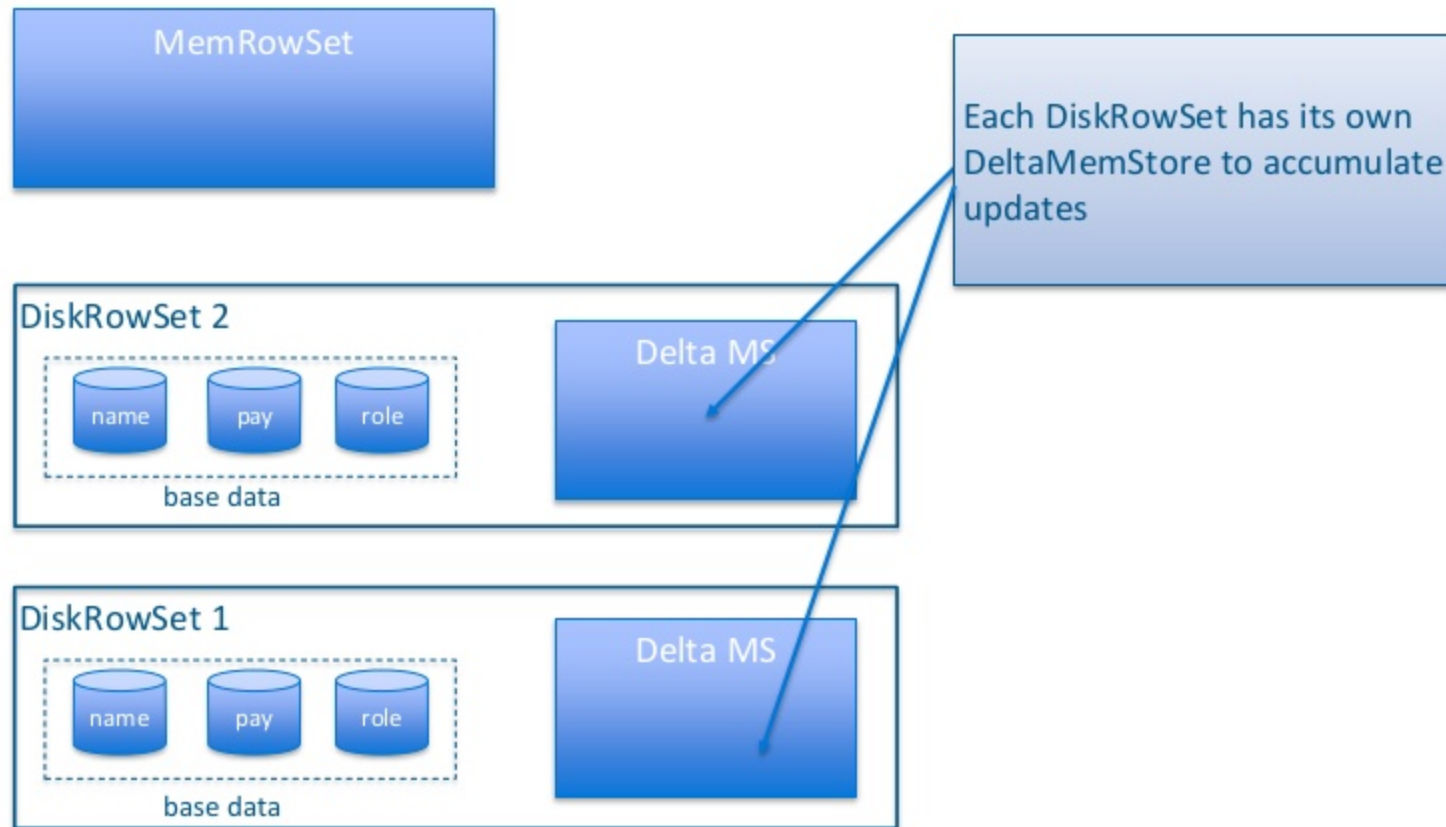
# Kudu storage – Inserts and Flushes



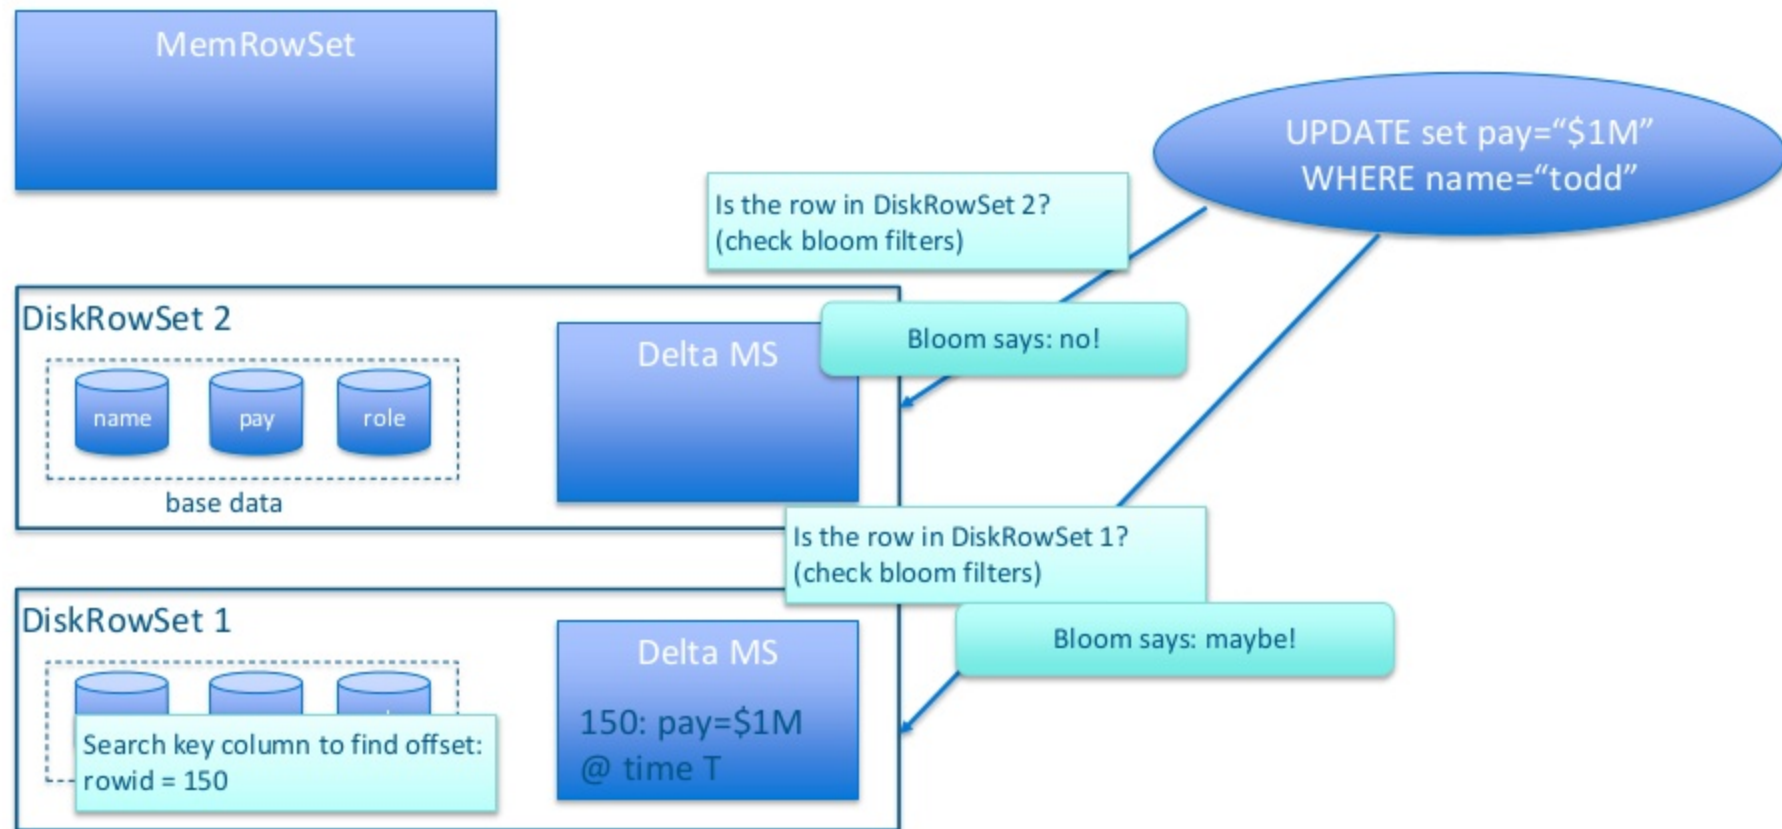
# Kudu storage – Inserts and Flushes




# Kudu storage - Updates




# Kudu storage - Updates



# Kudu storage – Read path





Read rows in DiskRowSet 2



Then, read DiskRow

Updates are applied based on ordinal offset within DRS: array indexing = fast

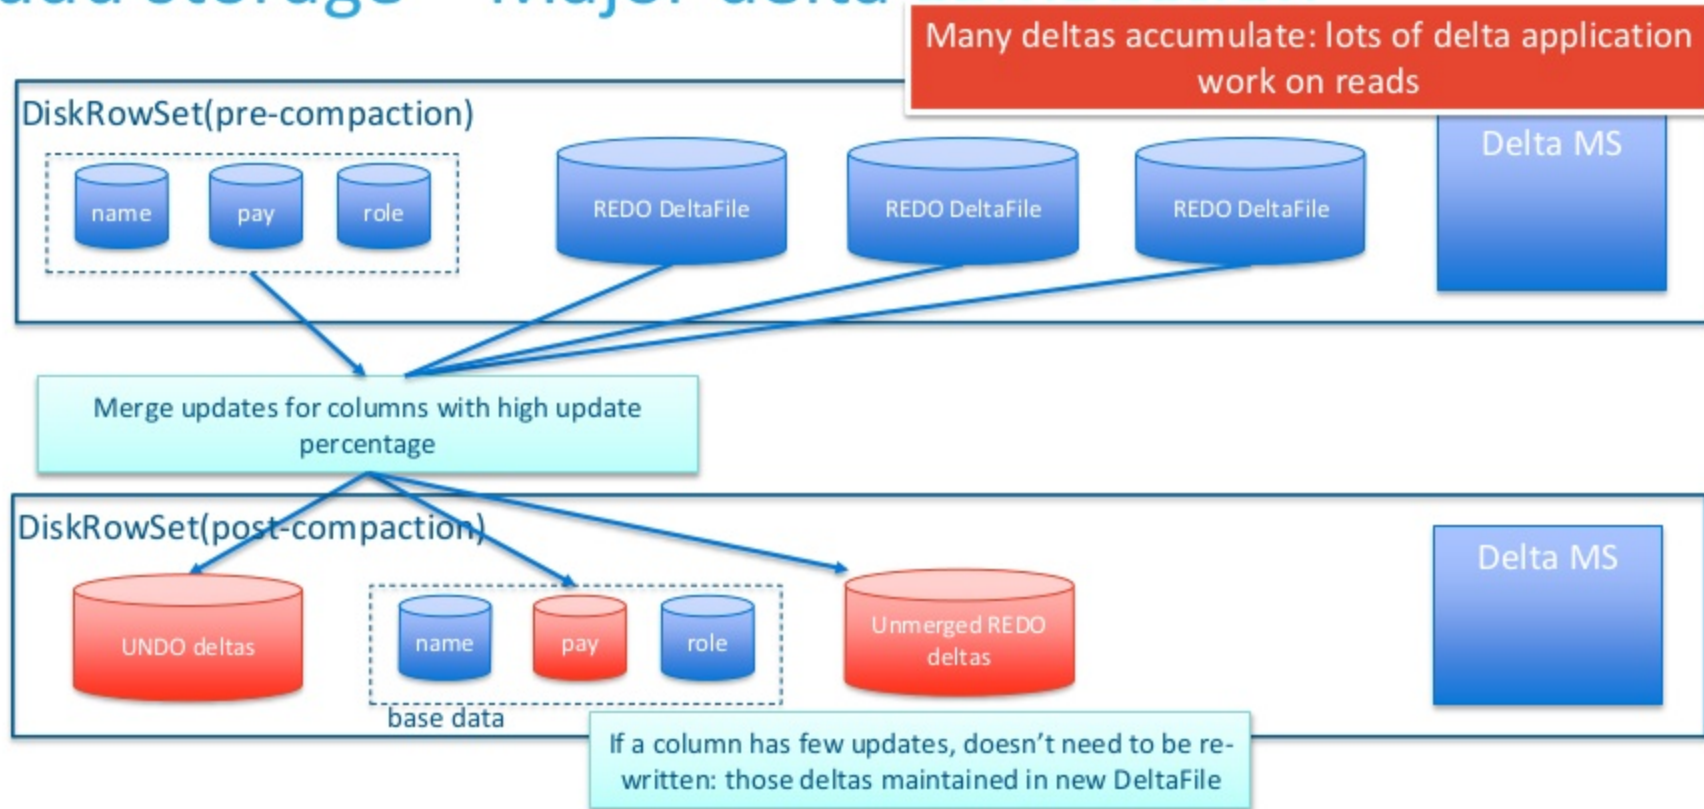
# Kudu storage – Delta flushes



A **REDO** delta indicates how to transform between the 'base data' (columnar) and a later version



# Kudu storage – Major delta compaction



# Kudu storage – RowSet Compactions

