AWS

SUMMIT

Building Your Data Lake on AWS

Ben Snively, Senior Solutions Architect, AWS

August 14, 2017

amazon web services

What is a Data Lake?

An architectural approach that allows you to store massive amounts of “raw” data into a central location

It’s readily available to be categorized, processed, analyzed, and consumed by diverse groups

Why use a Data Lake?

Leverage **all data** within your organization

- Customer centricity
- Business agility
- Better predictions
- Competitive advantage

Leads to…

Legacy data architectures exist as isolated data silos

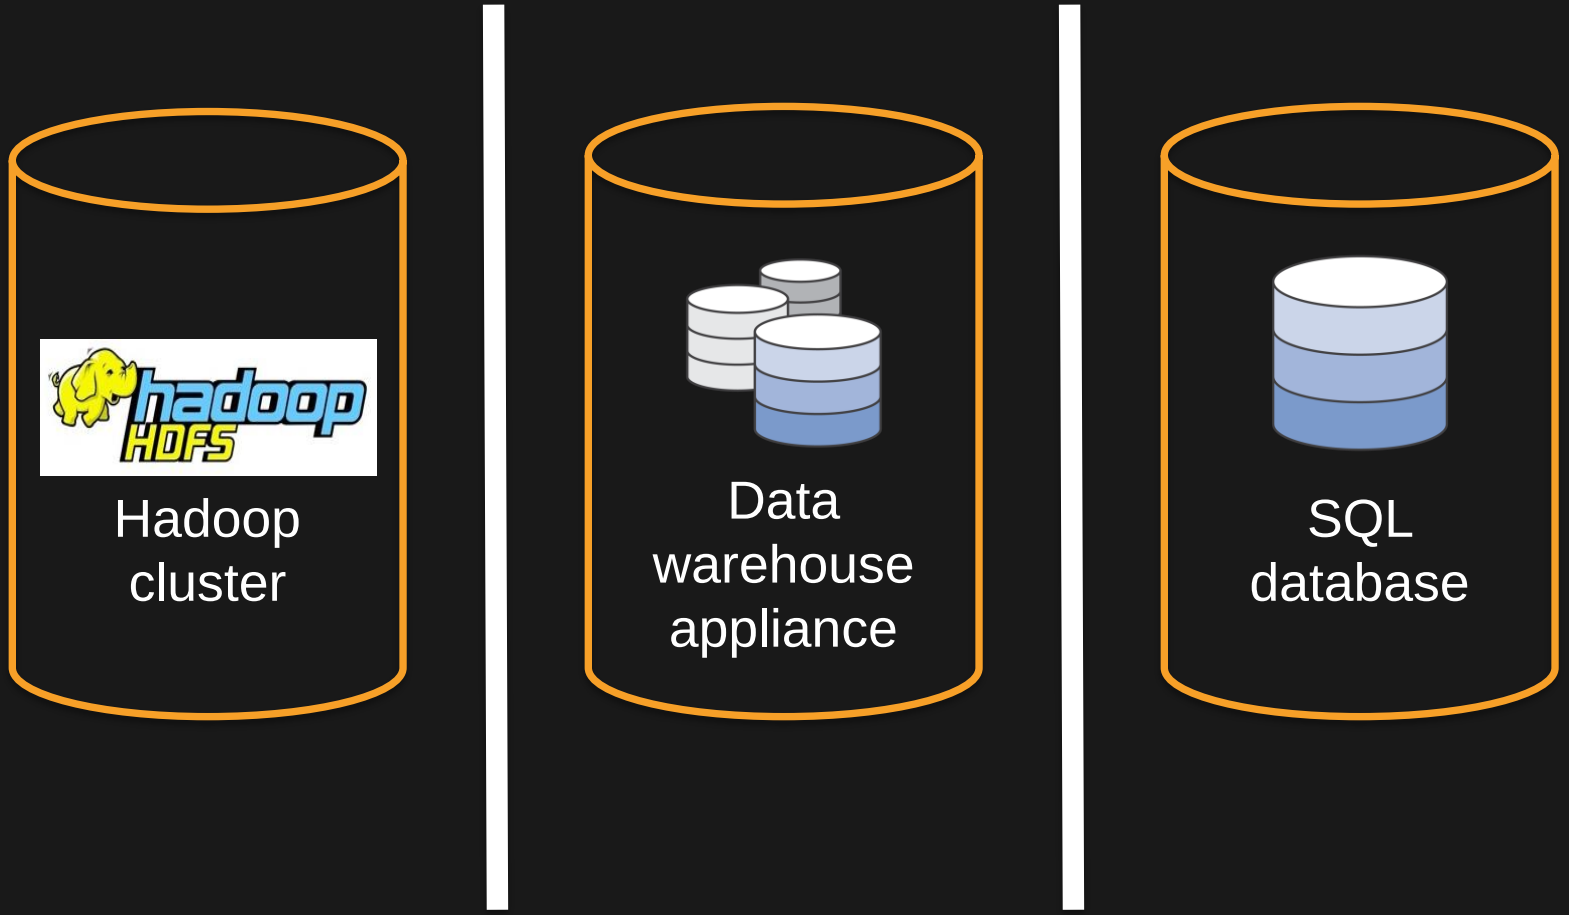

Navigating the Data Lake…

Data Lake is a new and increasingly popular architecture to store and analyze massive volumes and heterogenous types of data in a centralized repository

Building a Data Lake on AWS

Why AWS?

Implementing a Data Lake architecture requires a broad set of tools and technologies to serve an increasingly diverse set of applications and use cases.

Boils down to:

Picking the right tool for the right job…on a consumption-based model…

Data Lake reference architecture

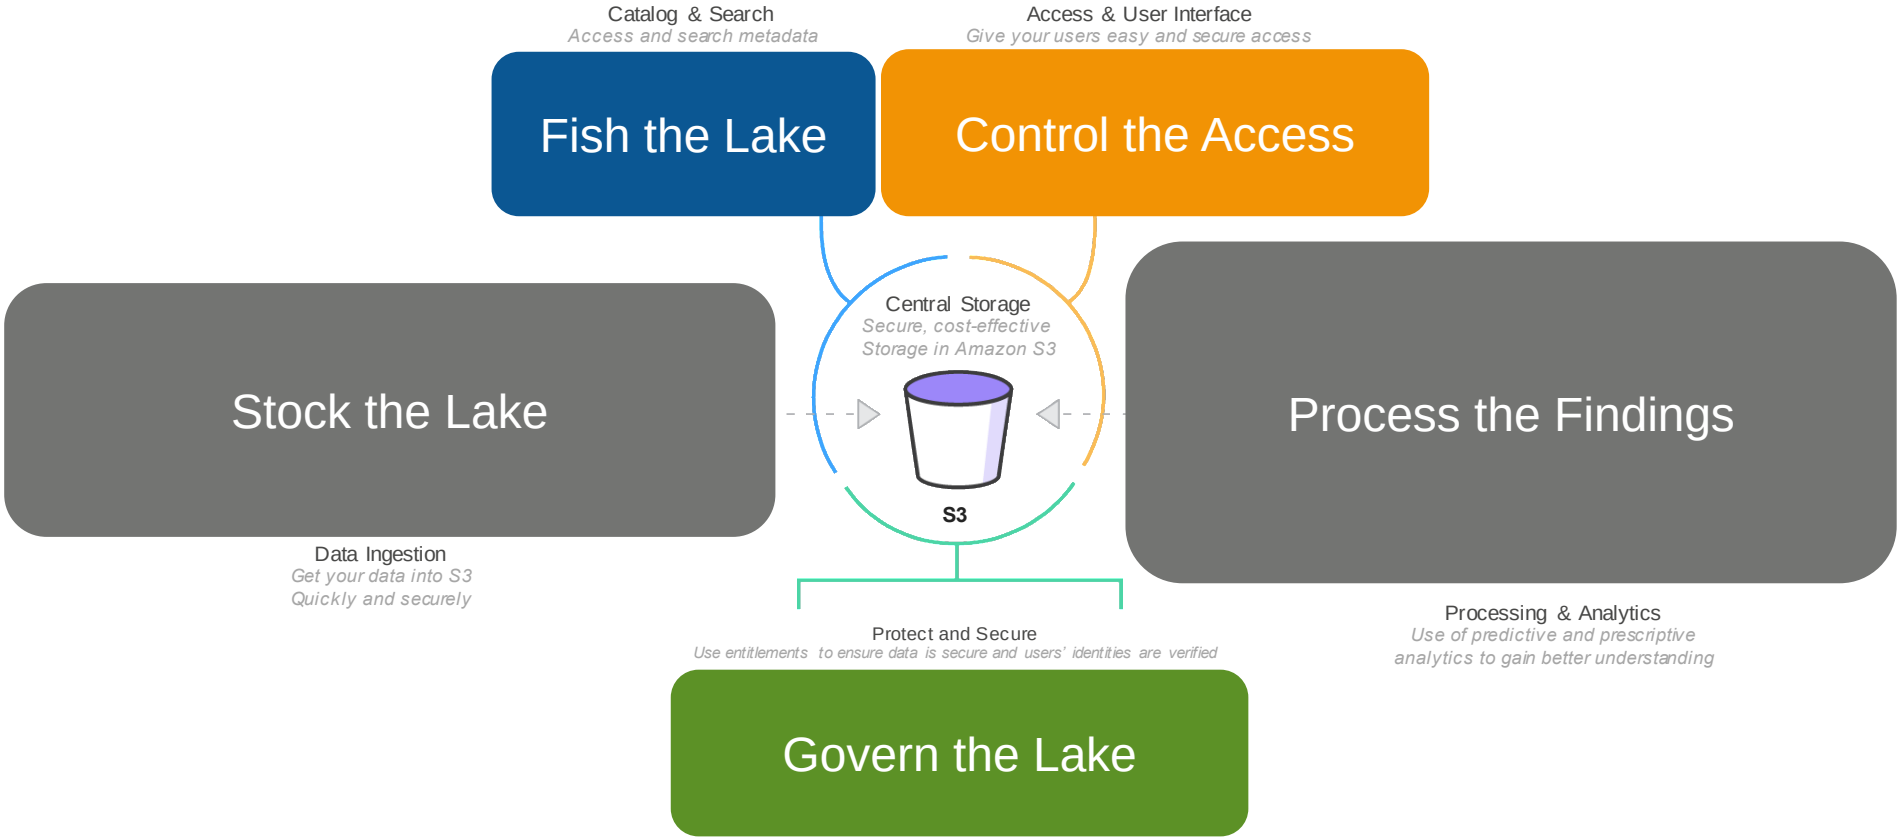

Data Lake reference architecture

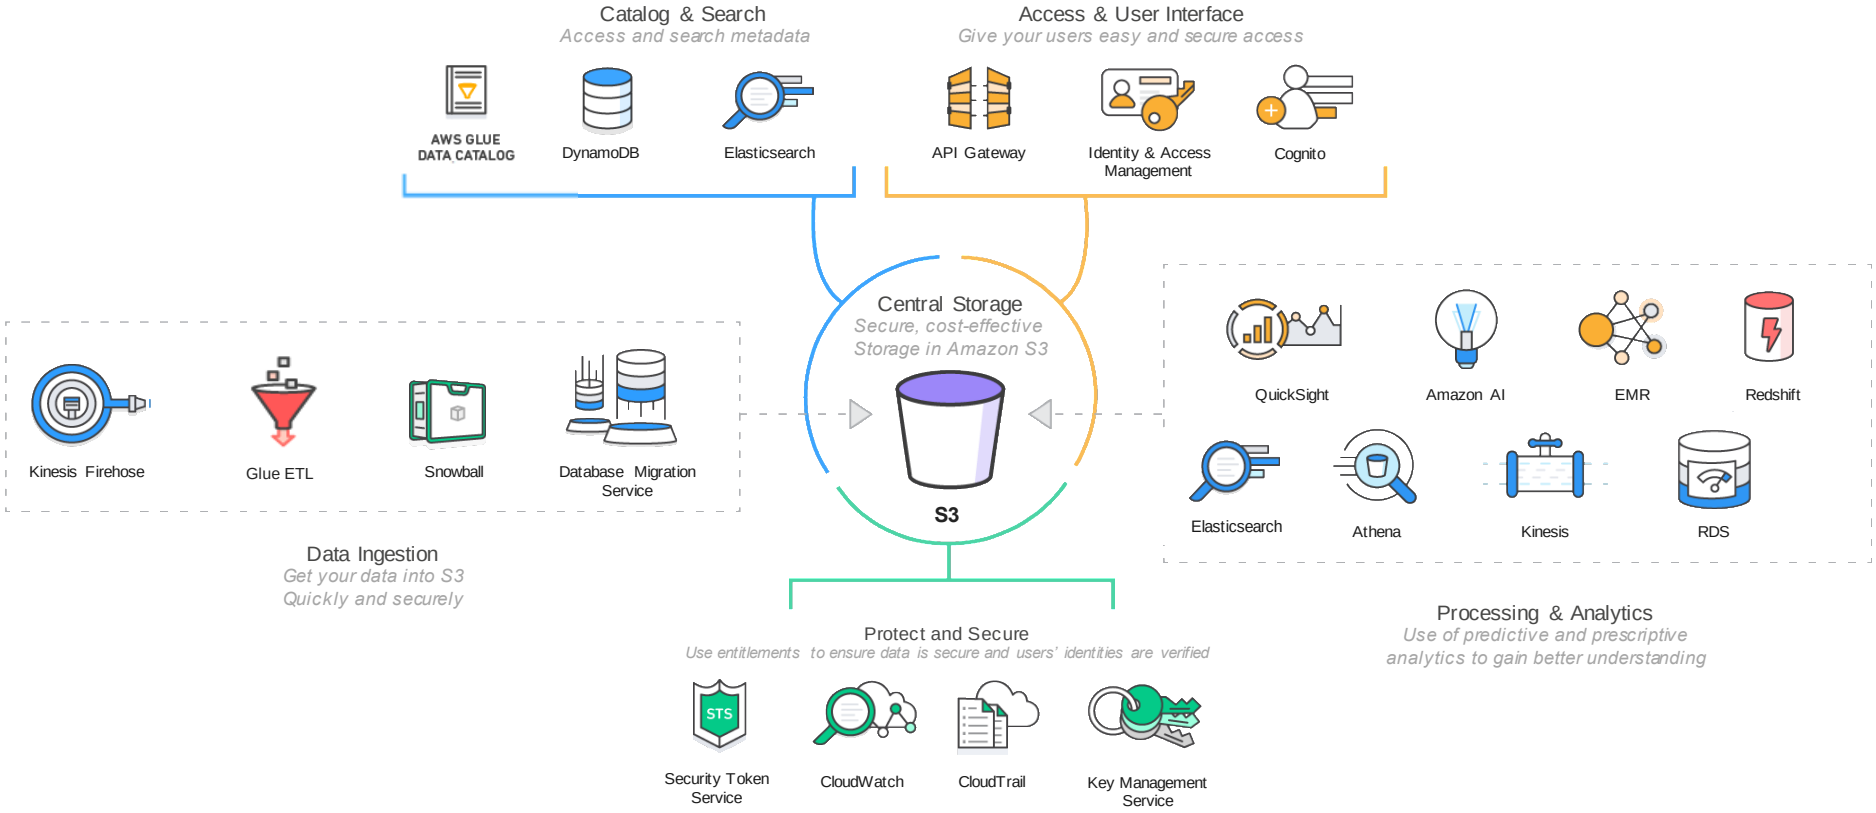

S3 – Center of the Data Lake

Why Amazon S3 for Data Lake?

Durable

Designed for 11 9s of durability

Available

Designed for **99.99**% availability

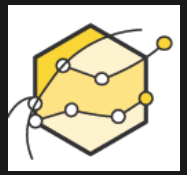

High performance

- Multiple Upload
- Range GET

Easy to use

- Simple REST API
- AWS SDKs
- Read-after-create consistency
- Event notification
- Lifecycle policies

Scalable

- Store as much as you need
- Scale storage and compute independently
- No minimum usage commitments

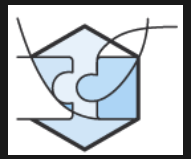

Integrated

- Amazon EMR
- Amazon Redshift
- Amazon DynamoDB

Stock the Lake – Data Ingestion

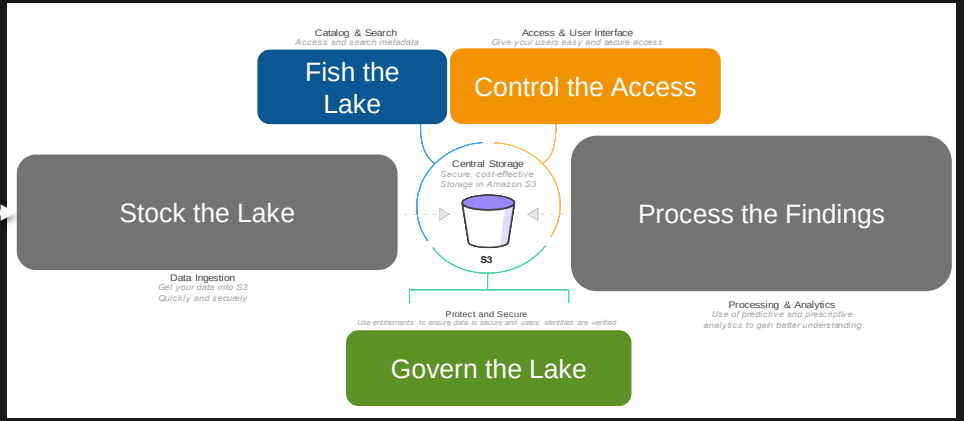

Amazon Glue ETL

- Serverless ETL engine generates Python code that is entirely customizable, reusable, and portable.

Database Migration Service

Migrate your RDBMS into S3 (as well as other targets)

Amazon Kinesis Streams

- Build your own custom applications that process or analyze streaming data

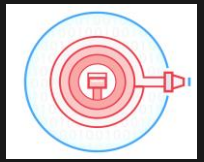

Amazon Kinesis Firehose

- Easily load massive volumes of streaming data into S3, Amazon Redshift, and Amazon Elasticsearch Service

Stock the Lake – Data Ingestion

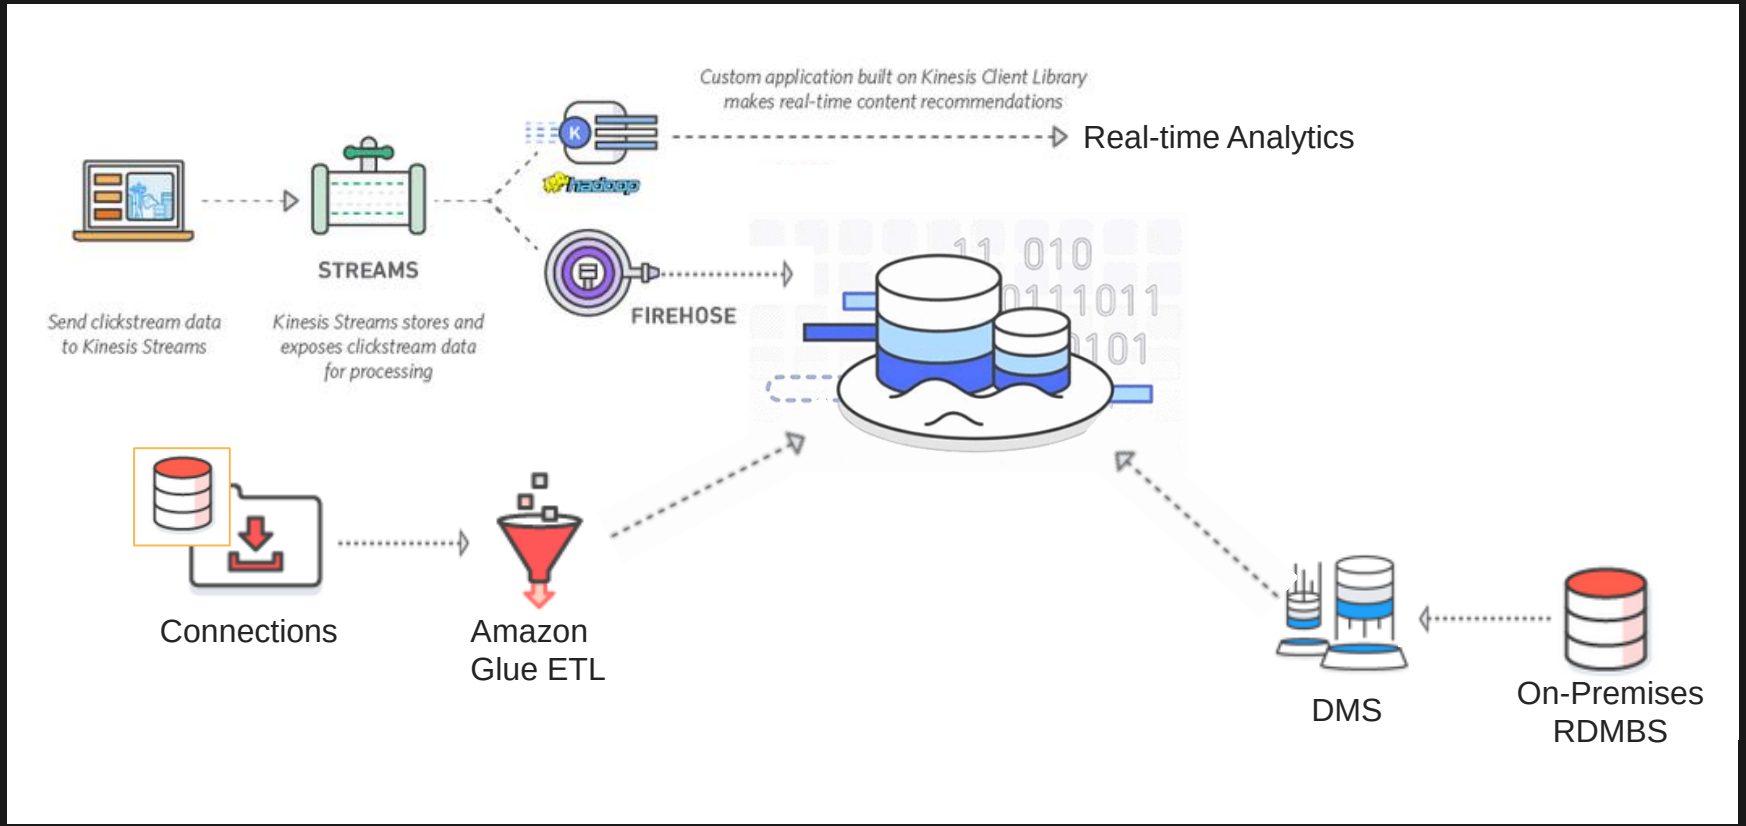

Demonstration

What is the speed of the data?

What is the source of the data?

Fishing the Lake – Catalog/Search

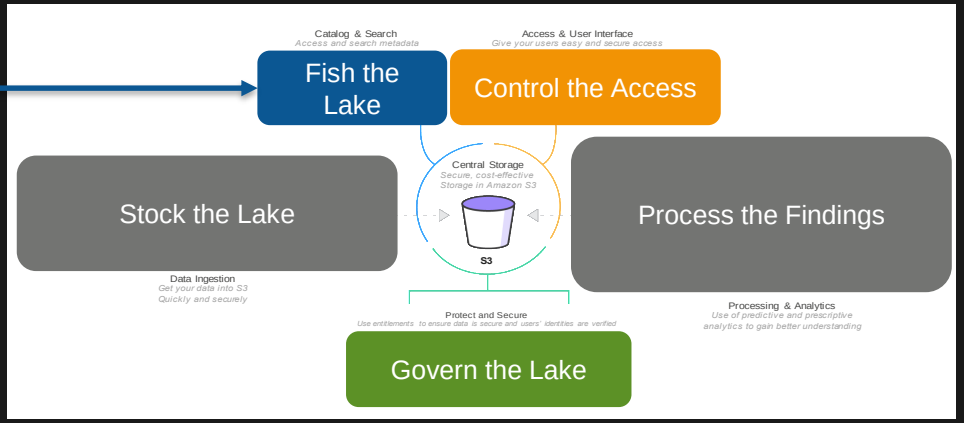

Glue Data Catalog

Structural and Operational metadata

Amazon DynamoDB

Metadata/Tag Lookup

Search Content

Demonstration

What type of data?

How is the data being queried?

Govern the Lake

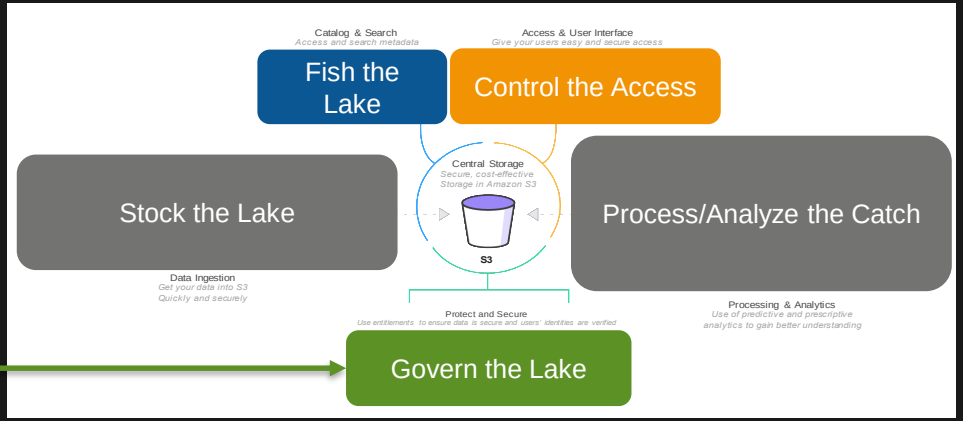

Temporary Tokens

Performance

Auditing

Encryption

Control the Access

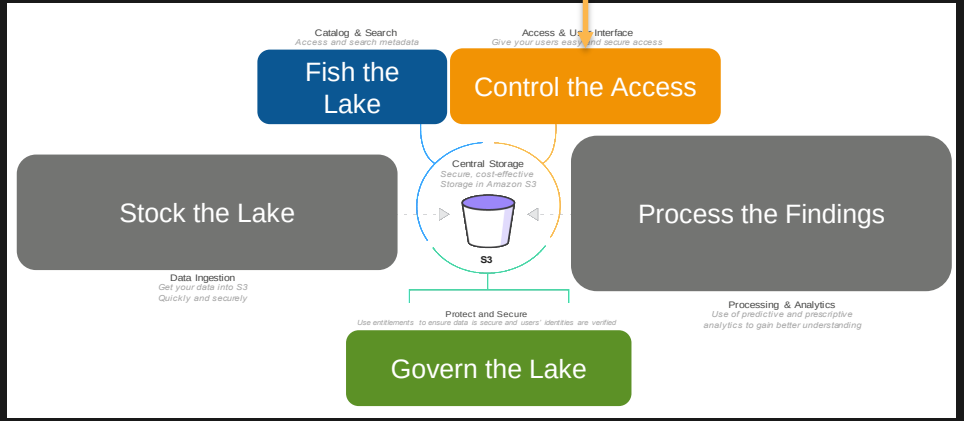

Interfaces

Identity/Access

User Authentication

Process/Analyze the Catch

Processing & Analytics

Real-time

- ElastiSearch Service
- Kinesis Analytics, Kinesis Streams
- Spark Streaming on EMR
- Apache Flink on EMR
- AWS Lambda
- Apache Storm on EMR

Batch

- EMR Hadoop, Spark, Presto
- Redshift Data Warehouse
- Athena Query Service

AI & Predictive

- Amazon Lex
- Amazon Polly
- Amazon Rekognition
- Machine Learning

Transactional & RDBMS

- DynamoDB, NoSQL DB
- Aurora Relational Database

BI & Data Visualization

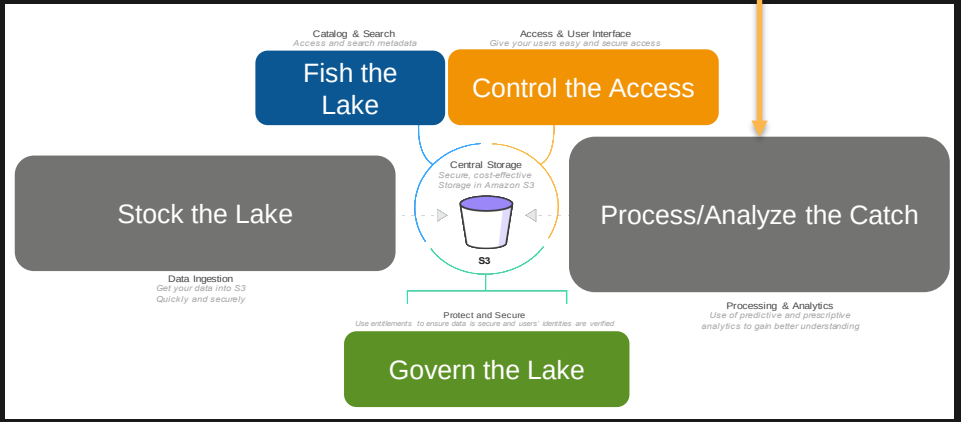

Process/Analyze the Catch

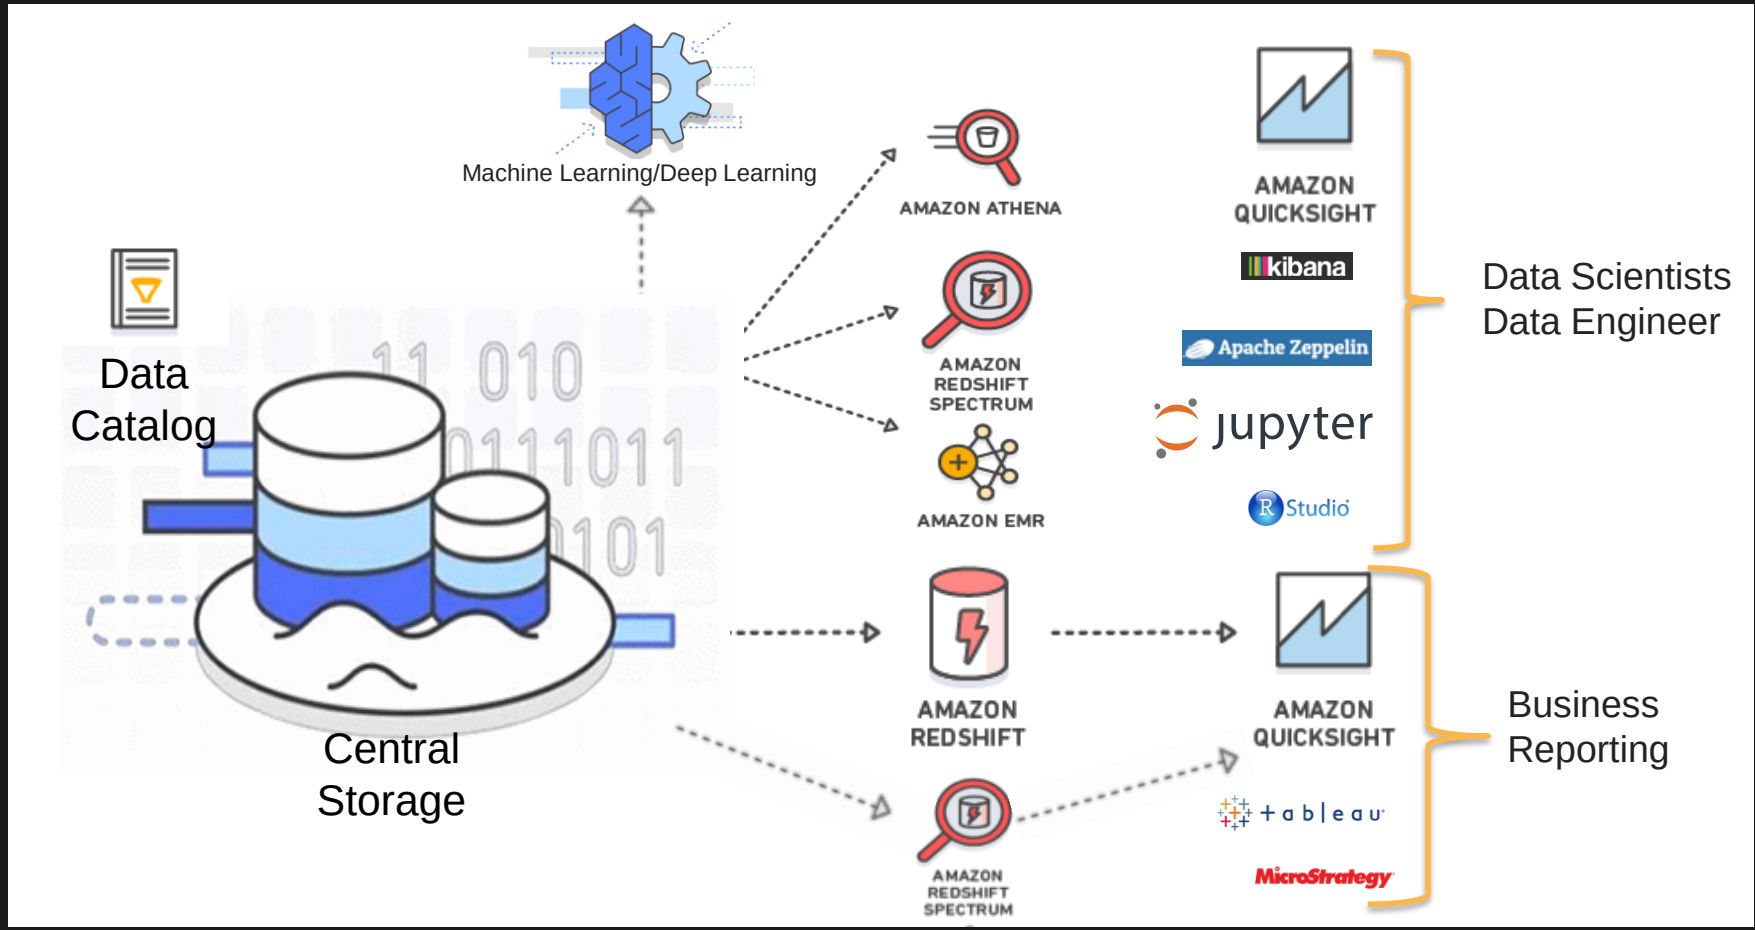

Interactive query service

Amazon Athena

- **Query directly from Amazon S3**
- **Use ANSI SQL**
- **Serverless**
- **Multiple data formats**
- **Pay per query**

Amazon
Elastic
MapReduce

Hadoop/HDFS clusters

Hive, Pig, Impala, Hbase, Spark, Presto

Easy to use, fully managed

On-demand, reserved instance, and Spot pricing

Tight integration with Amazon S3, DynamoDB, and Kinesis

Resizable clusters

Easy to add and remove compute capacity on your cluster.

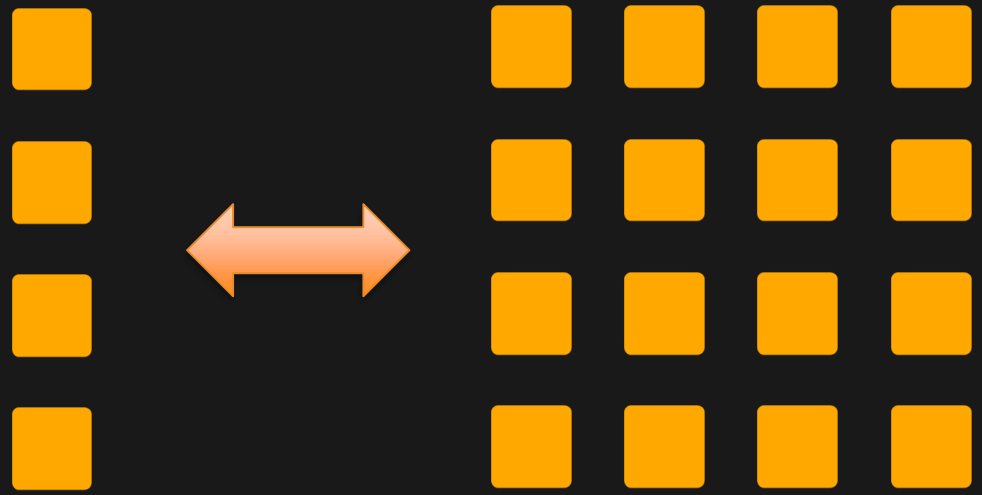

ON A SINGLE MACHINE

COST: 4h x $1.06 = **$4.24**
PROCESSING TIME: **4h**

ON MULTIPLE MACHINES

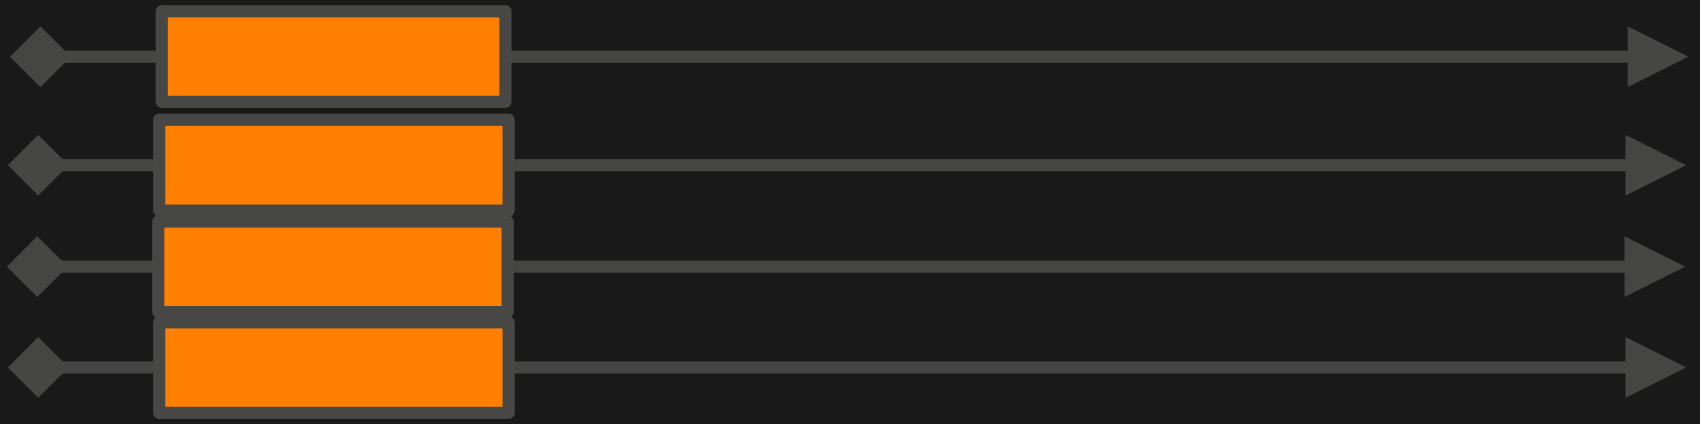

COST: 4 x 1h x $1.06 = **$4.24**
PROCESSING TIME: **1h**

Amazon Redshift works with third-party analysis tools

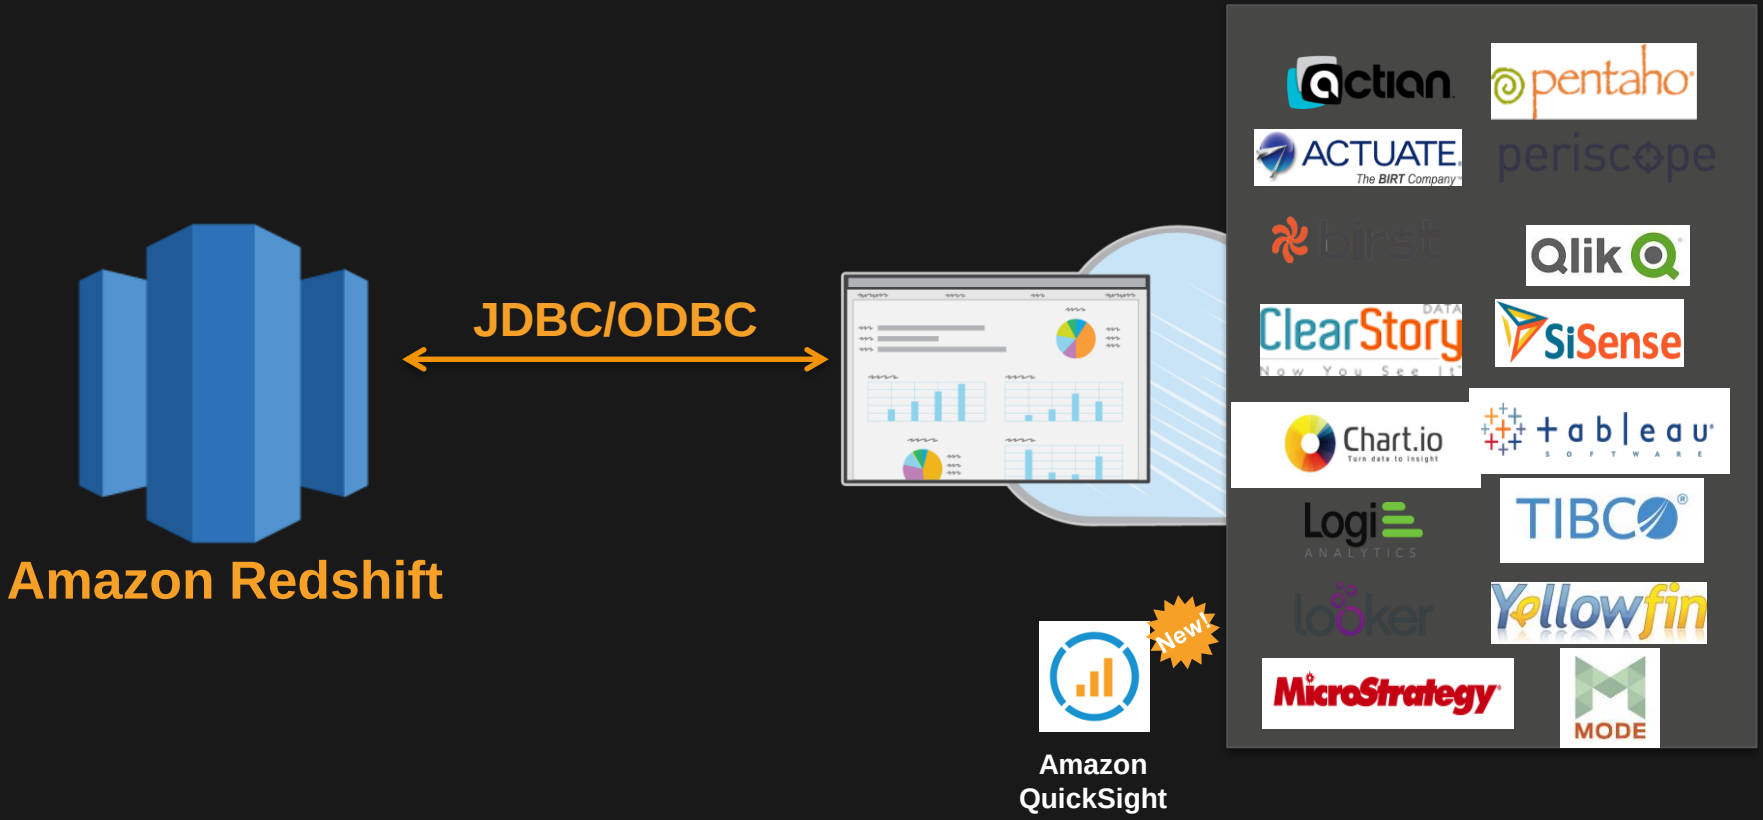

Amazon Redshift has security built in

SSL to secure data in transit

Encryption to secure data at rest

- AES-256; hardware accelerated
- All blocks on disks and in Amazon S3 encrypted
- HSM support

No direct access to compute nodes

Audit logging, AWS CloudTrail, AWS KMS integration

Amazon VPC support

SOC 1/2/3, PCI-DSS Level 1, FedRAMP, HIPAA

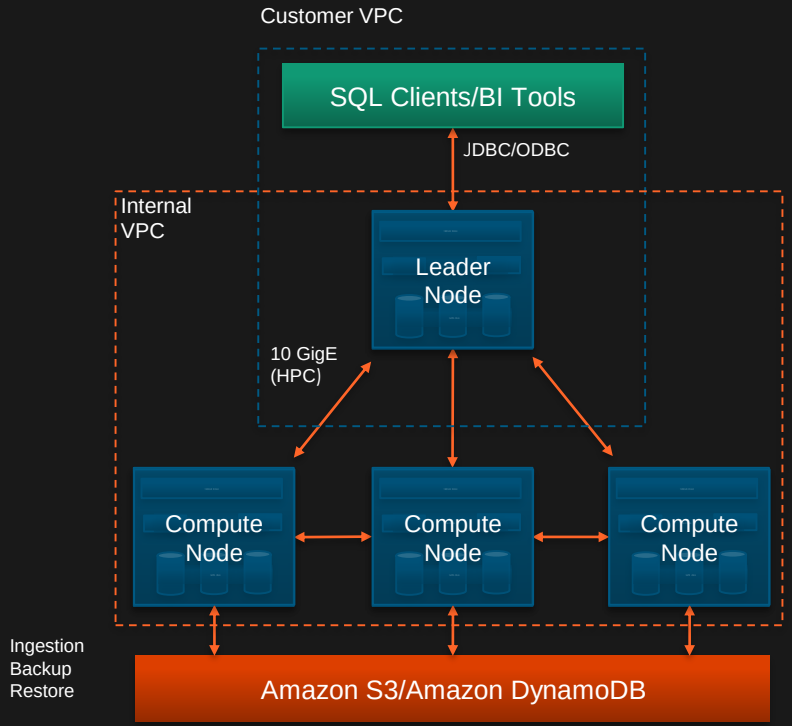

Redshift Spectrum

Leverages Amazon Redshift's advanced cost-based optimizer

Pushes down projections, filters, aggregations and join reduction

Dynamic partition pruning to minimize data processed

Automatic parallelization of query execution against Amazon S3 data

Efficient join processing within the Amazon Redshift cluster

Data Catalog

Apache Hive Metastore

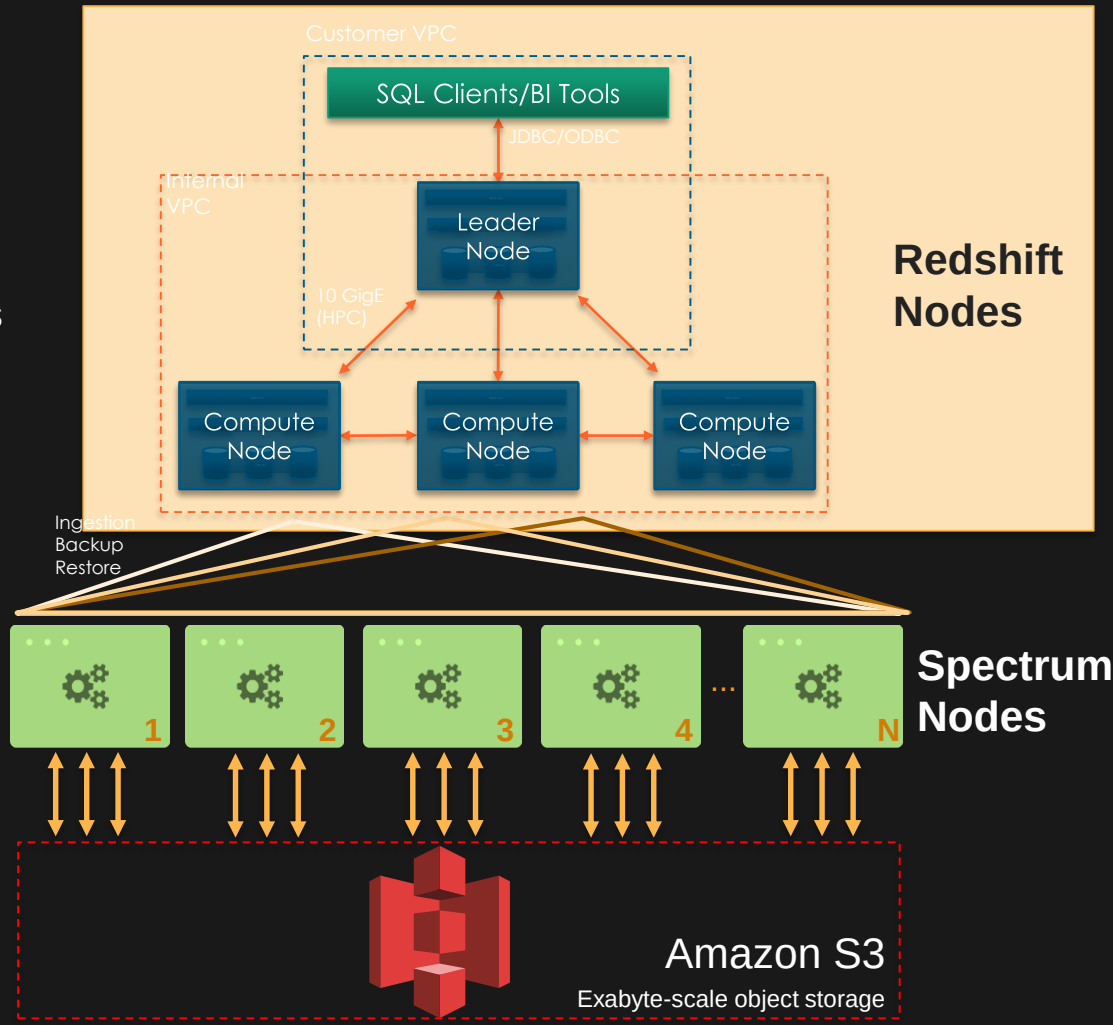

Amazon AI

Intelligent services powered by deep learning

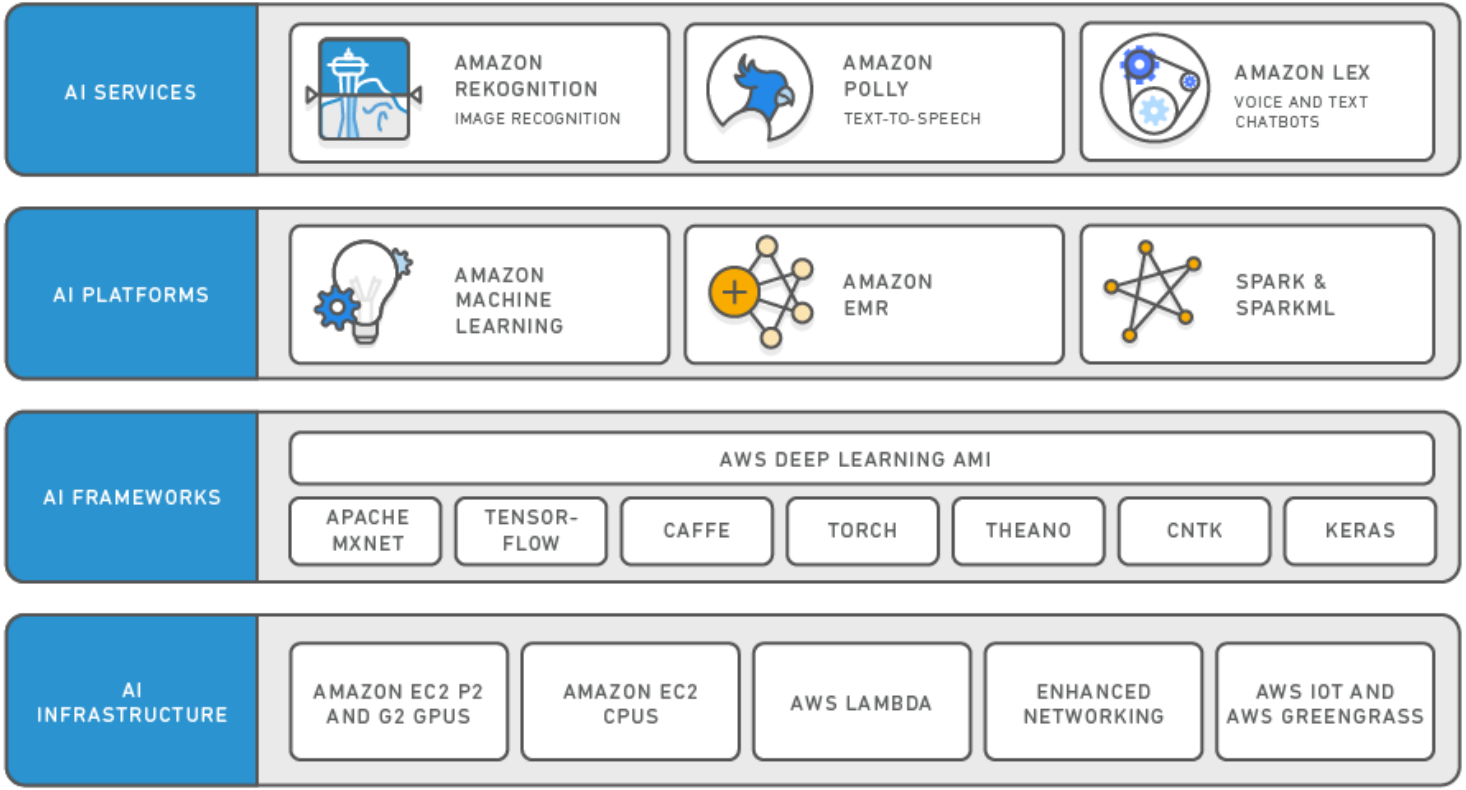

Proven customer success

The vast majority of big data use cases deployed in the cloud today run on AWS.

Case study: Re-architecting compliance

“For our market surveillance systems, we are looking at about 40% [savings with AWS], but the real benefits are the business benefits: We can do things that we physically weren’t able to do before, and that is priceless.”

- Steve Randich, CIO

FINRA

<u>What FINRA needed</u>

- Infrastructure for its market surveillance platform
- Support of analysis and storage of approximately 75 billion market events every day

<u>Why they chose AWS</u>

- Fulfillment of FINRA’s security requirements
- Ability to create a flexible platform using dynamic clusters (Hadoop, Hive, and HBase), Amazon EMR, and Amazon S3

<u>Benefits realized</u>

- Increased agility, speed, and cost savings
- Estimated savings of $10-20M annually by using AWS

Fraud detection

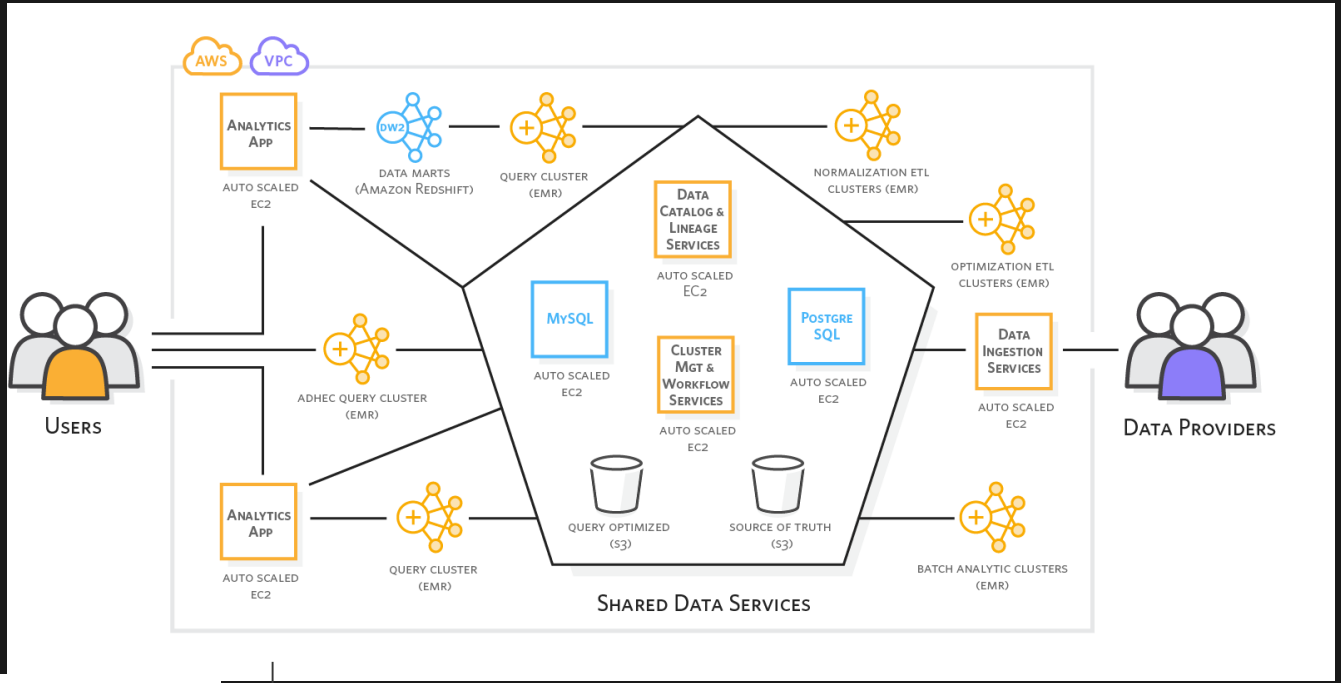

FINRA uses Amazon EMR and Amazon S3 to process up to 75 billion trading events per day and securely store over 5 petabytes of data, attaining savings of $10-20M per year.

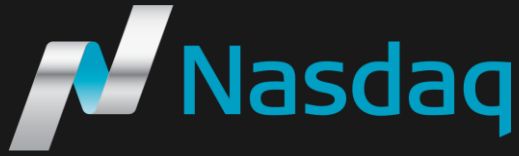

- Nasdaq implements an Amazon S3 data lake + Amazon Redshift data warehouse architecture
- Most recent two years of data is kept in the Amazon Redshift data warehouse and snapshotted into Amazon S3 for disaster recovery
- Data between two and five years old is kept in Amazon S3
- Presto on Amazon EMR is used to ad-hoc query data in Amazon S3
- Transitioned from an on-premises data warehouse to Amazon Redshift & Amazon S3 data lake architecture
- Over 1,000 tables migrated
- Average daily ingest of over 7B rows
- ***Migrated*** off legacy DW to AWS (start to finish) ***in 7 man-months***
- AWS costs were ***43%*** of legacy budget for the same data set (~1100 tables)

Building a Data Lake

An **architectural approach** that allows you to store massive amounts of "**raw**" data into a central location

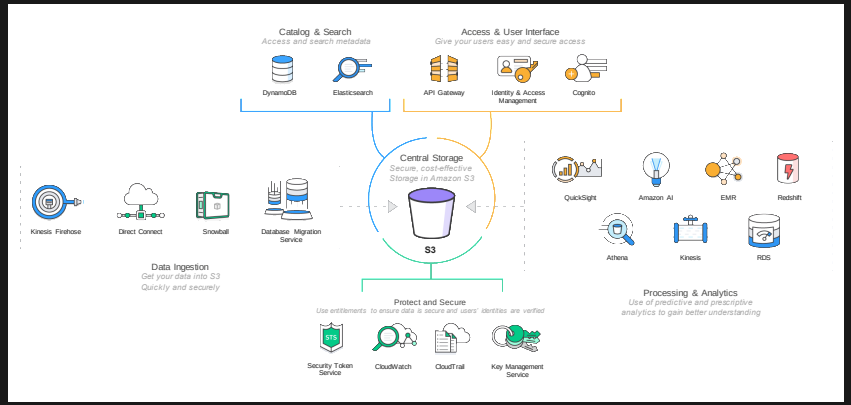

It's **readily available** to be **categorized, processed, analyzed, and consumed** by **diverse groups**


AWS
SUMMIT
Thank you!
amazon
web services