

Big Data Architectural Patterns and Best Practices on AWS

Keith Steward, Ph.D.
Specialist (EMR) Solution Architect
Amazon Web Services

October 2016

What we'll cover:

1. Big data challenges
2. How to simplify big data processing
3. What technologies should you use?
 - Why?
 - How?
4. Reference architecture
5. Design patterns

Ever Increasing Big Data



Volume



Velocity



Variety

Big Data Evolution

Batch → Real-time → Prediction

Reports



Alerts



Forecasts



Cloud Services Evolution

Virtual
machines



Managed
services



Serverless



Plethora of Tools



Big Data Challenges

Is there a reference architecture?

What tools should I use?

How?

Why?

Architectural Principles

Decoupled “data bus”

- Data → **Store** → **Process** → **Store** → **Analyze** → Answers

Use the right tool for the job

- Data structure, latency, throughput, access patterns

Leverage AWS managed services

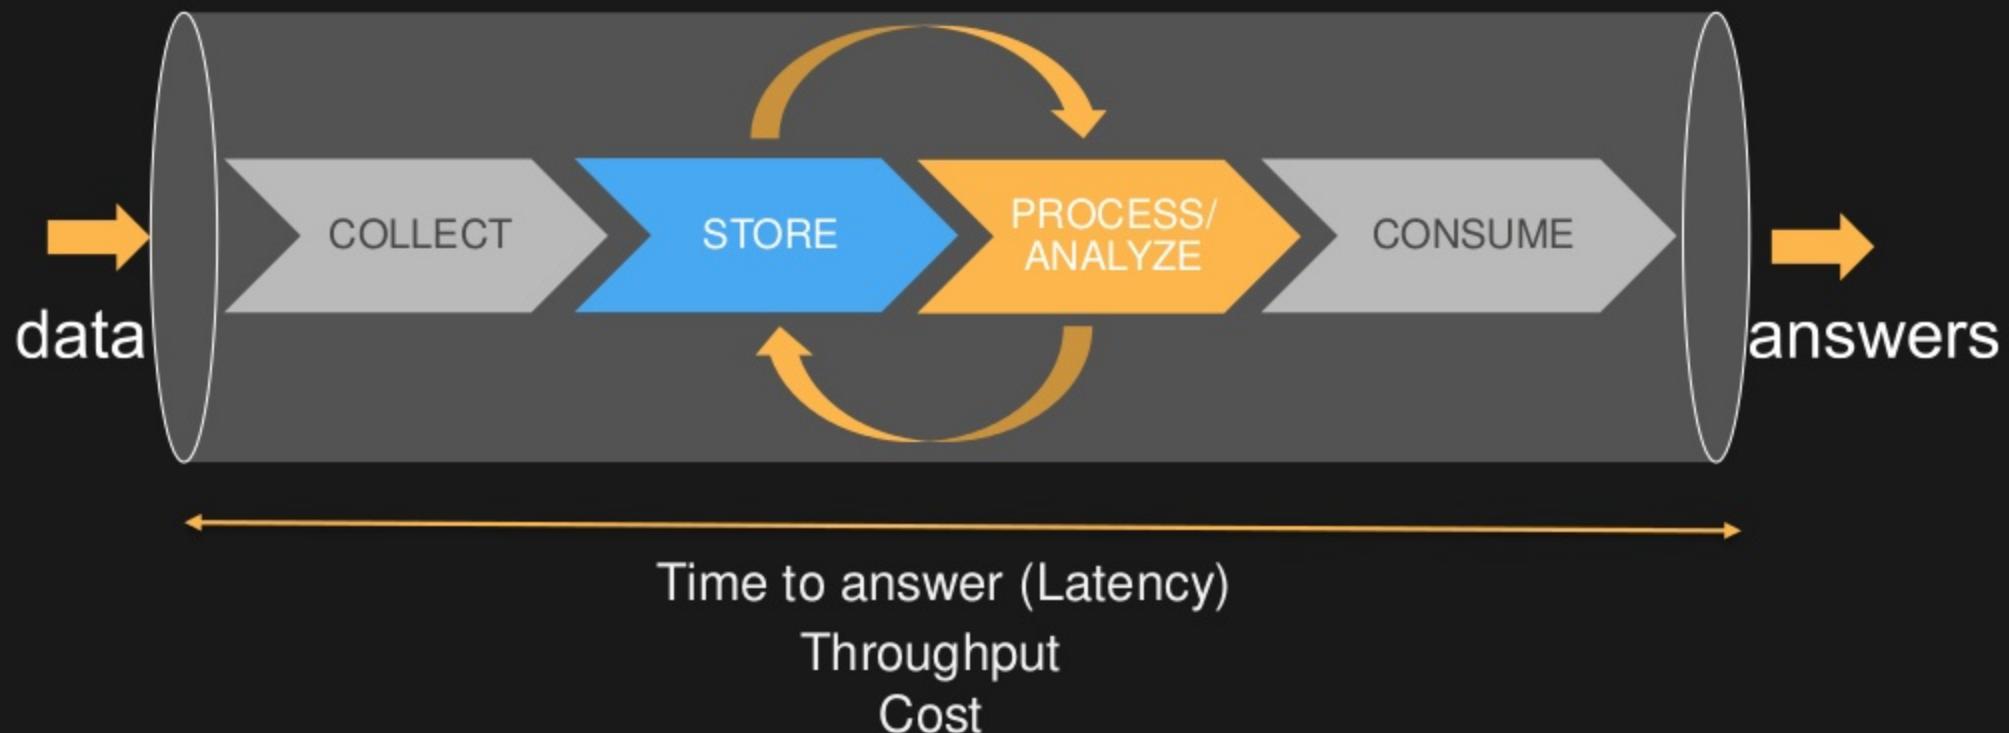
- Scalable / elastic, available, reliable, secure, no / low admin

Use Lambda architecture ideas

- Immutable (append-only) log, batch / real-time / serving layer

Big data ≠ big cost

Conceptual Framework to Simplify Big Data Processing

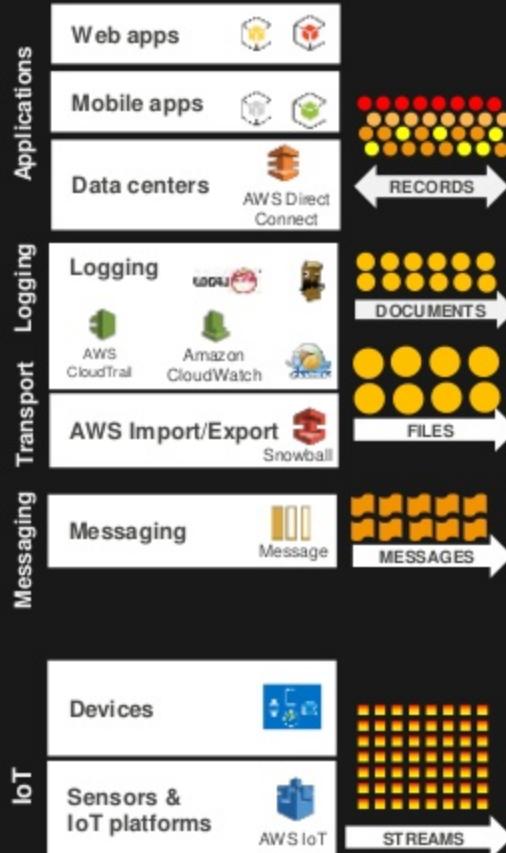




COLLECT

COLLECT

Types of Data



Data Structures

Database Records

Search Documents

Log Files

Messages

Data Streams

What Is the Temperature of Your Data ?

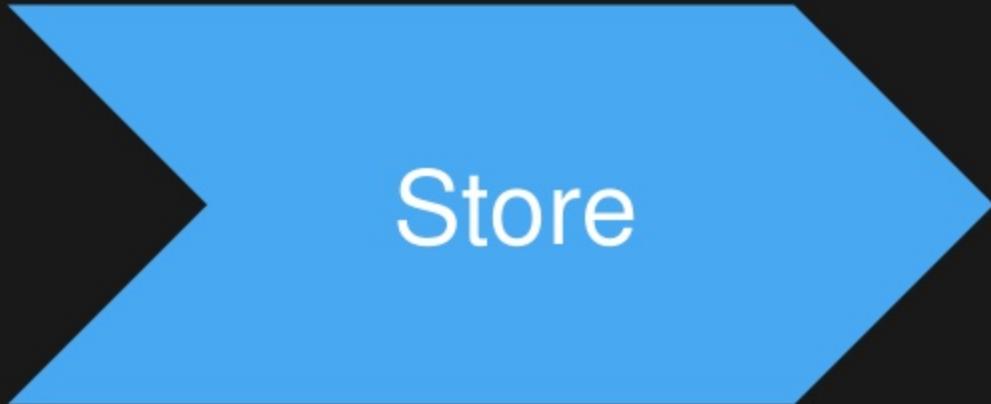


Data Characteristics: Hot, Warm, Cold

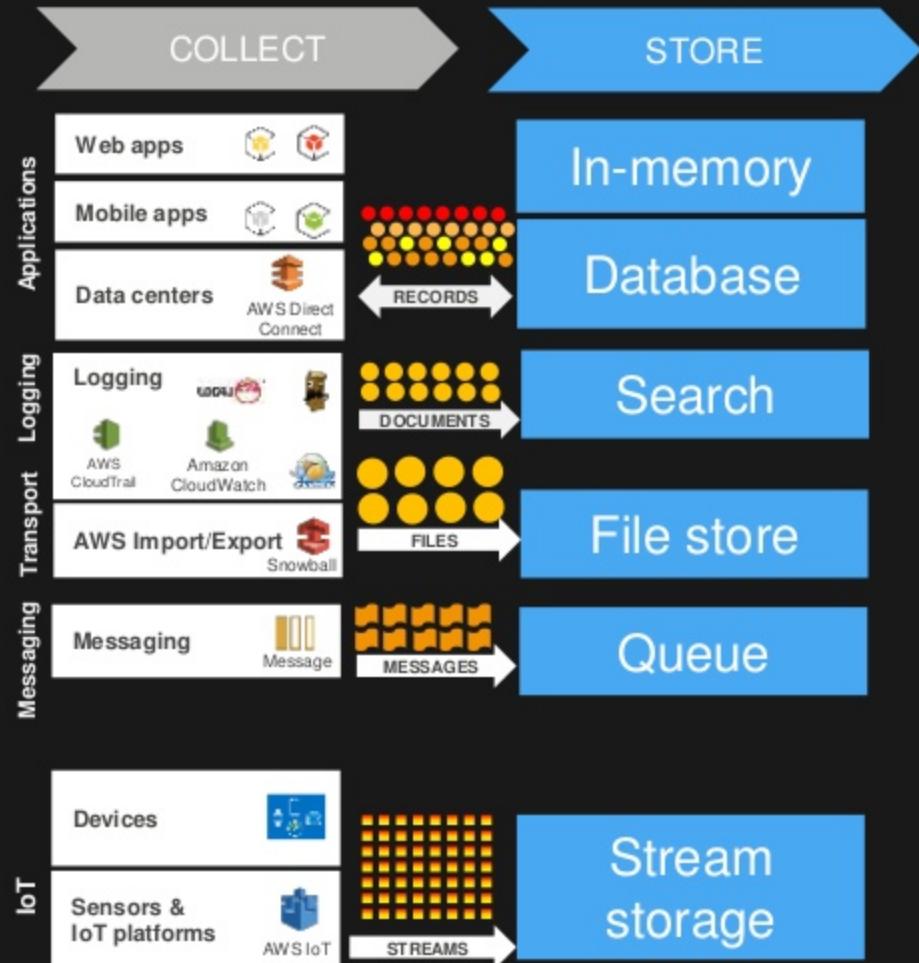
	Hot	Warm	Cold
Volume	MB–GB	GB–TB	PB–EB
Item size	B–KB	KB–MB	KB–TB
Latency	ms	ms, sec	min, hrs
Durability	Low–high	High	Very high
Request rate	Very high	High	Low
Cost/GB	\$\$-\$	\$-¢¢	¢



Hot data Warm data Cold data



Store



Types of Data Stores

Caches, Data Structure Servers

SQL & NoSQL Databases

Search Engines

File Systems

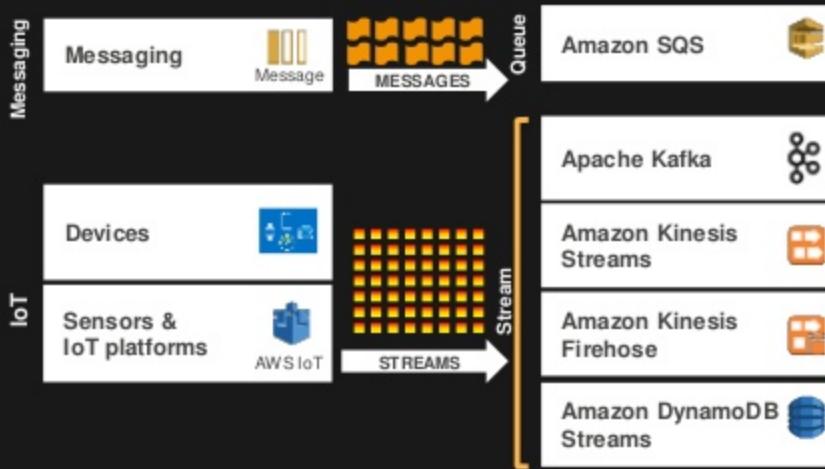
Message Queues

Pub/Sub Message Queues

COLLECT

STORE

Message & Stream Storage



Amazon SQS

- Managed message queue service

Apache Kafka

- High throughput distributed messaging system

Amazon Kinesis Streams

- Managed stream storage + processing

Amazon Kinesis Firehose

- Managed data delivery

Amazon DynamoDB

- Managed NoSQL database
- Tables can be stream-enabled

Why Stream Storage?

Decouple producers & consumers

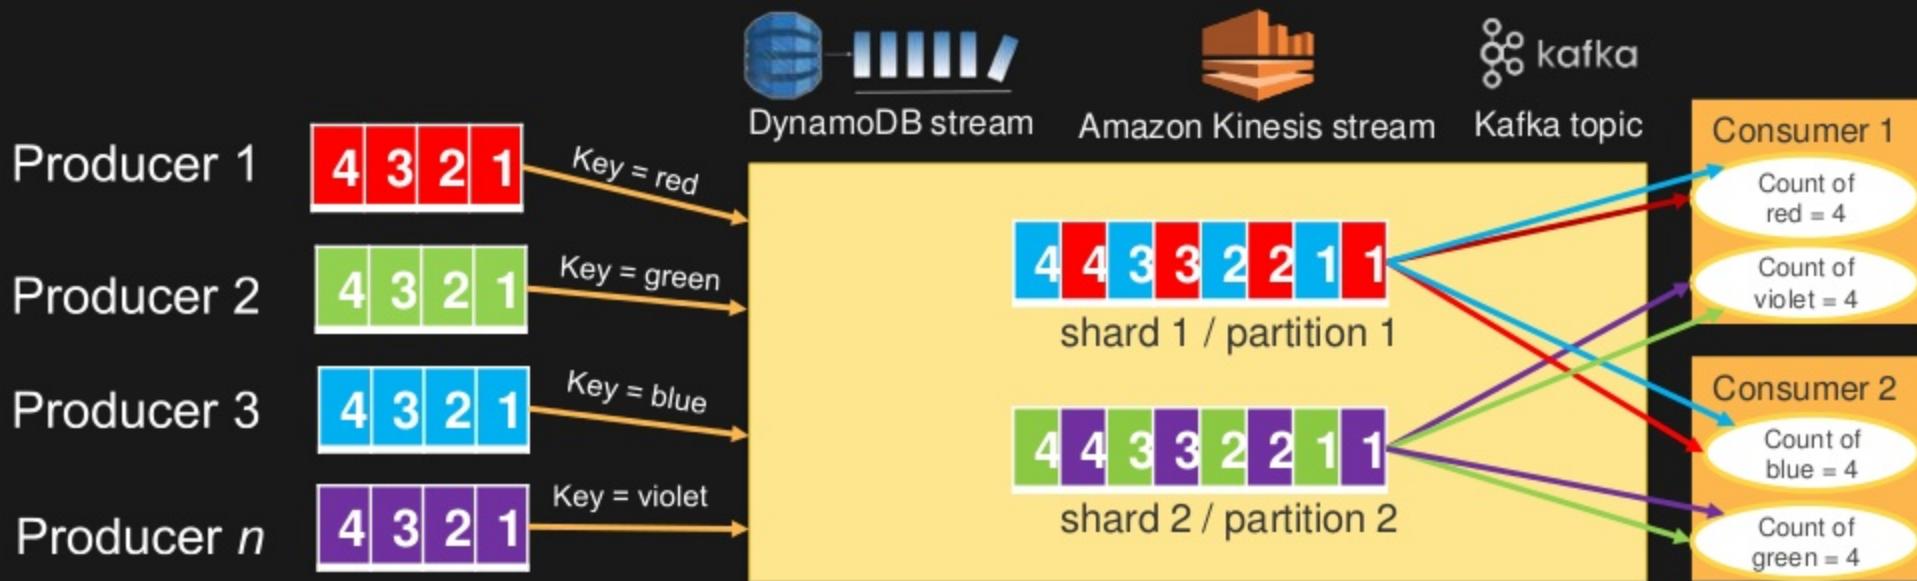
Persistent buffer

Collect multiple streams

Preserve client ordering

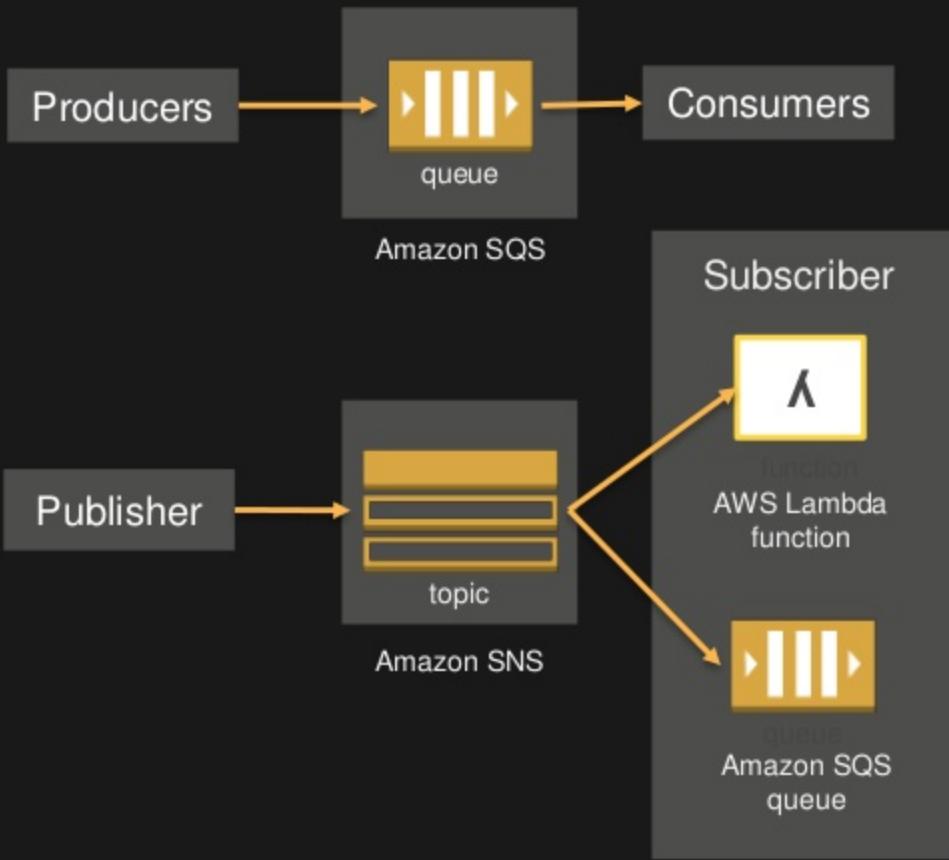
Parallel consumption

Streaming MapReduce



What About Messaging?

- Decouple producers & consumers
- Persistent buffer
- Collect multiple streams
- **No** client ordering
- **No** streaming MapReduce
- **No** parallel consumption for Amazon SQS consumers
 - Amazon SNS can publish to multiple SNS subscribers (queues or λ functions)



What Stream / Message Storage Should I Use?

	Amazon DynamoDB Streams	Amazon Kinesis Streams	Amazon Kinesis Firehose	Apache Kafka	Amazon SQS
AWS managed	Yes	Yes	Yes	No	Yes
Guaranteed ordering	Yes	Yes	No	Yes	No
Delivery (deduping)	Exactly-once	At-least-once	At-least-once	At-least-once	At-least-once
Data retention period	24 hours	7 days	N/A	Configurable	14 days
Availability	3 AZ	3 AZ	3 AZ	Configurable	3 AZ
Scale / throughput	No limit / ~ table IOPS	No limit / ~ shards	No limit / automatic	No limit / ~ nodes	No limits / automatic
Parallel clients	Yes	Yes	No	Yes	No
Stream MapReduce	Yes	Yes	N/A	Yes	N/A
Row/object size	400 KB	1 MB	Destination row/object size	Configurable	256 KB
Cost	Higher (table cost)	Low	Low	Low (+admin)	Low-medium

Hot

Warm

COLLECT

STORE

File Storage

Applications

Web apps



Mobile apps

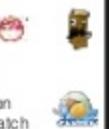


Data centers



Logging

Logging



AWS Import/Export



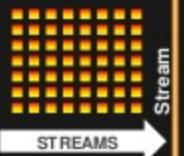
Transport

Messaging



Messaging

Devices



IoT

Sensors & IoT platforms



FILES

MESSAGES

STREAMS

Hot

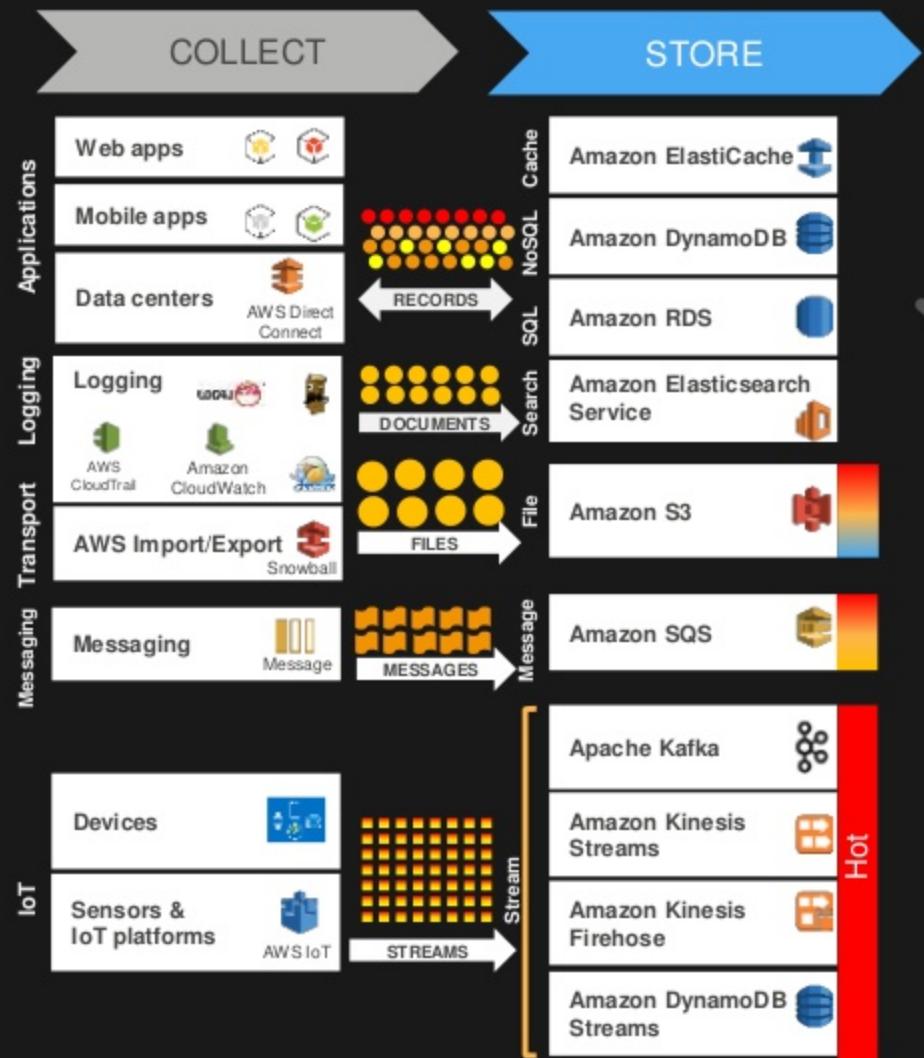
Why is Amazon S3 Good for Big Data?

- Native support by big data frameworks (Spark, Hive, Presto, etc.)
- No need to run compute clusters for storage (unlike HDFS)
- Supports transient Hadoop clusters & Amazon EC2 Spot Instances
- Multiple distinct (Spark, Hive, Presto) clusters can use the same data
- Elastic: Unlimited number of objects
- Very high bandwidth – no aggregate throughput limit
- Highly available – can tolerate AZ failure
- Designed for 99.99999999% durability
- Tired-storage (Standard, IA, Amazon Glacier) via life-cycle policy
- Secure – SSL, client/server-side encryption at rest
- Low cost

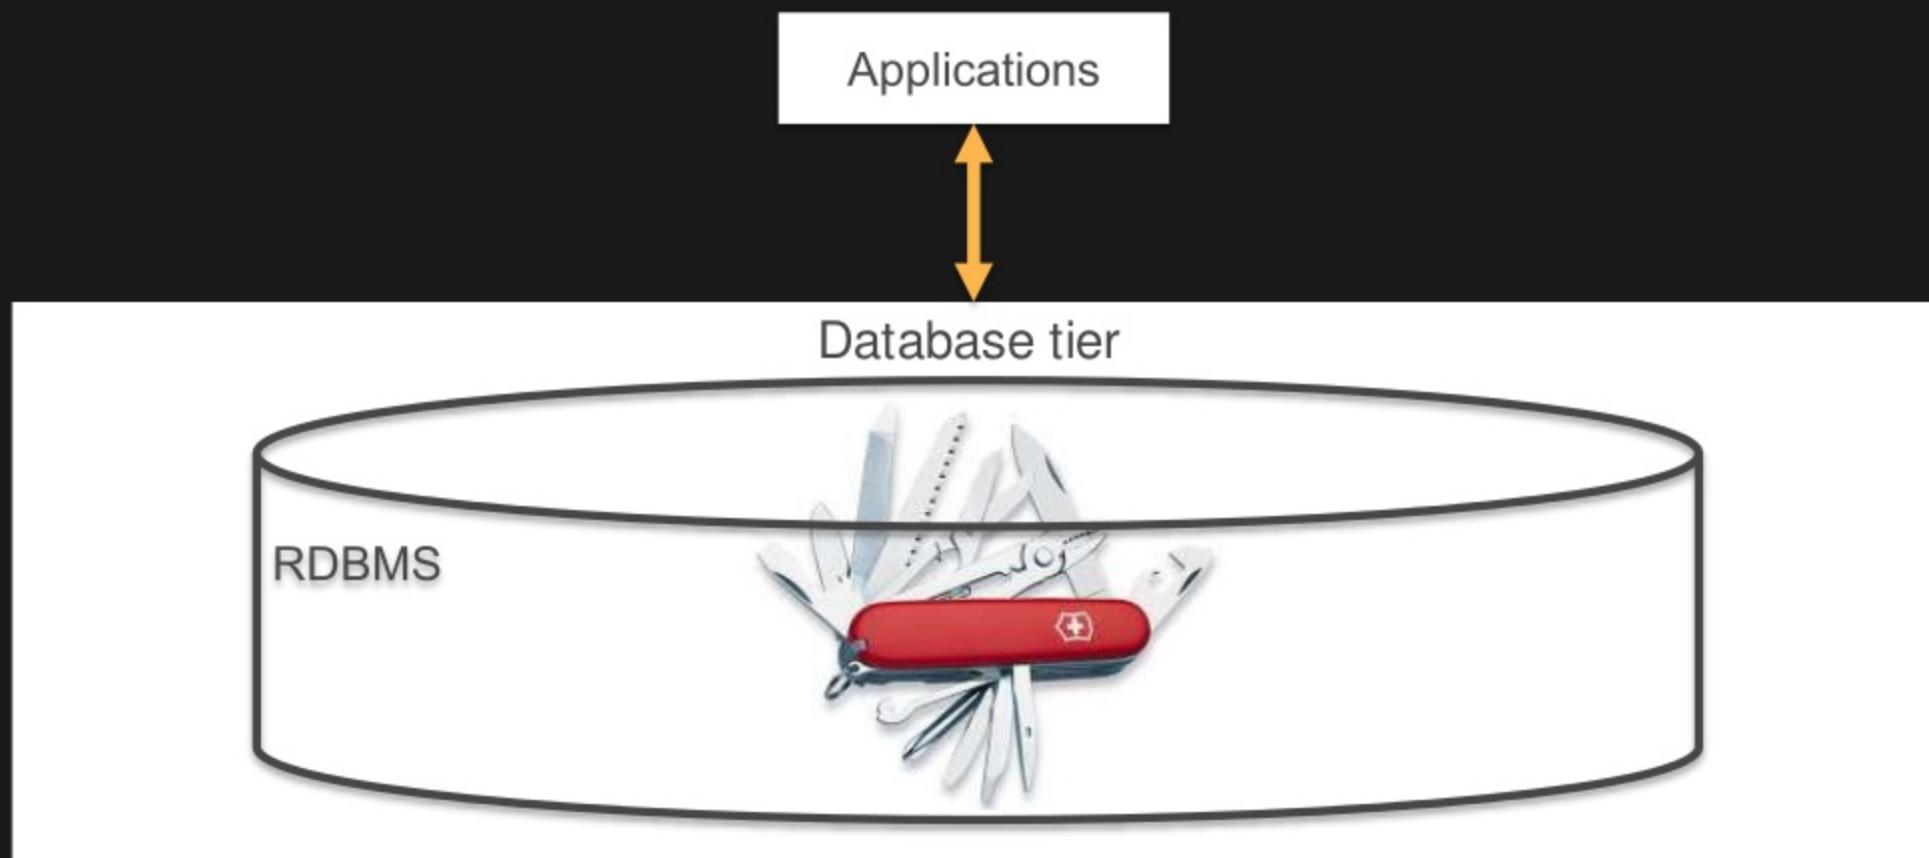
What about HDFS & Amazon Glacier?

- Use HDFS for very frequently accessed (hot) data
- Use Amazon S3 Standard for frequently accessed data
- Use Amazon S3 Standard – IA for infrequently accessed data
- Use Amazon Glacier for archiving cold data

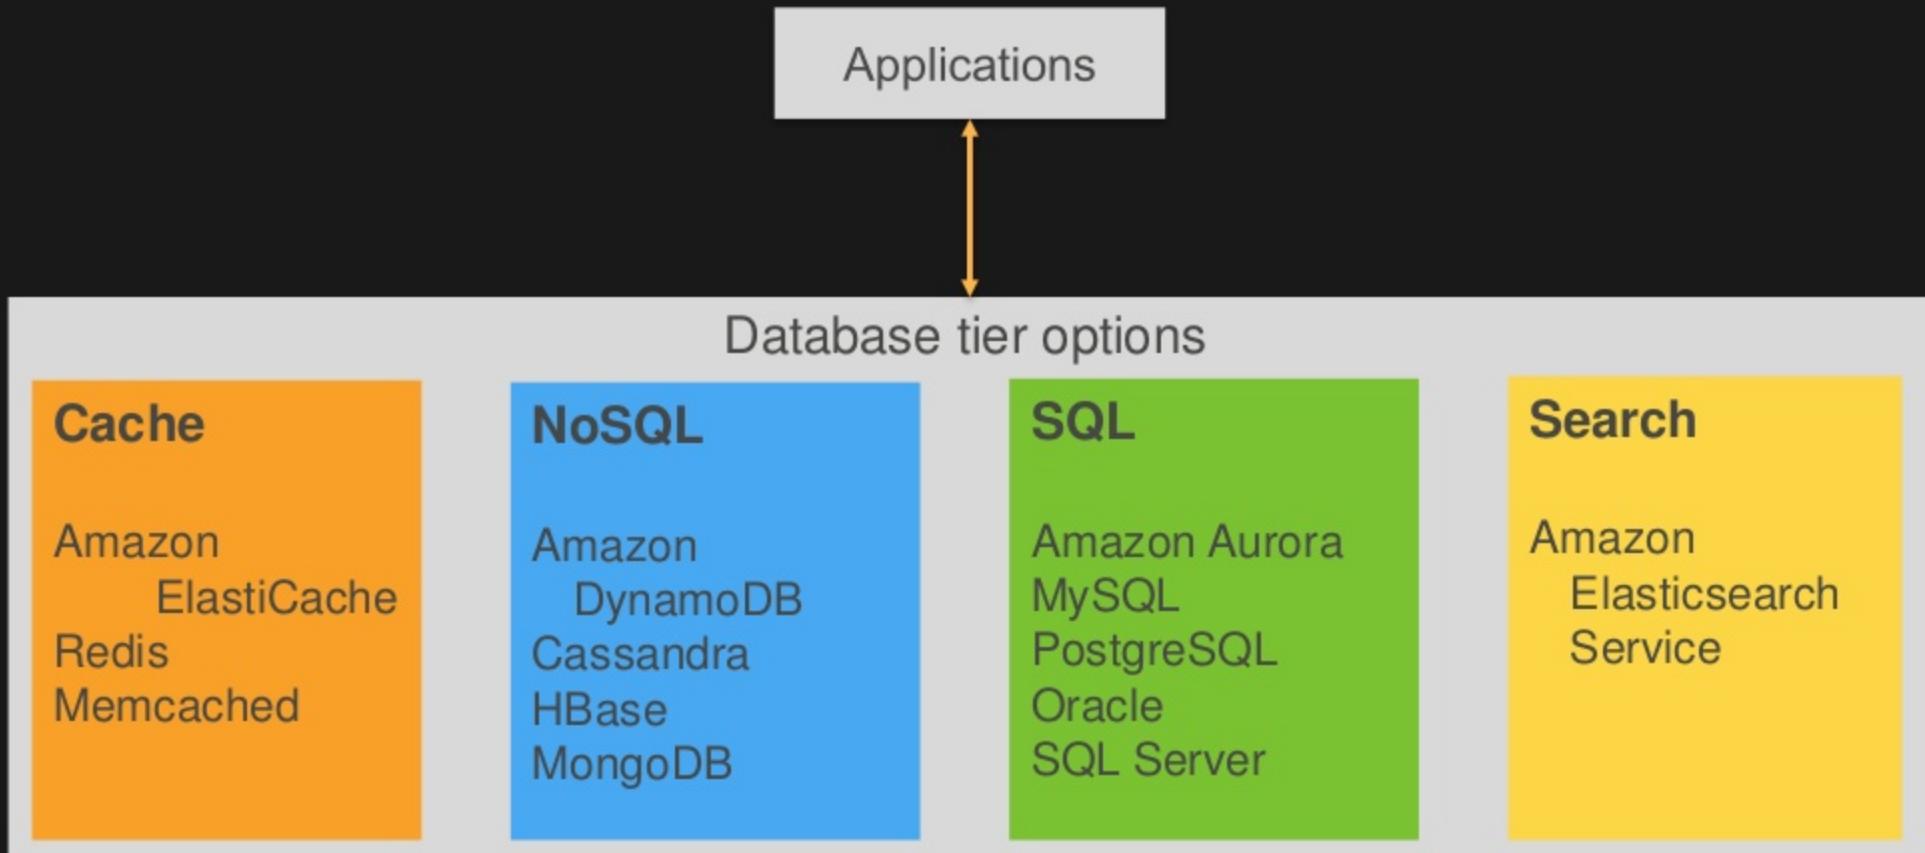




Database Anti-pattern



Best Practice - Use the Right Tool for the Job



What Data Store Should I Use?

Data structure? → Fixed schema, JSON, key-value

Access patterns? → Store data in format you will access it

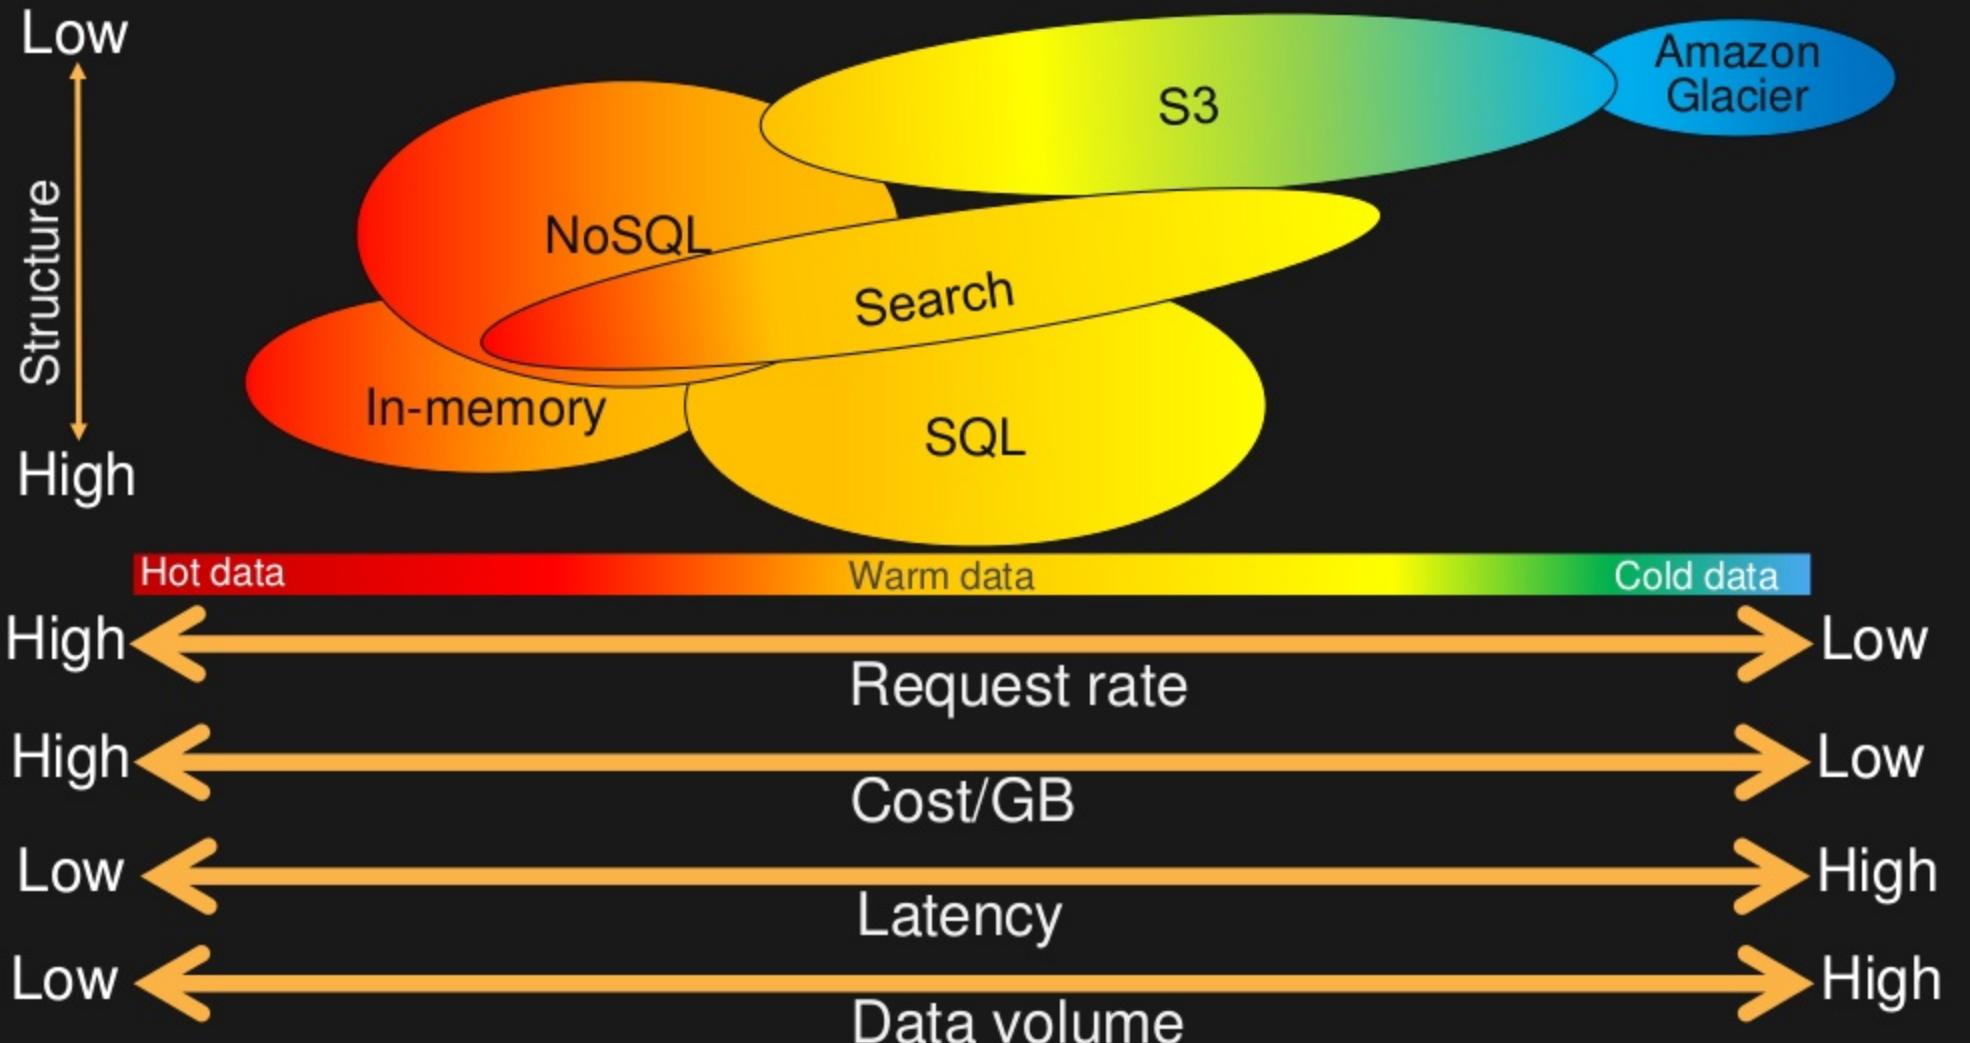
Data / access characteristics? → Hot, warm, cold

Cost? → Right cost

Data Structure and Access Patterns

Access Patterns	What to use?
Put/Get (key, value)	Cache, NoSQL
Simple relationships → 1:N, M:N	NoSQL
Cross table joins, transaction, SQL	SQL
Faceting, search	Search

Data Structure	What to use?
Fixed schema	SQL, NoSQL
Schema-free (JSON)	NoSQL, Search
(Key, value)	Cache, NoSQL



What Data Store Should I Use?

	Amazon ElastiCache	Amazon DynamoDB	Amazon RDS/Aurora	Amazon Elasticsearch	Amazon S3	Amazon Glacier
Average latency	ms	ms	ms, sec	ms,sec	ms,sec,min (~ size)	hrs
Typical data stored	GB	GB–TBs (no limit)	GB–TB (64 TB max)	GB–TB	MB–PB (no limit)	GB–PB (no limit)
Typical item size	B-KB	KB (400 KB max)	KB (64 KB max)	KB (2 GB max)	KB-TB (5 TB max)	GB (40 TB max)
Request Rate	High – very high	Very high (no limit)	High	High	Low – high (no limit)	Very low
Storage cost GB/month	\$\$	¢¢	¢¢	¢¢	¢	¢/10
Durability	Low - moderate	Very high	Very high	High	Very high	Very high
Availability	High 2 AZ	Very high 3 AZ	Very high 3 AZ	High 2 AZ	Very high 3 AZ	Very high 3 AZ

Hot data

Warm data

Cold data

Cost Conscious Design

Example: Should I use Amazon S3 or Amazon DynamoDB?

“I’m currently scoping out a project. The design calls for **many small files**, perhaps up to a **billion during peak**. The **total size** would be on the order of **1.5 TB per month...**”

Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
300	2048	1483	777,600,000

Cost Conscious Design

Example: Should I use Amazon S3 or Amazon DynamoDB?



Simple Monthly
Calculator

<https://calculator.s3.amazonaws.com/index.html>

Amazon S3 or DynamoDB?

Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
300	2,048	1,483	777,600,000

Amazon DynamoDB is a high performance non-relational database service that is easy to set up, operate, and scale. It is designed to address the core problems of database management, performance, scalability, and reliability. It also provides predictable high performance and low latency at scale.

Indexed Data Storage:

Dataset Size:

1483 GB

Provisioned Throughput Capacity *:

Item Size (All attributes):

2 KB

Number of items read per second:

0 Reads/Second

Read Consistency:

Strongly Consistent

Eventually Consistent
(cheaper)

Number of items written per second:

300 Writes/Second

 Amazon S3 is storage for the Internet. It is designed to make web-scale computing easier for developers.

Storage:

Storage:	1483 GB
Reduced Redundancy Storage:	0 GB

Requests:

PUT/COPY/POST/LIST Requests:	77760000 Requests
GET and Other Requests:	0 Requests

Amazon S3 Service (US-East)	\$ 3932.27
Storage:	\$ 44.27
Put/List Requests:	\$ 3888.00

Amazon DynamoDB Service (US-East)	\$ 644.30
Provisioned Throughput Capacity:	\$ 261.69
Indexed Data Storage:	\$ 382.61



SIMPLE MONTHLY CALCULATOR



use

Amazon DynamoDB

<u>Amazon S3 Service (US-East)</u>	\$ 3932.27
Storage:	\$ 44.27
Put/List Requests:	\$ 3888.00
<u>Amazon DynamoDB Service (US-East)</u>	\$ 644.30
Provisioned Throughput Capacity:	\$ 261.69
Indexed Data Storage:	\$ 382.61
DynamoDB Streams:	\$ 0.00

	Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
<u>Scenario 1</u>	300	2,048	1,483	777,600,000
<u>Scenario 2</u>	300	32,768	23,730	777,600,000



use

Amazon S3

<u>Amazon S3 Service (US-East)</u>	\$ 4588.55
Storage:	\$ 700.55
Put/List Requests:	\$ 3888.00
<u>Amazon DynamoDB Service (US-East)</u>	\$ 10131.40
Provisioned Throughput Capacity:	\$ 4187.04
Indexed Data Storage:	\$ 5944.36
DynamoDB Streams:	\$ 0.00

Cost Conscious Design

Example: Should I use Amazon S3 or Amazon DynamoDB?

“I’m currently scoping out a project that will greatly increase my team’s use of Amazon S3. Hoping you could answer some questions. The current iteration of the design calls for **many small files**, perhaps up to a **billion during peak**. The **total size** would be on the order of **1.5 TB per month...**”

Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
300	2048	1483	777,600,000

Cost Conscious Design

Example: Should I use Amazon S3 or Amazon DynamoDB?



Simple Monthly
Calculator

<https://calculator.s3.amazonaws.com/index.html>

Amazon S3 or Amazon DynamoDB?

Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
300	2,048	1,483	777,600,000

Amazon DynamoDB is a high performance non-relational database service that is easy to set up, operate, and scale. It is designed to address the core problems of database management, performance, scalability, and reliability. It also provides predictable high performance and low latency at scale.

Indexed Data Storage:

Dataset Size:

1483 GB

Provisioned Throughput Capacity *:

Item Size (All attributes):

2 KB

Number of items read per second:

0 Reads/Second

Read Consistency:

Strongly Consistent

Eventually Consistent
(cheaper)

Number of items written per second:

300 Writes/Second

Amazon S3 is storage for the Internet. It is designed to make web-scale computing easier for developers.

Storage:

Storage:

1483 GB

Reduced Redundancy Storage:

0 GB

Requests:

PUT/COPY/POST/LIST Requests:

77760000 Requests

GET and Other Requests:

0 Requests

Amazon DynamoDB Service (US-East)

Provisioned Throughput Capacity:

\$ 261.69

\$ 644.30

Indexed Data Storage:

\$ 382.61

Amazon S3 Service (US-East)

Storage:

\$ 3932.27

Put/List Requests:

\$ 44.27

\$ 3888.00



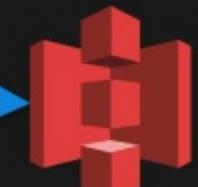
SIMPLE MONTHLY CALCULATOR



use

<u>Amazon S3 Service (US-East)</u>	\$ 3932.27
Storage:	\$ 44.27
Put/List Requests:	\$ 3888.00
<u>Amazon DynamoDB Service (US-East)</u>	\$ 644.30
Provisioned Throughput Capacity:	\$ 261.69
Indexed Data Storage:	\$ 382.61
DynamoDB Streams:	\$ 0.00

	Request rate (Writes/sec)	Object size (Bytes)	Total size (GB/month)	Objects per month
<u>Scenario 1</u>	300	2,048	1,483	777,600,000
<u>Scenario 2</u>	300	32,768	23,730	777,600,000



use

<u>Amazon S3 Service (US-East)</u>	\$ 4588.55
Storage:	\$ 700.55
Put/List Requests:	\$ 3888.00
<u>Amazon DynamoDB Service (US-East)</u>	\$ 10131.40
Provisioned Throughput Capacity:	\$ 4187.04
Indexed Data Storage:	\$ 5944.36
DynamoDB Streams:	\$ 0.00

Amazon S3





PROCESS /
ANALYZE

Analytics Types & Frameworks

Batch

Takes minutes to hours

Example: Daily/weekly/monthly reports

Amazon EMR (MapReduce, Hive, Pig, Spark)

Interactive

Takes seconds

Example: Self-service dashboards

Amazon Redshift, Amazon EMR (Presto, Spark)

Message

Takes milliseconds to seconds

Example: Message processing

Amazon SQS applications on Amazon EC2

Stream

Takes milliseconds to seconds

Example: Fraud alerts, 1 minute metrics

Amazon EMR (Spark Streaming), Amazon Kinesis Analytics, KCL, Storm, AWS Lambda

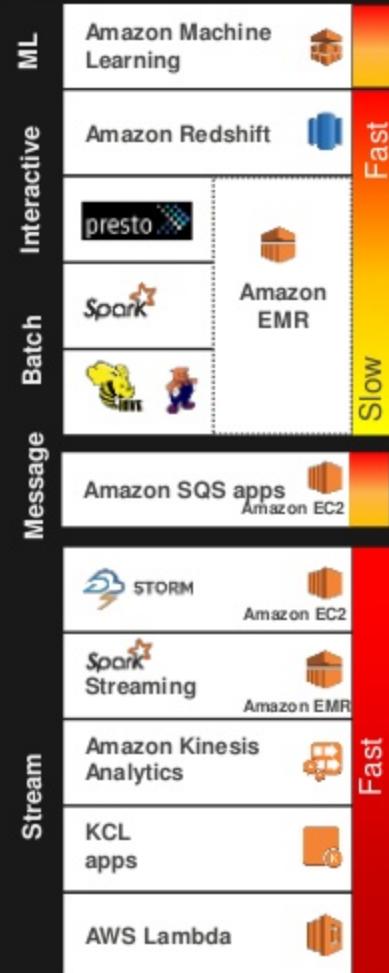
Machine Learning

Takes milliseconds to minutes

Example: Fraud detection, forecast demand

Amazon ML, Amazon EMR (Spark ML)

PROCESS / ANALYZE



What Stream & Message Processing Technology Should I Use?

	Amazon EMR (Spark Streaming)	Apache Storm	KCL Application	Amazon Kinesis Analytics	AWS Lambda	Amazon SQS Application
AWS managed	Yes (Amazon EMR)	No (Do it yourself)	No (EC2 + Auto Scaling)	Yes	Yes	No (EC2 + Auto Scaling)
Serverless	No	No	No	Yes	Yes	No
Scale / throughput	No limits / ~ nodes	No limits / ~ nodes	No limits / ~ nodes	Up to 8 KPU / automatic	No limits / automatic	No limits / ~ nodes
Availability	Single AZ	Configurable	Multi-AZ	Multi-AZ	Multi-AZ	Multi-AZ
Programming languages	Java, Python, Scala	Almost any language via Thrift	Java, others via MultiLangDaemon	ANSI SQL with extensions	Node.js, Java, Python	AWS SDK languages (Java, .NET, Python, ...)
Uses	Multistage processing	Multistage processing	Single stage processing	Multistage processing	Simple event-based triggers	Simple event based triggers
Reliability	KCL and Spark checkpoints	Framework managed	Managed by KCL	Managed by Amazon Kinesis Analytics	Managed by AWS Lambda	Managed by SQS Visibility Timeout

Fast

Which Analysis Tool Should I Use?

	Amazon Redshift	Amazon EMR		
		Presto	Spark	Hive
Use case	Optimized for Data Warehousing	Interactive query	General purpose (iterative ML, RT, ..)	Batch
Scale/throughput	~Nodes	~ Nodes		
Storage	Local Storage	Amazon S3, HDFS		
Optimization	Columnar storage, Data compression, and Zone maps	Framework dependent		
Metadata	Amazon Redshift Managed	Hive Meta-store		
BI tools supports	Yes (JDBC/ODBC)	Yes (JDBC/ODBC & Custom)		
Access controls	Users, Groups and Access Controls	Integration with LDAP		
UDF support	Yes (Scalar)	Yes		
Fast		Slow		

What About ETL?



Data Integration

Reduce the effort to move, cleanse, synchronize, manage, and automatize data related processes.



Attunity CloudBeam



Informatica Cloud



Matillion ETL for Redshift



snapLogic



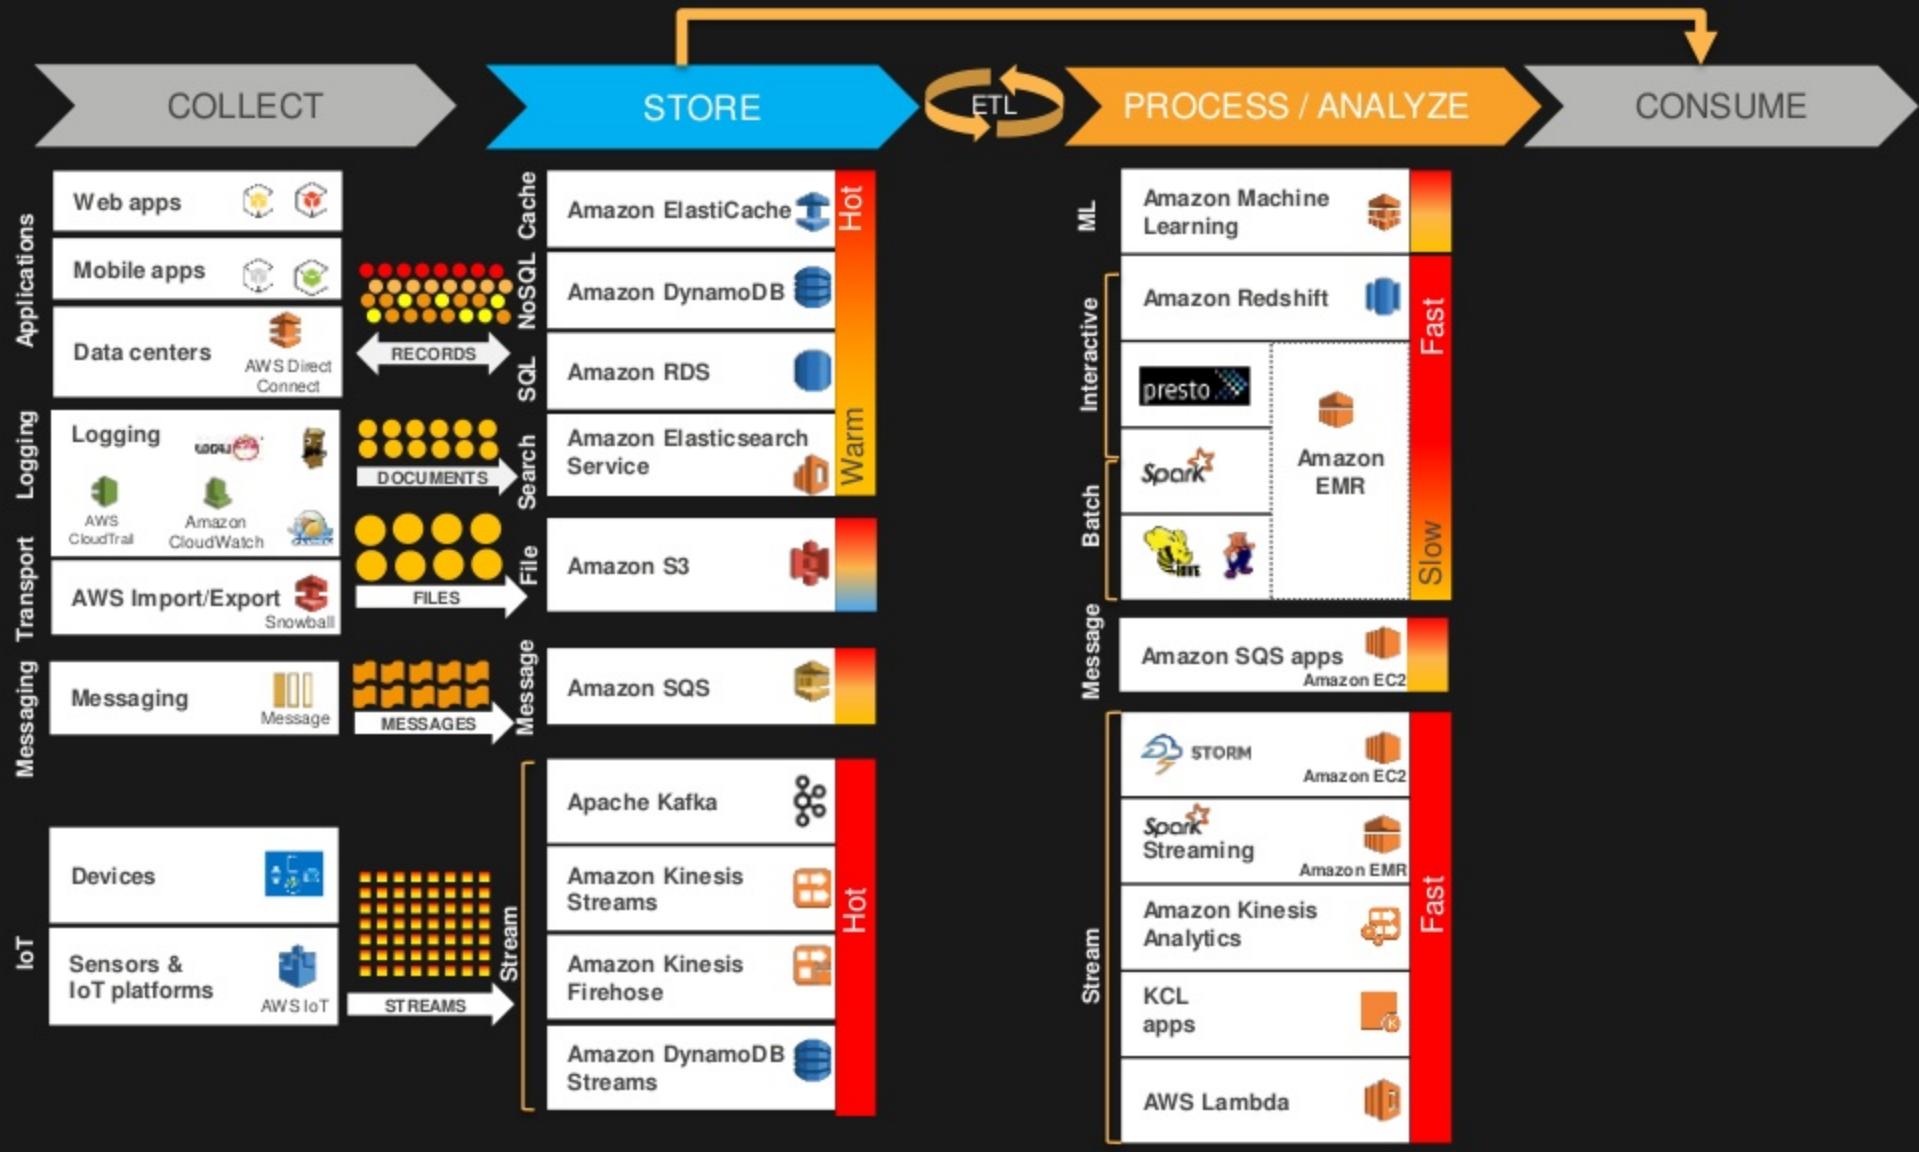
alteryx

• • • •

<https://aws.amazon.com/big-data/partner-solutions/>



CONSUME



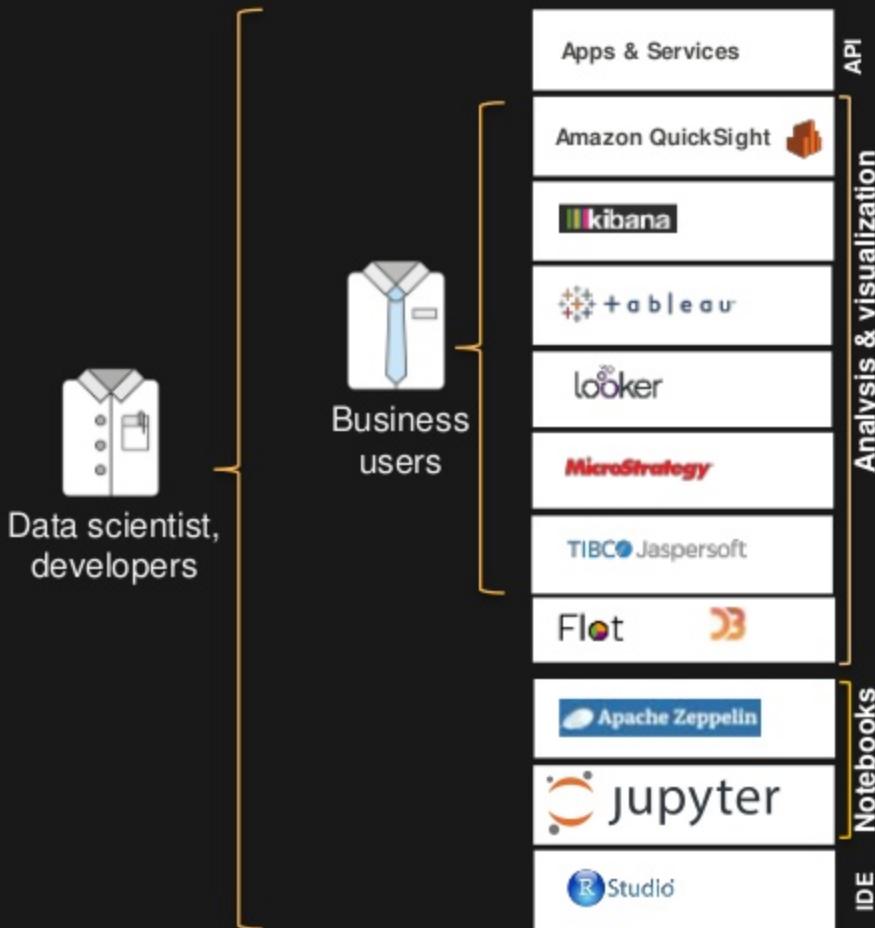


Applications & API

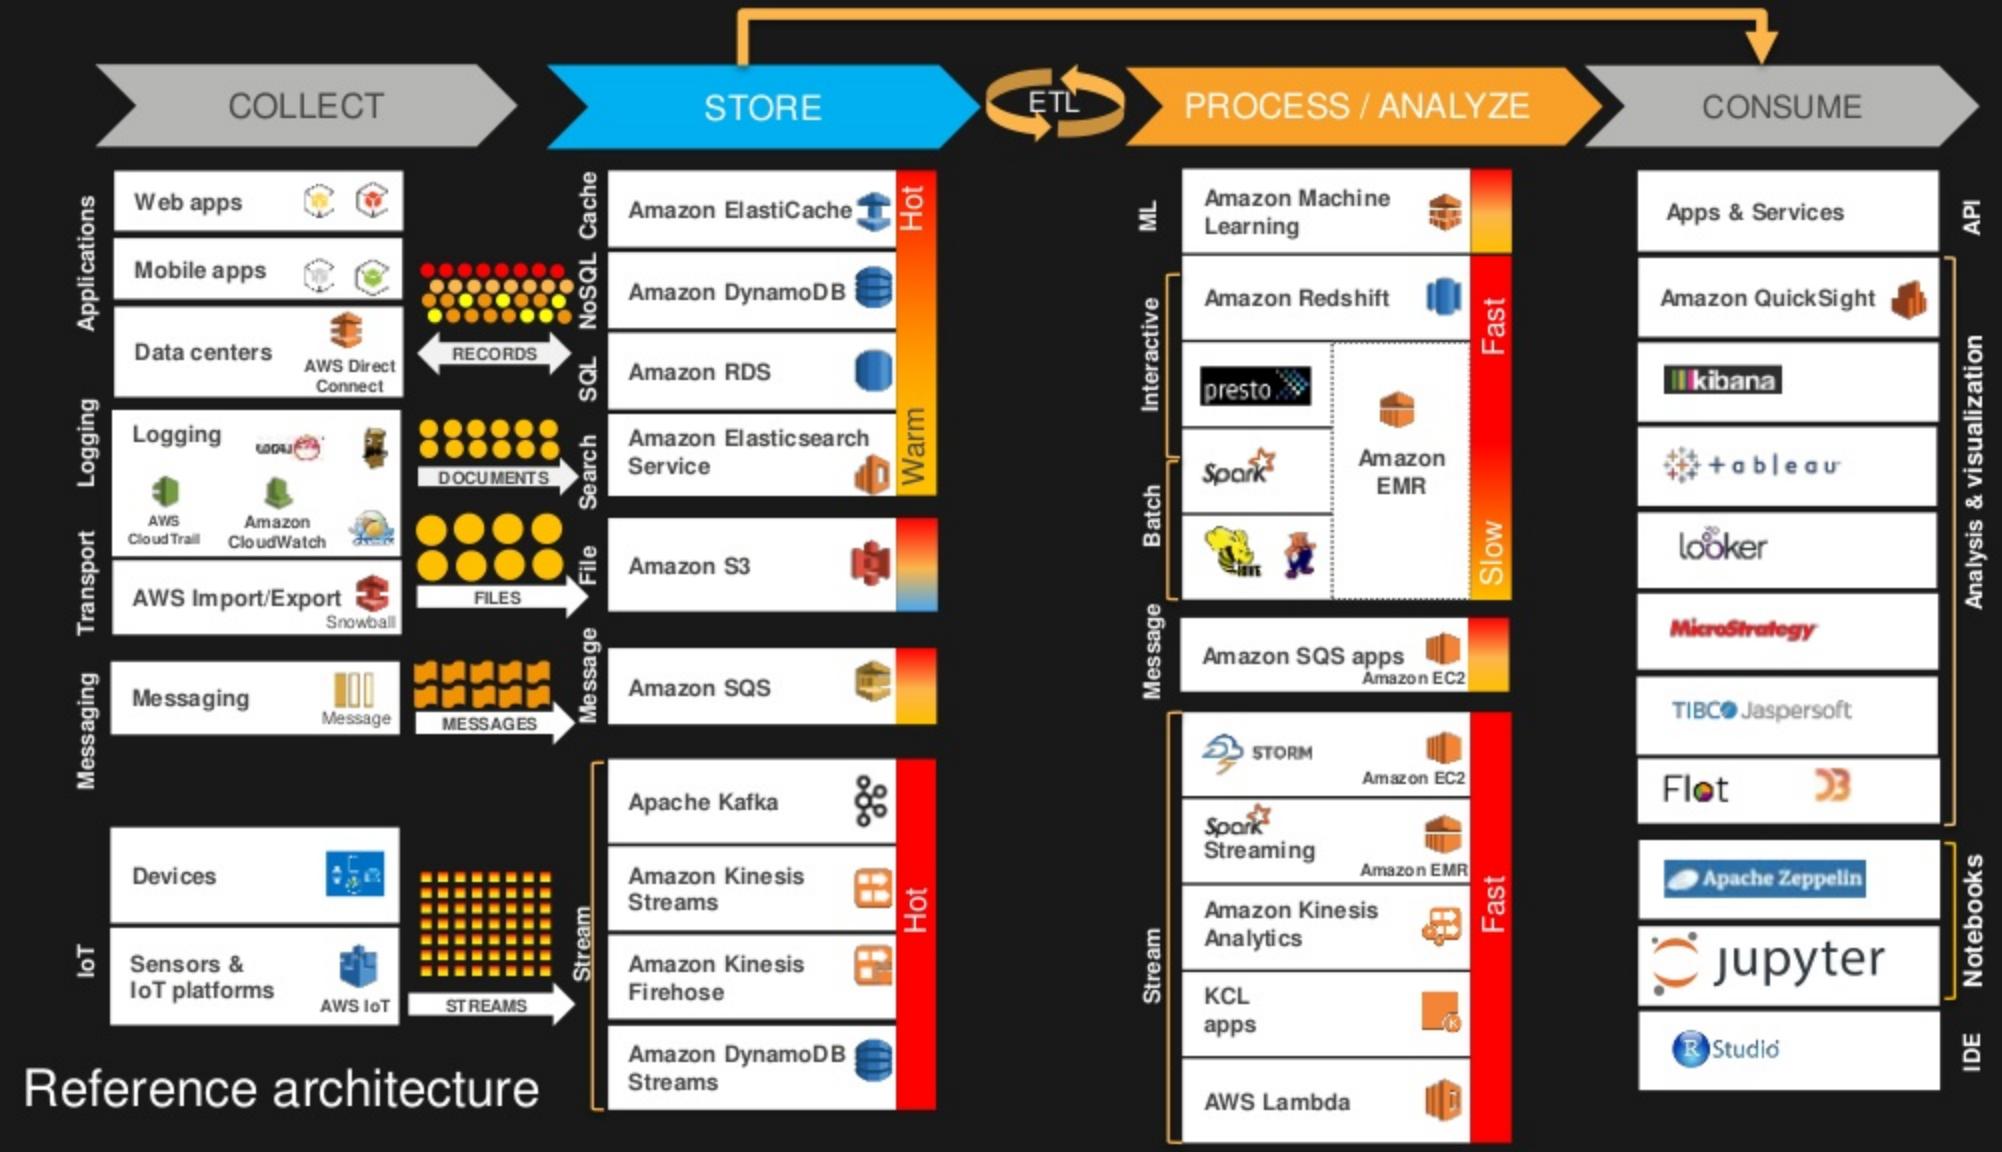
Analysis and visualization

Notebooks

IDE



Putting It All Together



Design Patterns

Primitive: Decoupled “Data Bus”

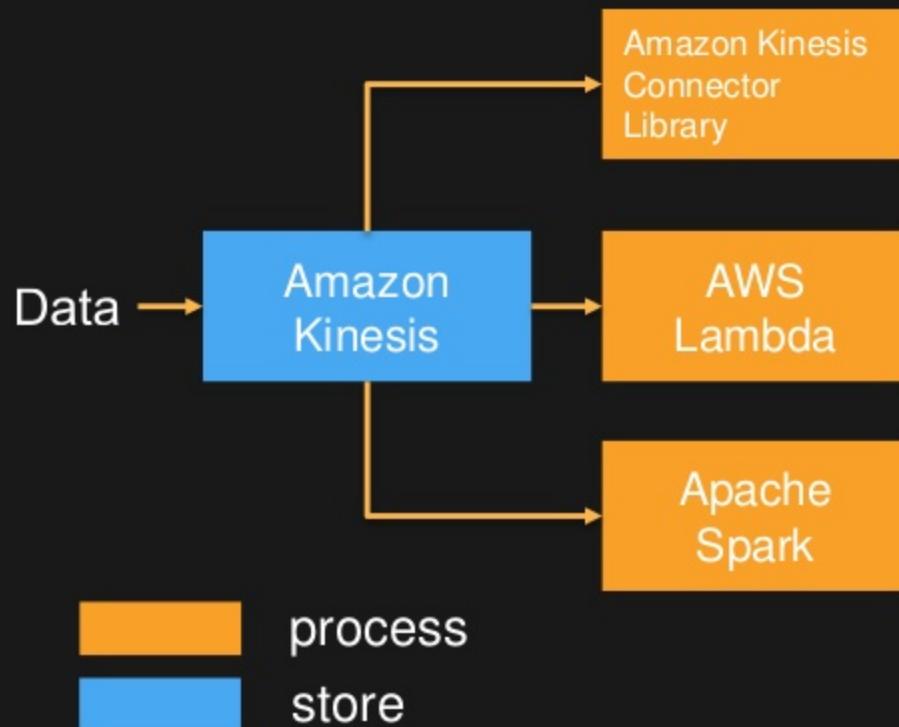
Storage decoupled from processing

Multiple stages



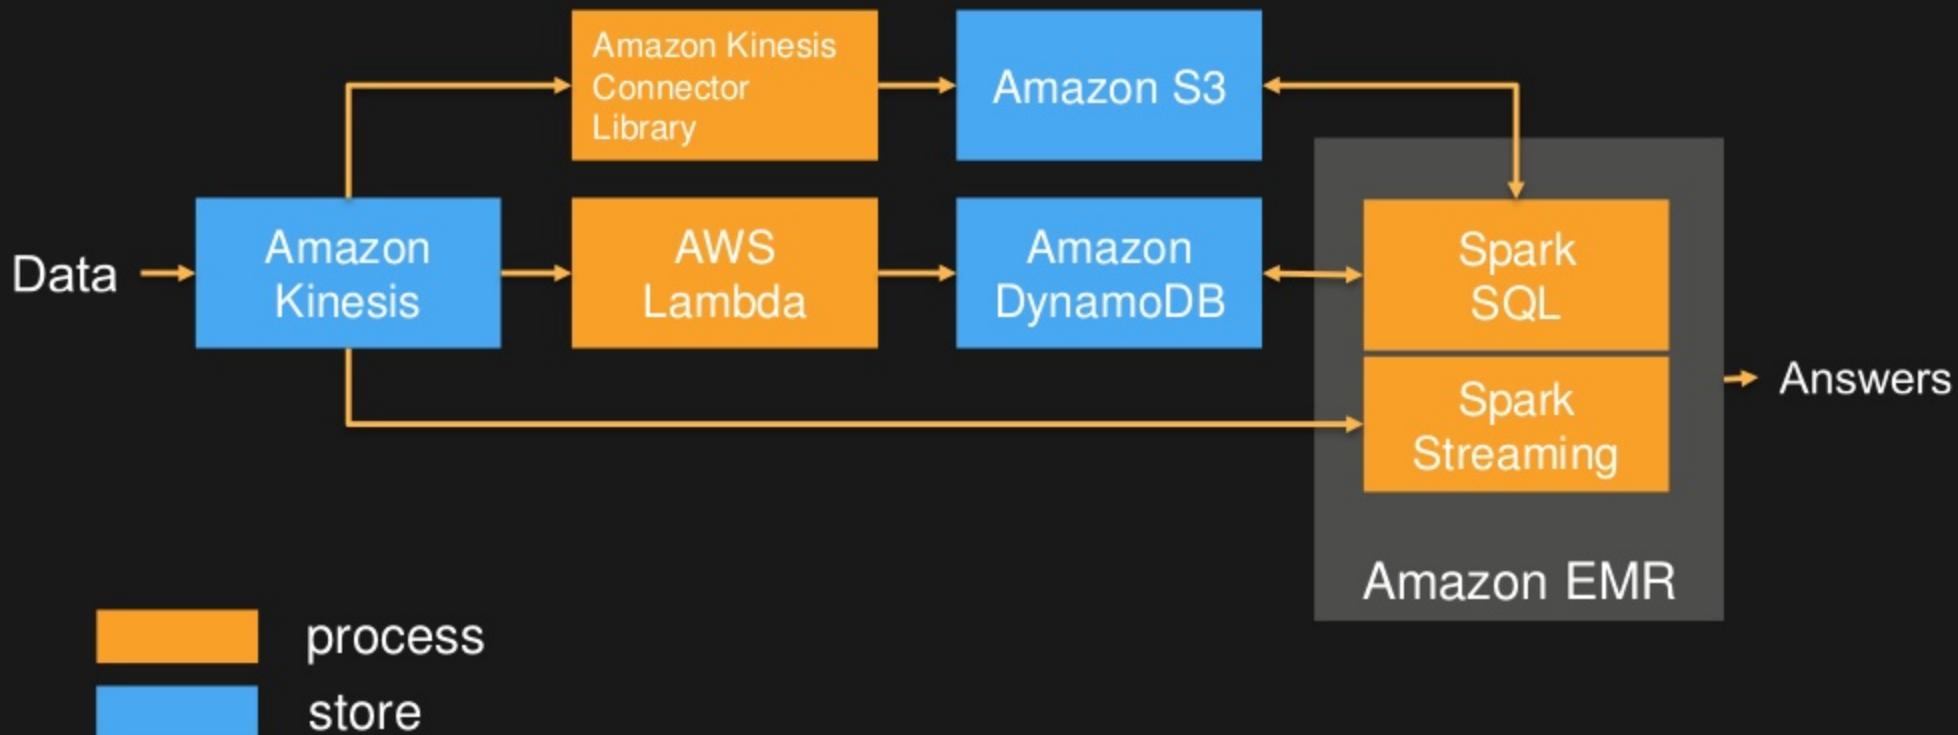
Primitive: Pub/Sub

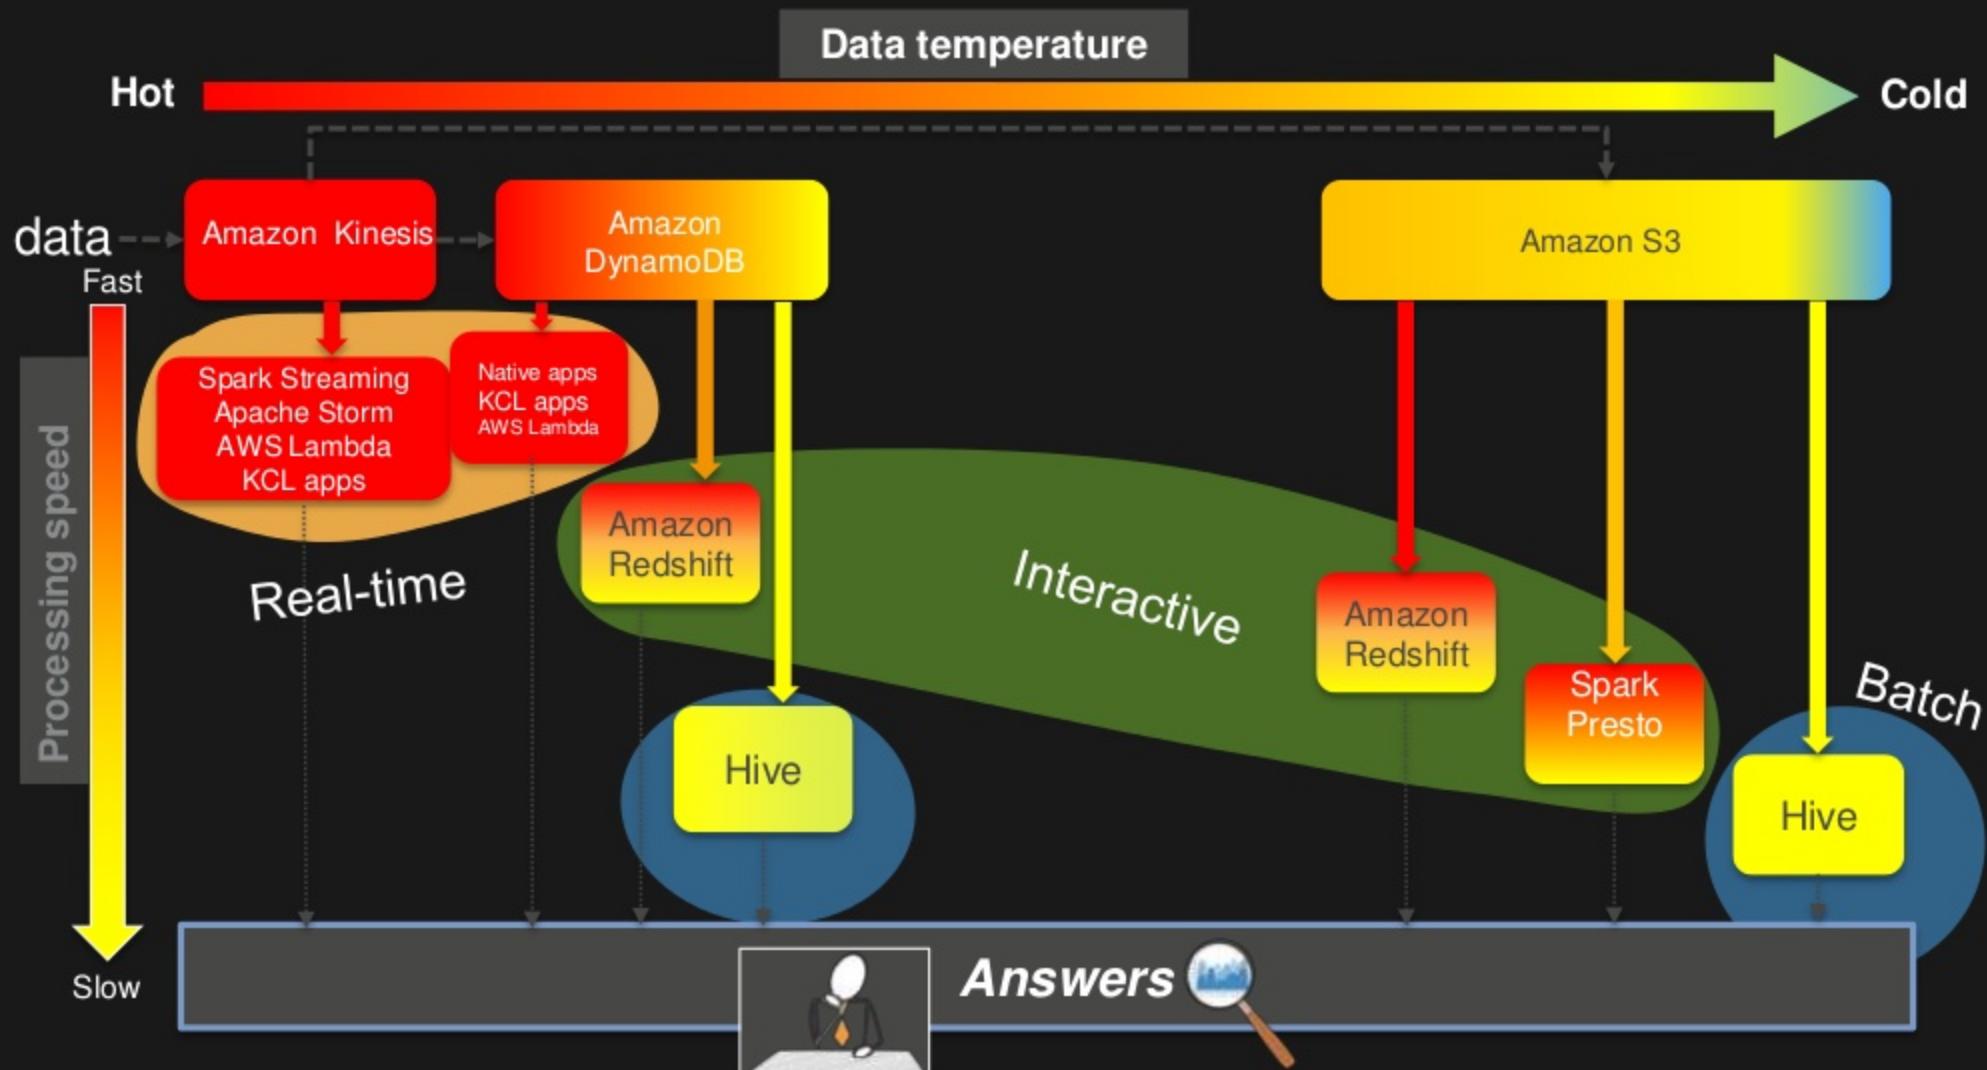
Parallel stream consumption/processing



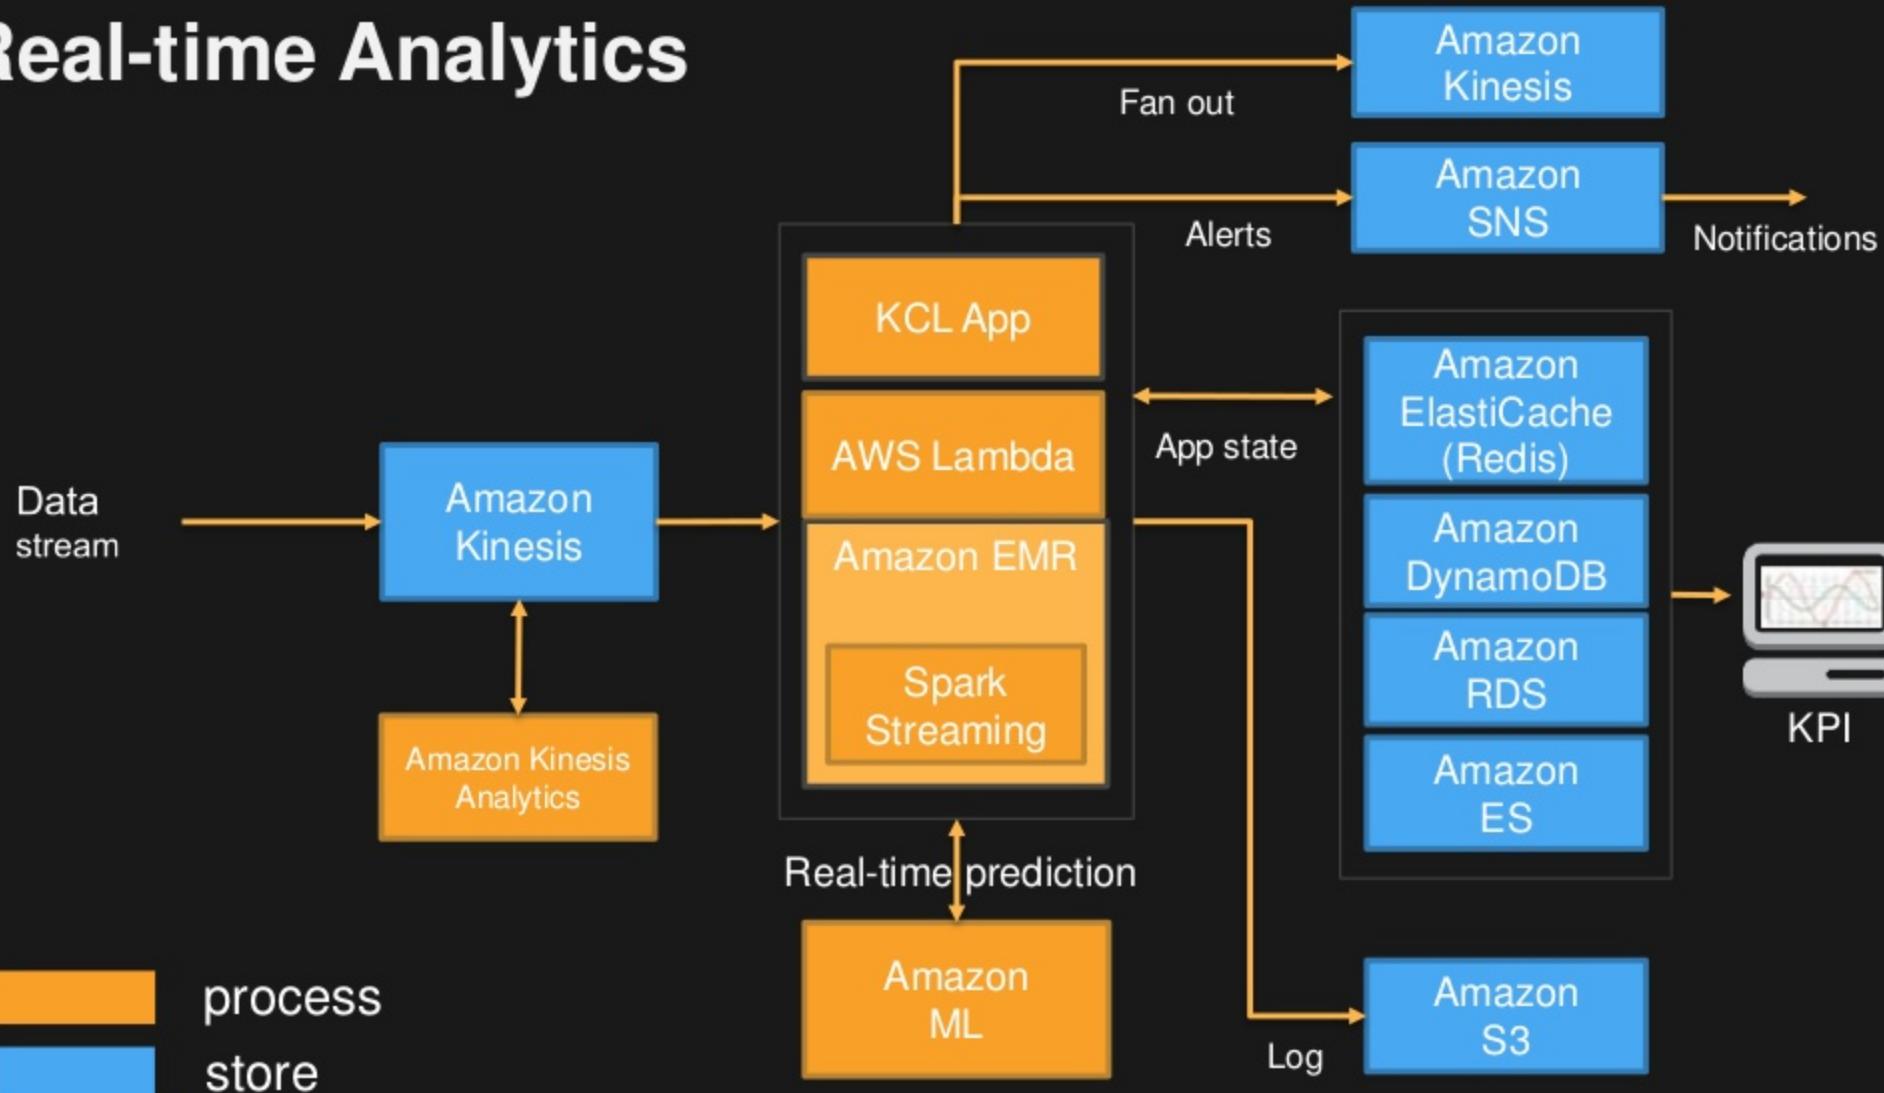
Primitive:

Analysis framework read from or write to multiple data stores





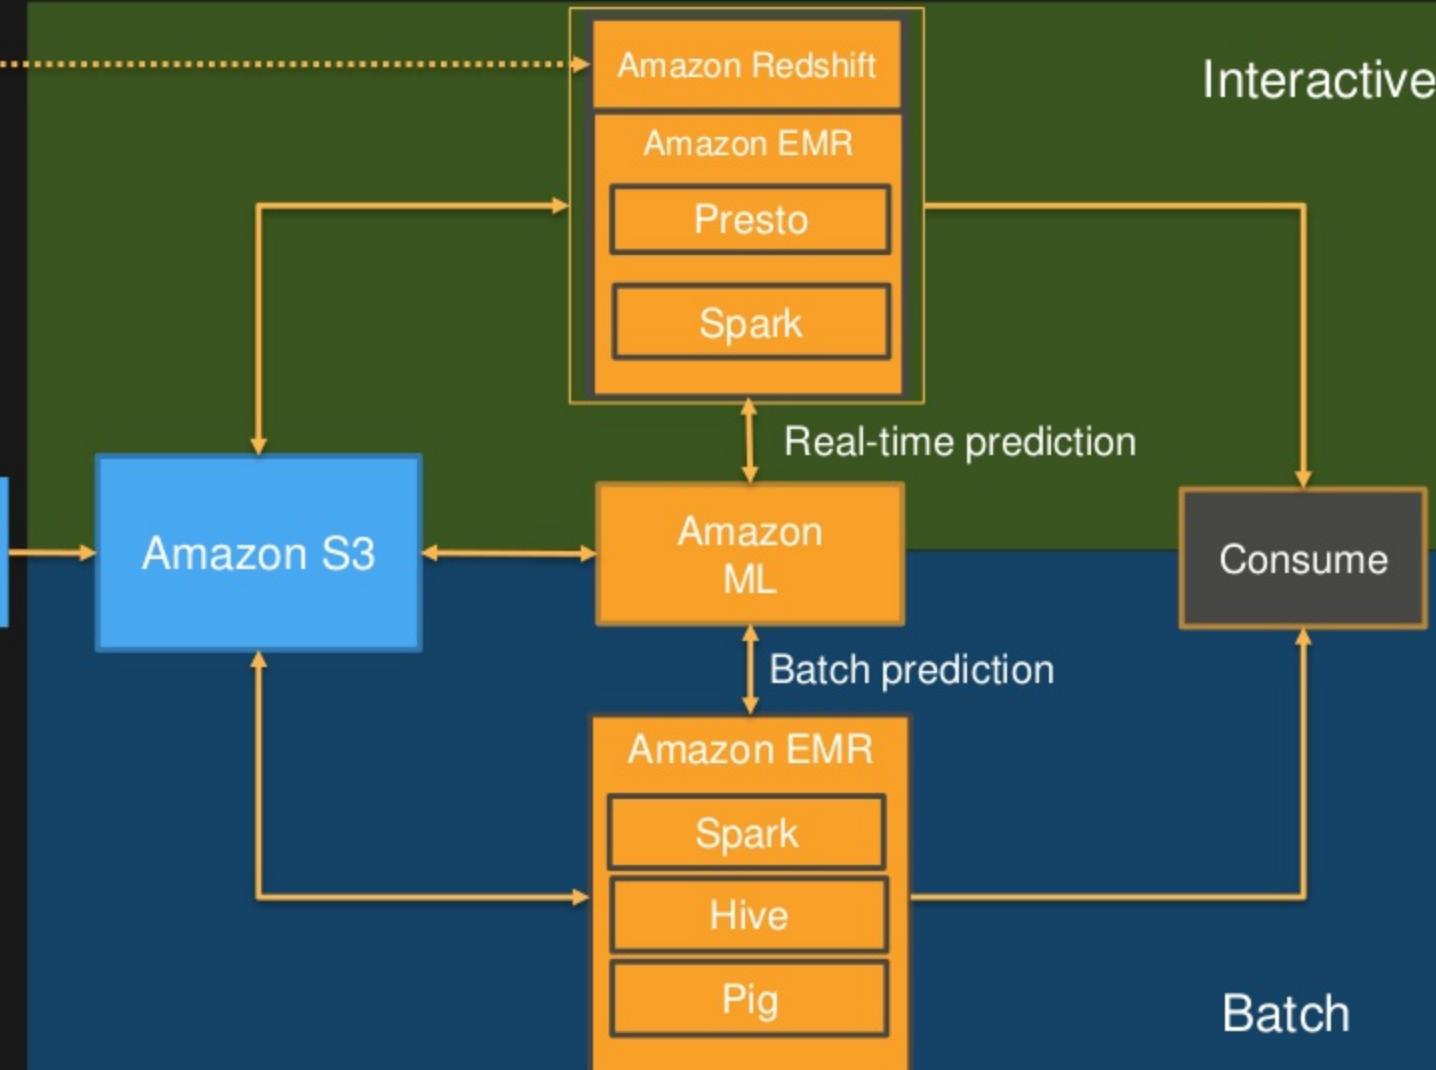
Real-time Analytics



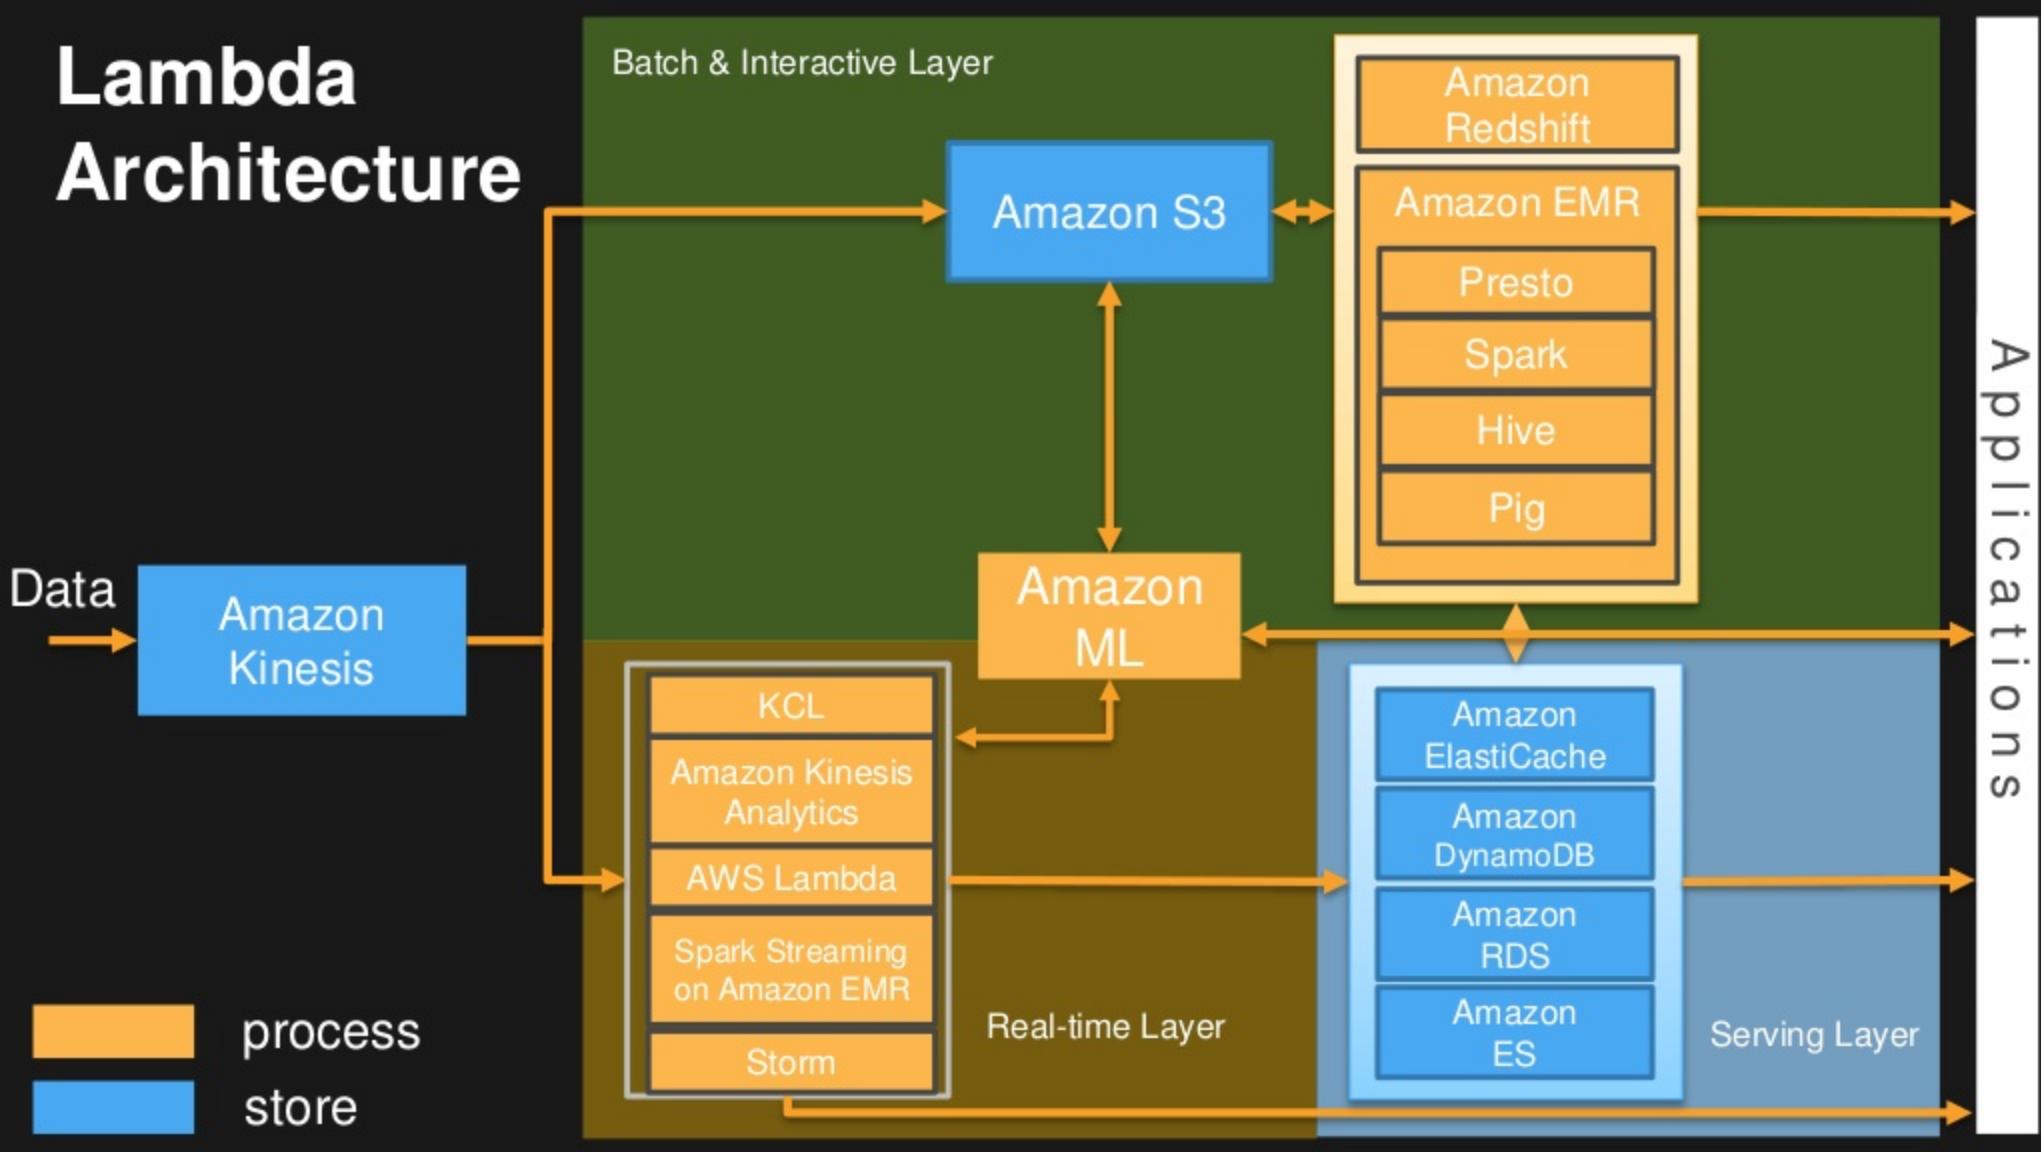
Interactive & Batch Analytics

Data stream →
Amazon Kinesis Firehose

process
store



Lambda Architecture



Summary

Decoupled “data bus”

- Data → **Store** → **Process** → **Store** → **Analyze** → Answers

Use the right tool for the job

- Data structure, latency, throughput, access patterns

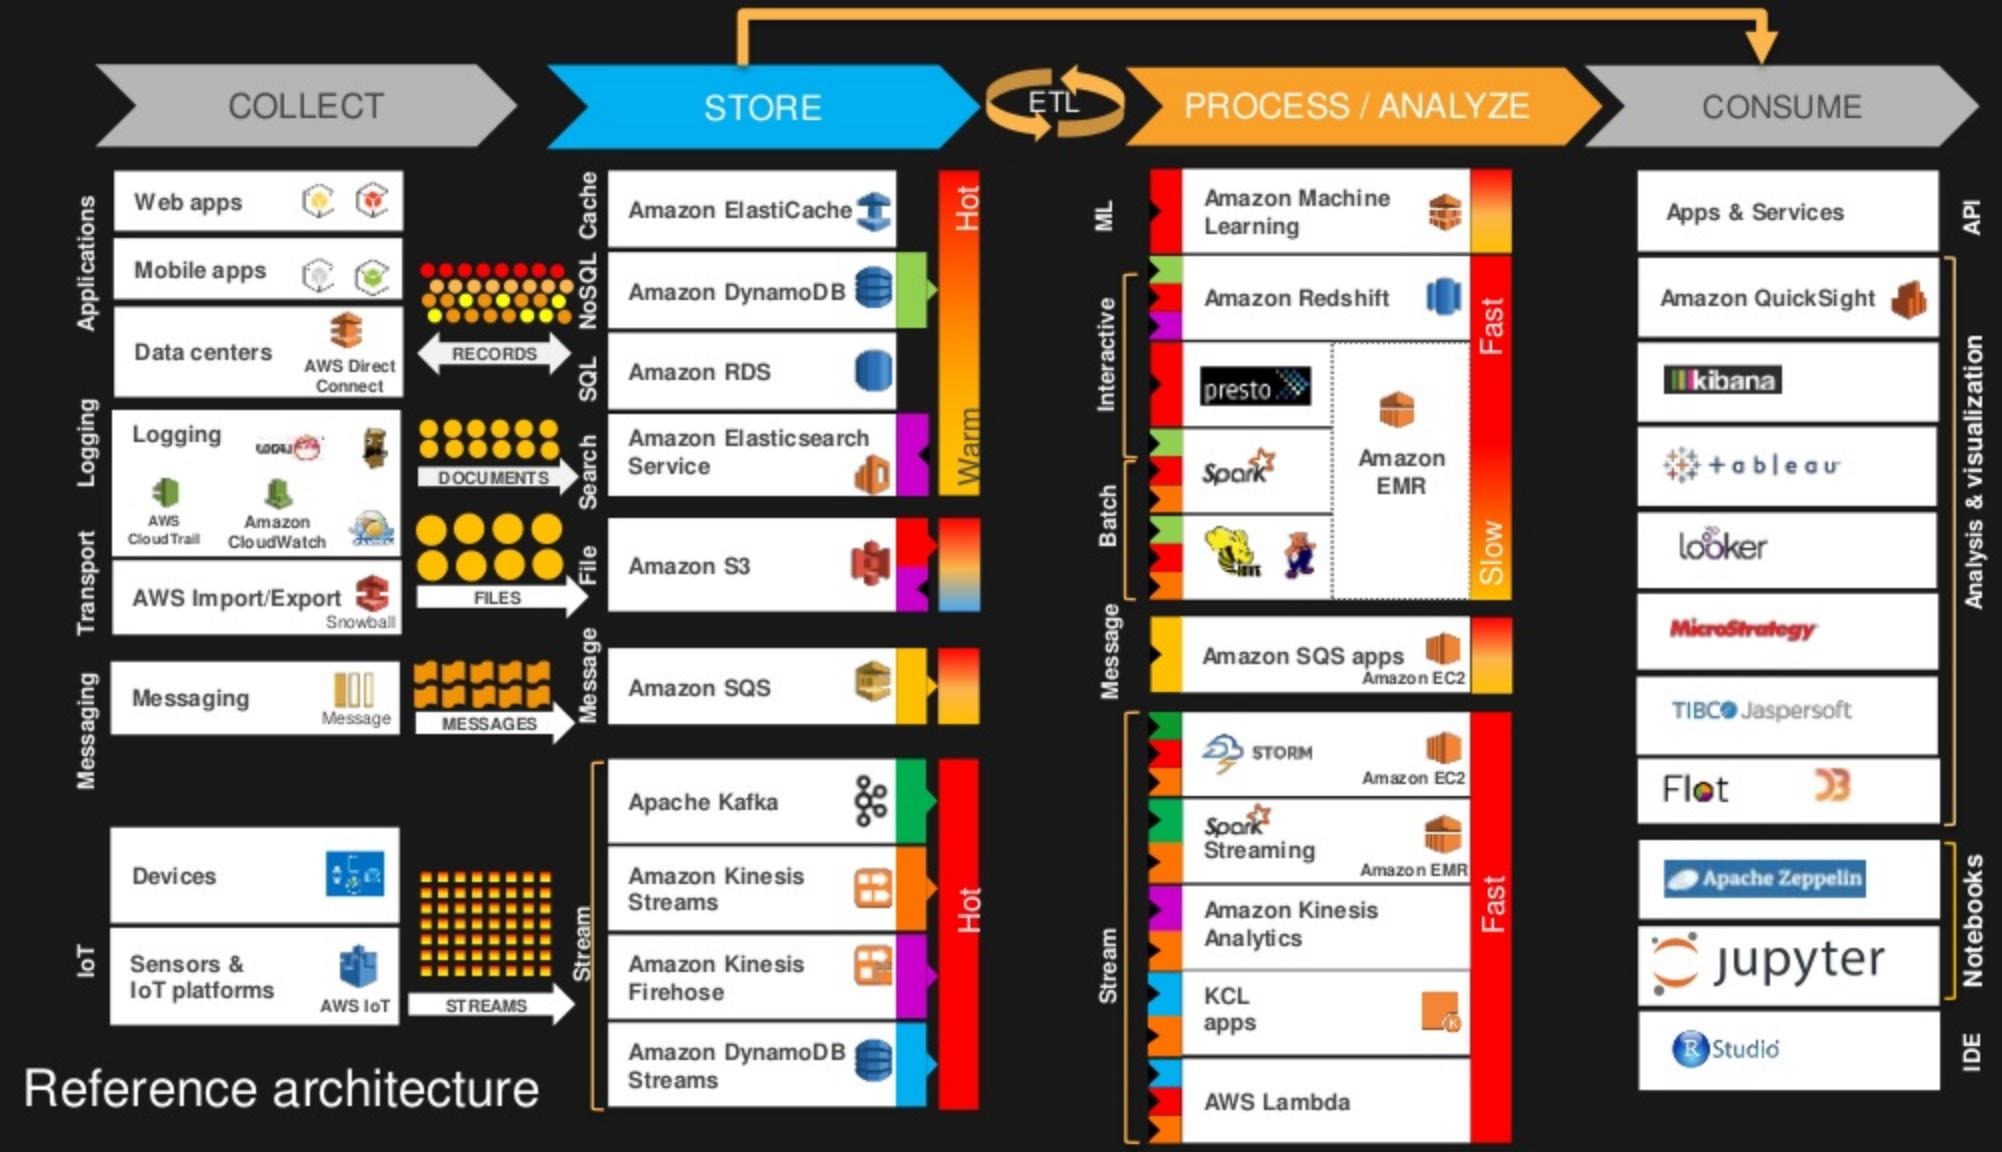
Leverage AWS managed services

- Scalable/elastic, available, reliable, secure, no/low admin

Use Lambda architecture ideas

- Immutable (append-only) log, batch/real-time/serving layer

Big data ≠ big cost



Thank you!
aws.amazon.com/big-data