

# AWS re:Invent


## Getting Maximum Performance from Amazon Redshift: High Concurrency, Fast Data Loads, and Complex Queries



November 13, 2013



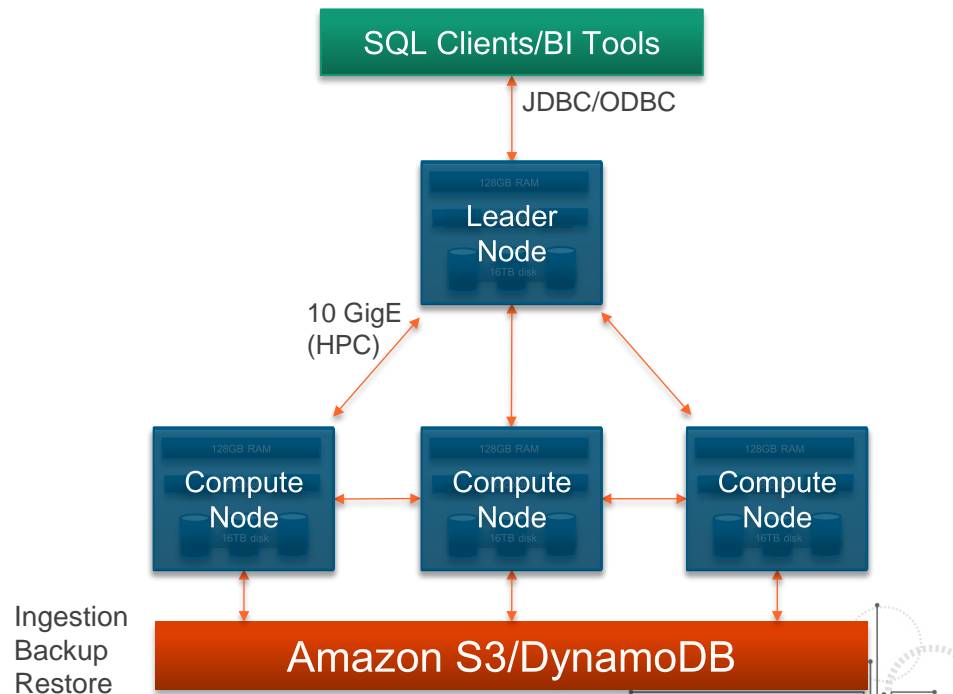
# Amazon Redshift



Fast, simple, petabyte-scale data warehousing for less than \$1,000/TB/Year

# Amazon Redshift architecture

- **Leader Node**
  - SQL endpoint
  - Stores metadata
  - Coordinates query execution
- **Compute Nodes**
  - Local, columnar storage
  - Execute queries in parallel
  - Load, backup, restore via Amazon S3
  - Parallel load from Amazon DynamoDB
- Single node version available



# Amazon Redshift is priced to let you analyze all your data

	Price Per Hour for HS1.XL Single Node	Effective Hourly Price per TB	Effective Annual Price per TB
On-Demand	\$ 0.850	\$ 0.425	\$ 3,723
1 Year Reservation	\$ 0.500	\$ 0.250	\$ 2,190
3 Year Reservation	\$ 0.228	\$ 0.114	\$ 999

## Simple Pricing

Number of Nodes x Cost per Hour

No charge for Leader Node

No upfront costs

Pay as you go





# Getting Maximum Performance from Amazon Redshift: High Concurrency

Ben Myles, Desk.com (@benmyles)

November 13, 2013

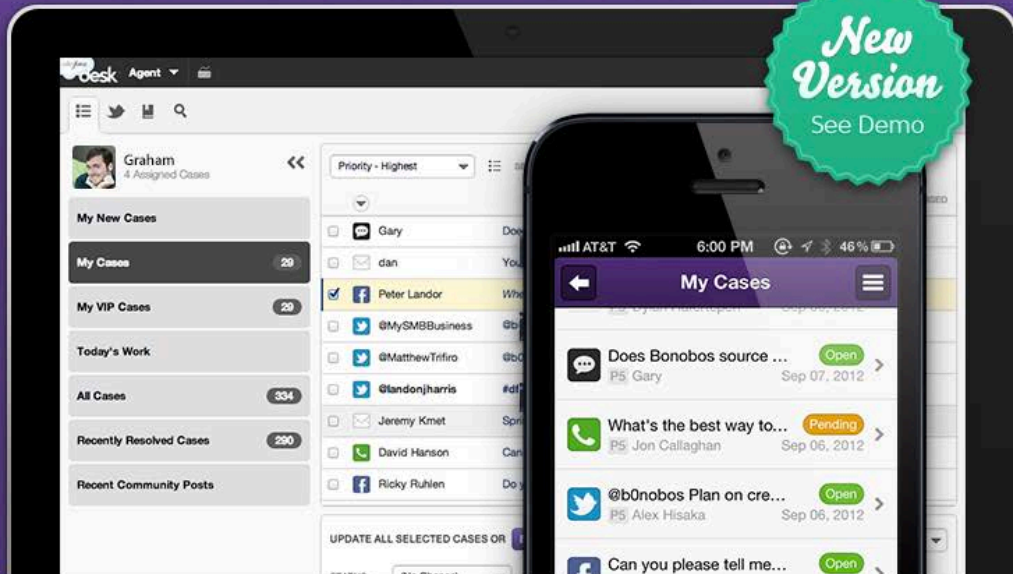


Delight your customers with  
awesome Customer Support

TRY IT FOR FREE TODAY No credit card required.

GET STARTED ▶

REQUEST A DEMO



*"With everything that we do, Desk.com makes us better."*



Fab.

Square

rdio

Eventbrite

TED

yelp

PANDORA

Ben  
4 Assigned Cases

All Cases58,903

Twitter512

Facebook63

EMAIL1,285

Twitter Search56,546

PHONE153

CHAT205

Q&A139

Outbox5

My Cases33

Dev Group45

WoW Group

QA group34

Updated - by Newest

SELECT: ALL, NONE

57,001NEW

1,077OPEN

4PENDING

11RESOLVED

810CLOSED

<input type="checkbox"/>		@Ruth_Parker	boom	Example	Facebook	social	TS	twitter	@ashahgill Is this whether you should	1	New	4 hours ago	General	5
<input type="checkbox"/>		@EBW	boom	Example	Facebook	social	TS	twitter	Admittedly it doesn't help that all the	1	New	4 hours ago	General	5
<input type="checkbox"/>		@FearVan	boom	Example	Facebook	social	TS	twitter	When you look me in the eyes, and tell	1	New	4 hours ago	General	5
<input type="checkbox"/>		@LoveClerMagis	boom	Example	Facebook	social	TS	twitter	@ygnrMISTRAL, @mChilliand@	1	New	4 hours ago	General	5
<input type="checkbox"/>		Bryon	Abandoned Chats	qna					Consequatur reiciendis est aperiam maiores.	1	New	4 hours ago	General	6
<input type="checkbox"/>		Bryon	Abandoned Chats	qna					Consequatur reiciendis est aperiam maiores.	1	New	4 hours ago	General	6
<input type="checkbox"/>		Bryon	Abandoned Chats	qna					Consequatur reiciendis est aperiam maiores.	1	New	4 hours ago	General	6
<input type="checkbox"/>		Desuku	Example	Facebook	Feedback				Hello, The sun is shining, a new day is upon us. Time is	1	New	4 hours ago	QA	5
<input type="checkbox"/>		Desuku	Example	Facebook	Feedback				Hello, The sun is shining, a new day is upon us. Time is	1	New	4 hours ago	QA	5
<input type="checkbox"/>		Desuku	Example	Facebook	Feedback				Message without logging in.	1	New	4 hours ago	QA	5
<input type="checkbox"/>		Desuku	Example	Facebook	Feedback				Boom Boom Pow	1	New	4 hours ago	QA	5
<input type="checkbox"/>		Desuku	Example	Facebook	Feedback				Hello dude, the time is 2013-10-10 16:27:21 -0700	1	New	4 hours ago	QA	5
<input type="checkbox"/>		Smart	Example	Facebook	Feedback				Hello, The sun is shining, a new day is upon us. Time is	1	New	4 hours ago	QA	5
<input type="checkbox"/>		glen	More Info						Quisquam dolor sapiente voluptatem.	1	New	5 hours ago	Chai	9
<input type="checkbox"/>		glen	More Info						Quisquam dolor sapiente voluptatem.	1	New	5 hours ago	Chai	9
<input type="checkbox"/>		glen	More Info						Quisquam dolor sapiente voluptatem.	1	Open	5 hours ago	Chai	9
<input type="checkbox"/>		glen	More Info						Quisquam dolor sapiente voluptatem.	1	Open	5 hours ago	Chai	9
<input type="checkbox"/>		glen	More Info						Quisquam dolor sapiente voluptatem.	1	New	5 hours ago	Chai	9
<input type="checkbox"/>		glen	More Info						Quisquam dolor sapiente voluptatem.	1	New	5 hours ago	Chai	9



CASE

CUSTOMER

COMPANY

INFO LINKS

SUBJECT

Nihil non eum optio hic.

STATUS

Open

PRIORITY

4

GROUP

Development

AGENT

(Unassigned)

DESCRIPTION

Asperiores ad sed nobis eum beatae. Expedita id consequatur voluptates enim aut quia. Dolore atque voluptatibus est. Illo non totam corrupti commodi nobis sit. Nisi magni aut velit.

LABELS

New Email

DIRECT CASE LINK

https://reporting.deskqa7.com/agent/issue/123456

0 ATTACHMENTS

Add

9 OTHER OPEN CASES FOR THIS CUSTOMER

View All

EMAIL

CUSTOMER HISTORY (10)

COMPANY HISTORY

Nihil non eum optio hic.

To: (Empty)

From: reporting.deskqa7.com@reporting.deskqa7.com

Asperiores ad sed nobis eum beatae. Expedita id consequatur voluptates enim aut quia. Dolore atque voluptatibus est. Illo non totam corrupti commodi nobis sit.

HIDE

You can format your message with Textile Syntax (formatting help here)

ADD CC ADD BCC ADD ATTACHMENT LESS

Add

SUBJECT

Re: Nihil non eum optio hic.

ATTACHMENTS

TO

reporting.deskqa7.com@reporting.deskqa7.com

CC

BCC

APPLY MACRO

Close Tab

Update



1st Contact Resolution ?

0%

high: 0%  
low: 0% start: 0%  
end: 0%

Case Reopen Rate ?

0%

high: 0%  
low: 0% start: 0%  
end: 0%

Average Handle Time ?

19.4min

high: 115.1  
low: 0 start: 0.1  
end: 0.1

Time to 1st Response ?

12.9d

high: 12.9  
low: 0 start: 0  
end: 0


Time to 1st Resolution ?

15.8hrs

high: 172.9  
low: 0 start: 172.9  
end: 0

## Case Volume Overview ?

by New Cases ▼



## Response Time vs. Resolution Time ?

Calendar Hours

Business Hours

Compare All Hours





Agents Report ▾

09/20/2013 - 09/24/2013 ▾

Active Segment: Channels

Chat


Twitter

Email

Q&A

Phone

Facebook



<div>▾</div>	Agents Report	0:26:49	444	0.2	0:00:10	279:15:33	23.5%	0.0%	
	Name	Time Online	Resolved Cases	Agent Replies Per Resolve	Average Handle Time	Average Time to Resolution	FCR Rate	% Cases Resolved with a Macro	
1	Eric	0:10:13	84	0.8	0:00:32	792:35:53	80.0%	0.0%	
2	Chai	8:40:25	53	0.3	0:00:04	469:49:28	50.0%	0.0%	
3	Jeremy	0:07:11	46	0.0	0:00:00	393:58:54	0.0%	0.0%	
4	A	0:07:55	25	0.0	0:00:00	002:11:31	50.0%	0.0%	



Overview Report ▾

09/20/2013 - 09/24/2013 ▾

Active Segment: Channels

Chat


Twitter

Email

Q&A

Phone

Facebook



Overview Report Date/Time	374 Created Cases	698 Resolved Cases	240:55:16 Average Time to First Response	0.1 Agent Replies Per Resolve	0:00:10 Average Handle Time	296:37:49 Average Time to Resolution	29.2% FCR Rate	
Sun, 09/22/2013	338	433	651:03:41	0.0	0:00:34	573:51:53	66.7%	
Thu, 09/19/2013	33	11	0:00:00	0.5	0:00:00	71:30:46	0.0%	
Mon, 09/23/2013	2	254	312:37:26	0.0	0:00:05	541:08:39	50.0%	
Fri, 09/20/2013	1	0	0:00:00	0.0	0:00:00	0:00:00	0.0%	

# Our Data

- Primarily raw events
  - E.g., 'case closed', 'case received', 'case resolved'
- Nearly 3B event rows
- About 200K+ new events per hour



# Amazon Redshift Environment


- **Read Cluster:** 2+1 Node 8XL
- **Write Cluster:** 2+1 Node XL
- < 1TB disk space used (compression ftw!)
- High CPU and I/O utilization

# Amazon Redshift Performance

- User-facing portal

## Requirements:


- All queries *must* execute in  $< 30$  seconds
- **Target is  $< 5$  seconds**



# Amazon Redshift Performance

## Technique 1: Less Processing at Query Time

- Generate *events from events* to simplify queries
- Example: `case_first_open` from `case_open`
- *However*, increases resources consumed and time taken by hourly data updates



# Amazon Redshift Performance

```
SELECT
```

```
DATE_TRUNC('day', e.created_at) AS t,
```

```
CASE WHEN event_type = 1 THEN
```

```
    COUNT(*) ELSE NULL END AS case_opens,
```

```
CASE WHEN event_type = /*case_first_open*/2 THEN
```

```
    COUNT(*) ELSE NULL END AS case_first_opens
```

```
FROM events e
```

```
WHERE ...
```



# Amazon Redshift Performance

## Technique 2: Read + Write Cluster

- Originally had single cluster
- Each hourly data update used significant CPU + I/O and ran for ~30mins
- Huge impact on query performance




# Amazon Redshift Performance

## Technique 2: Read + Write Cluster

- Added a 'Write' cluster just for processing data updates

**MySQL -> S3 ->**


**'Write' Cluster -> S3 -> 'Read' Cluster**



# Amazon Redshift Performance


## Technique 2: Read + Write Cluster

CPU Utilization



# Concurrency Challenges

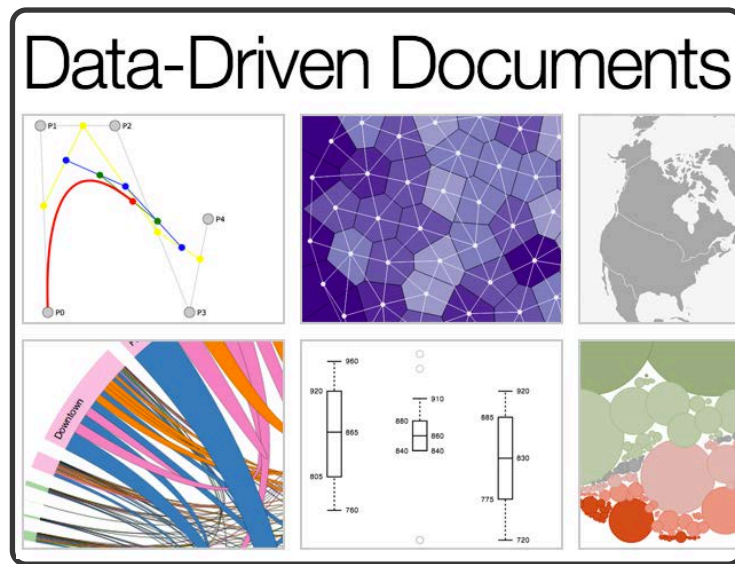
- Queries execute when end-users load reports
- Max of 15 concurrent queries in Amazon Redshift
- Single user rapidly hitting refresh could have big impact



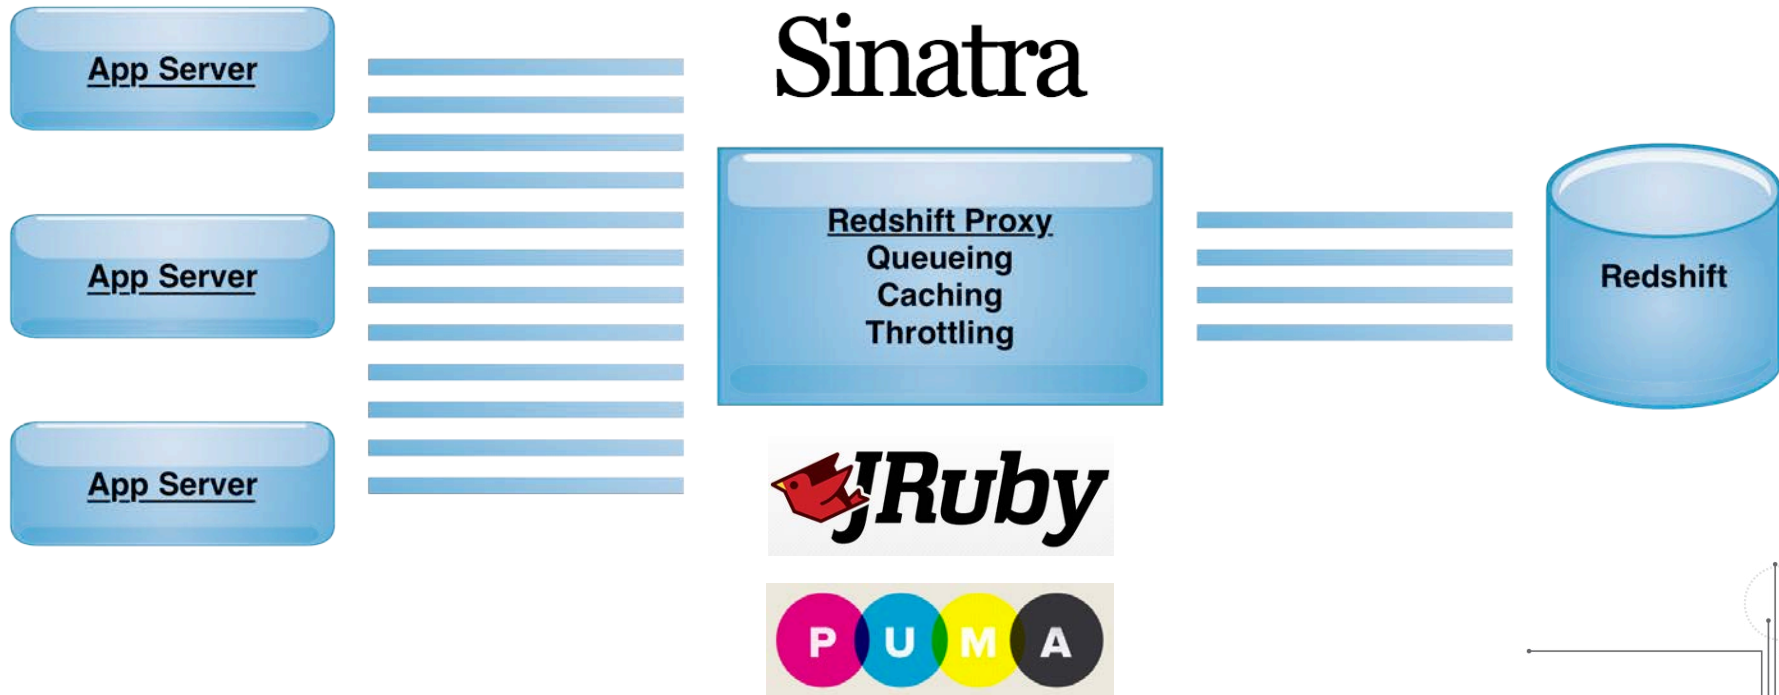


# Solution: Larger Datasets, Fewer Queries

- D3 allows binding MBs of data in browser
- Paging, sorting, filtering is done client-side – instantly



# Solution: Amazon Redshift Proxy



# Amazon Redshift Proxy: Queueing

- Queue requests outside Amazon Redshift
- Minimize (or control) contention
- Abandoned requests don't ever get to Amazon Redshift
- Future: provide ETAs and other queue statistics

# Amazon Redshift Proxy: Caching

- Data only updates once per hour
- Cache all reports (JSON) for duration of hour
- ***Every*** cache hit is a big win
- Just use memcached



# Amazon Redshift Proxy: Throttling

- We can rate limit reports on per-customer basis
- Ensures single customer cannot monopolize resources

# Amazon Redshift Proxy: Example

```
curl http://localhost:9292/query
  -H 'Accept: application/json'
  -H 'Content-Type: application/json'
  -d '{"sql": "SELECT * FROM events LIMIT
?", "params": [2]}'
```

# Amazon Redshift Proxy: Example

```
"header": [  
  ["site_id", "Int"],  
  ["event_type", "Int"],  
  ["created_at", "Timestamp"],  
  ...  
],  
"rows": [  
  [1, 1, "2013-07-05T19:01:32-07:00", ...],  
  [1, 1, "2013-07-05T19:01:37-07:00", ...]  
]
```

# Summary

- Amazon Redshift allowed us to rapidly build and ship our next-gen customer-facing reporting portal
- It's always getting better – improvements and new features nearly every week
- Awesome support and guidance from AWS team

# AWS re:Invent

## Getting Maximum Performance from Amazon Redshift: Fast Data Loads

Niek Sanders, HasOffers


November 13, 2013






- Attribution for web & mobile marketing
- Tons of data
- Ad-hoc analysis
- Near real-time expectation

# Live(ish) Amazon Redshift Data Loading






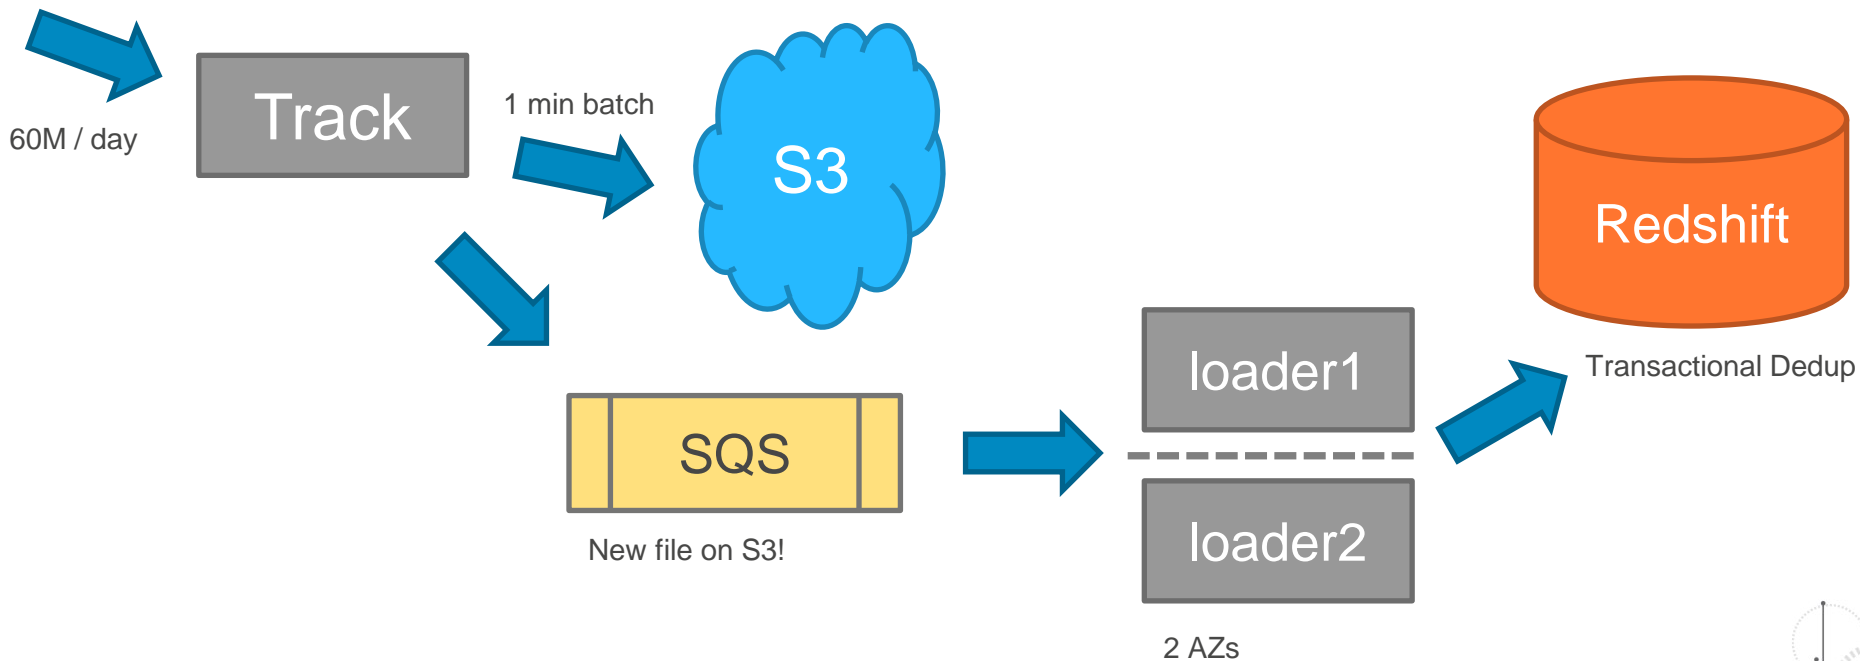
# Objective




# Requires Amazon S3, DML Inserts Slow



# Core Idea




# High Availability




- 2 loaders, separate AZs
- Shared-nothing on loaders
- Amazon SQS re-queue if loader dies
- Idempotent, transactional de-duplicator

# Four Big Ideas

1. Micro-batching and Amazon S3
2. Amazon SQS for work queue
3. Bulk Amazon S3 COPY
4. Transactional de-duplicator



# Four Big Ideas

1. ~~Micro-batching and Amazon S3~~
2. ~~Amazon SQS for work queue~~
3. Bulk Amazon S3 COPY 
4. Transactional de-duplicator

# Amazon S3 COPY

COPY clicks FROM 's3://foo/2014/01/' ...




Single file or path prefix.



# Bulk Amazon S3 COPY

Amazon SQS Work Queue




File Loads


Prefix Loads

# Bulk Amazon S3 COPY

- Bulk loads = 10x speed
- Cluster size, parallelism
  - Roughly equal file size
  - XL nodes → 2 files per node
  - 8XL nodes → 16 files per node



# Bulk Amazon S3 COPY: Common Prefixes




```
s3://foo/customerA/2013/03/clicks81.gz  
s3://foo/customerB/2013/03/clicks93.gz  
s3://foo/customerB/2013/03/clicks30.gz  
s3://foo/customerC/2013/03/clicks01.gz
```



```
s3://foo/common_prefix/
```

# Bulk Amazon S3 COPY: Common Prefixes




New!


COPY clicks  
FROM 's3://foo/files.manifest'  
MANIFEST ...

List of files

# Four Big Ideas

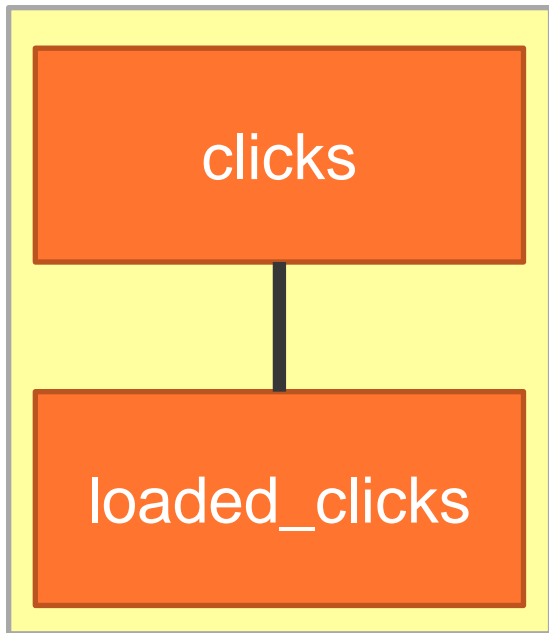
- ~~1. Micro-batching and Amazon S3~~
- ~~2. Amazon SQS for work queue~~
- ~~3. Bulk Amazon S3 COPY~~
4. Transactional de-duplicator 

# Transactional De-duplicator



- Prevent duplicate loads
- Redshift transactions
- Idempotent load flow

# Transactional Deduplicator




- Batch in loaded\_clicks? Stop
- Begin txn
  - COPY
  - Batch in loaded\_clicks? Rollback
  - Insert batch to loaded\_clicks
- Commit




# Four Big Ideas

1. Micro-batching and Amazon S3
2. Amazon SQS for work queue
3. Bulk Amazon S3 COPY
4. Transactional de-duplicator



# Results

Batch Create to Load Latency




min = 15  
max = 733  
avg = 129 sec  
stddev = 47



# Thank you!

[niek@hasoffers.com](mailto:niek@hasoffers.com)





# Getting Maximum Performance from Amazon Redshift: Complex Queries

Timon Karnezos, Aggregate Knowledge

November 13, 2013



## Study: Facebook Leads to 24% Sales Boost

Social network brings new cost, per Aggregate Knowledge

channel to see a user before they convert. That last-click attribution, for better or worse, is what helped Google establish a multibillion dollar advertising business. “Where Facebook gets last-touch credit is where it’s the only place on the planet that reached that user,” Jakubowski said. Facebook tends to factor in a couple days or weeks before a user converts, he noted. Smallwood emphasized that marketers shouldn’t look at Facebook in silo but consider an entire media mix and attribute success through multi-touch attribution.

# Meet the new boss

*Multi-touch Attribution*

## Examples and real world applications [\[edit\]](#)

Data shows that a large percentage of users using a certain eCommerce platform found it by searching for “Thai food” on Google Food” page and then logged off without placing an order. Looking at each of these events as separate data points does not represent the overall user behavior. However, viewing these data points as a representation of overall user behavior enables one to interpolate how and why users act.

Behavioral analytics looks at all site traffic and page views as a timeline of connected events that did not lead to orders. Since there is a disconnect between what they are searching for on Google and what the “Asian Food” page displays. Knowing this, a quick look at the data prominently and thus people do not think it is actually offered, even though it is.

Behavioral analytics is becoming increasingly popular in commercial environments. [Amazon.com](#) is a leader in using behavioral

# Same as the old boss

## *Behavioral Analytics*

## Examples [\[edit source\]](#)


Market basket analysis might tell a retailer that customers often purchase shampoo and conditioner **together**, so putting both items on promotion at the same time would not create a significant increase in profit, while a promotion involving just one of the items would likely drive sales of the other.

Market basket analysis may provide the retailer with information to understand the purchase behavior of a buyer. This information will enable the retailer to understand the buyer's needs and rewrite the store's layout accordingly, develop cross-promotional programs, or even capture new buyers (much like the **cross-selling** concept). An apocryphal early illustrative example for this was when one super market chain discovered in its analysis that customers that bought diapers often bought beer as well, have

# Same as the old boss

## *Market Basket Analysis*






We know how to do this  
*in SQL\**!

*\* SQL:2003*

# Here it is

```
SELECT record_date, user_id, action,  
       site, revenue,  
       SUM(1) OVER  
         (PARTITION BY user_id  
          ORDER BY record_date ASC)  
       AS position  
FROM user_activities;
```

So why is MTA  
*hard?*



# “Web Scale”

## Queries

- 30 queries
- 1700 lines of SQL
- 20+ logical phases
- ***GBs of output***

## Data

- $\sim 10^9$  daily impressions
- $\sim 10^7$  daily conversions
- $\sim 10^4$  daily sites
- ***x 90 days***

***per report***



# So, how do we deliver *complex reports* over *“web scale” data?*

*(Pssst. The answer's Amazon Redshift. Thanks, AWS.)*

***Write (good) queries***


***Organize the data***

***Optimize for the humans***



# *Write (good) queries*

Remember: SQL is **code**






# Software engineering rigor applies to SQL


*Factored*

*Concise*

*Tested*



# Common Table Expression



```
targeted_conversions AS (
  SELECT  record_date,
          ak_user_id,
          request_id,
          c.advertiser_id AS advertiser_id,
          attributed_revenue
  FROM    conversions_unified c
  JOIN    derived_filter_conversions_2013_07_17 fc ON fc.tag_id = c.tag_id AND
          fc.advertiser_id = c.advertiser_id
),
```

Factored  
Concise  
Tested


```
JOIN derived_filter_placements_2013_07_17 fp
  ON fp.inventory_placement_id = i.inventory_placement_id AND
  fp.campaign_id = i.campaign_id
JOIN campaign_metadata_2013_07_17 campaign_metadata
  ON campaign_metadata.campaign_id = i.campaign_id
JOIN tracking_campaign_metadata_2013_07_17 tracking_campaign_metadata
  ON tracking_campaign_metadata.campaign_id = campaign_metadata.campaign_id
JOIN inventory_placement_metadata_2013_07_17 inventory_placement_metadata
  ON inventory_placement_metadata.tracking_data_provider_id = tracking_campaign_metadata.tracking_data_provider_id AND
  inventory_placement_metadata.inventory_placement_id = i.inventory_placement_id
JOIN derived_inventory_provider_name_lookup_2013_07_17 inventory_provider_name_lookup
  ON inventory_provider_name_lookup.inventory_provider_name = inventory_placement_metadata.inventory_provider_name

targeted_conversions AS (
  SELECT  record_date,
          ak_user_id,
          request_id,
          c.advertiser_id AS advertiser_id,
          attributed_revenue
  FROM    conversions_unified c
  JOIN    derived_filter_conversions_2013_07_17 fc ON fc.tag_id = c.tag_id AND
          fc.advertiser_id = c.advertiser_id

  chains AS (
    SELECT  i.advertiser_id AS advertiser_id,
           DATE_TRUNC('day', c.record_date) AS conversion_record_date,
           i.ak_user_id AS ak_user_id,
           c.request_id AS conversion_request_id,
           i.request_id AS impression_request_id,
           i.inventory_provider_name_surrogate_id AS inventory_provider_name_surrogate_id,
           c.attributed_revenue AS conversion_attributed_revenue,
           -- DATETIME doesn't need a DATE_TRUNC because it computes number of day boundaries
           DATEDIFF(day, i.record_date, c.record_date) AS day_offset,
           SUM(1) OVER (PARTITION BY i.advertiser_id, i.ak_user_id, c.request_id ORDER BY i.record_date DESC ROWS UNBOUNDED PRECEDING) AS position
    FROM    targeted_impressions i
    JOIN    targeted_conversions c
           ON i.ak_user_id = c.ak_user_id AND
           i.advertiser_id = c.advertiser_id AND
           i.record_date < c.record_date AND
           i.record_date > (DATE_TRUNC('day', c.record_date) - interval '45 days')
  )
  SELECT * FROM chains;
```



# Window functions



-- Position in timeline

```
SUM(1) OVER (PARTITION BY user_id  
             ORDER BY record_date DESC ROWS UNBOUNDED PRECEDING)
```

-- Event count in timeline


```
SUM(1) OVER (PARTITION BY user_id  
             ORDER BY record_date DESC BETWEEN UNBOUNDED PRECEDING AND  
             UNBOUNDED FOLLOWING)
```

-- Transition matrix of sites

```
LAG(site_name) OVER (PARTITION BY user_id  
                    ORDER BY record_date DESC)
```

-- Unique sites in timeline, up to now

```
COUNT(DISTINCT site_name) OVER (PARTITION BY user_id  
                                ORDER BY record_date DESC  
                                ROWS UNBOUNDED PRECEDING)
```




# Window functions


*Scalable, combinable*

*Compact but expressive*

*Simple to reason about*



# *Organize the data*




# Leverage Amazon Redshift's MPP roots

*Fast, columnar scans, IO*

*Fast sort and load*

*Effective when work is distributable*






# Leverage Amazon Redshift's MPP roots

*Sort into multiple representations*

*Materialize shared views*

*Hash-partition by user\_id*

***Optimize for the humans***




# Operations should not be the bottleneck

*Develop without fear*

*Trade time for money*

*Scale with impunity*




# Operations should not be the bottleneck

*Fast Amazon S3 = scratch space for cheap*

*Linear query scaling = GTM quicker*

*Dashboard Ops = dev/QA envs, marts, clusters  
with just a click*

***But, be frugal***




# Quantify and control costs

*Test across different hardware, clusters*

*Shut down clusters often*

*Buy productivity, not bragging rights*



# *Thank you!*

## References

[http://bit.ly/rs\\_ak](http://bit.ly/rs_ak)

<http://www.adweek.com/news/technology/study-facebook-leads-24-sales-boost-146716>

[http://en.wikipedia.org/wiki/Behavioral\\_analytics](http://en.wikipedia.org/wiki/Behavioral_analytics)

[http://en.wikipedia.org/wiki/Market\\_basket\\_analysis](http://en.wikipedia.org/wiki/Market_basket_analysis)

# Amazon Redshift Customers at re:Invent

## DAT306 - How Amazon.com is Leveraging Amazon Redshift

Thursday 11/14: 3pm in Murano 3303



## DAT205 - Amazon Redshift in Action: Enterprise, Big Data, SaaS Use Cases

Friday 11/15: 9am in Lido 3006

**NASDAQ OMX**

**roundarchisobar**

**HAUTELOOK**  
A NORDSTROM COMPANY



# AWS re:Invent

Please give us your feedback on this presentation

DAT305

As a thank you, we will select prize winners daily for completed surveys!

Thank You