

# Data Pipeline at Tapad

@tobym

@TapadEng



# Who am I?

Toby Matejovsky


First engineer hired at Tapad 3+ years  
ago

Scala developer

@tobym



# What are we talking about?



# Outline

- What Tapad does
- Why bother with a data pipeline?
- Evolution of the pipeline
- Day in the life of a analytics pixel
- What's next



# What Tapad Does

Cross-platform advertising and analytics  
Process billions of events per day

A Unified View.  
The Tapad Difference.




# Cross platform?

Device Graph

Node=device  
edge=inferred connection

Billion devices  
Quarter billion edges

85+% accuracy



# Why a Data Pipeline?

- Graph building

- Sanity while processing big data

- Decouple components

- Data accessible at multiple stages



# Graph Building

Realtime mode, but don't impact bidding latency

Batch mode





# Sanity

Billions of events, terabytes of logs per day

Don't have NSA's budget

Clear data retention policy

Store aggregations



# Decouple Components

Bidder only bids, graph-building  
process only builds graph

Data stream can split and merge



# Data accessible at multiple stages

Logs on edge of system

Local spool of data

Kafka broker

Consumer local spool

HDFS



# Evolution of the Data Pipeline

**Dark Ages:** Monolithic process, synchronous process

**Renaissance:** Queues, asynchronous work in same process

**Age of Exploration:** Inter-process comm, ad hoc batching

**Age of Enlightenment:** Standardize on Kafka and Avro



# Dark Ages

**Monolithic process, synchronous process**

It was fast enough, and we had to start somewhere.



# Renaissance

**Queues, asynchronous work in same process**

No, it wasn't fast enough.



# Age of Exploration

**Inter-process communication, ad hoc batching**

Servers at the edge batch up events, ship them to another service.



# Age of Enlightenment

**Standardize on Kafka and Avro**

Properly engineered and supported, reliable





# Age of Enlightenment

**Standardize on Kafka and Avro**

Properly engineered and supported, reliable



# Tangent!

Batching, queues, and serialization



# Batching

**Batching** is great, will really help throughput

**Batching != slow**



# Queues

**Queues** are amazing, until they explode and destroy the Rube Goldberg machine.

“I’ll just increase the buffer size.”

- spoken one day before someone ended up on double PagerDuty rotation

# Care and feeding of your queue

Monitor

Back-pressure

Buffering

Spooling

Degraded mode



# Serialization - Protocol Buffers

Tagged fields

Sort of self-describing

required, optional, repeated fields in schema

“Map” type:

```
message StringPair {  
    required string key = 1;  
    optional string value = 2;  
}
```

# Serialization - Avro

Optional field: `union { null, long } user_timestamp = null;`

Splittable (Hadoop world)

Schema evolution and storage



**pixel server** - pixel ingress - kafka - consumer - hdfs - hadoop jobs

## Day in the life of a pixel

Browser loads pixel from pixel server

Pixel server immediately responds with 200 and transparent gif, then serializes requests into a batch file

Batch file ships every few seconds or when the file reaches 2K





pixel server - **pixel ingress** - kafka - consumer - hdfs - hadoop jobs

## Day in the life of a pixel

Pixel ingress server receives 2 kilobyte file containing serialized web requests.

Deserialize, process some requests immediately (update database), then convert into Avro records with schema hash header, and publish to various Kafka topics



pixel server - pixel ingress - **kafka** - consumer - hdfs - hadoop jobs

## Day in the life of a pixel

Producer client figures out where to publish via the broker they connect to

Kafka topics are partitioned into multiple chunks, each has a master and slave and are on different servers to survive an outage.

Configurable retention based on time

Can add topics dynamically



pixel server - pixel ingress - kafka - **consumer** - hdfs - hadoop jobs

# Day in the life of a pixel

Consumer processes are organized into groups

Many consumer groups can read from same Kafka topic

Plugins:

```
trait Plugin[A] {  
  def onStartup(): Unit  
  def onSuccess(a: A): Unit  
  def onFailure(a: A): Unit  
  def onShutdown(): Unit  
}
```

GraphitePlugin, BatchingLogfilePlaybackPlugin, TimestampDrivenClockPlugin,  
BatchingTimestampDrivenClockPlugin, ...



pixel server - pixel ingress - kafka - **consumer** - hdfs - hadoop jobs

## Day in the life of a pixel

```
trait Plugins[A] {  
  
  private val _plugins = ArrayBuffer.empty[Plugin[A]]  
  
  def plugins: Seq[Plugin[A]] = _plugins  
  
  def registerPlugin(plugin: Plugin[A]) = _plugins += plugin  
}
```



pixel server - pixel ingress - kafka - **consumer** - hdfs - hadoop jobs

## Day in the life of a pixel

```
object KafkaConsumer {  
  sealed trait Result {  
    def notify[A](plugins: Seq[Plugin[A]], a: A): Unit  
  }  
  
  case object Success extends Result {  
    def notify[A](plugins: Seq[Plugin[A]], a: A) {  
      plugins.foreach(_.onSuccess(a))  
    }  
  }  
}
```



## pixel server - pixel ingress - kafka - **consumer** - hdfs - hadoop jobs

*/\*\* Decorate a Function1[A, B] with retry logic \*/*

```
case class Retry[A, B](maxAttempts: Int, backoff: Long)(f: A => B){  
  def apply(a: A): Result[A, B] = {  
    def execute(attempt: Int, errorLog: List[Throwable]): Result[A, B] = {  
      val result = try {  
        Success(this, a, f(a))  
      } catch {  
        ... Failure(this, a, e :: errorLog) ...  
      }  
  
      result match {  
        case failure @ Failure(_, _, errorLog) if errorLog.size < maxAttempts =>  
          val _backoff = (math.pow(2, attempt) * backoff).toLong  
          Thread.sleep(_backoff) // wait before the next invocation  
          execute(attempt + 1, errorLog) // try again  
        case failure @ Failure(_, _, errorLog) =>  
          failure  
      }  
    }  
    execute(attempt = 0, errorLog = Nil)  
  }  
}
```

pixel server - pixel ingress - kafka - consumer - **hdfs** - hadoop jobs

## Day in the life of a pixel

Consumers log into “permanent storage” in HDFS.

File format is Avro, written in batches.

Data retention policy is essential.



pixel server - pixel ingress - kafka - consumer - hdfs - **hadoop jobs**

## Day in the life of a pixel

Hadoop 2 - YARN

Scalding to write map-reduce jobs easily

Rewrite Avro files as Parquet

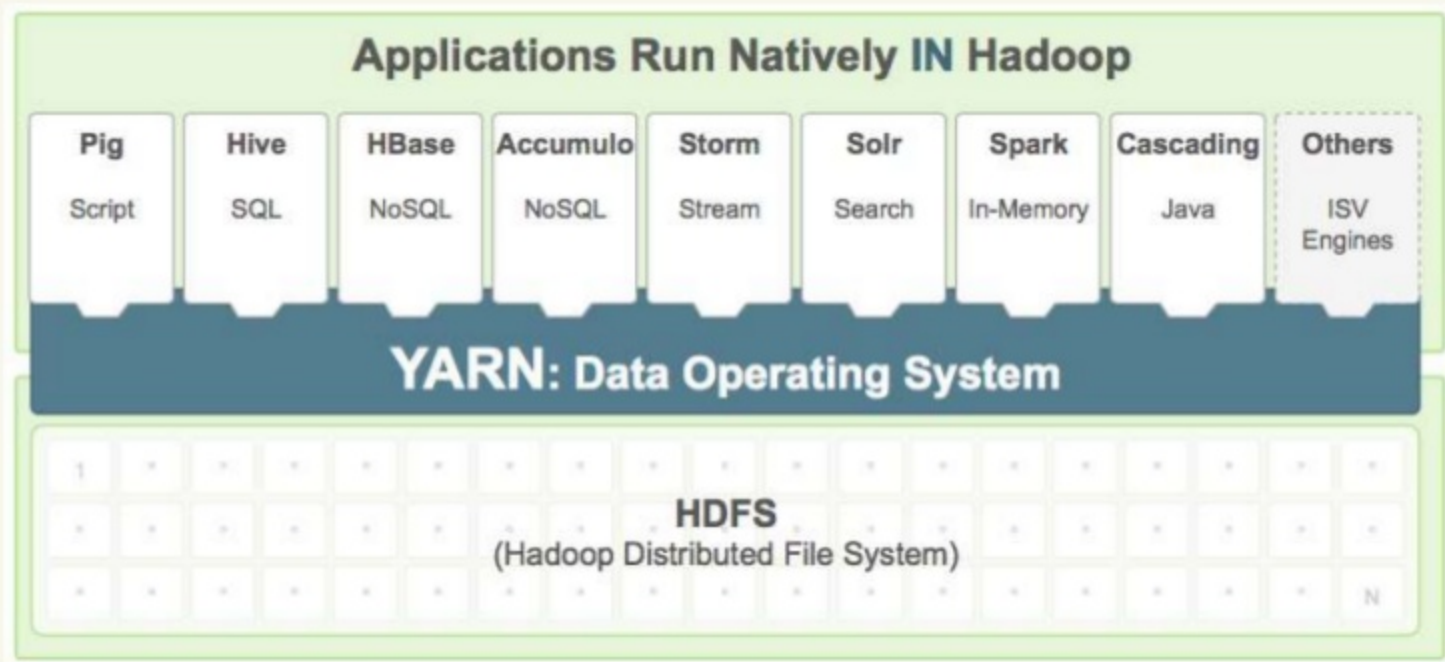
Oozie to schedule regular jobs





pixel server - pixel ingress - kafka - consumer - hdfs - **hadoop jobs**

# YARN



pixel server - pixel ingress - kafka - consumer - hdfs - **hadoop jobs**

# Scalding

```
class WordCountJob(args : Args) extends Job(args) {  
  TextLine( args("input") )  
    .flatMap('line -> 'word) { line : String => tokenize(line) }  
    .groupBy('word) { _.size }  
    .write( Tsv( args("output") ) )  
  
  // Split a piece of text into individual words.  
  def tokenize(text : String) : Array[String] = {  
    // Lowercase each word and remove punctuation.  
    text.toLowerCase.replaceAll("[^a-zA-Z0-9\\s]", "").split("\\s+")  
  }  
}
```



pixel server - pixel ingress - kafka - consumer - hdfs - **hadoop jobs**

# Parquet

Column-oriented storage for Hadoop

Nested data is okay

Projections

Predicates



pixel server - pixel ingress - kafka - consumer - hdfs - **hadoop jobs**

# Parquet

```
val requests = ParquetAvroSource
  .project[Request](args("requests"), Projection[Request]("header.query_params", "partner_id"))
  .read
  .sample(args("sample-rate").toDouble)
  .mapTo('Request -> ('queryParams, 'partnerId)) { req: TapestryRequest =>
    (req.getHeader.getQueryParams, req.getPartnerId)
  }
```



pixel server - pixel ingress - kafka - consumer - hdfs - **hadoop jobs**

# Oozie

```
<workflow-app name="combined_queries" xmlns="uri:oozie:workflow:0.3">
  <start to="devices-location"/>
  <!--<start to="export2db"/>-->

  <action name="devices-location">
    <shell xmlns="uri:oozie:shell-action:0.1">
      <job-tracker>${jobTracker}</job-tracker>
      <name-node>${nameNode}</name-node>

      <exec>hadoop</exec>
      <argument>fs</argument>
      <argument>-cat</argument>
      <argument>${devicesConfig}</argument>

      <capture-output/>
    </shell>

    <ok to="networks-location"/>
    <error to="kill"/>
  </action>
```



pixel server - pixel ingress - kafka - **consumer** - hdfs - **hadoop jobs**

## Day in the life of a pixel

Near real-time consumers and batch hadoop jobs generate data cubes from incoming events and save those aggregations into Vertica for fast and easy querying with SQL.



# Stack summary

Scala, Jetty/Netty, Finagle

Avro, Protocol Buffers, Parquet

Kafka

Zookeeper

Hadoop - YARN and HDFS

Vertica

Scalding

Oozie, Sqoop



# What's next?

Hive

Druid

Impala

Oozie alternative





A person is standing in a room with wood-paneled walls and floors. They are positioned next to a large window with a grid pattern, looking out. The person is wearing a dark long-sleeved shirt and dark pants. The overall lighting is warm and slightly dim.

# Thank You

@tobym

@TapadEng

yes, we're hiring! :)

Toby Matejovsky, Director of Engineering

[toby@tapad.com](mailto:toby@tapad.com)

@tobym

**TAPAD**