



Pop-up Loft

# Data warehousing in the era of Big Data: Intro to Amazon Redshift

Shree Kenghe, Solution Architect, AWS

# Agenda

- Introduction
- Benefits
- Use cases
- Getting started
- Q&A

# AWS big data portfolio

## Collect



AWS Direct Connect



AWS Import/Export



Amazon Kinesis Streams



Amazon Kinesis Firehose



AWS Database Migration Service

## Store



Amazon S3



Amazon RDS,  
Amazon Aurora



Amazon Glacier



Amazon DynamoDB



Amazon CloudSearch



Amazon Elasticsearch Service



AWS Data Pipeline

## Analyze



Amazon EMR



Amazon EC2



Amazon Redshift



Amazon Machine Learning



Amazon QuickSight



Amazon  
Redshift

Relational data warehouse

Massively parallel; petabyte scale

Fully managed

HDD and SSD platforms

\$1,000/TB/year; starts at \$0.25/hour

*a lot faster  
a lot simpler  
a lot cheaper*

# The Amazon Redshift view of data warehousing



## Enterprise

10x cheaper

Easy to provision

Higher DBA productivity



## Big data

10x faster

No programming

Easily leverage BI tools,  
Hadoop, machine learning,  
streaming



## SaaS

Analysis inline with process  
flows

Pay as you go, grow as you  
need

Managed availability and  
disaster recovery

# Forrester Wave™ Enterprise Data Warehouse Q4 '15

FIGURE 2 The Forrester Wave™: Enterprise Data Warehouse, Q4 '15



The Forrester Wave™ is copyrighted by Forrester Research, Inc. Forrester and Forrester Wave™ are trademarks of Forrester Research, Inc. The Forrester Wave™ is a graphical representation of Forrester's call on a market and is plotted using a detailed spreadsheet with exposed scores, weightings, and comments. Forrester does not endorse any vendor, product, or service depicted in the Forrester Wave. Information is based on best available resources. Opinions reflect judgment at the time and are subject to change.

# Selected Amazon Redshift customers



BEACHMINT



NOKIA

foursquare

Pinterest



sling



latentview  
Actionable Insights • Accurate Decisions

NTT docomo

NASDAQ OMX



amazon

etix

scopely

has offers

imshealth  
INTELLIGENCE APPLIED

euclid



Sansan

Schumacher group

Albert  
Optimization technology

spuul

peak  
GAMES



Book my Show

VIVAKI

DataXu

MINICLIP



UMUC  
University of Maryland University College

# Amazon Redshift architecture

## Leader node

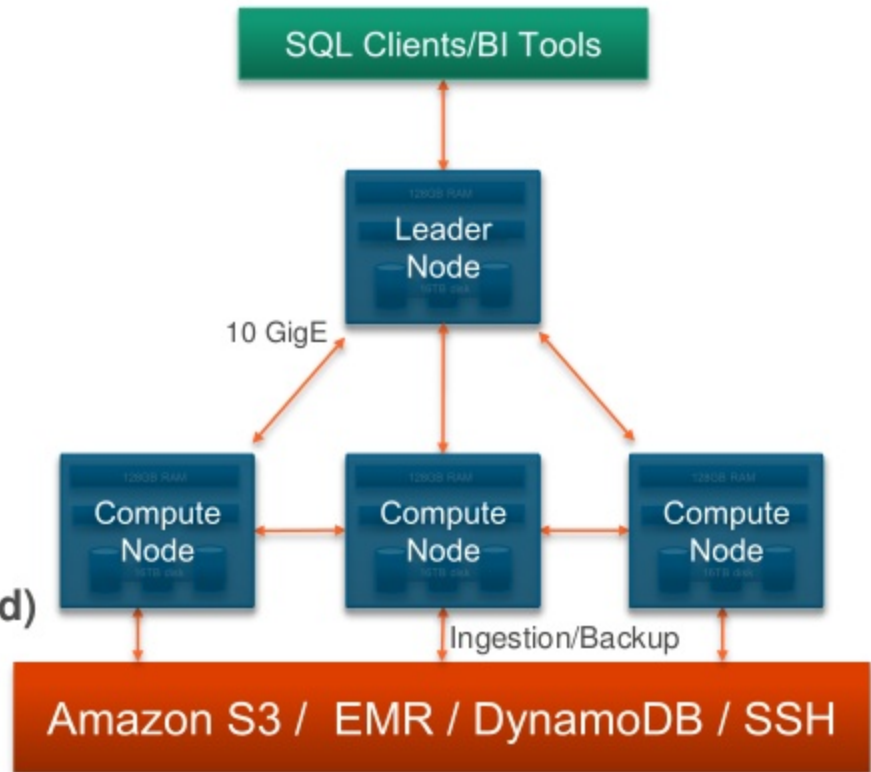
- Simple SQL endpoint
- Stores metadata
- Optimizes query plan
- Coordinates query execution

## Compute nodes

- Local columnar storage
- Parallel/distributed execution of all queries, loads, backups, restores, resizes

**Start at just \$0.25/hour, grow to 2 PB (compressed)**

- DC1: SSD; scale from 160 GB to 326 TB
- DS2: HDD; scale from 2 TB to 2 PB





# Benefit #1: Amazon Redshift is fast

## Dramatically less I/O

Column storage

Data compression


Zone maps

Direct-attached storage

Large data block sizes

```
analyze compression listing;
```

Table	Column	Encoding
listing	listid	delta
listing	sellerid	delta32k
listing	eventid	delta32k
listing	dateid	bytedict
listing	numtickets	bytedict
listing	priceperticket	delta32k
listing	totalprice	mostly32
listing	listtime	raw



# Benefit #1: Amazon Redshift is fast

## Parallel and distributed

Query

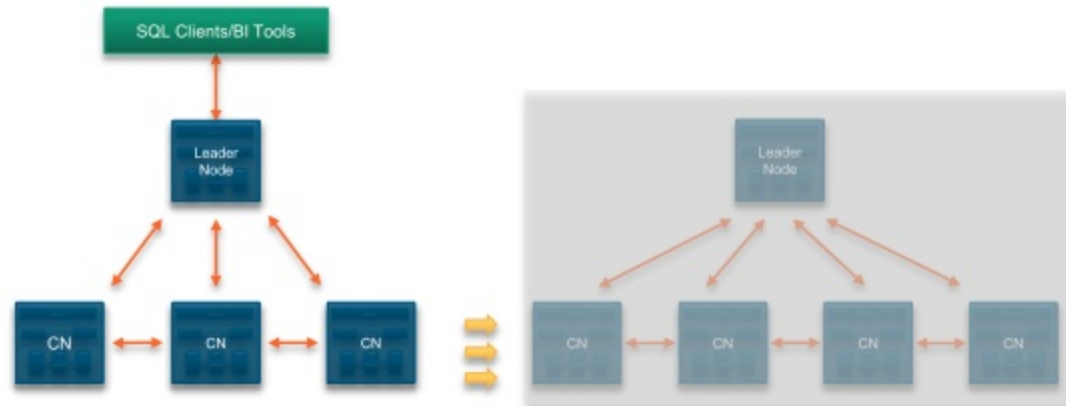

Load

Export

Backup

Restore

Resize




## Benefit #1: Amazon Redshift is fast

Hardware optimized for I/O intensive workloads, 4 GB/sec/node

Enhanced networking, over 1 million packets/sec/node

Choice of storage type, instance size

Regular cadence of auto-patched improvements



# Benefit #1: Amazon Redshift is fast

New Dense Storage (HDD) instance type (Jun 15)

Improved memory 2x, compute 2x, disk throughput 1.5x

Cost: Same as our prior generation!

Performance improvement: 50%


Enhanced I/O and commit improvements (Jan 16)

Reduce amount of time to commit data

Throughput performance improvement: 35%

Improved memory allocation for query processing (May 16)

Increased overall throughput by up to 60%



## Benefit #2: Amazon Redshift is inexpensive

DS2 (HDD)	Price per hour for DS2.XL single node	Effective annual price per TB compressed
On-demand	\$ 0.850	\$ 3,725
1 year reservation	\$ 0.500	\$ 2,190
3 year reservation	\$ 0.228	\$ 999

DC1 (SSD)	Price per hour for DC1.L single node	Effective annual price per TB compressed
On-demand	\$ 0.250	\$ 13,690
1 year reservation	\$ 0.161	\$ 8,795
3 year reservation	\$ 0.100	\$ 5,500

### Pricing is simple

Number of nodes x price/hour

No charge for leader node

No upfront costs

Pay as you go

# Benefit #3: Amazon Redshift is fully managed

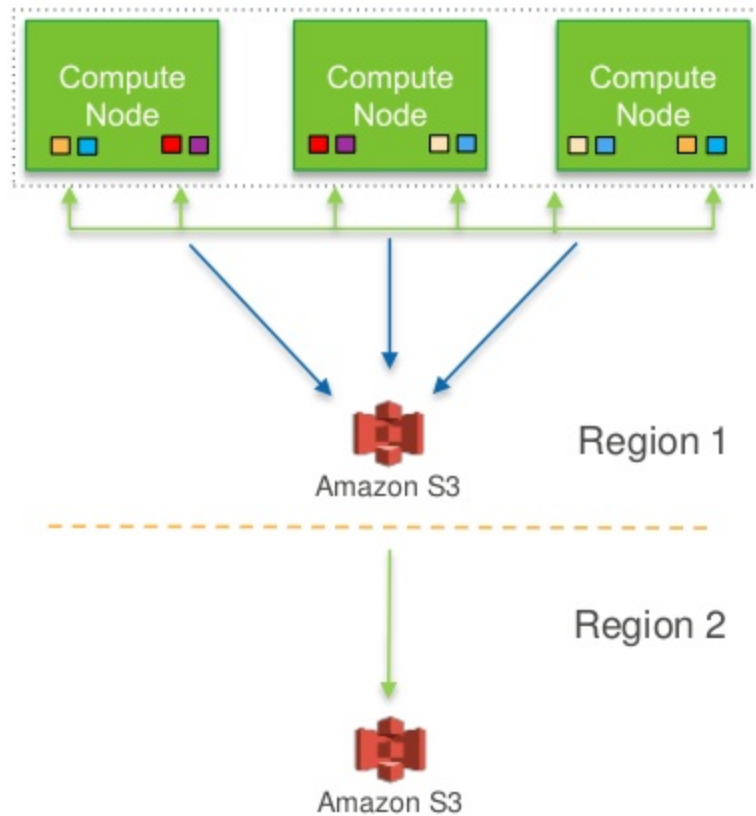
## Continuous/incremental backups

Multiple copies within cluster

Continuous and incremental backups to Amazon S3

Continuous and incremental backups across regions

Streaming restore



# Benefit #3: Amazon Redshift is fully managed


## Fault tolerance

Disk failures

Node failures

Network failures

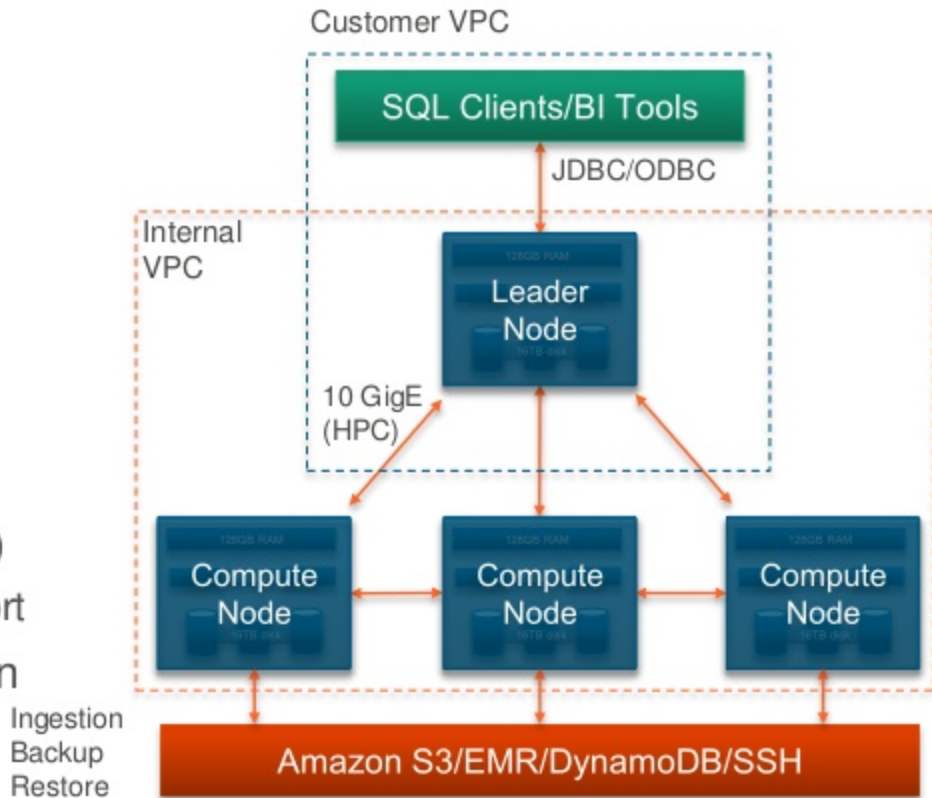
Availability Zone/region level disasters





# Benefit #4: Security is built-in

- Load encrypted from S3
- SSL to secure data in transit
  - ECDHE perfect forward security
- Amazon VPC for network isolation
- Encryption to secure data at rest
  - All blocks on disks and in S3 encrypted
  - Block key, cluster key, master key (AES-256)
  - On-premises HSM & AWS CloudHSM support
- Audit logging and AWS CloudTrail integration
- SOC 1/2/3, PCI-DSS, FedRAMP, BAA



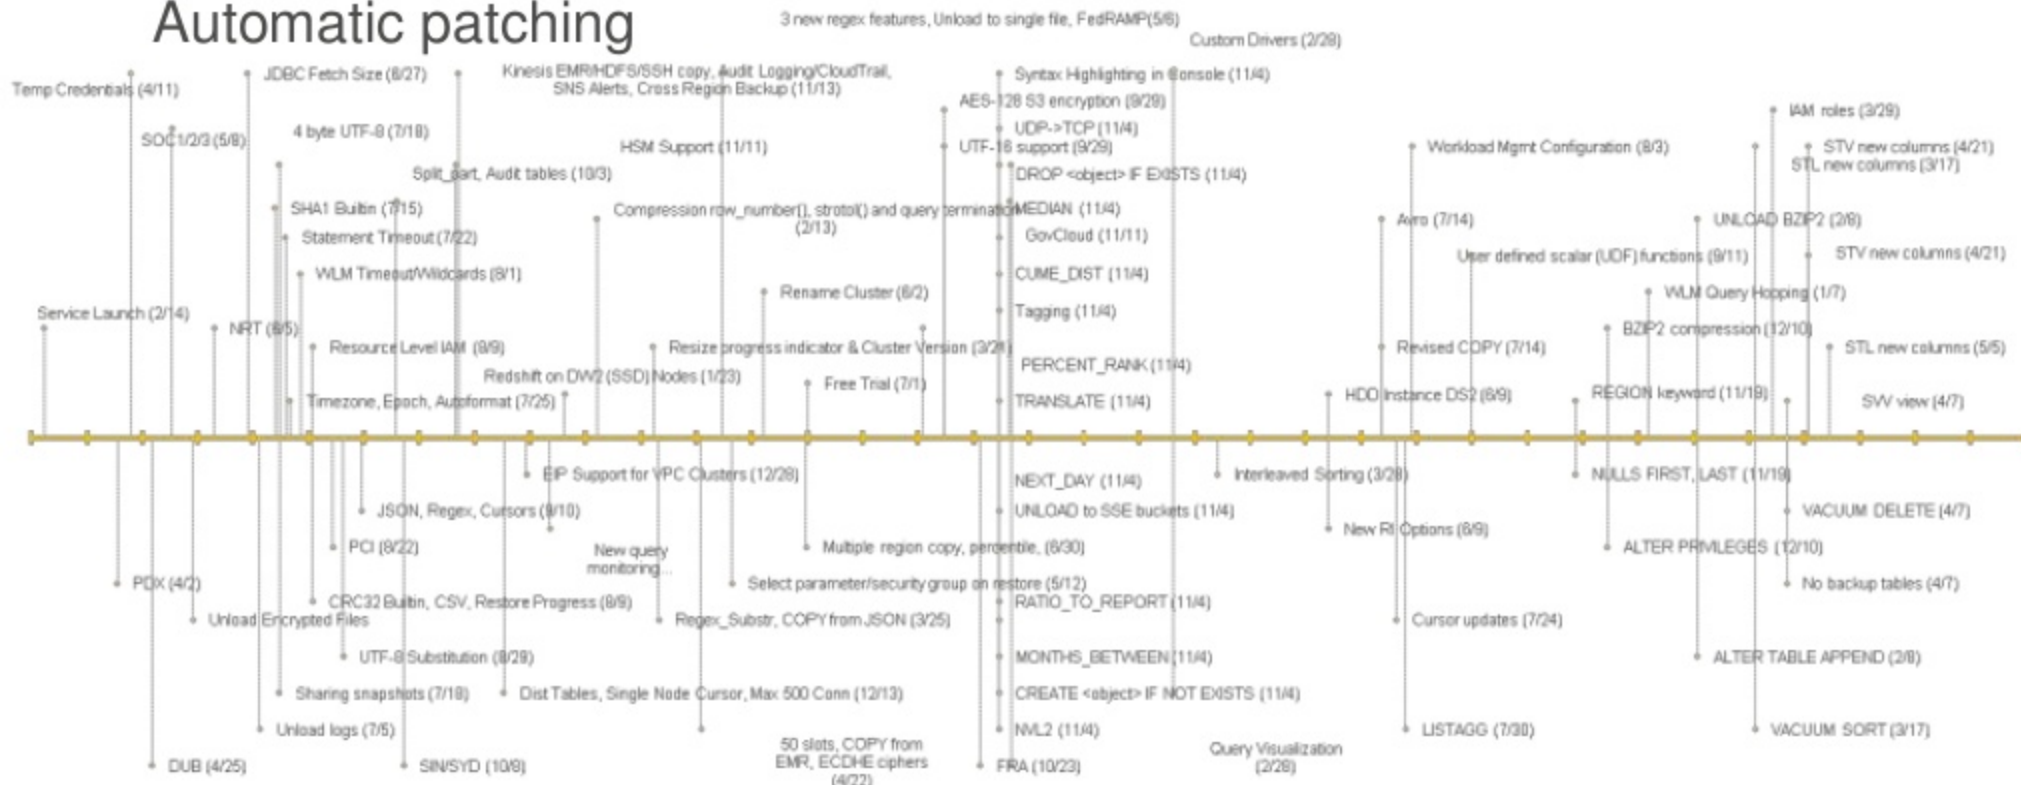


# Benefit #5: We innovate quickly

Well over 125 new features added since launch

Release every two weeks


Automatic patching



## Benefit #6: Redshift is powerful

- Approximate functions
- User defined functions
- Machine learning
- Data science

*HyperLogLog: analysis of a near-optimal cardinality algorithm*



# Benefit #7: Amazon Redshift has a large ecosystem

## Data integration




## Business intelligence



## Systems integrators



# Benefit #8: Service oriented architecture



# Use cases

# NTT Docomo: Japan's largest mobile service provider



68 million customers

Tens of TBs per day of data across a mobile network

6 PB of total data (uncompressed)

Data science for marketing operations, logistics, and so on

---

Greenplum on-premises

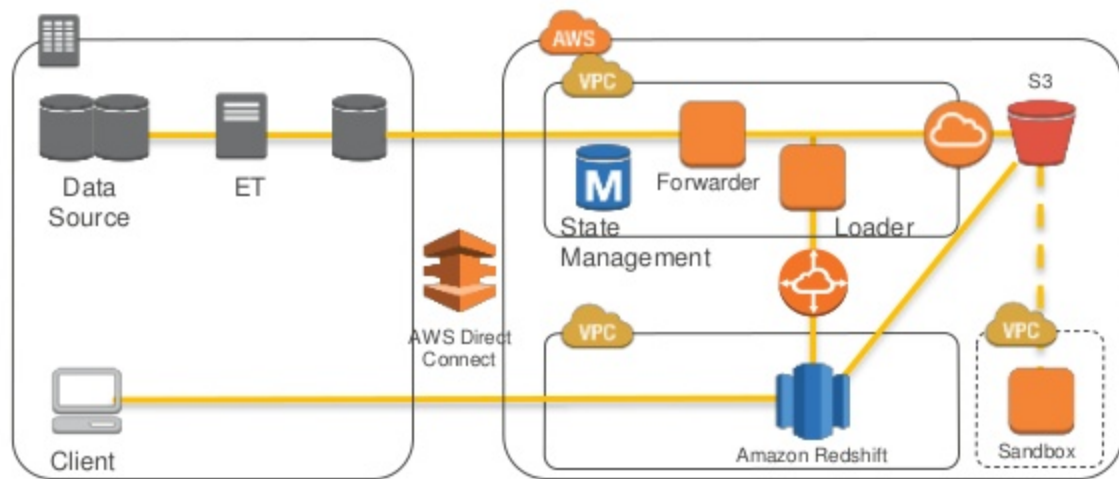
Scaling challenges

Performance issues

Need same level of security

Need for a hybrid environment

# NTT Docomo: Japan's largest mobile service provider



125 node DS2.8XL cluster  
4,500 vCPUs, 30 TB RAM  
2 PB compressed

10x faster analytic queries  
50% reduction in time for new  
BI application deployment  
Significantly less operations  
overhead



# Nasdaq: powering 100 marketplaces in 50 countries



Orders, quotes, trade executions,  
market “tick” data from 7 exchanges  
7 billion rows/day  
Analyze market share, client activity,  
surveillance, billing, and so on

---

Microsoft SQL Server on-premises

Expensive legacy DW  
(\$1.16 M/yr.)

Limited capacity (1 yr. of data  
online)

Needed lower TCO

Must satisfy multiple security  
and regulatory requirements

Similar performance



# Nasdaq: powering 100 marketplaces in 50 countries



23 node DS2.8XL cluster

828 vCPUs, 5 TB RAM

368 TB compressed

2.7 T rows, 900 B derived

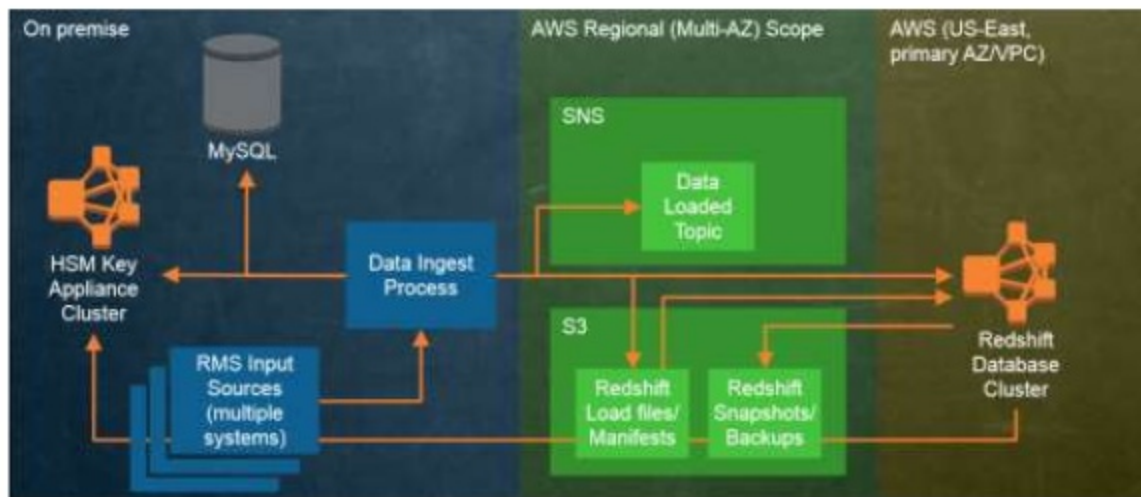
8 tables with 100 B rows

---

7 man-month migration

1/4 the cost, 2x storage, room to grow

Faster performance, very secure



# Getting started

# Provisioning

# Enter cluster details



Provide the details of your cluster. Fields marked with \* are required.

Cluster Identifier\*

This is the unique key that identifies a cluster.  
This parameter is stored as a lowercase string.  
(e.g. my-dw-instance)

Database Name

Optional. A default database named dev is created for the cluster. Optionally, specify a custom database name (e.g. mydb) to create an additional database.

Database Port\*

Port number on which the database accepts connections.

Master User Name\*

Name of master user for your cluster. (e.g. awsuser)

Master User Password\*

Password must contain 8 to 64 printable ASCII characters excluding: /, ", ', \, and @. It must contain 1 uppercase letter, 1 lowercase letter, and 1 number.

Confirm Password\*

Confirm Master User Password.

Cancel

Continue

# Select node configuration



Choose a number of nodes and Node Type below. Number of Compute Nodes is required for multi-node clusters.

Node Type **dc1.large** ▼

Specifies the compute, memory, storage, and I/O capacity of the cluster's nodes.

CPU 7 EC2 Compute Units (2 virtual cores) per node

Memory 15 GiB per node

Storage 160GB SSD storage per node

I/O Performance Moderate

Cluster Type **Single Node** ▼

Number of Compute Nodes\*

Single Node clusters consist of a single node which performs both leader and compute functions.

Maximum 1

Minimum 1

Cancel

Previous

Continue

# Select security settings and provision

Provide the optional additional configuration details below.

Cluster Parameter Group: **default-redshift-1.0** Parameter group to associate with this cluster.

Encrypt Database: ☐ None ☒ KMS ☐ HSM [Learn more about database encryption](#)

Master Key: (default) **aws/redshift**

Description: Default master key that protects my Redshift clusters when no other key is defined

Account: This account (052854472383)

KMS Key ID: alias/aws/redshift

Configure Networking Options:

Choose a VPC: **vpc-0544feb1** The identifier of the VPC in which you want to create your cluster

Cluster Subnet Group: **ppdsubnetgrp** Selected Cluster Subnet Group may limit the choice of Availability Zones

Publicly Accessible: **Yes** Select Yes if you want the cluster to be accessible from the public internet. Select No if you want it to be accessible only from within your private VPC network.

Choose a Public IP Address: **No** Select Yes if you want to select your own public IP address from a list of elastic IP (EIP) addresses that are already configured for your cluster's VPC. Select No if you want Amazon Redshift to provide an EIP for you instead.

Availability Zone: **No Preference** The EC2 Availability Zone that the cluster will be created in.

Optionally, associate your cluster with one or more security groups.

VPC Security Groups: **appservers-prodapp1 (sg-467e8129)**  
**rdc-dbservers-prodapp1 (sg-79cd3c16)**  
**quick-start-1 (sg-1e7e9171)**  
**quick-start-2 (sg-48e7b824)** List of VPC Security Groups to associate with this cluster.

Optionally, create a basic alarm for this cluster.

Create CloudWatch Alarm: ☐ Yes ☒ No Create a CloudWatch alarm to monitor the disk usage of your cluster.

Cancel

Previous

Continue



Cluster **pptest** is being created.

Note: Your cluster may take a few minutes to launch.

[View your cluster on the Clusters dashboard.](#)

# Point-and-click resize



### Resize Cluster

Choose the number of nodes and optionally a new node type for the resize operation. Note that the available node type and cluster type options may be limited by the cluster's current availability zone.

**Node Type:**

**Cluster Type:**

**Number Of Nodes:\***

Note: Resizing the cluster will cause it to be restarted into read-only mode for the duration of the resize operation. All currently executing queries and database connections on the cluster will be terminated when the resize operation begins and again when it is complete.

# Resize



- Resize while remaining online
- Provision a new cluster in the background
- Copy data in parallel from node to node
- Only charged for source cluster



# Data modeling

# Zone maps

SELECT COUNT(\*) FROM LOGS WHERE DATE = '09-JUNE-2013'

## Unsorted table



## Sorted by date



## Sort Keys

- Single column
- Compound
- Interleaved

# Single Column

[ SORTKEY ( date ) ]

Date	Region	Country
2-JUN-2015	Oceania	New Zealand
2-JUN-2015	Asia	Singapore
2-JUN-2015	Africa	Zaire
2-JUN-2015	Asia	Hong Kong
3-JUN-2015	Europe	Germany
3-JUN-2015	Asia	Korea

- Best for:
  - Queries that use 1<sup>st</sup> column (i.e. *date*) as primary filter
  - Can speed up joins and group bys
  - Quickest to VACUUM

# Compound

- Table is sorted by 1<sup>st</sup> column , then 2<sup>nd</sup> column etc.  
[ SORTKEY COMPOUND ( date, region, country) ]

Date	Region	Country
2-JUN-2015	Africa	Zaire
2-JUN-2015	Asia	Korea
2-JUN-2015	Asia	Singapore
2-JUN-2015	Europe	Germany
3-JUN-2015	Asia	Hong Kong
3-JUN-2015	Asia	Korea

- Best for:
  - Queries that use 1<sup>st</sup> column as primary filter, then other cols
  - Can speed up joins and group bys
  - Slower to VACUUM

# Interleaved

- Equal weight is given to each column.

[ SORTKEY INTERLEAVED ( date, region, country) ]

Date	Region	Country
2-JUN-2015	Africa	Zaire
3-JUN-2015	Asia	Singapore
2-JUN-2015	Asia	Korea
2-JUN-2015	Europe	Germany
3-JUN-2015	Asia	Hong Kong
2-JUN-2015	Asia	Korea

- Best for:
  - Queries that use different columns in filter
  - Queries get faster the more columns used in the filter
  - Slowest to VACUUM

- EVEN
- KEY
- ALL

# Distribution

ID	Gender	Name
101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	James White
306	F	Lisa Green



ID	Gender	Name
101	M	John Smith
306	F	Lisa Green



ID	Gender	Name
292	F	Jane Jones
209	M	James White



ID	Gender	Name
139	M	Peter Black
164	M	Brian Snail



ID	Gender	Name
446	M	Pat Partridge
658	F	Sarah Cyan



# Distribution

ID	Gender	Name
101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	James White
306	F	Lisa Green

KEY



ID	Gender	Name
101	M	John Smith
306	F	Lisa Green



ID	Gender	Name
292	F	Jane Jones
209	M	James White



ID	Gender	Name
139	M	Peter Black
164	M	Brian Snail



ID	Gender	Name
446	M	Pat Partridge
658	F	Sarah Cyan

# Distribution

ID	Gender	Name
101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	James White
306	F	Lisa Green



ID	Gender	Name
101	M	John Smith
139	M	Peter Black
446	M	Pat Partridge
164	M	Brian Snail
209	M	James White



ID	Gender	Name
292	F	Jane Jones
658	F	Sarah Cyan
306	F	Lisa Green

# Distribution

ID	Gender	Name
101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	James White
306	F	Lisa Green



101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	Lisa Green
306	F	James White



101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	Lisa Green
306	F	James White



101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	Lisa Green
306	F	James White

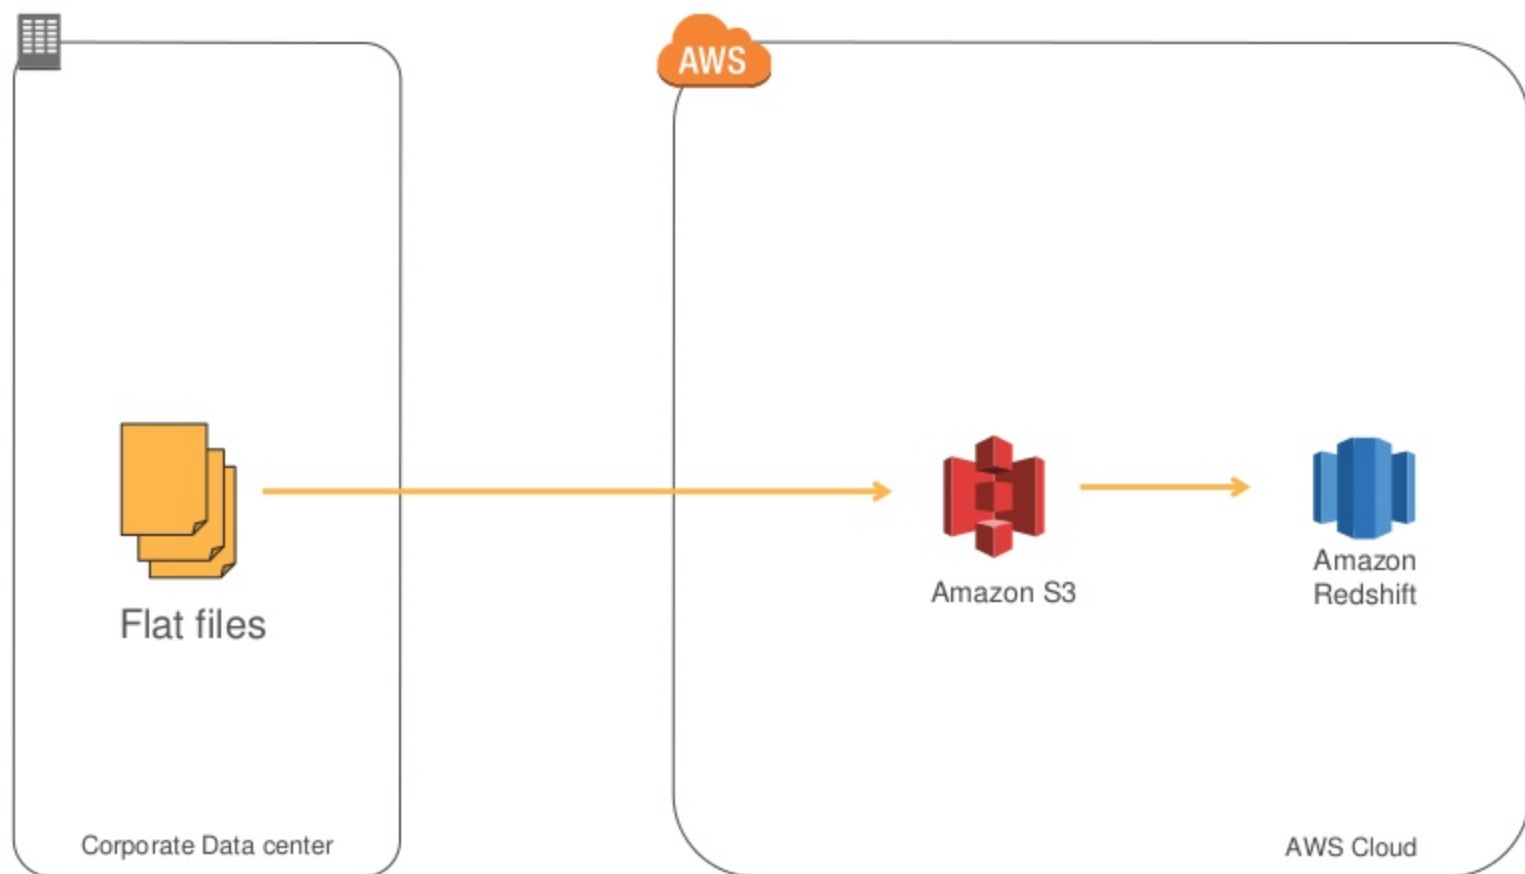


101	M	John Smith
292	F	Jane Jones
139	M	Peter Black
446	M	Pat Partridge
658	F	Sarah Cyan
164	M	Brian Snail
209	M	Lisa Green
306	F	James White

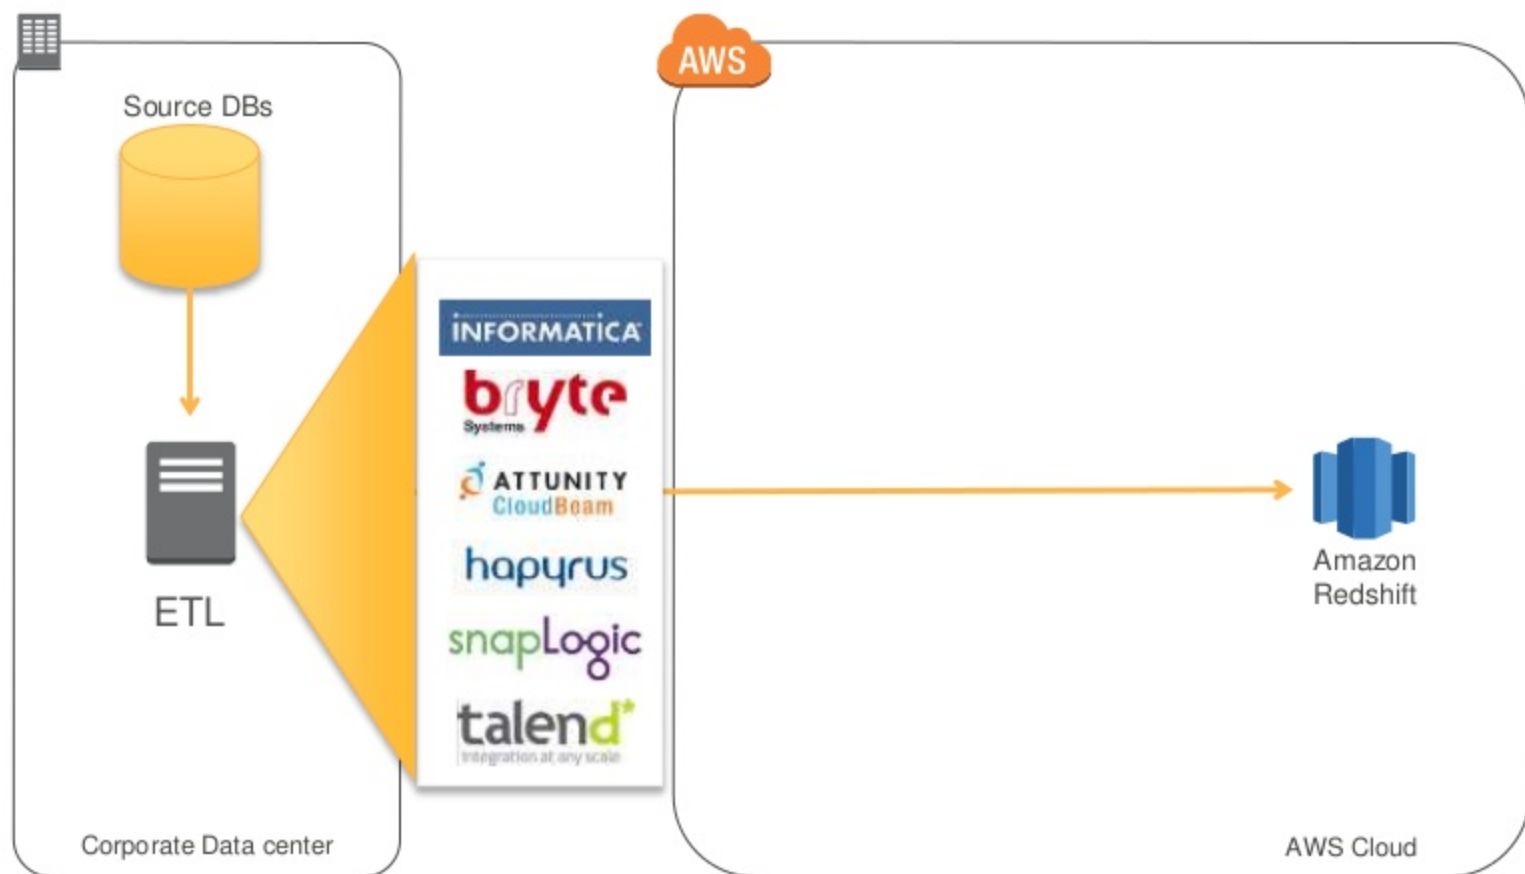
- EVEN
  - Tables with no joins or group by
- KEY
  - Large Fact tables
  - Large dimension tables
- ALL
  - Medium dimension tables (1K – 2M)
  - Small dimension tables

# Loading data


# Data loading options



# Data loading options



# Data loading options





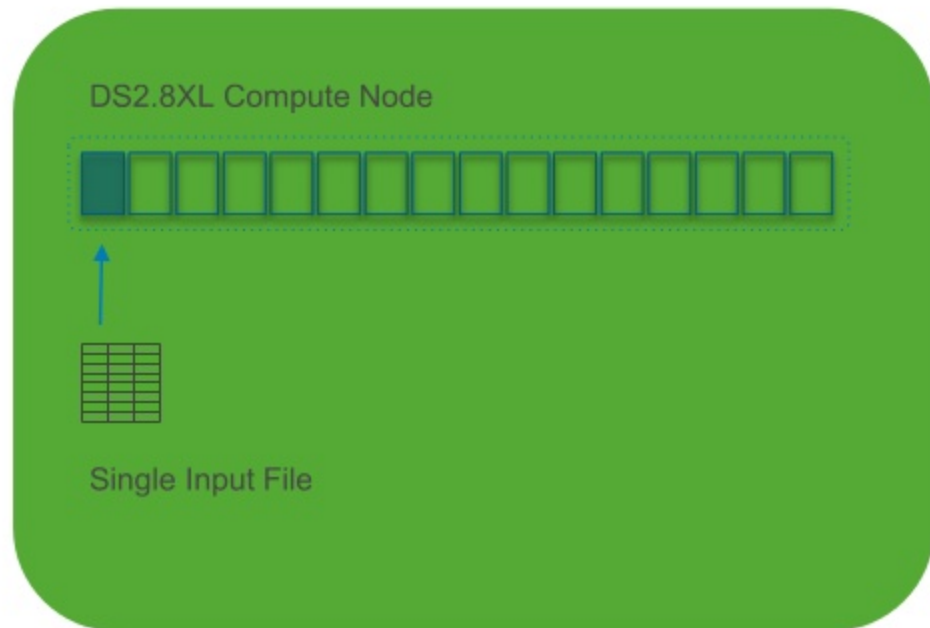
# Use multiple input files to maximize throughput

Use the COPY command

Each slice can load one file at a time

A single input file means only one slice is ingesting data

Instead of 100MB/s, you're only getting 6.25MB/s



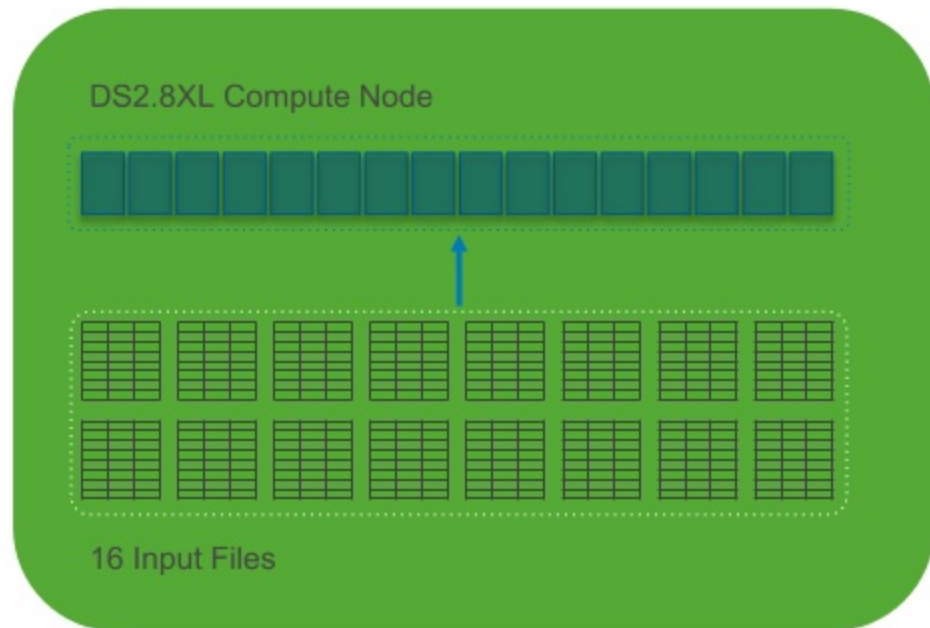
# Use multiple input files to maximize throughput

Use the COPY command

You need at least as many input files as you have slices

With 16 input files, all slices are working so you maximize throughput

Get 100MB/s per node; scale linearly as you add nodes



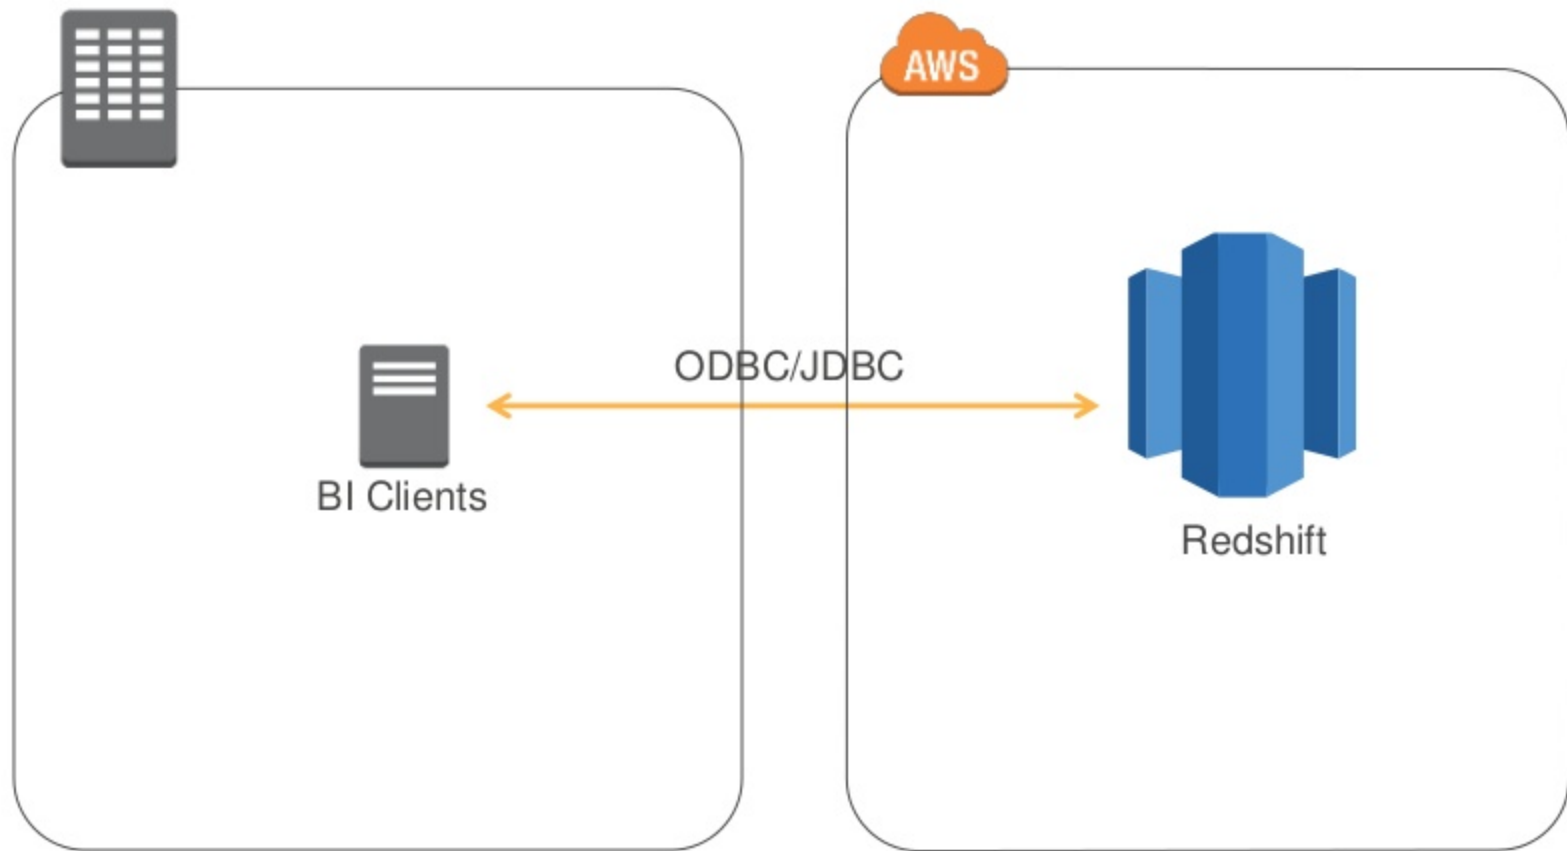
# Querying

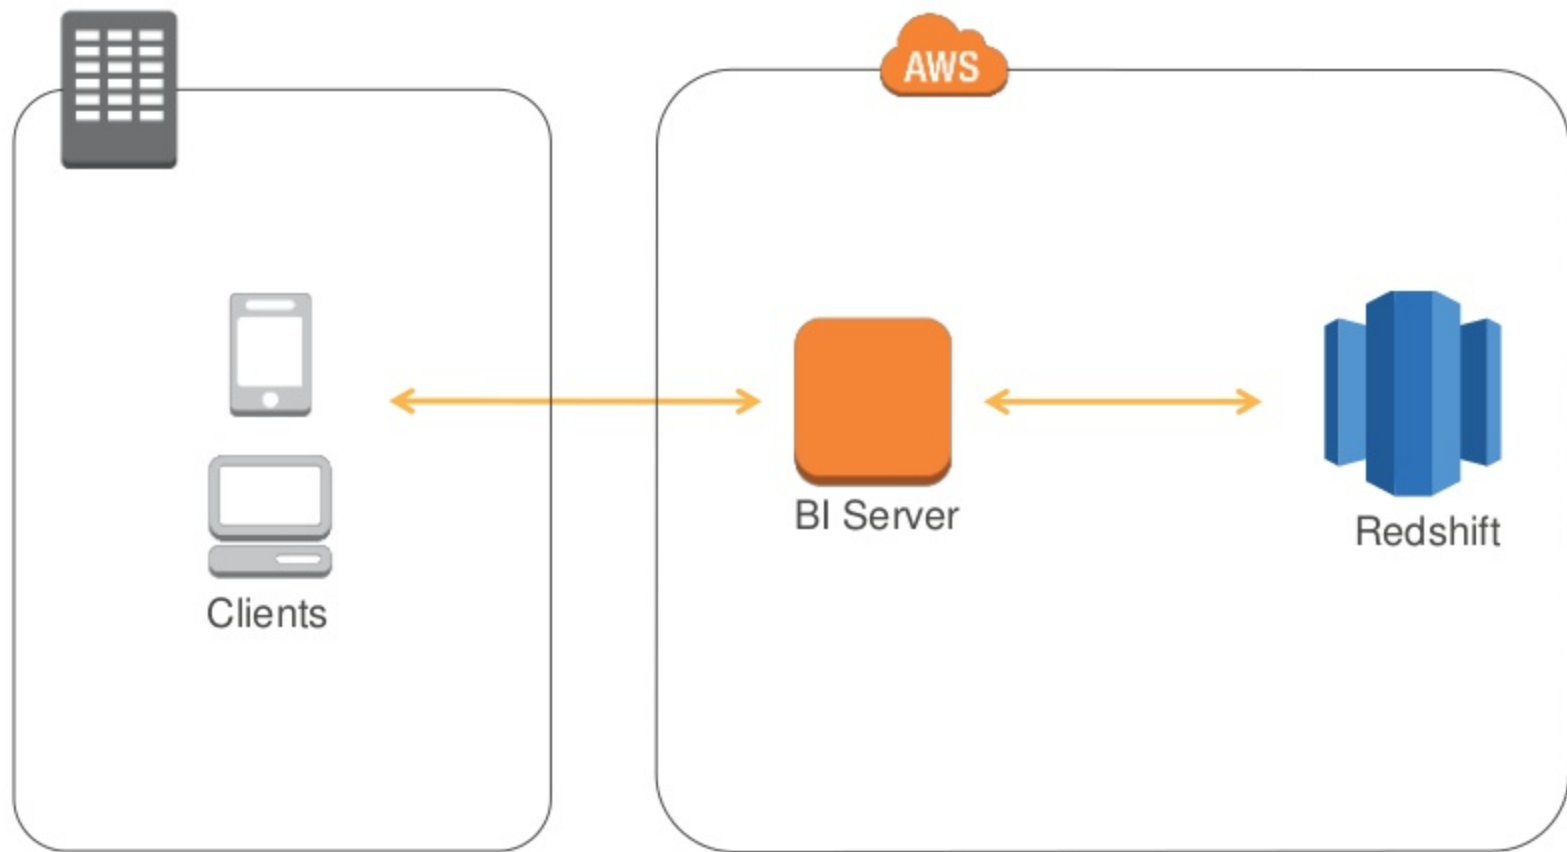
# Amazon Redshift works with your existing BI tools



JDBC/ODBC







## View explain plans

▼ Query Execution Details

Plan	Actual
------	--------

```

XN Hash Join DS_BCAST_INNER (cost=52602972108.91..68656434844.99 rows=23608 width=71)
-> XN Seq Scan on supplier s (cost=0.00..10000.00 rows=1000000 width=22)
-> XN Hash (cost=52602972049.89..52602972049.89 rows=23608 width=57)
    -> XN Hash Join DS_BCAST_INNER (cost=1752.83..52602972049.89 rows=23608 width=57)
        -> XN Hash Join DS_BCAST_INNER (cost=1718.72..52499650202.87 rows=70082 width=61)
            -> XN Hash (cost=31.95..31.95 rows=861 width=4)
                -> XN Hash Join DS_BCAST_INNER (cost=624.99..14000065191.80 rows=714354 width=57)
                    -> XN Hash (cost=874.99..874.99 rows=87499 width=33)
                        -> XN Seq Scan on lineorder lo (cost=0.00..149965.90 rows=14996590 width=19)
                            -> XN Hash (cost=499.99..499.99 rows=49999 width=20)
                                -> XN Seq Scan on part p (cost=0.00..499.99 rows=49999 width=2)

```

	Plan	Actual
Revenue	100,000	100,000
Variable costs	60,000	60,000
Contribution margin	40,000	40,000
Fixed costs	30,000	30,000
Operating income	10,000	10,000

Expand Level

Collapse Level

Expand All

- Hash Join (Broadcasted Inner Table)
  - Sequential Scan on supplier s
- Hash
  - Hash Join (Broadcasted Inner Table)
    - Hash Join (Broadcasted Inner Table)
      - Hash Join (Broadcasted Inner Table)
        - Sequential Scan on lineorder lo
      - Hash
        - Sequential Scan on part p
    - Hash
      - Sequential Scan on customer c
  - Hash
    - Sequential Scan on dwdate d



select

```
lo_orderkey,  
p_name,  
c_name,  
s_address,  
lo_quantity
```

from

```
lineorder lo,
part p,
supplier s,
customer c,
dwdate d
```

where

```
lo_custkey = c_custkey
and      lo_partkey = p_partkey
and      lo_suppkey = s_suppkey
and      lo_orderdate = d_datekey
and      d_sellingseason = 'Summer'
```


# Monitor query performance

Cluster: **pptest2** | Configuration | Status | Performance | Queries | Loads

Time Range: 06/16 17:53—06/16 18:53 UTC-7 | Period: 5 Minutes | Statistic: Average | All Metrics | Nodes | Refresh

## Queries

Hover over the queries graph or click on a query ID in the legend to inspect queries. Click and drag on any graph to zoom in.



## CPU Utilization



## NetworkReceiveThroughput



Query ID: 351

Type: Load

User: pavanpo

Run Time: 45.96s

Start Time: Tue Jun 16 18:43:38 GMT-700 2015

End Time: Tue Jun 16 18:44:17 GMT-700 2015

SQL:

```
copy lineorder from
's3://awsamplesh/load/lo/lineorder-
single.tbl' credentials '' gzip compute
off region 'us-east-1'
```



# Resources

## Detail Pages

- <http://aws.amazon.com/redshift>
- <https://aws.amazon.com/marketplace/redshift/>
- [Amazon Redshift Utilities - GitHub](#)

## Best Practices

- [http://docs.aws.amazon.com/redshift/latest/dg/c\\_loading-data-best-practices.html](http://docs.aws.amazon.com/redshift/latest/dg/c_loading-data-best-practices.html)
- [http://docs.aws.amazon.com/redshift/latest/dg/c\\_designing-tables-best-practices.html](http://docs.aws.amazon.com/redshift/latest/dg/c_designing-tables-best-practices.html)
- <http://docs.aws.amazon.com/redshift/latest/dg/c-optimizing-query-performance.html>



Pop-up Loft

**Thank you!**