



Real-time Streaming Data on AWS

Deep Dive & Best Practices Using Amazon Kinesis,
Spark Streaming, AWS Lambda, and Amazon EMR

Roy Ben-Alta, Sr. Business Development Manager, AWS
Orit Alul, Data and Analytics R&D Director, Sizmek

June-16-2016



Agenda

Real-time streaming overview

Use cases and design patterns

Amazon Kinesis deep dive

Streaming data ingestion & Stream processing

Sizmek Case Study

Q&A

It's All About the Pace

Batch Processing

Hourly server logs

Weekly or monthly bills

Daily web-site clickstream

Daily fraud reports

Stream Processing

Real-time metrics

Real-time spending alerts/caps

Real-time clickstream analysis

Real-time detection

Streaming Data Scenarios Across Verticals

Scenarios/ Verticals	Accelerated Ingest- Transform-Load	Continuous Metrics Generation	Responsive Data Analysis
Digital Ad Tech/Marketing	Publisher, bidder data aggregation	Advertising metrics like coverage, yield, and conversion	User engagement with ads, optimized bid/buy engines
IoT	Sensor, device telemetry data ingestion	Operational metrics and dashboards	Device operational intelligence and alerts
Gaming	Online data aggregation, e.g., top 10 players	Massively multiplayer online game (MMOG) live dashboard	Leader board generation, player-skill match
Consumer Online	Clickstream analytics	Metrics like impressions and page views	Recommendation engines, proactive care

Customer Use Cases

SONOS

Sonos runs near real-time streaming analytics on device data logs from their connected hi-fi audio equipment.



Glu Mobile collects billions of gaming events data points from millions of user devices in real-time every single day.

REDFIN.

One of the biggest online brokerages for real estate in US. Built Hot Homes feature.

NORDSTROM

Nordstrom recommendation team built online stylist using Amazon Kinesis Streams and AWS Lambda.

Streaming Data Challenges: Variety & Velocity

- Streaming data comes in different types and formats
 - Metering records, logs and sensor data
 - JSON, CSV, TSV
- Can vary in size from a few bytes to kilobytes or megabytes
- High velocity and continuous processing

```
{  
  "payerId": "Joe",  
  "productCode": "AmazonS3",  
  "clientProductCode": "AmazonS3",  
  "usageType": "Bandwidth",  
  "operation": "PUT",  
  "value": "22490",  
  "timestamp": "1216674828"  
}
```

Metering Record

```
{  
  127.0.0.1 user-  
  identifier frank  
  [10/Oct/2000:13:5  
  5:36 -0700] "GET  
  /apache_pb.gif  
  HTTP/1.0" 200  
  2326  
}
```

Common Log Entry

```
{  
  <165>1 2003-10-11T22:14:15.003Z  
  mymachine.example.com evntslog -  
  ID47 [exampleSDID@32473 iut="3"  
  eventSource="Application"  
  eventID="1011"] [examplePriority@  
  32473 class="high"]  
}
```

Syslog Entry

```
{  
  "SeattlePublicWa  
  ter/Kinesis/123/  
  Realtime" -  
  412309129140  
}
```

MQTT Record

Two Main Processing Patterns

Stream processing (real time)

- Real-time response to events in data streams

Examples:

- Proactively detect hardware errors in device logs
- Notify when inventory drops below a threshold
- Fraud detection

Micro-batching (near real time)

- Near real-time operations on small batches of events in data streams

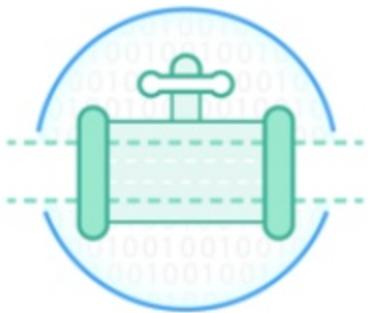
Examples:

- Aggregate and archive events
- Monitor performance SLAs

Amazon Kinesis Deep Dive

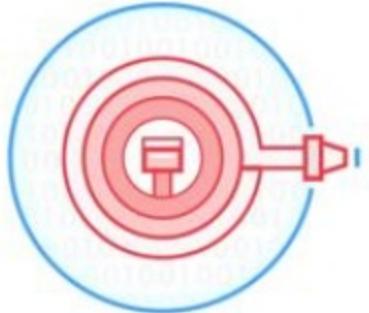
Amazon Kinesis: Streaming Data Made Easy

Services make it easy to capture, deliver and process streams on AWS



Amazon Kinesis Streams

- For Technical Developers
- Build your own custom applications that process or analyze streaming data



Amazon Kinesis Firehose

- For all developers, data scientists
- Easily load massive volumes of streaming data into S3, Amazon Redshift and Amazon Elasticsearch



Amazon Kinesis Analytics

- For all developers, data scientists
- Easily analyze data streams using standard SQL queries
- **Preview**

Amazon Kinesis Firehose

Load massive volumes of streaming data into Amazon S3, Amazon Redshift and Amazon Elasticsearch



*Capture and submit
streaming data to Firehose*

*Firehose loads streaming data
continuously into S3, Amazon Redshift
and Amazon Elasticsearch*

*Analyze streaming data using your
favorite BI tools*

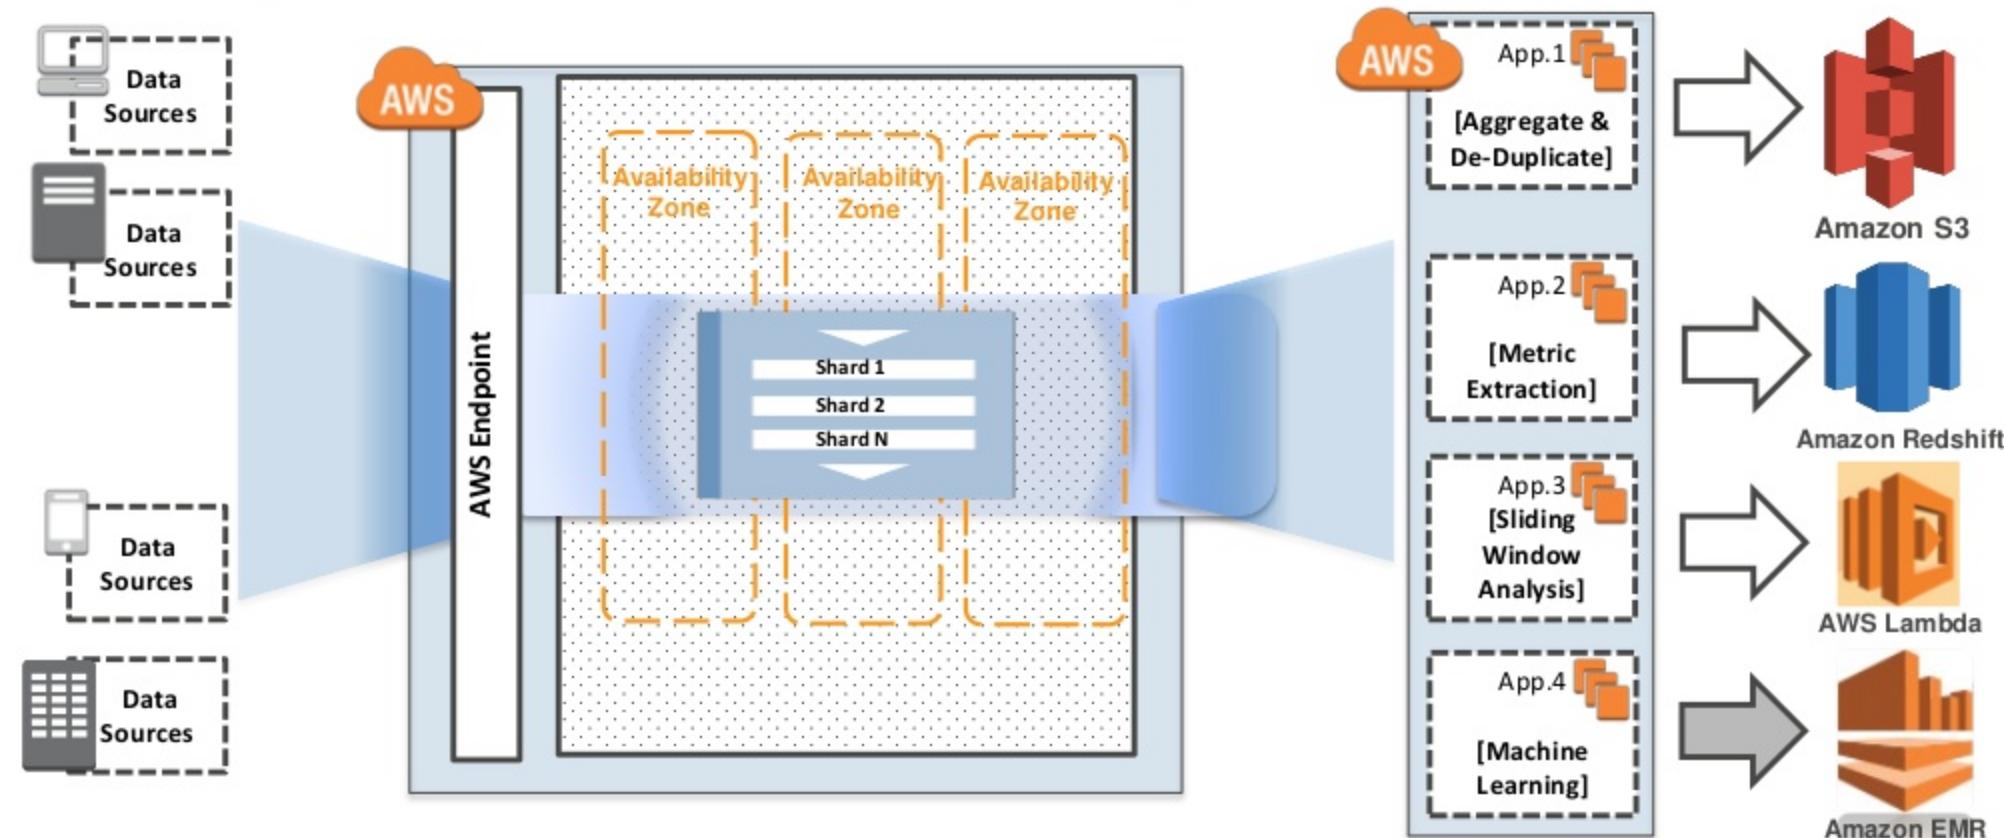
Zero administration: Capture and deliver streaming data into Amazon S3, Amazon Redshift and Amazon Elasticsearch **without writing an application or managing infrastructure.**

Direct-to-data store integration: **Batch, compress**, and **encrypt** streaming data for delivery into data destinations **in as little as 60 secs** using simple configurations.

Seamless elasticity: Seamlessly scales to match data throughput w/o intervention

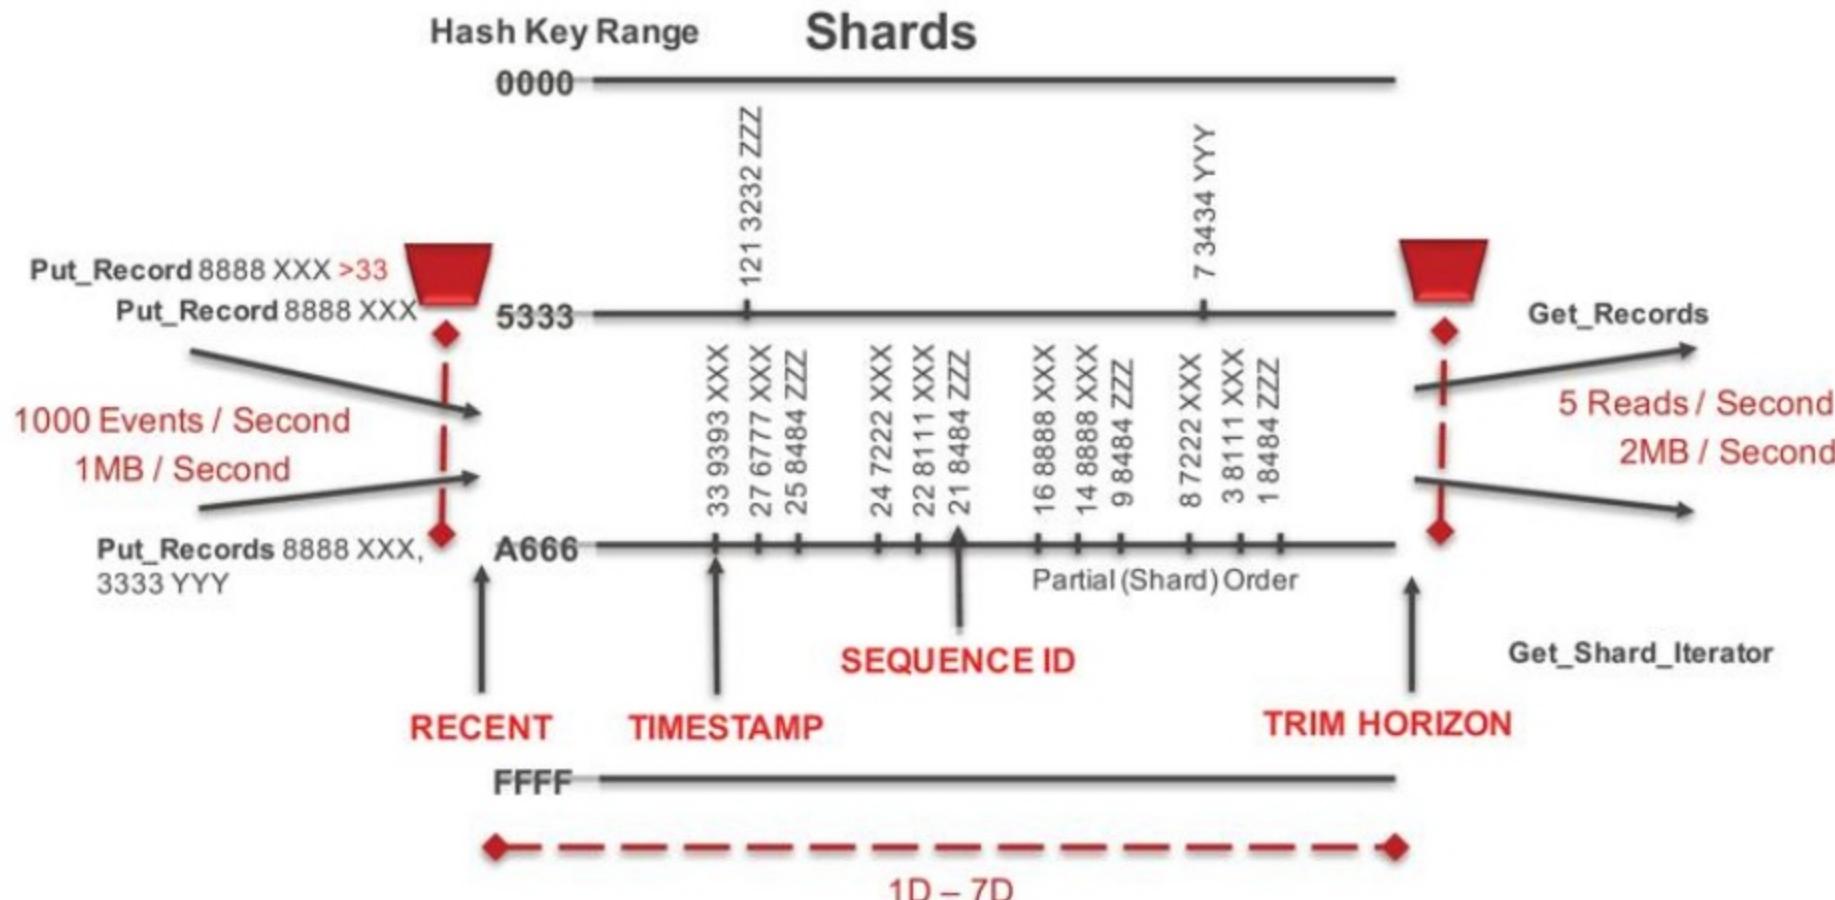
Amazon Kinesis Streams

Managed service for real-time streaming



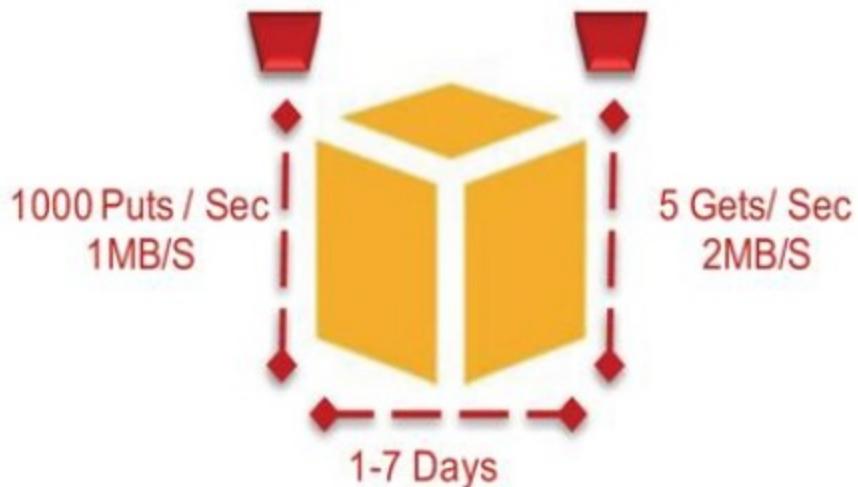
Amazon Kinesis Streams

Managed ability to capture and store data



Amazon Kinesis Streams

Managed ability to capture and store data



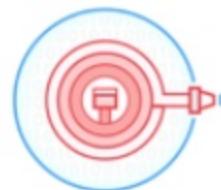
- Streams are made of **shards**
- Each shard ingests up to 1MB/sec, and 1000 records/sec
- Each shard emits up to 2 MB/sec
- All data is **stored for 24 hours by default**; storage can **be extended for up to 7 days**
- **Scale** Kinesis streams using scaling util
- **Replay** data

Amazon Kinesis Firehose vs. Amazon Kinesis Streams



Amazon Kinesis Streams

Amazon Kinesis Streams is for use cases that require **custom processing**, per incoming record, with sub-1 second processing latency, and a choice of stream processing frameworks.



Amazon Kinesis Firehose

Amazon Kinesis Firehose is for use cases that require zero administration, ability to **use existing analytics tools based on Amazon S3, Amazon Redshift and Amazon Elasticsearch**, and a data latency of 60 seconds or higher.

Streaming Data Ingestion and Stream Processing

Putting Data into Amazon Kinesis Streams

Determine your partition key strategy

- Managed buffer or streaming MapReduce job
- Ensure high cardinality for your shards

Provision adequate shards

- For ingress needs
- Egress needs for all consuming applications: if more than two simultaneous applications
- Include headroom for catching up with data in stream

Putting Data into Amazon Kinesis

Amazon Kinesis Agent – (supports pre-processing)

- <http://docs.aws.amazon.com/firehose/latest/dev/writing-with-agents.html>

Pre-batch before Puts for better efficiency

- Consider Flume, Fluentd as collectors/agents
- See <https://github.com/awslabs/aws-fluent-plugin-kinesis>

Make a tweak to your existing logging

- log4j appender option
- See <https://github.com/awslabs/kinesis-log4j-appender>

Amazon Kinesis Producer Library

- Writes to one or more Amazon Kinesis streams with automatic, configurable retry mechanism
- Collects records and uses PutRecords to write multiple records to multiple shards per request
- Aggregates user records to increase payload size and improve throughput
- Integrates seamlessly with KCL to de-aggregate batched records
- Use Amazon Kinesis Producer Library with AWS Lambda (**New!**)
- Submits Amazon CloudWatch metrics on your behalf to provide visibility into producer performance

Record Order and Multiple Shards

Unordered processing

- Randomize partition key to distribute events over many shards and use multiple workers

Exact order processing

- Control partition key to ensure events are grouped into the same shard and read by the same worker

Need both? Use global sequence number



Sample Code for Scaling Shards

```
java -cp  
KinesisScalingUtils.jar-complete.jar  
-Dstream-name=MyStream  
-Dscaling-action=scaleUp  
-Dcount=10  
-Dregion=eu-west-1 ScalingClient
```

Options:

- **stream-name** - The name of the stream to be scaled
- **scaling-action** - The action to be taken to scale. Must be one of "scaleUp", "scaleDown" or "resize"
- **count** - Number of shards by which to absolutely scale up or down, or resize

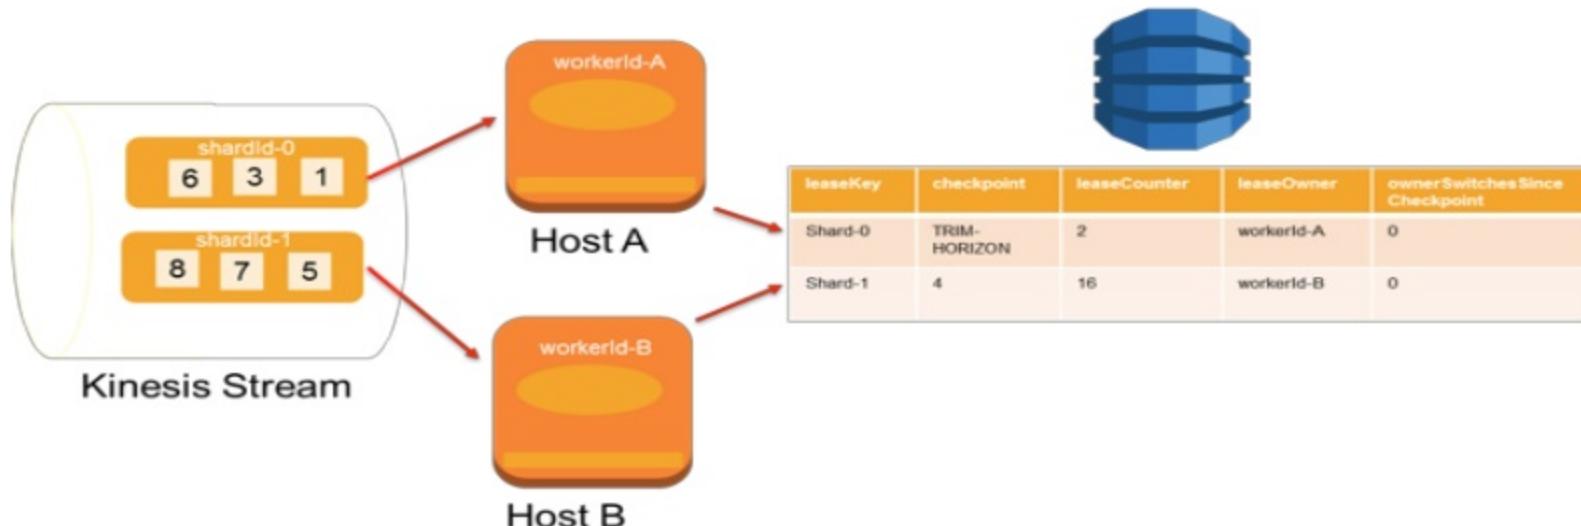
See <https://github.com/awslabs/amazon-kinesis-scaling-utils>

Amazon Kinesis Client Library

- Build Kinesis Applications with Kinesis Client Library (KCL)
- Open source client library available for Java, Ruby, Python, Node.JS dev
- Deploy on your EC2 instances
- KCL Application includes three components:
 1. **Record Processor Factory** – Creates the record processor
 2. **Record Processor** – Processor unit that processes data from a shard in Amazon Kinesis Streams
 3. **Worker** – Processing unit that maps to each application instance

State Management with Kinesis Client Library

- One record processor maps to one shard and processes data records from that shard
- One worker maps to one or more record processors
- Balances shard-worker associations when worker / instance counts change
- Balances shard-worker associations when shards split or merge



Other Options

- Third-party connectors(for example, Apache Storm, Splunk and more)
- AWS IoT platform
- **Amazon EMR with Apache Spark, Pig or Hive**
- **AWS Lambda**

Apache Spark and Amazon Kinesis Streams

Apache Spark is an in-memory analytics cluster using RDD for fast processing

Spark Streaming can read directly from an Amazon Kinesis stream

Amazon software license linking – Add ASL dependency to SBT/MAVEN project, `artifactId = spark-streaming-kinesis-asl_2.10`

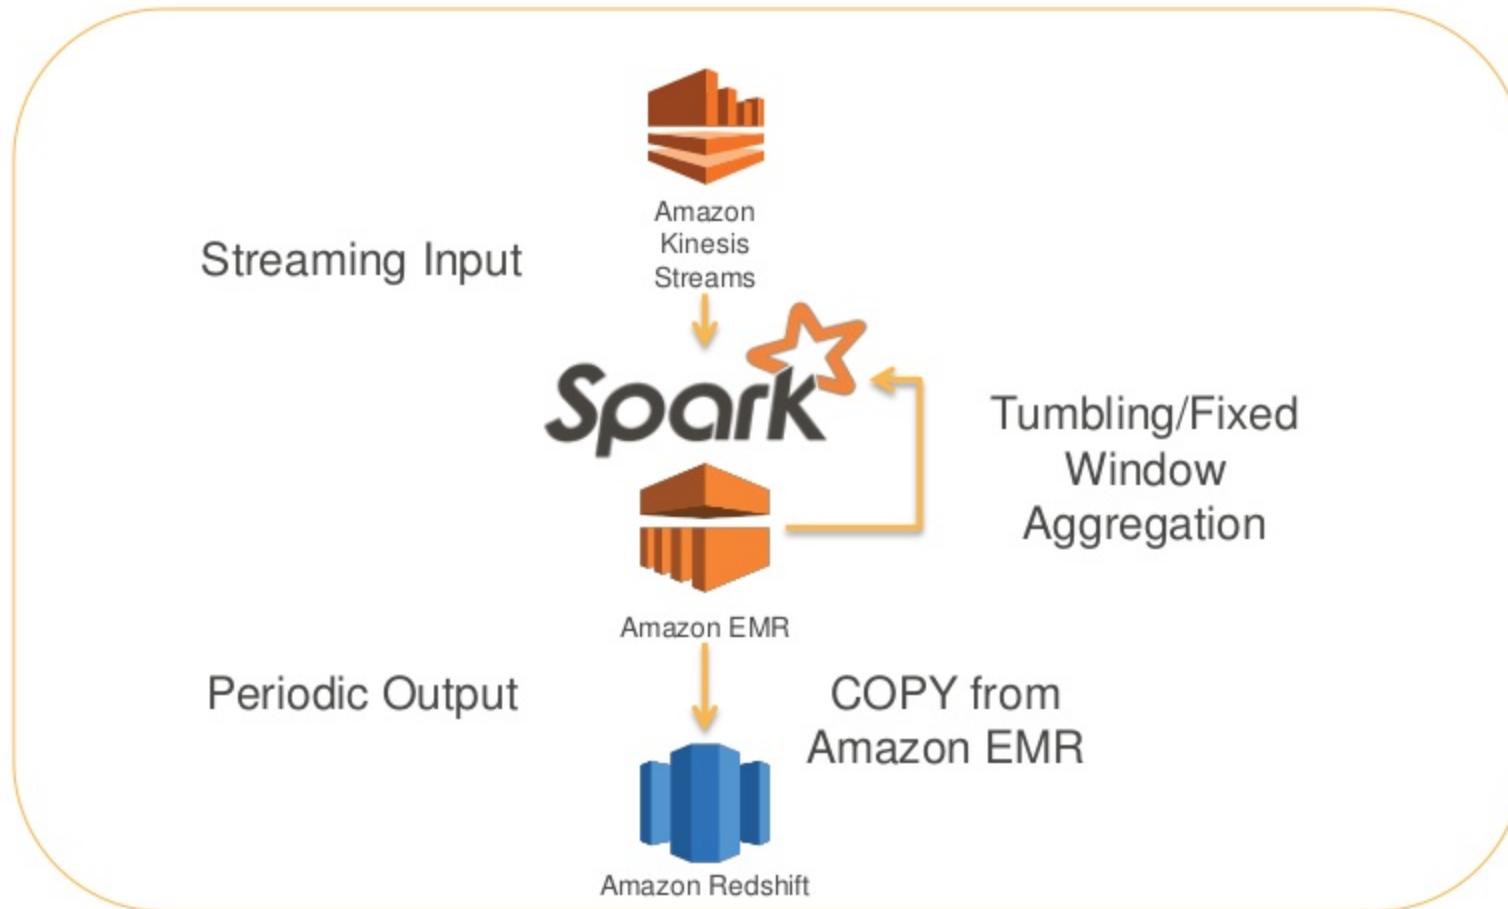
Example: Counting tweets on a sliding window

```
KinesisUtils.createStream('twitter-stream')
  .filter(_.getText.contains("Open-Source"))
  .countByWindow(Seconds(5))
```



Common Integration Pattern with Amazon EMR

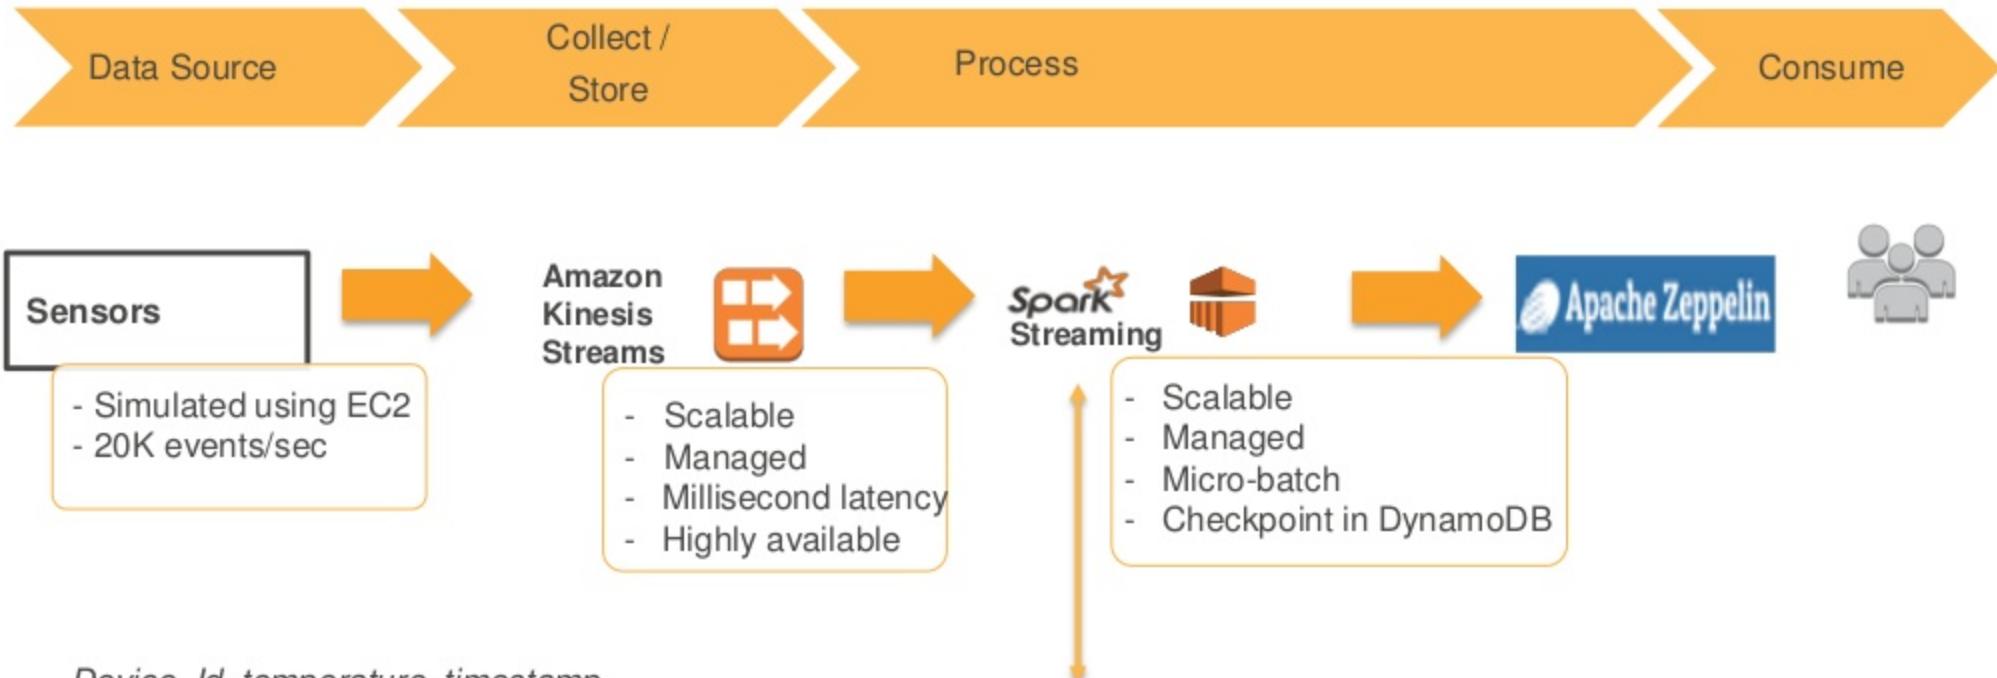
Tumbling Window Reporting



Using Spark Streaming with Amazon Kinesis Streams

1. **Use Spark 1.6+ with EMRFS consistent view option** – if you use Amazon S3 as storage for Spark checkpoint
2. **Amazon DynamoDB table name** – make sure there is only one instance of the application running with Spark Streaming
3. **Enable Spark-based checkpoints**
4. Number of Amazon Kinesis receivers is multiple of executors so they are load-balanced
5. Total processing time is less than the batch interval
6. Number of executors is the same as number of cores per executor
7. Spark Streaming uses default of 1 sec with KCL

Demo



Device_Id, temperature, timestamp

1,65,2016-03-25 01:01:20

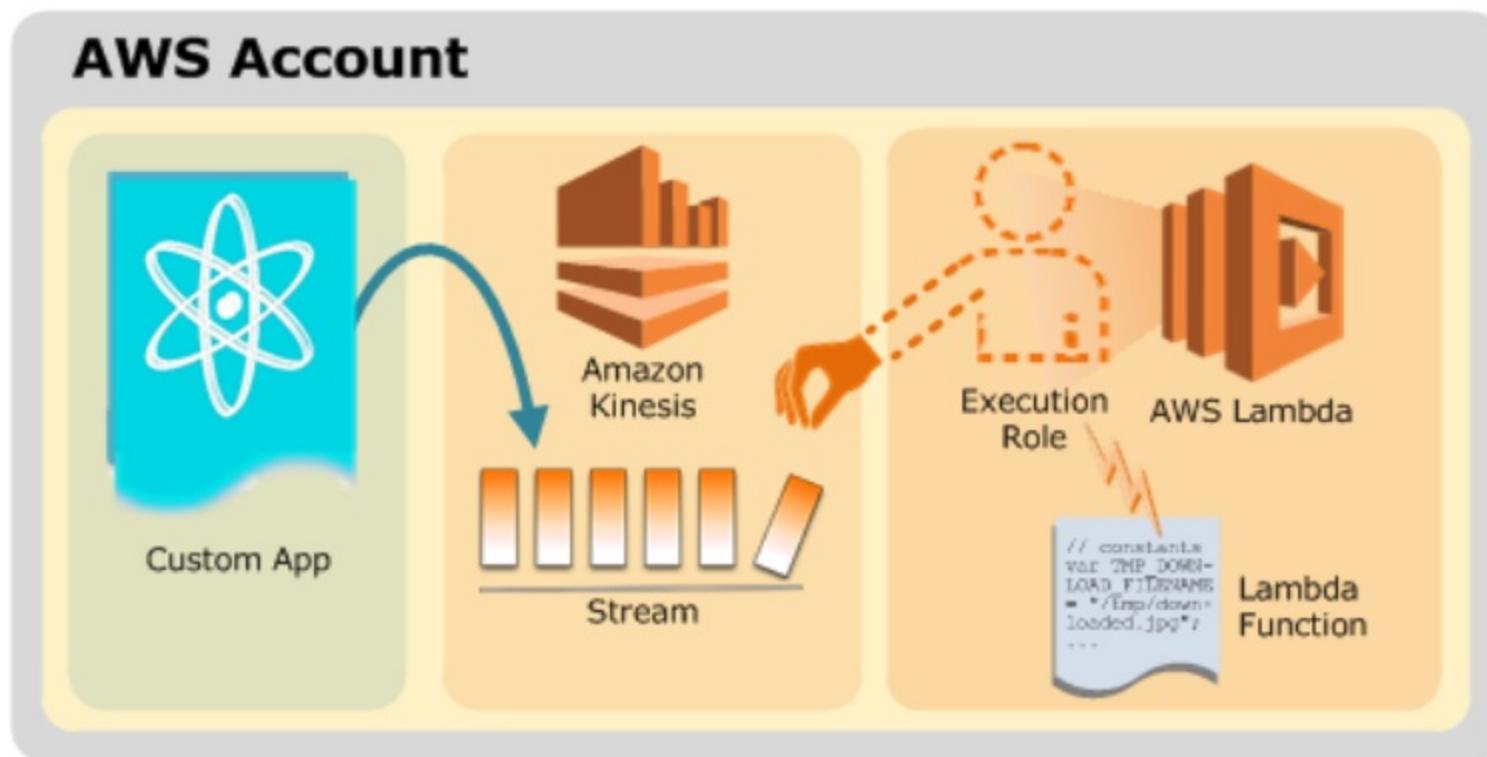
2,68,2016-03-25 01:01:20

3,67,2016-03-25 01:01:20

4,85,2016-03-25 01:01:20

Amazon
DynamoDB

Amazon Kinesis Streams with AWS Lambda

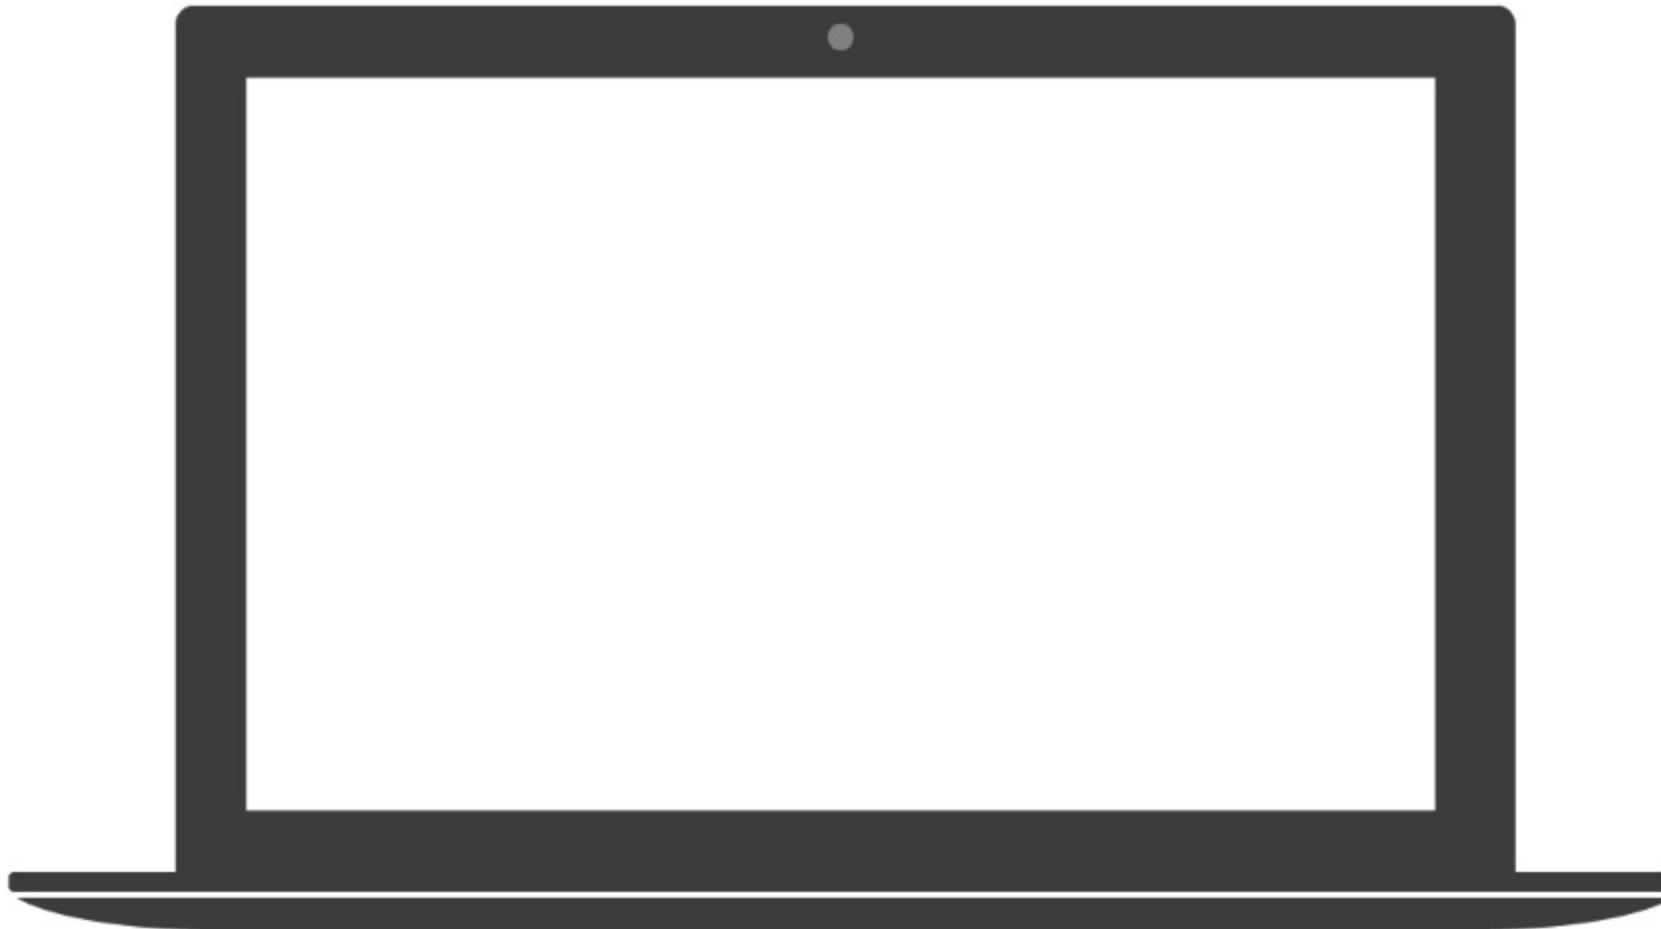


Hearst's wonderful world of streaming data

- Digital Media
- Publication brands
- Mobile Apps
- Re-targeting platforms

 <p>Car and Driver Car and Driver is a leading source of information for auto enthusiasts and in-market car buyers...</p> <p>Twitter Facebook LinkedIn</p>	 <p>Cosmopolitan Cosmopolitan is the best-selling young women's magazine, a bible for fun, fearless females...</p> <p>Twitter Facebook</p>	 <p>Country Living Country Living believes that how you live means a whole lot more than where you live...</p> <p>Twitter Facebook</p>
 <p>Dr. Oz THE GOOD LIFE Dr. Oz THE GOOD LIFE brings the upbeat, engaging personality and advice of Dr. Oz to life...</p> <p>Twitter</p>	 <p>ELLE ELLE inspires women to explore and celebrate their own style in all aspects of their lives...</p> <p>Twitter Facebook LinkedIn</p>	 <p>ELLE DECOR ELLE DECOR is where style lives. It opens the doors to the world's most stylish places...</p> <p>Twitter Facebook</p>
 <p>Esquire Esquire defines and celebrates what it means to be a man in contemporary American culture...</p> <p>Twitter Facebook</p>	 <p>Food Network Magazine Food Network Magazine brings passion, fun, and personalities around food to the table...</p> <p>Twitter Facebook</p>	 <p>Good Housekeeping Founded in 1885, Good Housekeeping magazine reaches nearly 24 million readers each month...</p> <p>Twitter Facebook</p>

Buzzing@Hearst Demo



The Business Value of Buzzing@Hearst

Real-Time Reactions

Instant feedback on articles from our audiences

Promoting Popular Content Cross-Channel

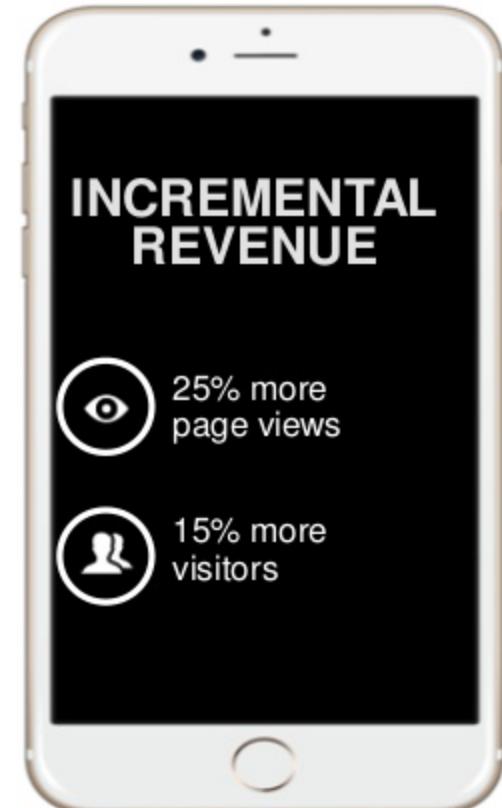
Incremental re-syndication of popular articles across properties
(e.g. trending newspaper articles can be adopted by magazines)

Authentic Influence

Inform Hearst editors to write articles that are more relevant to our audiences

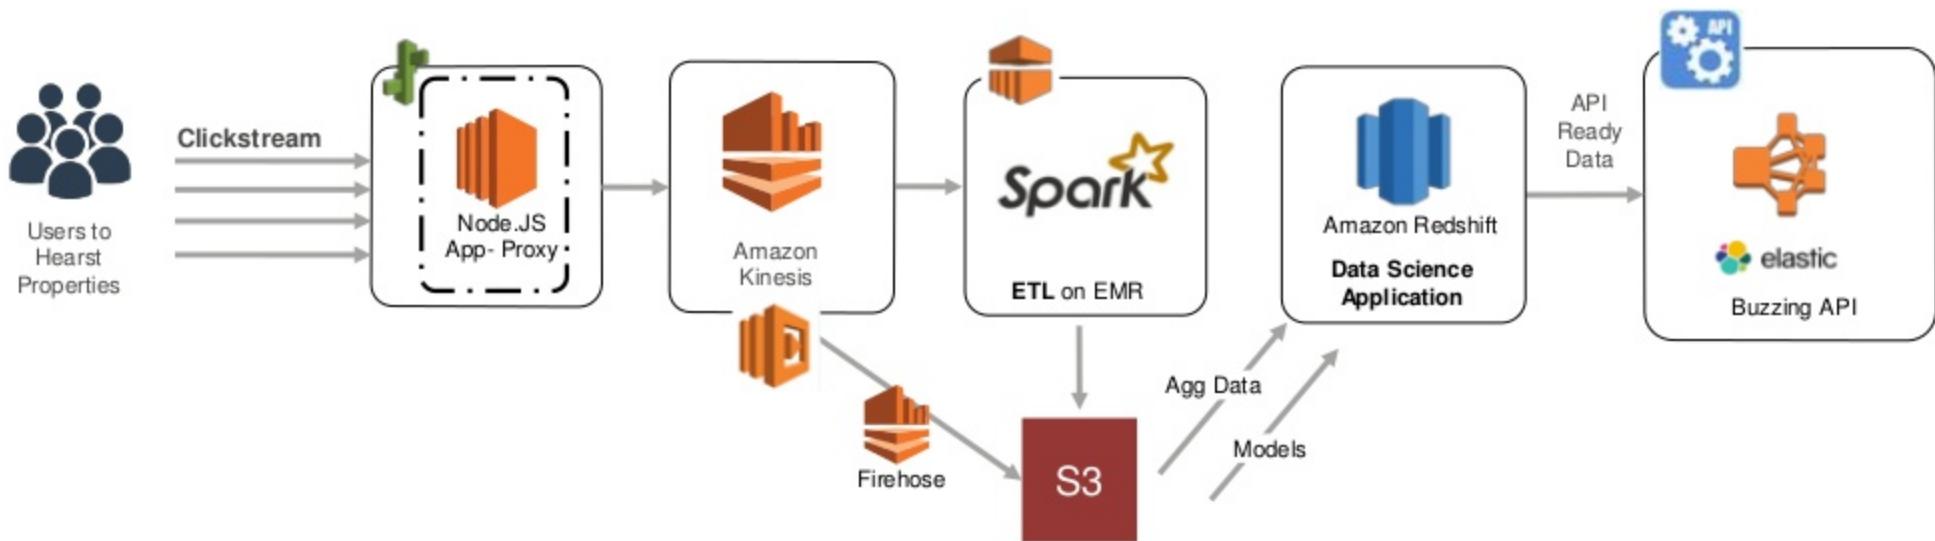
Understanding Engagement

Inform both editors what channels are our audiences leveraging to read Hearst articles





Final Hearst Data Pipeline



LATENCY

Milliseconds

30 Seconds

100 Seconds

5 Seconds

THROUGHPUT

100GB/Day

5GB/Day

1GB/Day

1GB/Day

Real-time streaming @sizmek

Orit Alul – R&D Director

Sizmek – Who? What? How?

▪ Who are we?

- **2nd** Largest Ad management company in the world operating globally in **50+** countries
 - Open Ad Management Platform

▪ Who are our customers?

- #### ▪ Advertising agencies and advertisers

▪ What is our value?

- We enable our clients to manage their cross-channel campaigns, and find their relevant audience seamlessly
 - We built a fast, high throughput infrastructure that enables real time analytics and decision making.

Sizmek
OPEN AD MANAGEMENT



Sizmek – Our customers

Over 13,000 brands
use our platform



Over 5,000 media
agencies use our
platform

group^m dentsu

HAVAS IPG

Omnicom WPP

Data processing at Sizmek



- 15B events per day
- 4 TB data per day
- 150K events per second
- Volatile data traffic load
 - During the day /week/year
 - During the special event(e.g. Black Friday)
- Increase demand for Realtime insights

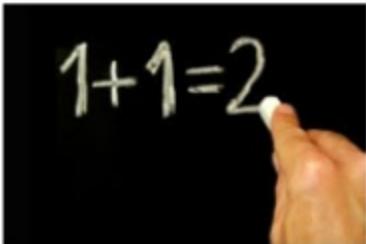
NXT generation of Analytics at Sizmek



Flexible

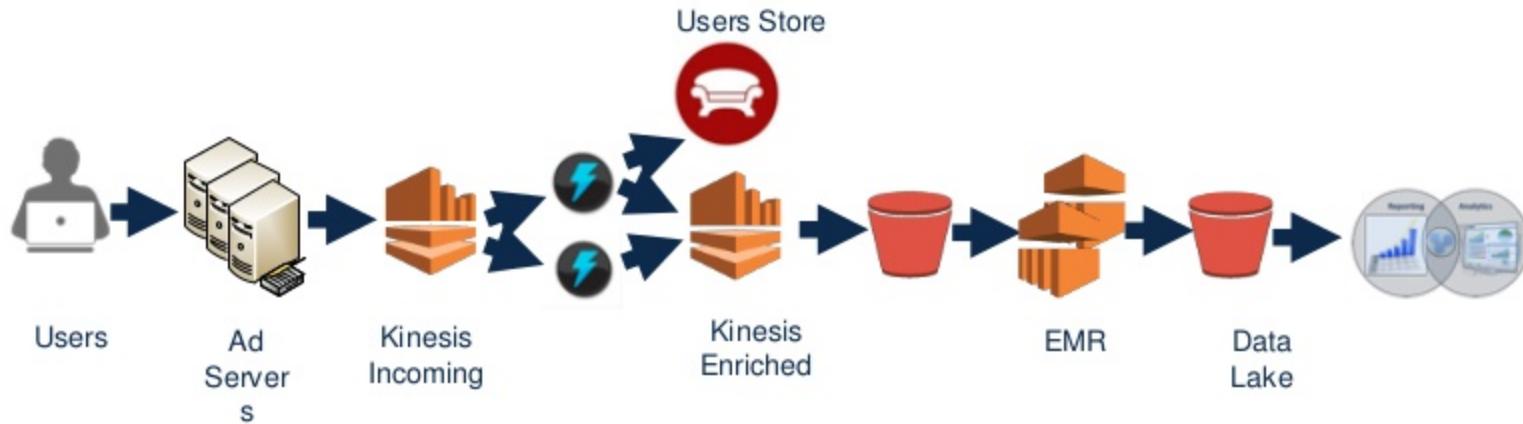


Faster



Simple

Sizmek Data Platform on AWS

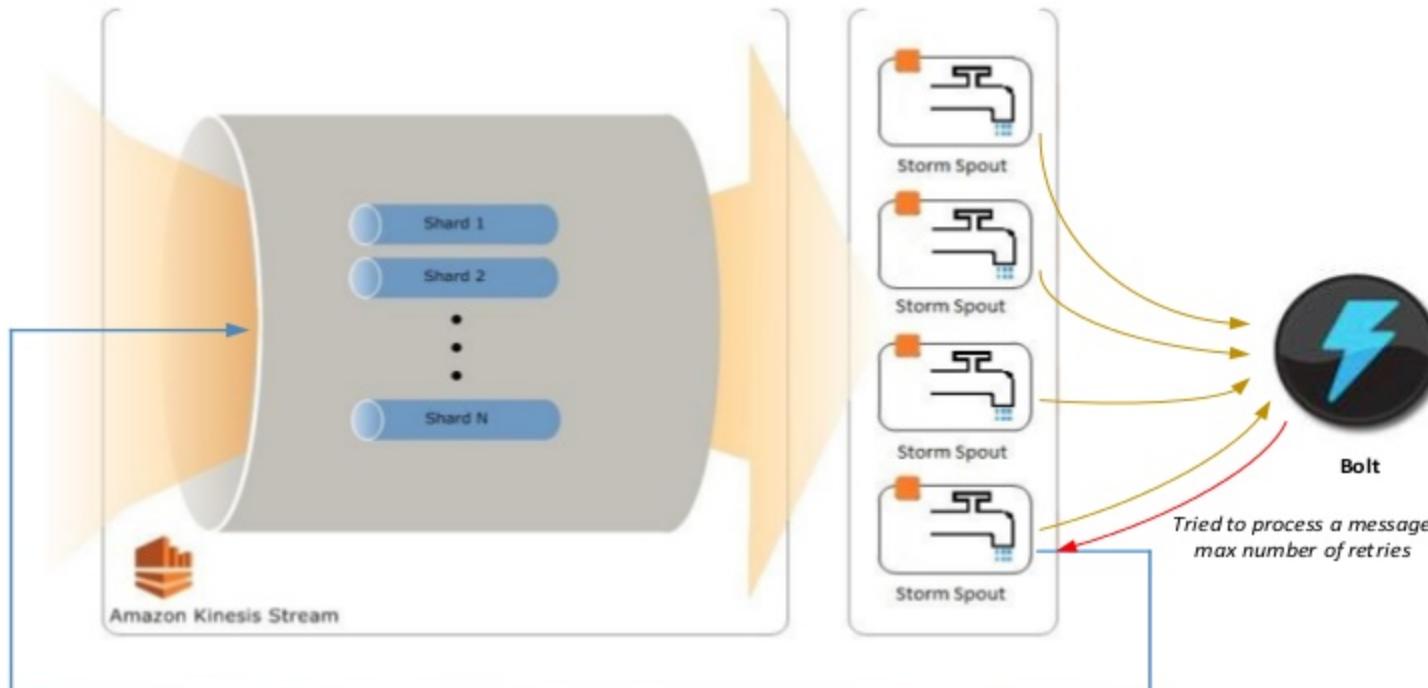


Our Kinesis Journey

Why Amazon Kinesis?

- Apache Kafka -> Amazon Kinesis = Managed Service = no Ops effort
- Integration with Apache Storm and AWS Lambda
- Integration with Amazon S3, Amazon Redshift and other AWS services
- Meets our throughput needs

Tip 1: Amazon Kinesis Apache Storm connector



Insert a non-processed message back to the stream as new message

Tip 2: putRecord = more shards!

- **Requirements:**

- Apache Storm requires processing event by event
- Average event size is 1KB
- We must use putRecord API

- **Testing results:**

- 75% of Amazon Kinesis maximum TP limit
- 300 - 500 shards

- **Conclusion**

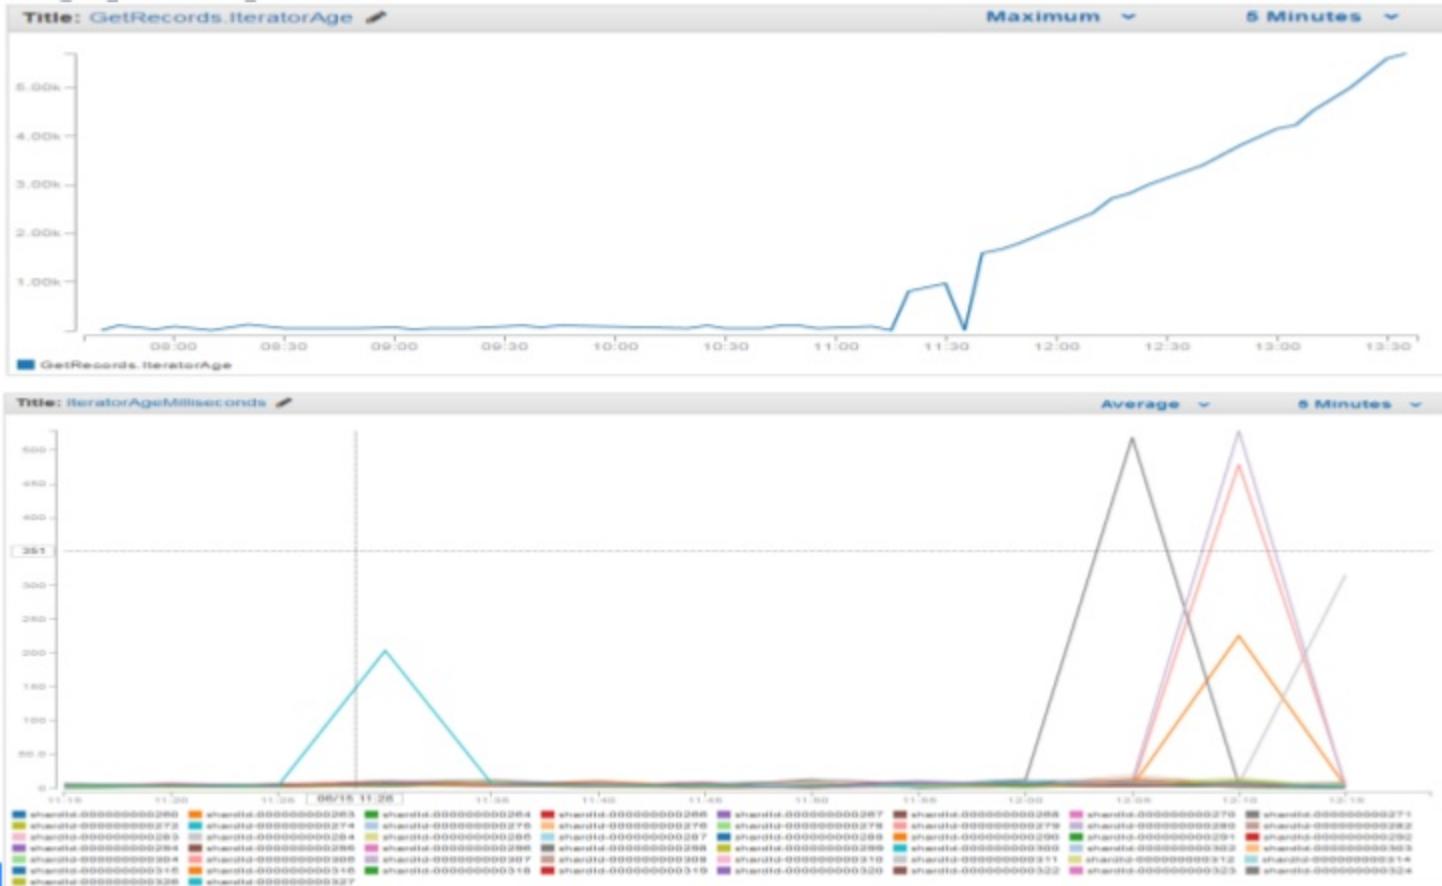
- Provision more shards (25%)
- Always test your workload

Tip 3: Error handling on record level

Kinesis putRecords (bulk insert) response

```
Date: <Date>
{
    "FailedRecordCount": 2,
    "Records": [
        {
            "SequenceNumber":
"49543463076548007577105092703039560359975228518395012686",
            "ShardId": "shardId-000000000000"
        },
        {
            "ErrorCode": "ProvisionedThroughputExceededException",
            "ErrorMessage": "Rate exceeded for shard shardId-000000000001 in
stream <StreamName> under account 111111111111."
        },
        {
            "ErrorCode": "InternalFailure",
            "ErrorMessage": "Internal service failure."
        }
    ]
}
```

Tip 4: Shard level metrics is your





A close-up photograph of a woman's face, partially obscured by several overlapping yellow circles. She has dark hair that is blowing in the wind, and she is smiling broadly. The background is a bright, slightly blurred teal color.

Thank You!

& Yes we're
hiring 😊

Conclusion

- Amazon Kinesis offers: managed service to build applications, streaming data ingestion, and continuous processing
- Ingest aggregate data using Amazon Producer Library
- Process data using Amazon Connector Library and open source connectors
- Determine your partition key strategy
- Try out Amazon Kinesis at <http://aws.amazon.com/kinesis/>

Reference

- **Technical documentations**
 - [Amazon Kinesis Agent](#)
 - [Amazon Kinesis Streams and Spark Streaming](#)
 - [Amazon Kinesis Producer Library Best Practice](#)
 - [Amazon Kinesis Firehose and AWS Lambda](#)
 - [Building Near Real-Time Discovery Platform with Amazon Kinesis](#)
- **Public case studies**
 - [Glu mobile – Real-Time Analytics](#)
 - [Hearst Publishing – Clickstream Analytics](#)
 - [How Sonos Leverages Amazon Kinesis](#)
 - [Nordstrom Online Stylist](#)



Got Feedback on the Session?



Or

Get the Summit App



Comment on the Summit App



AWS

S U M M I T

tel aviv

Thank you!
benaltar@amazon.com

