



Amazon DynamoDB and Amazon RDS Deep Dive

Sundar Raghavan

raghavas@amazon.com





Provides real-time weather updates



Makes video gamers happy



Scales to millions of users



Delivers billions of page flips

How?





Provides real-time weather updates



Makes video gamers happy



Scales to millions of users



Delivers billions of page flips

Amazon
DynamoDB

Amazon
RDS





YesSQL?

NoSQL?

(Hint)
Focus on your App





Factors	YesSQL (RDS)	NoSQL (DynamoDB)
Application	<ul style="list-style-type: none">• App with complex business logic?	<ul style="list-style-type: none">• Web app with lots of users?
Transactions	<ul style="list-style-type: none">• Complex txns, joins, updates?	<ul style="list-style-type: none">• Simple data model, updates, queries?
Scale	<ul style="list-style-type: none">• Developer managed	<ul style="list-style-type: none">• Automatic, on-demand scaling
Performance	<ul style="list-style-type: none">• Developer architected	<ul style="list-style-type: none">• Consistent, high performance at scale
Availability	<ul style="list-style-type: none">• Architected for fail-over	<ul style="list-style-type: none">• Seamless and transparent
Skills	<ul style="list-style-type: none">• Relational + Web expertise	<ul style="list-style-type: none">• Web expertise

Best of both worlds: Use SQL and NoSQL models in one app





Focus on **your App** with Amazon DynamoDB





Amazon DynamoDB

Quick overview





DynamoDB is a managed NoSQL database service.

Store and retrieve any amount of data.
Serve any level of request traffic.
Without any operational burden.





Consistent, predictable performance.

Single digit millisecond latencies.
Backed on solid-state drives.





Flexible data model

Key/attribute pairs.
No schema required.

Easy to create. Easy to adjust.





Seamless scalability

No table size limits. Unlimited storage.

No downtime.





Durable

Consistent, disk-only writes (not memory)

Replication across data centres and availability zones.





Without the
operational burden.

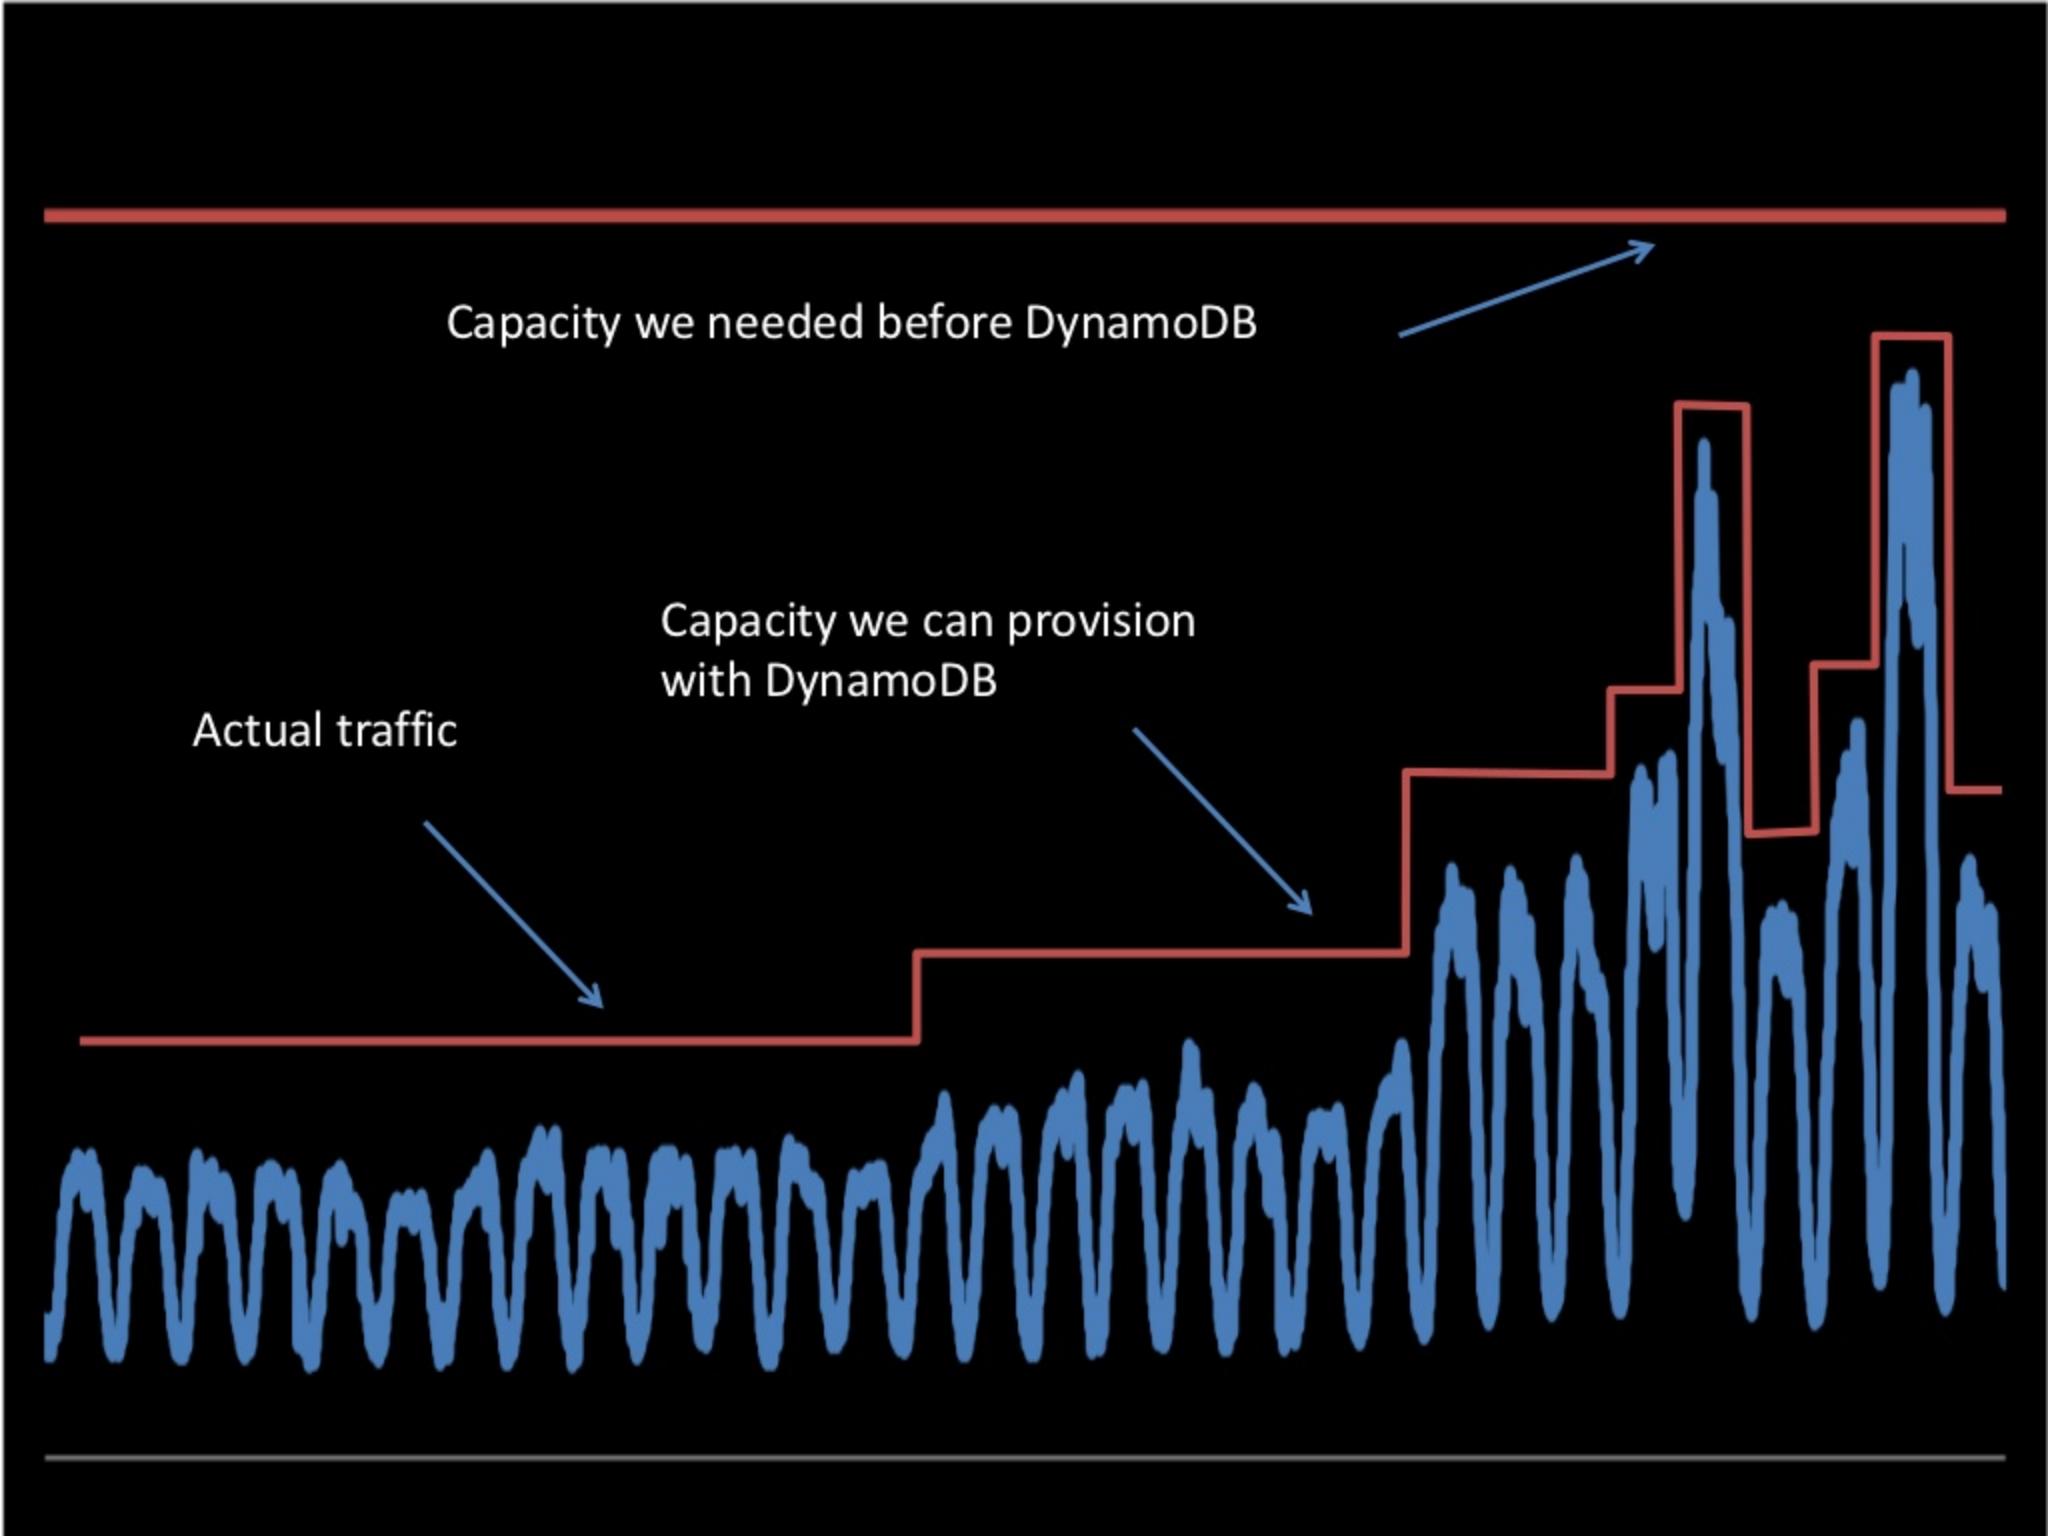
FOCUS ON YOUR APP



November traffic at Amazon.com

76%

24%



Actual traffic

Capacity we can provision
with DynamoDB

Capacity we needed before DynamoDB



Two decisions + three
clicks = ready for use





Primary keys +
Provisioned level of throughput (easy to change)

Two decisions + three
clicks = ready for use





Order Table
(collection of items)

Item (max size 64k, collection of attributes)

Attribute (key value pair)
String – unicode, utf8 binary
Number – 38 digit precision

id = 100	date = 2012-05-16-09-00-10	total = 25.00
id = 101	date = 2012-05-16-09-00-11	total = 50.00

Hash Key Range Key





Items are indexed by
primary key

Single hash keys and composite keys.





Think of database operations as reads and writes

Item read = 1 read

Item update = 2 reads and 1 write

Item insert = 1 write





Provisioned throughput

Reserve IOPS for reads and writes.

Scale up (or down) at any time.

Priced per hour of provisioned throughput.





Provisioned throughput

Write throughput units = item size * writes/sec

\$0.01 per hour for 10 write units

Read throughput units = item size * ½ * reads/sec

\$0.01 per hour for 100 read units





Consistent writes

Atomic increment/decrement.

Optimistic concurrency control.

aka: “conditional writes”.





Transactions

Item level transactions only.

Puts, updates and deletes are ACID.





strongly consistent

eventually consistent

Read throughput

Provisioned units =
size of item x reads/second

\$0.01 per hour for 50 read units





strongly consistent

eventually consistent

Read throughput

Same latency expectations.

Mix and match at “read time”.



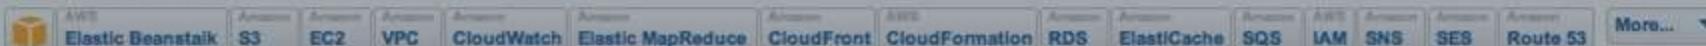


Two decisions + three
clicks = ready for use



AWS Management Console > Amazon DynamoDB

Matt Wood Help



Tables

Region: [US East \(N. Virginia\)](#)[Help](#)

Filter:

Name

images

items

tags

Create Table

Cancel

PROVISIONED
THROUGHPUT CAPACITYTHROUGHPUT ALARMS
(optional)

Table Name:

photos

Table will be created in us-east-1 region

Primary Key:

DynamoDB is a schema-less database. You only need to tell us your primary key attribute(s).

Primary Key Type: Hash Hash and Range String Number

Hash Attribute Name:

id

Choose a hash attribute that ensures that your workload is evenly distributed across hash keys.

For example, "Customer ID" is a good hash key, while "Game ID" would be a bad choice if most of your traffic relates to a few popular games.

[Learn more about choosing your primary key](#)

Cancel

Continue

Help

Tables

Region: [US East \(N. Virginia\)](#)

Filter:

Name

images

items

tags

Create Table

Cancel X

Help

PRIMARY KEY PROVISIONED THROUGHPUT CAPACITY

THROUGHPUT ALARMS (optional)

Provisioned Throughput Capacity:

 Help me calculate how much throughput capacity I need to provision

Throughput capacity to provision:

Amazon DynamoDB lets you specify how much read and write throughput capacity you wish to provision for your table. Using this information, Amazon will provision the appropriate resources to meet your throughput needs. [More Information](#)

Read Capacity Units:

5

Write Capacity Units:

5

⚠ If you exceed the free tier you are charged for the provisioned throughput capacity of your table **even if you do not actively use your provisioned capacity.** [Learn more](#) about DynamoDB's free tier and pricing.

Back

Continue Help 



Matt Wood | Help ▾

AWS Management Console > Amazon DynamoDB

[AWS Elastic Beanstalk](#) [Amazon S3](#) [Amazon EC2](#) [Amazon VPC](#) [Amazon CloudWatch](#) [Amazon Elastic MapReduce](#) [Amazon CloudFront](#) [Amazon CloudFormation](#) [Amazon RDS](#) [Amazon ElastiCache](#) [Amazon SQS](#) [Amazon IAM](#) [Amazon SNS](#) [Amazon SES](#) [Amazon Route 53](#) More...**Tables**Region: [US East \(N. Virginia\)](#)

Filter:

▲ Name

images

items

tags

Create Table

Cancel

PRIMARY KEY

PROVISIONED
THROUGHPUT CAPACITYTHROUGHPUT ALARMS
(optional)**Throughput Alarms (optional)** Use Basic Alarms

Notify me when my table's request rates exceed of Provisioned Throughput for 60 minutes.

Notification will be sent when:

- Read Capacity Units consumed > 4
or
- Write Capacity Units consumed > 4

Send notification to:

Additional charges may apply if you exceed the AWS Free Tier levels for CloudWatch or Simple Notification Service.

Advanced alarm settings are available in the CloudWatch Management Console.

[Back](#)[Create Table](#) [Help](#)



Two decisions + three
clicks = ready for use





Two decisions + one API call
= ready for use





```
$create_response = $dynamodb->create_table(array(  
    'TableName' => 'ProductCatalog',  
    'KeySchema' => array(  
        'HashKeyElement' => array(  
            'AttributeName' => 'Id',  
            'AttributeType' => AmazonDynamoDB::TYPE_NUMBER  
        )  
    ),  
    'ProvisionedThroughput' => array(  
        'ReadCapacityUnits' => 10,  
        'WriteCapacityUnits' => 5  
    )  
));
```





Two decisions + one API call
= ready for **use**





Two decisions + one API call
= ready for **development**





Two decisions + one API call
= ready for production





Two decisions + one API call
= ready for **scale**



Tables

Region: US East (N. Virginia)

Filter: tags

Name

tags

Modify Provisioned Throughput

Cancel

Modify throughput for: tags**Provisioned Throughput**

Throughput can be increased by a factor of two or less each time you modify your table.

Read Capacity Units: (allowed values: 5 - 20)Write Capacity Units: (allowed values: 5 - 20) **Use Basic Alarms**

Notify me when my table's request rates exceed of Provisioned Throughput for 60 minutes.

Notification will be sent when:

- Read Capacity Units consumed > 16
or
- Write Capacity Units consumed > 16

Send notification to:

Cancel **Save**

Recent Alarms

--- No alarms triggered ---

[Go to CloudWatch](#)**CloudWatch Metrics (for the past hour)**

Click to view in CloudWatch. Links open in new windows or tabs

Table Name: tags

Metric Name: ConsumedReadCapacityUnits

Operation: All



Authentication

Session based to minimize latency.

Uses Amazon Security Token Service.

Handled by AWS SDKs.

Integrates with IAM.





Monitoring

CloudWatch metrics:
latency, consumed read and write
throughput, errors and throttling.





Libraries, mappers & mocks

ColdFusion, Django, Erlang, Java, .Net,
Node.js, Perl, PHP, Python, Ruby

<http://j.mp/dynamodb-libs>





(And I can't
find my joins)

Ma, I lost my schema!

Tables do not require a formal schema.
Items are an arbitrary sized hash.
Just need to specify the primary key.





Programming DynamoDB

Small but perfectly formed.

Whole programming interface
fits on one slide.







Query patterns

Retrieve all items by hash key.

Range key conditions:

`==, <, >, >=, <=, begins with, between.`

Counts. Top and bottom n values.

Paged responses.





Modeling patterns





1. Mapping relationships with range keys.

No cross-table joins in DynamoDB.

Use composite keys to model
relationships.





2. Handling large items.

Unlimited attributes per item.

Unlimited items per table.

Max 64k per item.





Store a pointer to objects in Amazon S3

Large data stored in S3.
Location stored in DynamoDB.
99.999999999% data durability in S3.





3. Managing secondary indices.

Not currently supported by DynamoDB.

You can create on your own.





4. Time series data.

Logging, click through, ad views,
game play data, application usage.

Non-uniform access patterns.

Newer data is ‘live’.

Older data is read only.





Hot and cold tables.





Amazon CloudDrive

- Stores user objects in Cloud
 - Metadata in DynamoDB
 - Objects in S3
- Queries and object searches are served by DynamoDB
 - What are all the objects in my drive?
 - Find me all the music albums stored on my drive



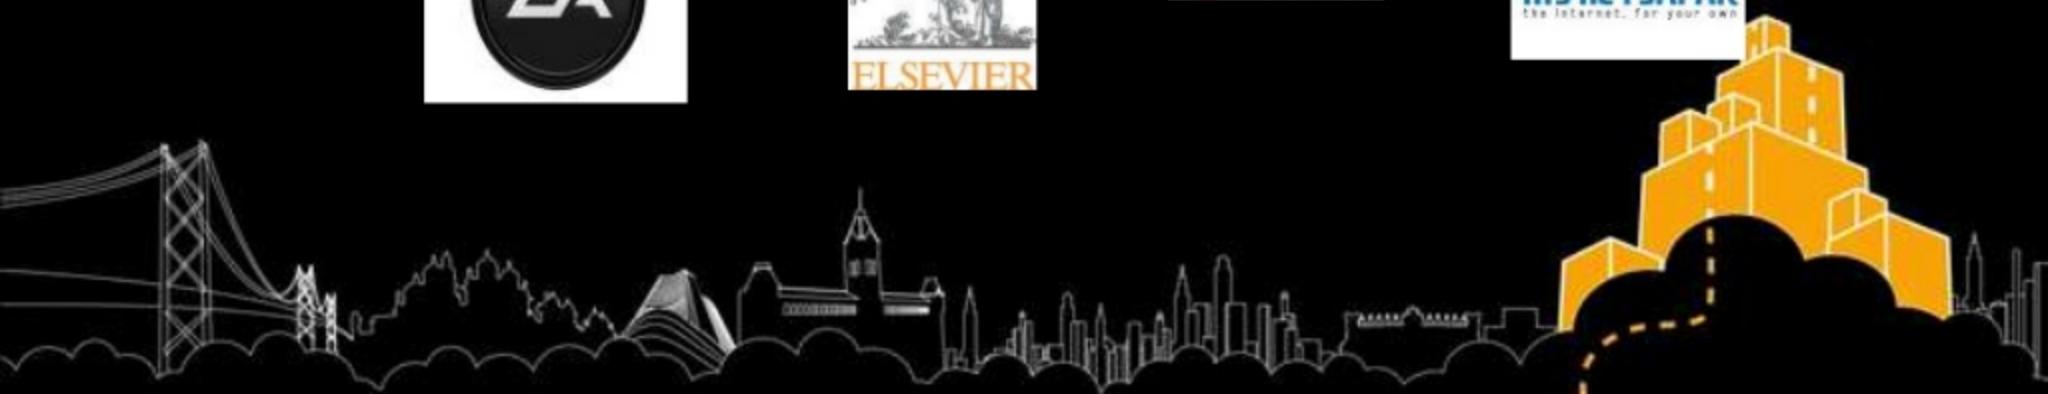
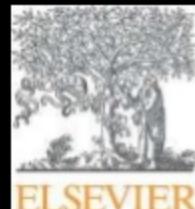
"We were excited by how fast we were able to put DynamoDB into production and how much developer time we saved. In addition, DynamoDB lets us scale up and down easily by simply reserving increased throughput capacity when we need it and dialing it back when we don't"

- Russel Dicker, Amazon CloudDrive





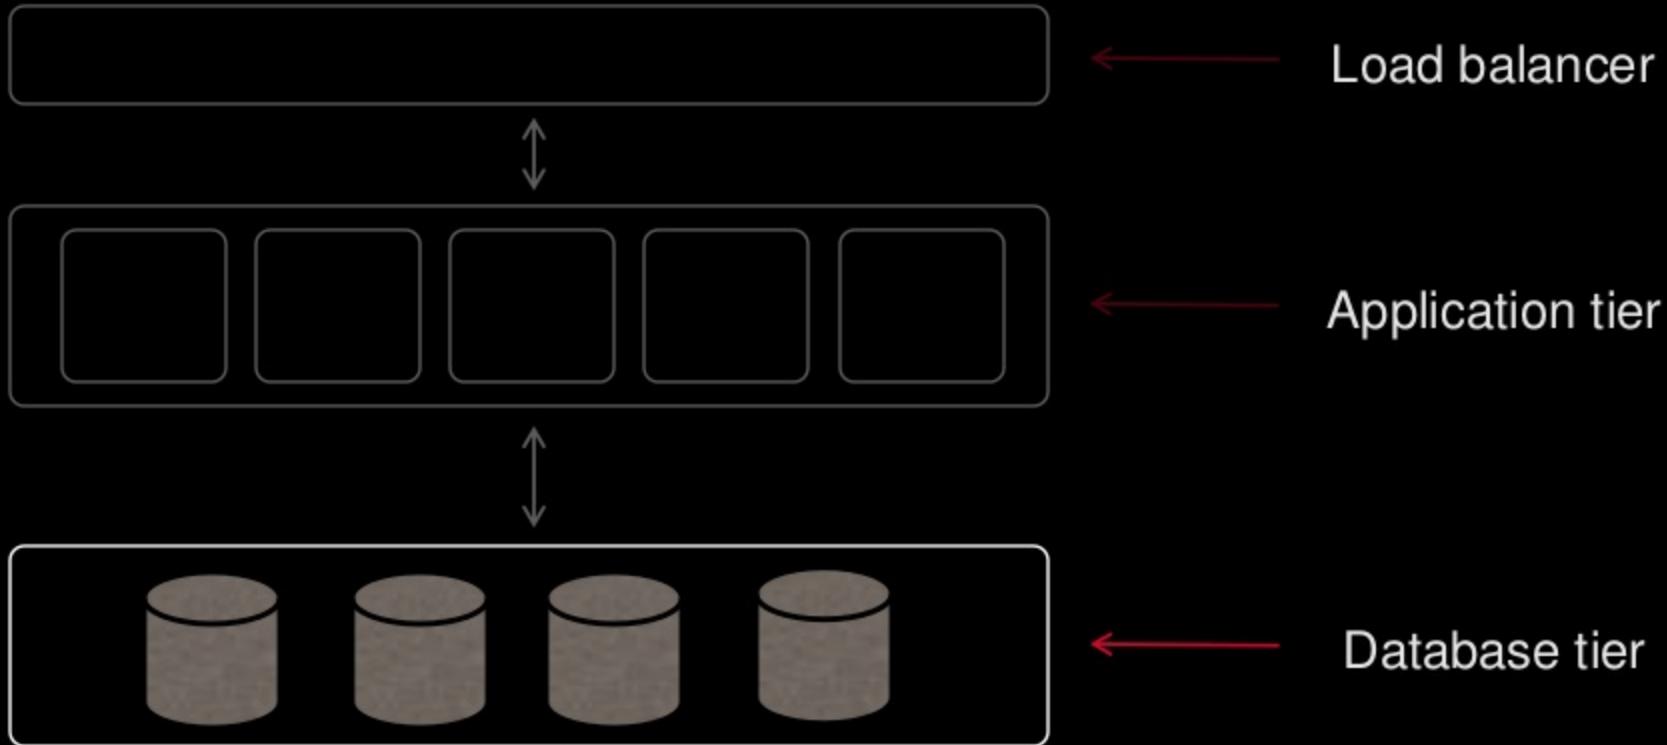
DynamoDB Customer Highlights

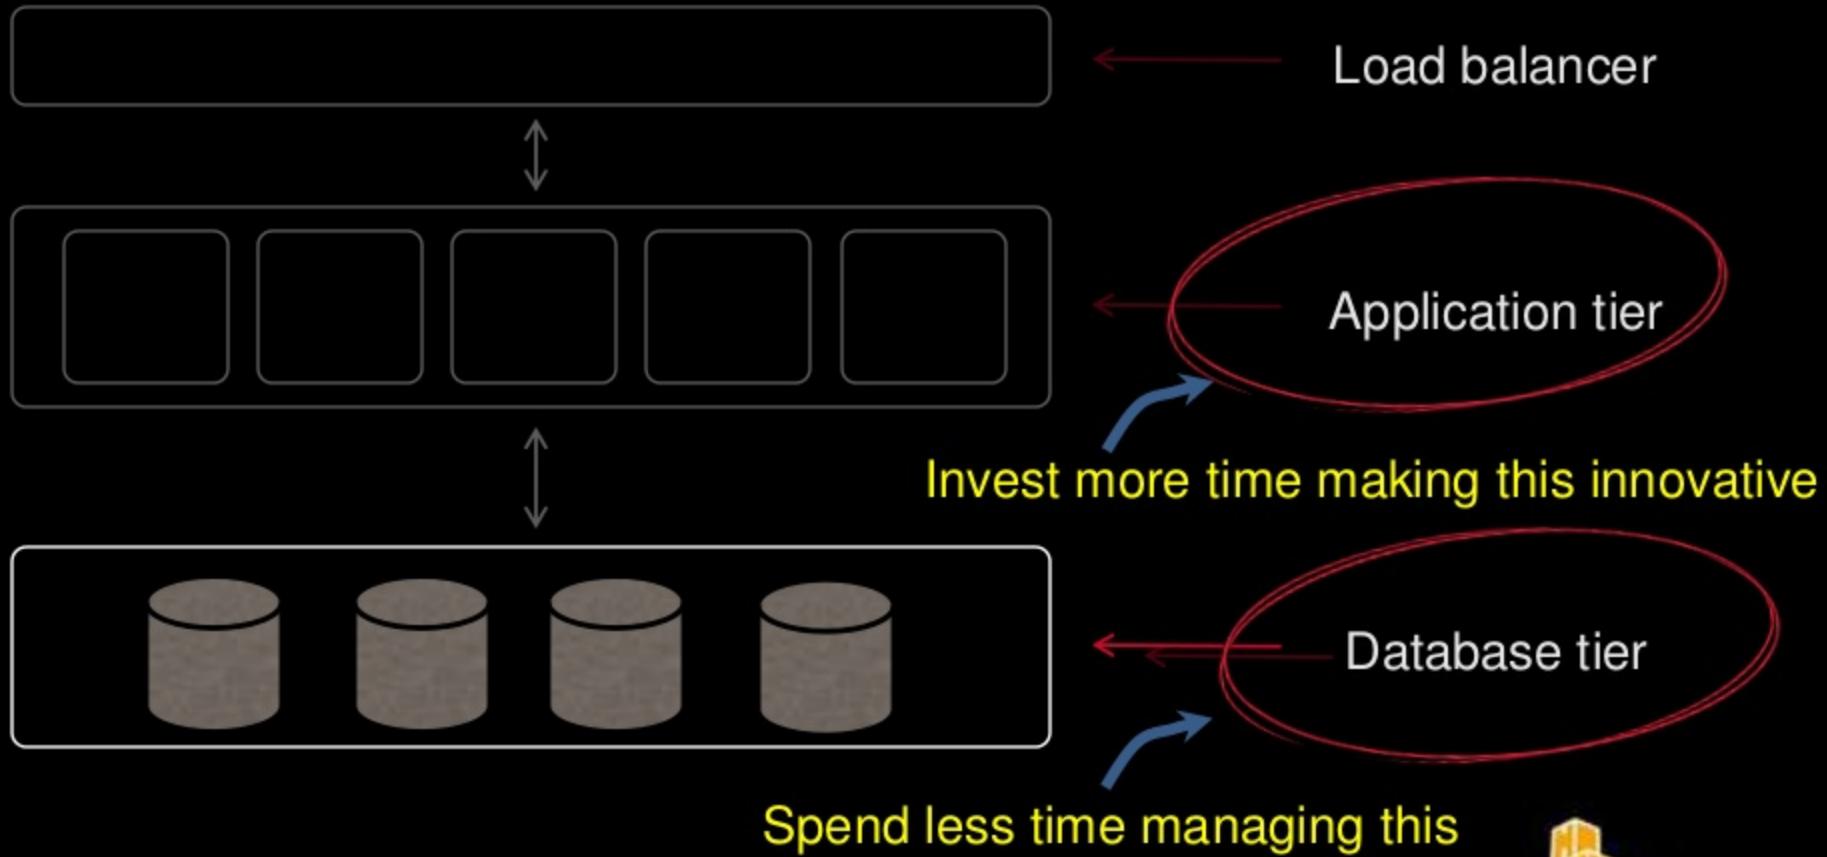




Focus on your App with Amazon RDS









Schema design

Patching

Configuration

Frequent server
upgrades

Query optimization

Migration

Backup and
recovery

Query construction

Software upgrades

Storage upgrades

Hardware crash





Schema design

Query construction

Query optimization

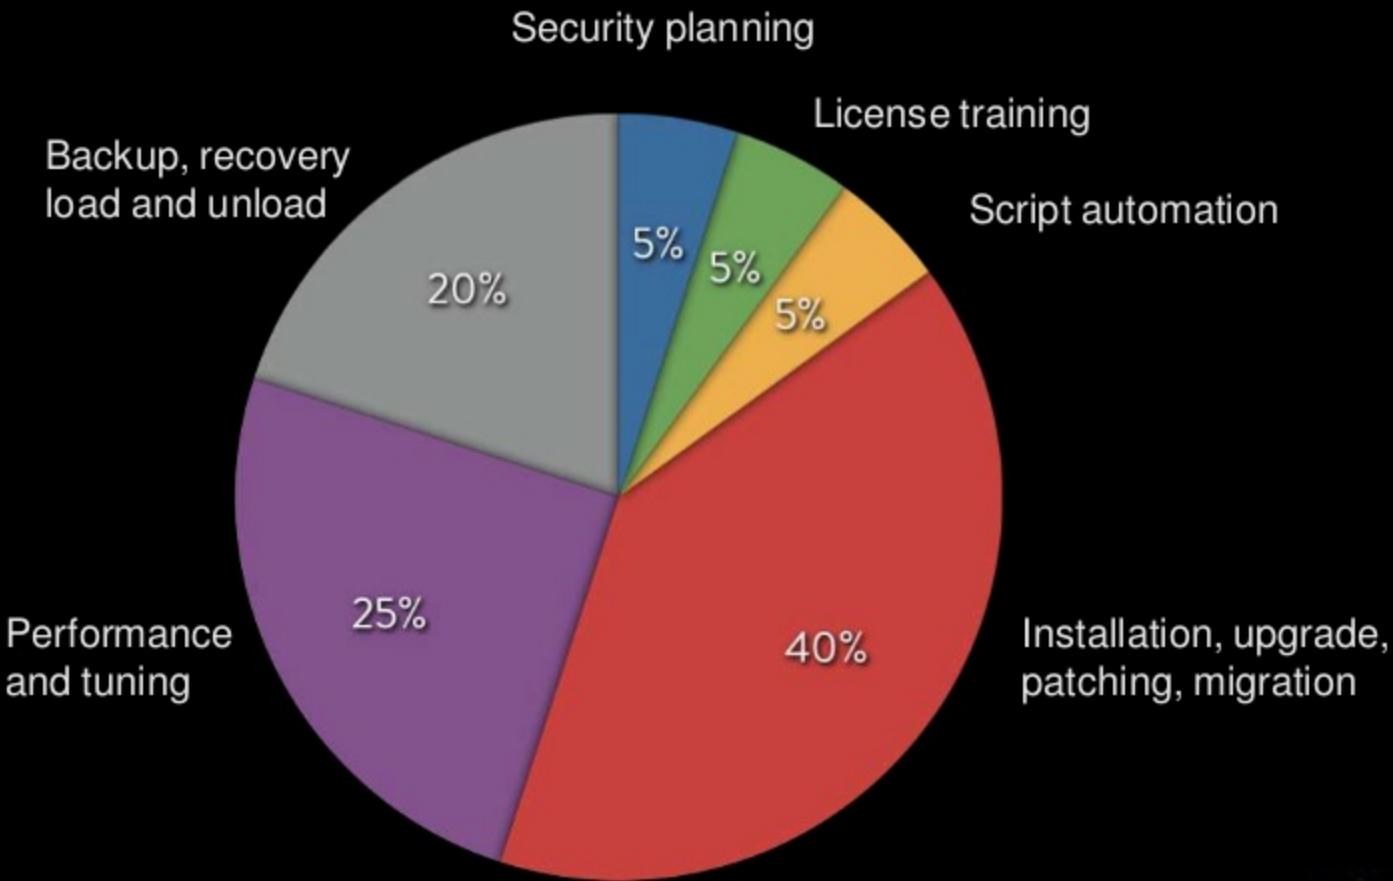
Focus on the “buck”

Off load the “muck”

Migration
Backup and recovery
Patching
Configuration

Software upgrades
Storage upgrades
Frequent server upgrades
Hardware crash





Source: Forrester



Amazon RDS

Relational Database Service

Quick introduction





i Amazon RDS Overview

Amazon RDS is a fully managed relational database service.

Choice of Database engines

Simple to deploy

Easy to operate and scale

Reliable

Cost effective





Choice of database engines

MySQL

Oracle

Microsoft SQL Server





Three decisions + a few clicks
= ready for use





Database engine +
Instance size & storage + } Easy to change
Availability

Three decisions + a few clicks
= ready for use





Database engine +
Instance size & storage + } Easy to change
Availability

Three decisions + a few clicks
= ready for use





Launch DB Instance Wizard

Cancel

ENGINE SELECTION

DB INSTANCE DETAILS

ADDITIONAL CONFIGURATION

MANAGEMENT OPTIONS

REVIEW

To get started, choose the DB Instance details below and click **Continue**

MySQL	mysql MySQL Community Edition	Select
ORACLE	oracle-ee Oracle Database Enterprise Edition	Select
Microsoft SQL Server	sqlserver-ex Microsoft SQL Server Express Edition <i>Note that SQL Server Express Edition limits the storage of per database to a maximum of 10GB. Refer to this link for more details.</i>	Select
Microsoft SQL Server	sqlserver-web Microsoft SQL Server Web Edition <i>Note that in accordance with Microsoft's licensing policies, SQL Server Web Edition can only be used to support public and</i>	Select



Launch DB Instance Wizard

[Cancel](#)[ENGINE SELECTION](#)[DB INSTANCE DETAILS](#)[ADDITIONAL CONFIGURATION](#)[MANAGEMENT OPTIONS](#)[REVIEW](#)

To get started, choose a DB engine below and click **Continue**

DB Engine: mysql

License Model: General Public License ▾

DB Engine Version: 5.5.27 (default) ▾

DB Instance Class: db.m2.4xlarge ▾

Multi-AZ Deployment: Yes ▾

Auto Minor Version Upgrade: Yes No

Provide the details for your RDS Database Instance.

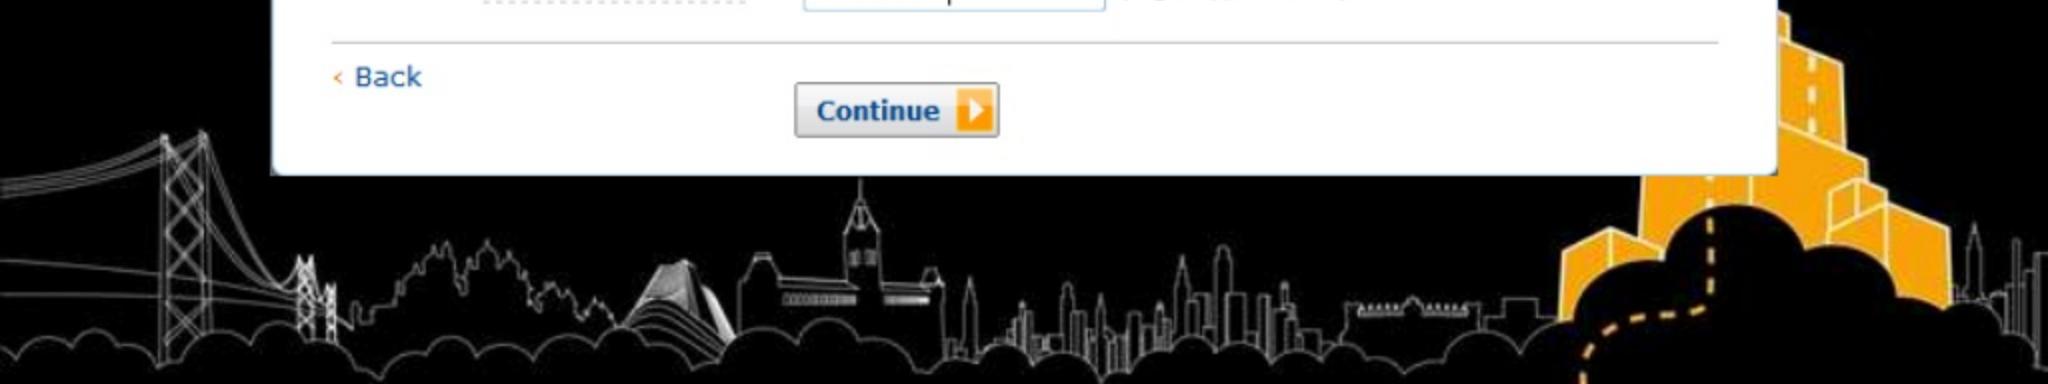
Allocated Storage:* GB (Minimum: 5 GB, Maximum: 1024 GB)

Higher allocated storage [may improve IOPS performance](#).

DB Instance Identifier:* (e.g. mydbinstance)

Master Username:* (e.g. awsuser)

Master Password:* (e.g. mypassword)

[Back](#)[Continue](#)

Launch DB Instance Wizard

[Cancel](#)

ENGINE SELECTION

DB INSTANCE DETAILS

ADDITIONAL CONFIGURATION

MANAGEMENT OPTIONS

REVIEW

Please review the information below, then click **Launch DB Instance**.

Engine: mysql

Engine Version: 5.5.27

License Model: general-public-license

Auto Minor Ver. Upgrade: Yes

DB Instance Class: db.m2.4xlarge

Multi-AZ Deployment: Yes

Allocated Storage: 100

DB Instance Identifier: mydb

Master User Name: greatguy

Master User Password: *****

Database Name: aaa

Database Port: 3306

Availability Zone: Using a Multi-AZ Deployment disables this preference.

Option Group: default:mysql-5-5

DB Parameter Group: default.mysql5.5

DB Security Group(s): default

DB Subnet Group:

Backup Retention Period: 1

Backup Window: No Preference

Maintenance Window: No Preference

[Back](#)[Launch DB Instance](#)



Three decisions + a few clicks
= ready for **use**





Three decisions + a few clicks
= ready for **development**





Three decisions + a few clicks
= ready for **scale**





Three decisions + a few clicks
= ready for production





Three ways to remove the “muck”





Productivity



Performance



Data Protection





⌚ Productivity

Most admin tasks are managed by RDS

Provisioning

Database backups

Patching

Performance management





⌚ Productivity

Database snapshots

- Automatic
- Point in time recovery
- Faster time to recover





⌚ Productivity

- ✓ Multiple databases per instance
- ✓ Standard user accounts
- ✓ Connect and query using common MySQL tools & drivers
- ✓ Tune engine parameters
- ✓ Import and export data using standard MySQL tools (mysqldump)
- ✓ Diagnostics
- ✓ Native MySQL replication
- ✓ SSL for encryption over the wire
- ✓ Monitor metrics



✗ Shell, super user or direct file system access (**Think security!**)



⌚ Productivity

Near zero administration

Painless patching, automatic upgrades

Cloudwatch monitoring, Metric alarms

One click. High Availability.





Modify DB Instance: test1

Cancel X

DB Engine Version: 5.5.20	Parameter Group: default.mysql5.5
DB Instance Class: db.m1.large	Security Group: default
Multi-AZ Deployment: Yes	Option Group: defaultmysql-5-5
Allocated Storage: 200 GB	New Master Password: <input type="text"/>
Backup Retention Period: 10 days	
Backup Window: Start Time 07 : 30 UTC	
Duration : 0.5 hours	
Maintenance Window: Start Time Saturday 09 : 00 UTC	
Duration : 0.5 hours	
Apply Immediately: <input type="checkbox"/>	Cancel Continue





⌚ Productivity

One click. High availability with Multi-AZ

Automated deployment across multiple AZs

Synchronous replication from master to replica

Automatic fail-over; replica promoted to master

Test fail-over





Performance

Push-button scale, high performance

Scale storage from 5Gb to 1Tb of storage

Scale instance from small to 4XL (better I/O)

Add read replicas with asynchronous replication

Add **ElastiCache** for performance





Performance

Amazon ElastiCache

In-memory cache service

Ideal front end to RDS for read-heavy applications

Low latency





Data Protection

Painless data protection

Daily database back-ups (save up to 35 days)

Snapshot any time. Point in time restore-use innodb

Recover everything up to final 5 minutes-innodb!

VPC, SSL, DB Security groups





Samsung saved \$34M

Problem:

- Needed to reduce IT costs and were looking to create a more flexible IT environment

Solution:

- AWS and RDS services. With every request, the application authenticates devices, delivers apps and content, and pushes notifications.

Business Benefits:

- Saved \$34M in hardware and maintenance expenses, 85% less than running on-premises





YesSQL

NoSQL

Focus on your App





Amazon RDS Customer Highlights





Thank you

Free tier or
Free trial

aws.amazon.com/dynamodb

aws.amazon.com/rds

raghavas@amazon

