

Getting Started with Amazon DynamoDB

Padma Malligarjuna
Sr. Technical Account Manager
AWS Enterprise Support



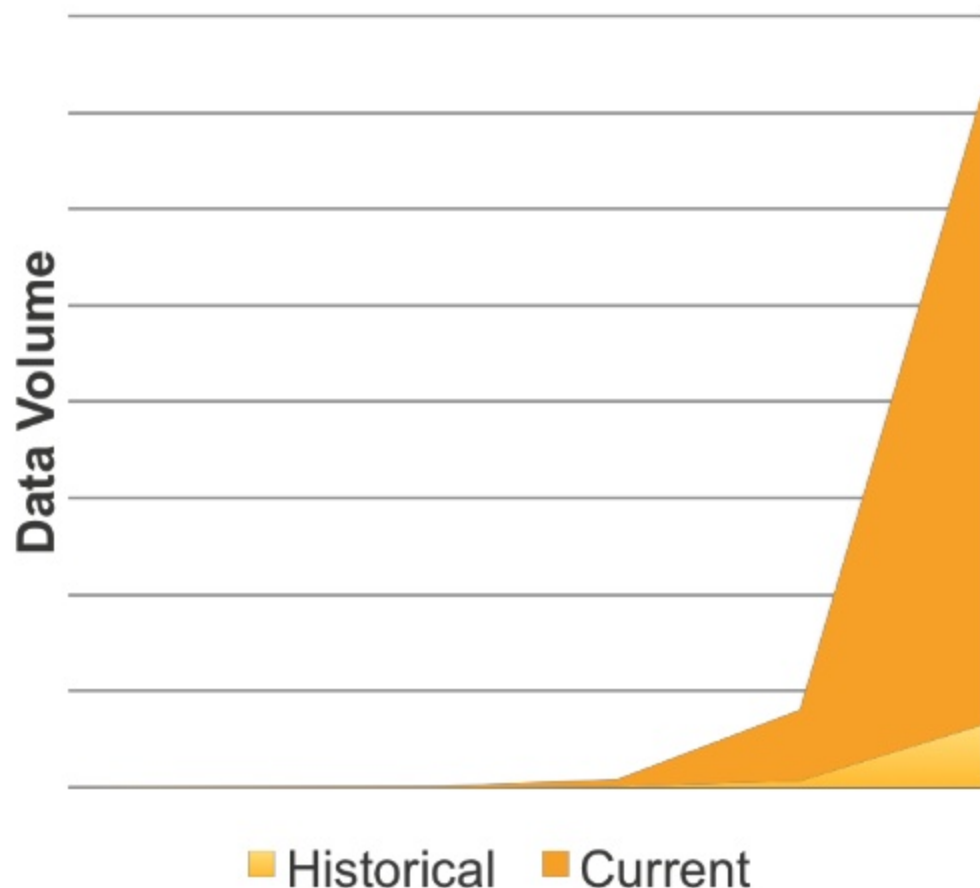
Pop-up Loft

November 02, 2016

Agenda

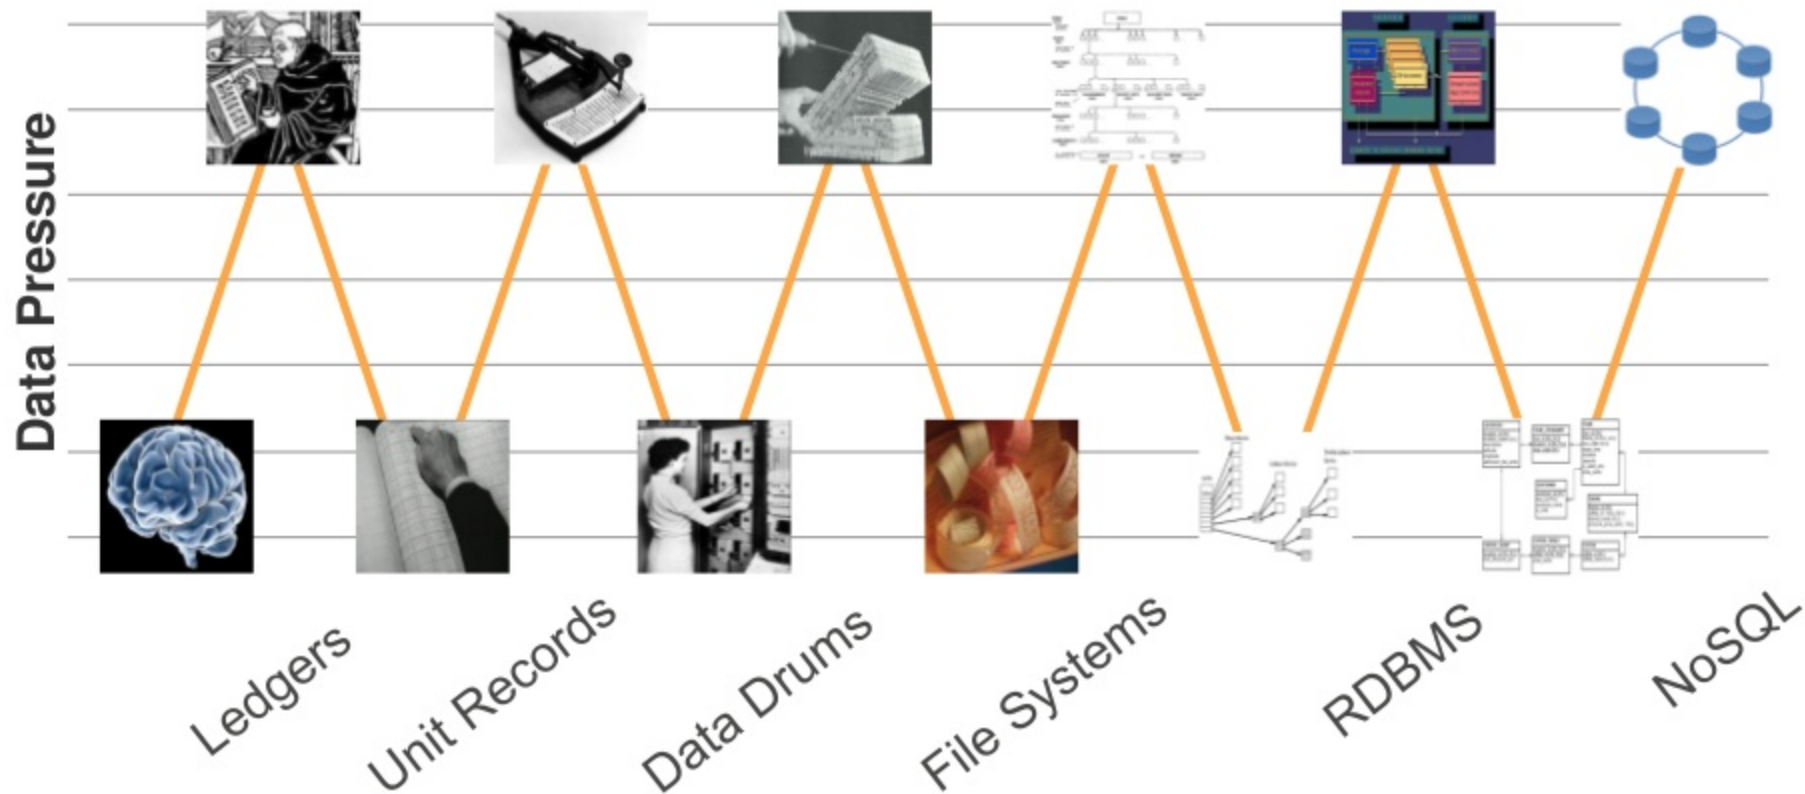
- Brief history of data processing
- Relational (SQL) vs. nonrelational (NoSQL)
- Fully managed features of DynamoDB
- Customer use cases
- Demo – serverless applications
- Pricing and Free Tier

Data volume since 2010




- 90% of stored data generated in last 2 years
- 1 terabyte of data in 2010 equals 6.5 petabytes today
- Linear correlation between data pressure and technical innovation
- No reason these trends will not continue over time

Timeline of database technology




Technology adoption and the hype curve




Relational (SQL) vs. nonrelational (NoSQL)

Relational vs. nonrelational databases

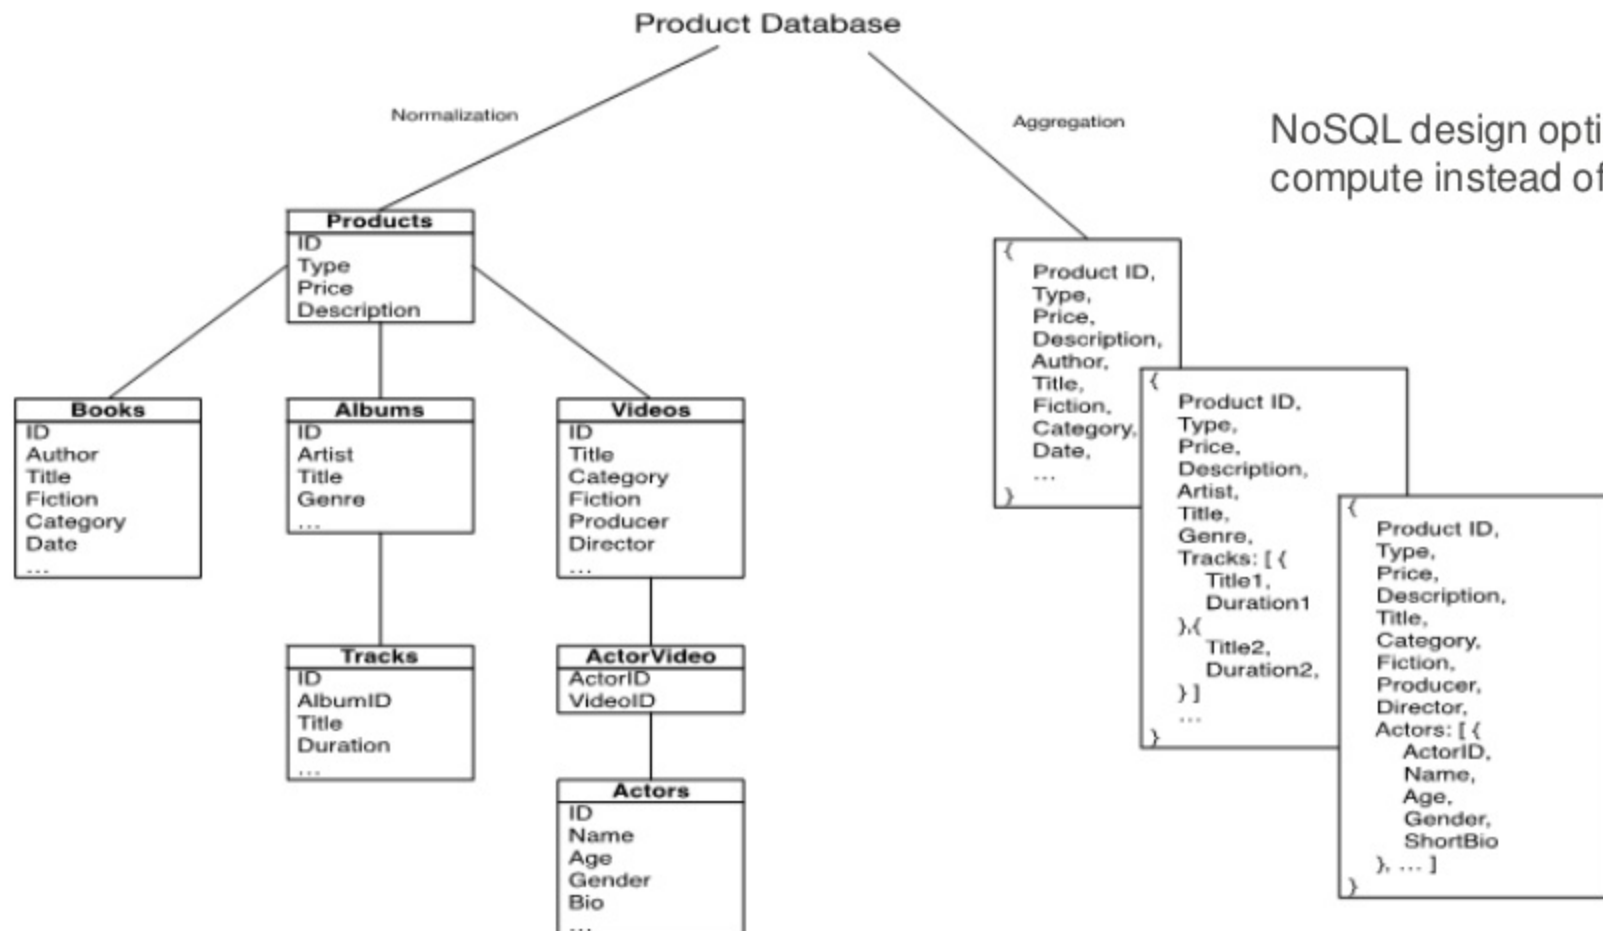
Traditional SQL



NoSQL



SQL vs. NoSQL schema design



Why NoSQL?

SQL

NoSQL

Optimized for storage	Optimized for compute
Normalized/relational	Denormalized/hierarchical
Ad hoc queries	Instantiated views
Scale vertically	Scale horizontally
Good for OLAP	Built for OLTP at scale

Amazon DynamoDB

Run your business, not your database



DynamoDB Benefits



Fully managed



Fast, consistent performance



Highly scalable



Flexible

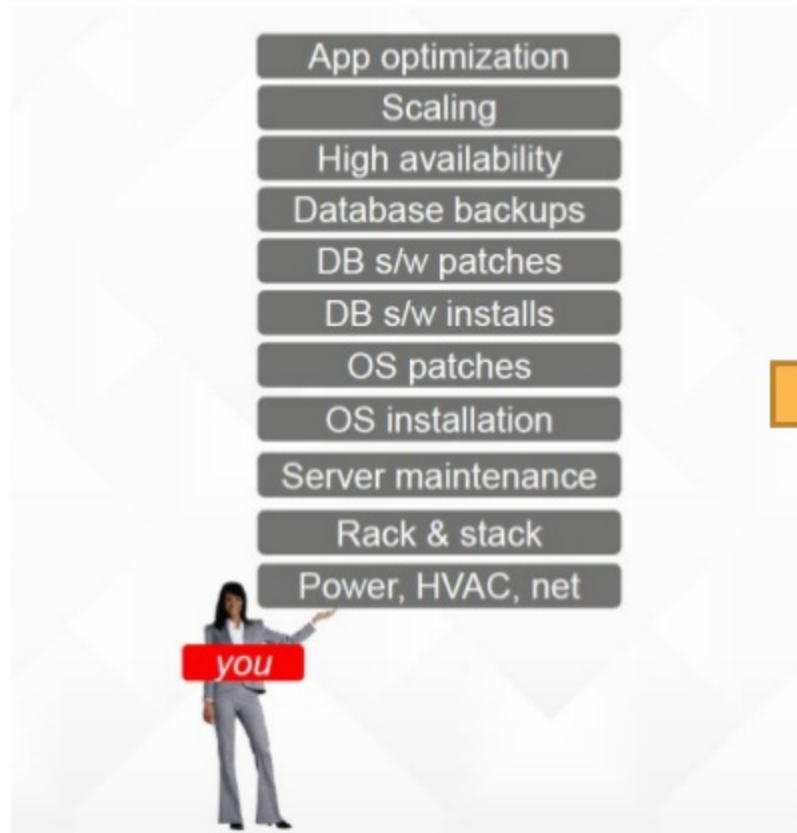


Event-driven programming

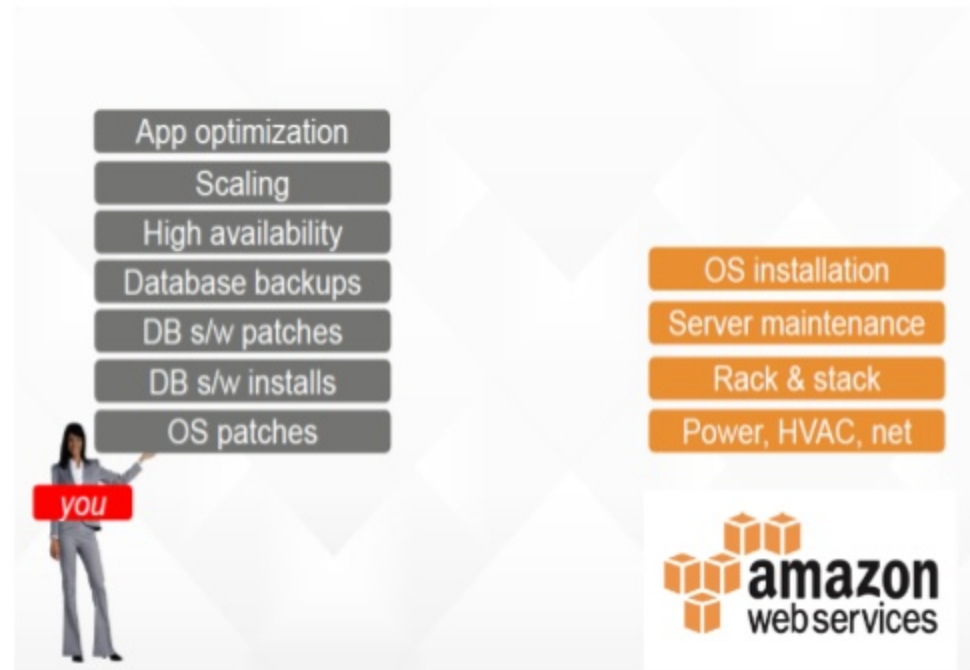


Fine-grained access control

Fully managed service = automated operations

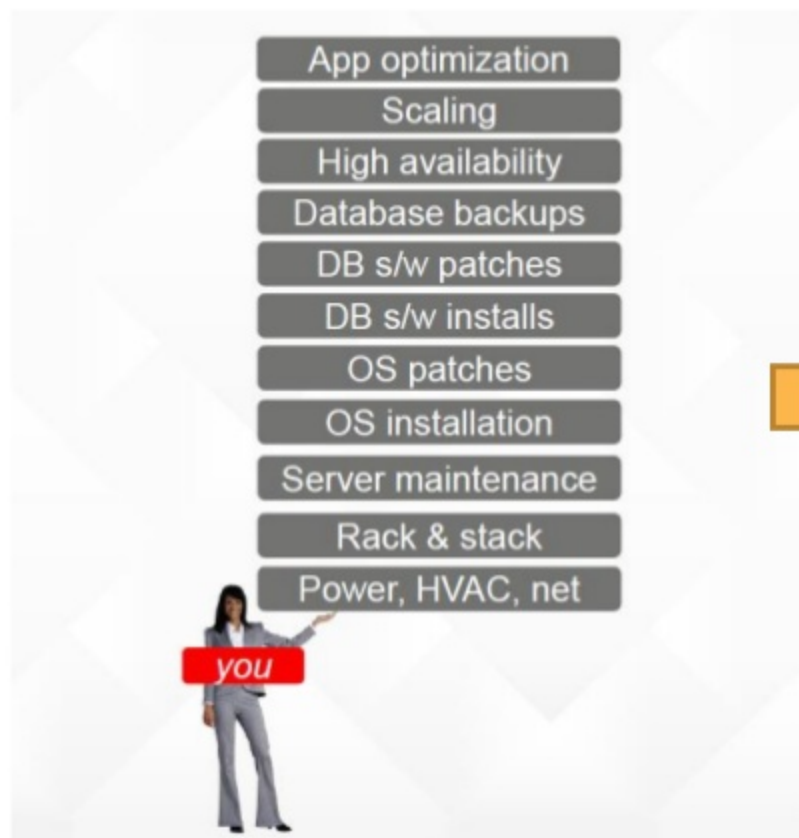


DB hosted on-premises



DB hosted on Amazon EC2

Fully managed service = automated operations




DB hosted on premise



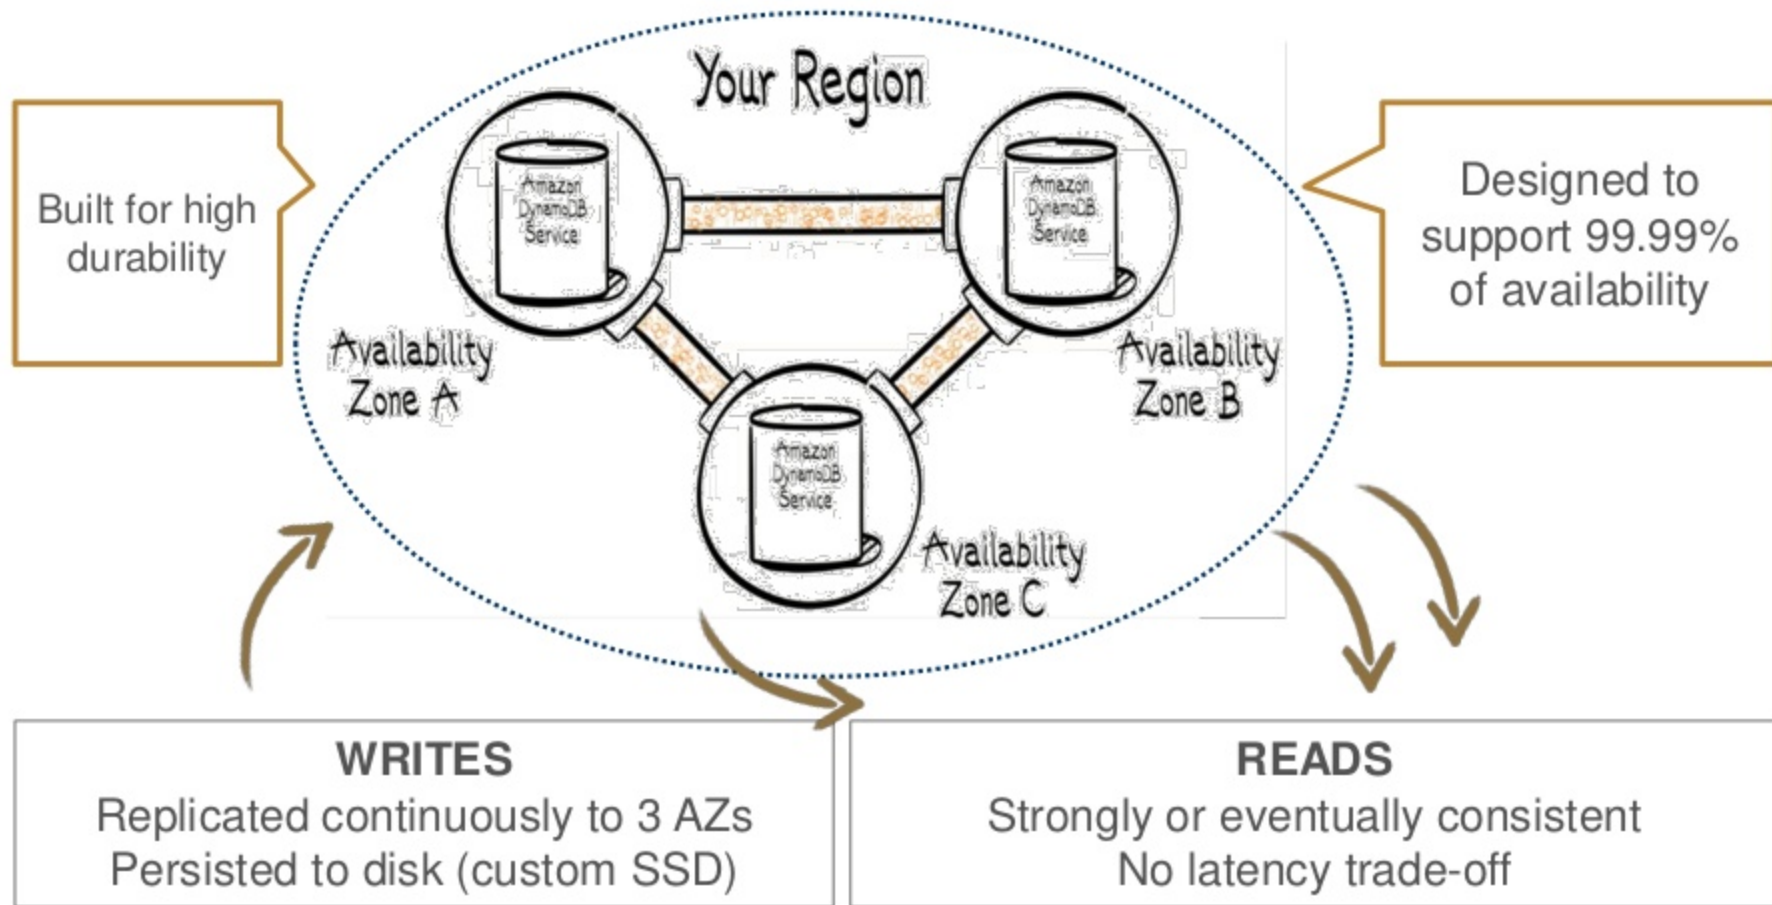
DynamoDB

Consistently low latency at scale




PREDICTABLE
PERFORMANCE!

High availability and durability



Customer use cases

Amazon's Path to DynamoDB



Dynamo: Amazon's Highly Available Key-value Store

Giuseppe DeCandia, Deniz Hastorun, Madan Jambani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Vosshall and Werner Vogels

Amazon.com

ABSTRACT

Reliability at massive scale is one of the biggest challenges we face at Amazon.com, one of the largest e-commerce operations in the world; even the slightest outage has significant financial consequences and impacts customer trust. The Amazon.com platform, which provides services for many web sites worldwide, is implemented on top of an infrastructure of tens of thousands of servers and network components located in many datacenters around the world. At this scale, small and large components fail continuously and the way persistent state is managed in the face of these failures drives the reliability and scalability of the software systems.

This paper presents the design and implementation of Dynamo, a highly available key-value storage system that some of Amazon's core services use to provide an "always-on" experience. To achieve this level of availability, Dynamo sacrifices consistency under certain failure scenarios. It makes extensive use of object versioning and application-assisted conflict resolution in a manner that provides a novel interface for developers to use.

One of the lessons our organization has learned from operating Amazon's platform is that the reliability and scalability of a system is dependent on how its application state is managed. Amazon uses a highly decentralized, loosely coupled, service oriented architecture consisting of hundreds of services. In this environment there is a particular need for storage technologies that are always available. For example, customers should be able to view and add items to their shopping cart even if disks are failing, network routes are flapping, or data centers are being destroyed by tornadoes. Therefore, the service responsible for managing shopping carts requires that it can always write to and read from its data store, and that its data needs to be available across multiple data centers.

Dealing with failures in an infrastructure composed of millions of components in our standard mode of operation; there are always a small but significant number of server and network components that are failing at any given time. As such Amazon's software systems need to be constructed in a manner that treats failure handling as the normal case without impacting availability or performance.

Major League Baseball Fields Big Data, Excitement with Amazon DynamoDB

“

For the first time, we can measure things we've never been able to measure before.

Joe Inzerillo

Executive Vice President and CTO, MLBAM



”

- MLBAM can scale to support many games on a single day.
- Amazon DynamoDB powers queries and supports the fast data retrieval required.
- MLBAM distributes 25,000 live events annually and 10 million streams daily.

MLBAM (MLB Advanced Media) is a full service solutions provider, operating a powerful content delivery platform.

Redfin Is Revolutionizing Home Buying and Selling with Amazon DynamoDB



We have billions of records on DynamoDB being refreshed daily or hourly or even by seconds.

Yong Huang
Director, Big Data Analytics, Redfin

REDFIN.



- Redfin provides property and agent details and ratings through its websites and apps.
- With DynamoDB, latency for “similar” properties improved from 2 seconds to just 12 milliseconds.
- Redfin stores and processes five billion items in DynamoDB.

Redfin is a full-service real estate company with local agents and online tools to help people buy & sell homes.

Expedia's Real-time Analytics Application Uses Amazon DynamoDB



With DynamoDB we were up and running in a less than day, and there is no need for a team to maintain.

Kuldeep Chowhan
Principal Engineer, Expedia



- Expedia's real-time analytics application collects data for its "test & learn" experiments on Expedia sites.
- The analytics application processes ~200 million messages daily.
- Ease of setup, monitoring, and scaling were key factors in choosing Amazon DynamoDB.

Expedia is a leader in the \$1 trillion travel industry, with an extensive portfolio that includes some of the world's most trusted travel brands.

Nexon Scales Mobile Gaming with Amazon DynamoDB

“

By using AWS, we decreased our initial investment costs, and only pay for what we use.

Chunghoon Ryu
Department Manager, Nexon



”

- Nexon used Amazon DynamoDB as its primary game database for a new blockbuster mobile game, HIT
- HIT became the #1 Mobile Game in Korea within the first day of launch and has > 2M registered users
- Nexon's HIT leverages DynamoDB to deliver steady latency of less than 10ms to deliver a fantastic mobile gaming experience for 170,000 concurrent players

Nexon is a leading South Korean video game developer and a pioneer in the world of interactive entertainment.

Scaling high-velocity use cases with DynamoDB

Ad Tech

AdRoll

DataXu

ADBRAIN

doapp
we do cool stuff

VidRoll

Gaming

SUPERCELL

zynga

NEXON

BATTLE CAMP

FRONTIER

IoT

mlbam

ACTi
Connecting Vision

canary

dropcam

MEDIATEK

Mobile

duolingo

Mapbox

REDFIN

remind

INFRAWARE

Web

Expedia

Nordstrom

NY Summit 2015

JustGiving

jobandtalent

CREATIVE LIFE STORE
TOKYU HANDS

That sounds really good. How do I get started?

Let's create a table..



Create DynamoDB table

Tutorial



Table name*

Products



Primary key*

Partition key

Product_Id

String



☐ Add sort key

Table settings

Default settings provide the fastest way to get started with your table. You can modify these default settings now or after your table has been created.

☒ Use default settings

- No secondary indexes
- Provisioned capacity set to 5 reads and 5 writes.
- Basic alarms with 80% upper threshold using SNS topic "dynamodb".

Additional charges may apply if you exceed the AWS Free Tier levels for CloudWatch or Simple Notification Service. Advanced alarm settings are available in the CloudWatch management console.

Cancel

Create



DynamoDB

Dashboard

Tables

Reserved capacity

Products [Close](#)



Overview

Items

Metrics

Alarms

Capacity

Indexes

Triggers

Access control

Create item

Actions



Scan: [Table] Products: Product_id

Viewing 1 to 3 items

Scan


[Table] Products: Product_id

+ Add filter

Start search

<input type="checkbox"/>	Product_id	Author	Book_Title	Product_name	Artist	Movie_Name
<input type="checkbox"/>	443			Movie	Clark Kent	Batman vs Superman
<input type="checkbox"/>	234			Movie	Tony Stark	Avengers
<input type="checkbox"/>	123	Douglas Pr...	Relic	Books		

DynamoDB table structure



Partition key Sort key

Mandatory
Key-value access pattern
Determines data distribution

Optional
Model 1:N relationships
Enables rich query capabilities

All items for key
==, <, >, >=, <=
"begins with"
"between"
"contains"
"in"
sorted results
counts
top/bottom N values

Local secondary index (LSI)

Alternate sort key attribute


Index is local to a partition key

Table



10 GB maximum per partition key; LSIs limit the number of range keys!

LSIs



Global secondary index (GSI)

Alternate partition and/or sort key

Index is across all partition keys

Online indexing

Read capacity units (RCUs) and write capacity units (WCUs) are provisioned separately for GSIs

Table



GSIs



KEYS_ONLY

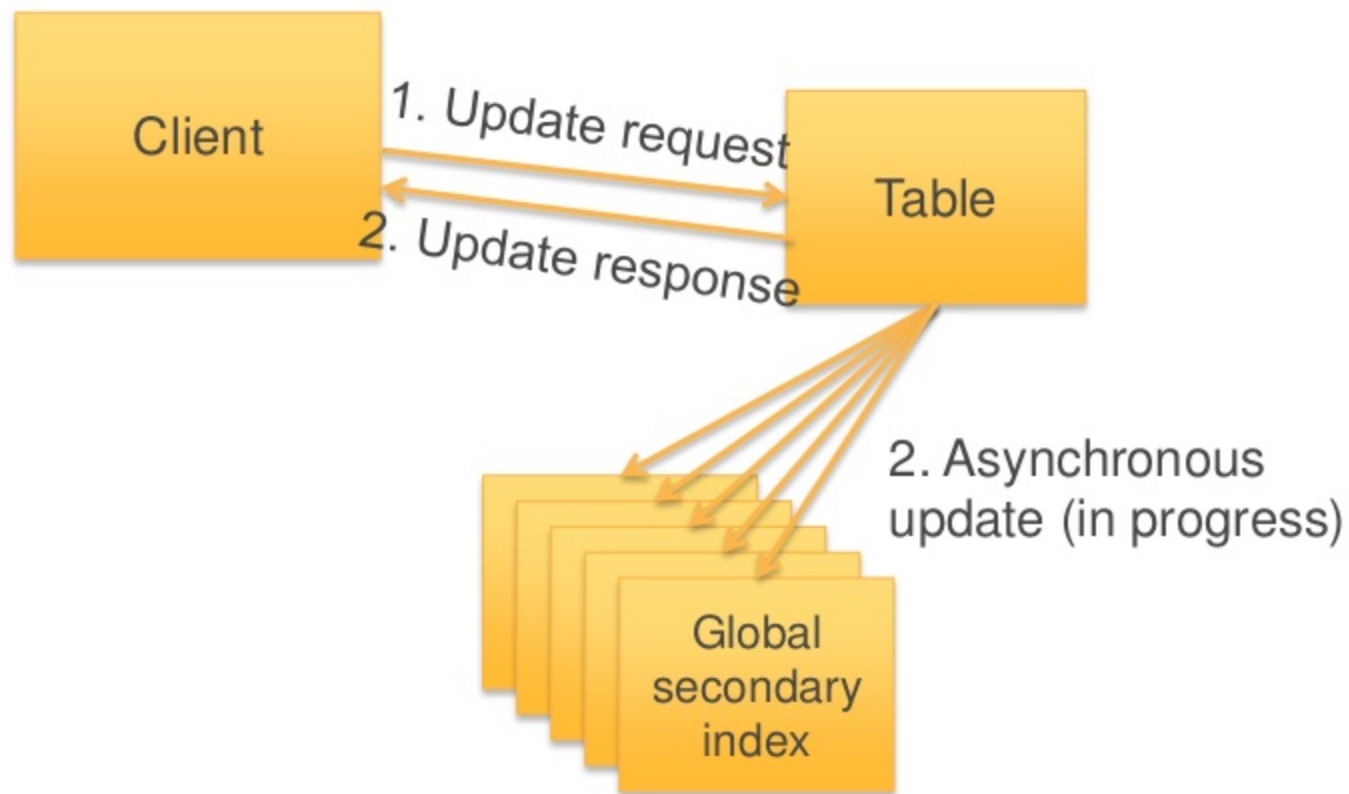


INCLUDE A3



ALL

How do GSI updates work?



If GSIs don't have enough write capacity, table writes will be throttled!

LSI or GSI?

LSI can be modeled as a GSI

If data size in an item collection > 10 GB, use GSI

If eventual consistency is okay for your scenario, use GSI!

Advanced topics in DynamoDB

- Data modeling
- Understanding Partitions
 - # of partitions depend on table throughput and size
- DynamoDB Scaling
- Design patterns and best practices

To learn more, please attend:

Deep Dive on Amazon DynamoDB

Sean Shriver, AWS NoSQL Solutions Architect

Hands on 90 minute lab (Serverless Web apps)

Sam Elmanak & Scott Kellish

Amazon DynamoDB: <https://aws.amazon.com/dynamodb/>

Demo

**Serverless Web Apps with
Amazon DynamoDB, API
Gateway, and AWS Lambda**

Simple serverless web application – use case

Today's Super Hero Mission

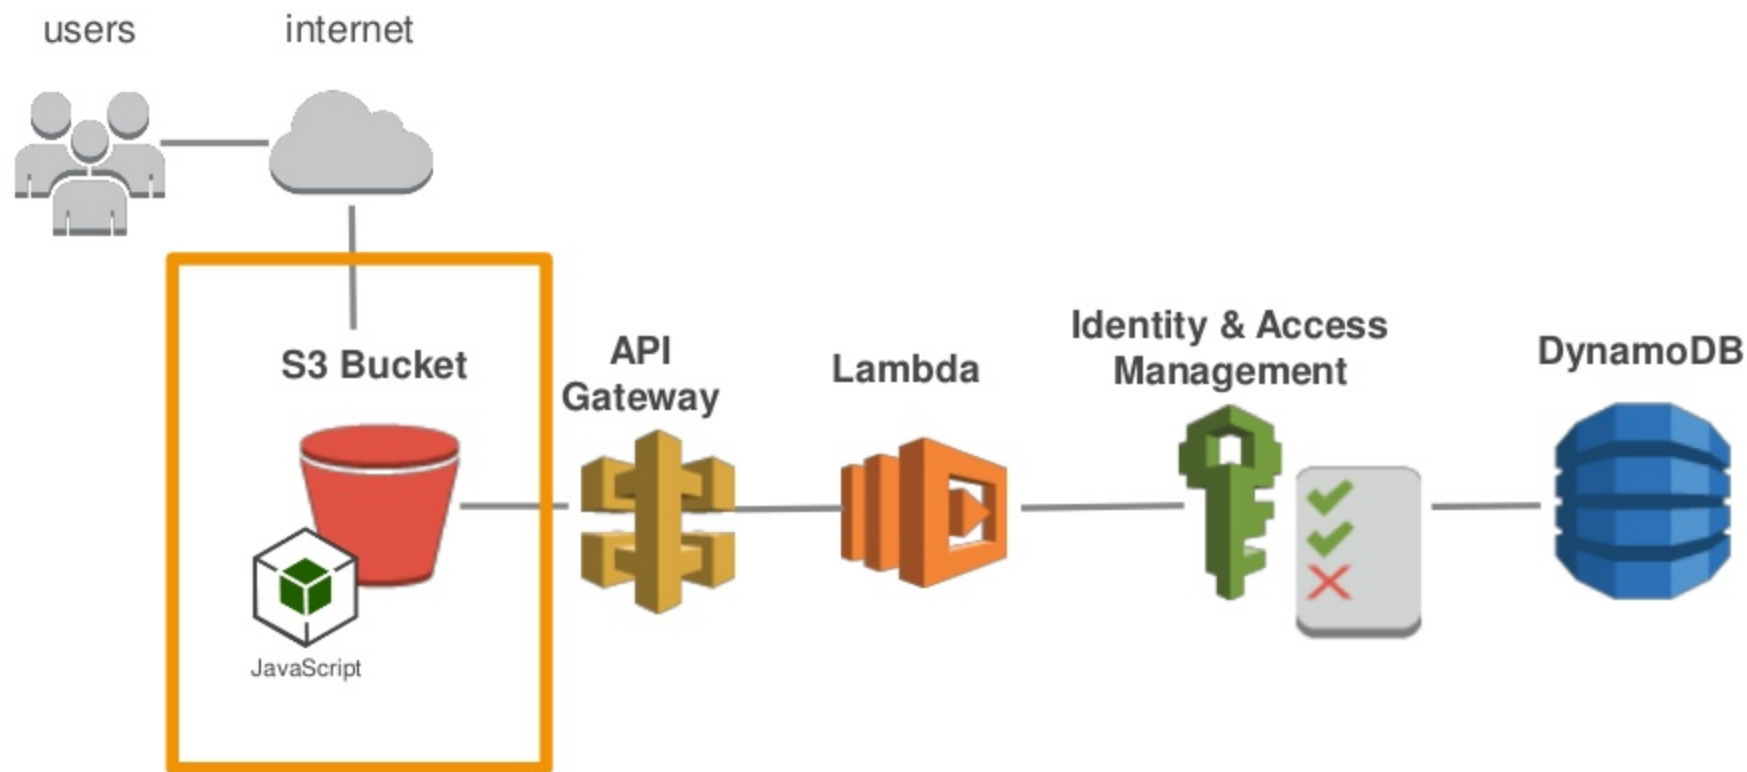
What's your Super Hero name? ▼

What's your Super Hero name?
Avengers
Batman
The Winchester Brothers
Superman

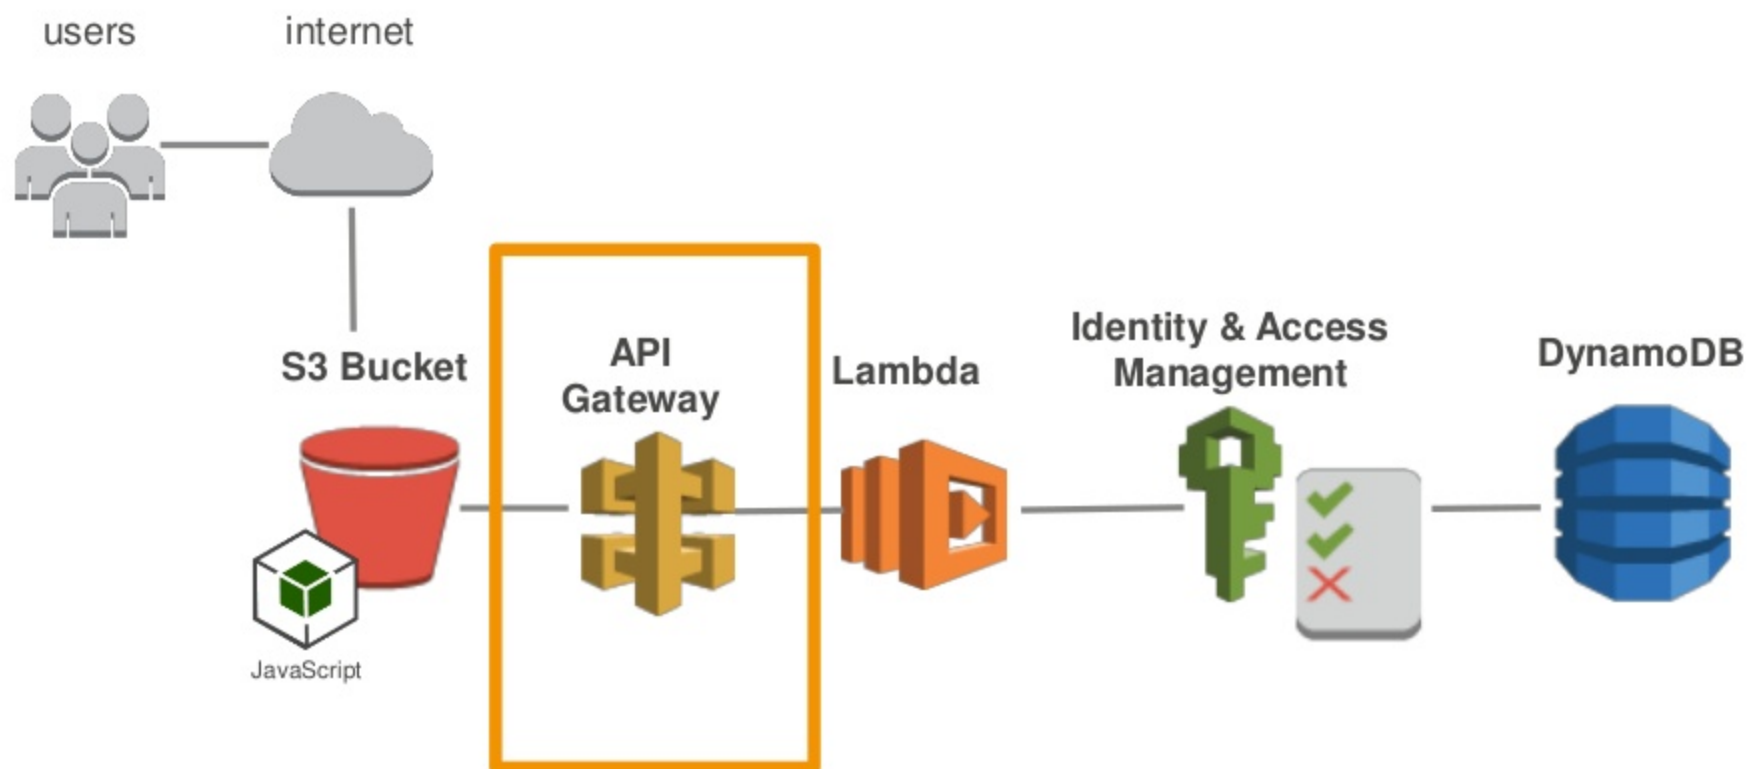
GO

Mission Dossier

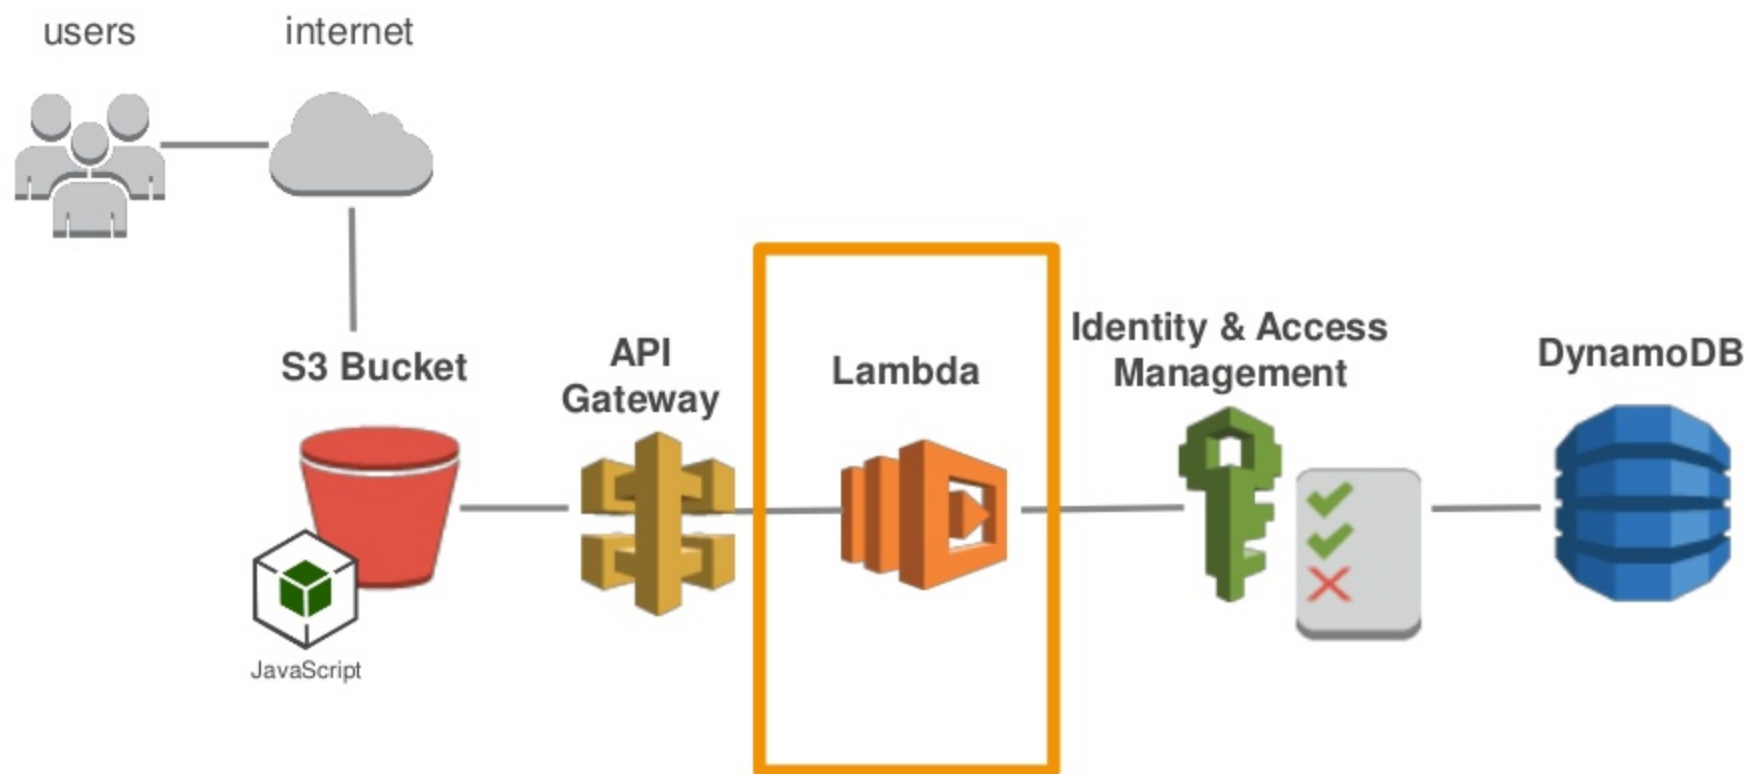
Architecture of a simple serverless web application



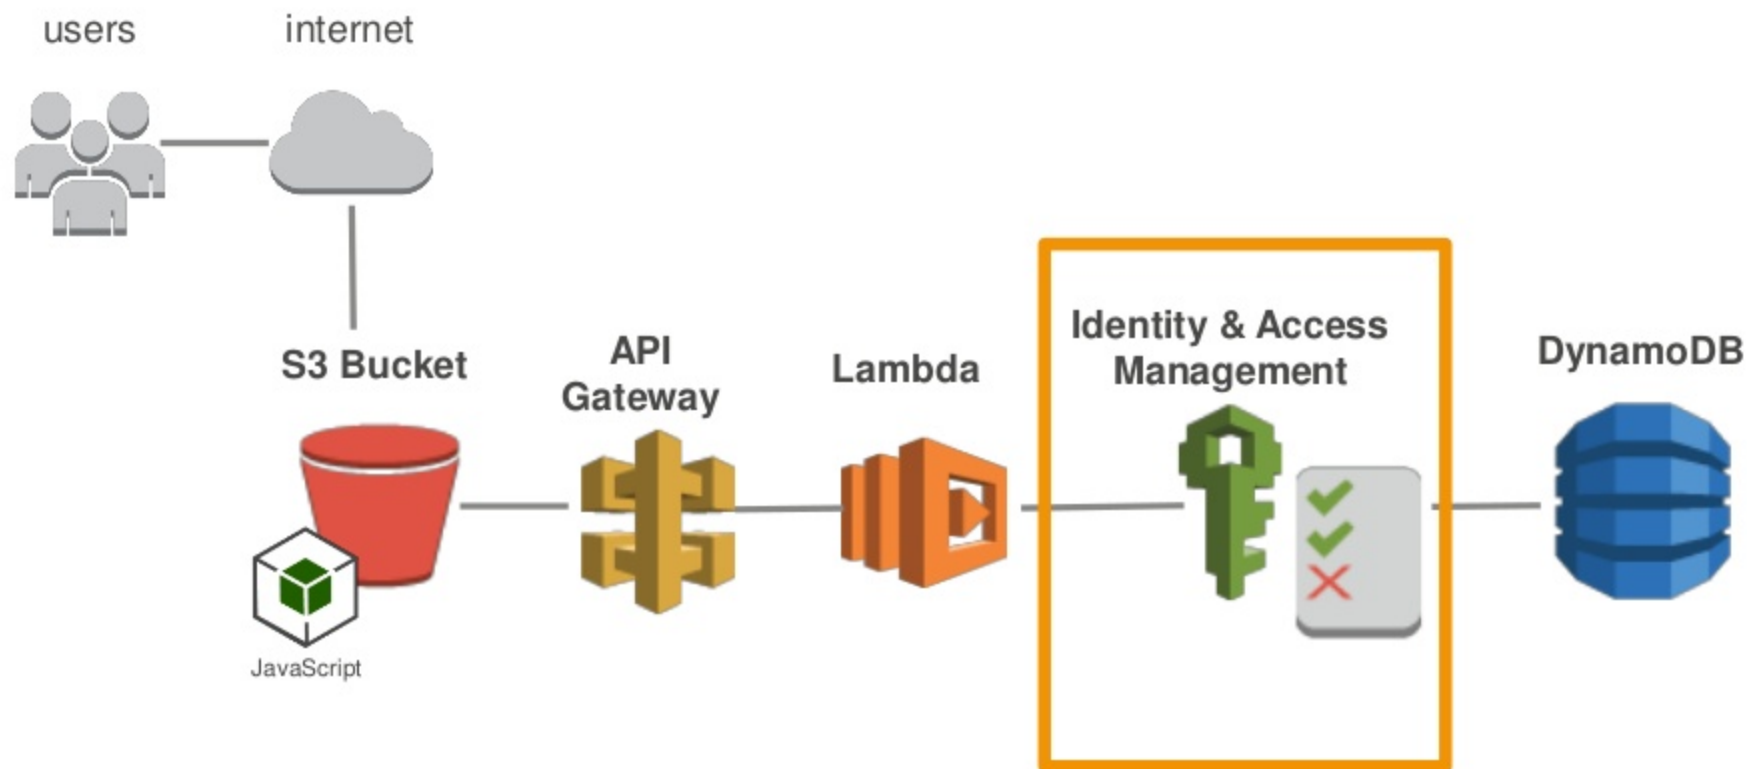
Architecture of a simple serverless web application



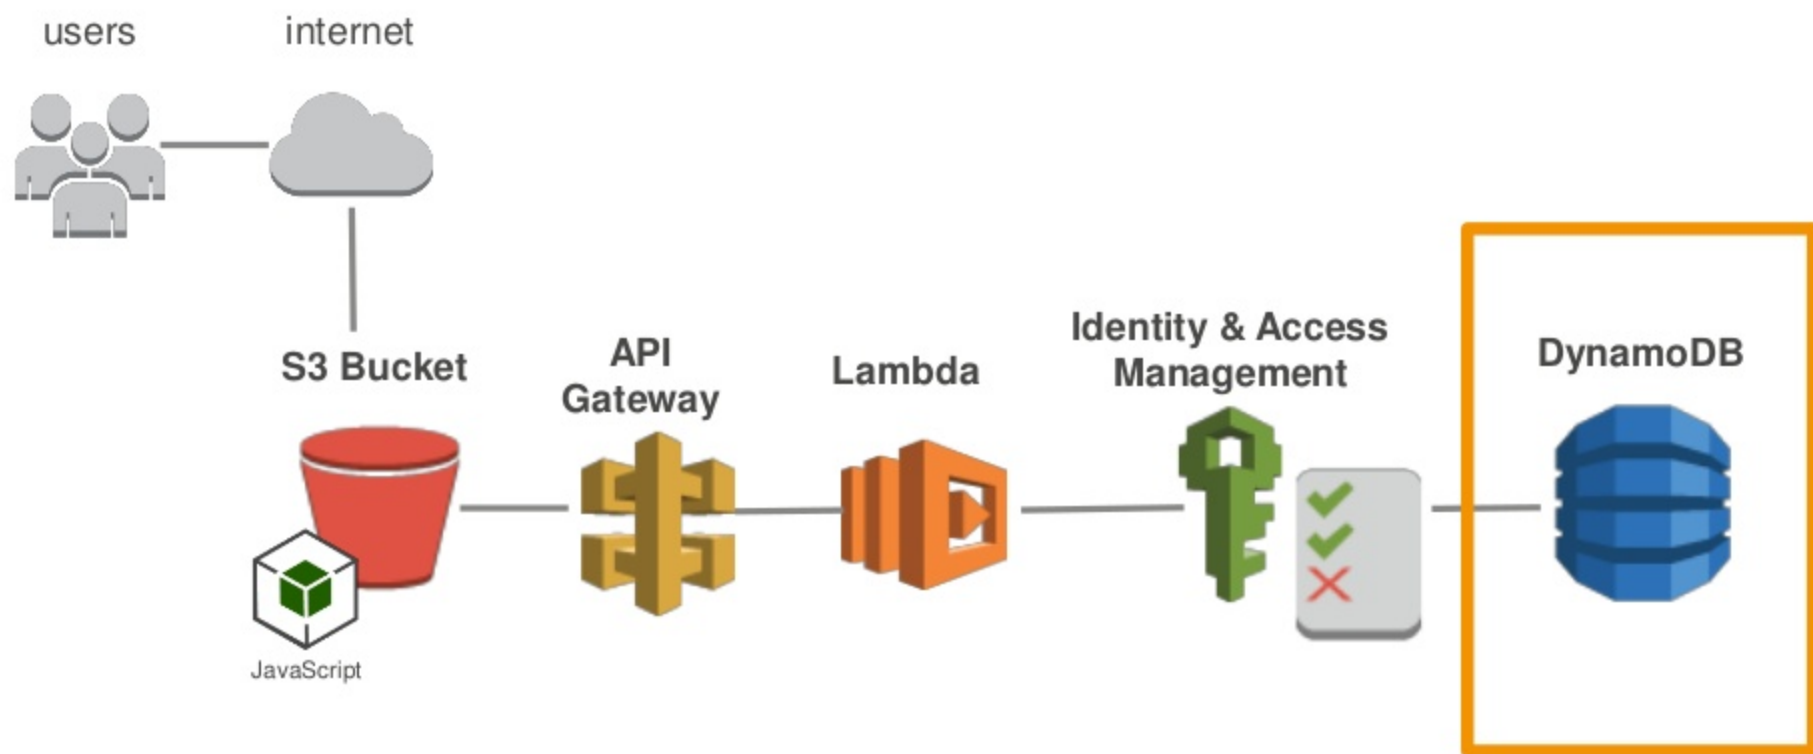
Architecture of a simple serverless web application



Architecture of a simple serverless web application



Architecture of a simple serverless web application



Demo

Amazon Web Services

Compute

- EC2**
Virtual Servers in the Cloud
- EC2 Container Service**
Run and Manage Docker Containers
- Elastic Beanstalk**
Run and Manage Web Apps
- Lambda**
Run Code in Response to Events

Storage & Content Delivery

- S3**
Scalable Storage in the Cloud
- CloudFront**
Global Content Delivery Network
- Elastic File System**
Fully Managed File System for EC2
- Glacier**
Archive Storage in the Cloud
- Snowball**
Large Scale Data Transport
- Storage Gateway**
Hybrid Storage Integration

Database

- RDS**
Managed Relational Database Service
- DynamoDB**
Managed NoSQL Database
- ElastiCache**
In-Memory Cache
- Redshift**
Fast, Simple, Cost-Effective Data Warehousing
- DMS**
Managed Database Migration Service

Networking

- VPC**
Isolated Cloud Resources
- Direct Connect**

Developer Tools

- CodeCommit**
Store Code in Private Git Repositories
- CodeDeploy**
Automate Code Deployments
- CodePipeline**
Release Software using Continuous Delivery

Management Tools

- CloudWatch**
Monitor Resources and Applications
- CloudFormation**
Create and Manage Resources with Templates
- CloudTrail**
Track User Activity and API Usage
- Config**
Track Resource Inventory and Changes
- OpsWorks**
Automate Operations with Chef
- Service Catalog**
Create and Use Standardized Products
- Trusted Advisor**
Optimize Performance and Security

Security & Identity

- Identity & Access Management**
Manage User Access and Encryption Keys
- Directory Service**
Host and Manage Active Directory
- Inspector**
Analyze Application Security
- WAF**
Filter Malicious Web Traffic
- Certificate Manager**
Provision, Manage, and Deploy SSL/TLS Certificates

Analytics

- EMR**
Managed Hadoop Framework
- Data Pipeline**
Data-Driven Workflows

Internet of Things

- AWS IoT**
Connect Devices to the Cloud

Game Development

- GameLift**
Deploy and Scale Session-based Multiplayer Games

Mobile Services

- Mobile Hub**
Build, Test, and Monitor Mobile Apps
- Cognito**
User Identity and App Data Synchronization
- Device Farm**
Test Android, iOS, and Web Apps on Real Devices in the Cloud
- Mobile Analytics**
Collect, View and Export App Analytics
- SNS**
Push Notification Service

Application Services

- API Gateway**
Build, Deploy and Manage APIs
- AppStream**
Low Latency Application Streaming
- CloudSearch**
Managed Search Service
- Elastic Transcoder**
Easy-to-Use Scalable Media Transcoding
- SES**
Email Sending and Receiving Service
- SQS**
Message Queue Service
- SWF**
Workflow Service for Coordinating Application Components

Enterprise Applications

- WorkSpaces**
Desktops in the Cloud
- WorkDocs**
Secure Enterprise Storage and Sharing Service

Resource Groups

[Learn more](#)

A resource group is a collection of resources that share one or more tags. Create a group for each project, application, or environment in your account.

[Create a Group](#)

[Tag Editor](#)

Additional Resources

Getting Started

[Read our documentation](#) or [view our training](#) to learn more about AWS.

AWS Console Mobile App

View your resources on the go with our AWS Console mobile app, available from Amazon Appstore, Google Play, or iTunes.

AWS Marketplace

Find and buy software, launch with 1-Click and pay by the hour.

AWS re:Invent Announcements

Explore the next generation of AWS cloud capabilities. [See what's new](#)

Service Health

All services operating normally.

Updated: Jul 15 2016 11:15:00 GMT-0700

[Service Health Dashboard](#)

DynamoDB Pricing & Free Tier



- Free Tier
 - ☐ 25GB of storage
 - ☐ 25 Reads per second
 - ☐ 25 Writes per second
- Pricing for additional usage in US East (N. Virginia)
 - ☐ \$0.25 per GB per month
 - ☐ Write throughput: \$0.0065 per hour for every 10 units of Write Capacity
 - ☐ Read throughput: \$0.0065 per hour for every 50 units of Read Capacity

THANK YOU!

Don't forget...

- **SURVEY:** emailed to you within 12 hours. We value your feedback!
- **SLIDES:** links will display on screen when you complete the survey
- **MORE QUESTIONS?** Visit the Ask an Architect bar on the 1st Floor





Pop-up Loft

Everything and Anything Startups Need to Get Started on AWS

aws.amazon.com/activate