

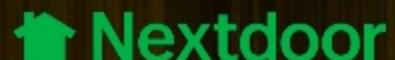
SRV319

How Nextdoor Built a Scalable, Serverless Data Pipeline for Billions of Events per Day

Prakash Janakiraman · Co-Founder & Chief Architect

Matt Wise · Sr. Systems Architect

Slava Markeyev · Software Engineer



November 29, 2017

Speakers



Prakash Janakiraman
Co-Founder & Chief Architect



Matt Wise
Sr. Systems Architect



Slava Markeyev
Systems Engineer

What to Expect from the Session



- About Nextdoor
- Overview of our legacy pipeline
- Requirements for new pipeline
- Recap of AWS Serverless offerings
- Tool we built to help standardize Serverless ETL
- Deployment example via Cloudformation
- Q&A

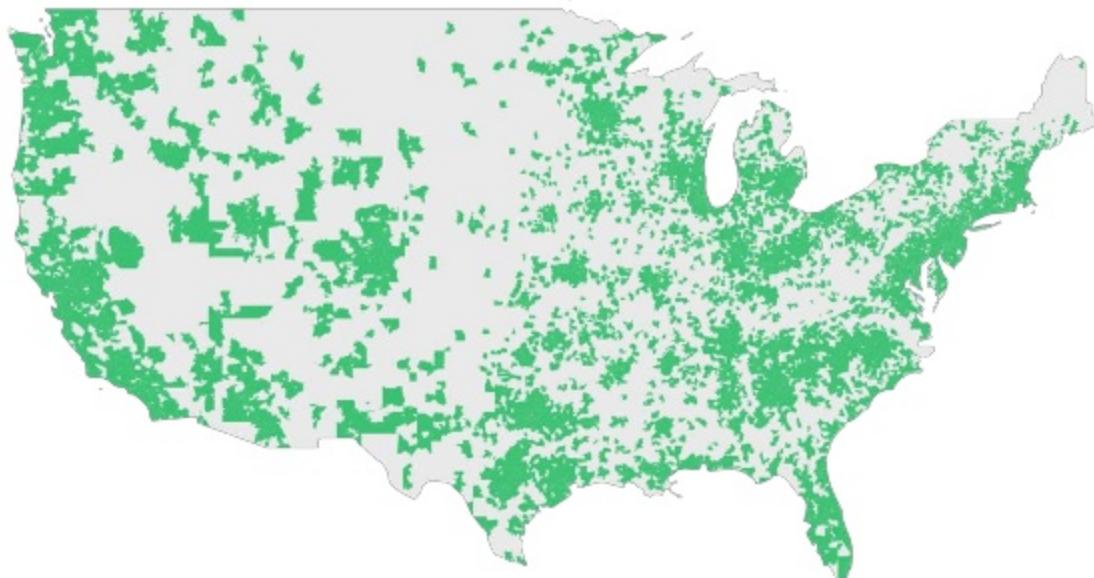


AWS
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Used by 80% of US neighborhoods



Data drives products at Nextdoor



Central Fillmore

339 neighbors

Invite

18% of 1,269 households

20,034 neighbors in
19 nearby neighborhoods



P Progressive, Sponsored

Protect the home you love

With our all-new HomeQuote Explorer, it's easy to protect your home at a price that fits your budget.

As the only insurer that lets you compare home insurance rates and coverages from See more...

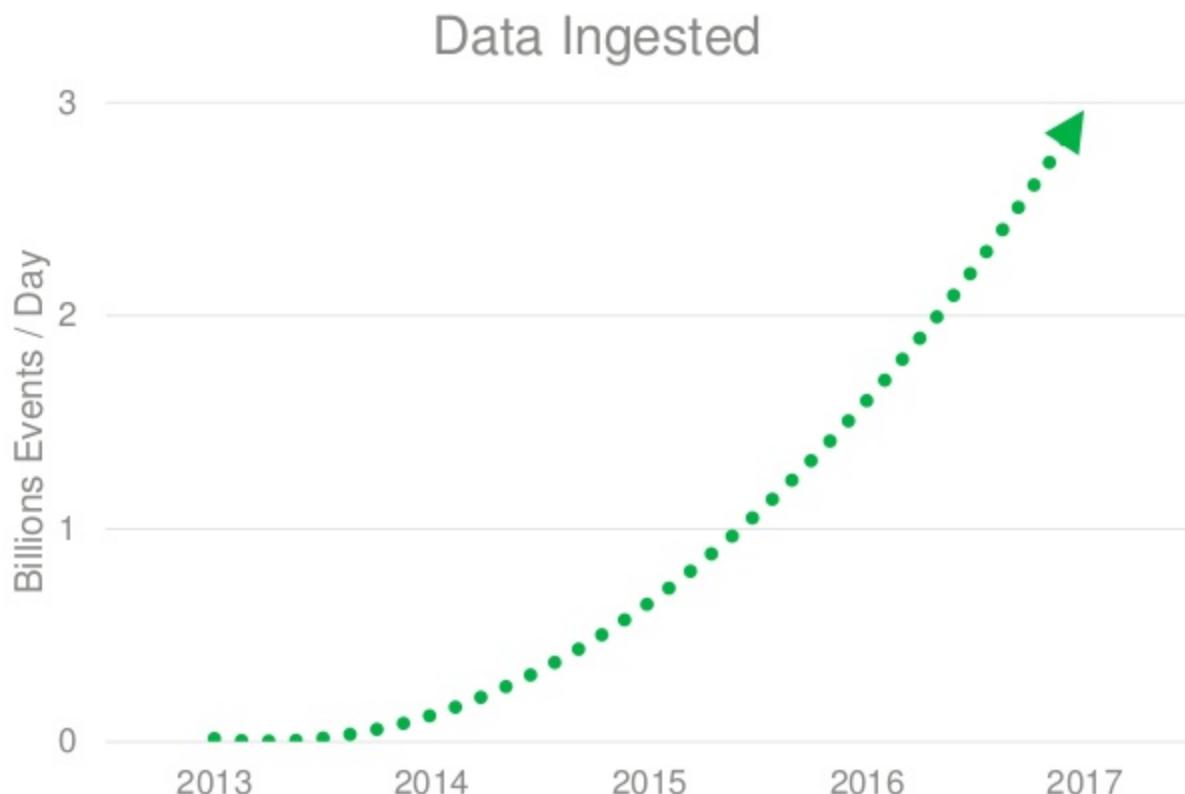


Compare rates easily with HomeQuote Explorer

Get quote

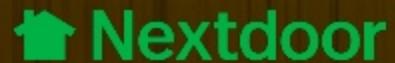
Thank Reply

Data at Nextdoor



Realtime Data Ingestion Pipeline

Slava Markeyev · Software Engineer



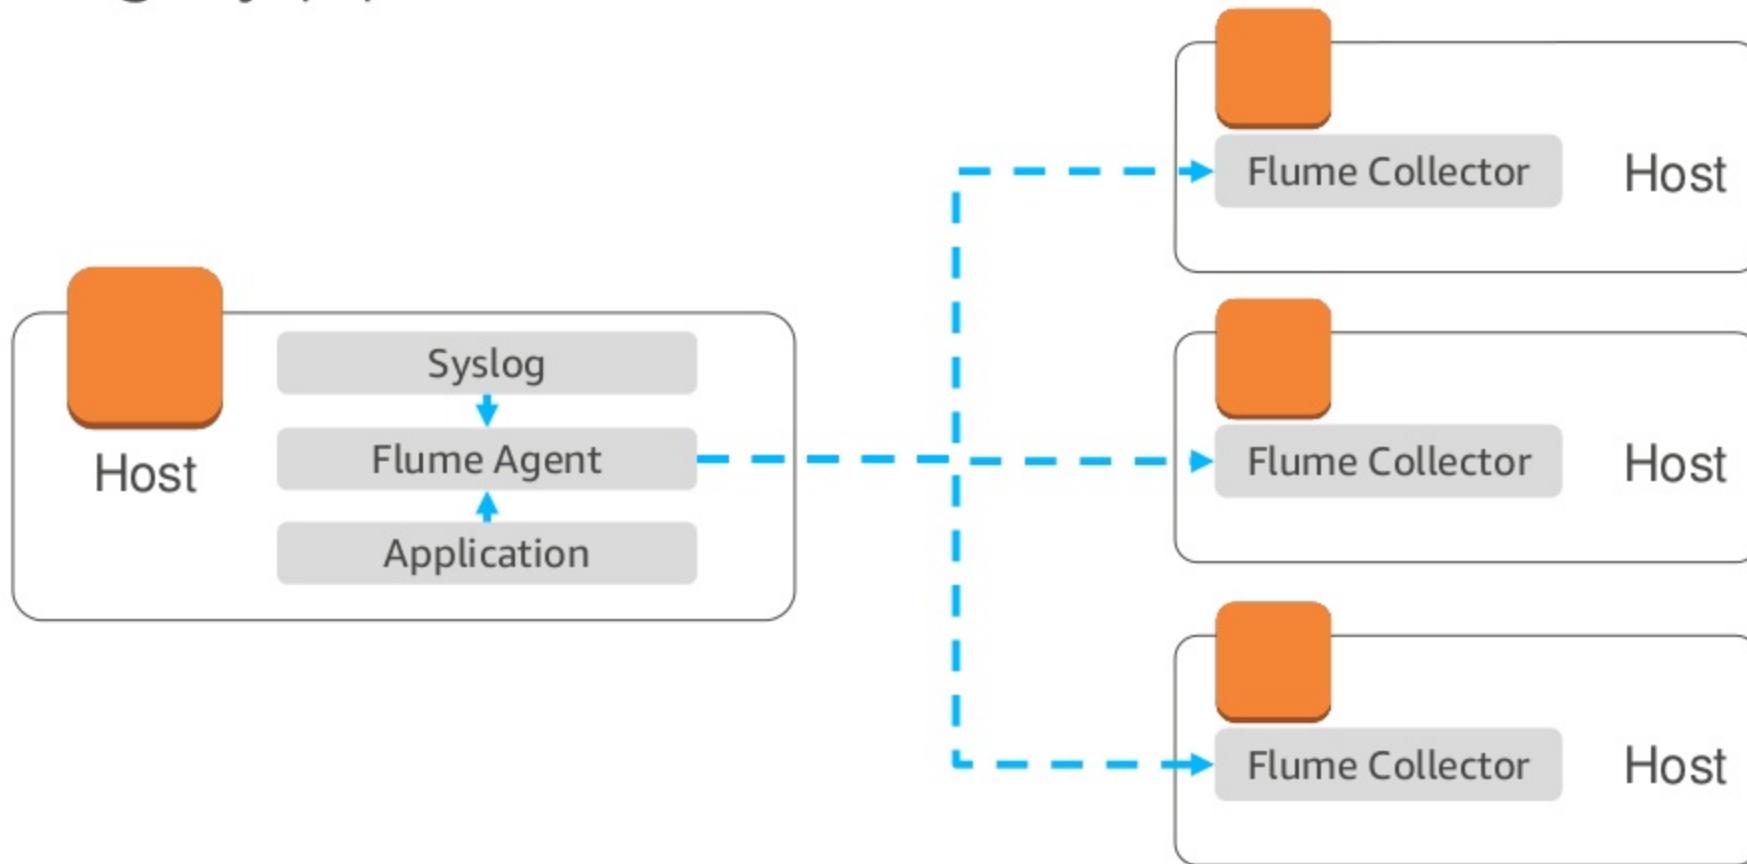
Legacy pipeline overview



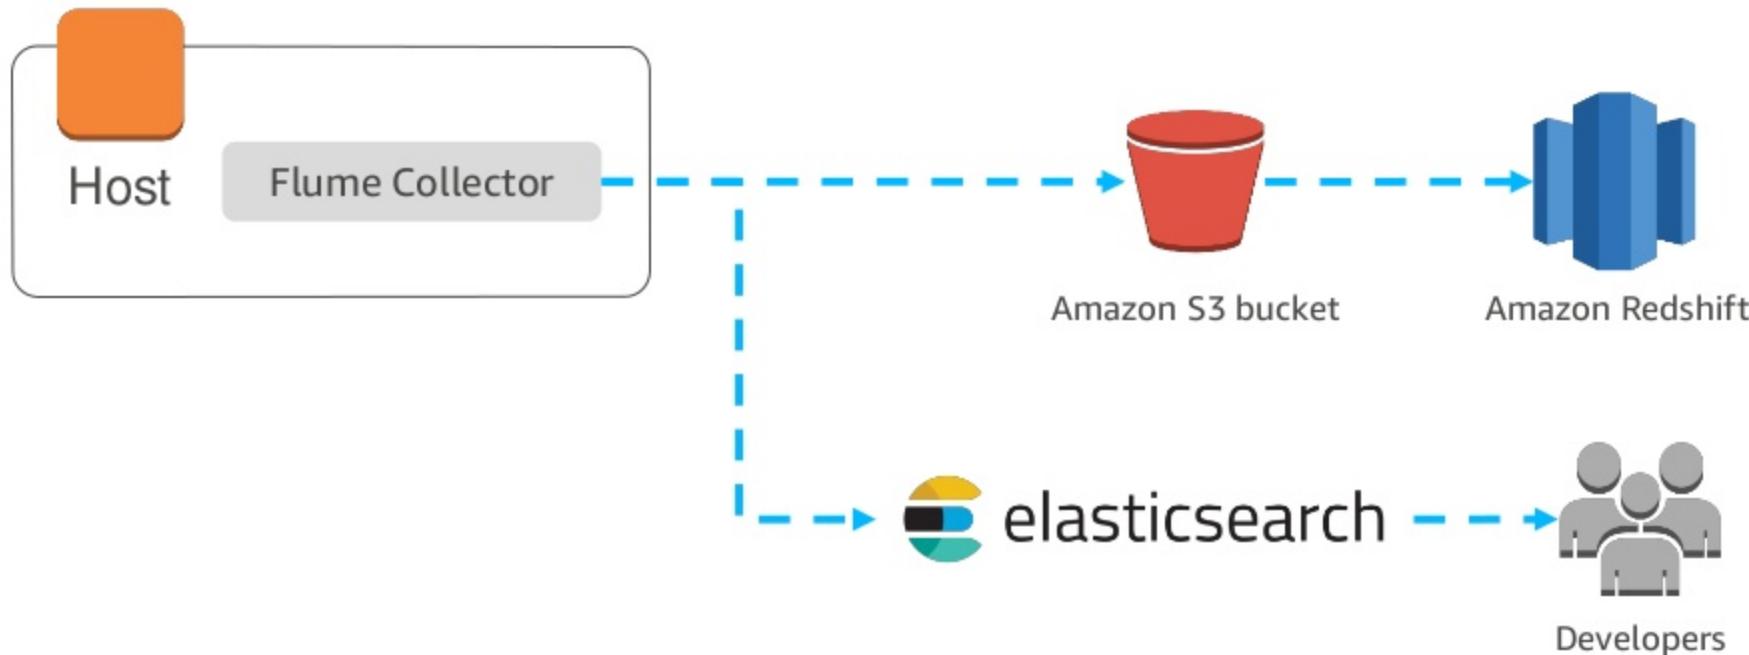
- Built using Apache Flume
- A distributed self-hosted service for collecting and aggregating log data (similar to AWS Kinesis Streams, Apache Kafka, and Fluentd)
- Served us well for 4 years
- Scaled to handle ~1.5B events per day



Legacy pipeline overview



Legacy pipeline overview



Legacy pipeline pain points



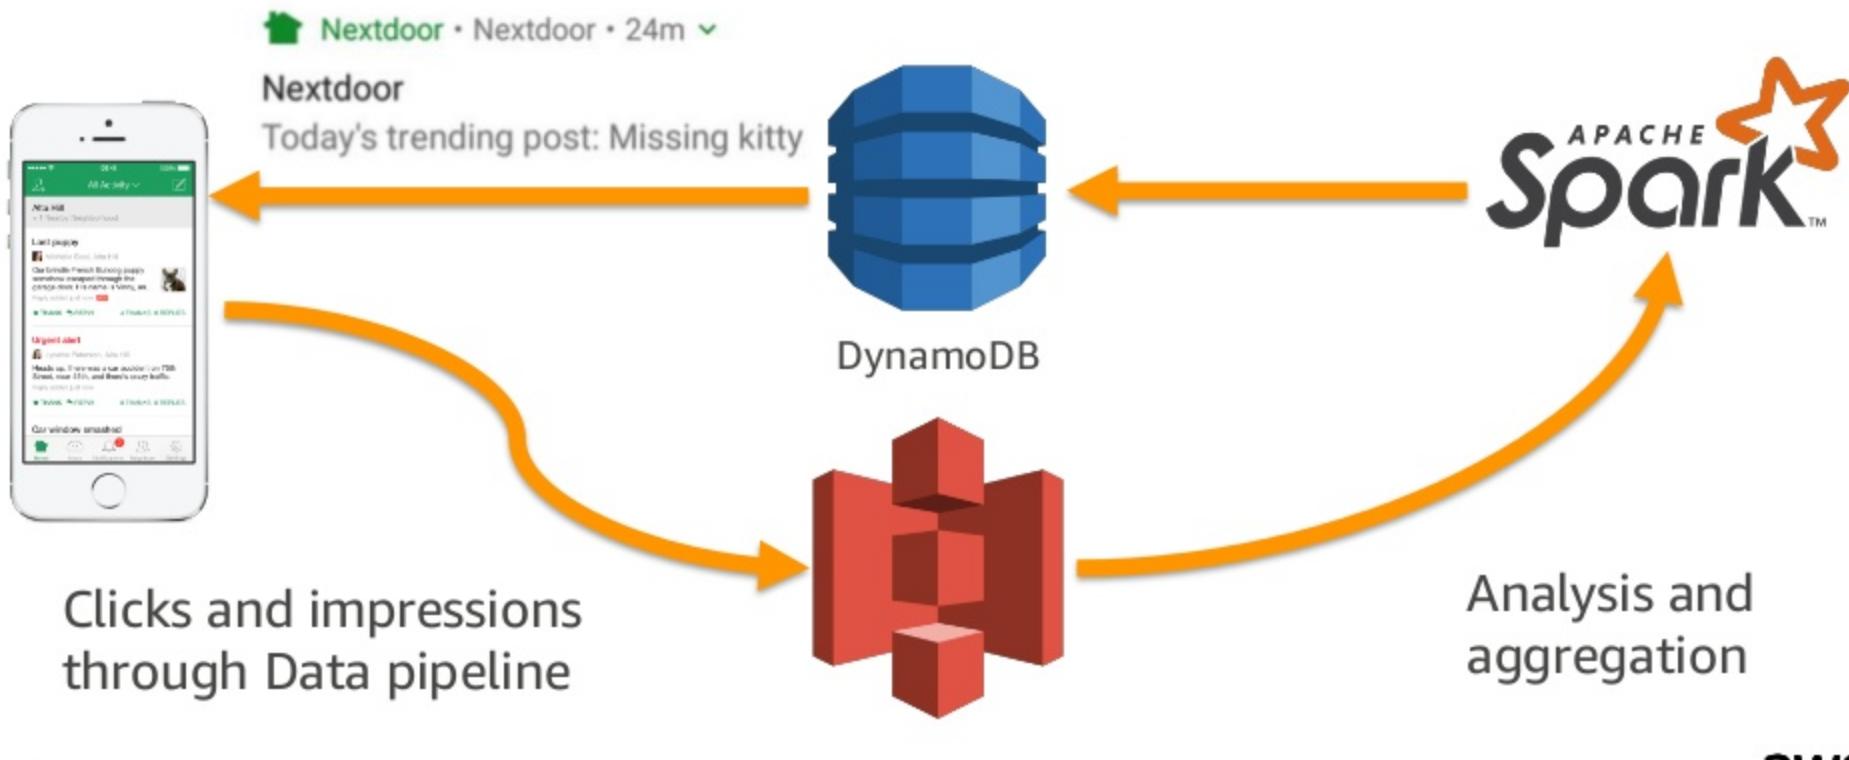
- Highly susceptible to back pressure
 - Slow destination
 - Influx of data
 - Occasional slow Flume Collector (process, IO, CPU, bugs?)
 - Restarts were very slow (tens of minutes)
- Hard to apply changes to the pipeline
- Proprietary storage format made it difficult to near-impossible to recover corrupted data
- Flume agents were CPU hungry and this caused issues on critical infrastructure
- Hard to scale quickly with influxes of log data

What we wanted in a new pipeline



- Spend less (or no) time managing the pipeline so we could focus on other responsibilities
- Leaner server infrastructure
- Easy (re)configuration of the pipeline
- Tighter SLAs
 - Streaming logs into Elasticsearch with an SLA of 1 minute
 - Writing partitioned batches of data into S3 with an SLA of 5 minutes
- Automatically handle influx in data and normal cyclical patterns

Example: Trending post notifications



AWS managed offerings



Kinesis Streams



Lambda

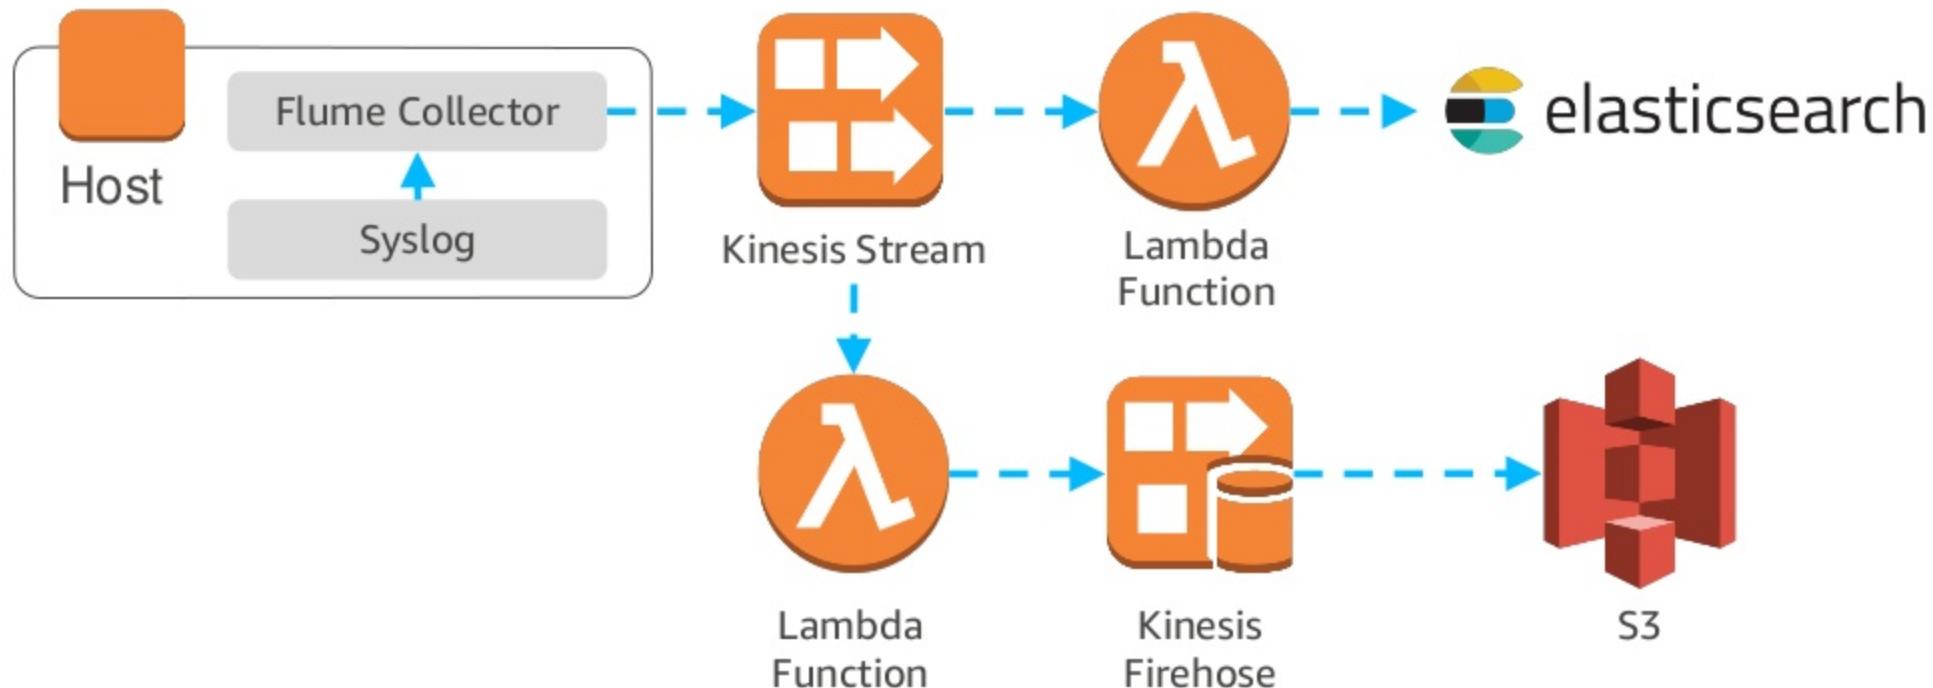


Kinesis Firehose



S3

What your pipeline might look like



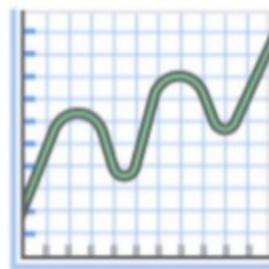
AWS Lambda functions



- Easy to write lightweight processing functions in Node, Python, Java, and C#
- However, it's ETL and there's a lot of boilerplate



No Servers to Manage



Continuous Scaling



Subsecond Metering

What is ETL?

Extract

- Read data from a source and make available for consumption
- Consume the data into a data structure readable by a machine language (for example Java Objects)

Transform

- Lowercase, Redact, Augment, Join, and so on

Load

- Write object back to binary form (maybe compress)
- Transport the binary data to a destination

Where is the boilerplate?

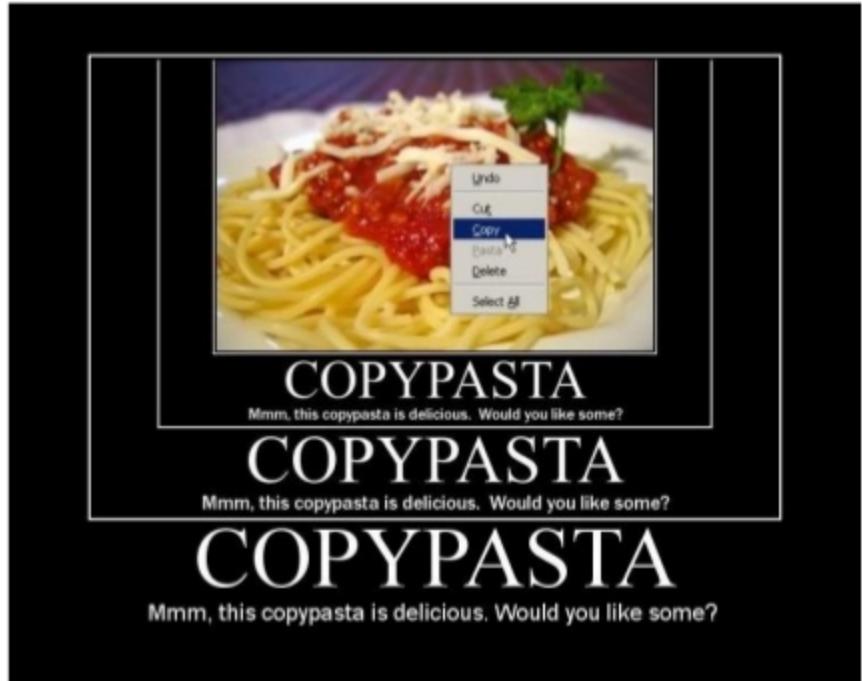


	Step	Example
Extract	Reading data source	Reading file from S3 in batches
	Deserialization	Binary → Object
Transform	Filtering	Remove trace level logging
	Transformation	Lowercase, redact, augment, join
Load	Serialization	Object → Binary
	Transport	Binary → Destination format → Destination
	Insight	Counting, logging, monitoring, and alerting
	Error handling	Retrying, exception handling, error reporting

Standardizing Lambda functions



- What if tomorrow we wanted to ingest logs from ALB, ELB, VPC...
- How do we reuse existing code while adding features and minimizing spaghetti code and copypasta?



Design considerations



- CPU cycles
- Memory footprint
- Network I/O
- Concurrency and thread pooling
- Buffering and compressing
- Retry logic
- Graceful error handling
- Metric reporting
- Configuration

Our solution



Bender\

github.com/nextdoor/bender



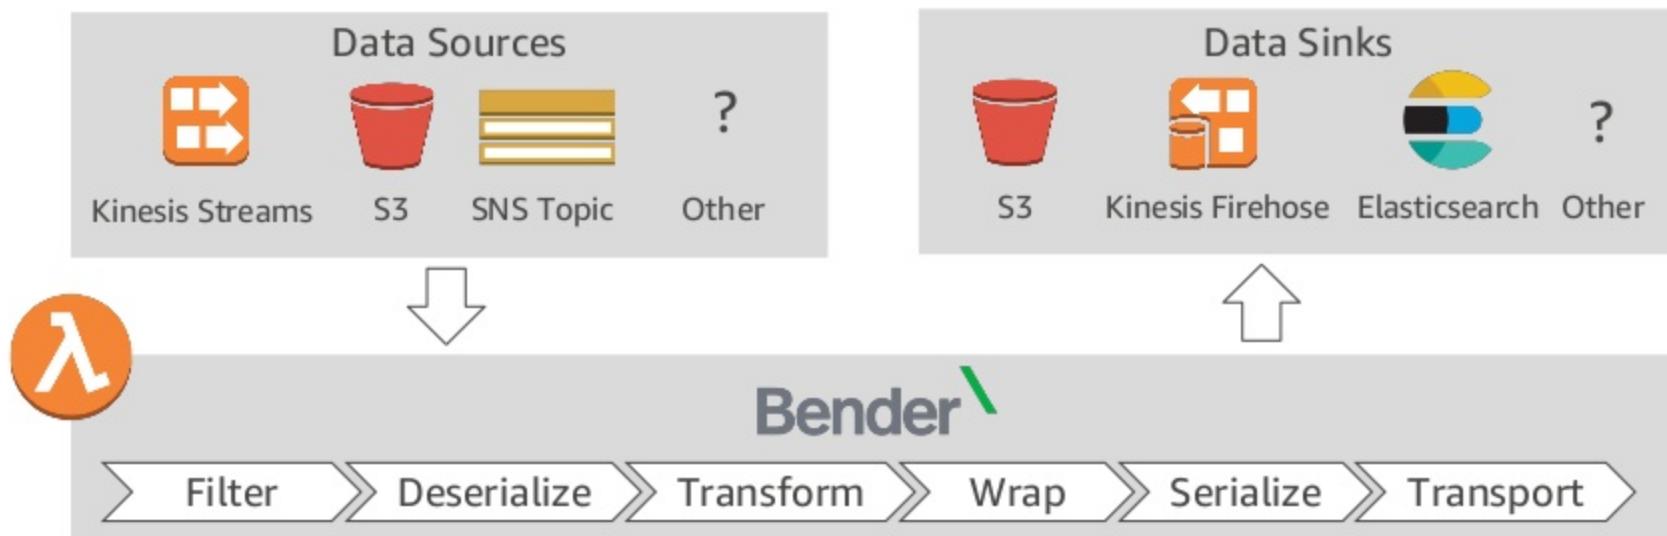
© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



No more spaghetti and copypasta!



Bender handles the complex plumbing and provides the interfaces necessary to build your own business logic for any step of the ETL process.



What ETL steps are supported?



ETL Step	
Extract	Sources
	Deserialization
Transform	Filtering
	Operations
Load	Wrappers
	Transporters

Additional functionality



Other		
Useful	Error handling	Graceful retries and robust exception handling
	Reporting	CloudWatch metrics and Datadog
	Configuration	Config validation before deployment
	Documentation	Auto generated documentation from code

Bender: Configuration



What does the configuration look like?



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```

What does the configuration look like?



```
handler:  
  type: KinesisHandler
```

Type of Lambda trigger this function is using

```
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```



Kinesis: prod-syslog

arn:aws:kinesis:us-east-1:123456789:stream/prod-syslog

Last processing result: **OK** Batch size: **10000**

What does the configuration look like?



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```

How to treat data coming from the source

What does the configuration look like?



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```



What does the configuration look like?



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson ← How to read data coming from the source  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```

What does the configuration look like?



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```

A large black arrow points from the text "Operations to perform on the data" to the **operations:** section of the configuration code.

Operations to perform on the data



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



What does the configuration look like?



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper ← Add a wrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```

```
{  
  "partition_key": "1234.456",  
  "sequence_number": "12345",  
  "source_arn": "us-east-1:1234:stream/logs.",  
  "event_source": "aws:kinesis",  
  "function_name": "log-pipeline...",  
  "function_version": "97",  
  "processing_time  "arrival_time  "processing_delay  "timestamp": 1507823272930,  
  "payload": {  
    // JSON data from original message  
}
```

What does the configuration look like?



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```

Write it back to JSON string or maybe something else?

What does the configuration look like?



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```

Send it off somewhere

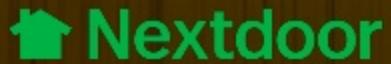
What does the configuration look like?



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
  regex_patterns: [".*TRACE.*"]  
deserializer:  
  type: GenericJson  
operations:  
- ...  
wrapper:  
  type: KinesisWrapper  
serializer:  
  type: Json  
transport:  
  type: ElasticSearch  
  index_time_format: yyyy-MM-dd  
reporters:  
- type: Cloudwatch
```



Extending Bender



Operation to transform JSON



If key exists replace with value

```
if payload.has_key(key):  
    payload[key] = new_value
```

Something like this...

What the config code will look like



```
@JsonTypeName("JsonReplaceKVOperation")
@JsonSchemaDescription("Replaces a given key with specified value.")
public class JsonReplaceKVOperationConfig extends OperationConfig {

    @JsonSchemaDescription("Key name to find")
    @JsonProperty(required = true)
    private String key;

    @JsonSchemaDescription("New string value")
    @JsonProperty(required = true)
    private String value;

    @Override
    public Class<JsonReplaceKVOperationFactory> getFactoryClass() {
        return JsonReplaceKVOperationFactory.class;
    }

    // getters and setters
}
```

What the config file will look like



```
handler:  
  type: KinesisHandler  
sources:  
- name: Events  
  source_regex: .*:stream/my-stream  
deserializer:  
  type: GenericJson  
operations:  
- type: JsonReplaceKVOperation  
  key: foo  
  value: bar  
...
```

What the documentation will look like



JsonReplaceKVOperation
Replaces a given key with specified value.

<u>type</u>	req		
string	constant	JsonReplaceKVOperation	default JsonReplaceKVOperation
<u>key</u>	req	string	
Key name to find			
<u>value</u>	req	string	
New string value			

Operation code



```
public class JsonReplaceKVOperation extends PayloadOperation {  
    private String key;  
    private String value;  
  
    public JsonReplaceKVOperation(String key, String value) {  
        this.key = key;  
        this.value = value;  
    }  
  
    @Override  
    protected void perform(JsonObject payload) {  
        if (payload.has(this.key))  
            // Note: addProperty replaces the current value  
            payload.addProperty(this.key, this.value);  
    }  
}
```

Code on GitHub



Bender\

github.com/nextdoor/bender



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Why did we choose AWS serverless offerings?

Matt Wise · Sr. Systems Architect



Kinesis for data ingestion



- Operational overhead is near zero
- Scaling streams up and down is a simple API call
- Scaling up gives you more performance in every dimensions
- Data retention of 24 hours to 7 days
- Stream access is fast – delays measured in milliseconds
- AWS provides tools...
 - Log collection agent
 - Stream Auto Scaling agent
 - KCL for in-app direct stream publishing

Lambda for our execution layer



- Scales automatically with Kinesis and your data
- Manages Kinesis Triggers, Function Iterators all on its own
- Provides plenty of metrics that are easy to alarm on
- Blazing fast ... backlogs are worked through as fast as your target can handle it
- Built in easy to use management operations
 - Pausing and enabling triggers
 - Releasing by updating a function to the latest version
 - Rollbacks by just rolling back to the last known good version

Firehose for data aggregation



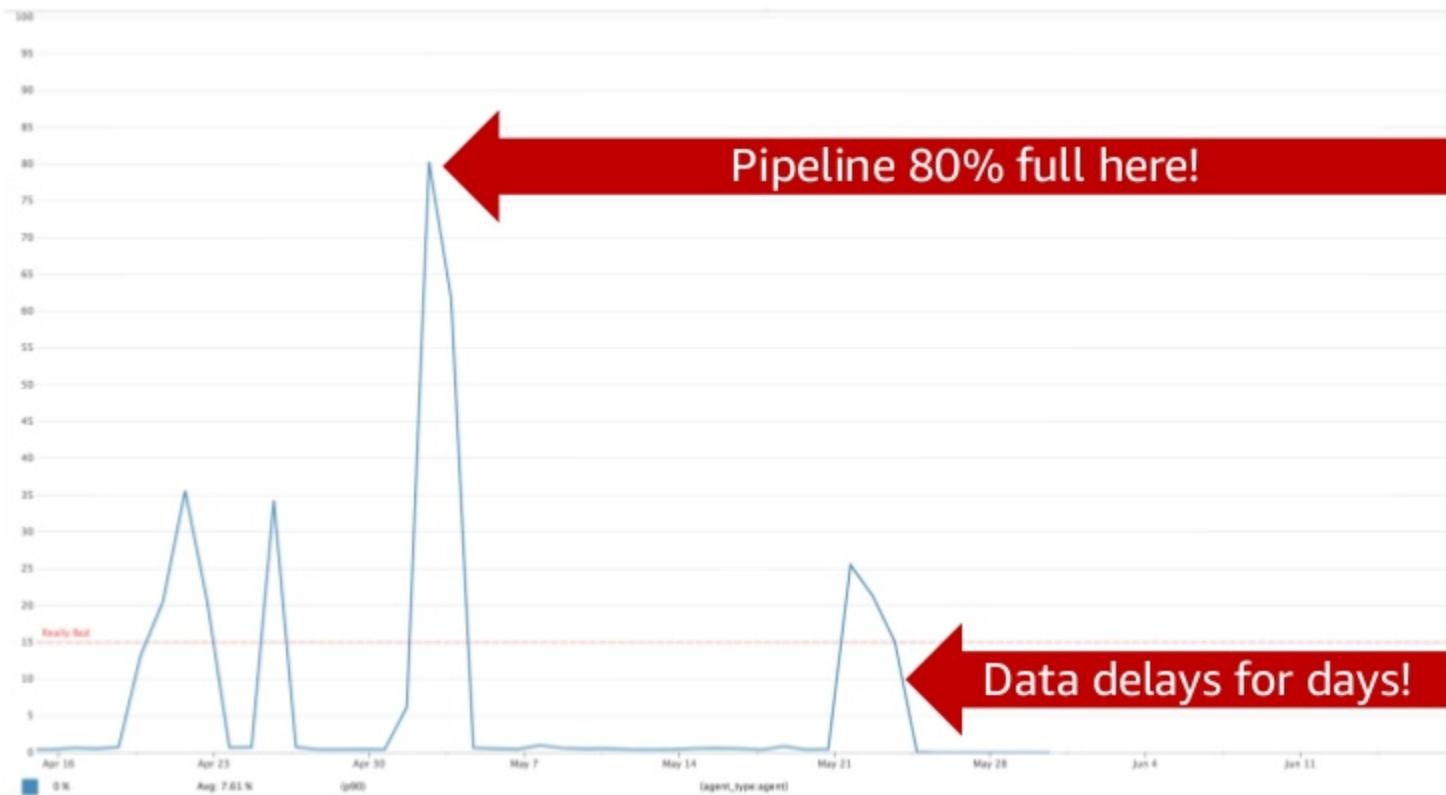
- Lambda+Kinesis Streams = Lots of small operations!
- Kinesis -> Lambda -> S3 = Tons of tiny files
- Bulk-uploading files to S3 leads to faster downstream analytics performance
- Bender can batch multiple individual events up into a single Firehose Record, keeping the costs sane and maximizing performance!
- Max-file-size and Max-event-age settings allow you to control your uploads and your data delay

S3 for short term and long term storage

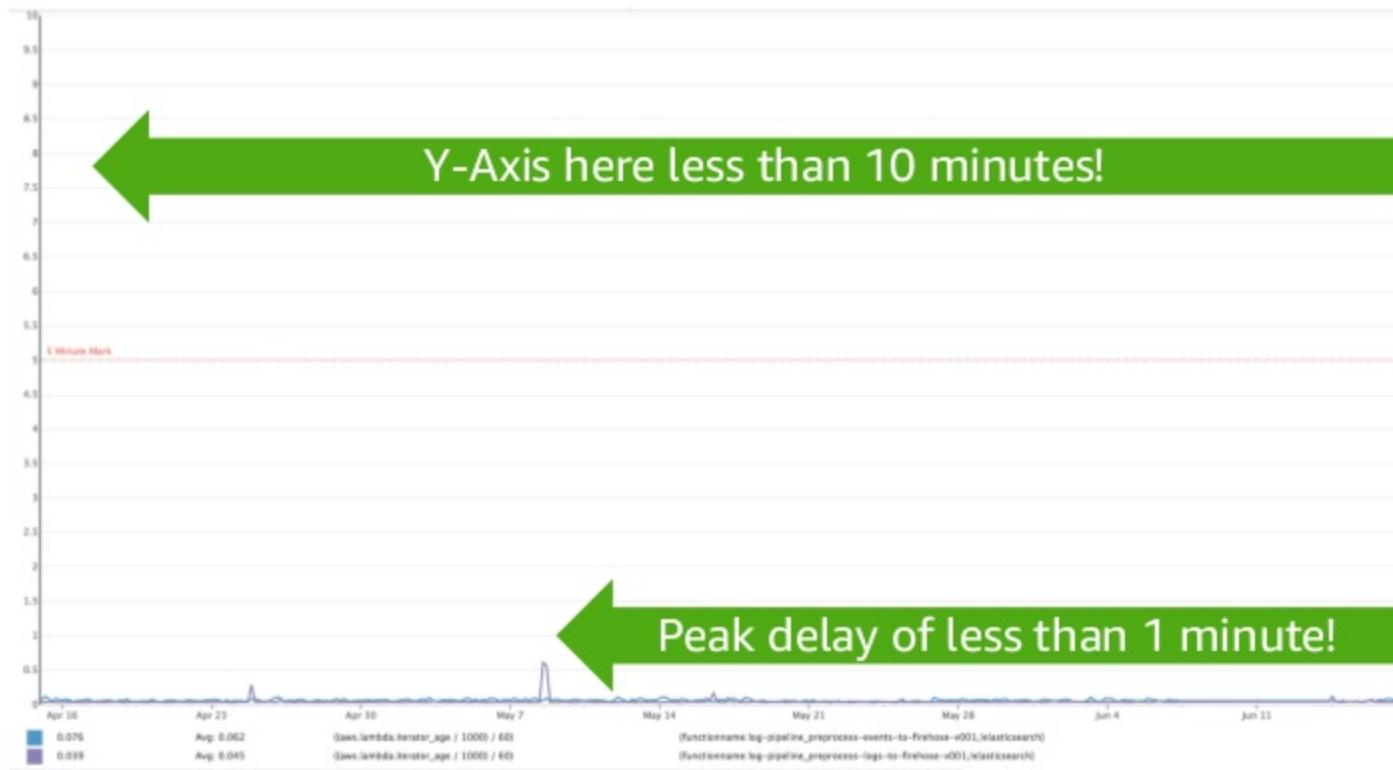


- Fast, cost effective, and stable
- Lifecycle policies allow for automatic purging or migration of data
- Hands-off operation
- Fine grained access controls
- And so on, and so on, and so on...

Final months of the old pipeline...



Well that looks better!

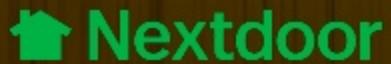


Totally worth it!

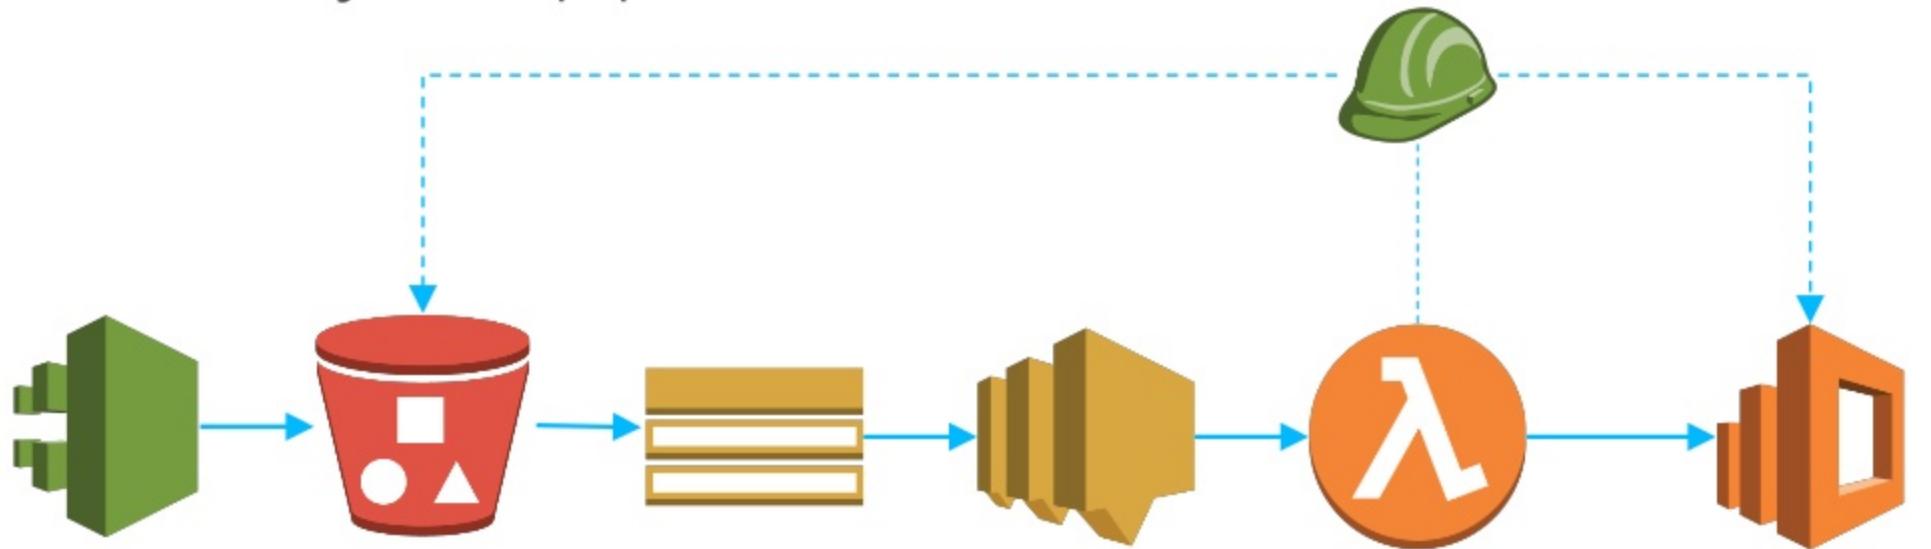


Flume	Kinesis/Lambda/Firehose/S3
@ 1.5B events/day, was dropping 0.02%	@ 3B events/day, all events make it
Required significant architectural work to scale beyond ~2B events/day	Built to scale on day one, all major scaling components are AWS-managed
Near weekly outages that lasted for hours	In 12 months, 1 outage that lasted more than an hour.. Automatically resolved by AWS
Significant operational knowledge required for operation and scaling of the system	Java knowledge required to add features – AWS operates the rest
~15-20% of an engineer's time to manage	Literally almost zero management in months
Combined ingest/storage/forwarding make changes hard and reduce stability	Separated responsibilities are easier to scale and tweak, improve reliability

A real pipeline!



Anatomy of a pipeline



CloudTrail

S3 Bucket

SNS Topic

Subscription

Bender

ElasticSearch
Domain

cloudtrail-bender-demo-Cloudtrail-1U3VDHDF3JUJL ▼ Trail settings 

When a trail applies to all regions, the trail exists in all regions and delivers log files for all regions to one Amazon S3 bucket and an optional CloudWatch Logs log group. To see all of your trails, click [Trails](#).

Apply trail to all regions Yes

► Management events 

► Data events 

▼ Storage location 

S3 bucket [cloudtrail-bender-demo](#)

 Last log file delivered 2017-10-20, 3:17 pm

Encrypt log files No

Enable log file validation No

Publish to SNS No

Overview

 Type a prefix and press Enter to search. Press ESC to clear.

 Upload

 + Create folder

 More

US East (N. Virginia) 

Viewing 1 to 300 >

<input type="checkbox"/>	Name	Last modified	Size	Storage class
<input type="checkbox"/>	 364942603424_CloudTrail_us-east-1_20171020T0000Z_UIR0fJ...	Oct 19, 2017 5:02:32 PM	60.1 KB	Standard
<input type="checkbox"/>	 364942603424_CloudTrail_us-east-1_20171020T0000Z_uyotu4...	Oct 19, 2017 5:02:27 PM	38.3 KB	Standard
<input type="checkbox"/>	 364942603424_CloudTrail_us-east-1_20171020T0005Z_7tqlnn...	Oct 19, 2017 5:07:31 PM	26.5 KB	Standard
<input type="checkbox"/>	 364942603424_CloudTrail_us-east-1_20171020T0005Z_CaZ7Q...	Oct 19, 2017 5:07:40 PM	39.7 KB	Standard
<input type="checkbox"/>	 364942603424_CloudTrail_us-east-1_20171020T0010Z_BvhYfJ...	Oct 19, 2017 5:12:34 PM	21.9 KB	Standard
<input type="checkbox"/>	 364942603424_CloudTrail_us-east-1_20171020T0010Z_frlhQ3t...	Oct 19, 2017 5:12:38 PM	17.0 KB	Standard
<input type="checkbox"/>	 364942603424_CloudTrail_us-east-1_20171020T0015Z_LetfqJ...	Oct 19, 2017 5:17:38 PM	34.8 KB	Standard

cloudtrail-bender-demo-

Function-DGYQVDG82HG2

Qualifiers ▾

Actions ▾

Select a test event.. ▾

Test

Configuration

Triggers

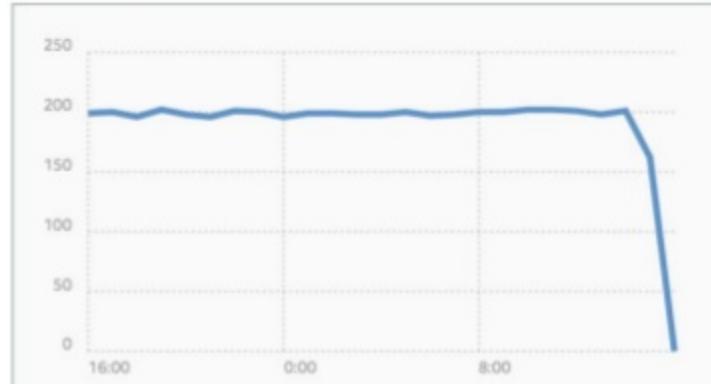
Monitoring

CloudWatch metrics at a glance (last 24 hours) 

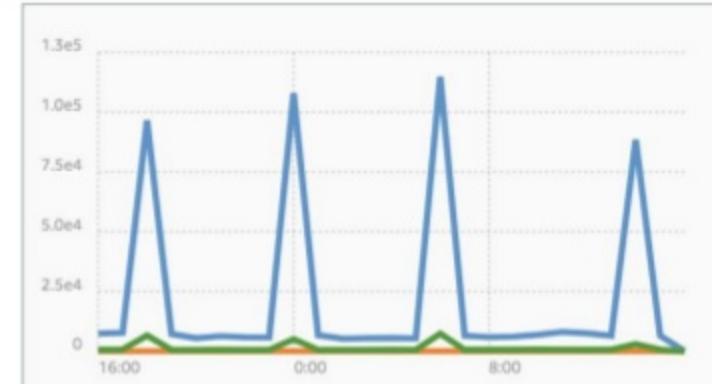
[View traces in X-Ray](#)

[View logs in CloudWatch](#)

Invocations



Duration



cloudtrail-bender-demo

[Configure cluster](#)[Modify access policy](#)[Manage tags](#)[Delete domain](#)[Overview](#)[Cluster health](#)[Indices](#)[Monitoring](#)[Logs](#)

Time range

Last 2 weeks ▾

ClusterStatus.green (Count)
Statistic: Minimum

Green

**ClusterStatus.yellow** (Count)
Statistic: Maximum

Yellow

**ClusterStatus.red** (Count)
Statistic: Maximum

Red

**Nodes** (Count)
Statistic: Average

3

2

1

0

**SearchableDocuments** (Count)
Statistic: Average

12000000

10000000

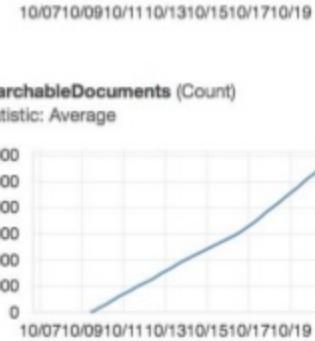
8000000

6000000

4000000

2000000

0

**DeletedDocuments** (Count)
Statistic: Average

5000

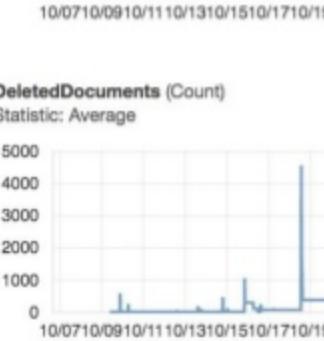
4000

3000

2000

1000

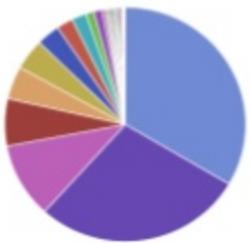
0

**aws**
re:Invent

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

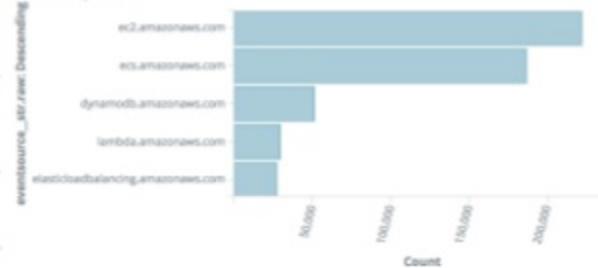
aws

API Calls by IAM Name

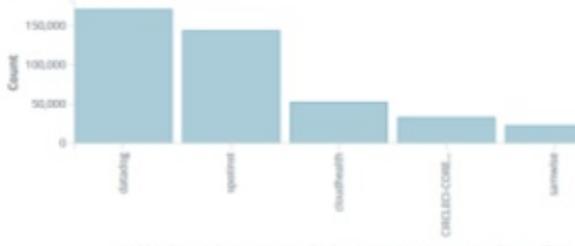


- datadog
- spotinst
- cloudehealth
- CIRCLECI-CORE-SERV...
- samwise
- puppet_profile_ecs
- taskworker_junitor
- feedstore
- CIRCLECI-CORE-SERV...
- preview-helper
- rds_badger_runner
- notifications
- nextdoor-api-staging...
- nextdoor-fa-staging...

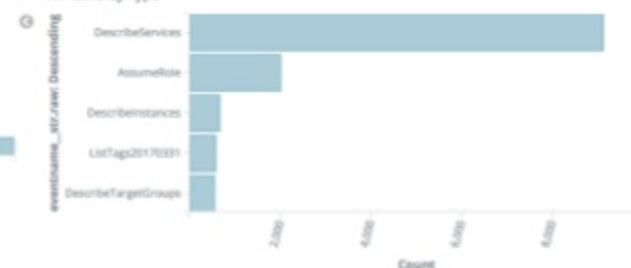
API Calls by Service



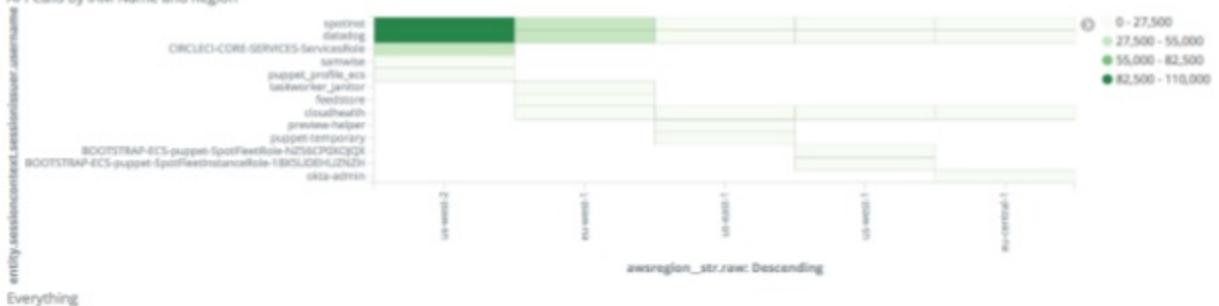
API Calls by IAM Name - Histogram



API Calls by Type



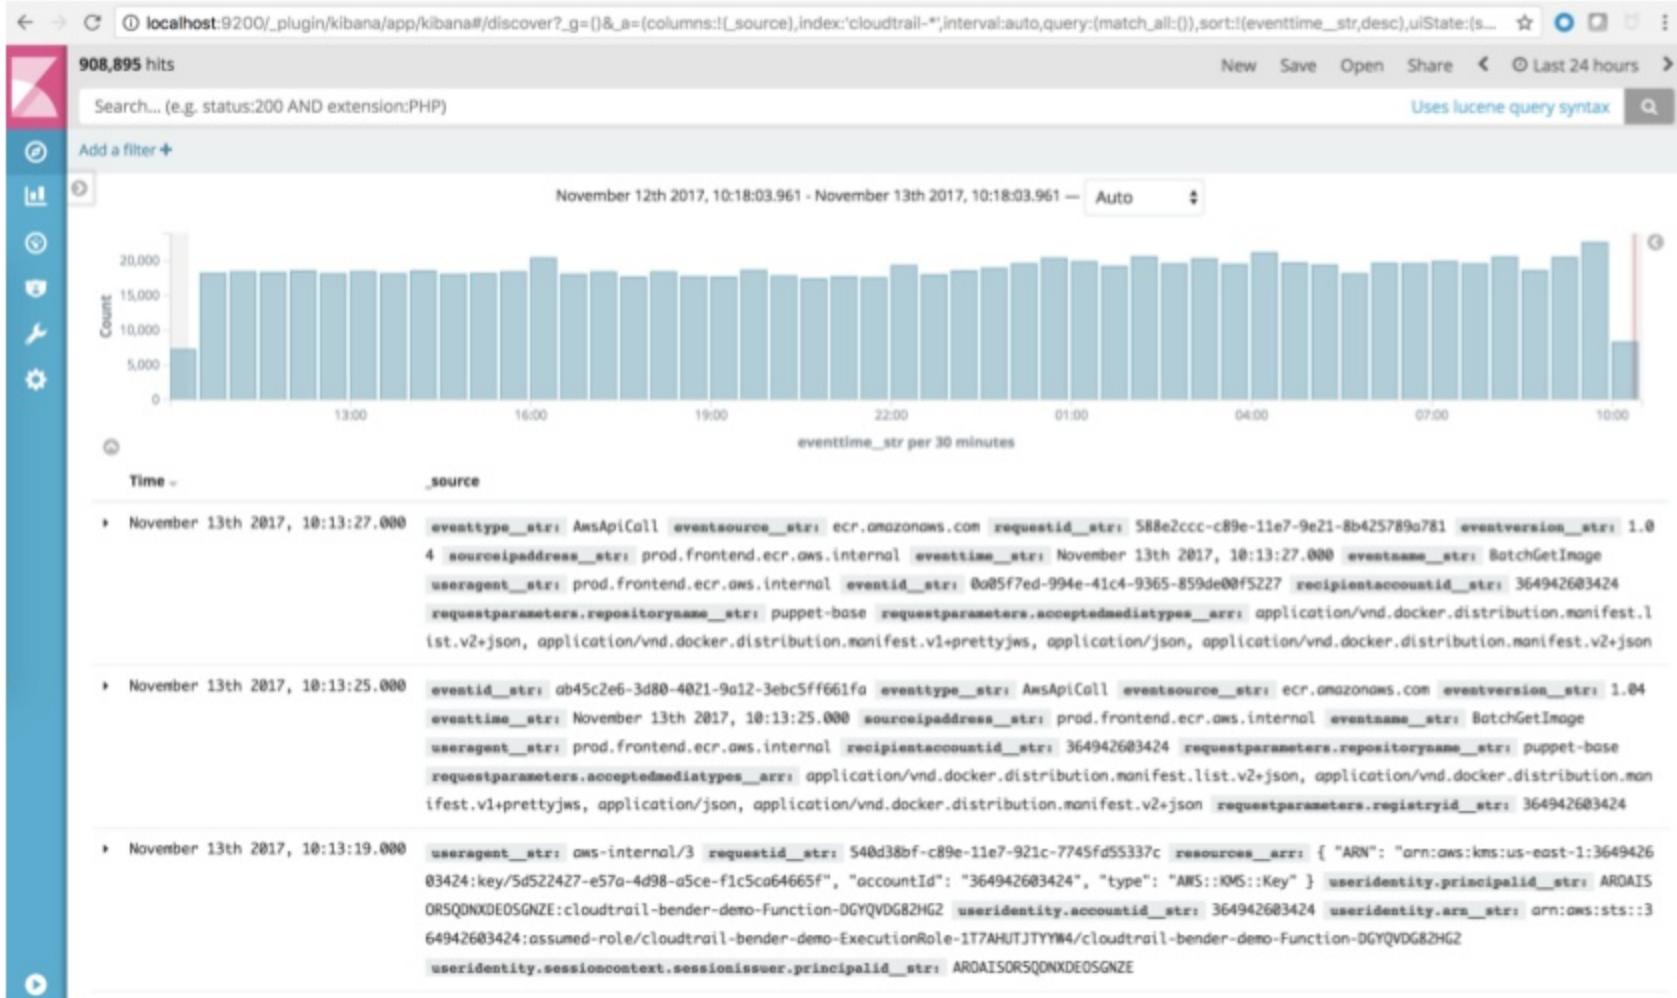
API Calls by IAM Name and Region



Everything

1-50 of 648,415





t _id	Q Q □ * 85bb6f92dba44ba0a043a3a856c62292a273ffb6
t _index	Q Q □ * cloudtrail-2017-11-13
# _score	Q Q □ * -
t _type	Q Q □ * cloudtrail
t awsregion__str	Q Q □ * us-east-1
t eventid__str	Q Q □ * 37a93267-b673-4777-b8de-f65b209bdc2a
t eventname__str	Q Q □ * DescribeReservedInstances
t eventsource__str	Q Q □ * ec2.amazonaws.com
⌚ eventtime__str	Q Q □ * November 13th 2017, 10:13:13.000
t eventtype__str	Q Q □ * AwsApiCall
t eventversion__str	Q Q □ * 1.05
t recipientaccountid__str	Q Q □ * 364942603424
t requestid__str	Q Q □ * 18661cd4-3813-4e3e-b50c-f9ed413a4b48
? requestparameters.filterset.items__arr	Q Q □ * { "name": "state", "valueSet": { "items": [{ "value": "active" }] } }
▲ t sourceipaddress__str	Q Q □ * 34.231.218.242
t useragent__str	Q Q □ * aws-sdk-nodejs/2.148.0 linux/v5.11.1
t useridentity.accesskeyid__str	Q Q □ * ASIAIYVFVB4VRJ3LGG6Q
t useridentity.accountid__str	Q Q □ * 364942603424
t useridentity.arn__str	Q Q □ * arn:aws:sts::364942603424:assumed-role/spotinst/spotinst.session.1510596017001
t useridentity.principalid__str	Q Q □ * AR0AI324345A5RR742AKY:spotinst.session.1510596017001

Search... (e.g. status:200 AND extension:PHP)

Uses lucene query syntax



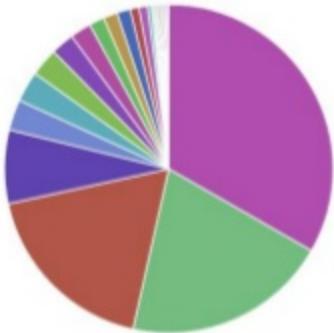
ecs.amazonaws.com

3rd Party

Add a filter +

Actions ➔

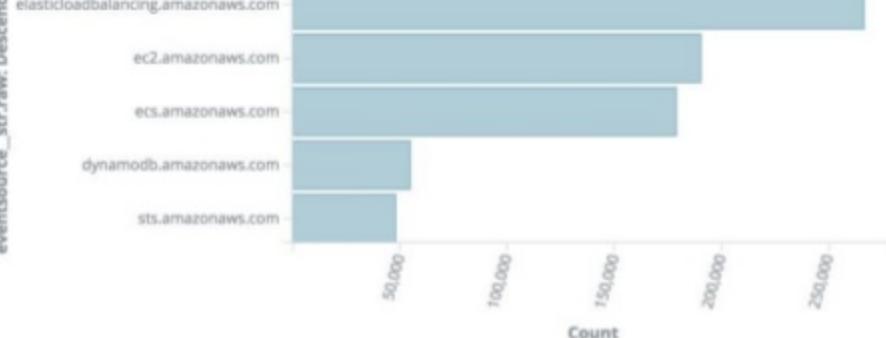
API Calls by IAM Name



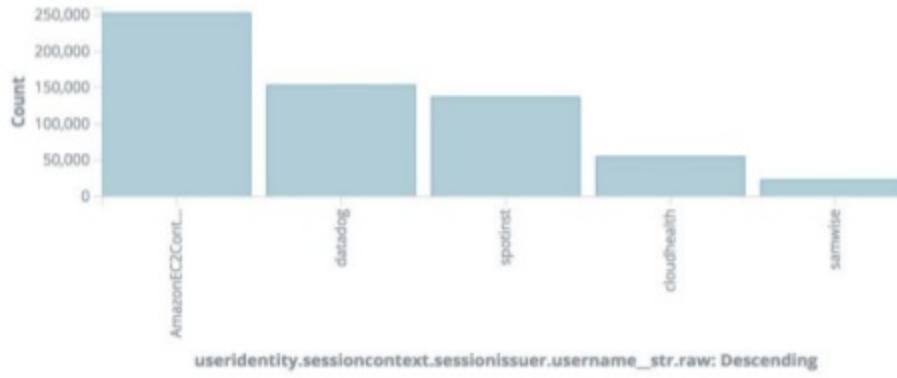
API Calls by Service

- AmazonEC2ContainerServiceRole
- datadog
- spotinst
- cloudhealth
- samwise
- feedstore
- BOOTSTRAP-ECS-puppet
- puppet_profile_ecs
- taskworker_janitor
- preview-helper
- CIRCLECI-CORE-SERVICE
- CIRCLECI-CORE-SERVICE
- rds_badger_runner
- okta-poweruser

eventsource_str.raw: Descending



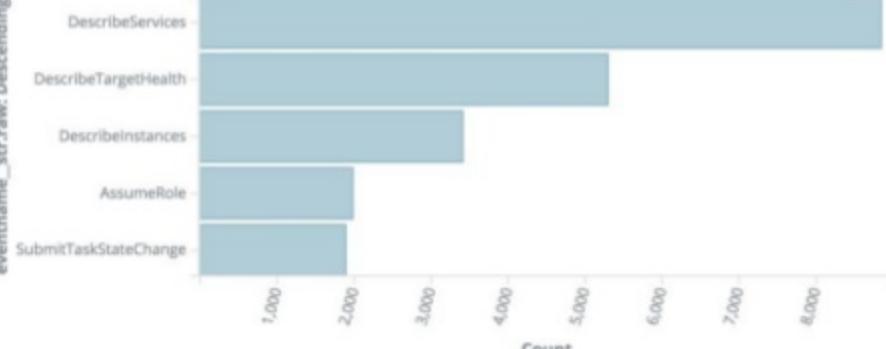
API Calls by IAM Name - Histogram



API Calls by Type

- DescribeServices
- DescribeTargetHealth
- DescribeInstances
- AssumeRole
- SubmitTaskStateChange

eventname_str.raw: Descending



Search... (e.g. status:200 AND extension:PHP)

Uses lucene query syntax



ecs.amazonaws.com

3rd Party

useridentity.sessioncontext.sessionissuer.username_str.raw: "AmazonEC2ContainerServiceRole"

useridentity.sessioncontext.sessionissuer.username_str.raw: "datadog"

useridentity.sessioncontext.sessionissuer.username_str.raw: "spotinst"

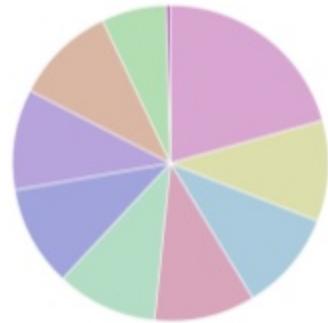
useridentity.sessioncontext.sessionissuer.username_str.raw: "cloudhealth"

eventname_str.raw: "DescribeServices"

Add a filter +

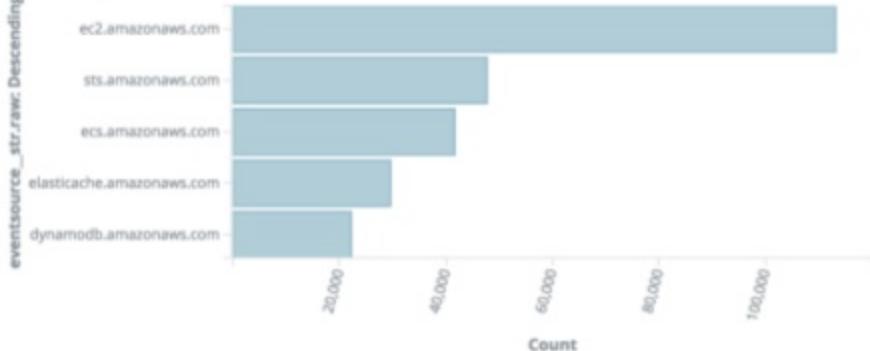
Actions ▾

API Calls by IAM Name

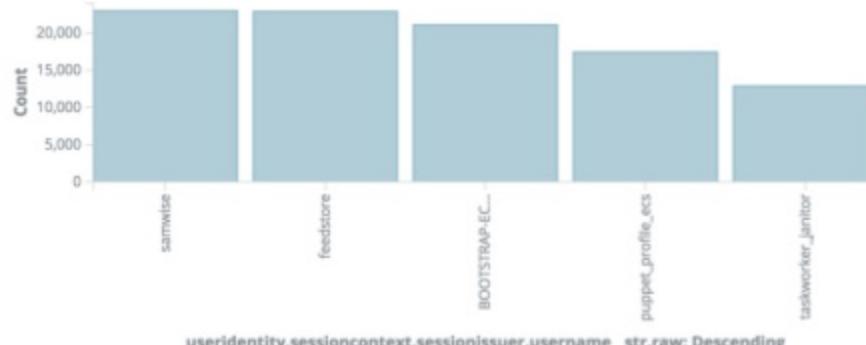


API Calls by Service

- taskworker_janitor
- go-photo-staging-eu...
- nextdoor-api-staging...
- nextdoor-fe-staging...
- go-photo-staging-us1...
- nd-sockets-staging-e...
- nd-sockets-staging-e...
- preview-helper
- okta-poweruser



API Calls by IAM Name - Histogram



API Calls by Type

- AssumeRole
- SubmitTaskStateChange
- RegisterTaskDefinition
- DescribeLoadBalancers
- DeregisterTaskDefinition

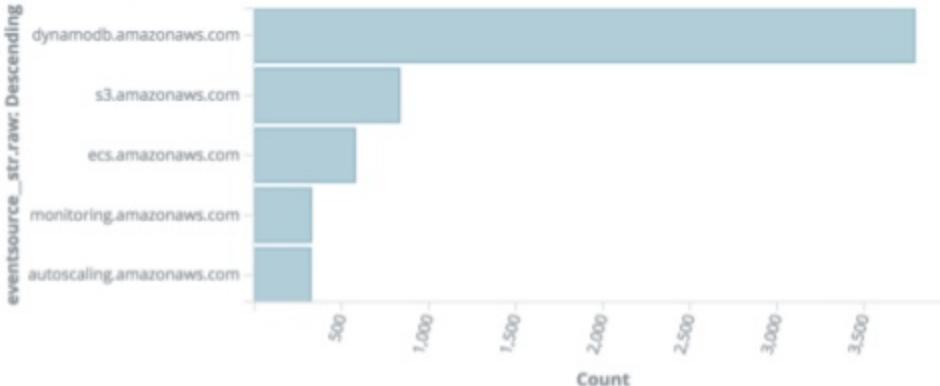


API Calls by IAM Name



okta-poweruser

API Calls by Service

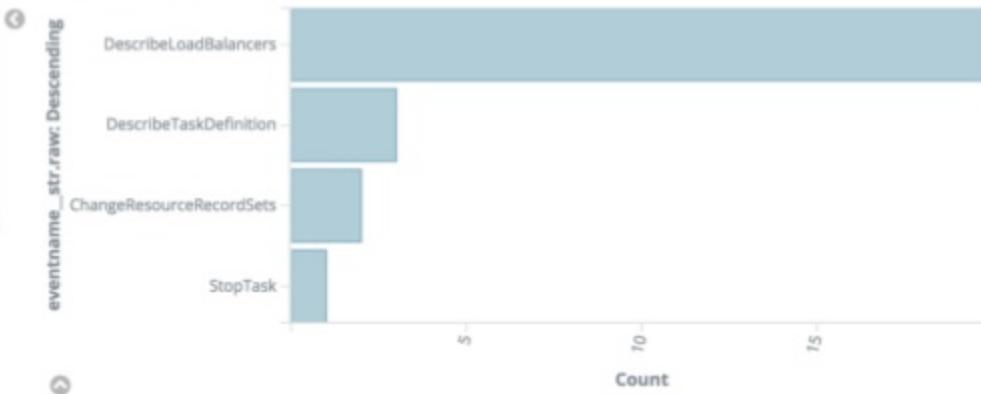


API Calls by IAM Name - Histogram



useridentity.sessioncontext.sessionissuer.username__str.raw: Descending

API Calls by Type



"okta-"

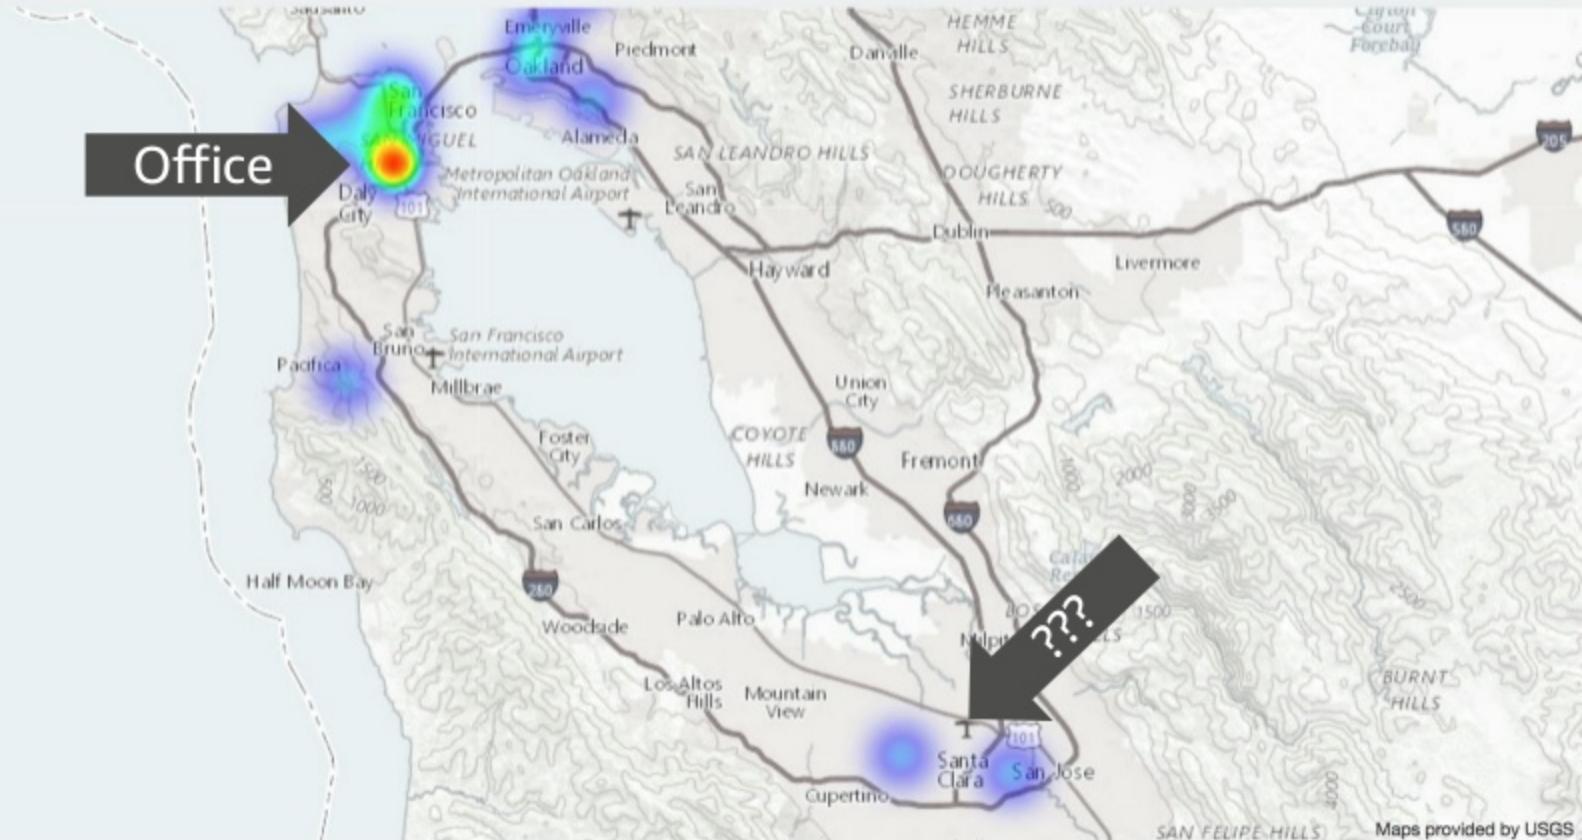
Uses lucene query syntax



Add a filter +



Office



Maps provided by USGS

"okta-"

Uses lucene query syntax



sourcelpaddress_geo.location: "{"lat": 37.41783605625168, "lon": -122.04437255859375 } to { "lat": 37.296670230081766, "lon": -121.7889404296875 }"

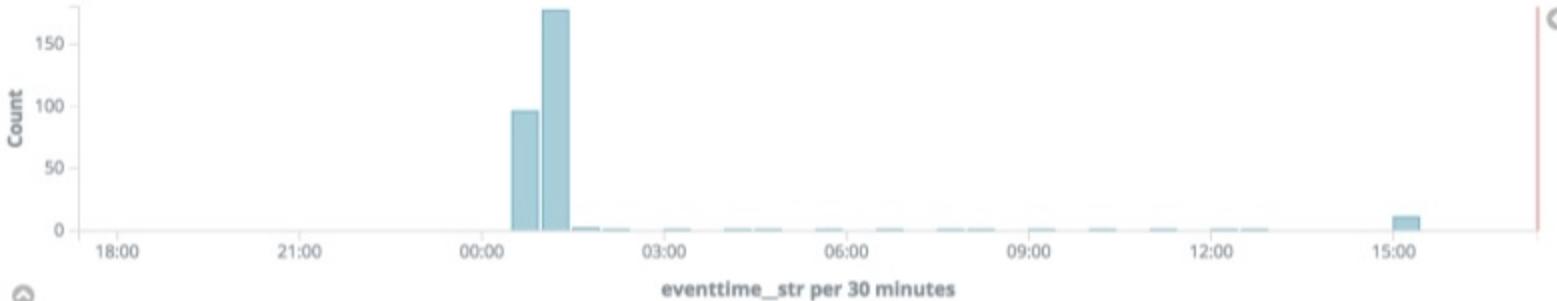
Geo Filter

Actions >

cloudtrail*

November 15th 2017, 17:22:13.594 - November 16th 2017, 17:22:13.594 —

Auto



Selected Fields

t eventname_str
t useridentity.arn...

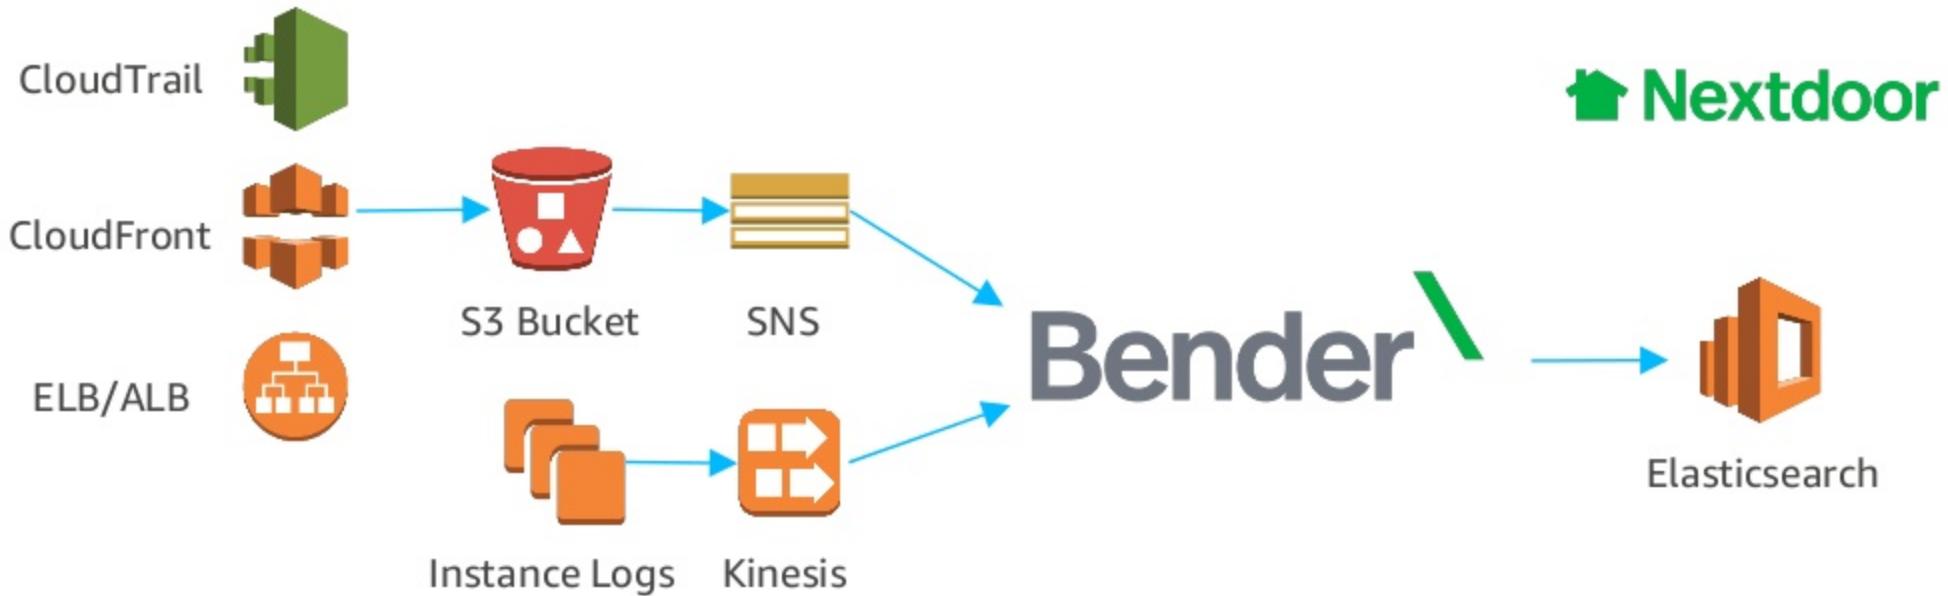
Available Fields

Popular

t requestparameters....

t _id
t _index
_score
t _type
t apiversion_str
t awsregion_str
t errorcode_str
t errormessage str

Time	eventname_str	useridentity.arn_str
▶ November 16th 2017, 15:13:34.000	DescribeVolumes	arn:aws:sts::364942603424:assumed-role/okta-admin/matt@nextdoor.com
▶ November 16th 2017, 15:13:34.000	DescribeInstanceStatus	arn:aws:sts::364942603424:assumed-role/okta-admin/matt@nextdoor.com
▶ November 16th 2017, 15:13:34.000	DescribeVolumeStatus	arn:aws:sts::364942603424:assumed-role/okta-admin/matt@nextdoor.com
▶ November 16th 2017, 15:13:34.000	DescribeAvailabilityZones	arn:aws:sts::364942603424:assumed-role/okta-admin/matt@nextdoor.com
▶ November 16th 2017, 15:13:34.000	DescribeVolumes	arn:aws:sts::364942603424:assumed-role/okta-admin/matt@nextdoor.com



Final thoughts...



- The Systems Team @ Nextdoor has 3 people...
- We manage hundreds of servers, and we *automate everything*
- Hands-off auto-scaling is critical to keeping our team small and efficient
- Lambda/Kinesis costs are a drop in the bucket compared to engineering productivity costs when things don't work

We're hiring!



nextdoor.com/jobs

<https://s3.amazonaws.com/bender-presentation-assets-reinvent-2017/bender-stack.yaml>



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





Thank you!