

# Hadoop & Spark – Using Amazon EMR

Michael Hanisch, AWS Solutions Architecture

# Agenda

Why did we build Amazon EMR?

What is Amazon EMR?

How do I get started using Amazon EMR?

Q&A

# Why did we build Amazon EMR?

# Ingest



Amazon Mobile  
Analytics



AWS Import/  
Export

# Store



Amazon RDS



Amazon  
Kinesis



Amazon  
DynamoDB



Amazon  
CloudSearch



Amazon S3



Amazon  
Glacier

# Process



Amazon EMR



Amazon Machine  
Learning



Amazon Redshift



AWS Data  
Pipeline



Amazon  
Lambda



Amazon  
EC2

# Visualize



# Hadoop: The Framework for Big Data Processing

Can process huge amounts of data

Scalable and fault tolerant

Respects multiple data formats

Supports an ecosystem of tools with a robust community

Executes batch and real-time analytics

# Common Customer Challenges

Must purchase, provision, deploy and manage the infrastructure for Hadoop

Deploying, configuring and administering the Hadoop clusters can be difficult, expensive and time-consuming

# What is Amazon Elastic MapReduce?

# Amazon EMR Benefits

## Simplifies Big Data Processing on a Managed Hadoop Framework

### Low Cost

Pay an hourly rate for every instance hour you use

### Elastic

Easily add or remove capacity

### Easy to use

Launch a cluster in minutes

### Reliable

Spend less time tuning and monitoring your cluster

### Secure

Managed firewall settings

### Flexible

You control every cluster

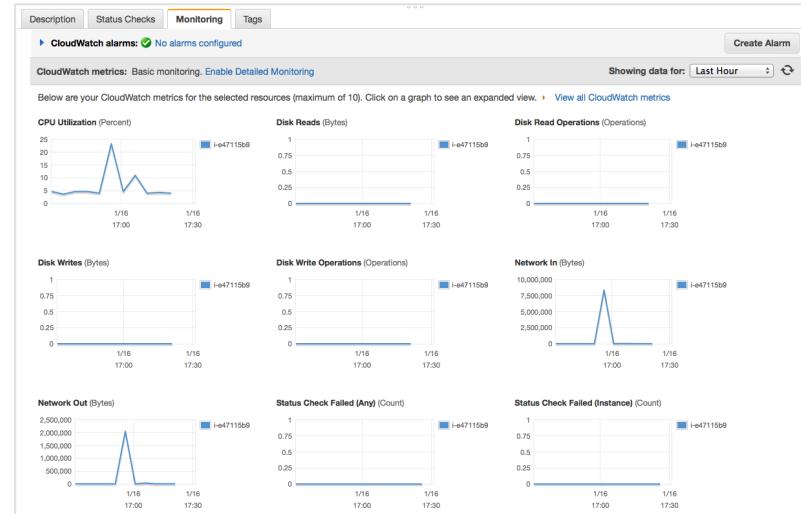


# Amazon EMR is Elastic

Dynamically change number of compute resources

Provision one or thousands of instances to Process Data at any scale

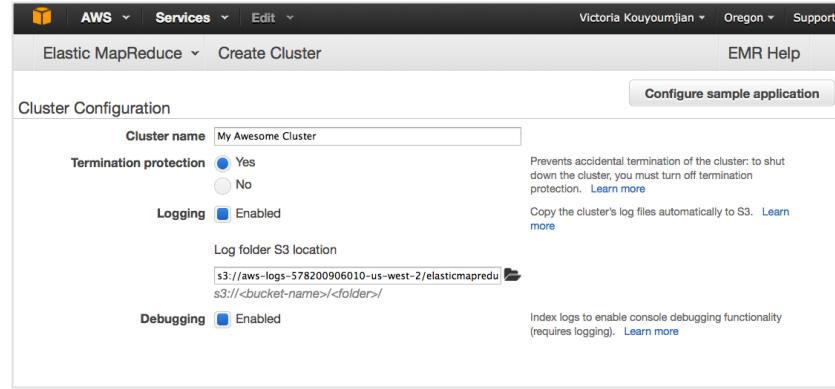
Use CloudWatch to alert for scaling needs



# Easy to Use

Launch a cluster in minutes  
through console

EMR sets up your cluster,  
provisions nodes, and configures  
Hadoop so you can focus on  
your analysis



# Low Cost

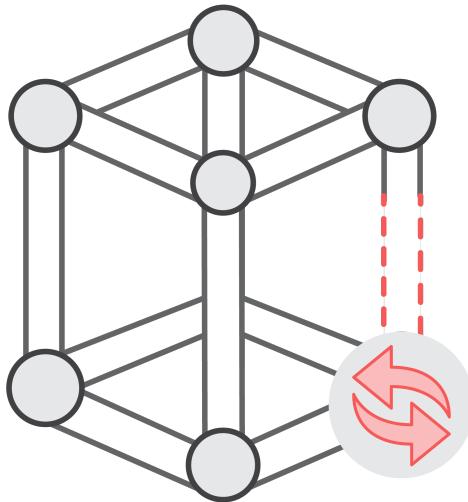


Low Hourly Pricing for On-demand Instance Hours

Name your Price with Amazon EC2 Spot Instance Integration

Use EC2 Reserved Instances to reserve capacity and further lower your costs

# Reliable



Monitors cluster health

Retries failed tasks

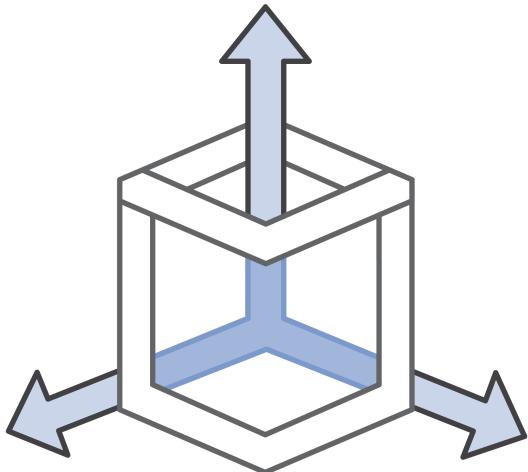
Automatically replaces  
under-performing instances

# Secure



- Automatically configures firewall settings to control network access
- Managed security groups for instances or specify your own
- Launch in the Amazon Virtual Private Cloud (VPC)
- Support for server-side & client-side encryption

# Flexible



Completely control your cluster including customizations

Use data sources from Amazon S3, HDFS, Redshift, RDS, Kinesis and DynamoDB

Supports Apache & MapR Hadoop distributions

Runs popular Hadoop tools and frameworks (Hive, Hbase, Pig, Mahout, Spark, Presto)

**How do I get started  
using Amazon EMR?**



# Easy to start

## AWS Management Console

The screenshot shows the AWS Management Console with the navigation bar at the top. Under 'Services', 'Elastic MapReduce' is selected, and 'Create Cluster' is chosen. The 'Cluster Configuration' section is displayed, containing the following fields:

- Cluster name:** My Awesome Cluster
- Termination protection:** Yes (radio button selected)
- Logging:** Enabled
- Log folder S3 location:** s3://aws-logs-578200905010-us-west-2/elasticmapred/
- Debugging:** Enabled

Below these fields, there is explanatory text and links for termination protection and log copying.

## Command Line

The screenshot shows a terminal window titled 'bash - 64x21'. The command entered is:

```
685b358705ce:~ [user] $ aws emr create-cluster --ami-version 3.2.1 --instance-groups InstanceGroupType=MASTER,InstanceCount=1,InstanceType=m3.xlarge InstanceGroupType=CORE,InstanceCount=2,InstanceType=m3.xlarge --region us-west-2
```

Below the terminal, a faint watermark or background image of an architecture diagram is visible.

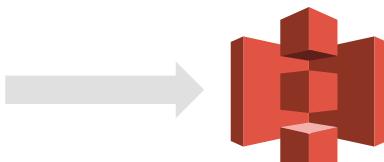
Or use the Amazon EMR API with your favorite SDK.

# How to use Amazon EMR

## 1. Upload

your application  
and data to S3

App &  
Data



## 2. Configure

your cluster: Choose Hadoop  
distribution, number and type  
of nodes, applications (Hive/  
Pig/Hbase)

## 3. Launch

your cluster using the  
console, CLI, SDK, or APIs



## 4. Retrieve

your output  
results from S3

# Configuration & Customization

# Choose your instance types

Try different configurations to find your optimal architecture.

**General**  
m1 family  
m3 family

**CPU**  
c3 family  
cc1.4xlarge  
cc2.8xlarge

**Memory**  
m2 family  
r3 family

**Disk/IO**  
d2 family  
i2 family

**Batch process**

**Machine learning**

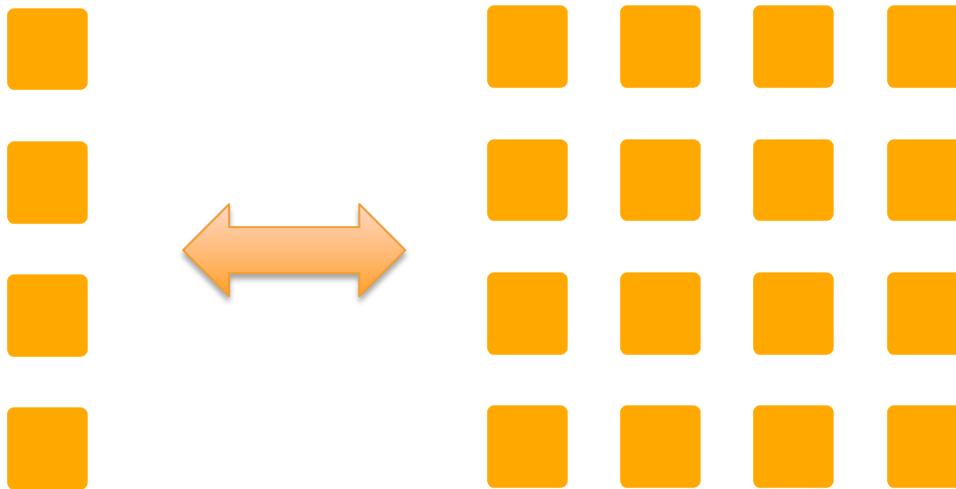
**Spark and interactive**

**Large HDFS**



# Resizable clusters

Easy to add and remove compute capacity on your cluster.

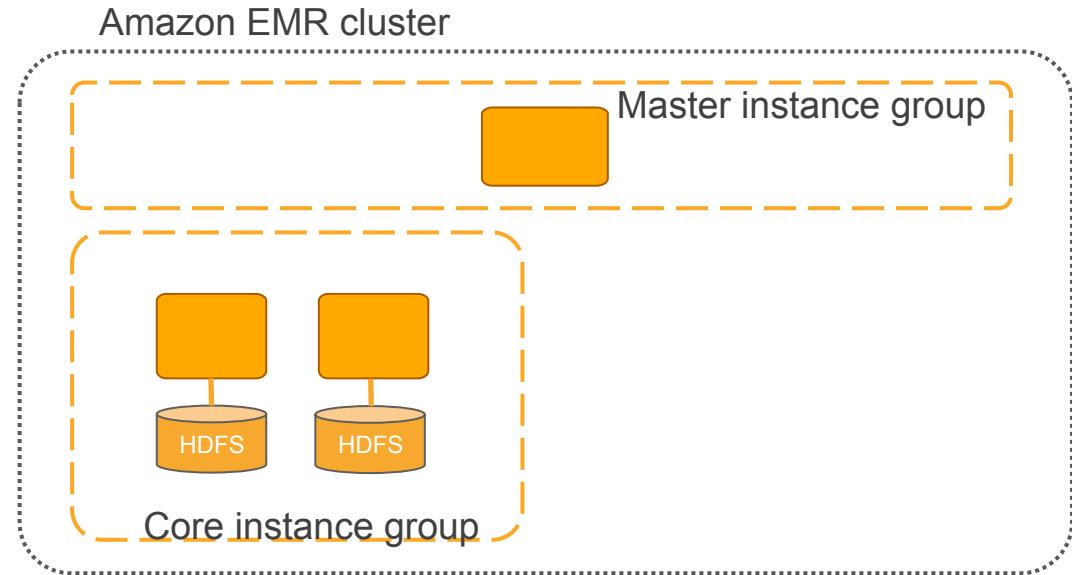


Match compute demands with cluster sizing.

# Core Nodes

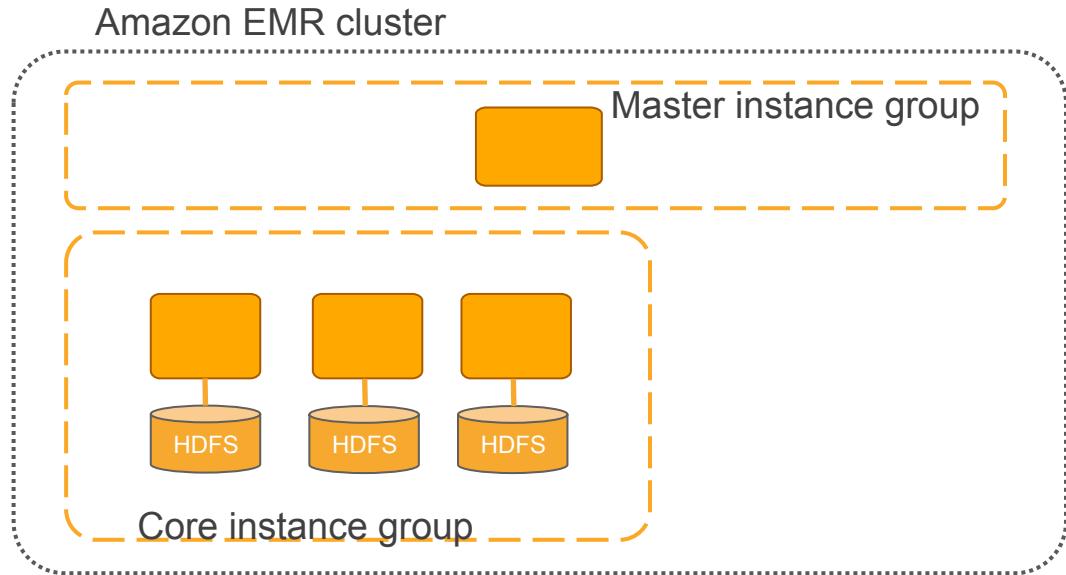
Run TaskTrackers  
(Compute)

Run DataNode  
(HDFS)



# Core Nodes

- Can add core nodes
- More HDFS space
- More CPU/mem

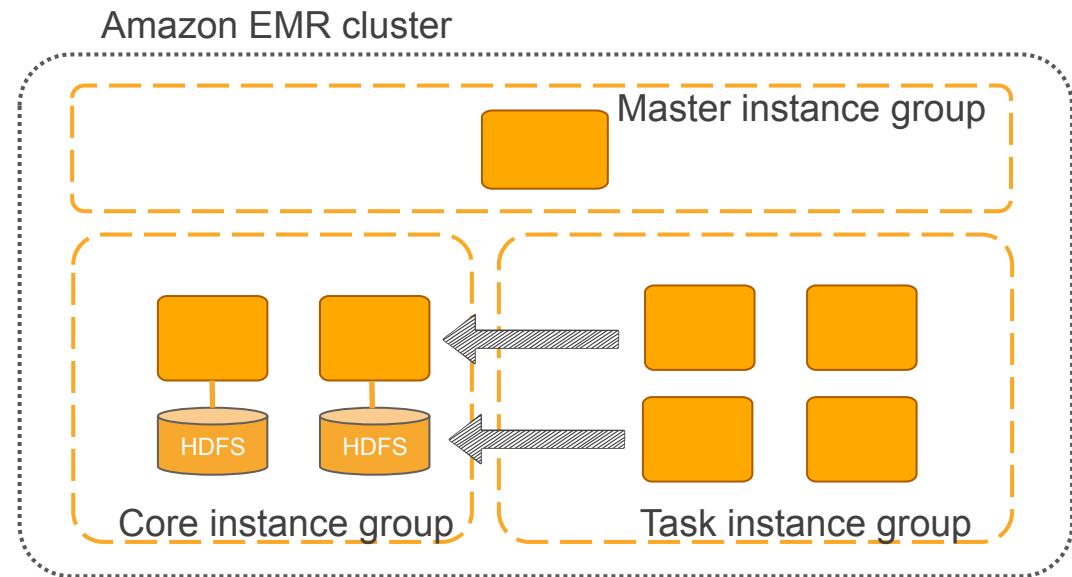


# Amazon EMR Task Nodes

Run TaskTrackers

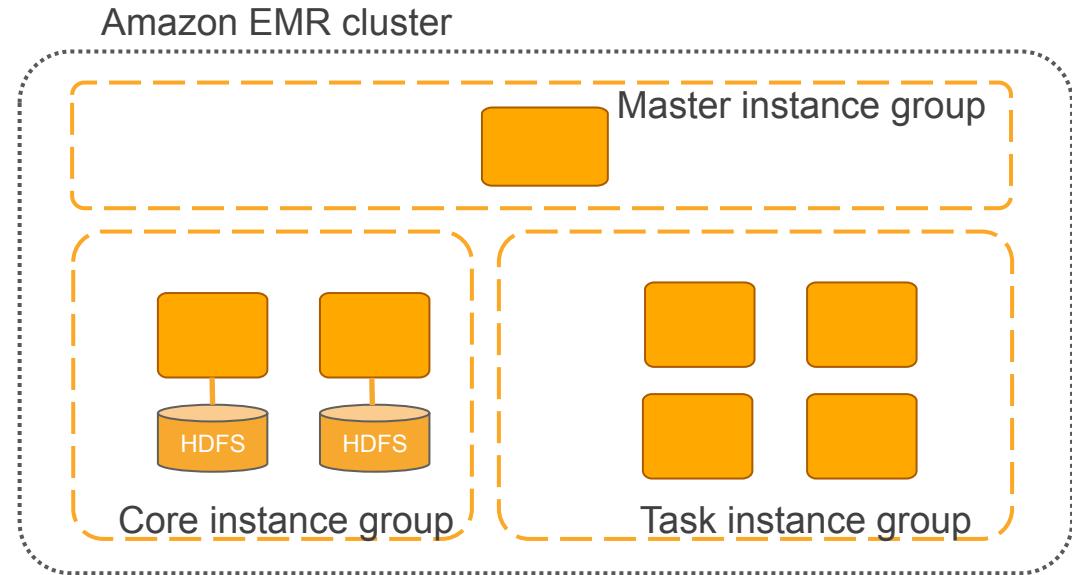
No HDFS

Reads from core node  
HDFS



# Amazon EMR Task Nodes

Can Add Task Nodes

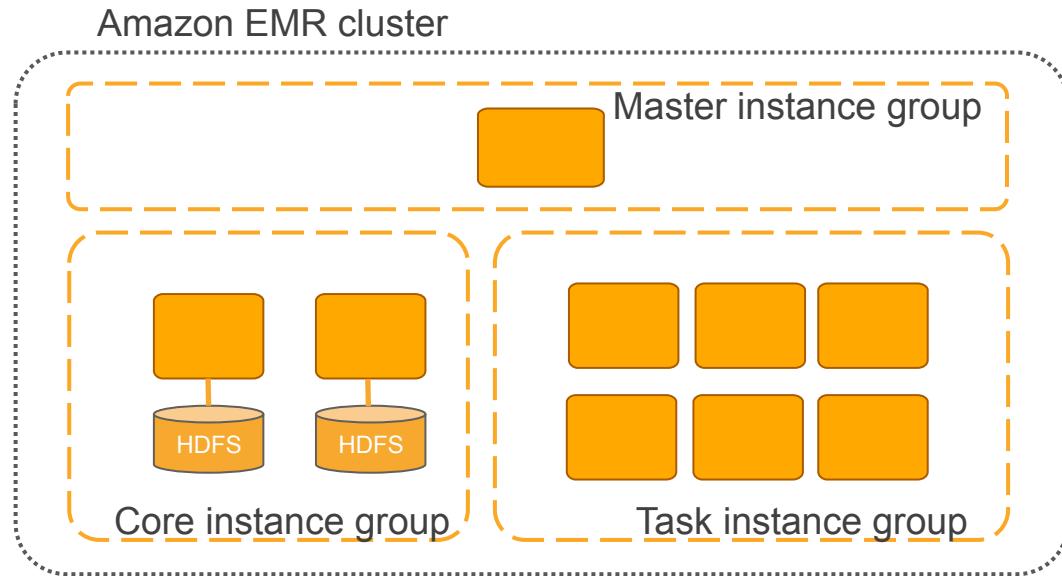


# Amazon EMR Task Nodes

More CPU power

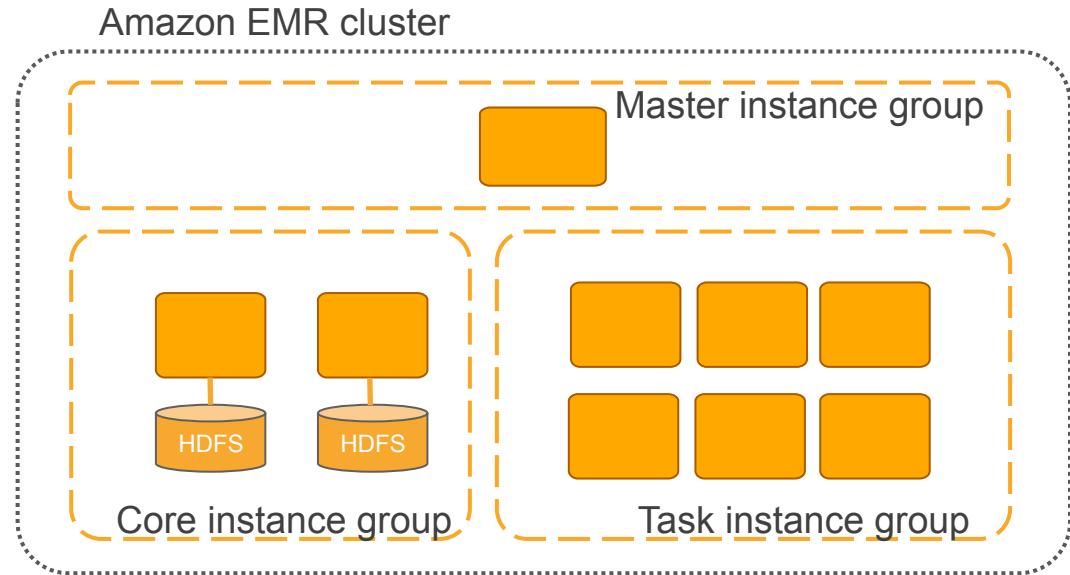
More memory

More network  
throughput



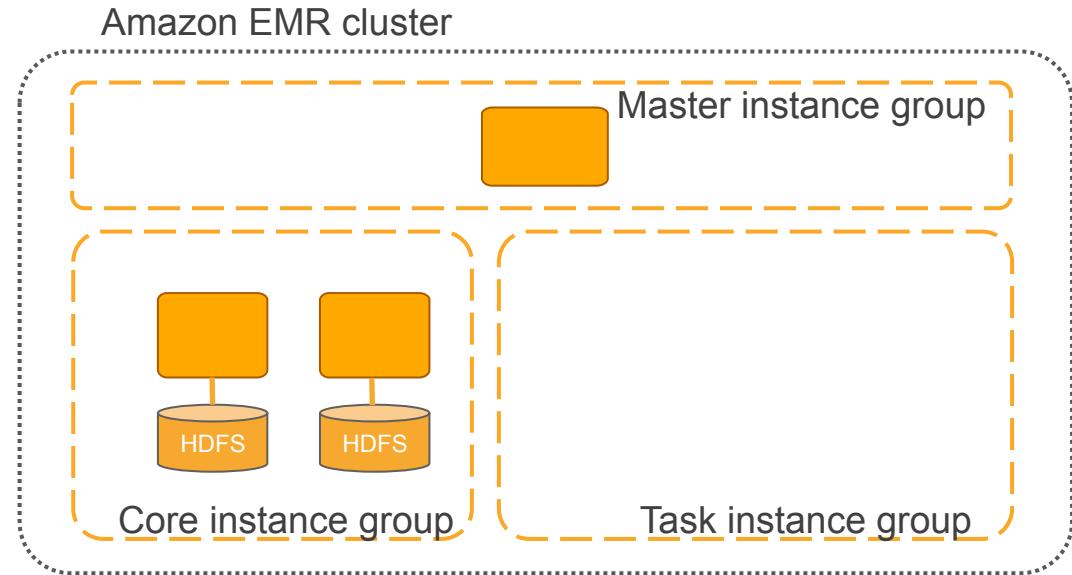
# Amazon EMR Task Nodes

You can remove Task  
Nodes



# Amazon EMR Task Nodes

Save Cost when  
Processing power is not  
required



# Easy to use Spot Instances

Meet SLA at predictable cost

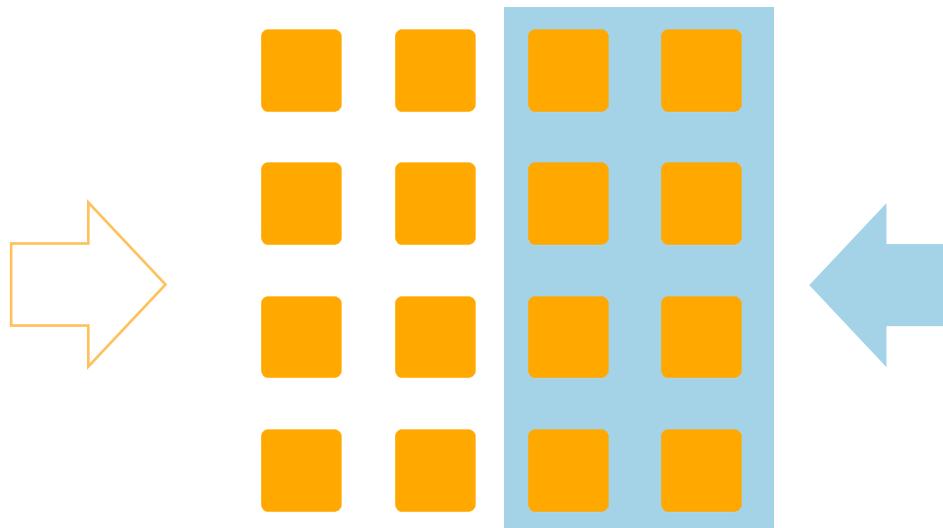
On-demand for  
core nodes

Standard  
Amazon EC2  
pricing for  
on-demand  
capacity

Exceed SLA at lower cost

Spot Instances  
for task nodes

Up to 86% lower  
on average  
off  
on-demand  
pricing



# Serial Processing



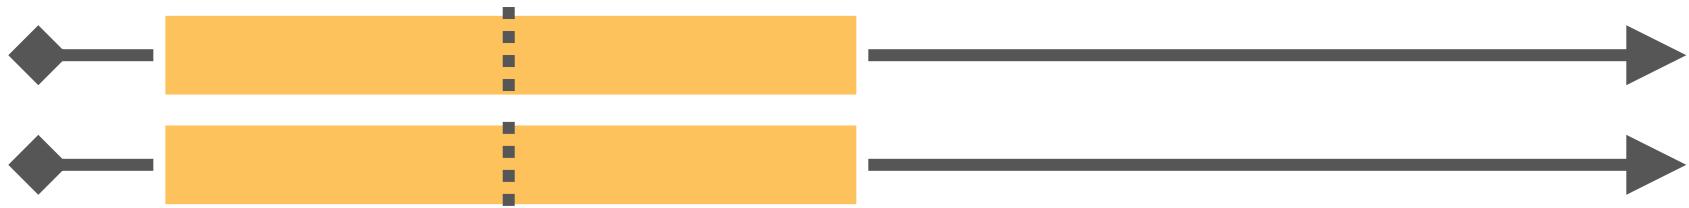
**COST:** 4h x \$2.1 = **\$8.4**  
**RENDERING TIME:** 4h

**1 Machine for 10 Hours**

$\approx$

**10 Machines for 1 Hour**

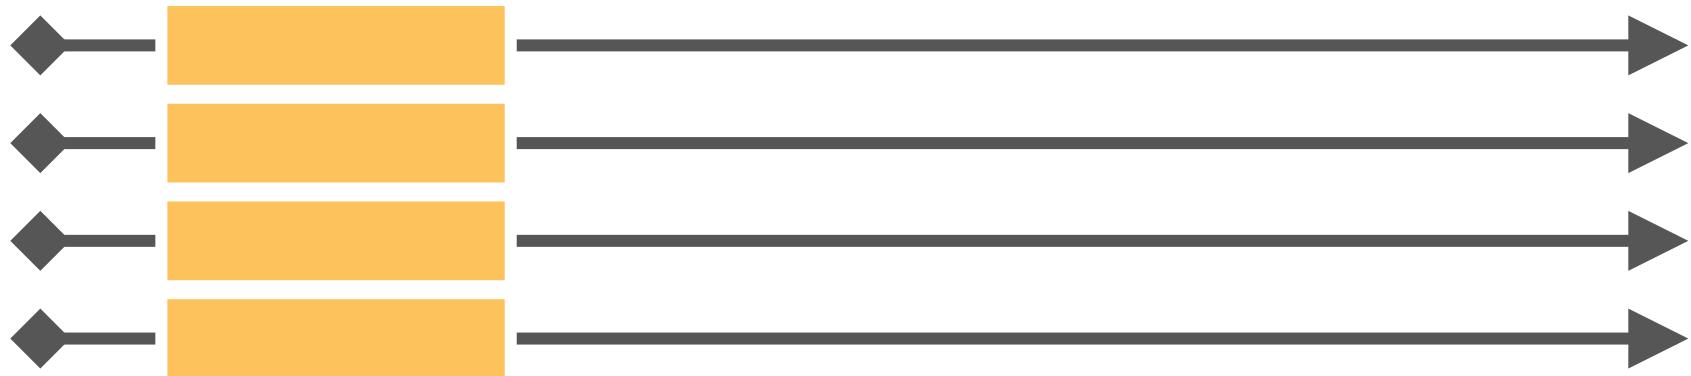
# Parallel Processing



COST:  $2 \times 2\text{h} \times \$2.1 = \$8.4$

RENDERING TIME: **2h**

# Embarrassingly Parallel Processing



COST:  $4 \times 1\text{h} \times \$2.1 = \$8.4$

RENDERING TIME: **1h**

# Pick the right version

# Pick your Hadoop Version & Applications

EMR < 4.x:

- AMI version had fixed Hadoop version, Hive version etc.
- Additional applications via **Bootstrap Actions**  
Lots of different ones available from EMR & on github
- Configuration via Bootstrap Actions

<https://github.com/awslabs/emr-bootstrap-actions>



# Pick your Hadoop Version & Applications

Starting with EMR 4.x:

- Release tags are independent of AMIs
- Pick from supported applications directly
  - Based on Apache BigTop
- Configuration objects
  - unified JSON syntax, mapped to configuration files

**Recommendation: Use EMR 4.x by default!**



# Transient or persistent cluster?

# Transient or persistent cluster?

Storing data on Amazon S3 give you more flexibility to run & shutdown Hadoop clusters when you need to.

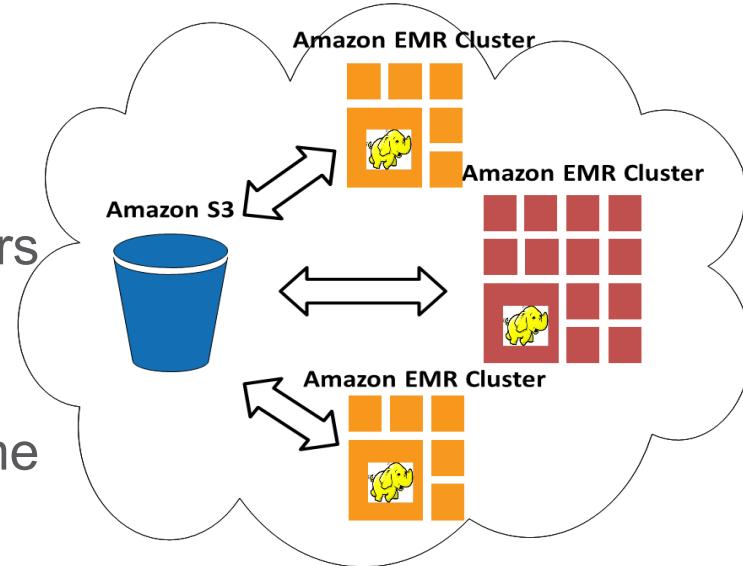
# Amazon S3 as your persistent data store

## Amazon S3

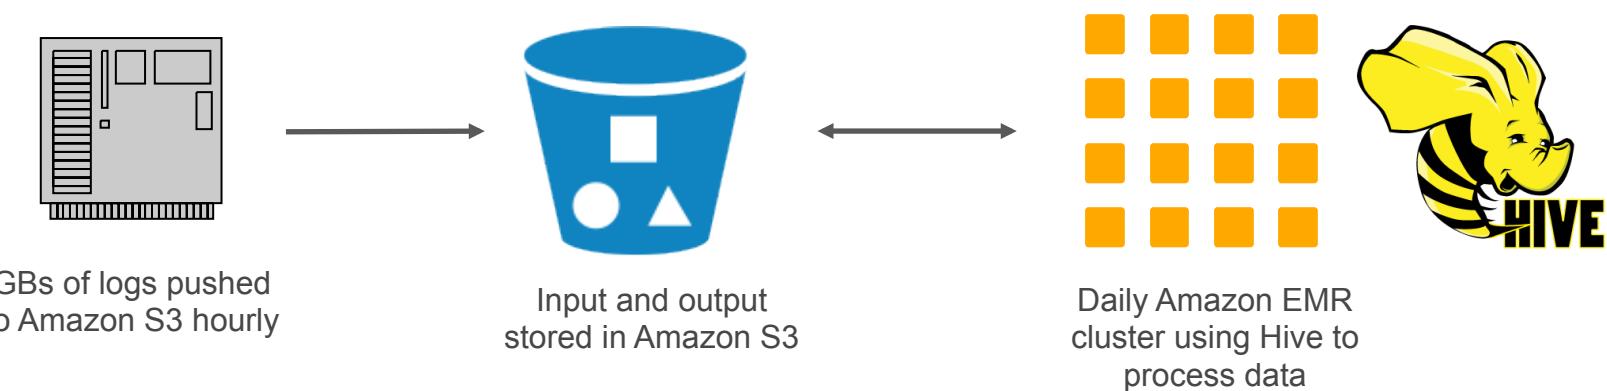
- Designed for 99.999999999% durability
- Separate compute and storage

Resize and shut down Amazon EMR clusters  
with no data loss

Point multiple Amazon EMR clusters at same  
data in Amazon S3



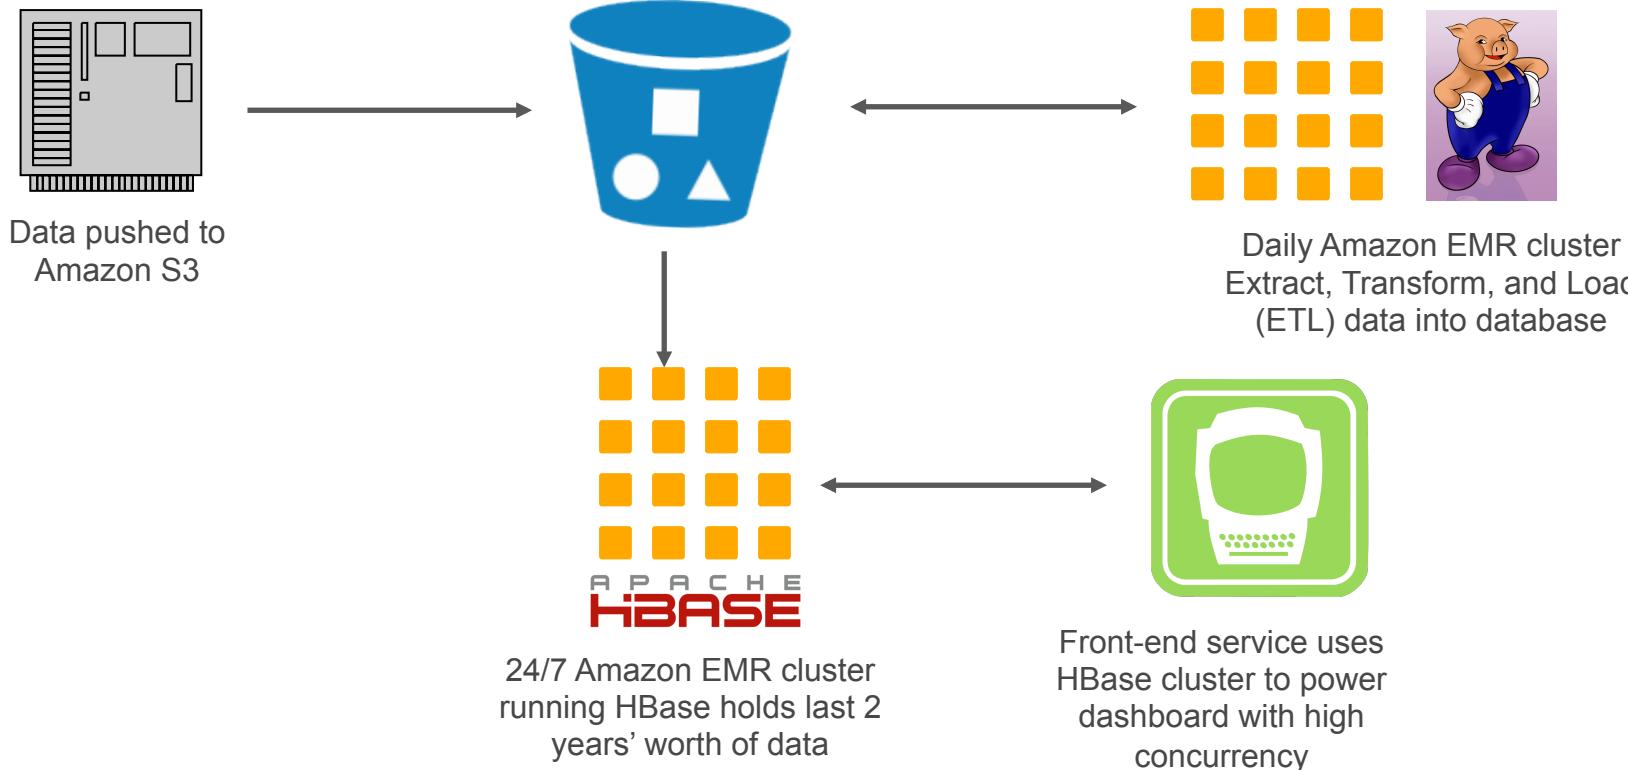
# Amazon EMR example #1: Batch processing



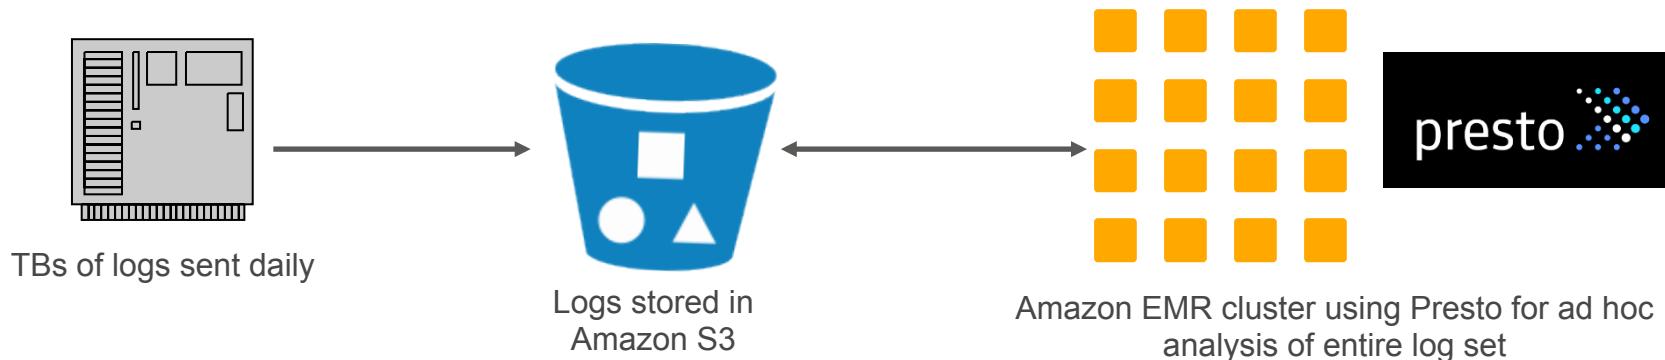
250 Amazon EMR jobs per day, processing 30 TB of data

<http://aws.amazon.com/solutions/case-studies/yelp/>

# Amazon EMR example #2: Long-running cluster



# Amazon EMR example #3: Interactive query



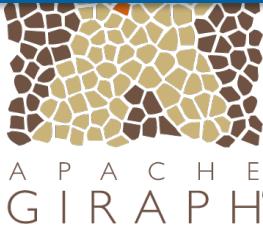
Interactive query using Presto on multipetabyte warehouse

<http://techblog.netflix.com/2014/10/using-presto-in-our-big-data-platform.html>

# Steps & Clusters

# Applications

# Vibrant Ecosystem



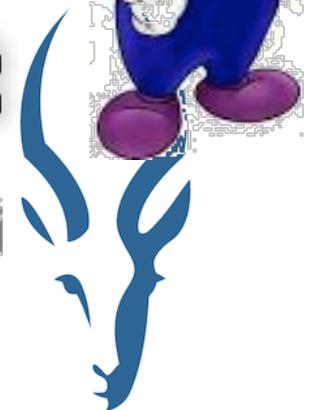
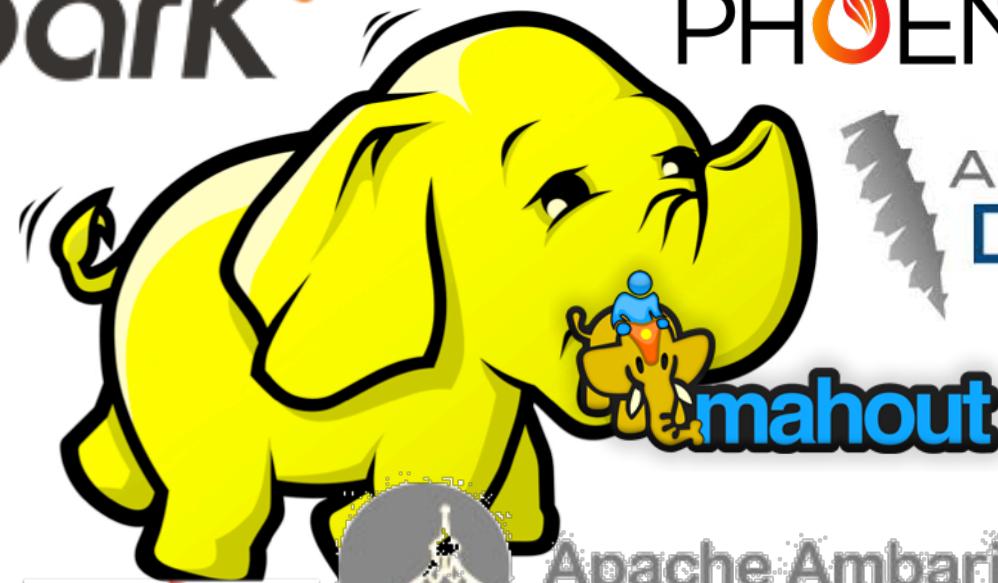
Apache  
Spark



TACHYON

APACHE  
PHOENIX

APACHE  
HBASE

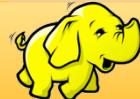


Apache Ambari™  
ACCUMULO

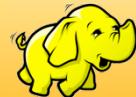
amazon  
web services



MR job



MapReduce



HDFS

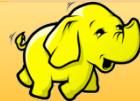


Amazon S3

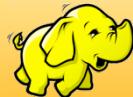
Local filesystem



MR job



MapReduce v2: YARN

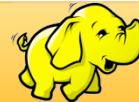


HDFS



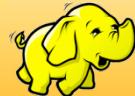
Amazon S3

Local filesystem



MR job

MapReduce v2: YARN

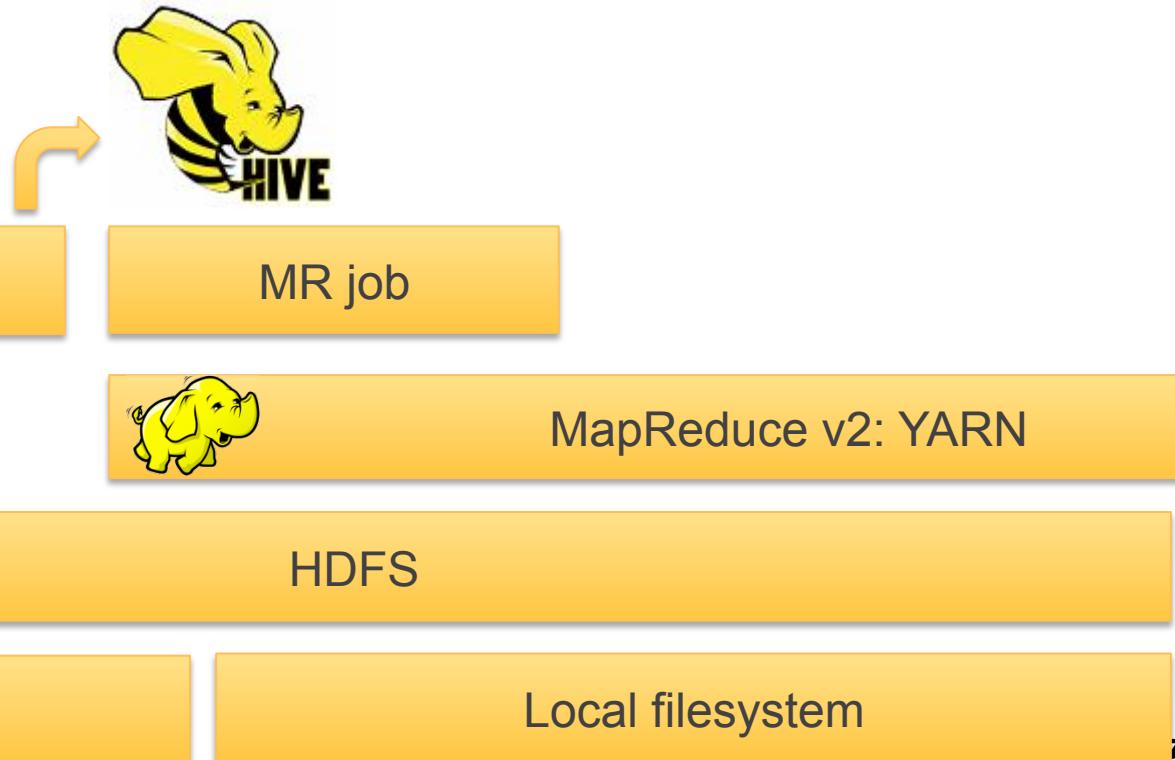


HDFS



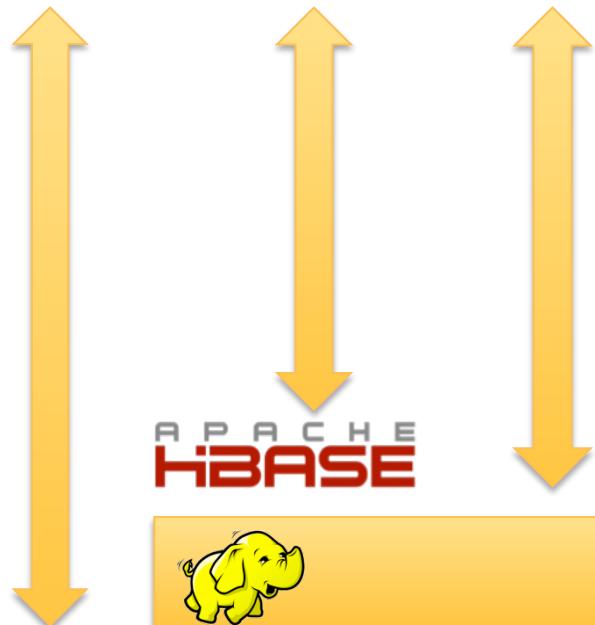
Amazon S3

Local filesystem

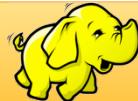




Presto



MR job



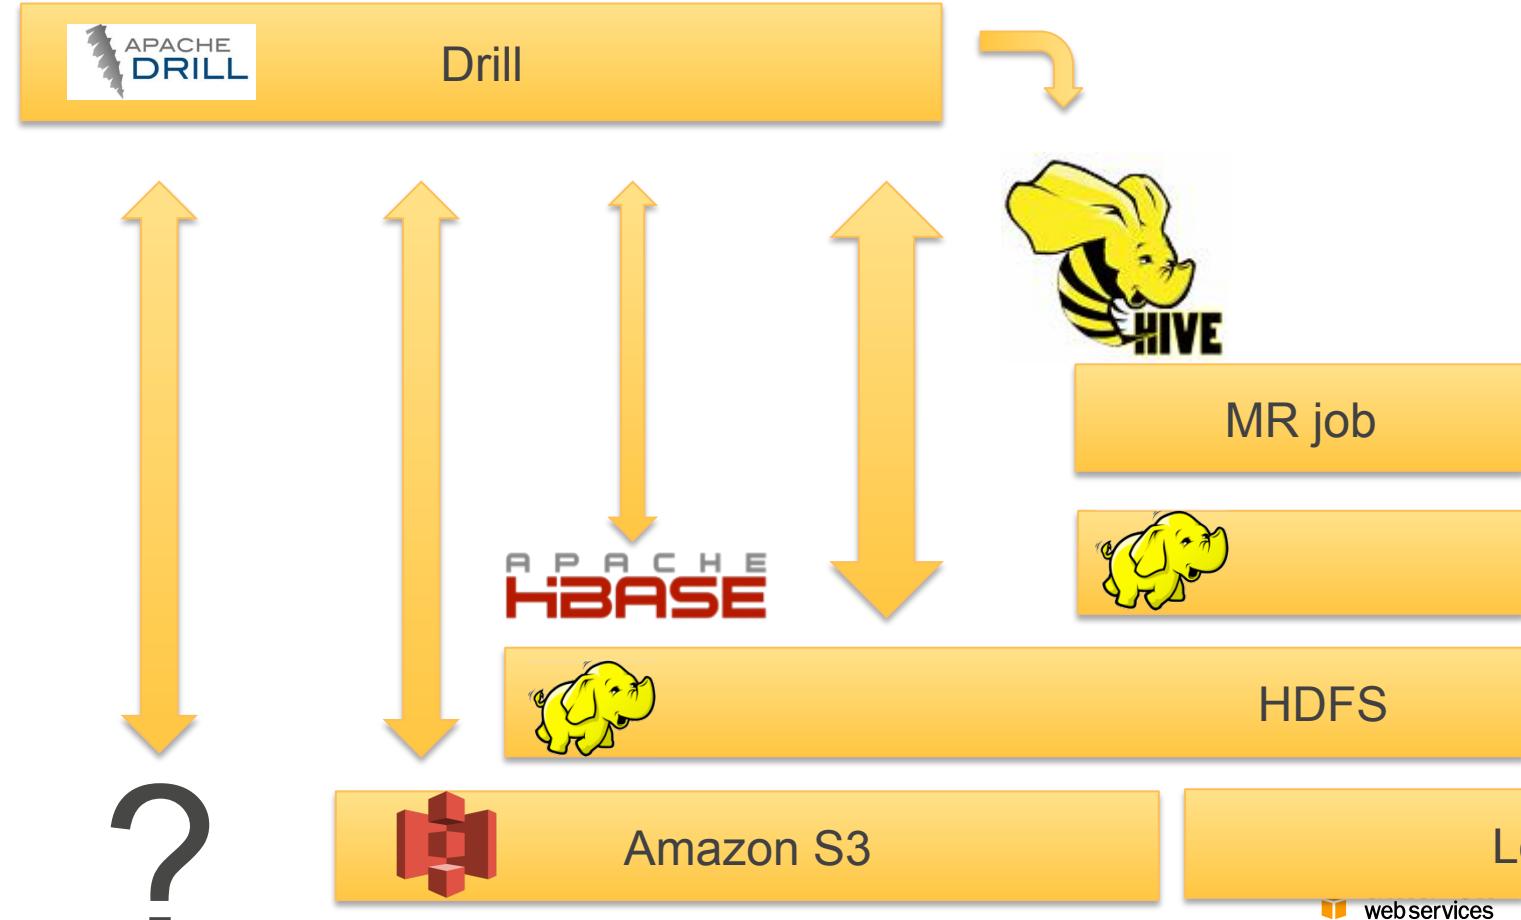
MapReduce v2: YARN

HDFS



Amazon S3

Local filesystem

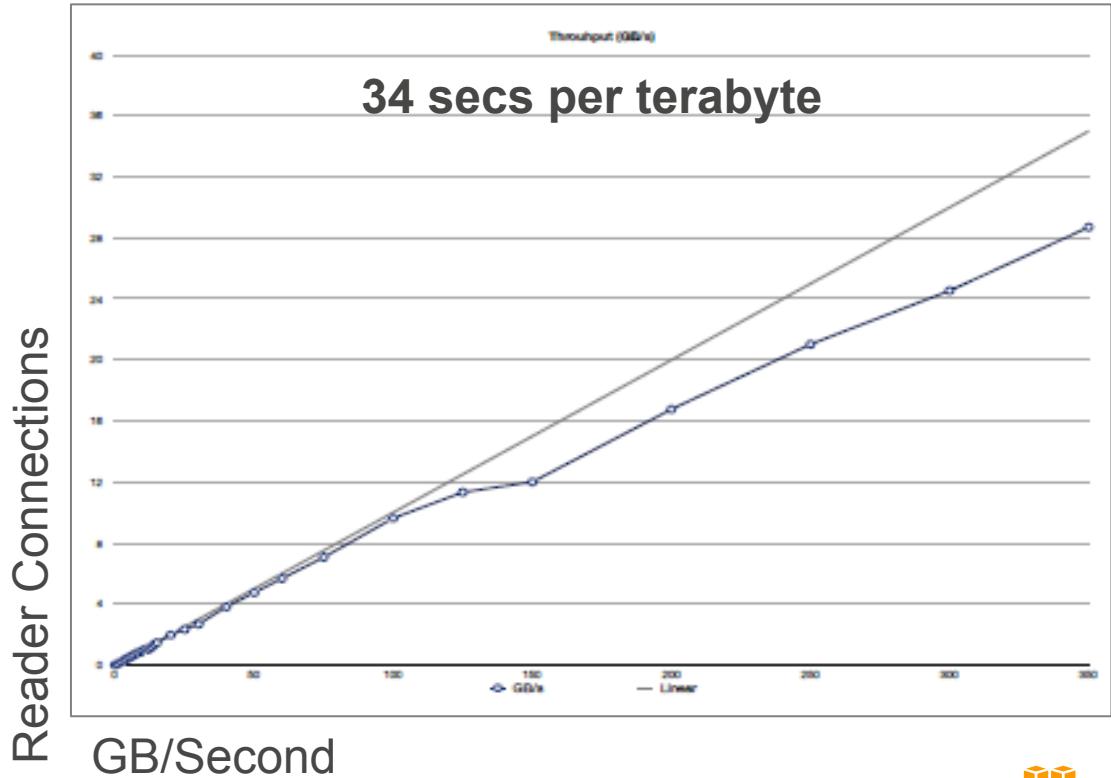


# Amazon S3 and HDFS

# Amazon S3 has near linear scalability



S3 Streaming Performance  
100 VMs; 9.6GB/s; \$26/hr  
350 VMs; 28.7GB/s; \$90/hr



# EMRFS makes it easier to leverage Amazon S3

Better performance and error handling options

Transparent to applications – just read/write to “s3://”

Consistent view

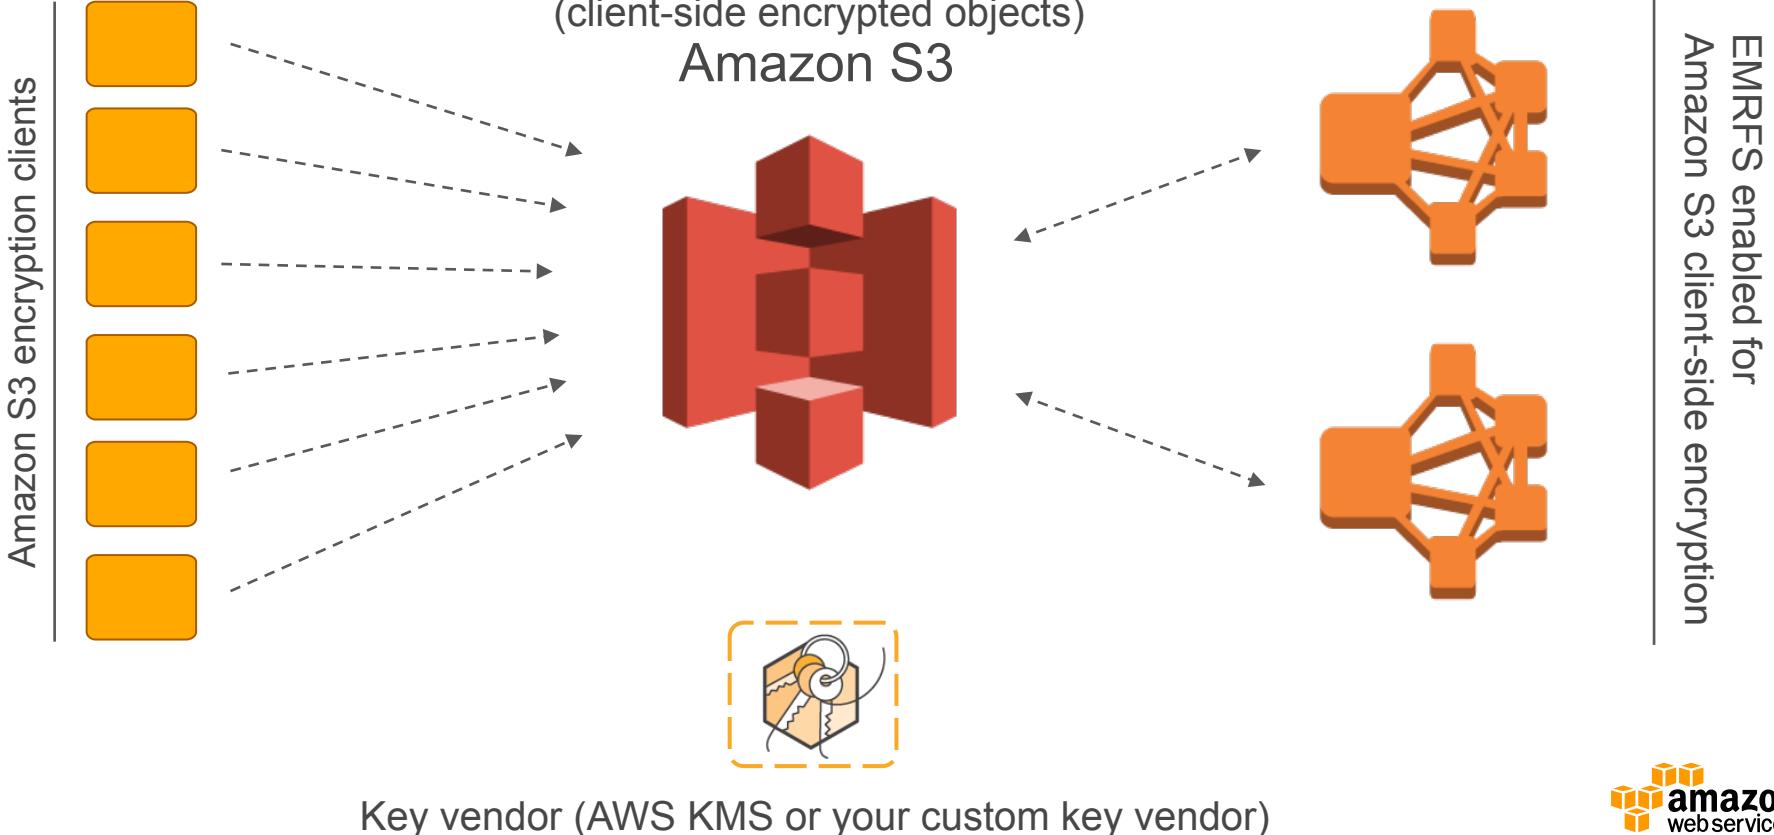
- For consistent list and read-after-write for new puts

Support for Amazon S3 server-side and client-side encryption

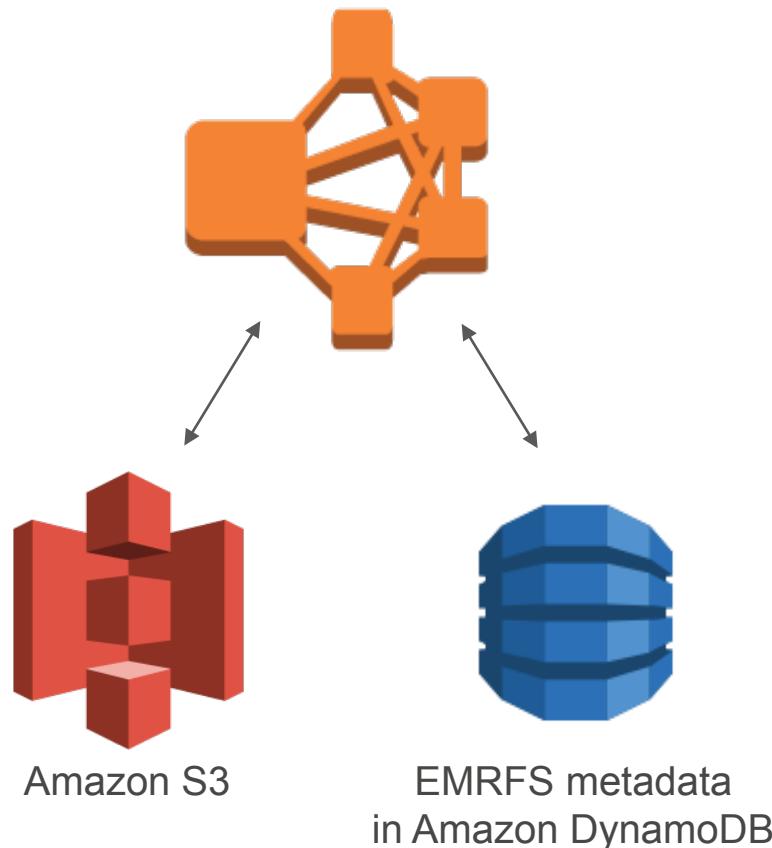
Faster listing using EMRFS metadata



# EMRFS support for Amazon S3 client-side encryption



# Fast listing of Amazon S3 objects using EMRFS metadata



- List and read-after-write consistency
- Faster list operations

Number of objects	Without Consistent Views	With Consistent Views
1,000,000	147.72	29.70
100,000	12.70	3.69

\*Tested using a single node cluster with a m3.xlarge instance.

# Optimize to leverage HDFS

## Iterative workloads

- If you're processing the same dataset more than once

## Disk I/O intensive workloads

Persist data on Amazon S3 and use S3DistCp to copy to HDFS for processing.

# Running Spark on Amazon EMR



# Launch the latest Spark version

July 15 – Spark 1.4.1 GA release

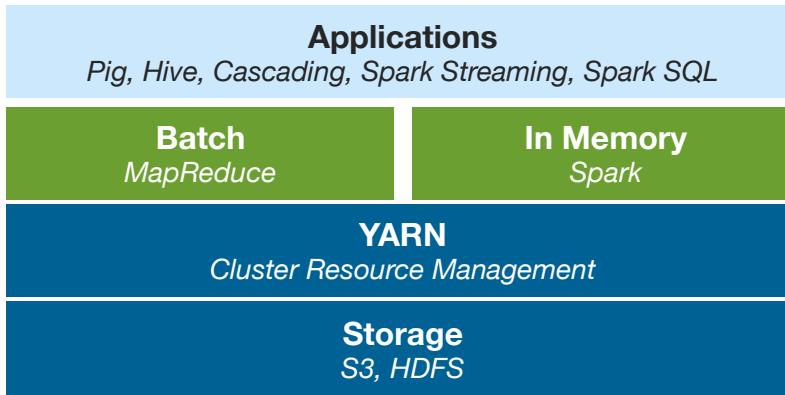
July 24 – Spark 1.4.1 available on Amazon EMR

September 9 – Spark 1.5.0 GA release

September 30 – Spark 1.5.0 available on Amazon EMR

**< 3 week cadence with latest open source release**

# Amazon EMR runs Spark on YARN



- Dynamically share and centrally configure the same pool of cluster resources across engines
- Schedulers for categorizing, isolating, and prioritizing workloads
- Choose the number of executors to use, or allow YARN to choose (dynamic allocation)
- Kerberos authentication

# Create a fully configured cluster in minutes

AWS Management  
Console

Software configuration

Vendor  Amazon  MapR

Release emr-4.1.0 A release installed

Applications  Spark: Spark 1.5.0 on Hadoop 2.6.0 YARN  
 All Applications: Hadoop 2.6.0, Hive 1.0.0, Hue 3.7.1, Mahout 0.11.0, Pig 0.14.0, and Spark 1.5.0  
 Core Hadoop: Hadoop 2.6.0, Hive 1.0.0, and Pig 0.14.0  
 Presto-Sandbox: Presto 0.119 with Hadoop 2.6.0 HDFS and Hive 1.0.0 Metastore

AWS Command Line  
Interface (CLI)

```
jonfritz@...: ~
$ aws emr create-cluster --release-label emr-4.1.0 --instance-type r3.xlarge --instance-count 8 --applications Name=Spark --ec2-attributes KeyName=MyKey --use-default-roles
```

Or use an AWS SDK directly with the Amazon EMR API



# Or easily change your settings

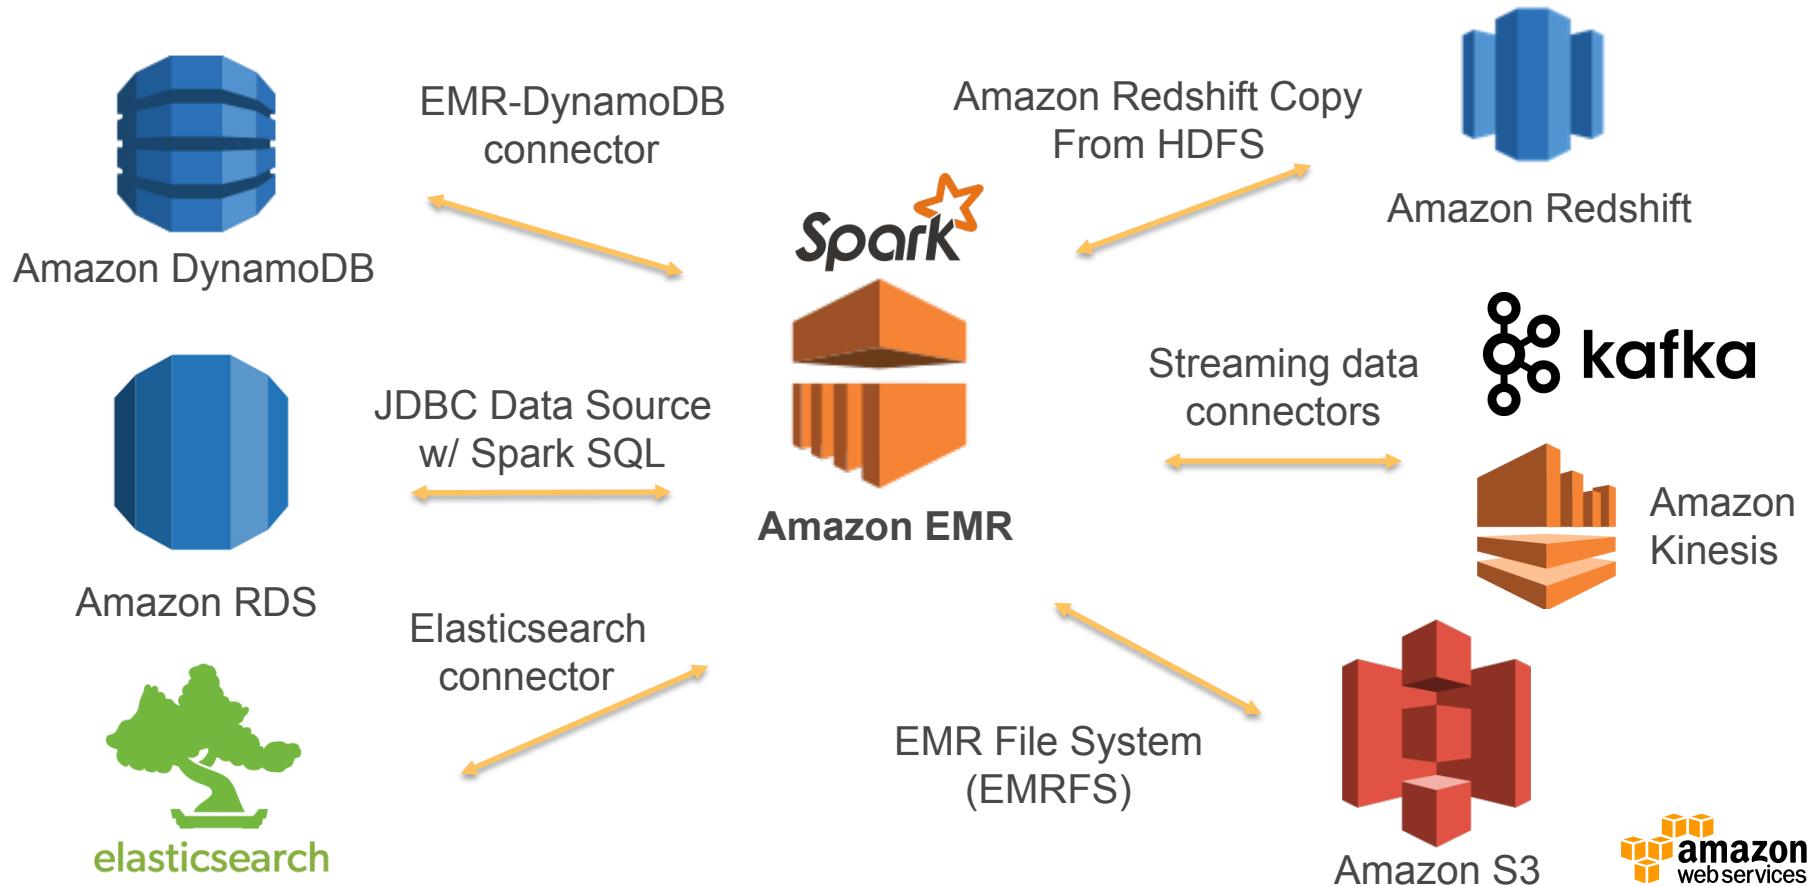
## ▼ Edit software settings (optional)

- i** Specify the new configuration values for the applications on your cluster. Available configuration files to edit: capacity-scheduler, core-site, hadoop-env, hadoop-log4j, hdfs-site, httpfs-env, https-site, mapred-env, mapred-site, yarn-env, yarn-site, hive-env, hive-exec-log4j, hive-log4j, hive-site, pig-properties, pig-log4j

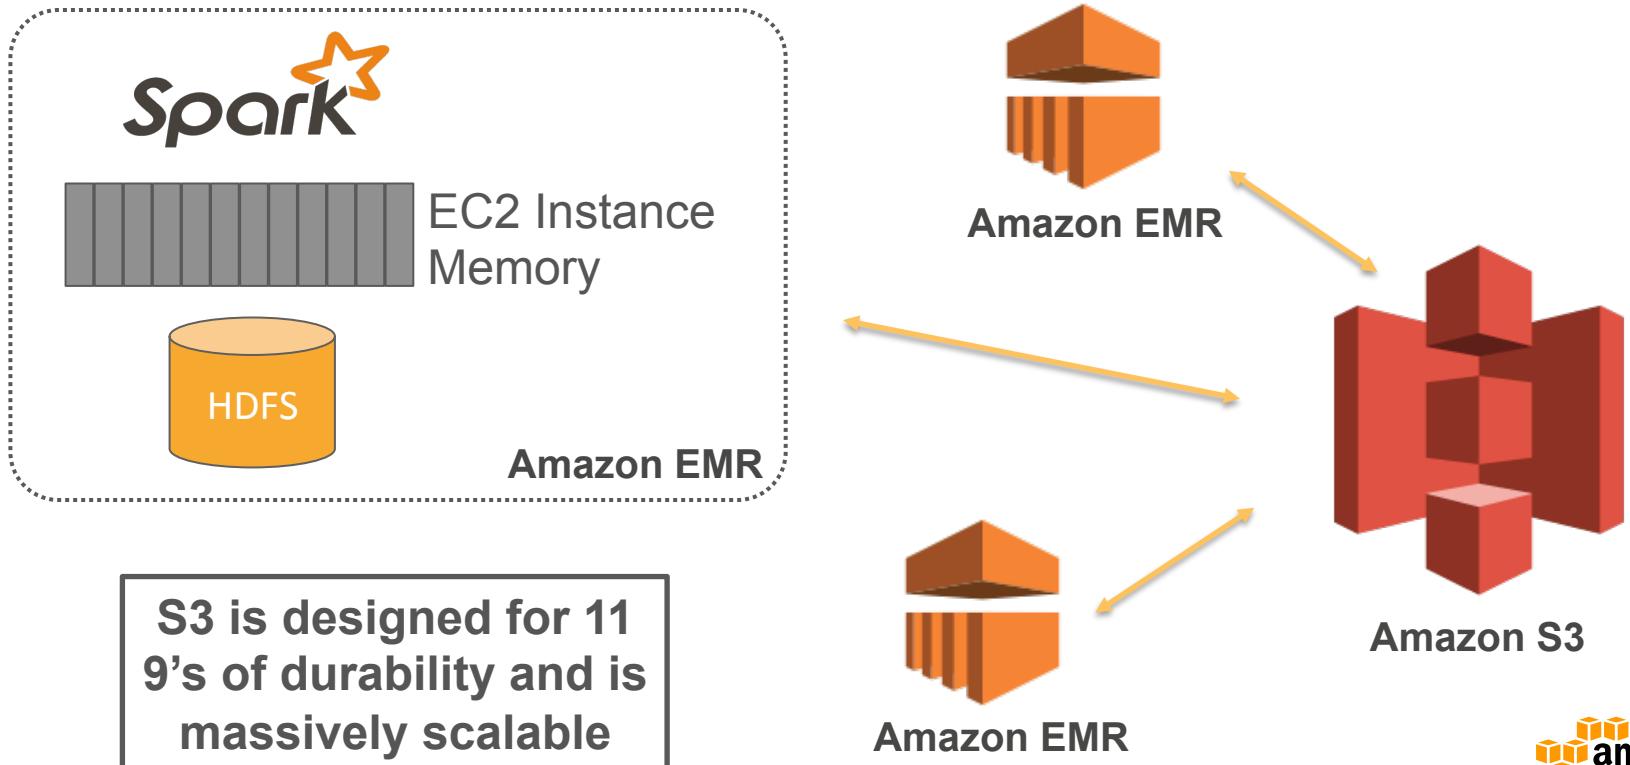
**Change settings**  Enter configuration  Load JSON from S3

```
classification=yarn-site,properties=[yarn.nodemanager.resource.cpu-
vcores=4,yarn.nodemanager.resource.memory-mb=1024,yarn.scheduler.maximum-
allocation-mb=512,yarn.scheduler.minimum-allocation-mb=256] classification=spark-
defaults,properties=[spark.executor.memory=4G,spark.driver.cores=2]
```

# Many storage layers to choose from

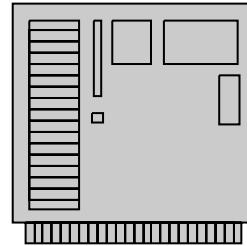


# Decouple compute and storage by using S3 as your data layer



# Easy to run your Spark workloads

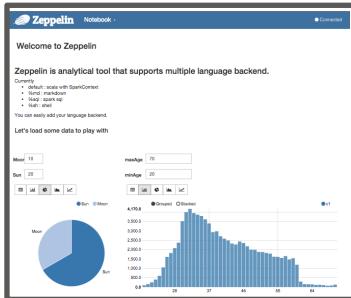
Submit a Spark application



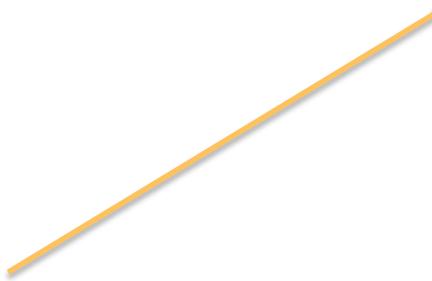
Amazon EMR Step API



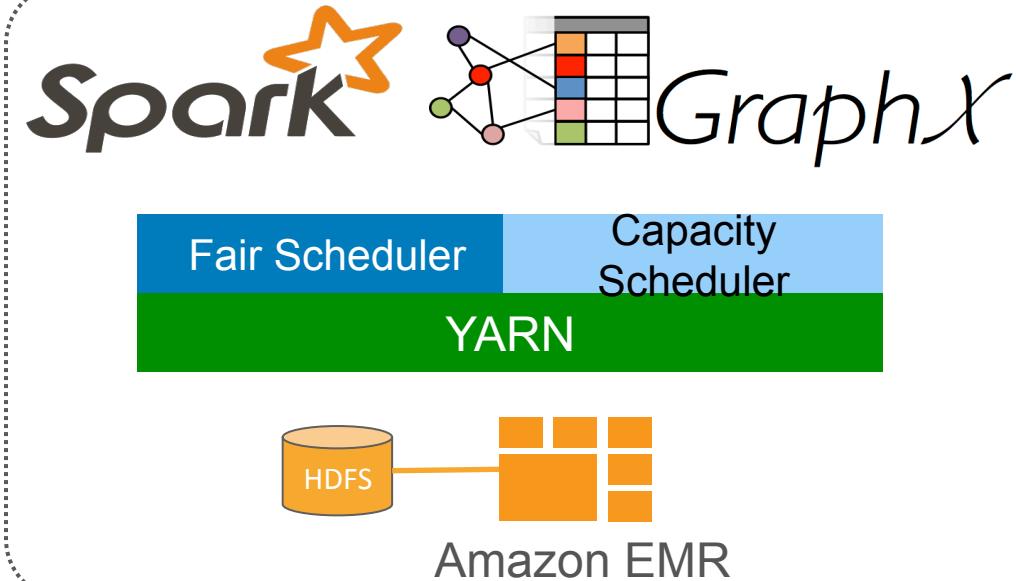
Amazon EMR



SSH to master node  
(Spark Shell)



# Spark on YARN

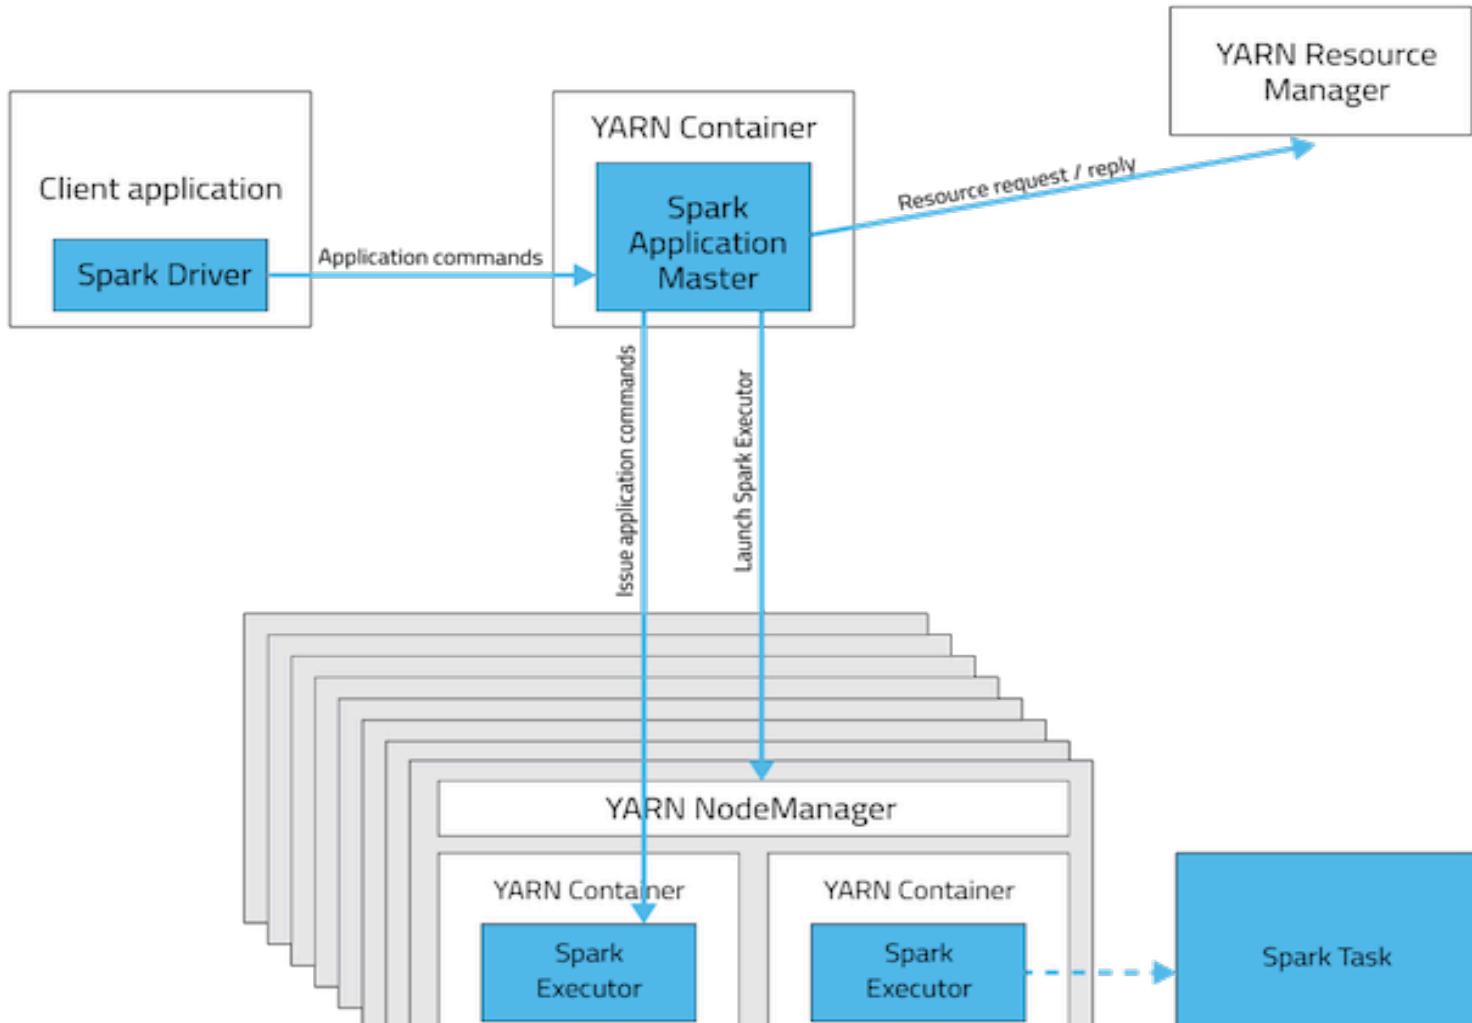


YARN is Hadoop 2's Cluster Manager  
Spark Standalone Cluster Manager is FIFO, 100% Core Allocation  
Not ideal for multi-tenancy  
YARN Default is 'Capacity Scheduler'  
Focus on Throughput  
Configure Queues  
'Fair Scheduler' also available  
Focus on even distribution

# Why run Spark on YARN?

- YARN allows you to dynamically share and centrally configure the same pool of cluster resources between all frameworks that run on YARN. You can throw your entire cluster at a MapReduce job, then use some of it on an Impala query and the rest on Spark application, without any changes in configuration.
- You can take advantage of all the features of YARN schedulers for categorizing, isolating, and prioritizing workloads.
- Spark standalone mode requires each application to run an executor on every node in the cluster, whereas with YARN you choose the number of executors to use.
- YARN is the only cluster manager for Spark that supports security. With YARN, Spark can run against Kerberized Hadoop clusters and uses secure





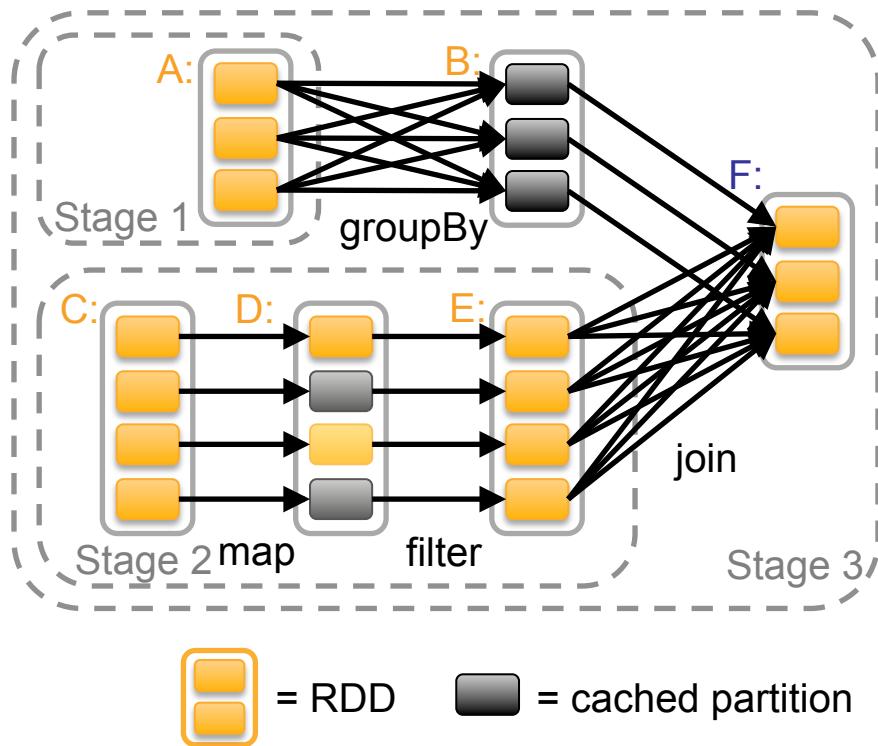
# Task Scheduler

Supports general task graphs

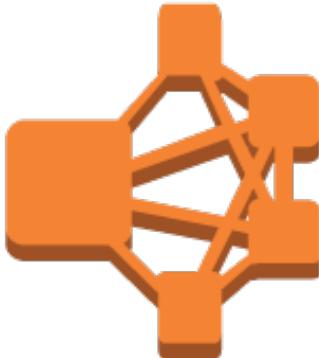
Pipelines functions where possible

Cache-aware data reuse & locality

Partitioning-aware to avoid shuffles



# Secure Spark clusters – encryption at rest



Amazon S3

## On-Cluster

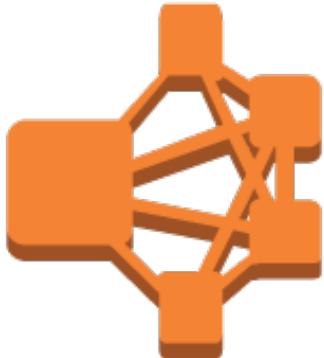
HDFS transparent encryption (AES 256)  
[new on release emr-4.1.0]

Local disk encryption for temporary files  
using LUKS encryption via bootstrap action

## Amazon S3

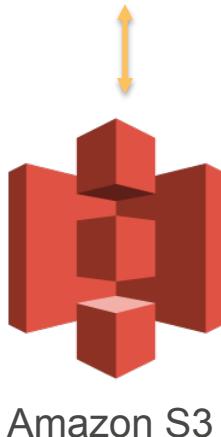
EMRFS support for Amazon S3 client-side  
and server-side encryption (AES 256)

# Secure Spark clusters – encryption in flight



**Internode communication on-cluster**  
Blocks are encrypted in-transit in HDFS when using transparent encryption

Spark's Broadcast and FileServer services can use SSL. BlockTransferService (for shuffle) can't use SSL (SPARK-5682).

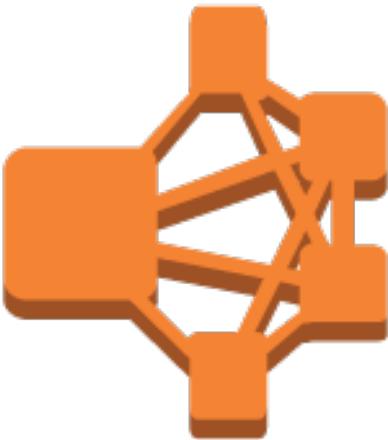


**S3 to Amazon EMR cluster**  
Secure communication with SSL

Objects encrypted over the wire if using client-side encryption



# Secure Spark clusters – additional features



## Permissions:

- **Cluster level:** IAM roles for the Amazon EMR service and the cluster
- **Application level:** Kerberos (Spark on YARN only)
- **Amazon EMR service level:** IAM users

**Access:** VPC, security groups

**Auditing:** AWS CloudTrail

# Thank you!