




IOT DATA STREAMING WITH APACHE KAFKA

From the data to the wisdom

Paolo Patierno
Principal Software Engineer @ Red Hat
15/12/2018

WHO AM I ?

 @ppatierno

- Principal Software Engineer @ Red Hat
 - Messaging & IoT team
- Lead/Committer @ Eclipse Foundation
 - Hono, Paho and Vert.x projects
- Microsoft MVP Azure/IoT
- Dad of two
- VR46's fan (I don't like MM93)
- Try to be a runner ...



IOT USE CASES

- Smart city
- Healthcare
- Connected cars
- Logistics
- Home automation
- Airlines
- Farmers
- You!



WHAT THEY HAVE IN COMMON?

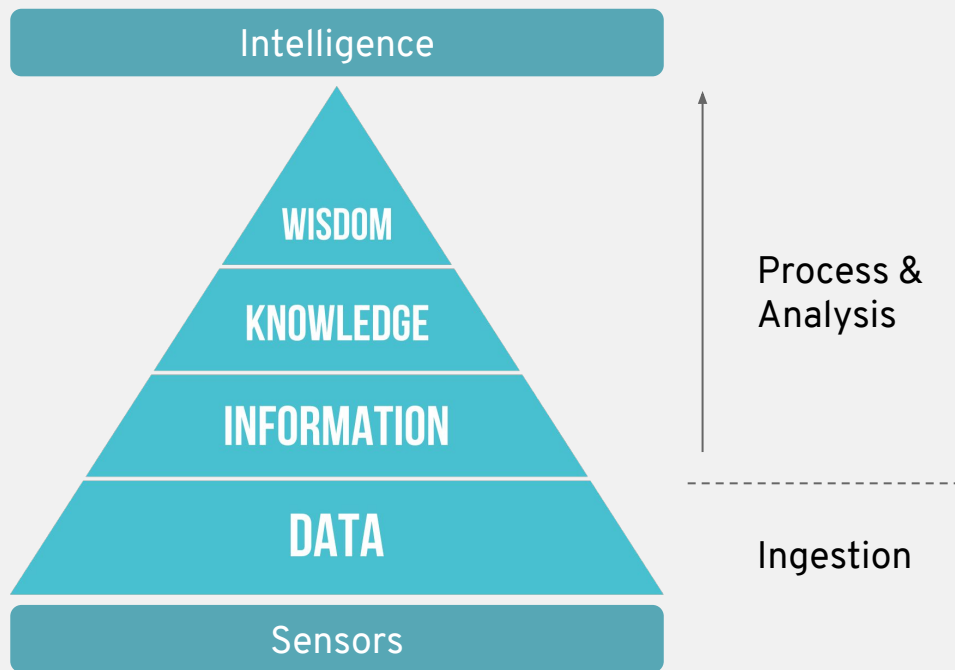
IOT WORKLOAD

- Huge amount of data to ingest
- Unbounded dataset
- Need for real time analytics
- Need for real time reactions
- Sometimes the need to join :
 - with static data
 - with historical data



IOT WORKLOAD

From raw data to the intelligence



THE NEED FOR A PARADIGM DIFFERENT FROM BATCH PROCESSING

STREAM PROCESSING

- A type of data processing engine that is designed with infinite datasets in mind
- Working with data as they arrive
- Working with an infinite stream of data
 - Moving boundaries with data in “windows”
- Tradeoffs between latency/cost/correctness

STREAM PROCESSING

Versus batch processing

Compared to batch based flows, streaming flows ...

- ... are often more time-sensitive
- ... could be larger in volume
- ... may or may not be of long-term value after processing

So that you can ...


- ... capture and aggregate millions of events per second
- ... instantly take action
- ... respond to changing conditions

STREAM PROCESSING

Because time matters

Time

- Event time, which is the time at which events actually occurred
- Processing time, which is the time at which events are observed in the system



STREAM PROCESSING & IOT

What are the good parts for IoT?

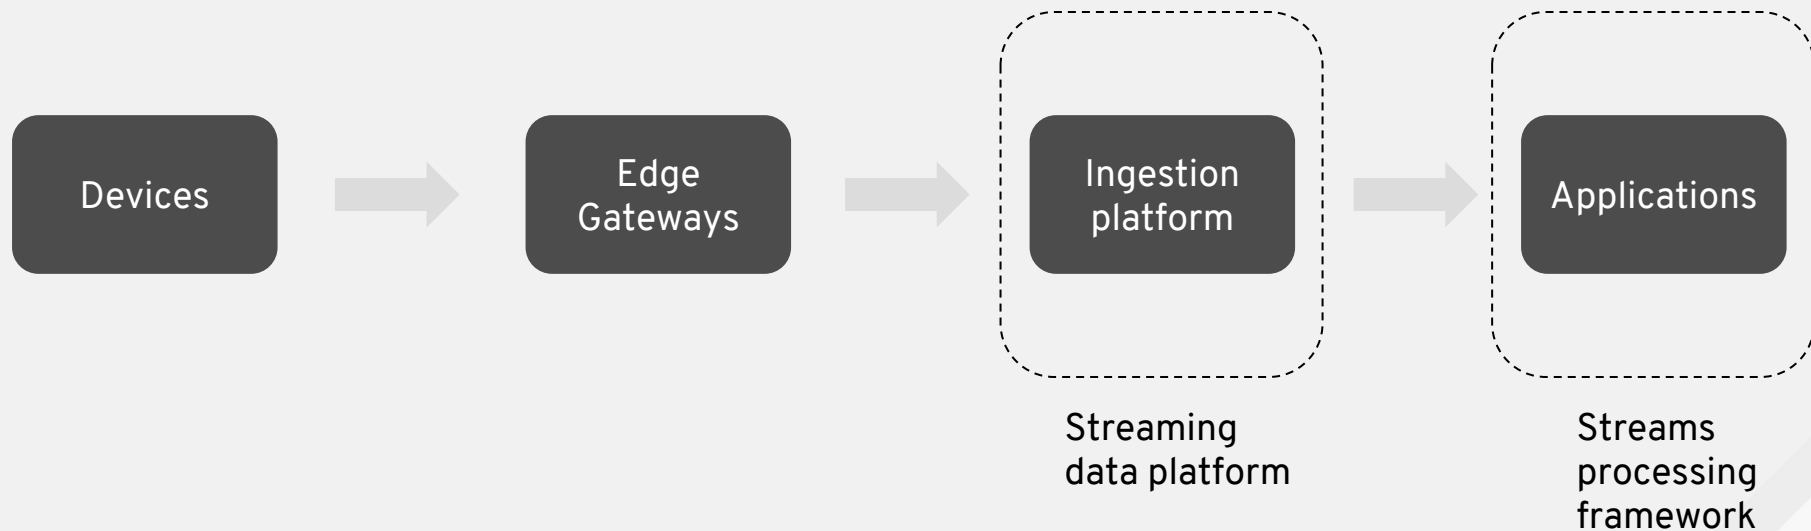
- Ordering
 - Handling devices data in the order they are sent
 - Handling “out of order” data when the device is not able to guarantee order
- Timing
 - Allowing “event time” which makes much sense in IoT scenarios
 - Supporting “windowing” for getting analytics in a specific time window
- State management
 - Be fast on joining real time data with metadata to enrich information content
 - Be fast on data aggregation
 - Be able to join different streams of data, from different devices

STREAM PROCESSING & IOT

What are the good parts for IoT?

- Re-processing
 - Supporting the possibility to re-read the stream of data
- Fault tolerance
 - Allowing to get data from devices continuously even in case of consumers crash
- Scalability/Partitioning
 - Being able to handle bursts of traffic
 - Being able to handle more and more devices data

STREAM PROCESSING & IOT



STREAM PROCESSING & IOT


It's a wild west out there

- Streaming data platform
 - Apache Kafka
- Streams processing framework
 - Apache Spark (streaming)
 - Apache Samza
 - Apache Flink
 - ...



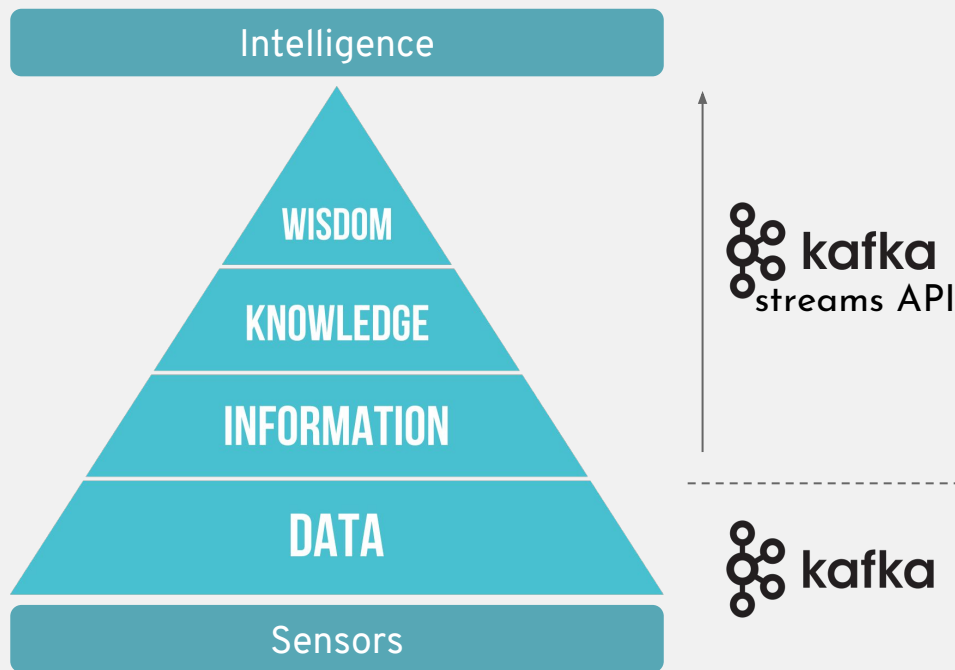
**DO WE REALLY NEED MORE THAN
ONE PLATFORMS IN PLACE?**

STREAM PROCESSING & IOT



IOT WORKLOAD

From raw data to the intelligence



APACHE KAFKA

WHAT IS APACHE KAFKA



A publish/subscribe messaging system?

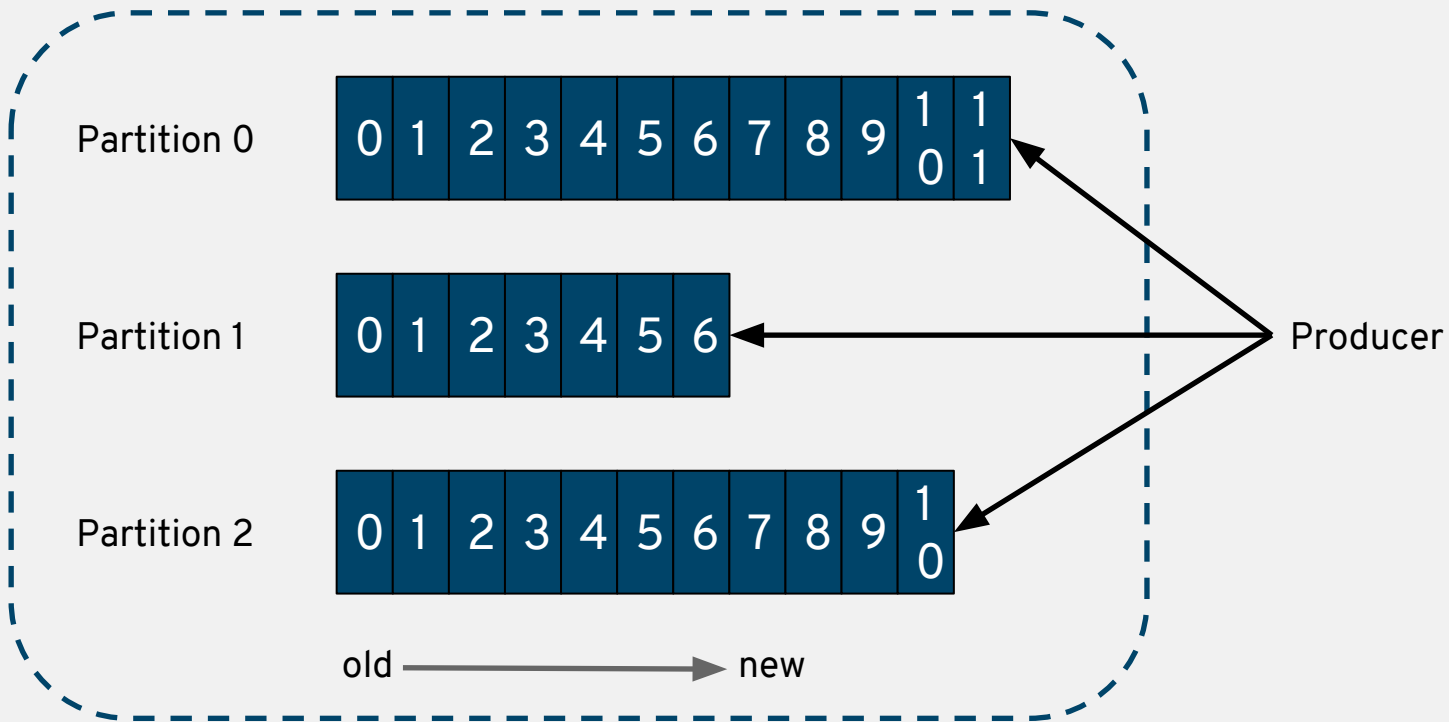
A streaming data platform?

A distributed, horizontally-scalable, fault-tolerant, commit log?

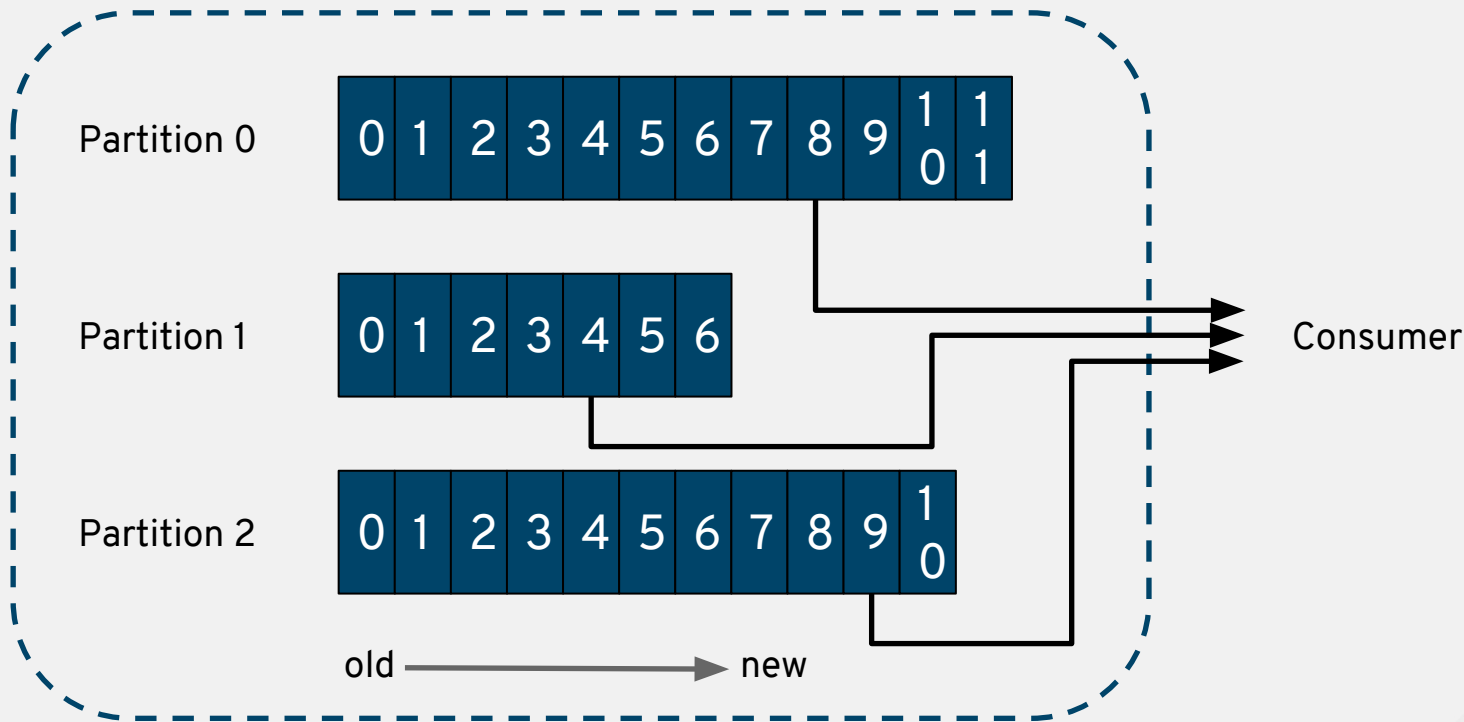
APACHE KAFKA CONCEPTS

- Messages are sent to and received from a topic
 - Topics are split into one or more partitions (aka shards)
 - All actual work is done on partition level, topic is just a virtual object
- Each message is written only into a one selected partition
 - Partitioning is usually done based on the message key
 - Message ordering within the partition is fixed
- Retention
 - Based on size / message age
 - Compacted based on message key
- Replication
 - Each partition can exist in one or more copies to achieve high availability


TOPIC & PARTITIONS



TOPICS & PARTITIONS

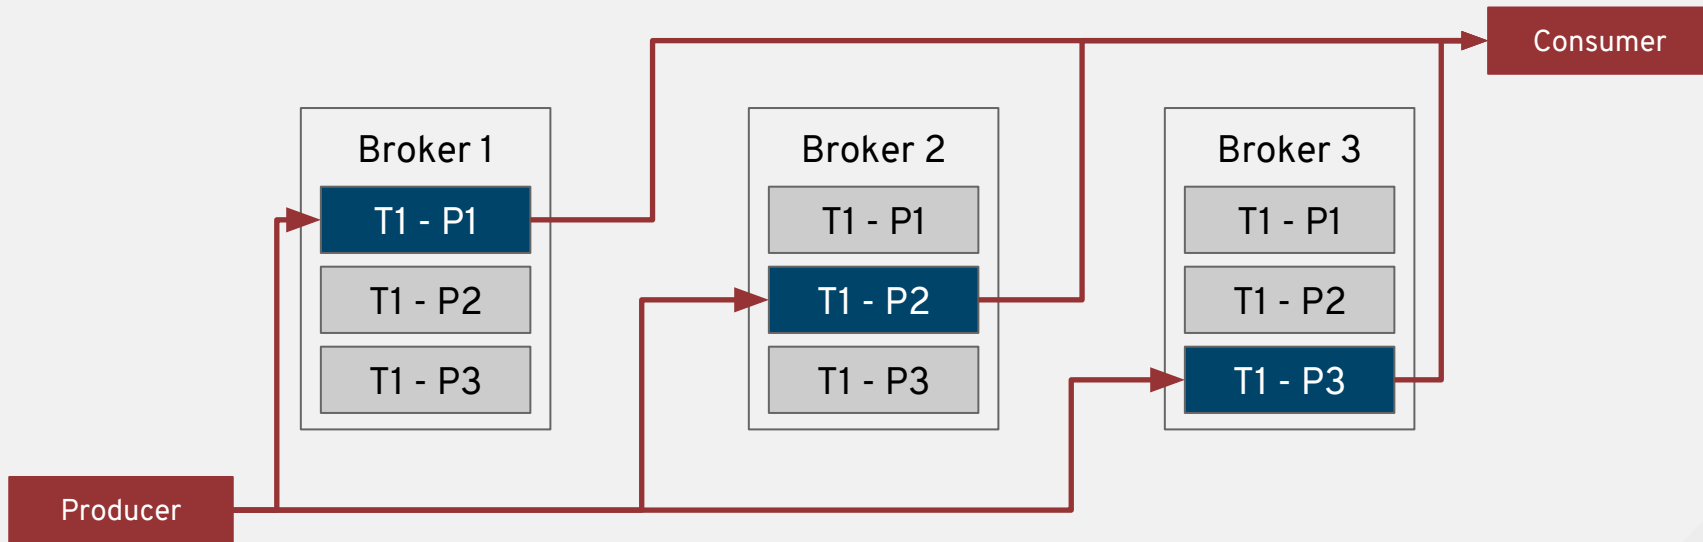


HIGH AVAILABILITY



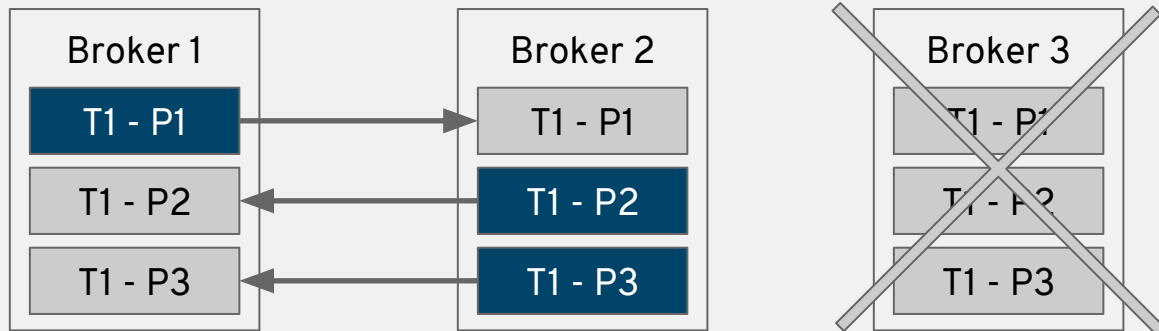
Leaders and followers spread across the cluster.

HIGH AVAILABILITY



Clients always connect only to the leaders.

HIGH AVAILABILITY

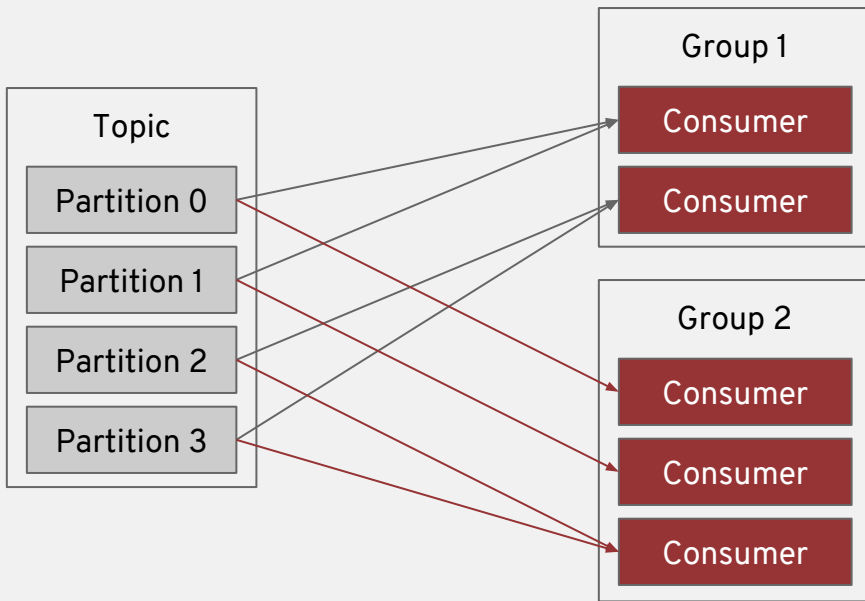


If a broker with leader partition goes down, a new leader partition is elected on different node

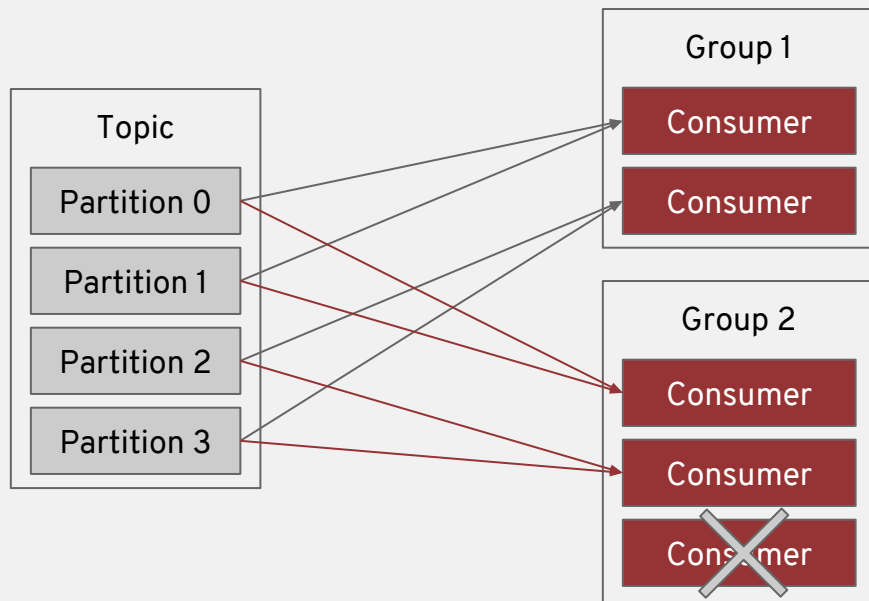
CONSUMER GROUPS

- Consumer Group
 - Grouping multiple consumers
 - Each consumer reads from a “unique” subset of partition → max consumers = num partitions
 - They are “competing” consumers on the topic, each message delivered to one consumer
 - Messages with same “key” delivered to same consumer
- More consumer groups
 - Allows publish/subscribe
 - Same messages delivered to different consumers in different consumer groups


CONSUMER GROUPS



CONSUMER GROUPS



CONSUMER GROUPS



APACHE KAFKA FOR IOT

Pros & Cons

Pros

- “Big” buffer for handling back pressure
- High throughput
- High availability
- Long term storage
- Reprocessing devices data

Cons

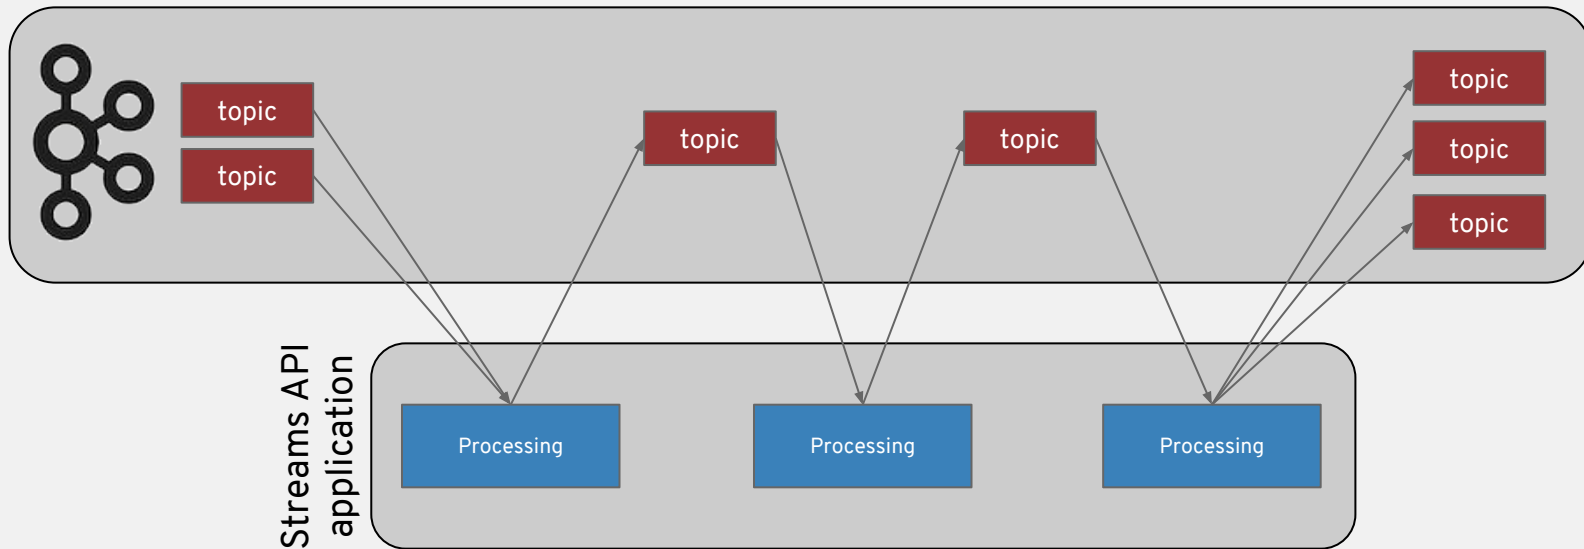
- Cannot handle a huge number of device connections
- No standard IoT devices oriented protocol
- Network needs to be stable

KAFKA STREAMS


KAFKA STREAMS API

- Stream processing framework
- Streams are Kafka topics (as input and output)
- It's really just a Java library to include in your application
- Scaling the stream application horizontally
- Creates a topology of processing nodes (filter, map, join etc) acting on a stream
 - Low level processor API
 - High level DSL
 - Using “internal” topics (when re-partitioning is needed or for “stateful” transformations)

KAFKA STREAMS API




PROCESSORS TOPOLOGY



STATE MANAGEMENT

Stateless vs Stateful processing

- Stateless (filter, map) vs Stateful (join, aggregate)
- State handled locally
 - Faster
 - Better isolation
 - Flexible
- It's based on RocksDB and backed by a Kafka topic as well
 - For rebuilding state on a different Streams node



EVENTS TIMESTAMP


Always because time matters

Time


- Event time, which is the time at which events actually occurred
 - As part of the message payload
 - Timestamp on the message (set by the producer)
- Ingestion time, which is the time when the data enters the processing pipeline
 - The broker sets the timestamp on the message
- Processing time, which is the time at which events are observed in the system
 - It's also called “wall-clock” time

EVENTS TIMESTAMP


Event time



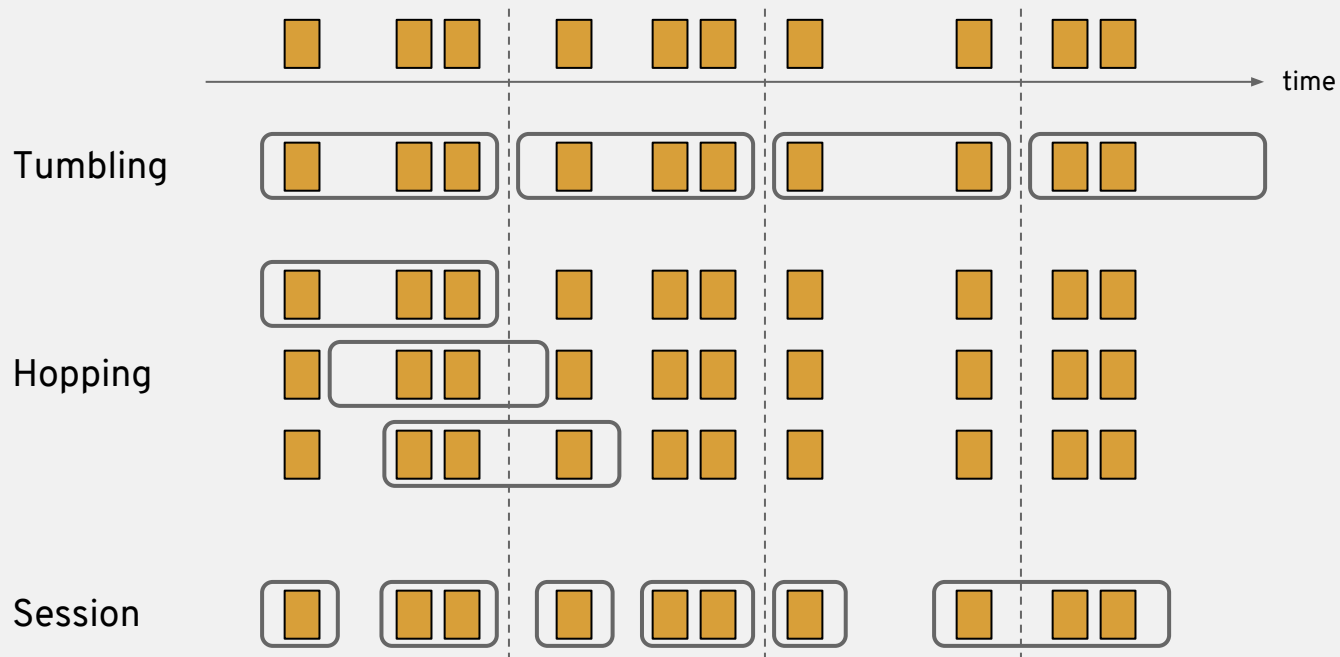
Ingestion time



Processing time

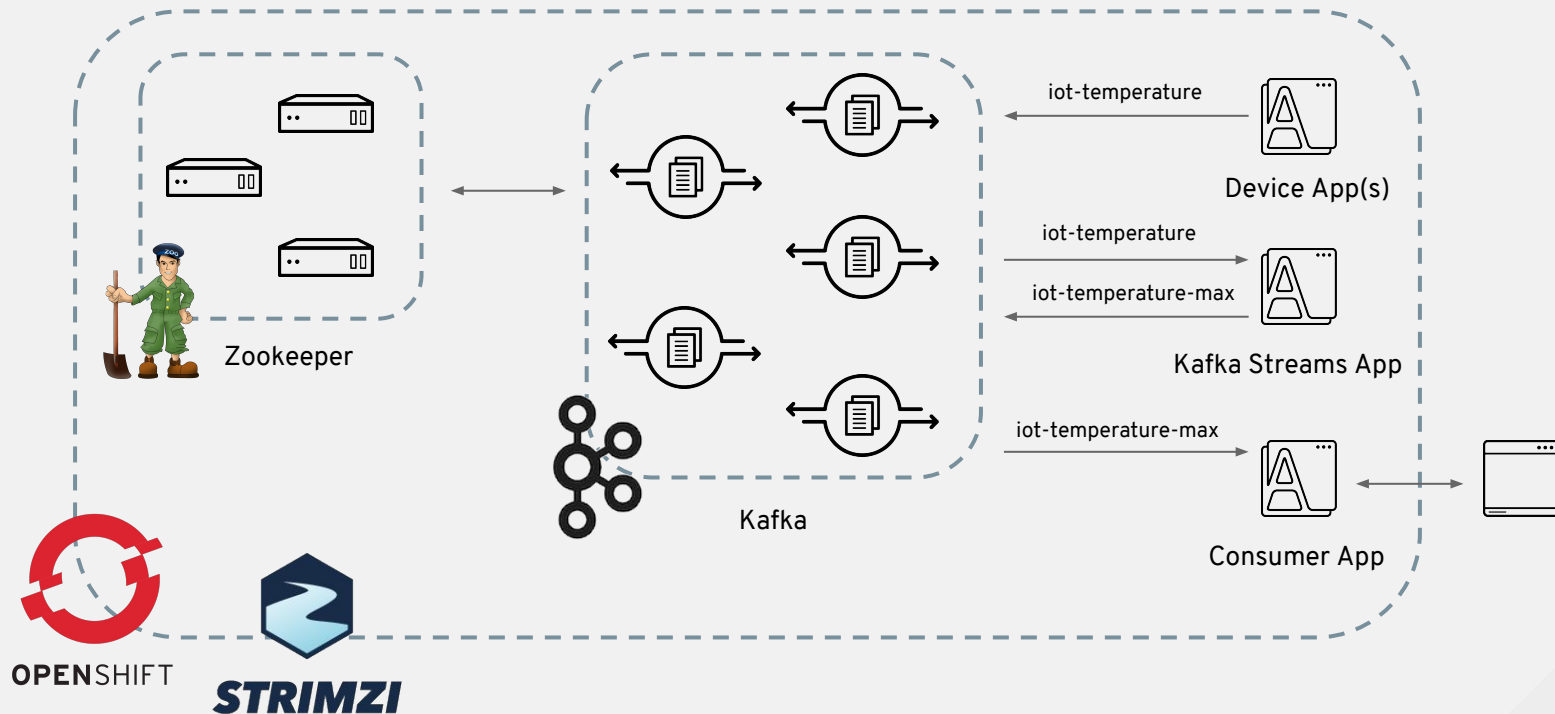


WINDOWING



DEMO

DEMO ARCHITECTURE





THANK YOU



plus.google.com/+RedHat



facebook.com/redhatinc



linkedin.com/company/red-hat



twitter.com/RedHat



youtube.com/user/RedHatVideos