



Masterclass

....

Amazon S3 – beyond simple storage

Ian Massingham – Technical Evangelist

 @IanMmmm

Masterclass

A technical deep dive beyond the basics

Help educate you on how to get the best from AWS technologies

Show you how things work and how to get things done

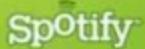
Broaden your knowledge in ~45 minutes

Amazon S3

It's more than just a 'simple' storage platform

A sophisticated 21st century distributed system

A bedrock architectural component for many applications



Listen to millions of songs for free.

Play, discover and share with your friends.

[Download Spotify](#)

By clicking download, you agree to Spotify's [terms & conditions](#) and [privacy policy](#).



Case Study



Listen to millions of songs for free.

Play, discover and share with your friends.

[Download Spotify](#)

By clicking download, you agree to Spotify's [terms & conditions](#) and [privacy policy](#).



Case Study

“Spotify needed a storage solution that could scale very quickly without incurring long lead times for upgrades. This led us to cloud storage, and in that market, Amazon Simple Storage Service (Amazon S3) is the most mature large-scale product.

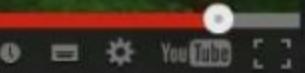
Amazon S3 gives us confidence in our ability to expand storage quickly while also providing high data durability.”

Emil Fredriksson, Operations Director

Find out more at : aws.amazon.com/solutions/case-studies/spotify/

MINECRAFT

► 0:57 / 1:00



Minecraft is a game about breaking and placing blocks. At first, people built structures to protect against local monsters, but as the game grew players worked together to create wonderful, imaginative things.

About adventuring with friends or watching the sun rise over a blocky ocean. It's pretty. Brave new things in The Nether, which is more scary than pretty. You can also visit a land of

more like your cup of tea.

Case Study

You can play Minecraft on PC/Mac. Minecraft: Pocket Edition is available for iOS and Android, and Minecraft: Xbox 360 Edition is available for the Microsoft XBLA Marketplace. So far 14,037,668 people bought the PC/Mac version of the game. Plus,

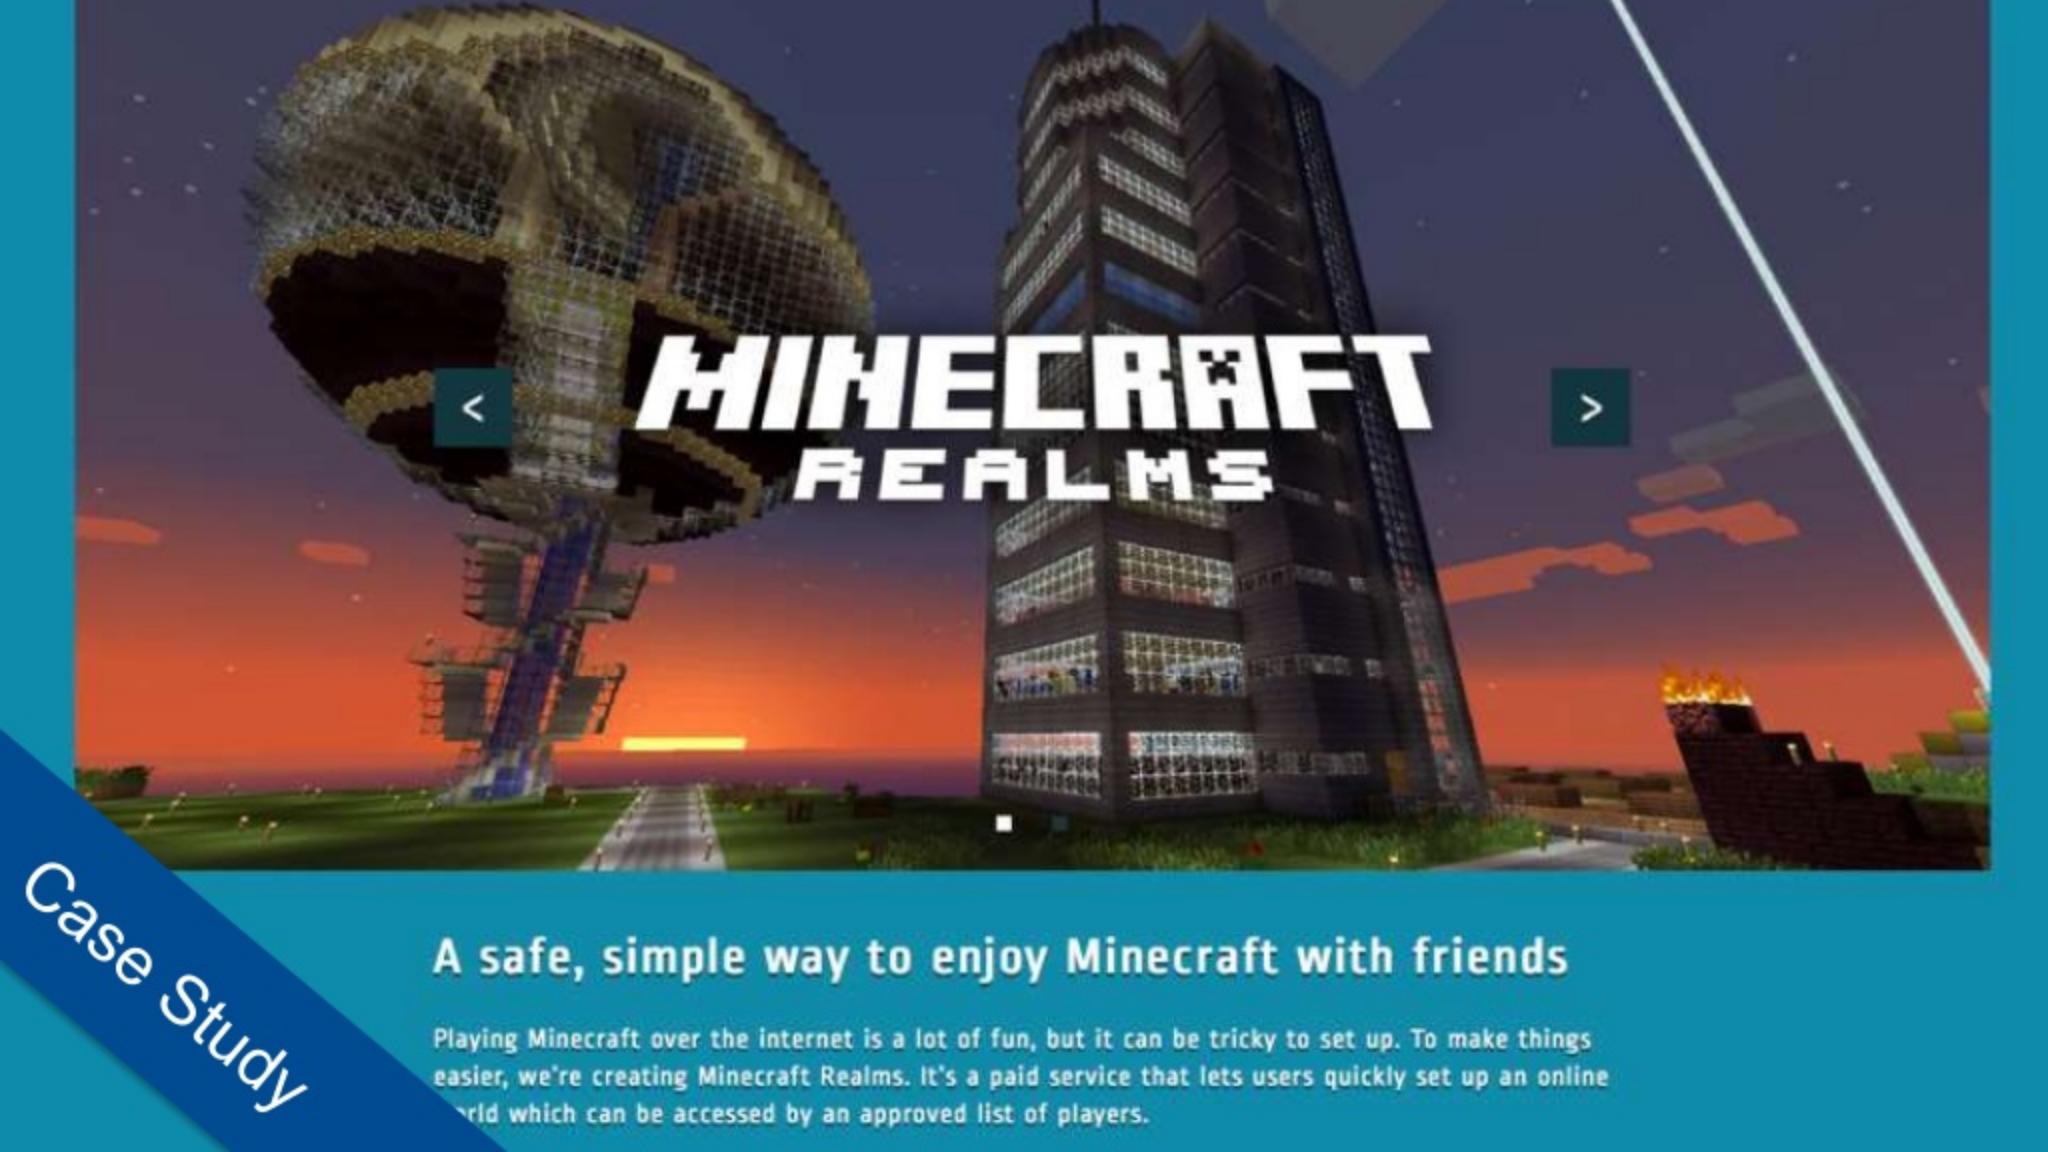


Get Minecraft!

or play the demo

Already bought the game?

Download it here



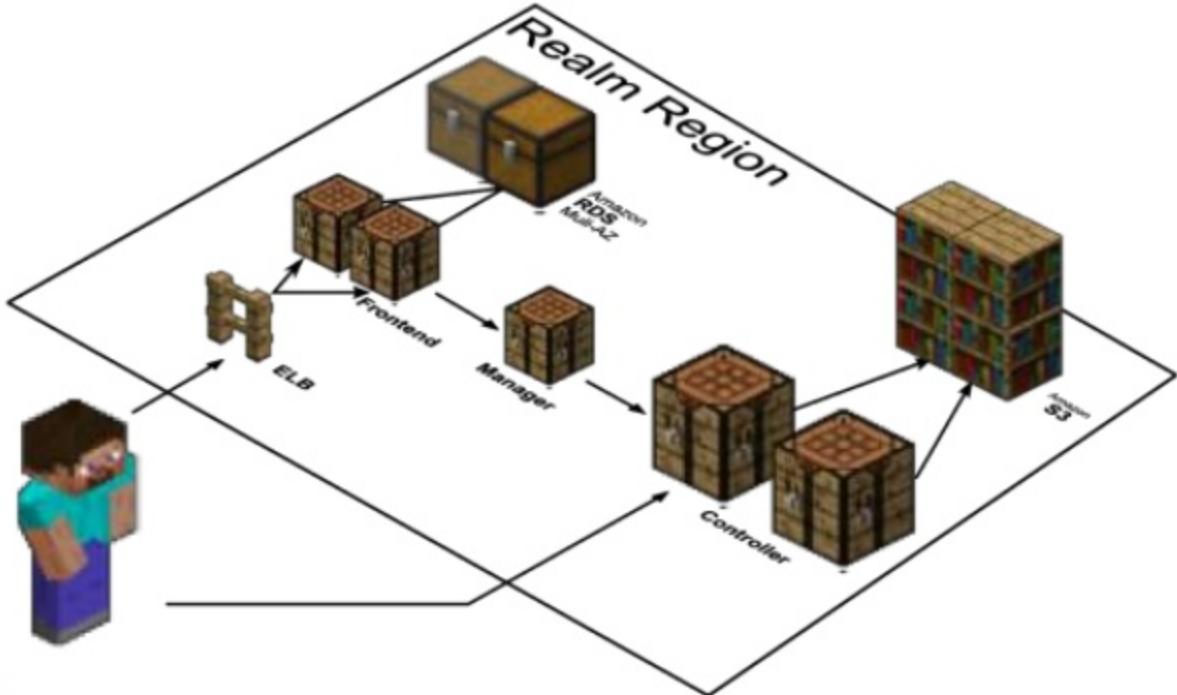
MINECRAFT REALMS

Case Study

A safe, simple way to enjoy Minecraft with friends

Playing Minecraft over the internet is a lot of fun, but it can be tricky to set up. To make things easier, we're creating Minecraft Realms. It's a paid service that lets users quickly set up an online world which can be accessed by an approved list of players.

Case Study



Minecraft Realms and AWS S3

Minecraft worlds and game state are stored in Amazon S3. The system takes advantage of S3 versioning, giving the owner / administrator of a realm the ability to roll back to a previous state. The team implemented efficient transfers to S3 by using S3's multipart upload capability.

Find out more at : aws.typepad.com/aws/2014/01/hosting-minecraft-realms-on-aws.html

You put it in S3

AWS stores with 99.99999999% durability

Highly scalable web
access to objects



You put it in S3

AWS stores with 99.99999999% durability



Multiple redundant
copies in a region

but it's more than just a
simple storage service

Objects in S3

Trillions of Objects
(000,000,000,000s)

Servicing over 2 million
requests per Second

Highly scalable data storage

A web store, not a file system

Access via APIs

Fast

Economical

Highly available & durable

What is S3?

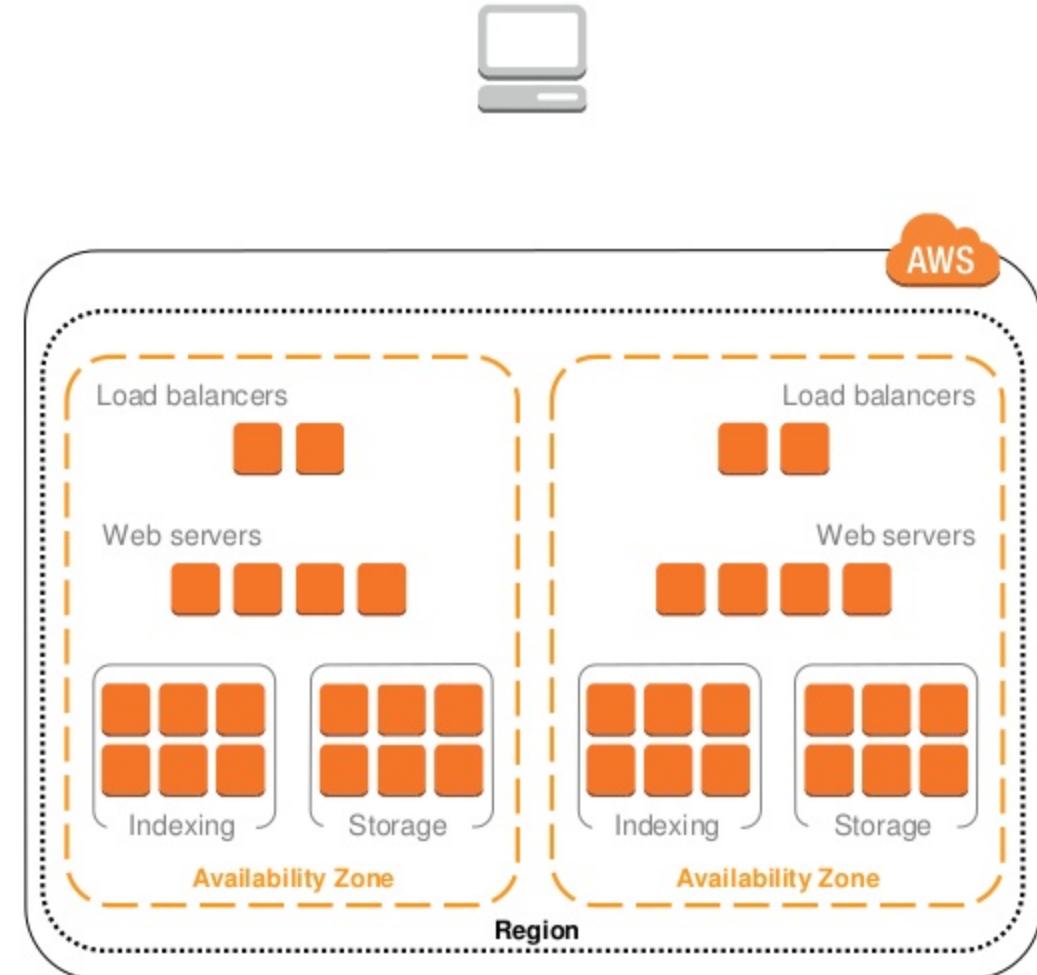
A web store, not a file system



Write once,
read many
(WORM)

Eventually
consistent

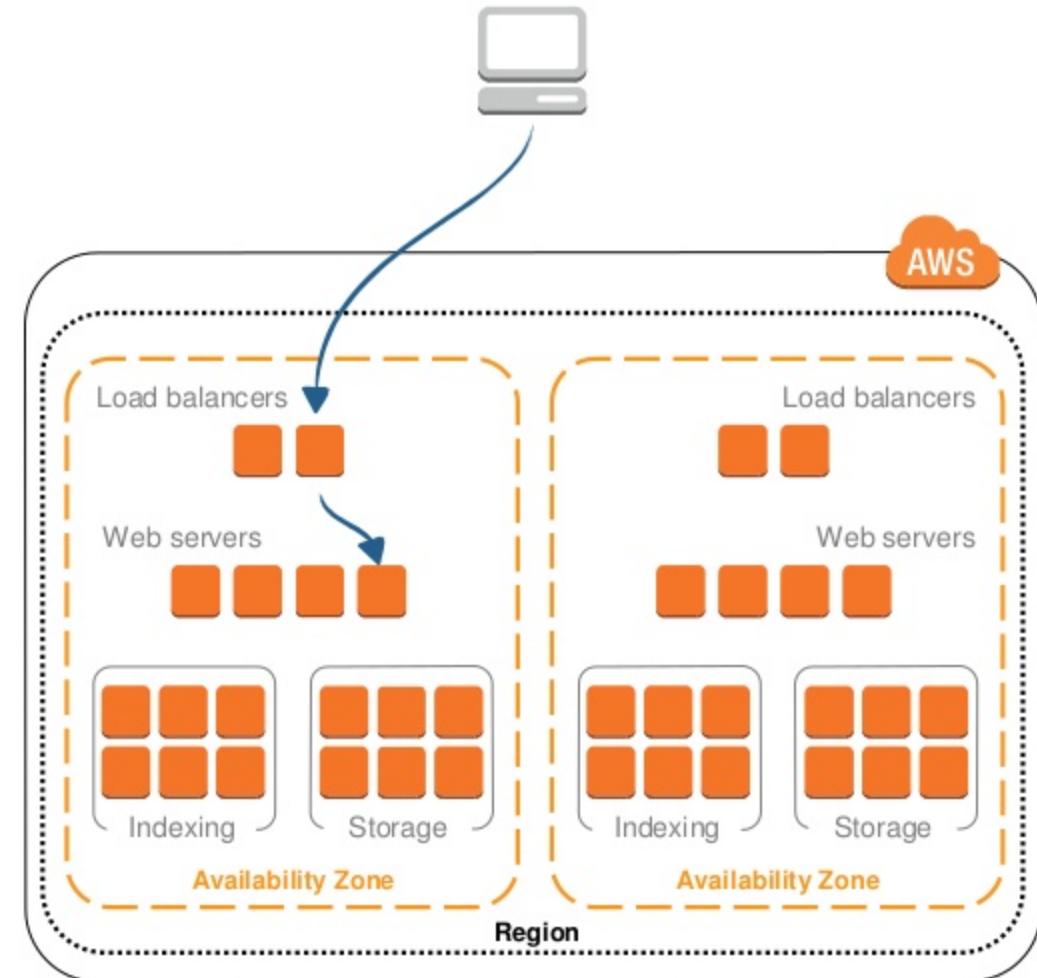
A web store, not a file system



A web store, not a file system



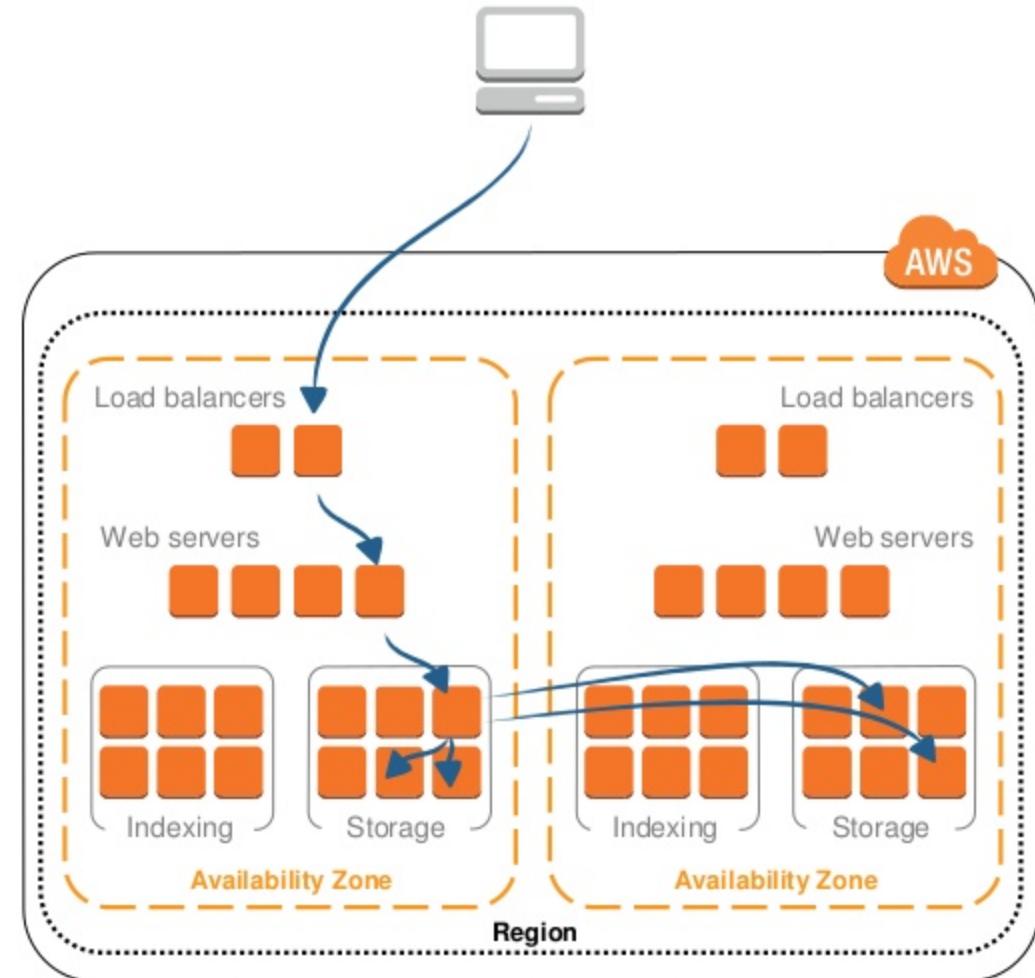
Write once,
read many
(WORM)



A web store, not a file system



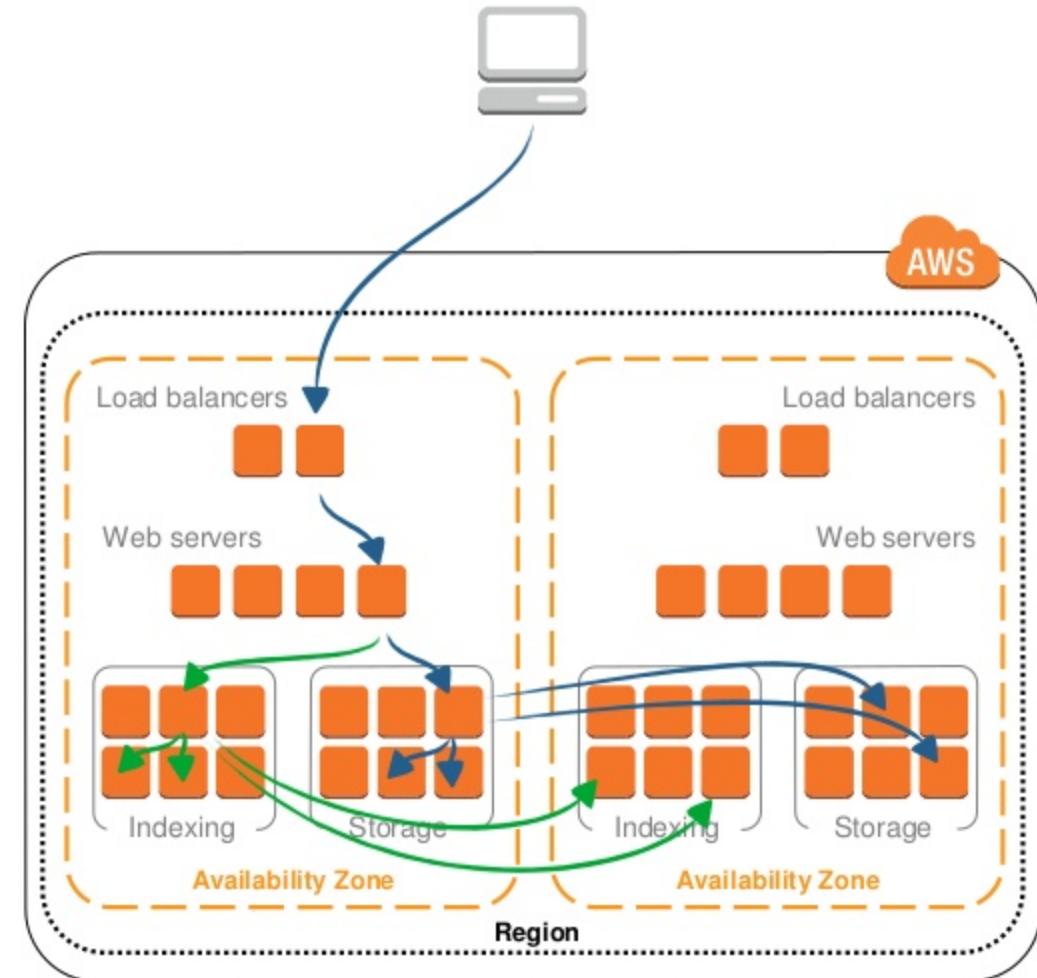
Write once,
read many
(WORM)



A web store, not a file system



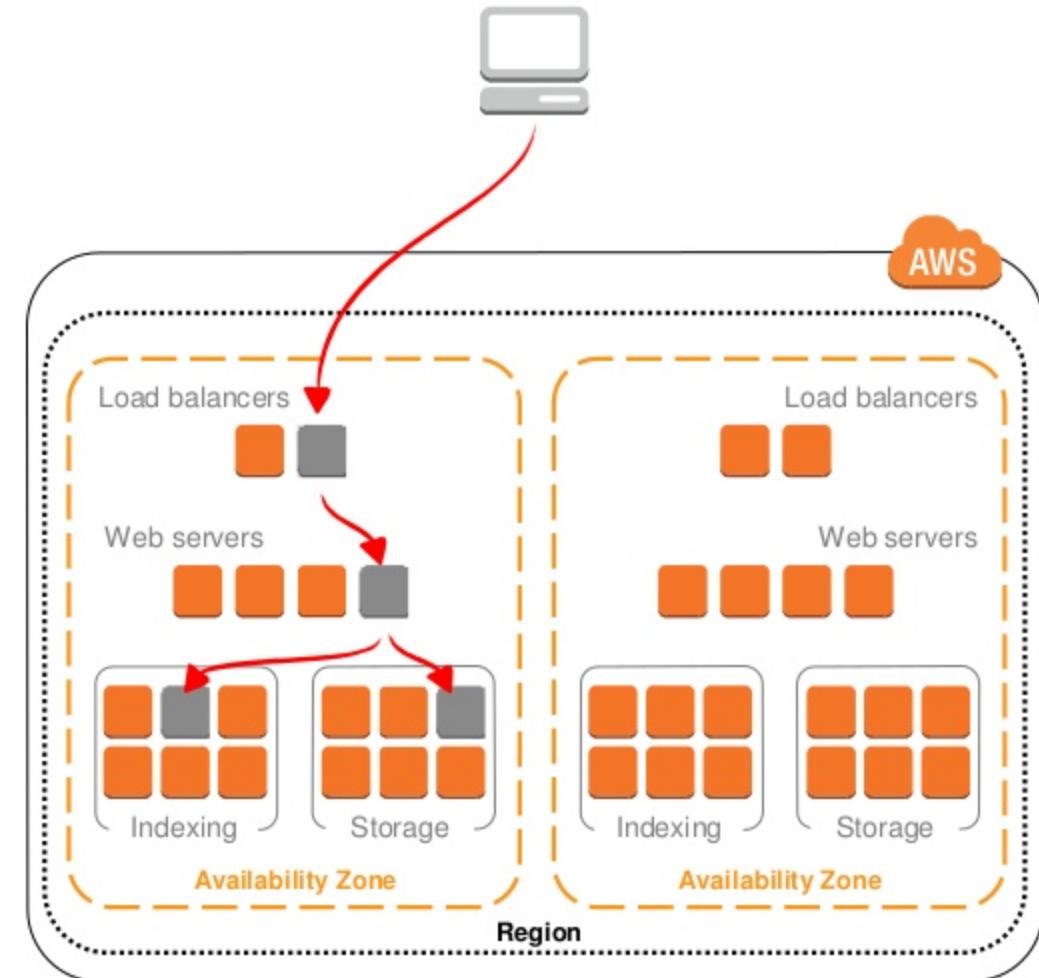
Write once,
read many
(WORM)



A web store, not a file system



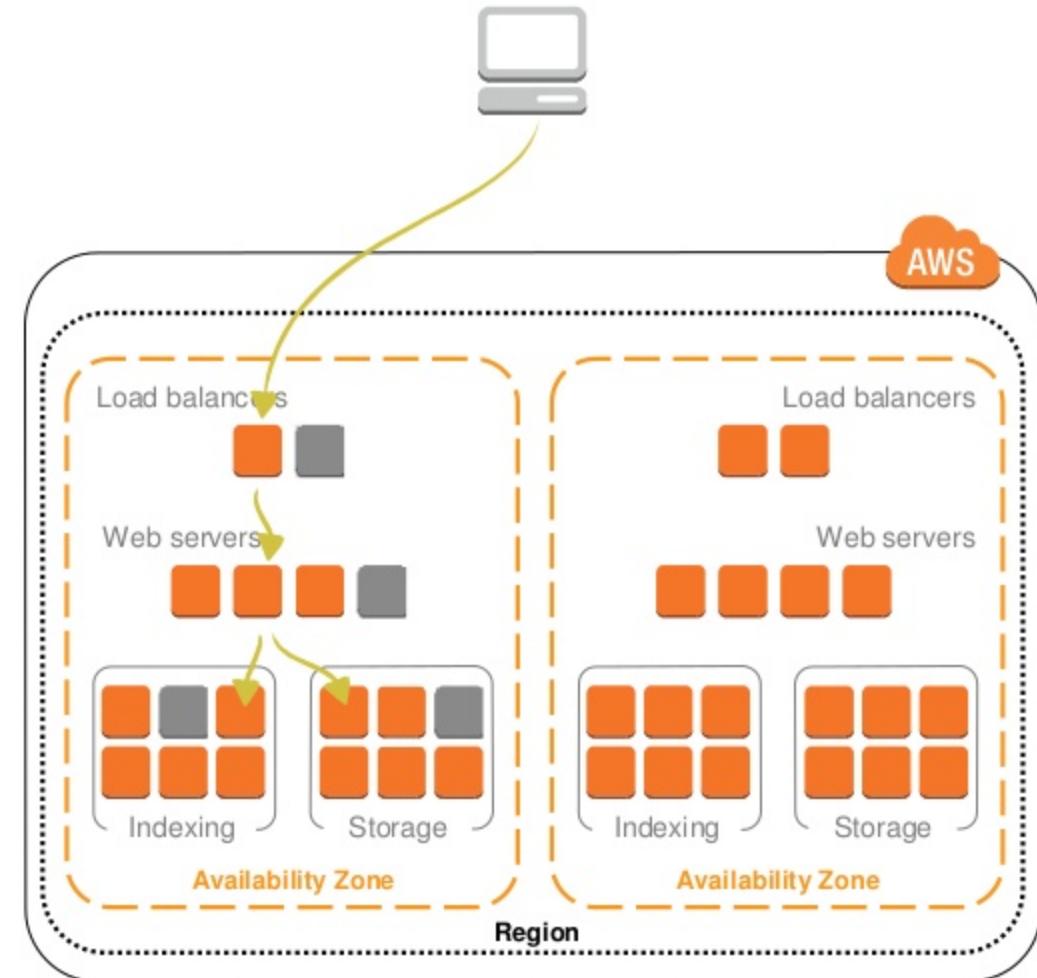
Write once,
read many
(WORM)



A web store, not a file system



Write once,
read many
(WORM)

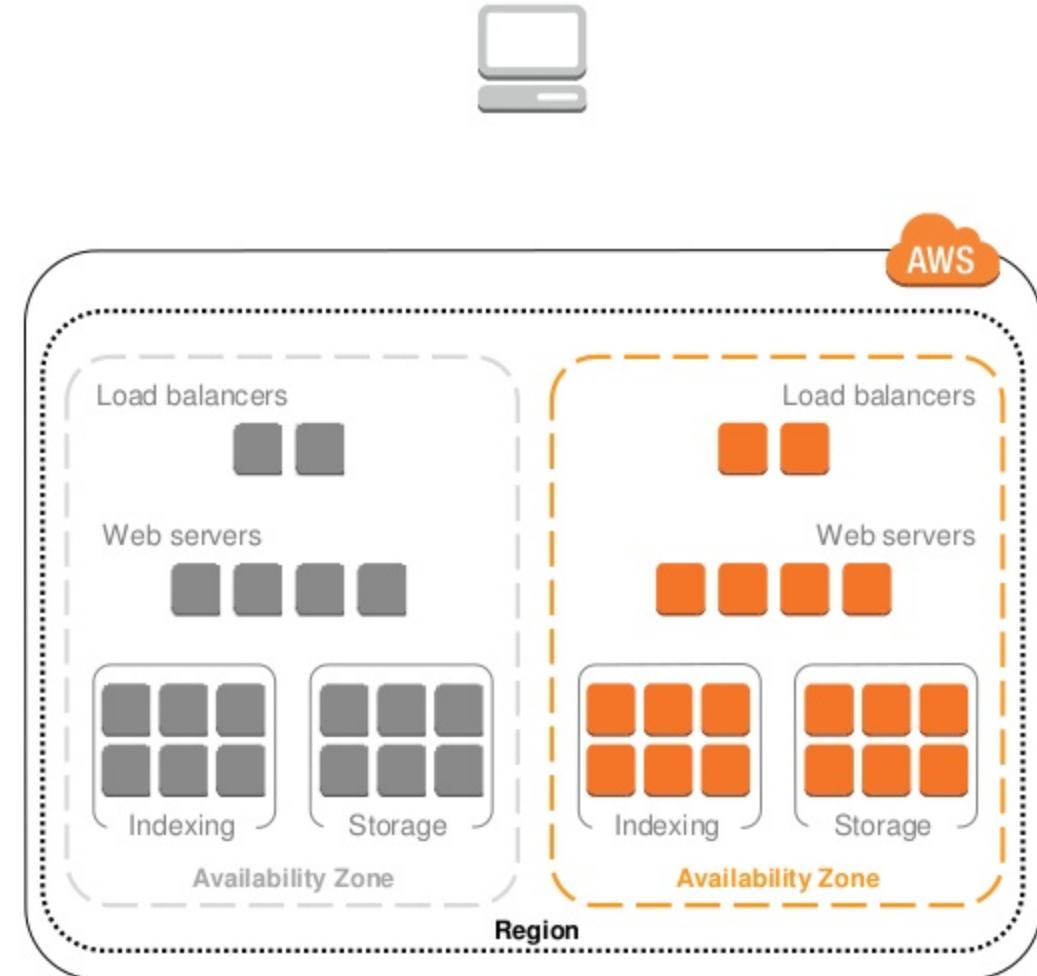


A web store, not a file system



Eventually
consistent

Write once,
read many
(WORM)

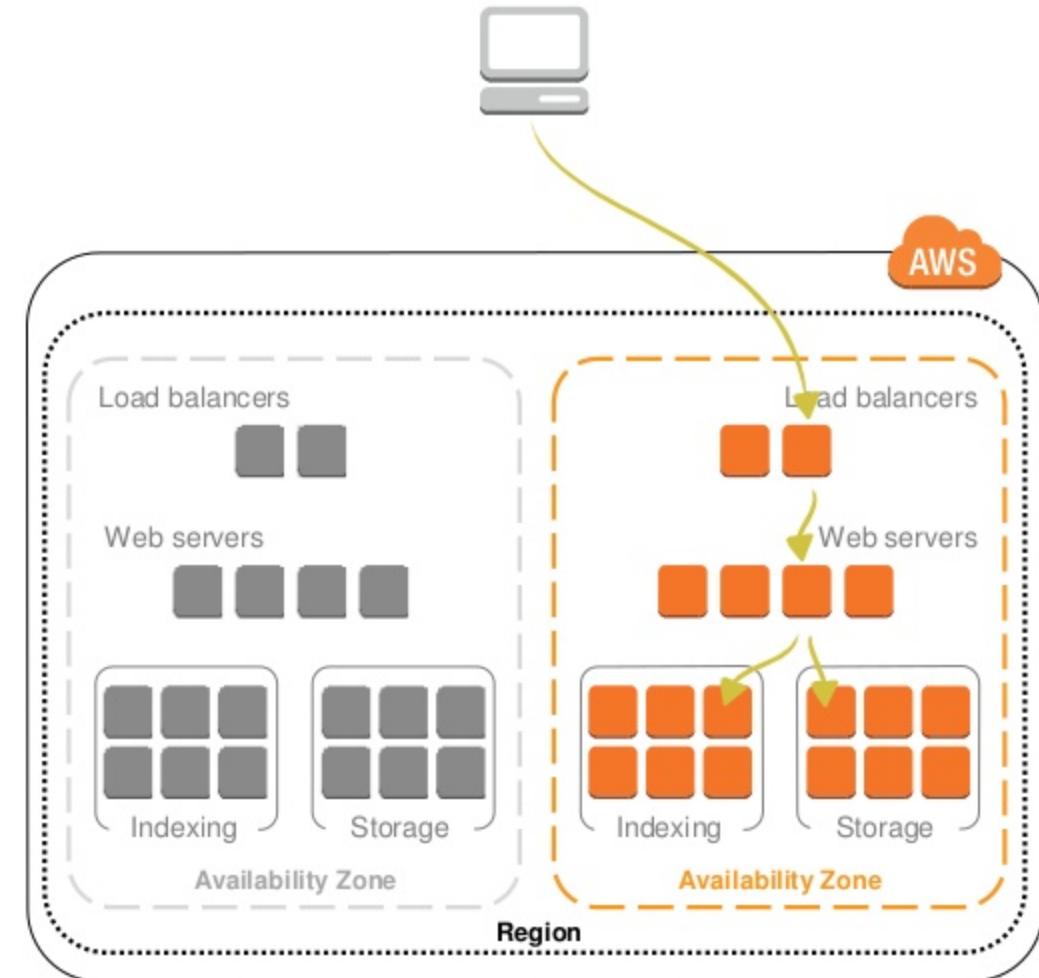


A web store, not a file system



Write once,
read many
(WORM)

Eventually
consistent



A web store, not a file system



Eventually
consistent

A web store, not a file system



New objects

Synchronously stores your data across multiple facilities before returning SUCCESS

*Read-after-write consistency**

Updates

Write then read: could report key does not exist

Write then list: might not include key in list

Overwrite then read: old data could be returned

Deletes

Delete then read: could still get old data

Delete then list: deleted key could be included in list

A regional service

Your data never leaves a region unless you move it

Storage classes

Controlling the way S3 holds your data

Amazon S3 storage classes

Standard

Designed to provide
99.99999999% durability
and 99.99% availability of
objects over a given year

Designed to sustain the
concurrent loss of data in two
facilities

Amazon S3 storage classes

Standard

Designed to provide 99.99999999% durability and 99.99% availability of objects over a given year

Designed to sustain the concurrent loss of data in two facilities

Reduced Redundancy Storage

Reduces costs by storing data at lower levels of redundancy than the Standard storage

Designed to provide 99.99% durability and 99.99% availability of objects over a given year

Amazon S3 storage classes

Standard

Designed to provide 99.99999999% durability and 99.99% availability of objects over a given year

Designed to sustain the concurrent loss of data in two facilities

Reduced Redundancy Storage

Reduces costs by storing data at lower levels of redundancy than the Standard storage

Designed to provide 99.99% durability and 99.99% availability of objects over a given year

Glacier

Suitable for archiving data, where data access is infrequent and a retrieval time of several hours is acceptable

Uses the very low-cost Amazon Glacier storage service, but managed through Amazon S3

Amazon S3 storage classes

Standard

Objects you want to have high durability

e.g. master copy of movie media

Reduced Redundancy

Objects you can afford to lose or can recreate

e.g. different encodings of movie media

Glacier

Objects you want to keep in archive for a long time

e.g. digital archive of old movies & broadcasts

Namespaces

Object naming, buckets and keys

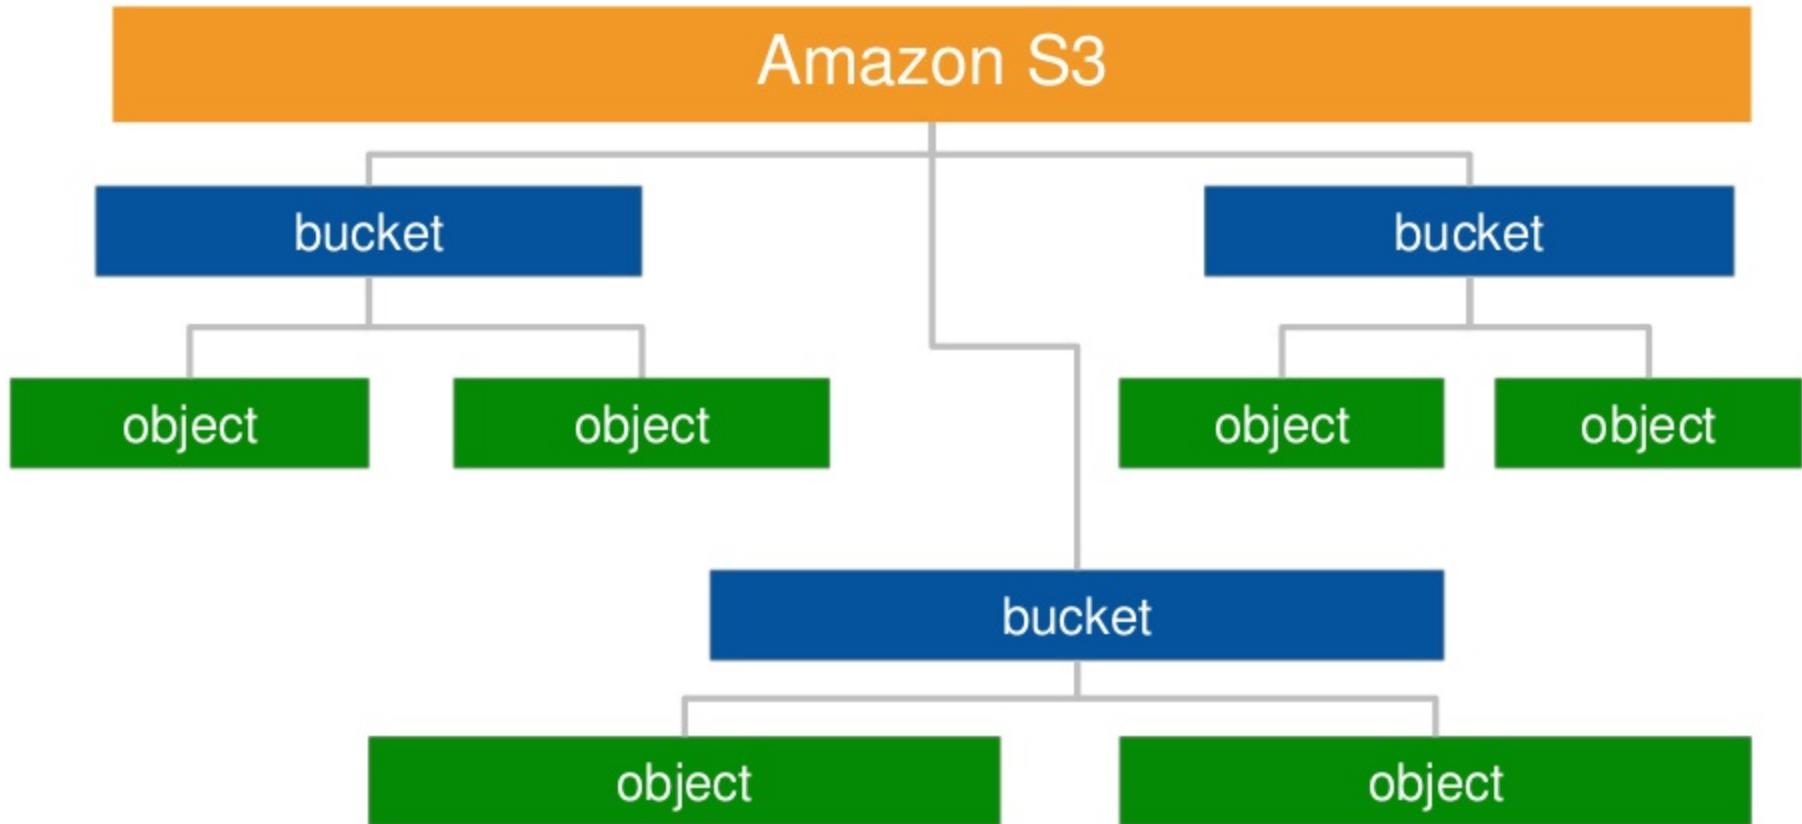
Amazon S3 namespace

Globally unique

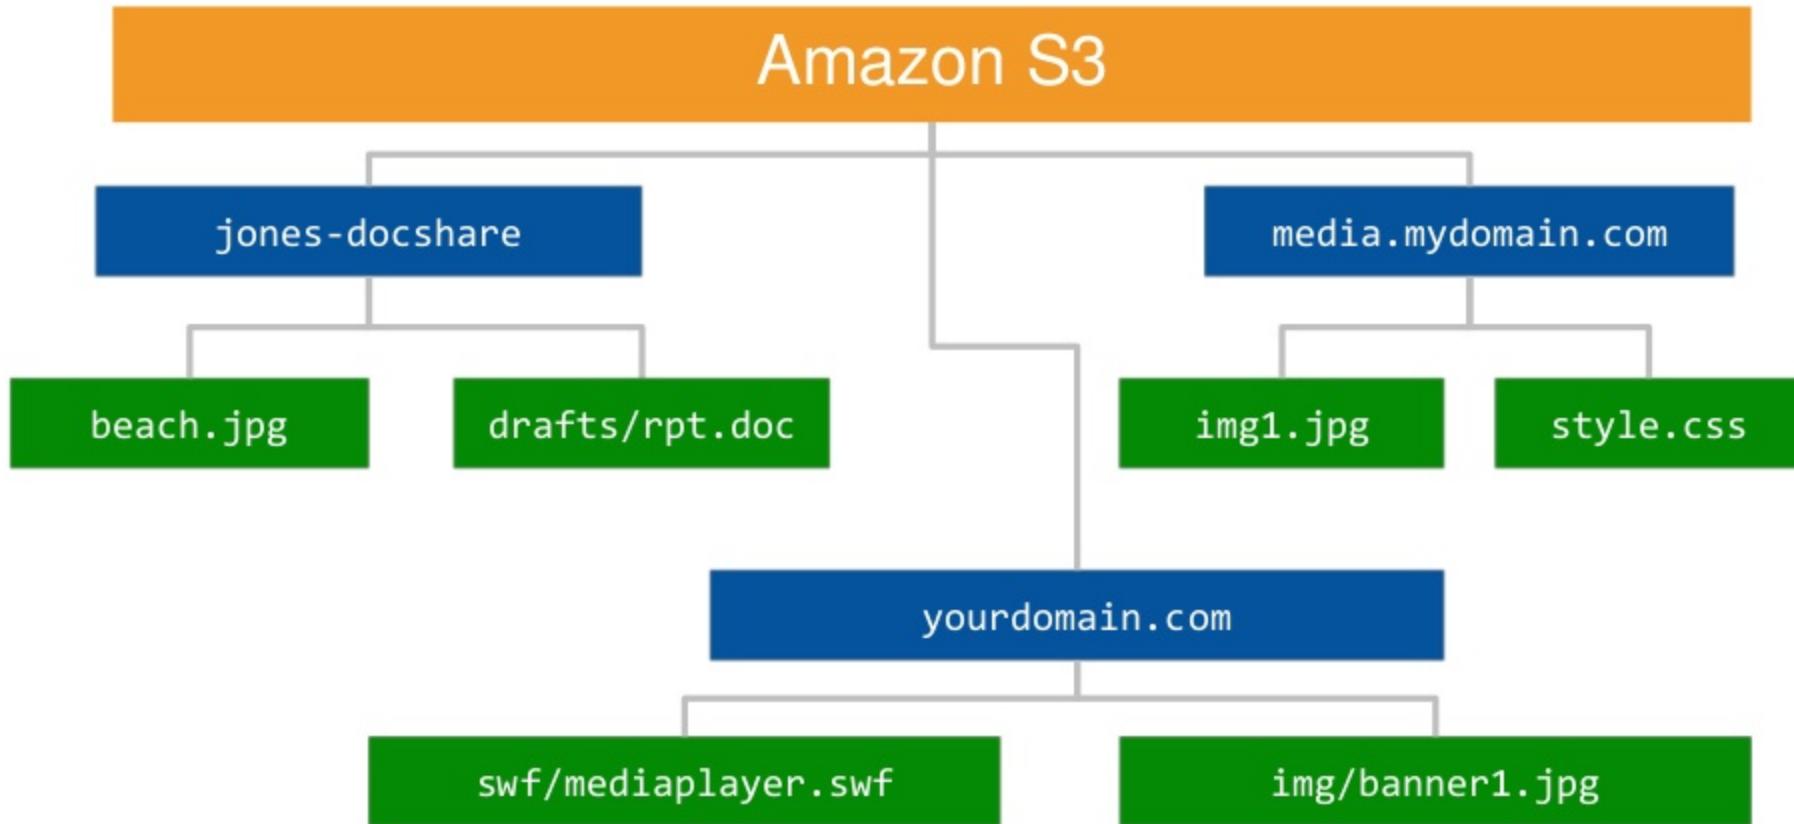


bucket name + object name (key)

Amazon S3 namespace



Amazon S3 namespace



Amazon S3 namespace

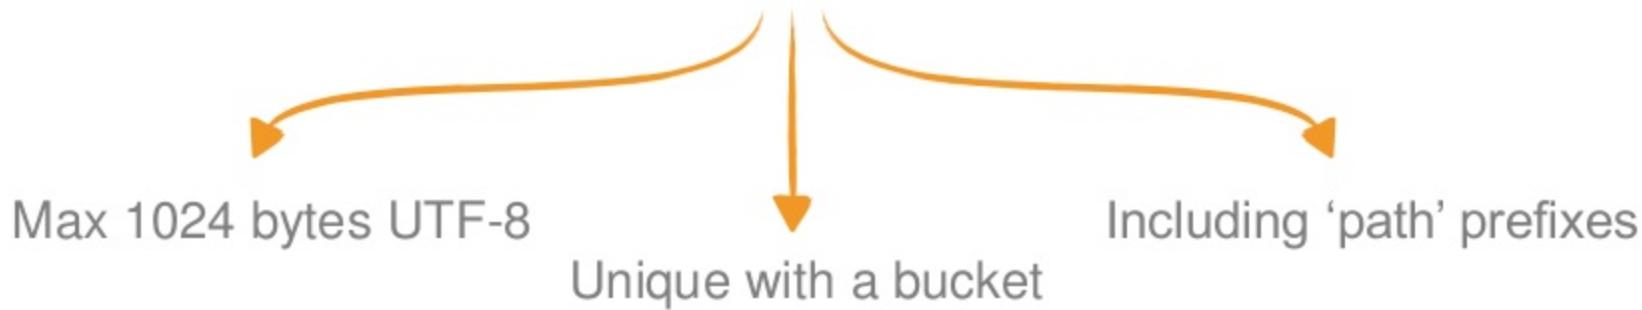
Object key



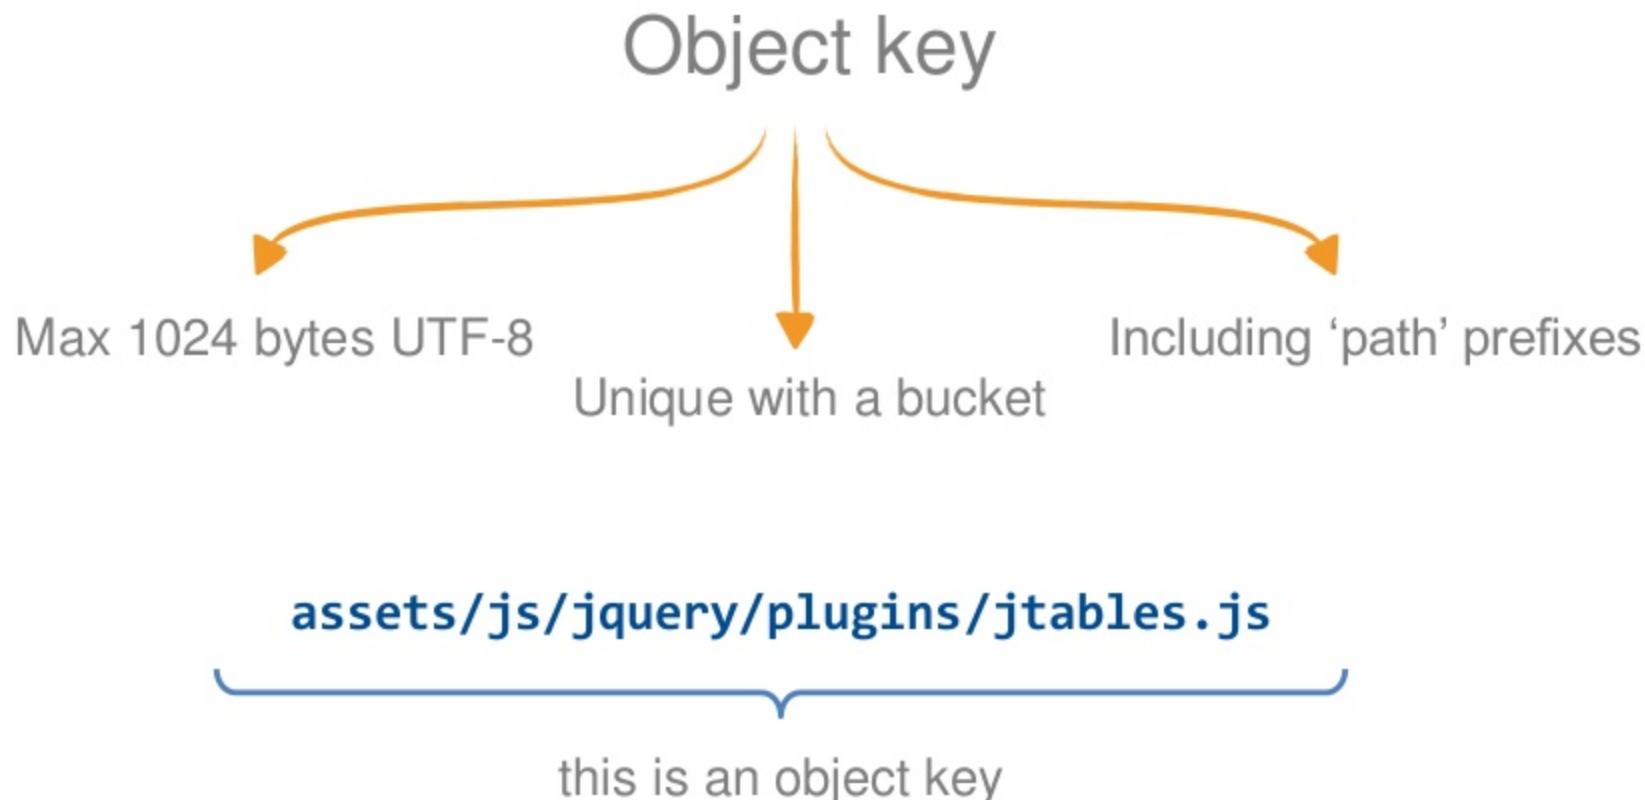
Unique with a bucket

Amazon S3 namespace

Object key



Amazon S3 namespace



Throughput optimisation

S3 automatically partitions based upon key prefix:

Bucket:

mynewgame

Object keys:

2134857/gamedata/start.png
2134857/gamedata/resource.rsrc
2134857/gamedata/results.txt
2134858/gamedata/start.png
2134858/gamedata/resource.rsrc
2134858/gamedata/results.txt
2134859/gamedata/start.png
2134859/gamedata/resource.rsrc
2134859/gamedata/results.txt

Throughput optimisation

S3 automatically partitions based upon key prefix:

Bucket:

mynewgame

Incrementing
game id

Object keys:

2134857/gamedata/start.png
2134857/gamedata/resource.rsrc
2134857/gamedata/results.txt
2134858/gamedata/start.png
2134858/gamedata/resource.rsrc
2134858/gamedata/results.txt
2134859/gamedata/start.png
2134859/gamedata/resource.rsrc
2134859/gamedata/results.txt



Throughput optimisation

S3 automatically partitions based upon key prefix:

Bucket:

mynewgame

Partition:

mynewgame/2

Object keys:



- 2134857/gamedata/start.png
- 2134857/gamedata/resource.rsrc
- 2134857/gamedata/results.txt
- 2134858/gamedata/start.png
- 2134858/gamedata/resource.rsrc
- 2134858/gamedata/results.txt
- 2134859/gamedata/start.png
- 2134859/gamedata/resource.rsrc
- 2134859/gamedata/results.txt

Throughput optimisation

S3 automatically partitions based upon key prefix:

Bucket:

mynewgame

Object keys:

7584312/gamedata/start.png

7584312/gamedata/resource.rsrc

7584312/gamedata/results.txt

8584312/gamedata/start.png

8584312/gamedata/resource.rsrc

8584312/gamedata/results.txt

9584312/gamedata/start.png

9584312/gamedata/resource.rsrc

9584312/gamedata/results.txt

Reversed
game ID



Throughput optimisation

S3 automatically partitions based upon key prefix:

Bucket:

mynewgame

Partitions:

mynewgame/7

mynewgame/8

mynewgame/9

Object keys:



7584312/gamedata/start.png
7584312/gamedata/resource.rsrc
7584312/gamedata/results.txt
8584312/gamedata/start.png
8584312/gamedata/resource.rsrc
8584312/gamedata/results.txt
9584312/gamedata/start.png
9584312/gamedata/resource.rsrc
9584312/gamedata/results.txt

Encryption

Securing data at rest

Server side encryption

*Automatic encryption
of data at rest*

↑

↓

Self managed
No need to manage a key store

Simple
Additional PUT
header

←

Strong
AES-256

←

Durable
S3 key storage

→

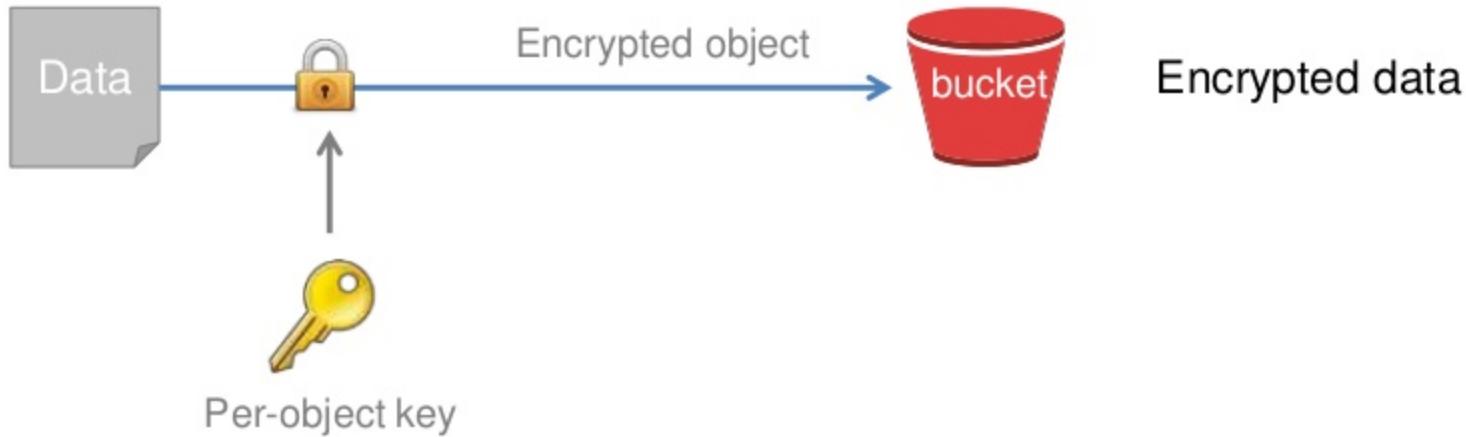
Secure
3-way
simultaneous
access

→

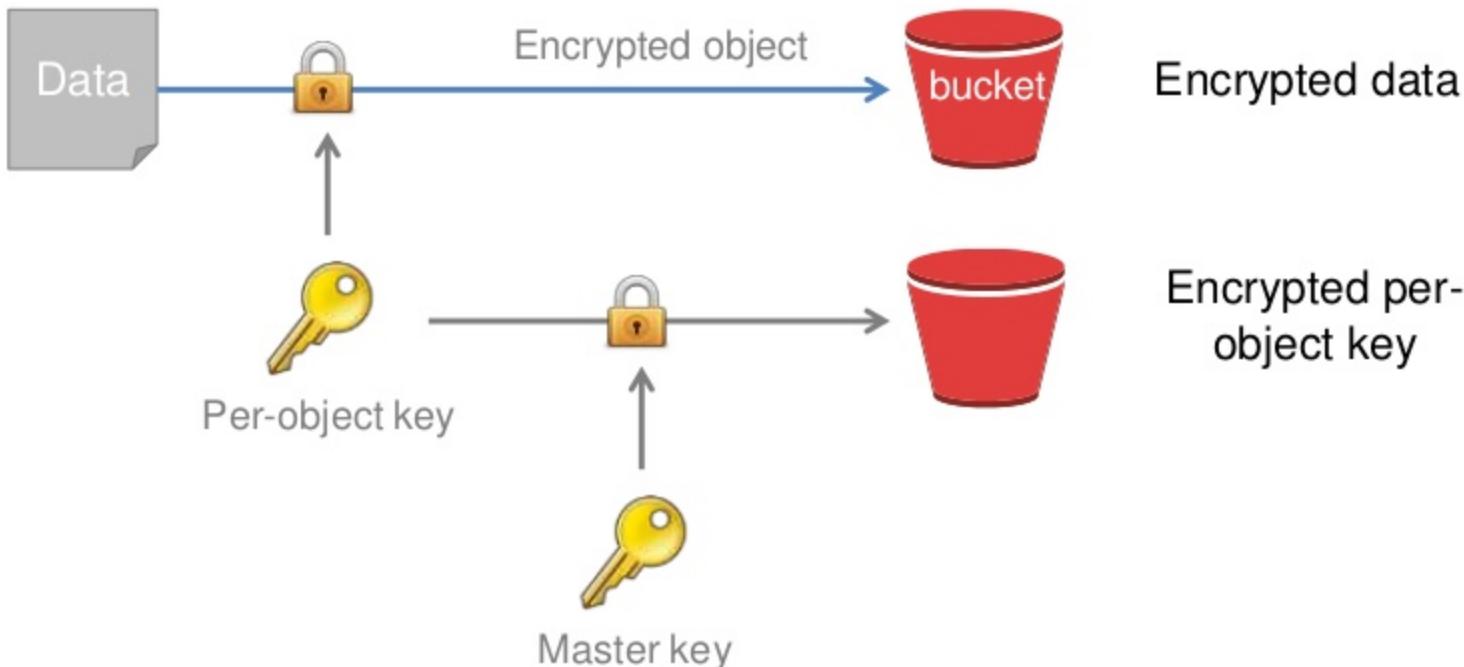
Server side encryption



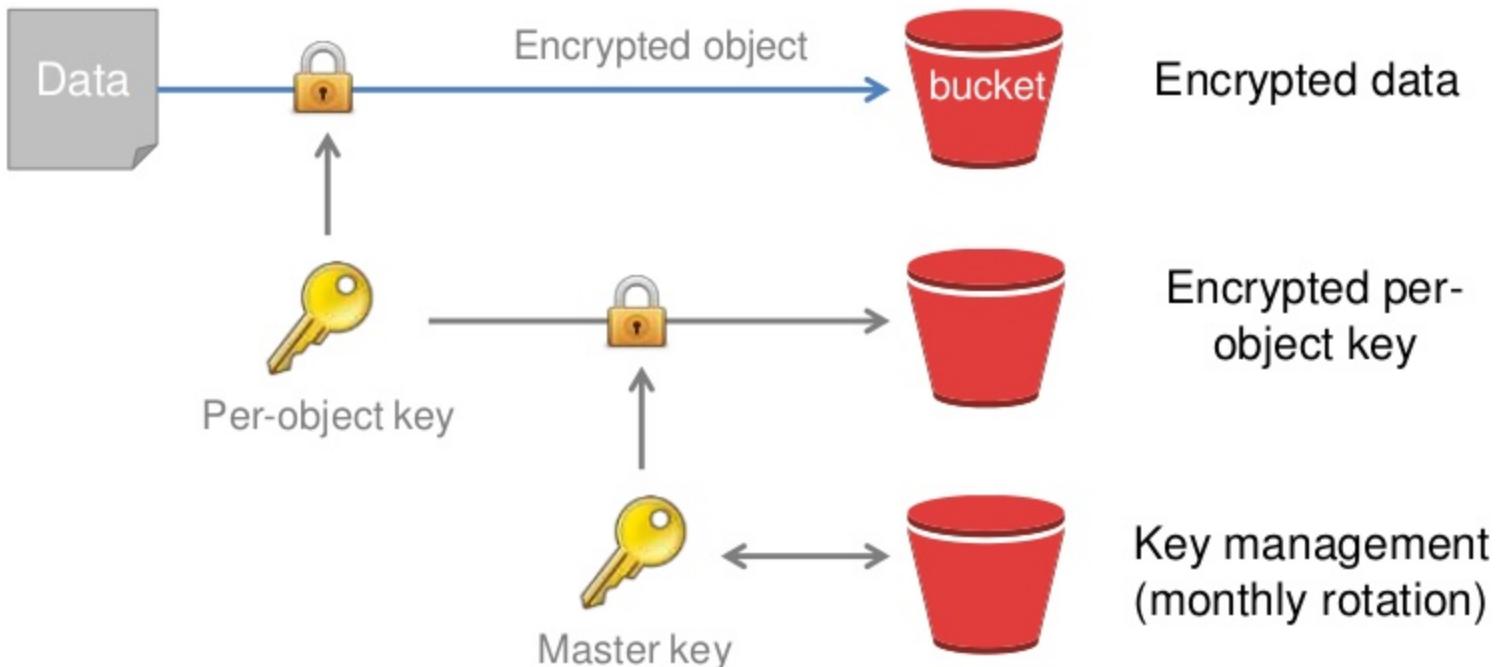
Server side encryption



Server side encryption



Server side encryption



Set Details

Cancel 

Upload to: Buckets / mystaticssite

Details: Set additional details for all of the objects you upload. You can choose between Standard Storage and [Reduced Redundancy Storage](#). You can also choose whether or not to [encrypt your files](#) on the server.

- Use Reduced Redundancy Storage
- Use Server Side Encryption



< Select Files Set Permissions > Start Upload Cancel

Access controls

You're in control of who does what

You decide what to share
Apply policies to buckets and objects

Secure by default



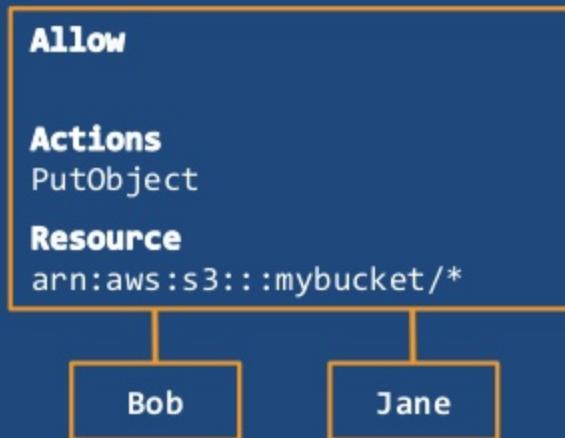
Policies, ACLs & IAM
Use S3 policies, ACLs or IAM to define rules

IAM

Fine grained

Administer as part of
role based access

Apply policies to S3 at
role, user & group level



IAM vs Bucket Policies

Fine grained

Administer as part of
role based access

Apply policies to S3 at
role, user & group level

Fine grained

Apply policies at the bucket
level in S3

Incorporate user restrictions
without using IAM

Allow

Actions

PutObject

Resource

arn:aws:s3:::mybucket/*

Allow

Bob, Jane

Actions

PutObject

Resource

arn:aws:s3:::mybucket/*

Bob

Jane

mybucket

IAM

vs

Bucket Policies

vs

ACLs

Fine grained

Administer as part of
role based access

Apply policies to S3 at
role, user & group level

Fine grained

Apply policies at the bucket
level in S3

Incorporate user restrictions
without using IAM

Coarse grained

Apply access control
rules at the bucket and/or
object level in S3

Allow

Actions

PutObject

Resource

arn:aws:s3:::mybucket/*

Bob

Jane

Allow

Bob, Jane

Actions

PutObject

Resource

arn:aws:s3:::mybucket/*

mybucket

Allow

Everyone, Bob, Jane

Actions

Read

mybucket

myobject

```
{"Statement": [ {  
    "Effect": "Allow",  
    "Principal": {"AWS": ["4649-6425", "5243-0045"]},  
    "Action": "*",  
    "Resource": "/mybucket/*",  
    "Condition": {  
        "IpAddress": {"AWS:SourceIp": "176.13.0.0/12"}  
    }]  
}]}
```

Bucket policy

```
{"Statement": [ { "Effect": "Allow", "Principal": {"AWS": ["4649-6425", "5243-0045"]}, "Action": "*", "Resource": "/mybucket/*", "Condition": { "IpAddress": {"AWS:SourceIp": "176.13.0.0/12"} } }]}
```

Accounts to allow



Bucket policy

```
{"Statement": [{

    "Effect": "Allow",

    "Principal": {"AWS": ["4649-6425", "5243-0045"]},

    "Action": "*",

    "Resource": "/mybucket/*"

    "Condition": {

        "IpAddress": {"AWS:SourceIp": "176.13.0.0/12"}


    }]]}
```

Resource

Bucket policy

```
{"Statement": [{

    "Effect": "Allow",

    "Principal": {"AWS": ["4649-6425", "5243-0045"]},

    "Action": "*",

    "Resource": "/mybucket/*"

    "Condition": {

        "IpAddress": {"AWS:SourceIp": "176.13.0.0/12"}


    }
}]]}
```

Source address to allow

Bucket policy

Transitions & Lifecycle Management

Automated management of objects

Lifecycle management

Object deletion

Permanently delete objects from S3



Lifecycle management

Object deletion

Permanently delete objects from S3



Lifecycle management



Object archiving

Move objects to Glacier and out of S3 storage

Glacier

Long term durable archive

Long term Glacier archive

Durable

Designed for 99.99999999%
durability of archives



Cost effective

Write-once, read-never. Cost effective
for long term storage. Pay for
accessing data

Expiry



accessible from S3



Expiry



accessible from S3



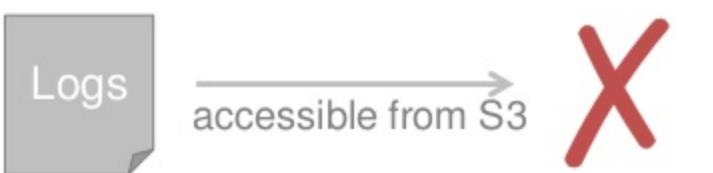
Objects expire
and are
deleted

time

Transition

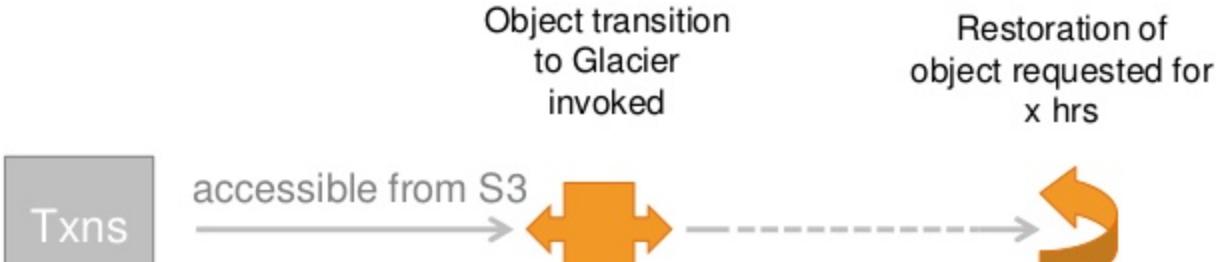


Expiry

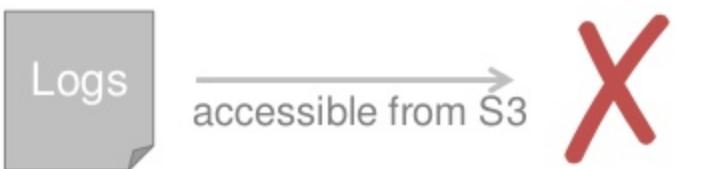


time

Transition



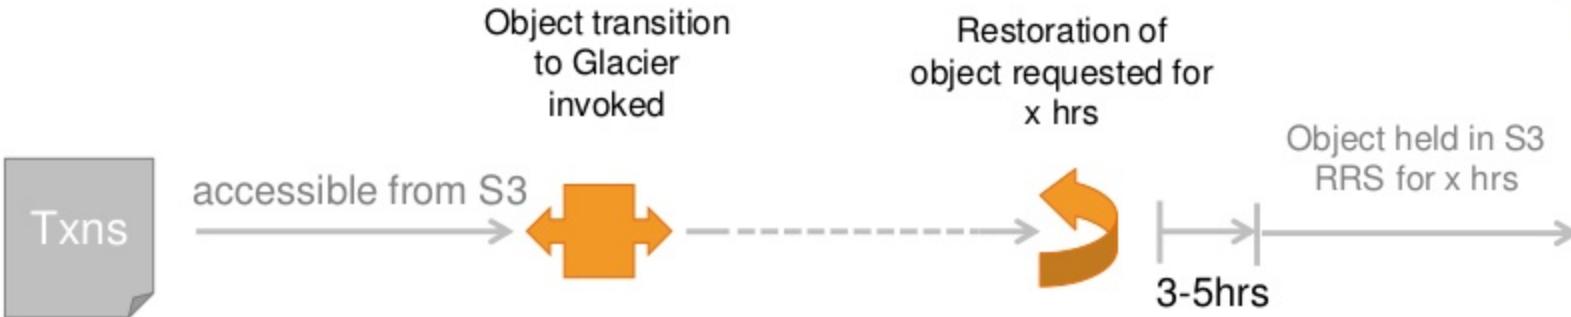
Expiry



time

Lifecycle Mgmt

Transition



Expiry



3-5 hour retrieval time

We assume you won't access often

Lifecycle Rule

Create a lifecycle rule to schedule the archival of objects to Glacier and/or permanent removal of objects. Objects transitioned to Glacier will no longer be immediately accessible. Most restores for transitioned objects will take 3 to 5 hours. [Learn more.](#)

Enabled:

Name (Optional):

Apply to Entire Bucket:

Prefix:

Time Period Format: Days from the creation date Date

Action	Time Period	X
Transition to Glacier	<input type="text" value="365"/> days from object's creation date	
Expiration (Delete Objects)	<input type="text" value="3650"/> days from object's creation date	

[Add Transition](#) [Add Expiration](#)

```
using (client = new AmazonS3Client()){
    var lifeCycleConfiguration = new LifecycleConfiguration()
    {
        Rules = new List<LifecycleRule>
        {
            new LifecycleRule
            {
                Id = "Archive and delete rule",
                Prefix = "projectdocs/",
                Status = LifecycleRuleStatus.Enabled,
                Transition = new LifecycleTransition()
                {
                    Days = 365,
                    StorageClass = S3StorageClass.Glacier
                },
                Expiration = new LifecycleRuleExpiration()
                {
                    Days = 3650
                }
            }
        }
    };
}
```

```
using (client = new AmazonS3Client()) {
    var lifeCycleConfiguration = new LifecycleConfiguration()
    {
        Rules = new List<LifecycleRule>
        {
            new LifecycleRule
            {
                Id = "Archive and delete rule",
                Prefix = "projectdocs/",
                Status = LifecycleRuleStatus.Enabled,
                Transition = new LifecycleTransition()
                {
                    Days = 365,
                    StorageClass = S3StorageClass.Glacier
                },
                Expiration = new LifecycleRuleExpiration()
                {
                    Days = 3650
                }
            }
        }
    };
}
```

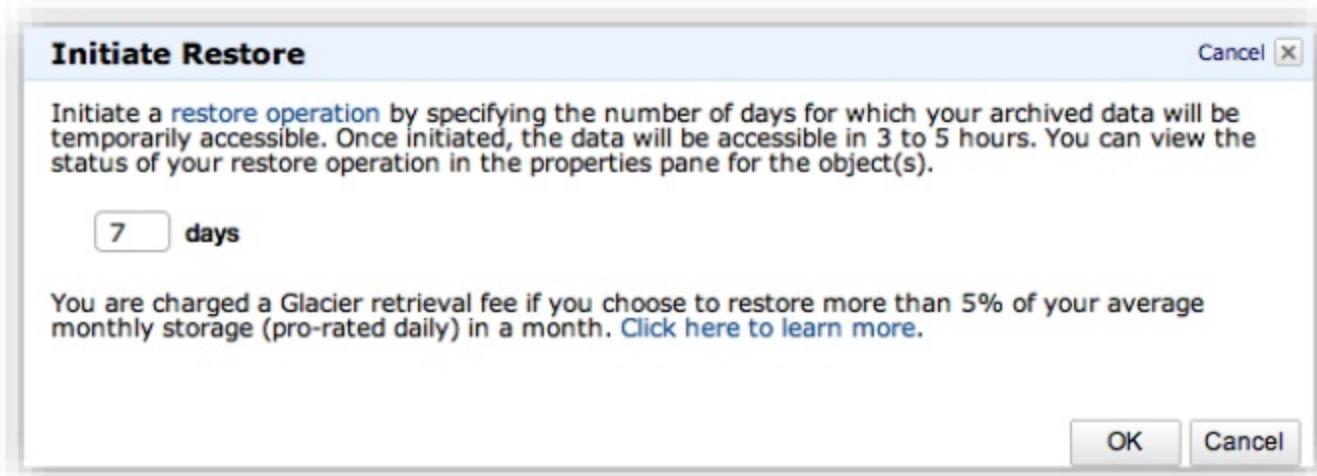
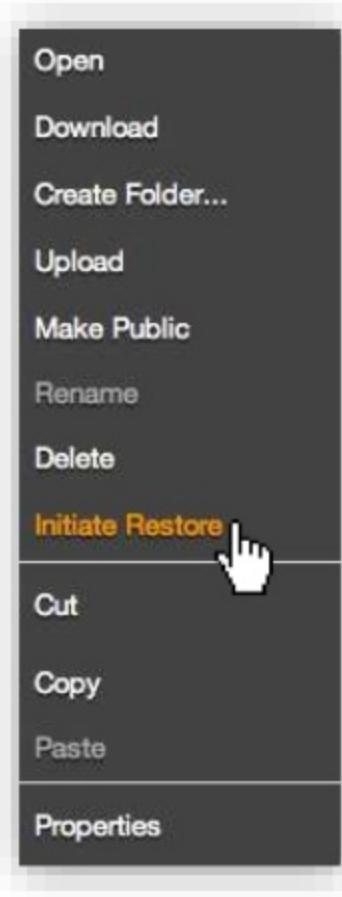
Transition to
Glacier after 1 year



```
using (client = new AmazonS3Client()) {
    var lifeCycleConfiguration = new LifecycleConfiguration()
    {
        Rules = new List<LifecycleRule>
        {
            new LifecycleRule
            {
                Id = "Archive and delete rule",
                Prefix = "projectdocs/",
                Status = LifecycleRuleStatus.Enabled,
                Transition = new LifecycleTransition()
                {
                    Days = 365,
                    StorageClass = S3StorageClass.Glacier
                },
                Expiration = new LifecycleRuleExpiration()
                {
                    Days = 3650
                }
            }
        }
    };
}
```

Delete object after
10 years





POST /*ObjectName*?restore HTTP/1.1
Host: *BucketName*.s3.amazonaws.com
Date: *date*
Authorization: *signatureValue*
Content-MD5: *MD5*

```
<RestoreRequest xmlns="http://s3.amazonaws.com/doc/2006-3-01">
  <Days>NumberOfDays</Days>
</RestoreRequest>
```

POST /*ObjectName*?restore HTTP/1.1
Host: *BucketName*.s3.amazonaws.com
Date: *date*
Authorization: *signatureValue*
Content-MD5: *MD5*

```
<RestoreRequest xmlns="http://s3.amazonaws.com/doc/2006-3-01">
  <Days>NumberOfDays</Days>
</RestoreRequest>
```

Response codes:

202 Accepted *Restore request accepted*

200 OK *Object already restored, number of days updated*

409 Conflict *Restoration already in progress*

Website hosting

Static sites straight from S3

It's a web service...

...so serve up web content

Enable website hosting

Index Document: index.html

Error Document: error.html

Edit Redirection Rules: You can set custom rules for web page requests for static website hosting.

Static Website Hosting

You can host your static website entirely on Amazon S3. Once you enable your bucket for static website hosting, all your content is accessible to web browsers via the Amazon S3 website endpoint for your bucket.

Endpoint: www.aws-exampl.es.s3-website-eu-west-1.amazonaws.com

Each bucket serves a website namespace (e.g. "www.example.com"). Requests for your host name (e.g. "example.com" or "www.example.com") can be routed to the contents in your bucket. You can also redirect requests to another host name (e.g. redirect "example.com" to "www.example.com"). See our walkthrough for how to set up an Amazon S3 static website with your host name.

Do not enable website hosting

Enable website hosting

Redirect all requests to another host name

To redirect requests to another bucket, enter the name of the target bucket below. If you are redirecting to a root domain address (e.g. example.com), see our walkthrough for configuring root domain website hosting.

Redirect all requests to

Save **Cancel**

Setting default documents

Redirecting requests

```
{  
    "Version": "2008-10-17",  
    "Statement": [{  
        "Sid": "PublicReadGetObject",  
        "Effect": "Allow",  
        "Principal": {  
            "AWS": "*"  
        },  
        "Action": ["s3:GetObject"],  
        "Resource": [ "arn:aws:s3:::example-bucket/*"  
        ]  
    }]  
}
```

Bucket policy

Website addressing

{bucket-name}.s3-website-{region}.amazonaws.com

e.g. mybucket.s3-website-eu-west-1.amazonaws.com

Normal addressing

s3-{region}.amazonaws.com/{bucket-name}/{object-key}

e.g. s3-eu-west-1.amazonaws.com/mybucket/img.png

{bucket-name}.s3-{region}.amazonaws.com/{object-key}

e.g. mybucket.s3-eu-west-1.amazonaws.com/img.png

Record set for:
aws-examples



Record set for:
aws-exampl.es



Website bucket name:

www.aws-exampl.es



Website bucket name:

aws-exampl.es

Record set for:
aws-exampl.es



Website redirect to:



Website bucket name:

www.aws-exampl.es

Website bucket name:

aws-exampl.es

Record set for:

aws-exampl.es



A Record 'Alias' to S3 website:

aws-exampl.es @
s3-website-eu-west-1.amazonaws.com

Website redirect to:

aws-exampl.es



Website bucket name:

www.aws-exampl.es

Website bucket name:

aws-exampl.es

Record set for:

aws-exampl.es

CNAME for www. to:

www.aws-exampl.es.s3-website-eu-west-1.amazonaws.com



A Record 'Alias' to S3 website:

aws-exampl.es @
s3-website-eu-west-1.amazonaws.com

Website redirect to:

aws-exampl.es

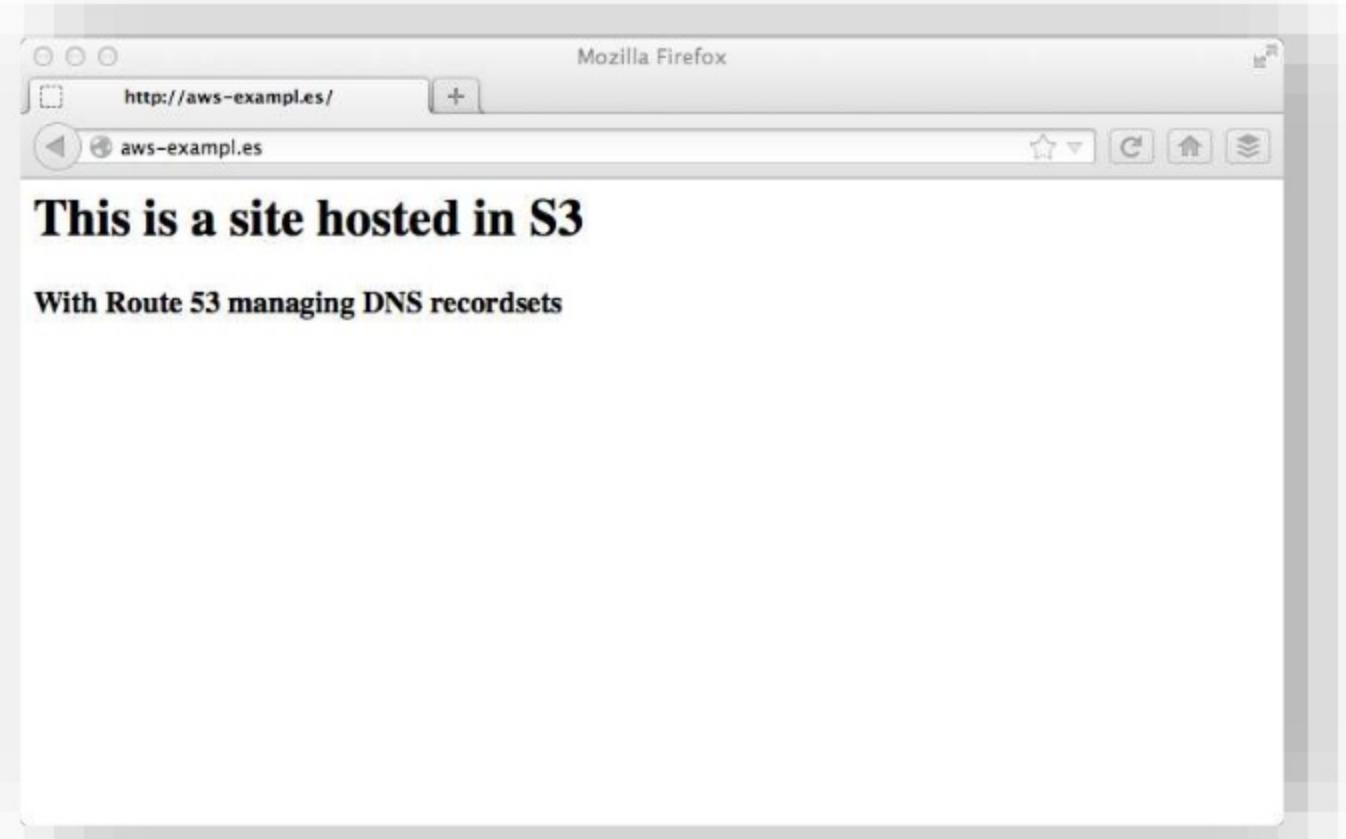


Website bucket name:

www.aws-exampl.es

Website bucket name:

aws-exampl.es



With Route 53 managing DNS recordsets

Time-Limited URLs

Control how long objects can be accessed for

Signed URLs

Provide time-limited access to specific objects that expires after a set period

Access Permissions

Use on objects in non-public buckets to prevent access once the signed URL has expired



```
https://ianmas-aws.testbucket.s3.amazonaws.com/testfile.txt  
?Signature=JHCa39GV1fKRKkEnAwzI88lH7f8%3D  
&Expires=1391425438  
&AWSAccessKeyId=AKIAIRBKB3ZAYAXFC2Q
```



>>>

Python boto

```
>>> import boto  
>>> conn = boto.connect_s3()  
>>> conn.generate_url(300, 'GET', bucket='ianmas-aws.testbucket',  
key='testfile.txt')
```

1st parameter is link
lifetime in seconds

```
'https://ianmas-  
aws.testbucket.s3.amazonaws.com/testfile.txt?Signature=QUWA%2BG0ohFeJ  
5pdEzxtdaIFI6w%3D&Expires=1391425142&AWSAccessKeyId=AKIAIRBKBJ3ZAYAX  
FC2Q'
```

Force a non-
SSL link

```
>>> conn.generate_url(300, 'GET', bucket='ianmas-aws.testbucket',  
key='testfile.txt', force_http=True)
```

```
'http://ianmas-  
aws.testbucket.s3.amazonaws.com/testfile.txt?Signature=tALx9KeeSisDS  
0N7K1M%2BIDFZXI%3D&Expires=1391425562&AWSAccessKeyId=AKIAIRBKBJ3ZAYAX  
FC2Q'
```

Error response: link
expired



```
<?xml version="1.0" encoding="UTF-8"?>
<Error>
  <Code>AccessDenied</Code>
  <Message>Request has expired</Message>
  <RequestId>70297390BE427DC7</RequestId>
  <Expires>2014-02-03T11:03:58Z</Expires>
  <HostId>I0rI00WUCnBttFSpEw6Mx4u8uRHgtOSw9k2euDW37skFCU7HH0u1SkUGGaUbn2vg</HostId>
  <ServerTime>2014-02-03T11:08:22Z</ServerTime>
</Error>
```

Object versioning

Preserving object histories

Bucket level

Automatically preserves
all copies of objects

Persistent

Even deleted object
history is held

Versioning

Versioning allows you to preserve, retrieve, and restore every version of every object stored in this bucket. This provides an additional level of protection by providing a means of recovery for accidental overwrites or deletions.

Once enabled, Versioning cannot be disabled and you will not be able to add Lifecycle Rules for this bucket.

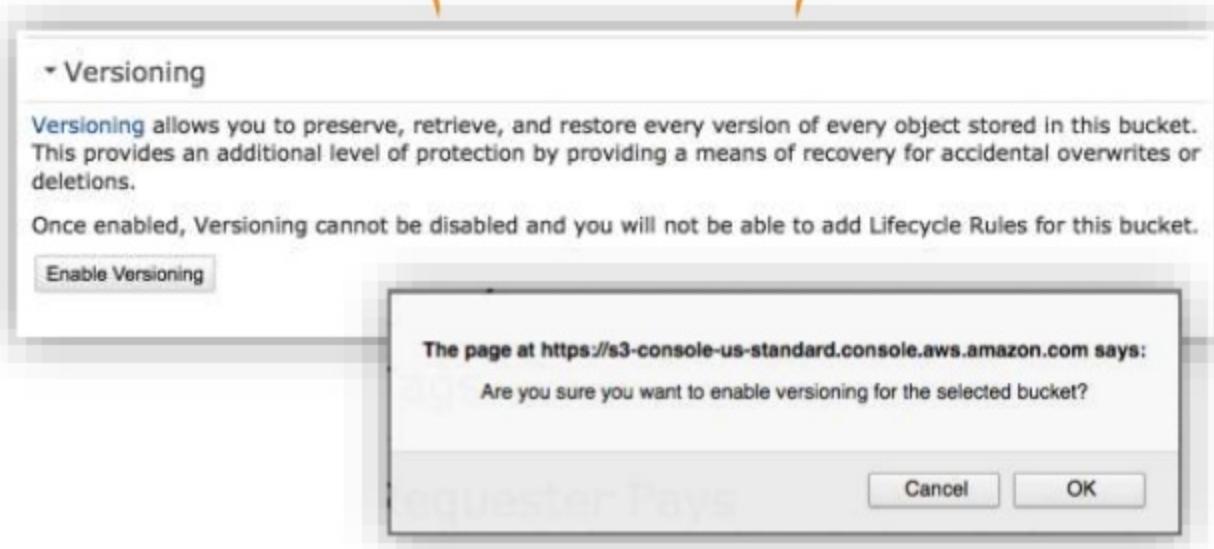
[Enable Versioning](#)

Bucket level

Automatically preserves
all copies of objects

Persistent

Even deleted object
history is held



Bucket level

Automatically preserves all copies of objects

Persistent

Even deleted object history is held

Versioning

Versioning allows you to preserve, retrieve, and restore every version of every object stored in this bucket. This provides an additional level of protection by providing a means of recovery for accidental overwrites or deletions.

Once enabled, Versioning cannot be disabled and you will not be able to add Lifecycle Rules for this bucket.

[Enable Versioning](#)

Versioning

Versioning allows you to preserve, retrieve, and restore every version of every object stored in this bucket. This provides an additional level of protection by providing a means of recovery for accidental overwrites or deletions.

Once enabled, Versioning cannot be disabled.

Enabled Suspended

[Save](#)

[Cancel](#)



>>>

Python boto



```
>>> import boto  
>>> conn = boto.connect_s3()  
>>> bucket = conn.get_bucket('mybucket')  
  
>>> versions = bucket.list_versions()  
>>> for version in versions:  
...     print version.name + version.version_id  
...
```

myfile.txt jU9eVv8000lP4PQx6zskMEyPIoExne57
myfile.txt x0JzMvMmGv0Bx2v4QpIypbkkH2XE2yyq
myfile.txt 8cjozv9Hmkzum8xj.8q8BZxR5CuXnzon

Object
version
IDs



Python boto

```
>>> key = bucket.get_key('myfile.txt',  
version_id='8cjozv9Hmkzum8xj.8q8BZxR5CuXnzon')  
>>> key.get_contents_as_string()
```

'this is version 1 of my file'



Grabbing the
contents of a
version

```
>>> key = bucket.get_key('myfile.txt',  
version_id='8cjozv9Hmkzum8xj.8q8BZxR5CuXnzon')  
>>> key.get_contents_as_string()
```

'this is version 1 of my file'

```
>>> key = bucket.get_key('myfile.txt',  
version_id='xOJzMvMmGv0Bx2v4QpIypbkkH2XE2yyq')  
>>> key.get_contents_as_string()
```

'this is version 2 of my file'

```
>>> key = bucket.get_key('myfile.txt',  
version_id='8cjozv9Hmkzum8xj.8q8BZxR5CuXnzon')  
>>> key.get_contents_as_string()
```

'this is version 1 of my file'

```
>>> key = bucket.get_key('myfile.txt',  
version_id='x0JzMvMmGv0Bx2v4QpIypbkkH2XE2yyq')  
>>> key.get_contents_as_string()
```

'this is version 2 of my file'

```
>>> key.generate_url(600)
```

'https://mybucket.s3.amazonaws.com/myfile.txt?Signature=ABCD&
Expires=1358857379&AWSAccessKeyId=AB&
versionId=x0JzMvMmGv0Bx2v4QpIypbkkH2XE2yyq'

Generating a '10
minute time bombed'
url for an older version

Metadata

System and user generated

System metadata

Name	Description	Editable?
Date	Object creation date	No
Content-Length	Object size in bytes	No
Content-MD5	Base64 encoded 128bit MD5 digest	No
x-amz-server-side-encryption	Server side encryption enabled for object	Yes
x-amz-version-id	Object version	No
x-amz-delete-marker	Indicates a version enabled object is deleted	No
x-amz-storage-class	Storage class for the object	Yes
x-amz-website-redirect-location	Redirects request for the object to another object or external URL	Yes

User metadata

Key-value pairs stored with objects and returned with requests

x-amz-meta-{your metadata key in here}

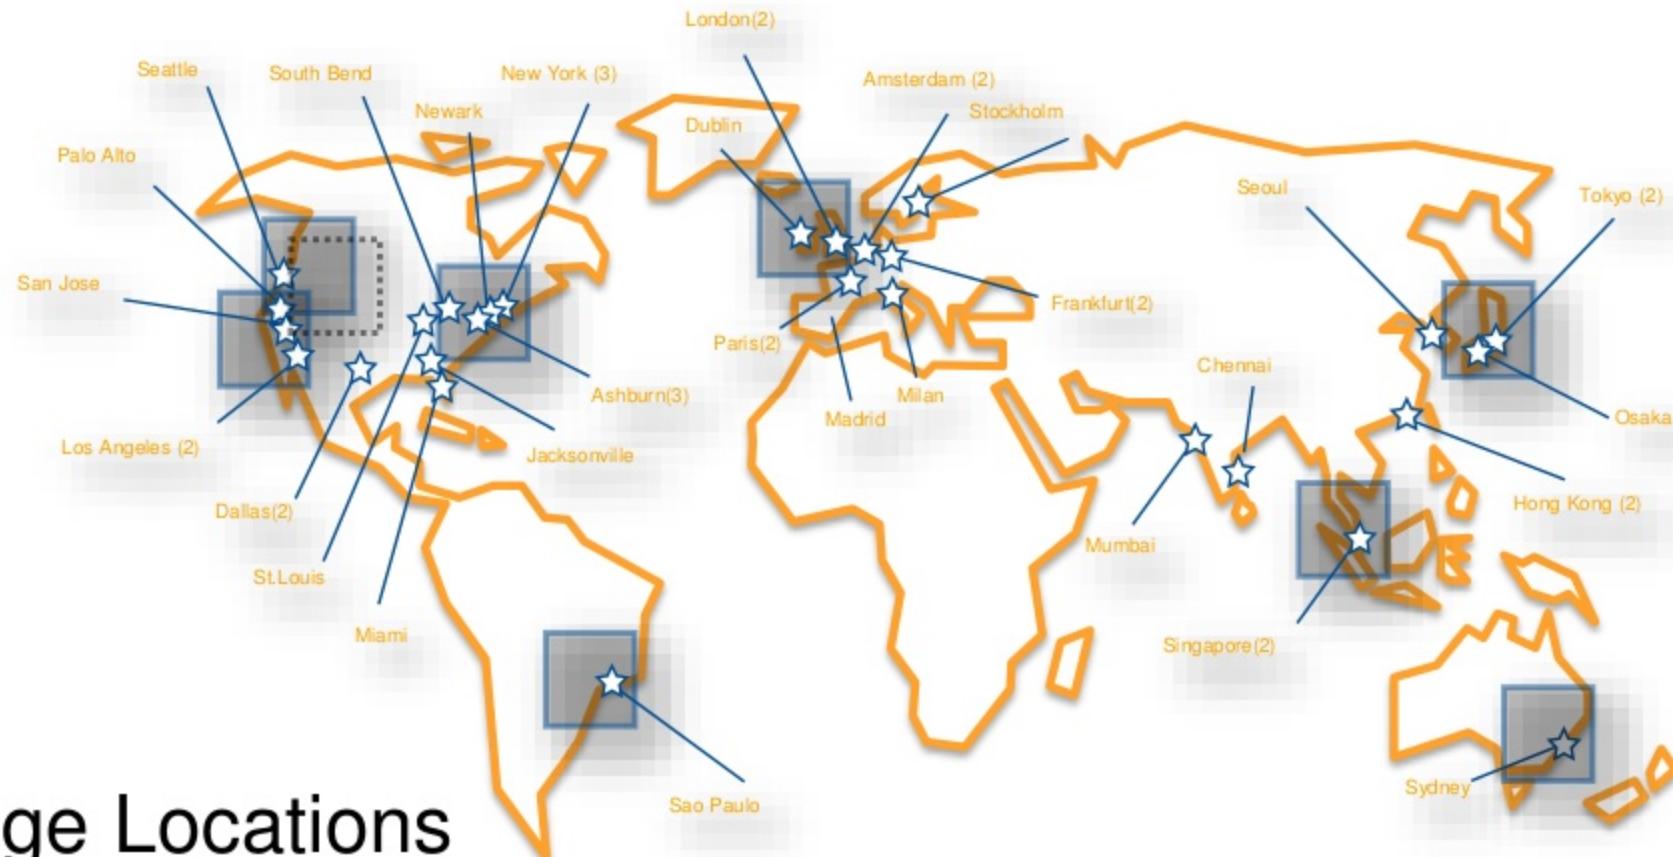
```
>>> key.set_metadata('my_tag', 'my metadata')

>>> key.get_metadata('my_tag')
'my metadata'
```

CloudFront

Content delivery from global edge locations

CloudFront Edge Locations



★ Edge Locations

To deliver content to end users with lower latency

A global network of edge locations Supports global DNS infrastructure (Route53) and Cloud Front CDN

Global content distribution

Download

Enable static and dynamic assets
to be served from edge locations



Streaming

Serve RTMP directly from media files
in buckets

CloudFront Management Console

Services EC2 S3 RDS Route 53 Edit lanmas @ lanmas-aws Global Help

CloudFront > Create Distribution

Step 1: Select delivery method

Step 2: Create distribution

Select a delivery method for your content. [Learn More](#)

Web

Create a web distribution if you want to:

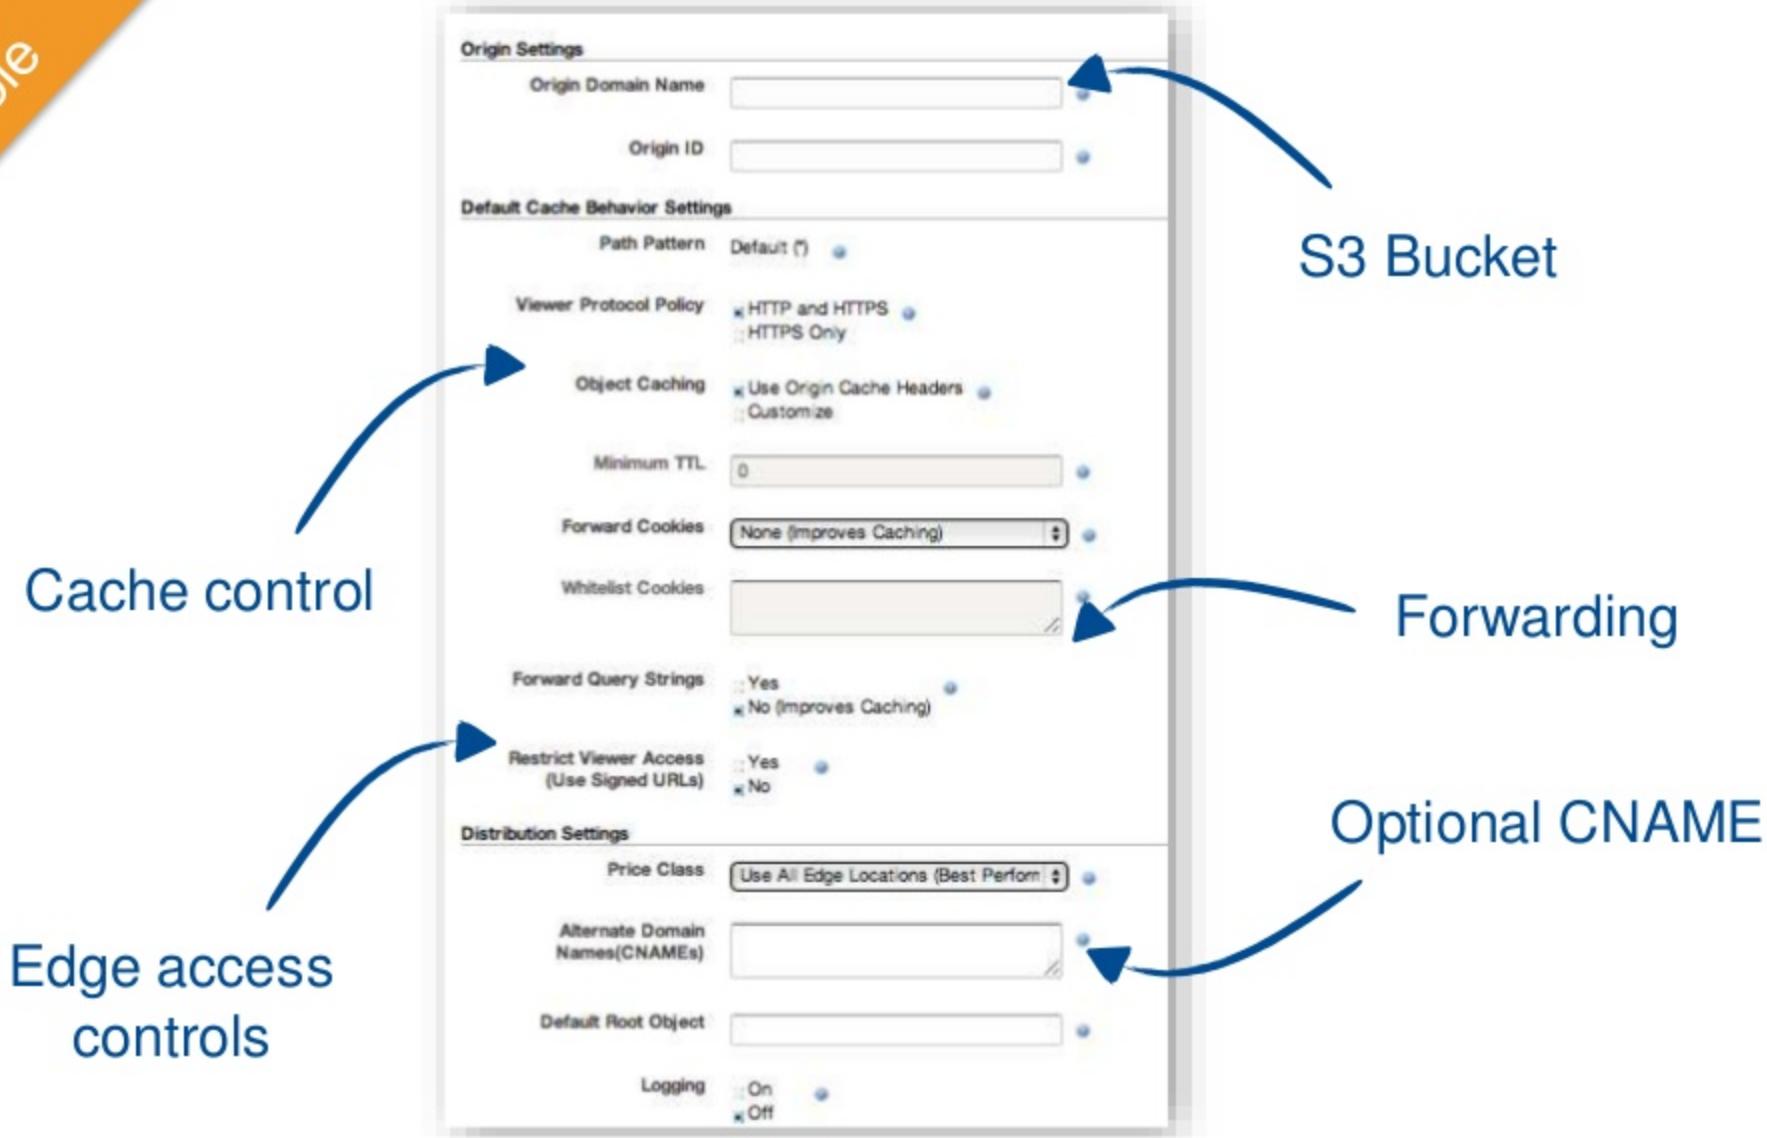
- Speed up distribution of static and dynamic content, for example, .html, .css, .php, and graphics files.
- Distribute media files using HTTP or HTTPS.
- Add, update, or delete objects, and submit data from web forms.
- Use live streaming to stream an event in real time.

You store your files in an origin — either an Amazon S3 bucket or a web server. After you create the distribution, you can add more origins to the distribution.

Cancel Continue

© 2008 - 2014, Amazon Web Services, Inc. or its affiliates. All rights reserved. [Privacy Policy](#) [Terms of Use](#) [Feedback](#)

Console



CloudFront Management Console

Services EC2 S3 RDS Route 53 Edit ianmas @ ianmas-aws

CloudFront > Create Distribution

Step 1: Select delivery method

Select a delivery method for your content. Learn More

Step 2: Create distribution

Web

RTMP

Create an RTMP distribution to speed up distribution of your streaming media files using Adobe Flash Media Server's RTMP protocol. An RTMP distribution allows an end user to begin playing a media file before the file has finished downloading from a CloudFront edge location. Note the following:

- To create an RTMP distribution, you must store the media files in an Amazon S3 bucket.
- To use CloudFront live streaming, create a web distribution.

Cancel Continue

CloudFront > Create Distribution

Step 1: Select delivery method

Step 2: Create distribution

RTMP Distribution Settings

Origin Domain Name: ianmas-media.testbucket.s3.amazonaws.com

Restrict Bucket Access: No

Restrict Viewer Access (Use Signed URLs): No

Price Class: Use All Edge Locations (Best Performance)

Alternate Domain Names (CNAMEs):

Logging: Off

Bucket for Logs:

Log Prefix:

Comment:

Distribution State: Enabled

S3 Bucket

Optional CNAME

Logging to S3

Cancel Back Create Distribution

Example

```
<html>
<script type='text/javascript'
src='http://d2ew7gdzogp20x.cloudfront.net/jwplayer/jwplayer.js'></script>

<body>
<div id='player'></div>
<script type='text/javascript'>
jwplayer('player').setup({
  file: "rtmp://s1eat02wfxn38u.cloudfront.net/cfx/st/montage-medium.mp4",
  width: "480",
  height: "270",
});
</script>
</body>
</html>
```

Download distribution

Streaming distribution

Summary

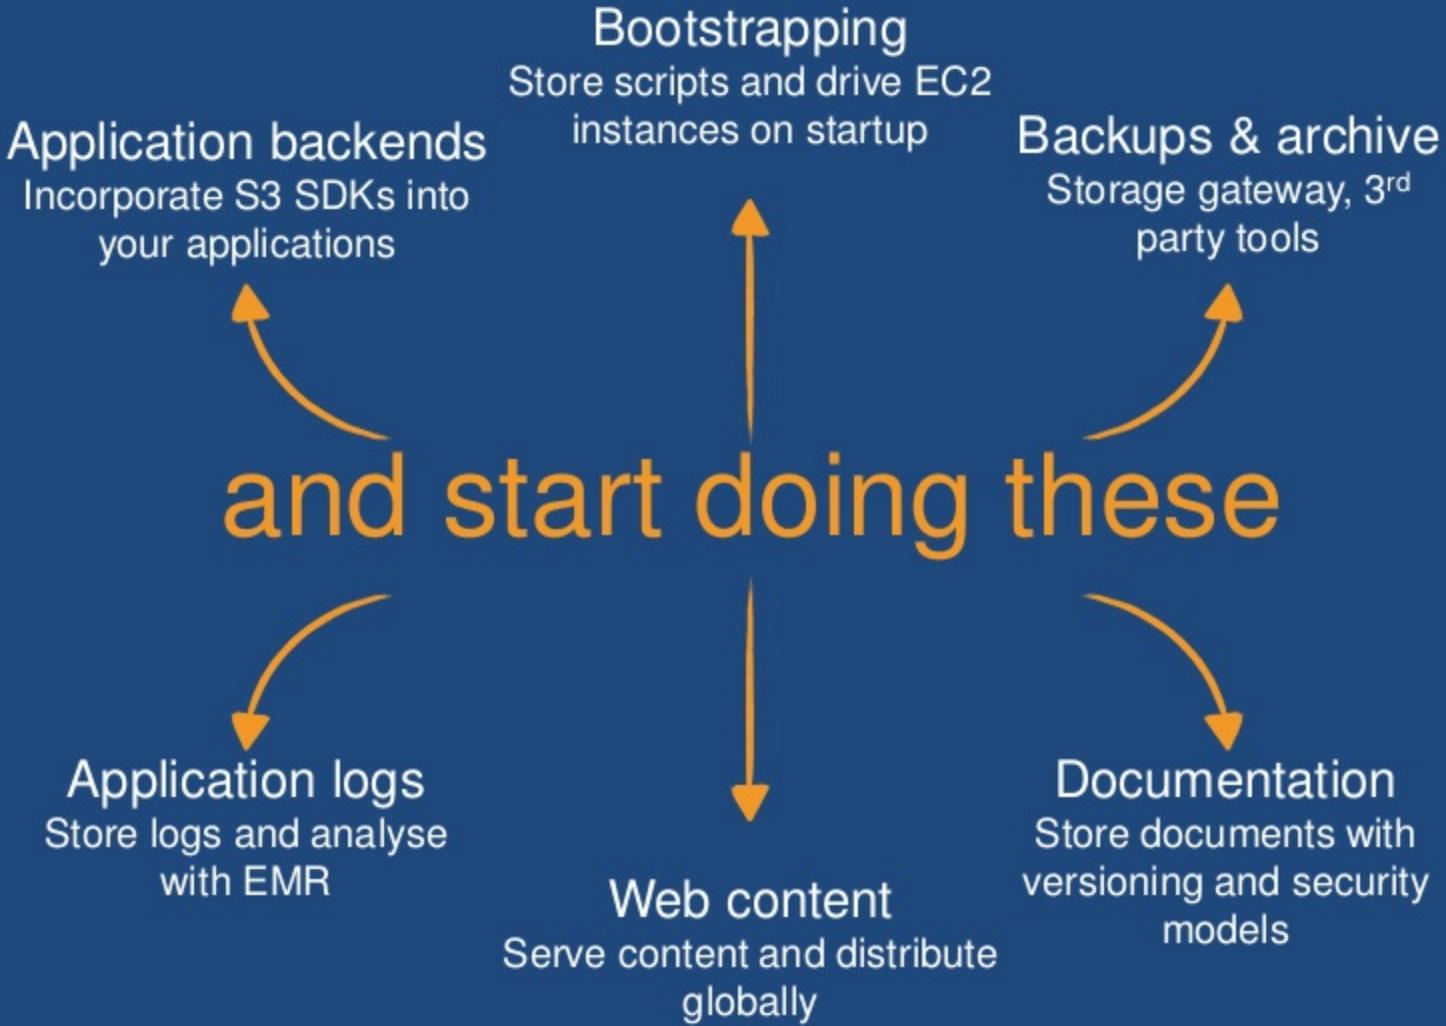
Stop doing these:

Capacity planning

Management of storage

Worrying about backing up the backup

Fixing broken hardware



aws.amazon.com

Follow us for more
events & webinars



Ian Massingham – Technical Evangelist

 @IanMmmm



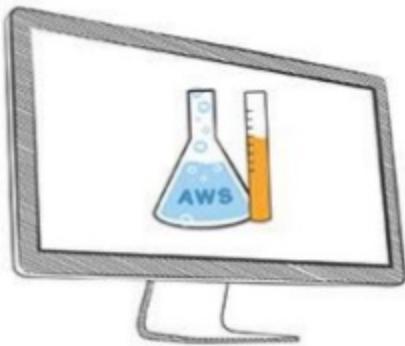
@AWS_UK for local AWS events & news



@AWScloud for Global AWS News and Announcements

AWS Training & Certification

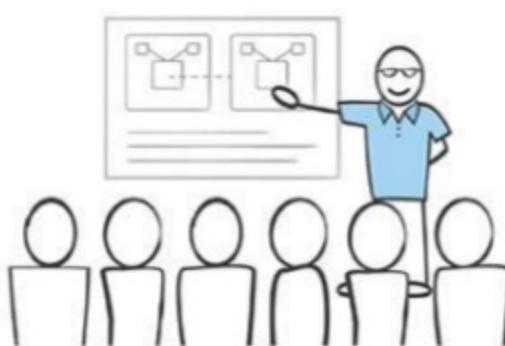
Self-Paced Labs



Try products, gain new skills, and get hands-on practice working with AWS technologies

[aws.amazon.com/training/
self-paced-labs](https://aws.amazon.com/training/self-paced-labs)

Training



Skill up and gain confidence to design, develop, deploy and manage your applications on AWS

aws.amazon.com/training

Certification



Demonstrate your skills, knowledge, and expertise with the AWS platform

aws.amazon.com/certification

We typically see customers start by trying our services

All Products

Compute & Networking

Storage

Database

Application Services

Development & Management

AWS Marketplace Software

FAQ »

Find answers to common questions about the AWS Free Tier.



Amazon EC2 »

Web service that provides resizable compute capacity in the cloud.



Amazon S3 »

Highly-scalable, reliable, and low-latency data storage.



Amazon RDS »

Managed MySQL, Oracle and SQL Server databases.



Amazon CloudWatch »

Monitoring for AWS cloud resources and applications.



AWS Data Pipeline »

Orchestration for data-driven workflows.



Amazon DynamoDB »

Fully managed NoSQL database service with seamless scalability.



Amazon EBS »

Highly available, highly reliable, predictable storage volumes.



Amazon ELB »

Web service that provides scalability and high availability.



Amazon ElastiCache »

Managed scale-out caching.



Amazon SNS »

Web service to set up, operate, and send notifications from the cloud.



Amazon Elastic Transcoder »

Convert your media files easily, at low cost and at scale.



Amazon SQS »

Scalable queue for storing messages as they travel between computers.



Amazon SWF »

Workflow service for building scalable, resilient applications.



AWS Marketplace »

Partner software pre-configured to run on AWS.

Get started now at : aws.amazon.com/getting-started



Design your application for the AWS Cloud

AWS Reference Architectures

The flexibility of AWS allows you to design your application architectures the way you like. AWS Reference Architecture Datasheets provide you with the architectural guidance you need in order to build an application that takes full advantage of the AWS cloud. Each datasheet includes a visual representation of the architecture and basic description of how each service is used.



Web Application Hosting
Build highly-scalable and reliable web or mobile-web applications ([PDF](#))



Content and Media Serving
Build highly reliable systems that serve massive amounts of content and media ([PDF](#))



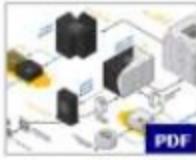
Batch Processing
Build auto-scalable batch processing systems like video processing pipelines ([PDF](#))



Fault tolerance and High Availability
Build systems that quickly failover to new instances in an event of failure ([PDF](#))



Large Scale Processing and Huge Data sets
Build high-performance computing systems that involve Big Data ([PDF](#))



Ad Serving
Build highly-scalable online ad serving solutions ([PDF](#))



Disaster Recovery for Local Applications
Build cost-effective Disaster Recovery solutions for on-premises applications ([PDF](#))



File Synchronization
Build simple file synchronization service ([PDF](#))

More details on the AWS Architecture Center at : aws.amazon.com/architecture

