

AWS

S U M M I T

# Migrating your Databases to AWS: Deep Dive on Amazon RDS and AWS Database Migration Service

Blair Layton, Business Development, Database  
Services, APAC

June 21, 2017





Amazon  
RDS

Relational Databases

Managed Service

Simple and fast to scale

Fast, predictable performance

Low cost, pay for what you use



ORACLE®



Amazon  
Aurora

# Key Features

Provision a database in 6 minutes

Provision a MAZ database with a few mouse clicks

Scale a database up/down with 60-90 seconds downtime

Apply patches with 60-90 seconds downtime

Add read replicas with a few mouse clicks

Protect your backups and logs with 11 9's of durability

Recover to any point in time from nightly backups + logs

Detailed metrics, down to 1 second intervals


Secure your data with single click encryption at rest

**Start/Stop your dev/test databases for up to 7 days**



Amazon  
RDS

# Amazon RDS Customers



# What is Amazon Aurora?

**MySQL-compatible and PostgreSQL-compatible  
relational database platform**

**Performance and availability of  
commercial databases**

**Simplicity and cost-effectiveness of  
open source databases**

**Delivered as a managed service**



# The Amazon Aurora Database Family



# Recent Announcements

**Performance enhancements:** Fast DDL, fast index build, spatial indexing, hot row contention

**Availability features:** Zero-downtime patching, database cloning (Q2), database backtrack (Q2)

**Eco-system integration:** Load from S3, IAM integration (Q2), select into S3 (Q2), log upload to CloudWatch Logs & S3 (Q2)

**Cost reduction:** t2.small – cuts cost of entry by half – you can run Aurora for \$1 / day

**Growing footprint:** London, Montreal, Ohio, and San Francisco – now available in all 3AZ regions

**2/3** of top 100 AWS customers  
**8** of top 10 gaming customers



**Fastest growing service in AWS history**



# RDS Engine Capability Matrix

Feature	Aurora	MySQL	MariaDB	PostgreSQL	Oracle	SQL Server
VPC	✓	✓	✓	✓	✓	✓
High availability	✓	✓	✓	✓	✓	✓
Instance Scaling	✓	✓	✓	✓	✓	✓
Encryption	✓	✓	✓	✓	✓	✓
Read replicas	✓	✓	✓	✓	Oracle Golden Gate / DMS	DMS
Cross region replicas	✓	✓	✓	✓		
Max Storage	64 TB	6 TB	6 TB	6 TB	6 TB	4 TB
Scale Storage	Auto scaling	✓	✓	✓	✓	
Provisioned IOPS	NA	30,000	30,000	30,000	30,000	20,000
Largest Instance	R3.8XL	R3.8XL M4.10XL	R3.8XL M4.10XL	R3.8XL M4.10XL	R3.8XL M4.10XL	R3.8XL M4.10XL

## **Customers Want to Migrate to AWS, but...**

They can't afford long periods of application downtime

Tools that enable minimal downtime are expensive

It seems too complex and expensive to migrate

They still need a copy of the data on-premise

They want to migrate to an open source database

Sending large volumes of data to AWS requires an expensive international network link

They don't have the skills inside their organization

# Traditional Approach to Migrate to AWS

1. Create your AWS account
2. Setup your Virtual Private Cloud (VPC) in AWS
3. Connect to AWS with a VPN or Direct Connect
4. Shutdown and backup your database
5. Transmit the backup to S3
6. Configure an EC2 instance with the DB software
7. Restore the backup
8. Configure EC2 instances for the application
9. Switch the users to use AWS

# Traditional Approach to Migrate to AWS

1. Create your AWS account
2. Setup your Virtual Private Cloud (VPC) in AWS
3. Connect to AWS with a VPN or Direct Connect
4. Shutdown and backup your database
5. Transmit the backup to S3
6. Configure an EC2 instance with the DB software
7. Restore the backup
8. Configure EC2 instances for the application
9. Switch the users to use AWS

Steps 4-9 could take a week or more!

**Now There is a Better Way!**



# AWS Database Migration Service (AWS DMS)

*DMS migrates databases to AWS easily and securely with minimal downtime. It can migrate your data to and from most widely used commercial and open-source databases.*



ORACLE



# New NoSQL support

## Migrate to AWS

- Move from MongoDB to Amazon DynamoDB
- Move from MongoDB to relational db's




## Move between NoSQL and SQL

- Change technologies



# Keep Your Apps Running During the Migration

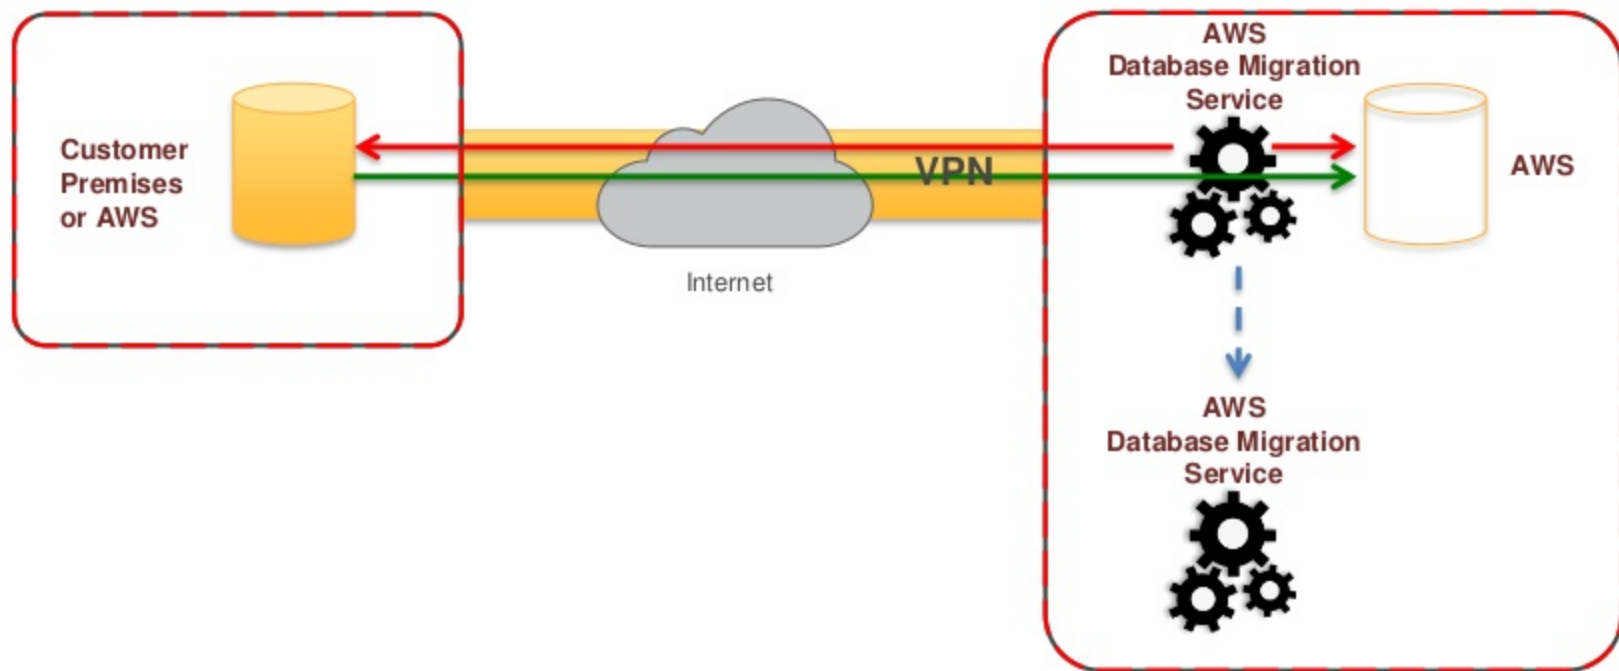


Start a replication instance  
Connect to source and target  
databases  
Select tables, schemas, or  
databases




Let AWS DMS create tables,  
load data, and keep them in  
sync  
Switch applications over to the  
target at your convenience

# Multi-AZ Option for High Availability




# Change Data Capture (CDC) and Apply







# Load is Table by Table




# Multiple Targets



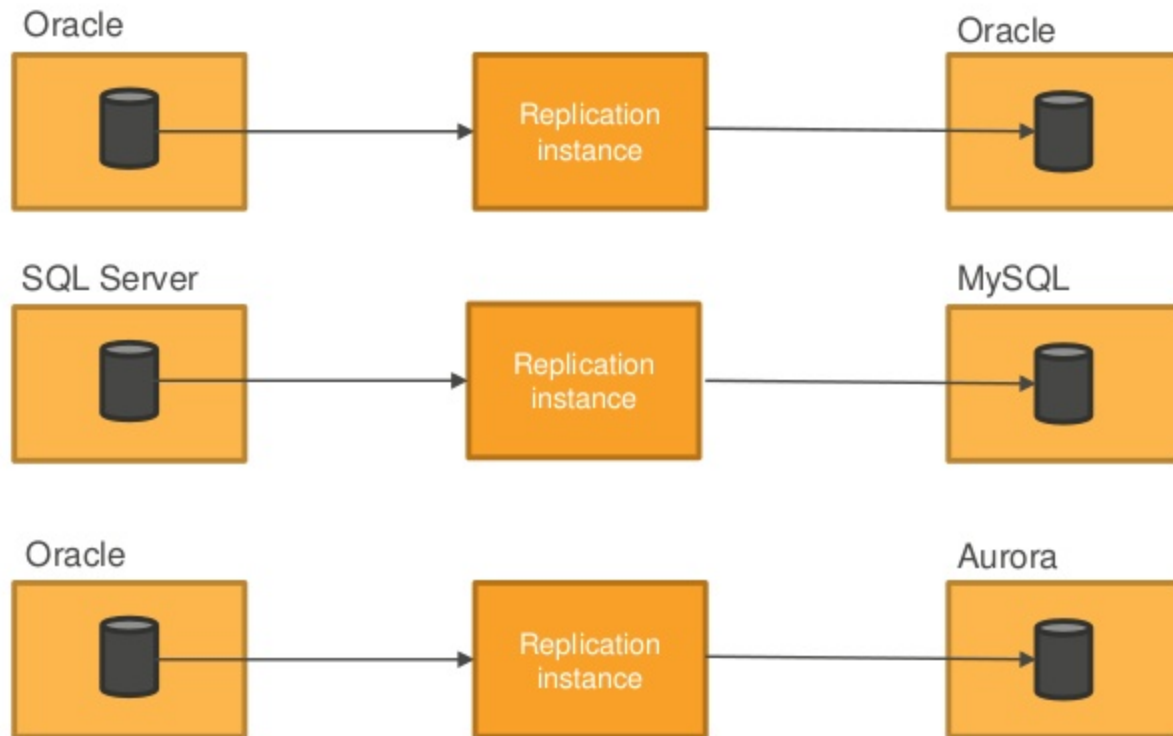
# Multiple Sources



# Customers Don't Have to Take Everything



# Homogenous or Heterogeneous





**For less than \$10 per TB!**

# Strengths and Focus Areas

DMS can act as a replication/migration Swiss Army knife, but is not a magic wand.



## Use It

- Heterogeneous migrations
- Minimal downtime required
- No native solution

## Caution

- Some tables with large LOBs
- Complex data types
- High load database

## Don't Use It

- Native no downtime solution exists
- Can take downtime + native
- > 5 TB + slow Internet\*\*\*\*


# AWS Schema Conversion Tool (AWS SCT)

*SCT helps automate many database schema and code conversion tasks when migrating between database engines or data warehouse engines*



# AWS Schema Conversion Tool


Supports Schema Copy and RDS Recommendations



# SCT Data Extractors

**Extract Data from your data warehouse\* and migrate to Amazon Redshift**

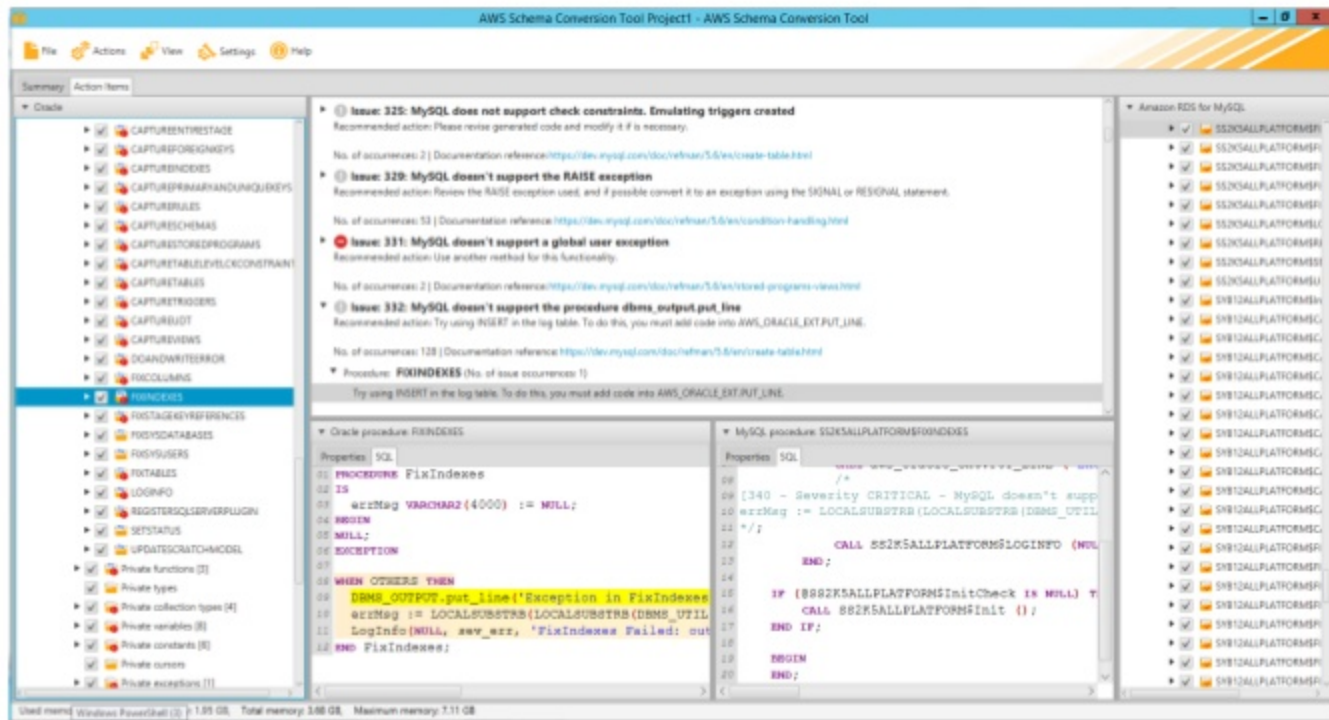
- Extracts through local migration agents
- Data is optimized for Redshift and Saved in local files
- Files are loaded to an Amazon S3 bucket (through network or Amazon Snowball) and then to Amazon Redshift



\* 1<sup>st</sup> release supports Oracle v11 and up and Teradata v14 and up



# SCT Helps with Converting Tables, Views & Code



Sequences  
User-Defined Types  
Synonyms  
Packages  
Stored Procedures  
Functions  
Triggers  
Schemas  
Tables  
Indexes  
Views  
Sort and distribution keys

# SCT can tell you how hard the migration will be

## Database Migration Assessment Report

Source Database: RDS, ADMINISTRATION (db, administration) 14 (17) 36-46 complete 1.amazonaws.com:81  
90 (90%)  
Oracle Database 12c Enterprise Edition (12.1.0.2.0 (64bit Production))



### Executive Summary

We completed the analysis of your Oracle source database and estimate that 91% of the database storage objects and 100% of database code objects can be converted automatically or with minimal changes if you select Amazon Aurora as your migration target. Database storage objects include schemas, tables, columns, constraints, indexes, sequences, synonyms, user-defined types and types. Database code objects include functions, procedures, packages, triggers, views, materialized views, events, SQL, scalar functions, SQL, table functions, SQL, table functions, attributes, variables, constraints, table types, public types, private types, cursors, exceptions, parameters and other objects. Based on our analysis of SQL system elements of your source database schema, we estimate that 99.9% of your entire database schema can be converted automatically to Amazon Aurora. To complete the migration, we recommend 297 conversion action(s) ranging from simple tasks to medium-complexity actions to significant conversion actions.

### Database Objects with Conversion Actions for Amazon Aurora


Of the total 1,276 database storage object(s) and 125 database code object(s) in the source database, we were able to identify 1,427 (91%) database storage object(s) and 125 (100%) database code object(s) that can be converted automatically or with minimal changes to Amazon Aurora.

149 (9%) database storage object(s) required 149 significant user action(s) to complete the conversion.

Figure: Conversion statistics for database storage objects



Figure: Conversion statistics for database code objects



### Detailed Recommendations for Amazon Aurora Migrations

If you choose to migrate your Oracle database to Amazon Aurora, we recommend the following actions.

1. Connect SCT to Source and Target databases.

2. Run Assessment Report.

3. Read Executive Summary.

4. Follow detailed instructions.

## Database Migration Assessment Report

Source Database: RDS, ADMINISTRATION (db, administration) 14 (17) 36-46 complete 1.amazonaws.com:81  
90 (90%)  
Oracle Database 12c Enterprise Edition (12.1.0.2.0 (64bit Production))



### Storage Object Actions

#### Sequence Changes

Some changes are required to sequences that cannot be converted automatically. You'll need to address these issues manually.

#### Issue 341: MySQL doesn't support sequences

Recommended Action: Try developing a system for sequences in your application.

Issue Code: 341 | No. of Occurrences: 134 | Estimated Complexity: Significant  
Schemas: RDS\_ADMINISTRATION Sequences: BACKUP\_ID\_SEQUENCE  
Schemas: RDS\_ADMINISTRATION Sequences: CERTIFICATE\_ID\_SEQUENCE  
Schemas: RDS\_ADMINISTRATION Sequences: CHARACTER\_SET\_ID\_SEQ  
Schemas: RDS\_ADMINISTRATION Sequences: CUSTOMER\_SUBNET\_GROUP\_ID\_SEQ  
Schemas: RDS\_ADMINISTRATION Sequences: CUSTOMER\_SUBNET\_ID\_SEQ  
+129 more

#### Index Changes

Some changes are required to indexes that cannot be converted automatically. You'll need to address these issues manually.

#### Issue 207: MySQL doesn't support function indexes

Recommended Action: Review your code and try to use simple indexes.

Issue Code: 207 | No. of Occurrences: 3 | Estimated Complexity: Significant  
Documentation References: <https://dev.mysql.com/doc/refman/5.6/en/create-table.html>  
Schemas: RDS\_ADMINISTRATION Tables: DBL\_ENGINE\_SEEDS Indexes: I\_DBL\_ENG\_SEED\_DBL\_ENG\_CONF\_ID  
Schemas: RDS\_ADMINISTRATION Tables: RDS\_SYSTEM\_ACCOUNTS Indexes: I\_SYS\_ACCOUNT\_DEFAULT  
Schemas: RDS\_ADMINISTRATION Tables: RUNNABLE\_DBL\_CONFIG Indexes: I\_RUNNABLE\_DBL\_CFG\_PREFERRERRED

#### Constraint Changes

Some changes are required to constraints that cannot be converted automatically. You'll need to address these issues manually.

#### Issue 210: MySQL doesn't support FUNCTION AS DEFAULT VALUE

Recommended Action: Try using a trigger.

Issue Code: 210 | No. of Occurrences: 2 | Estimated Complexity: Simple  
Documentation References: <https://dev.mysql.com/doc/refman/5.6/en/create-table.html>  
Schemas: RDS\_ADMINISTRATION Tables: CUSTOMERS Constraints: CK\_CUSTOMER\_TRUST\_LEVEL\_STATE: 0:30  
Schemas: RDS\_ADMINISTRATION Tables: STORAGE\_VOLUMES Constraints: CK\_SV\_EFFECTIVE: 0:8

#### Issue 325: MySQL does not support check constraints. Enabling triggers created

Recommended Action: Please review generated code and modify it if necessary.

Issue Code: 325 | No. of Occurrences: 283 | Estimated Complexity: Simple  
Documentation References: <https://dev.mysql.com/doc/refman/5.6/en/create-table.html>

# Pricing and Terms and Conditions

**\$0**

**for software license**


## Pricing

- Free software license
- For active AWS customers with accounts in good standing

## Allowed Use

- Use SCT to migrate database schemas to Amazon RDS, Amazon Redshift, or Amazon EC2-based databases
- To use SCT to migrate schemas to other destinations, contact for special pricing

# Customer Successes



# Split migration



35 million members on it's site

Many interdependent applications built over the last 15 years with unique High IOPS requirements

Migration from legacy cloud service to AWS

Combination of migration to MySQL on EC2 and Amazon Aurora

Benefits:

- DMS reduced the time required to migrate our databases by 40 percent
- Realized 55 percent cost savings by moving some db's to Amazon Aurora

# **Case Study – Migration from Microsoft SQL Server to RDS PostgreSQL**

# Uberfusion Case Study – Data Migration for SSO Media

## Company Profile

- ASEAN Media Company
- 4.7 million residential customers
- Approx. 66% penetration of TV households
- 183 TV channels including 53 HD channels

>2 million

Approximate Registered SSO Users  
(June 2016)



## On Premises Database

MSSQL Server 2008 R2 Enterprise Edition

Database Stats (April 2016):

- Number of Tables: 102
- Database Size: ~ 55GB
- Server Specs: 128GB RAM, Intel Xeon CPU with 24 Cores



~37

Integrated Applications  
(June 2016)

## Ideal State

- database in AWS caters for SSO's migration to the cloud
- able to manage current on-premises database load
- able to scale with the application
- must be highly available
- maintain cost efficiency while performing all of the above



# Problem Statement, Mandate and Challenge

## State of On-premise MSSQL Sever Database

### Limited Scalability

As the on-premise MSSQL Server Database is a physical machine, scaling requires purchasing and installing additional hardware.

### Higher Costs

Hardware upgrades are costly with a high hardware delivery turn-around time. Additionally, MSSQL Server Enterprise Edition licenses are high in costs.

### Shared Infrastructure

The on-premises database infrastructure is shared across multiple applications. A problem caused by another application could potentially affect SSO.

## New Mandate for SSO

To migrate the on-premise SSO system into Amazon Web Services (AWS) cloud to reduce cost and improve scalability.



UEFA  
EURO2016  
FRANCE



### Urgency

To be ready to handle high load during Euro and Olympics 2016

## Migration into the Cloud



**No data loss and completed in a short time frame at minimal costs**

The decision was made to go with RDS PostgreSQL, as:

- MSSQL on RDS does not support Multi Availability Zone in AWS Singapore and hence cannot be highly available
- No need to maintain a VPN tunnel back to premise
- Lower OPEX costs compared to running MSSQL on RDS or managing our own DB servers

**~11.5x cheaper**  
Compared to RDS MSSQL Server Enterprise

## The Challenge

### Zero Data Loss

#### Zero Data Loss

All user-related data needed to be migrated with zero data loss up to the point of migration.

### Transparent

#### Transparent to Users

There should be no visible impact to SSO users. Encrypted data such as passwords needed to be migrated seamlessly.

### Time

#### Short Migration Window

Data migration had to be completed within a 90 minute window.

### Legacy

#### Legacy Data

There is a lot of legacy data in the on-premises database which is not clean. Dummy data and special characters exist where they should not.

### Managed Service

#### Leverage Managed Services

To leverage Amazon's relational database service (RDS), instead of managing our own database servers.

### Testing

#### System Needed to Be Testable

System needed to be testable on migration night by a small team of testers while ensuring the general public cannot change the data.



# The Solution

## Uberfusion's Delivery for Data Migration for SSO System

### Challenge: Zero Data Loss

- All user-related data tables were identified
- System was put into read-only mode before starting migration
- Views were created in the schema of the new PostgreSQL database, and exported to flat files
- Number of records were verified for pre- and post-migration
- Data was verified for accuracy from a sample of users post-migration



### Challenge: Short Migration Window & Legacy Data

- Data cleansing was performed
- Instance created with access to both source and target database for ease of migration
- Export scripts automated for speed and accuracy
- Multiple scripts to cater for special characters in data
- Logs offloaded from database to Elasticsearch and Kibana



### Challenge: Leverage Managed Service

- Amazon's RDS service for PostgreSQL was used
- System code base was updated to integrate with the new database



Total Migration Time

<1hr

### Challenge: Transparent to Users

- PostgreSQL Membership library was modified to mimic ASP.NET Membership in regards to password encryption
- Migration was tested beforehand to ensure users would still be able to login with their existing passwords



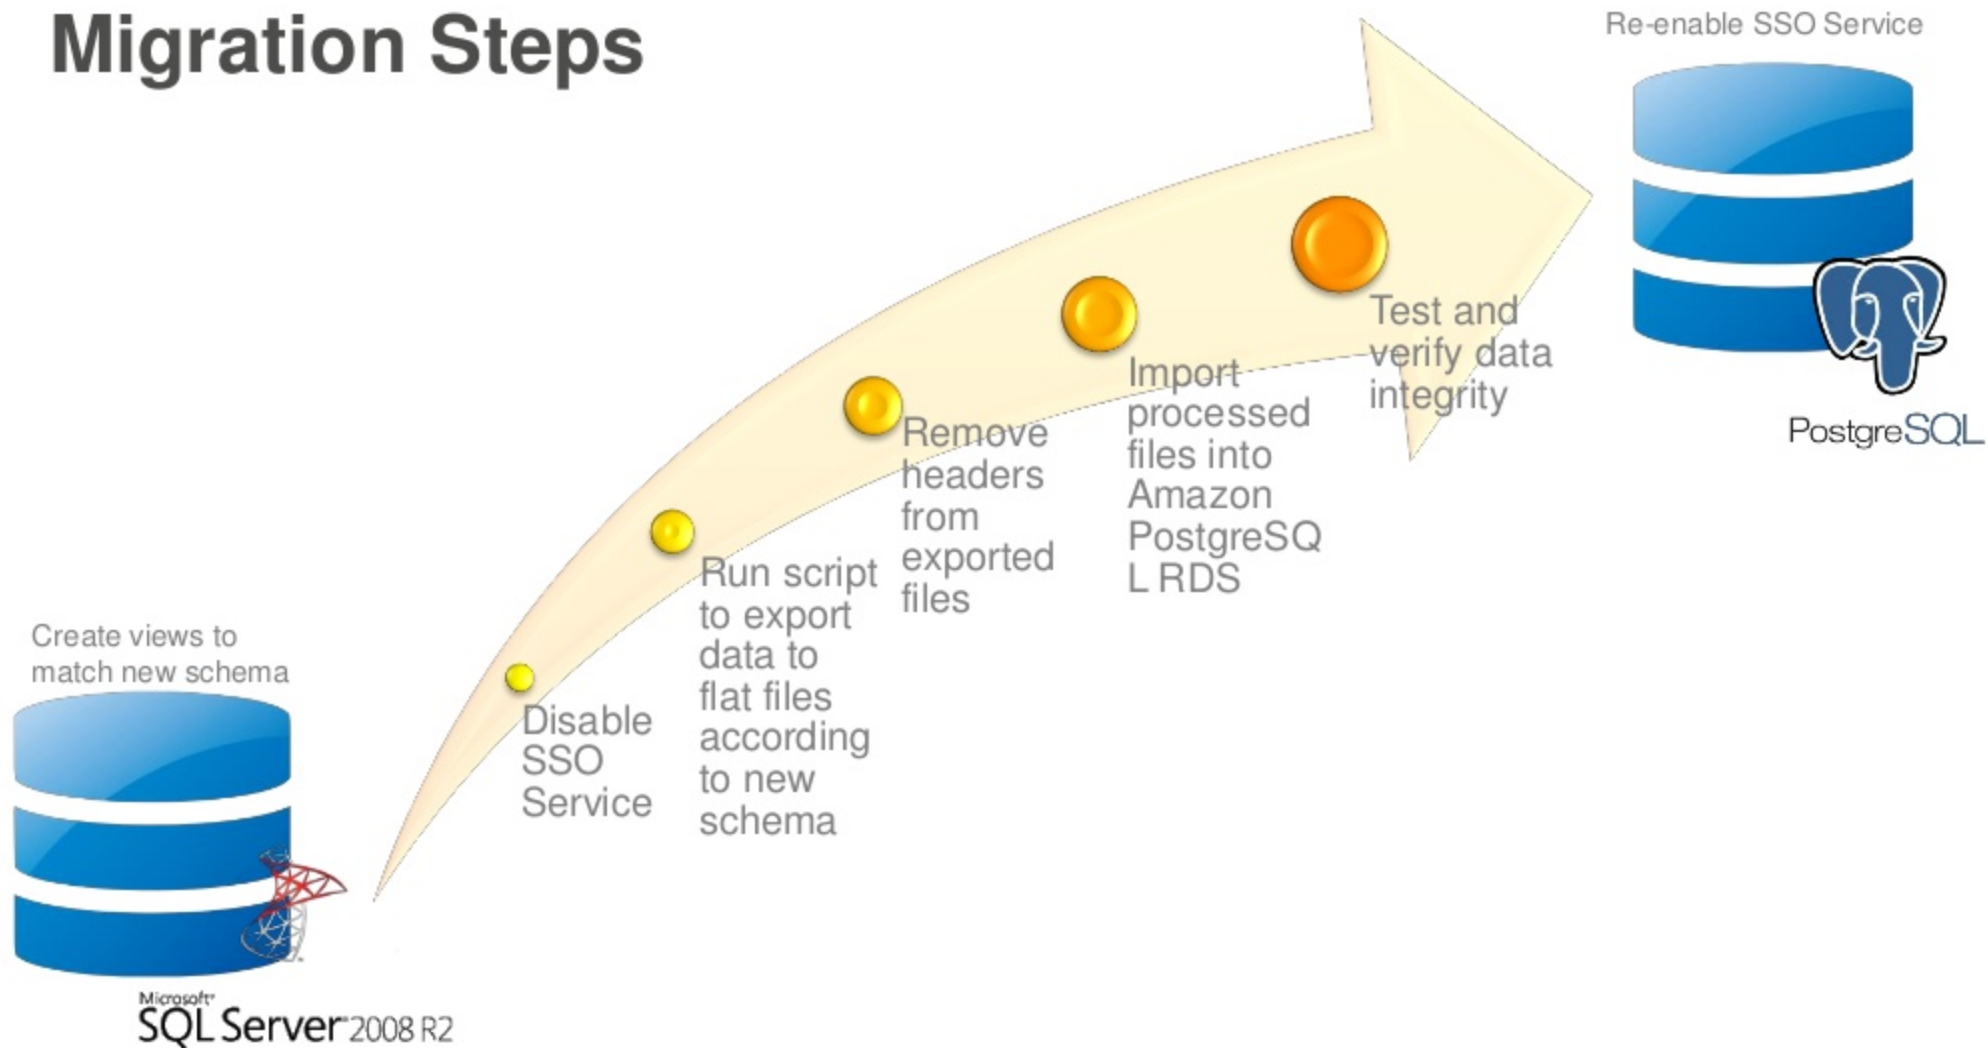
### Challenge: Testing

- Services were disabled pre-DNS migration for the SSO System
- Whitelist for specific users were created for testing and verification of data

~1.8 million

SSO users migrated

# Migration Steps




# The Solution and Benefits

## Uberfusion's Delivery for SSO Migration to AWS

### Challenge: Improve System Performance

- Use of Redis cache instead of DB for sessions and time-based tickets
- Frequently accessed data is cached to reduce repeated DB reads
- Compression and Browser Caching enabled on Amazon CloudFront, improving page load times by 50%
- TPS improved to 179 TPS for login on 4 auto-scaled servers vs 40 TPS during on-premise



### Challenge: Maintaining Connectivity

- Citrix CloudBridge is leveraged to ensure connectivity between AWS Cloud and on-premise DC



### Challenge: DB Migration from MSSQL to Postgres

- Data exported into flat files with same schema as new Postgres DB
- Membership libraries modified to replicate same encryption in PostgreSQL



> 500,000 sessions in June 2016)

### Challenge: End-to-End Monitoring

- Silicone and service metrics monitored with Amazon Cloudwatch
- Runscope used for API uptime and response time monitoring
- Extensive logs input into Amazon Elasticsearch via Amazon Firehose, visualised via Kibana



### Challenge: Scaling and Simplified Provisioning

- Separate applications on Amazon EC2
- Individual scaling groups according to load



## The Benefits



### Fast Provisioning of Infrastructure

- Scales independently and only pay for what you use.
- <1 day to duplicate environments vs >5 days for on-premise.



### Improved Development Efficiency

By leveraging AWS PaaS components, man-days can be focused on other parts of the system.



### Improved Availability and System Performance

- 50% improvement in page load times.



### Increased Visibility Allows for Proactive Teams

- Detected McAfee FIM tool causing CPU spikes.
- Detected iOS production application hitting SSO Staging.



### Sustain High Loads

- Anticipating Euro 2016 and Olympics 2016 load.
- 179 TPS in AWS vs 40 TPS on-premise.

# **Understanding Database Migration Projects**

# It's Not Just the Database!

Application code

Scripts and integration points

Backup and recovery

General management procedures

- Monitoring
- Notifications
- Auditing
- Tuning and Diagnostics

People!

# Project Planning

Project Assessment

Database Assessment

Application Assessment

Skill Sets Required

Hiring and Developing Talent

Choosing the right POCs

Project Execution

# Migration Effort



# Database migration

## Lift & shift

- Like to like

## Write from scratch

- Switch engine

## Modernize

- Fan out or consolidate
- Re-architect



# Database migration – multi phase process

Phase	Description	Automation	Effort (%)
1	Assessment	SCT	2
2	Database Schema Conversion	SCT/DMS	14
3	Application Conversion/Remediation	SCT	25
4	Scripts Conversion	SCT	7
5	Integration with 3 <sup>rd</sup> party applications		3
6	Data Migration	DMS	4
7	Functional testing of the entire system		29
8	Performance tuning	SCT	2
9	Integration and deployment		7
10	Training and knowledge		2
11	Documentation and version control		2
12	Post production support		3

# Database Migration Process

## STEP 1:



## STEP 2:



AWS

S U M M I T

