

Netflix Data Pipeline with Kafka

Allen Wang & Steven Wu

Agenda

- Introduction
- Evolution of Netflix data pipeline
- How do we use Kafka

What is Netflix?

Netflix is a logging company



that occasionally streams video



Numbers

- 400 billion events per day
- 8 million events & 17 GB per second during peak
- hundreds of event types

Agenda


- Introduction
- Evolution of Netflix data pipeline
- How do we use Kafka

Mission of Data Pipeline

Publish, Collect, Aggregate, Move Data @
Cloud Scale

In the old days ...

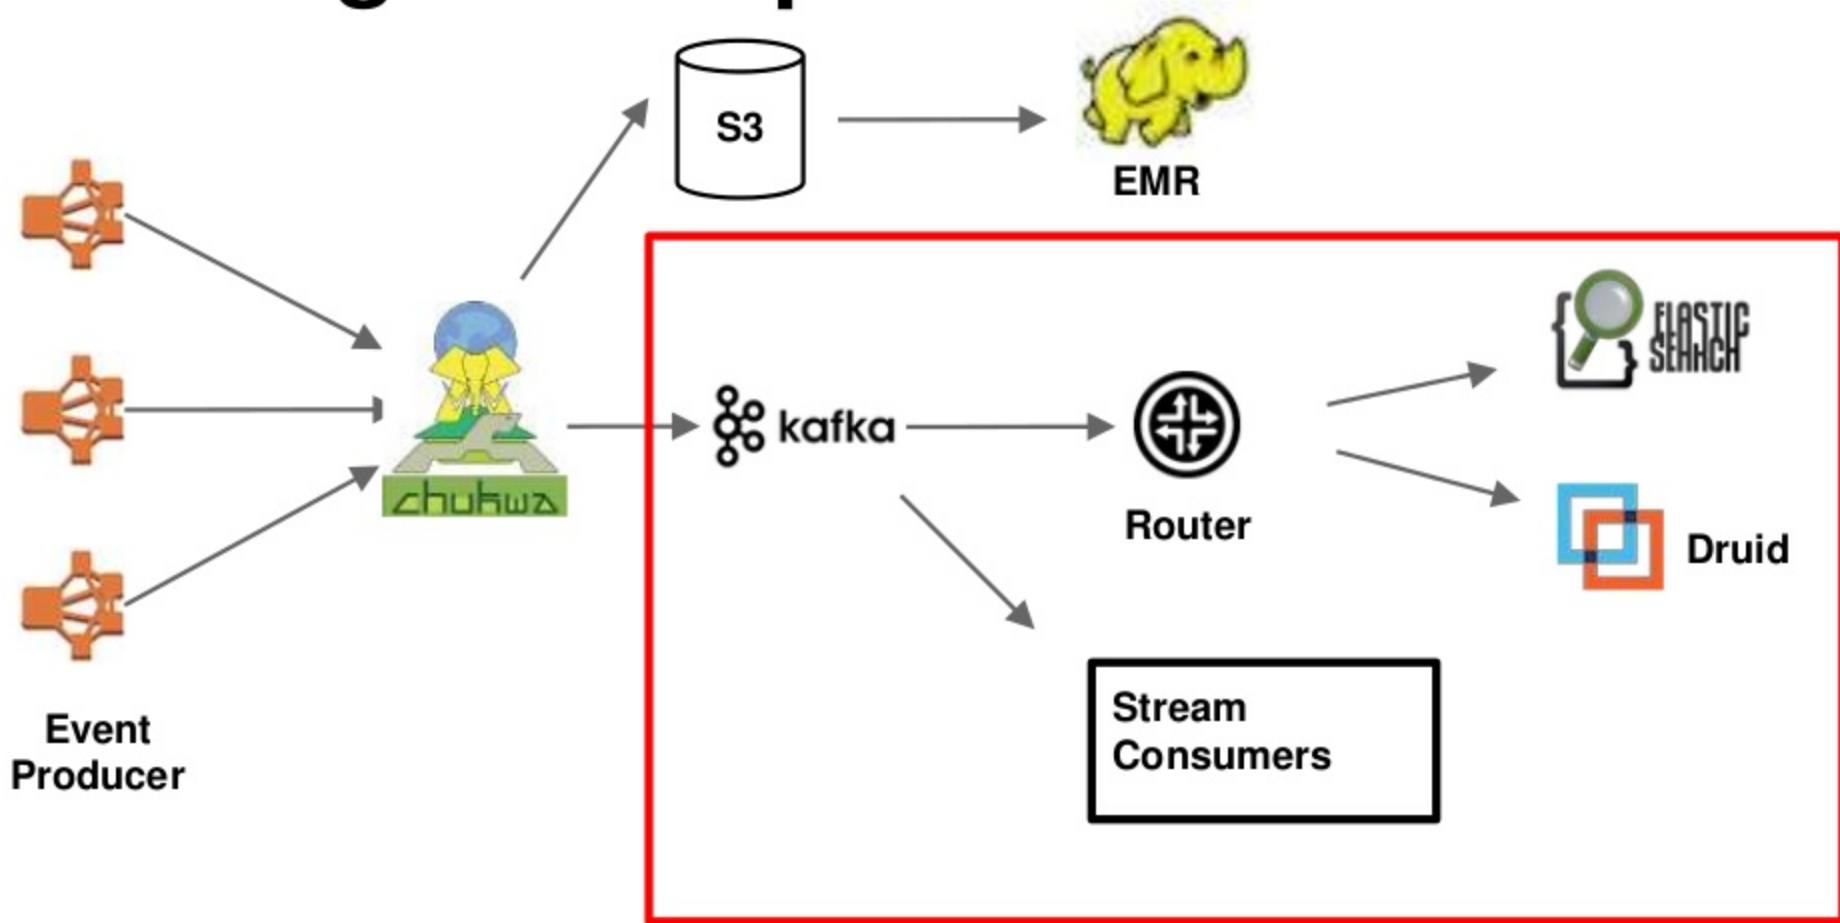




Nowadays ...



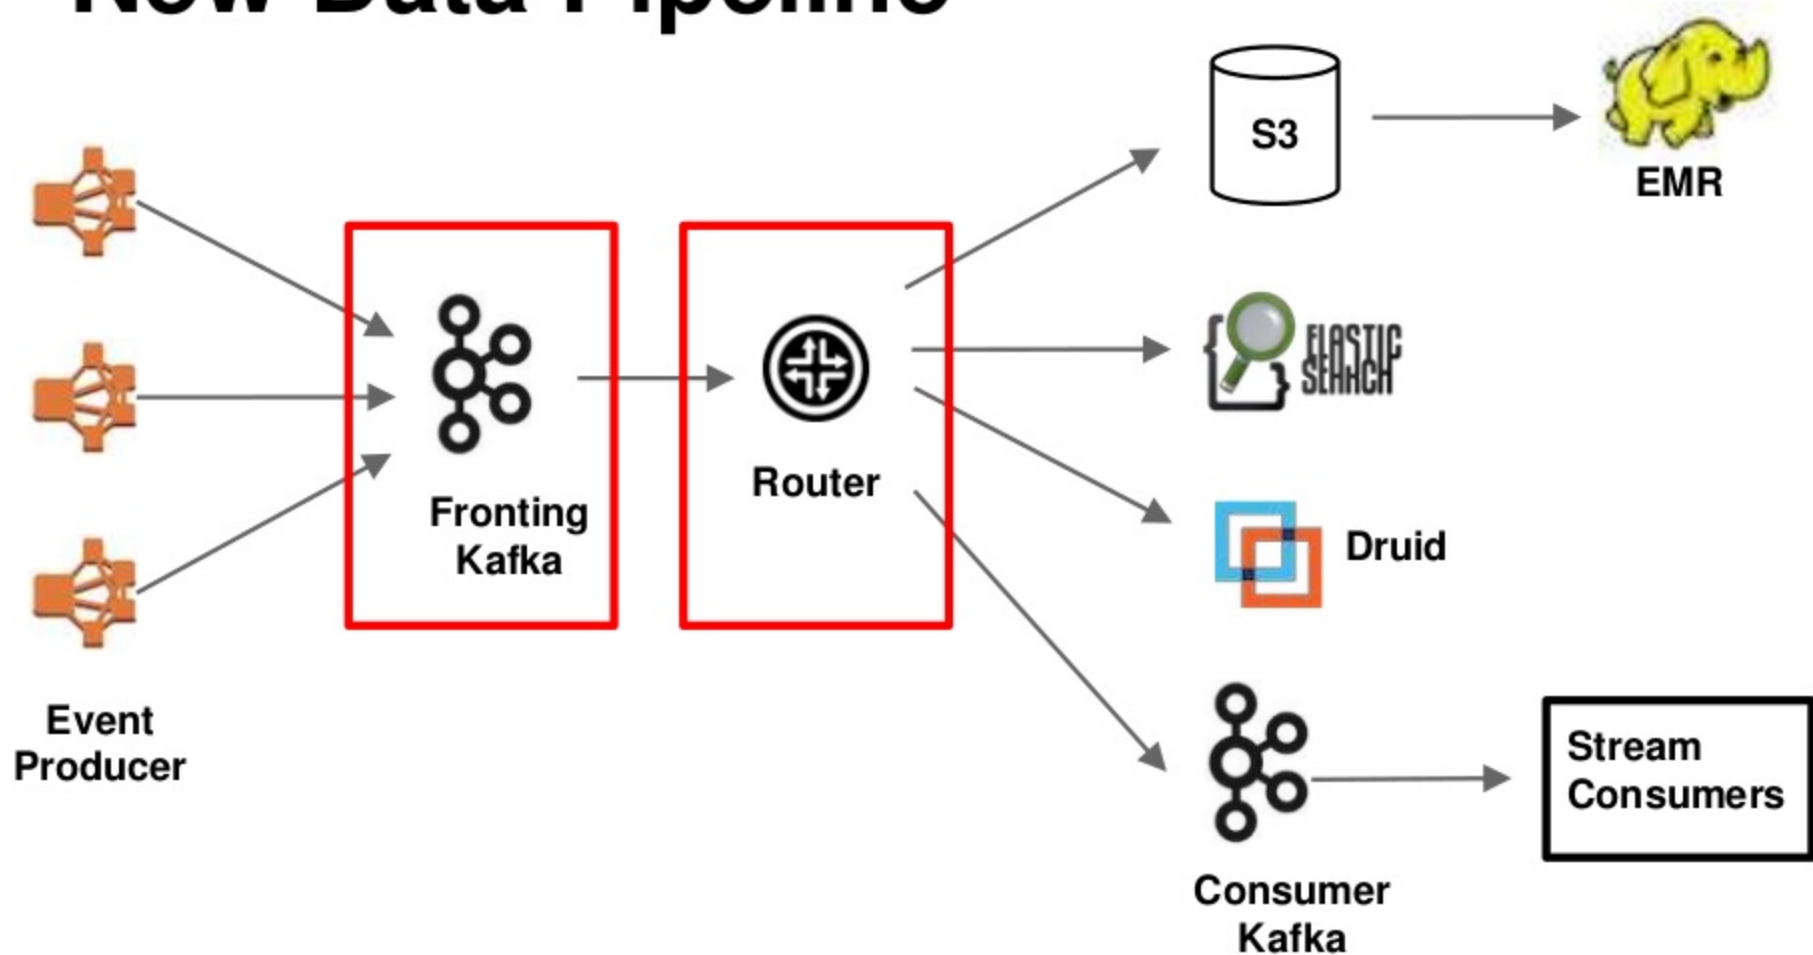
Existing Data Pipeline



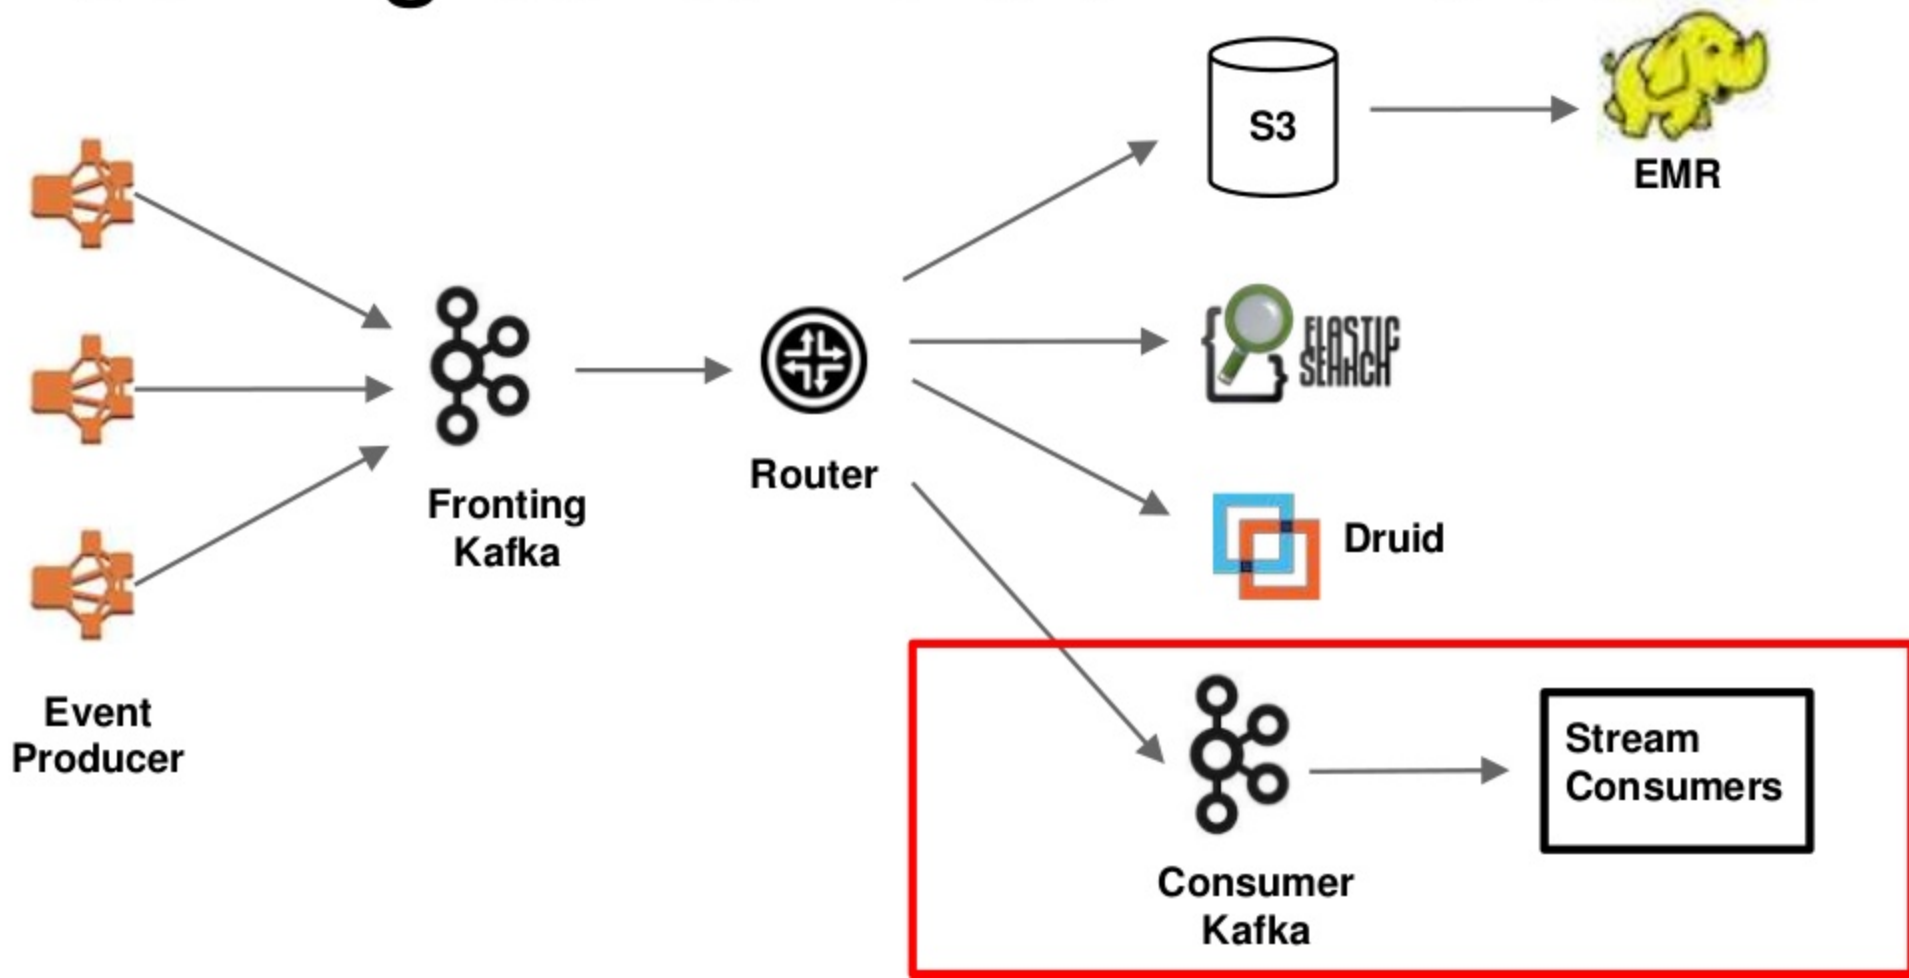
In to the Future ...



New Data Pipeline



Serving Consumers off Diff Clusters



Split Fronting Kafka Clusters

- Low-priority (error log, request trace, etc.)
 - 2 copies, 1-2 hour retention
- Medium-priority (majority)
 - 2 copies, 4 hour retention
- High-priority (streaming activities etc.)
 - 3 copies, 12-24 hour retention

Producer Resilience

- Kafka outage should never disrupt existing instances from serving business purpose
- Kafka outage should never prevent new instances from starting up
- After kafka cluster restored, event producing should resume automatically

Fail but Never Block

- `block.on.buffer.full=false`
- handle potential blocking of first meta data request
- Periodical check whether `KafkaProducer` was opened successfully

Agenda

- Introduction
- Evolution of Netflix data pipeline
- How do we use Kafka

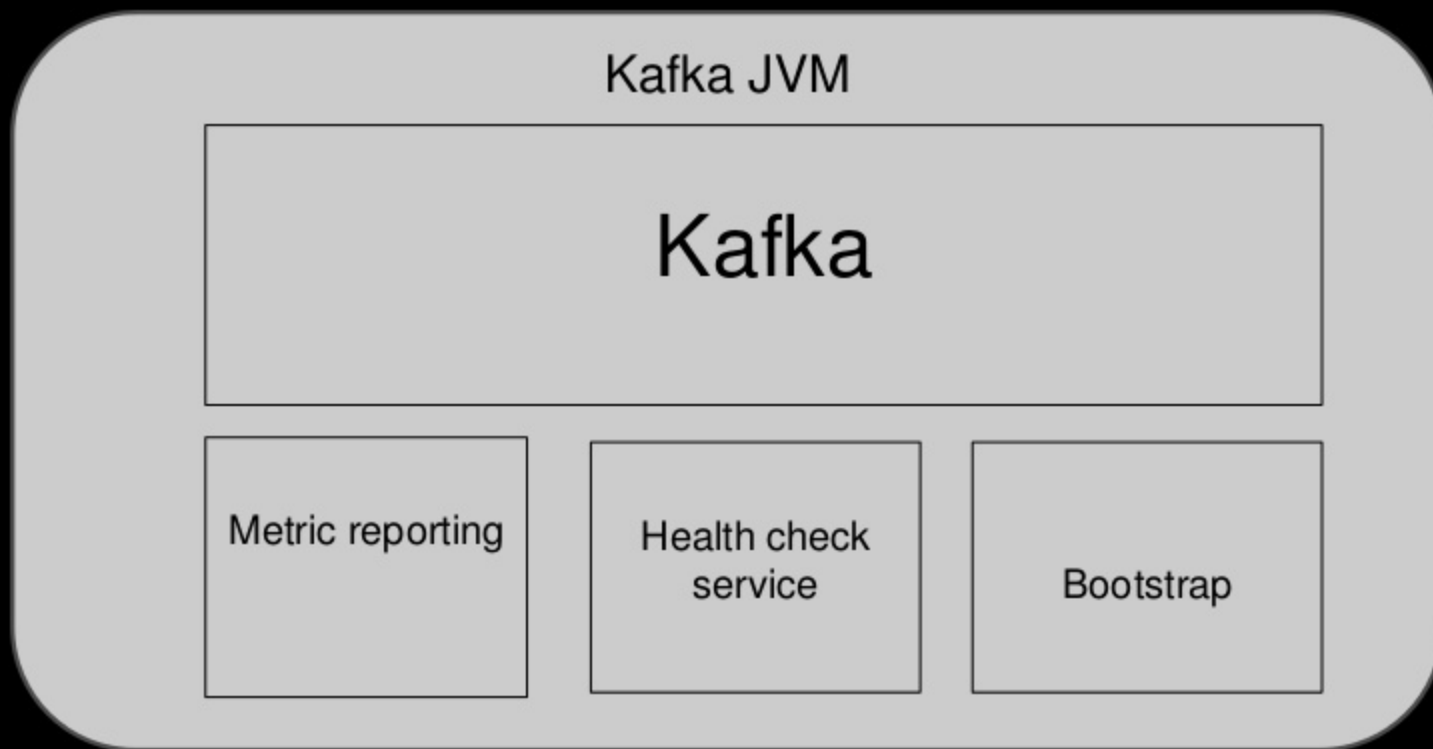
What Does It Take to Run In Cloud

- Support elasticity
- Respond to scaling events
- Resilience to failures
 - Favors architecture without single point of failure
 - Retries, smart routing, fallback ...

Kafka in AWS - How do we make it happen

- Inside our Kafka JVM
- Services supporting Kafka
- Challenges/Solutions
- Our roadmap

Netflix Kafka Container

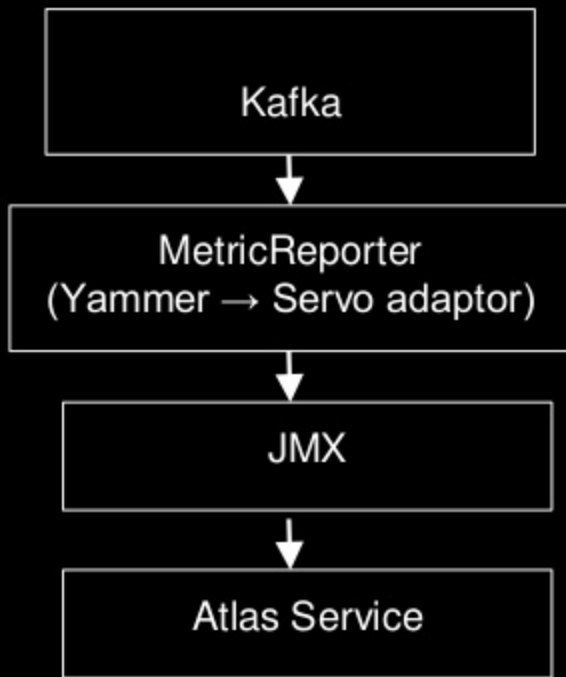


Bootstrap

- Broker ID assignment
 - Instances obtain sequential numeric IDs using Curator's locks recipe persisted in ZK
 - Cleans up entry for terminated instances and reuse its ID
 - Same ID upon restart
- Bootstrap Kafka properties from Archaius
 - Files
 - System properties/Environment variables
 - Persisted properties service
- Service registration
 - Register with Eureka for internal service discovery
 - Register with AWS Route53 DNS service

Metric Reporting

- We use Servo and Atlas from NetflixOSS



Kafka Atlas Dashboard



Health check service

- Use Curator to periodically read ZooKeeper data to find signs of unhealthiness
- Export metrics to Servo/Atlas
- Expose the service via embedded Jetty

Kafka in AWS - How do we make it happen

- Inside our Kafka JVM
- Services supporting Kafka
- Challenges/Solutions
- Our roadmap

ZooKeeper

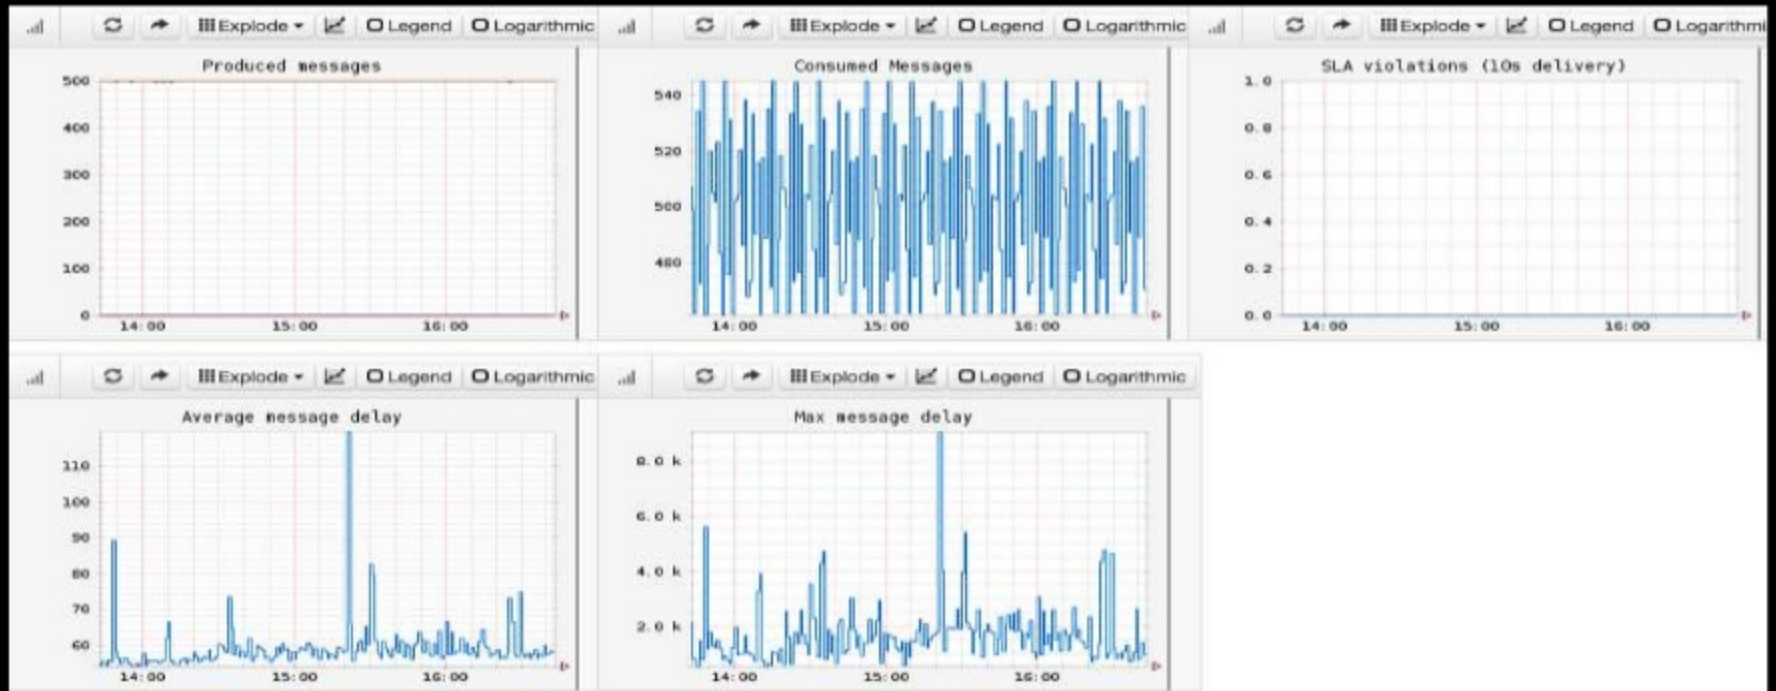
- Dedicated 5 node cluster for our data pipeline services
- EIP based
- SSD instance

Auditor

- Highly configurable producers and consumers with their own set of topics and metadata in messages
- Built as a service deployable on single or multiple instances
- Runs as producer, consumer or both
- Supports replay of preconfigured set of messages

Auditor

- Broker monitoring (Heartbeating)



Auditor

- Broker performance testing
 - Produce tens of thousands messages per second on single instance
 - As consumers to test consumer impact

Kafka admin UI

- Still searching ...
- Currently trying out KafkaManager



Clusters / kafkabroker-default / Topics / non_member_event

← non_member_event

Topic Summary

| | |
|-----------------------------|-----|
| Replication | 2 |
| Number of Partitions | 18 |
| Total number of Brokers | 12 |
| Number of Brokers for Topic | 12 |
| Preferred Replicas % | 72 |
| Brokers Skewed % | 41 |
| Brokers Spread % | 100 |

Generate Partition Assignments

Reassign Partitions

Partitions by Broker

| Broker | # of Partitions | Partitions | Skewed? |
|--------|-----------------|-------------------|---------|
| 0 | 1 | {0} | false |
| 1 | 1 | {1} | false |
| 2 | 2 | {0,2} | false |
| 3 | 1 | {1} | false |
| 4 | 2 | {2,3} | false |
| 5 | 1 | {4} | false |
| 6 | 6 | {5,9,3,16,11,15} | true |
| 7 | 6 | {10,6,17,12,16,4} | true |
| 8 | 5 | {5,13,17,7,11} | true |
| 9 | 4 | {14,6,12,8} | true |
| 10 | 3 | {9,13,7} | false |
| 11 | 4 | {10,14,8,15} | true |

Partition Information

| Partition | Leader | Replicas | In Sync Replicas | Preferred Leader? | Under Replicated? |
|-----------|--------|----------|------------------|-------------------|-------------------|
| 0 | 0 | {0,2} | {0,2} | true | false |
| 1 | 1 | {1,3} | {1,3} | true | false |

Kafka in AWS - How do we make it happen

- Inside our Kafka JVM
- Services supporting Kafka
- Challenges/Solutions
- Our roadmap

Challenges

- ZooKeeper client issues
- Cluster scaling
- Producer/consumer/broker tuning

ZooKeeper Client

- Challenges

- Broker/consumer cannot survive ZooKeeper cluster rolling push due to caching of private IP
- Temporary DNS lookup failure at new session initialization kills future communication

ZooKeeper Client

- Solutions
 - Created our internal fork of Apache ZooKeeper client
 - Periodically refresh private IP resolution
 - Fallback to last good private IP resolution upon DNS lookup failure

Scaling

- Provisioned for peak traffic
 - ... and we have regional fail-over

Strategy #1 Add Partitions to New Brokers

- Caveat
 - Most of our topics do not use keyed messages
 - Number of partitions is still small
 - Require high level consumer

Strategy #1 Add Partitions to new brokers

- Challenges: existing admin tools does not support atomic adding partitions and assigning to new brokers

Strategy #1 Add Partitions to new brokers

- Solutions: created our own tool to do it in one ZooKeeper change and repeat for all or selected topics
- Reduced the time to scale up from a few hours to a few minutes

Strategy #2 Move Partitions

- Should work without precondition, but ...
- Huge increase of network I/O affecting incoming traffic
- A much longer process than adding partitions
- Sometimes confusing error messages
- Would work if pace of replication can be controlled

Scale down strategy

- There is none
- Look for more support to automatically move all partitions from a set of brokers to a different set

Client tuning

- Producer

- Batching is important to reduce CPU and network I/O on brokers
- Stick to one partition for a while when producing for non-keyed messages
- “linger.ms” works well with sticky partitioner

- Consumer

- With huge number of consumers, set proper `fetch.wait.max.ms` to reduce polling traffic on broker

Effect of batching

| partitioner | batched records per request | broker cpu util [1] |
|-----------------------------|------------------------------------|-------------------------------|
| random without lingering | 1.25 | 75% |
| sticky without lingering | 2.0 | 50% |
| sticky with 100ms lingering | 15 | 33% |

[1] 10 MB & 10K msgs / second per broker, 1KB per message

Broker tuning

- Use G1 collector
- Use large page cache and memory
- Increase max file descriptor if you have thousands of producers or consumers

Kafka in AWS - How do we make it happen

- Inside our Kafka JVM
- Services supporting Kafka
- Challenges/Solutions
- Our roadmap

Road map

- Work with Kafka community on rack/zone aware replica assignment
- Failure resilience testing
 - Chaos Monkey
 - Chaos Gorilla
- Contribute to open source
 - Kafka
 - Schlep -- our messaging library including SQS and Kafka support
 - Auditor



Thank you!

<http://netflix.github.io/>

<http://techblog.netflix.com/>

@NetflixOSS

@allenxwang

@stevenzwu

NETFLIX