

AWS
re:Invent

NET 402

Transit Gateway and Transit VPCs

Reference Architectures for Many VPCs

Nick Matthews
Principal Solutions Architect
AWS



@nickpowpow

AWS
re:Invent

© 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.



What to expect


How it works

Transit VPC

Transit Gateway

New


Build out a reference architecture:




Account
Strategy



Segmentation
Model




Shared
Services



Connectivity



Network
Services



Multi-Region
Options

Challenges with many VPCs


VPC management differences



Ease of creation




Access models

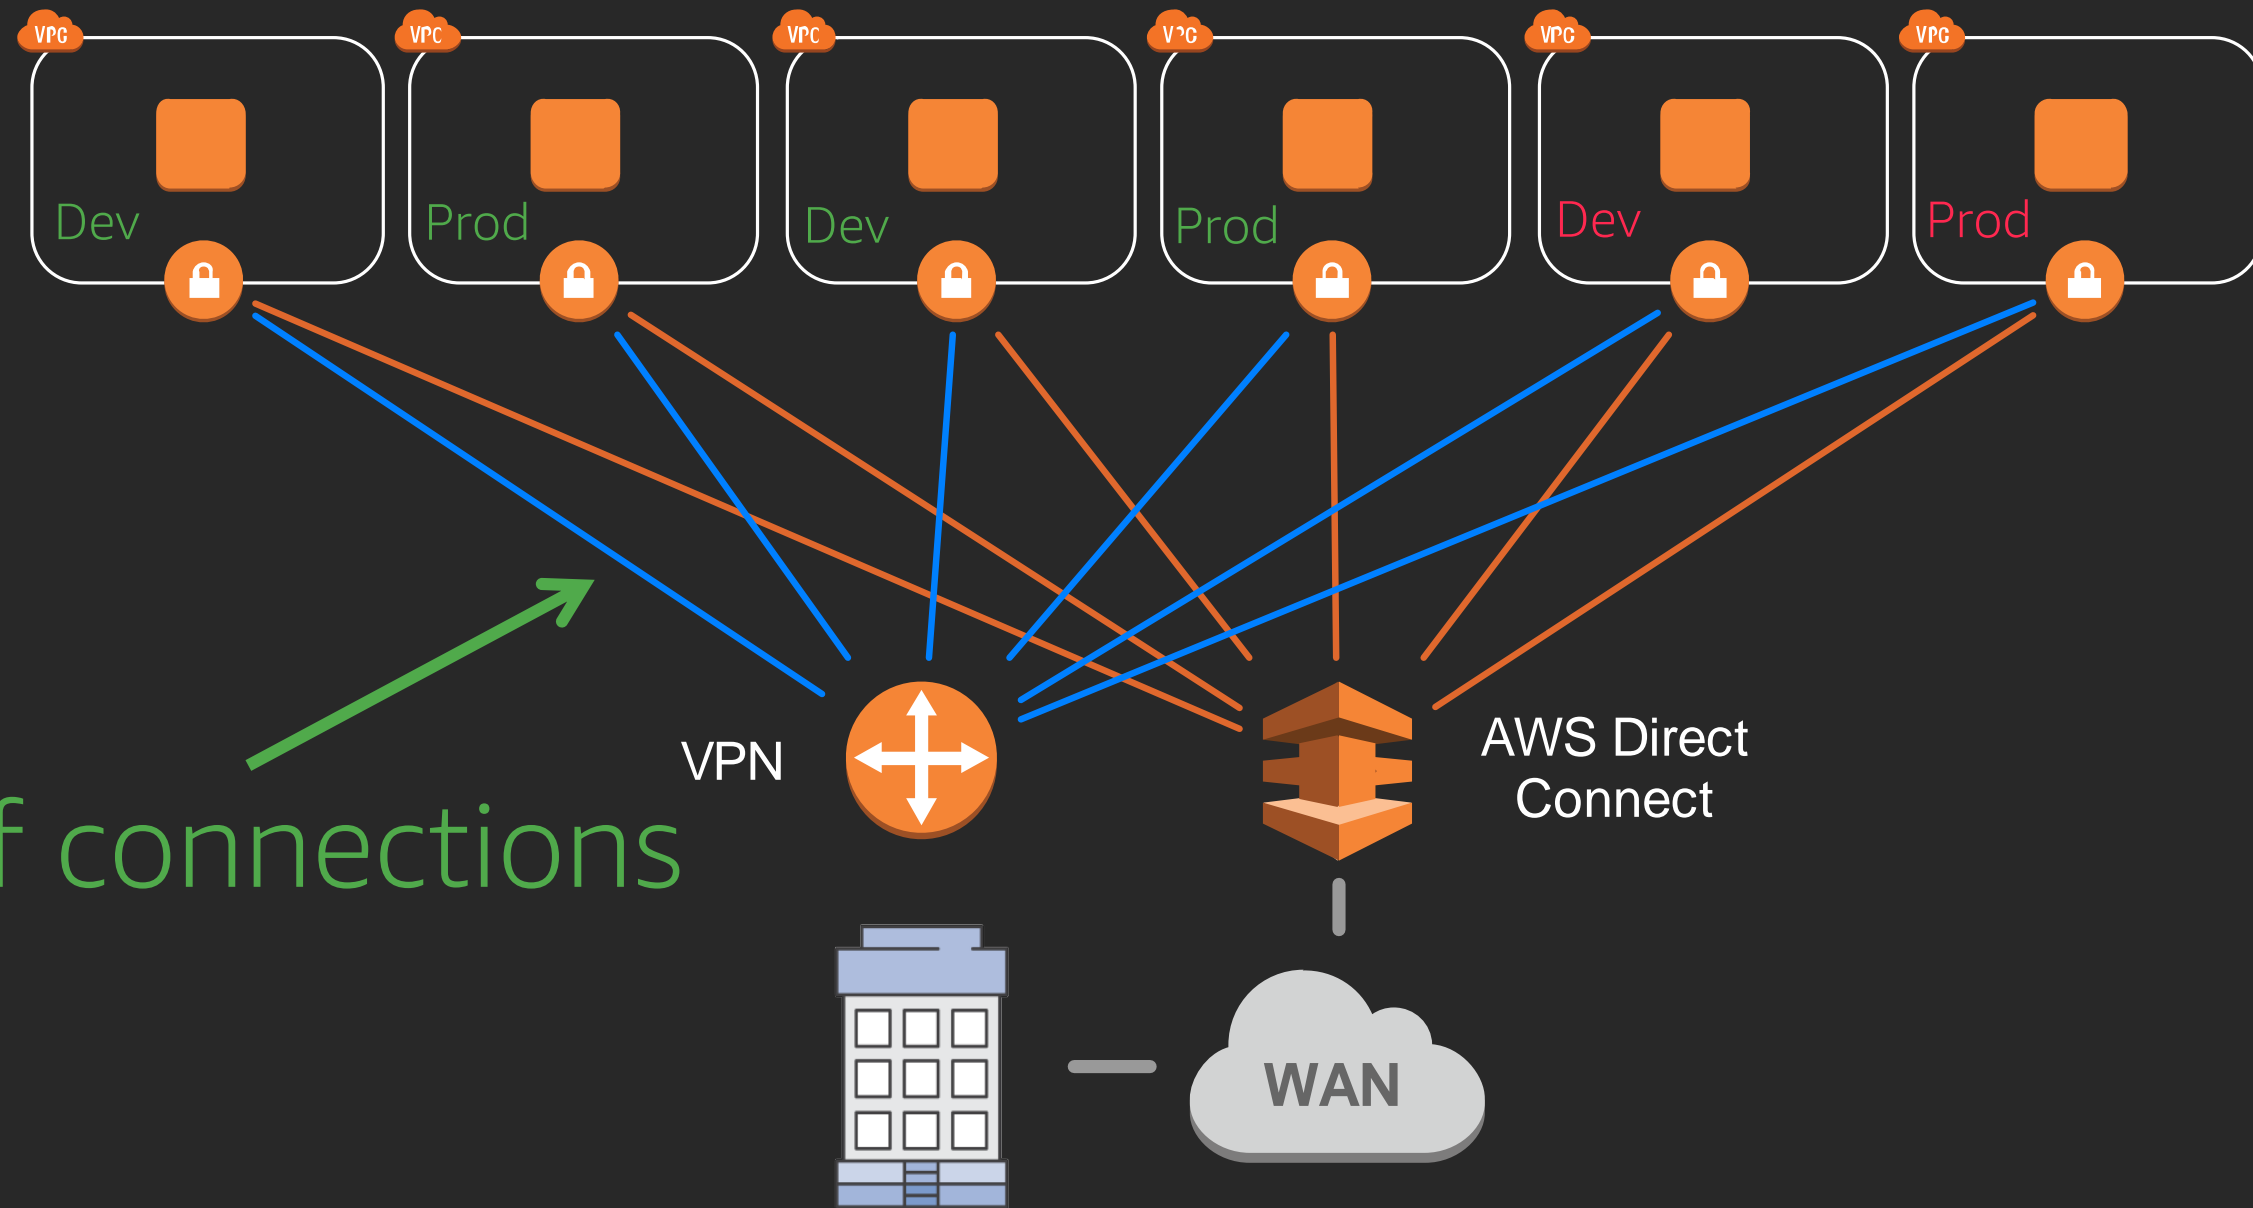


Diverse ownership

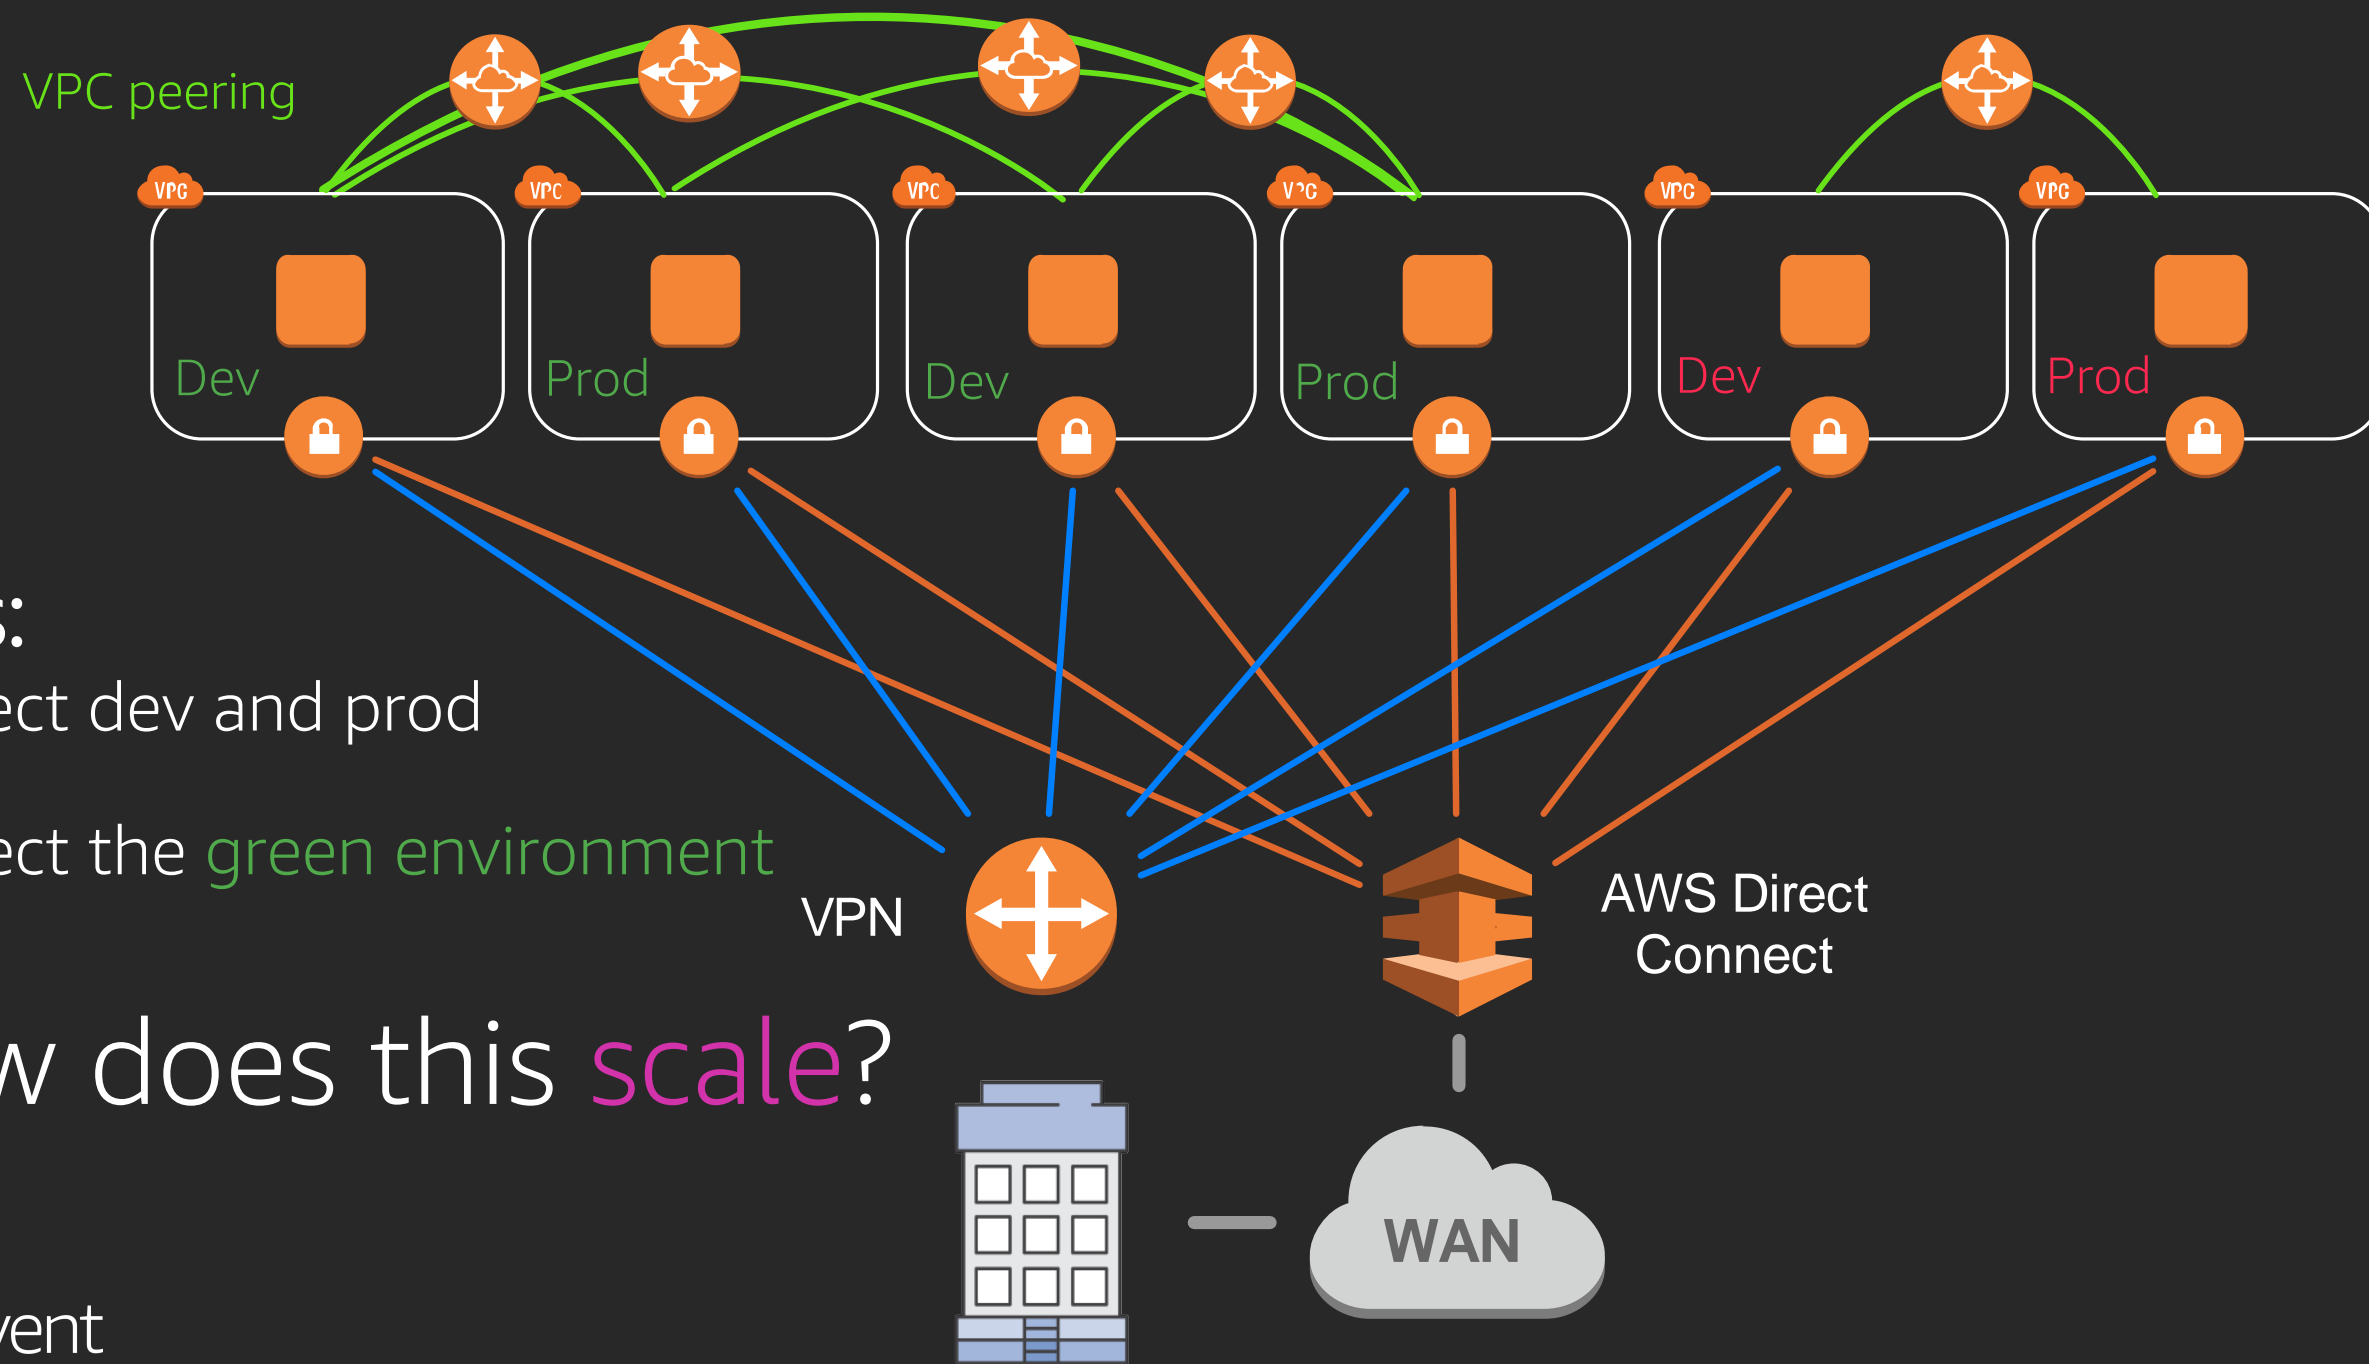
Our starting point



Challenge: Adding more VPCs



Challenge: Peering VPCs

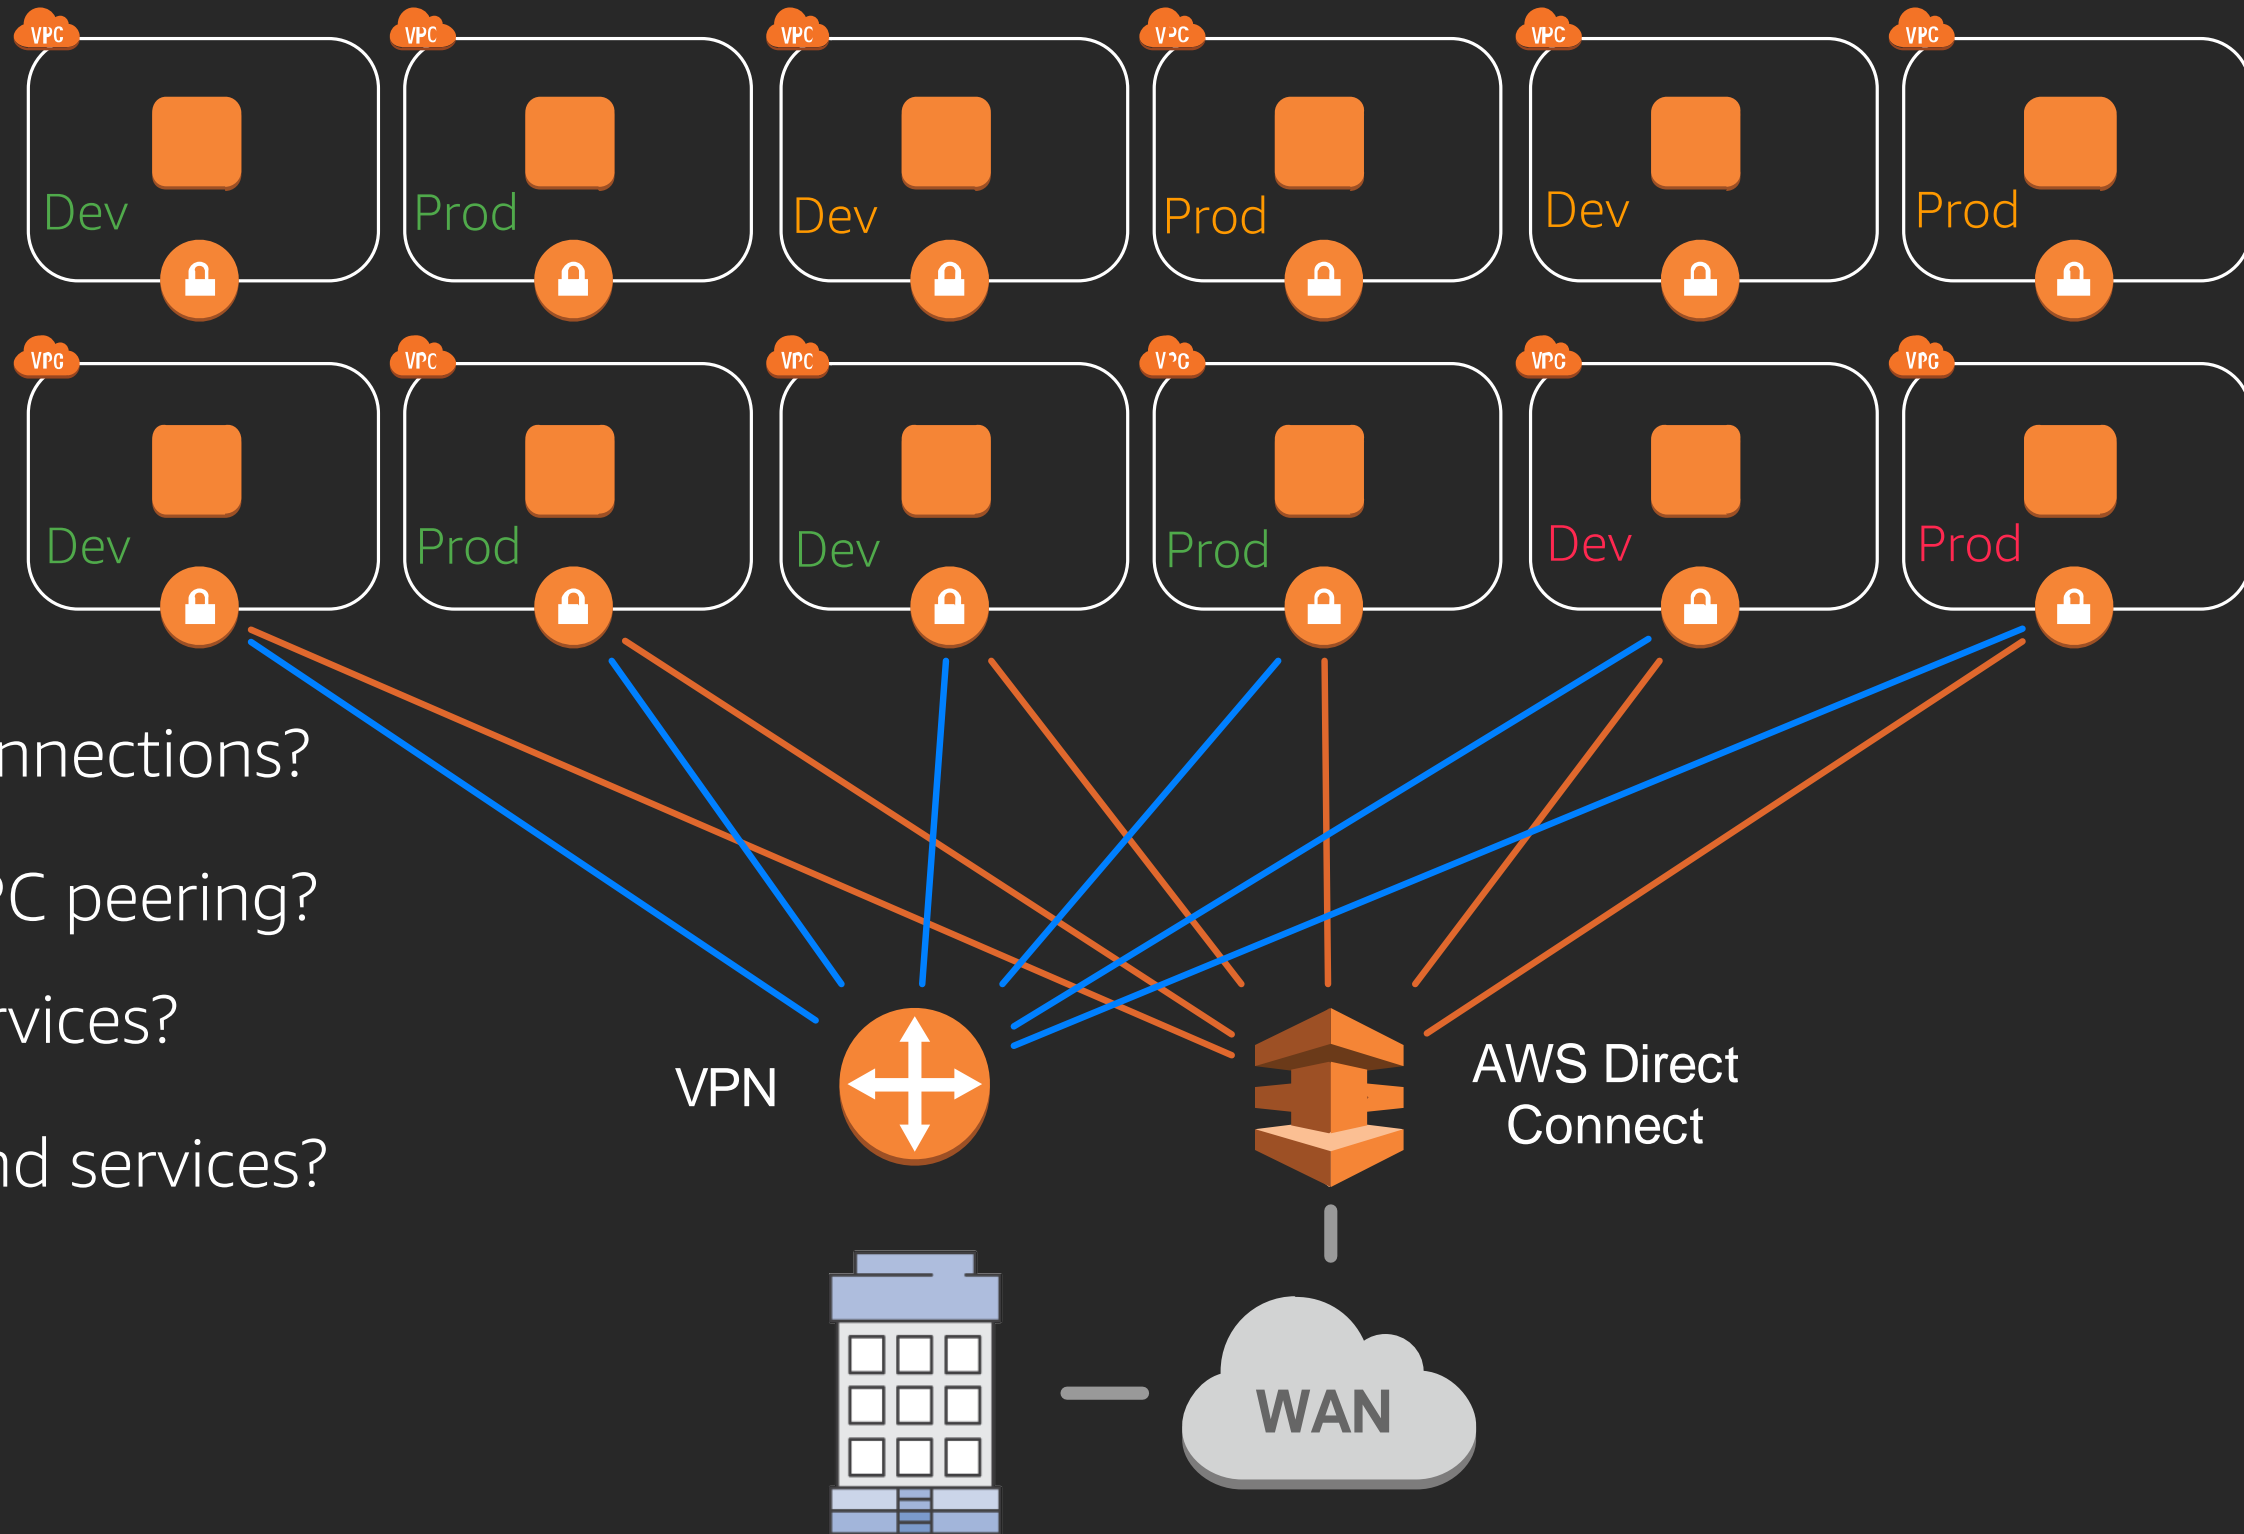


Let's:

Connect dev and prod

Connect the green environment

How does this **scale**?

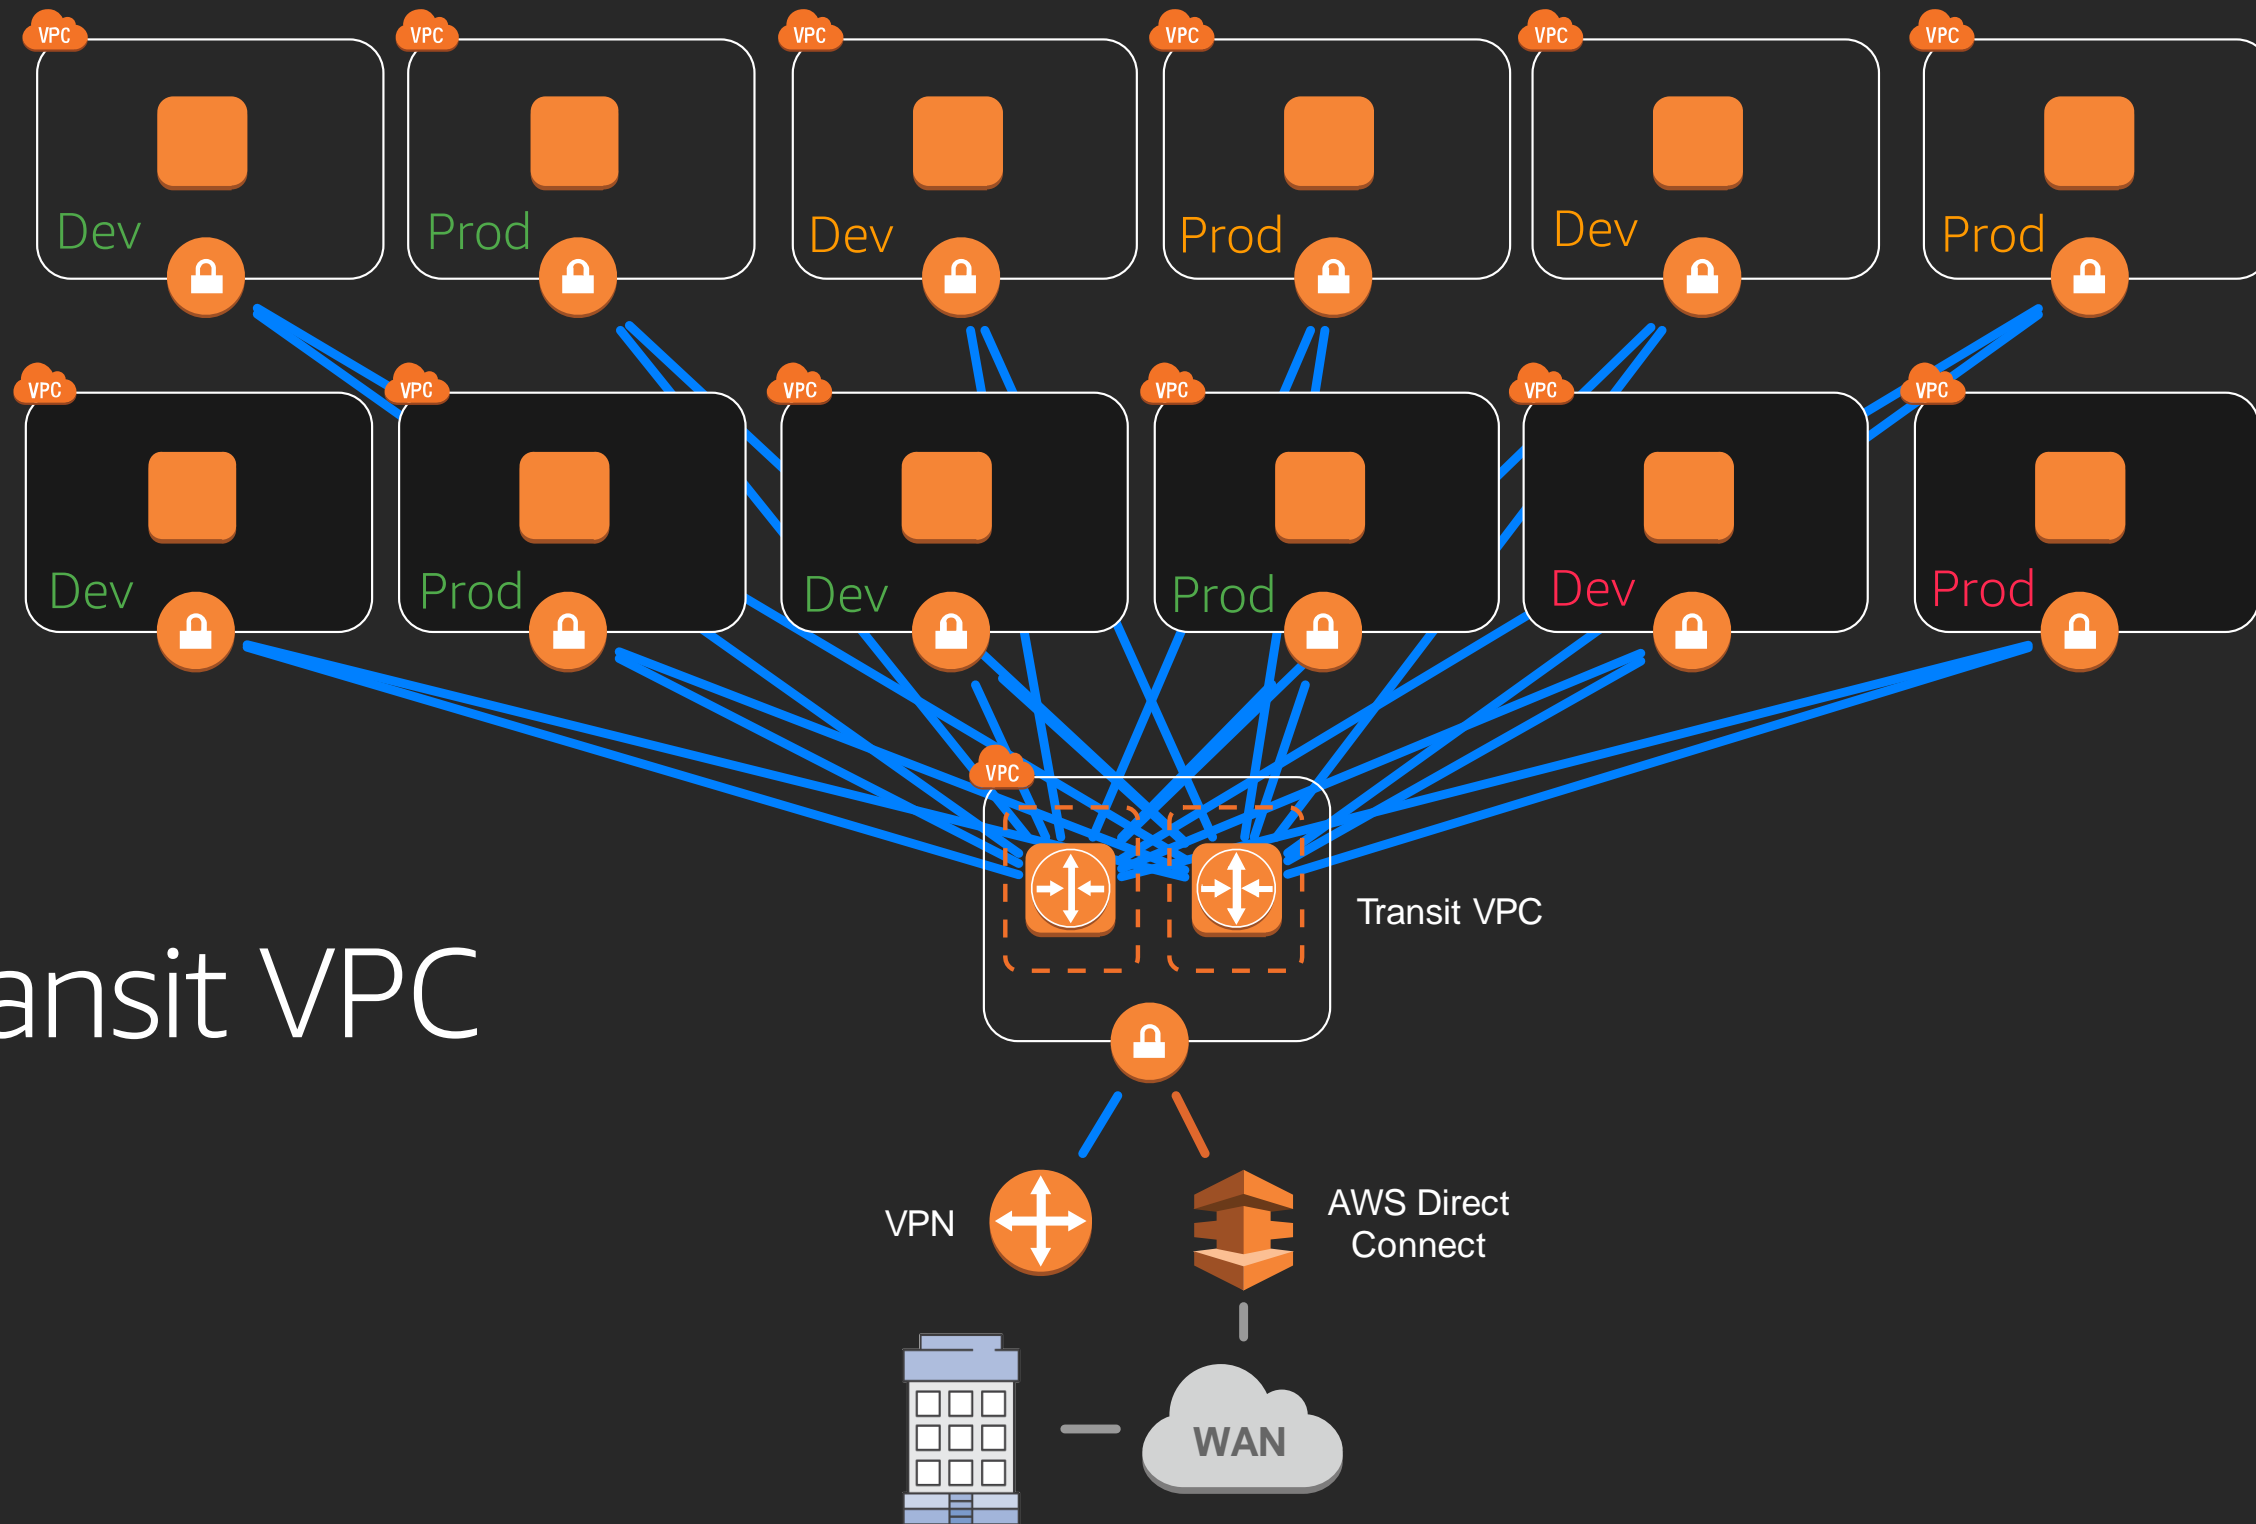


Scaling connections?

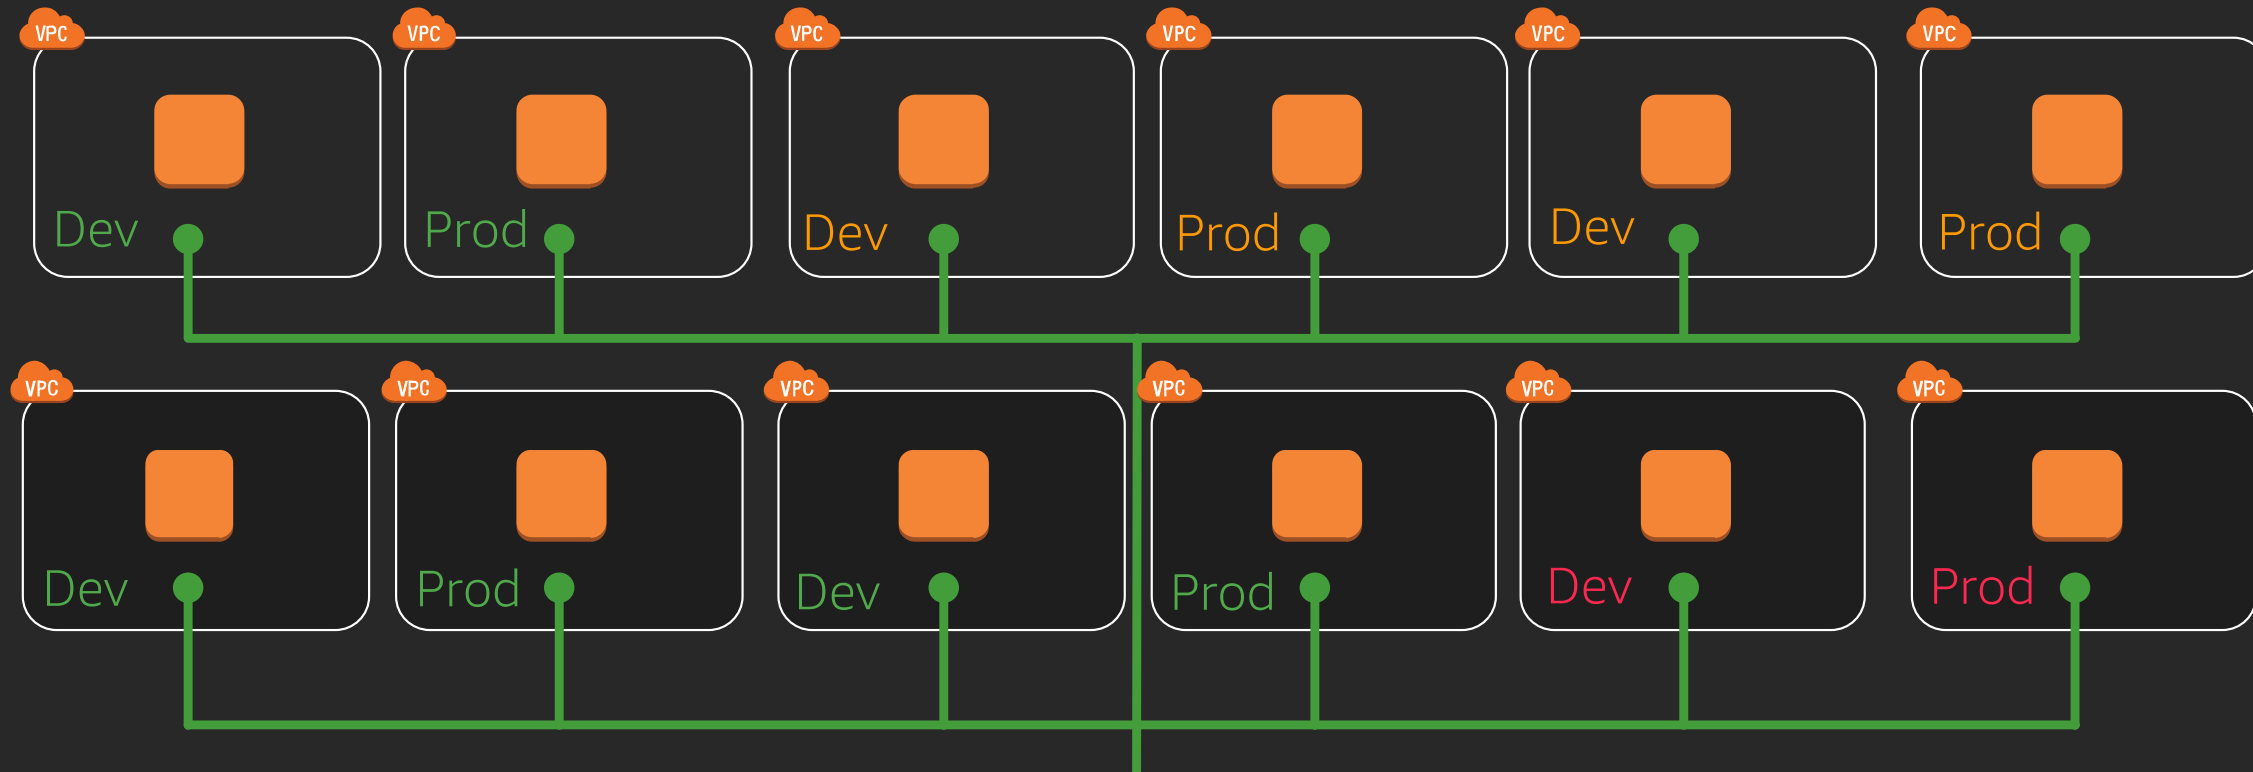
Scaling VPC peering?

Shared services?

Firewall and services?





Transit VPC



AWS Transit Gateway

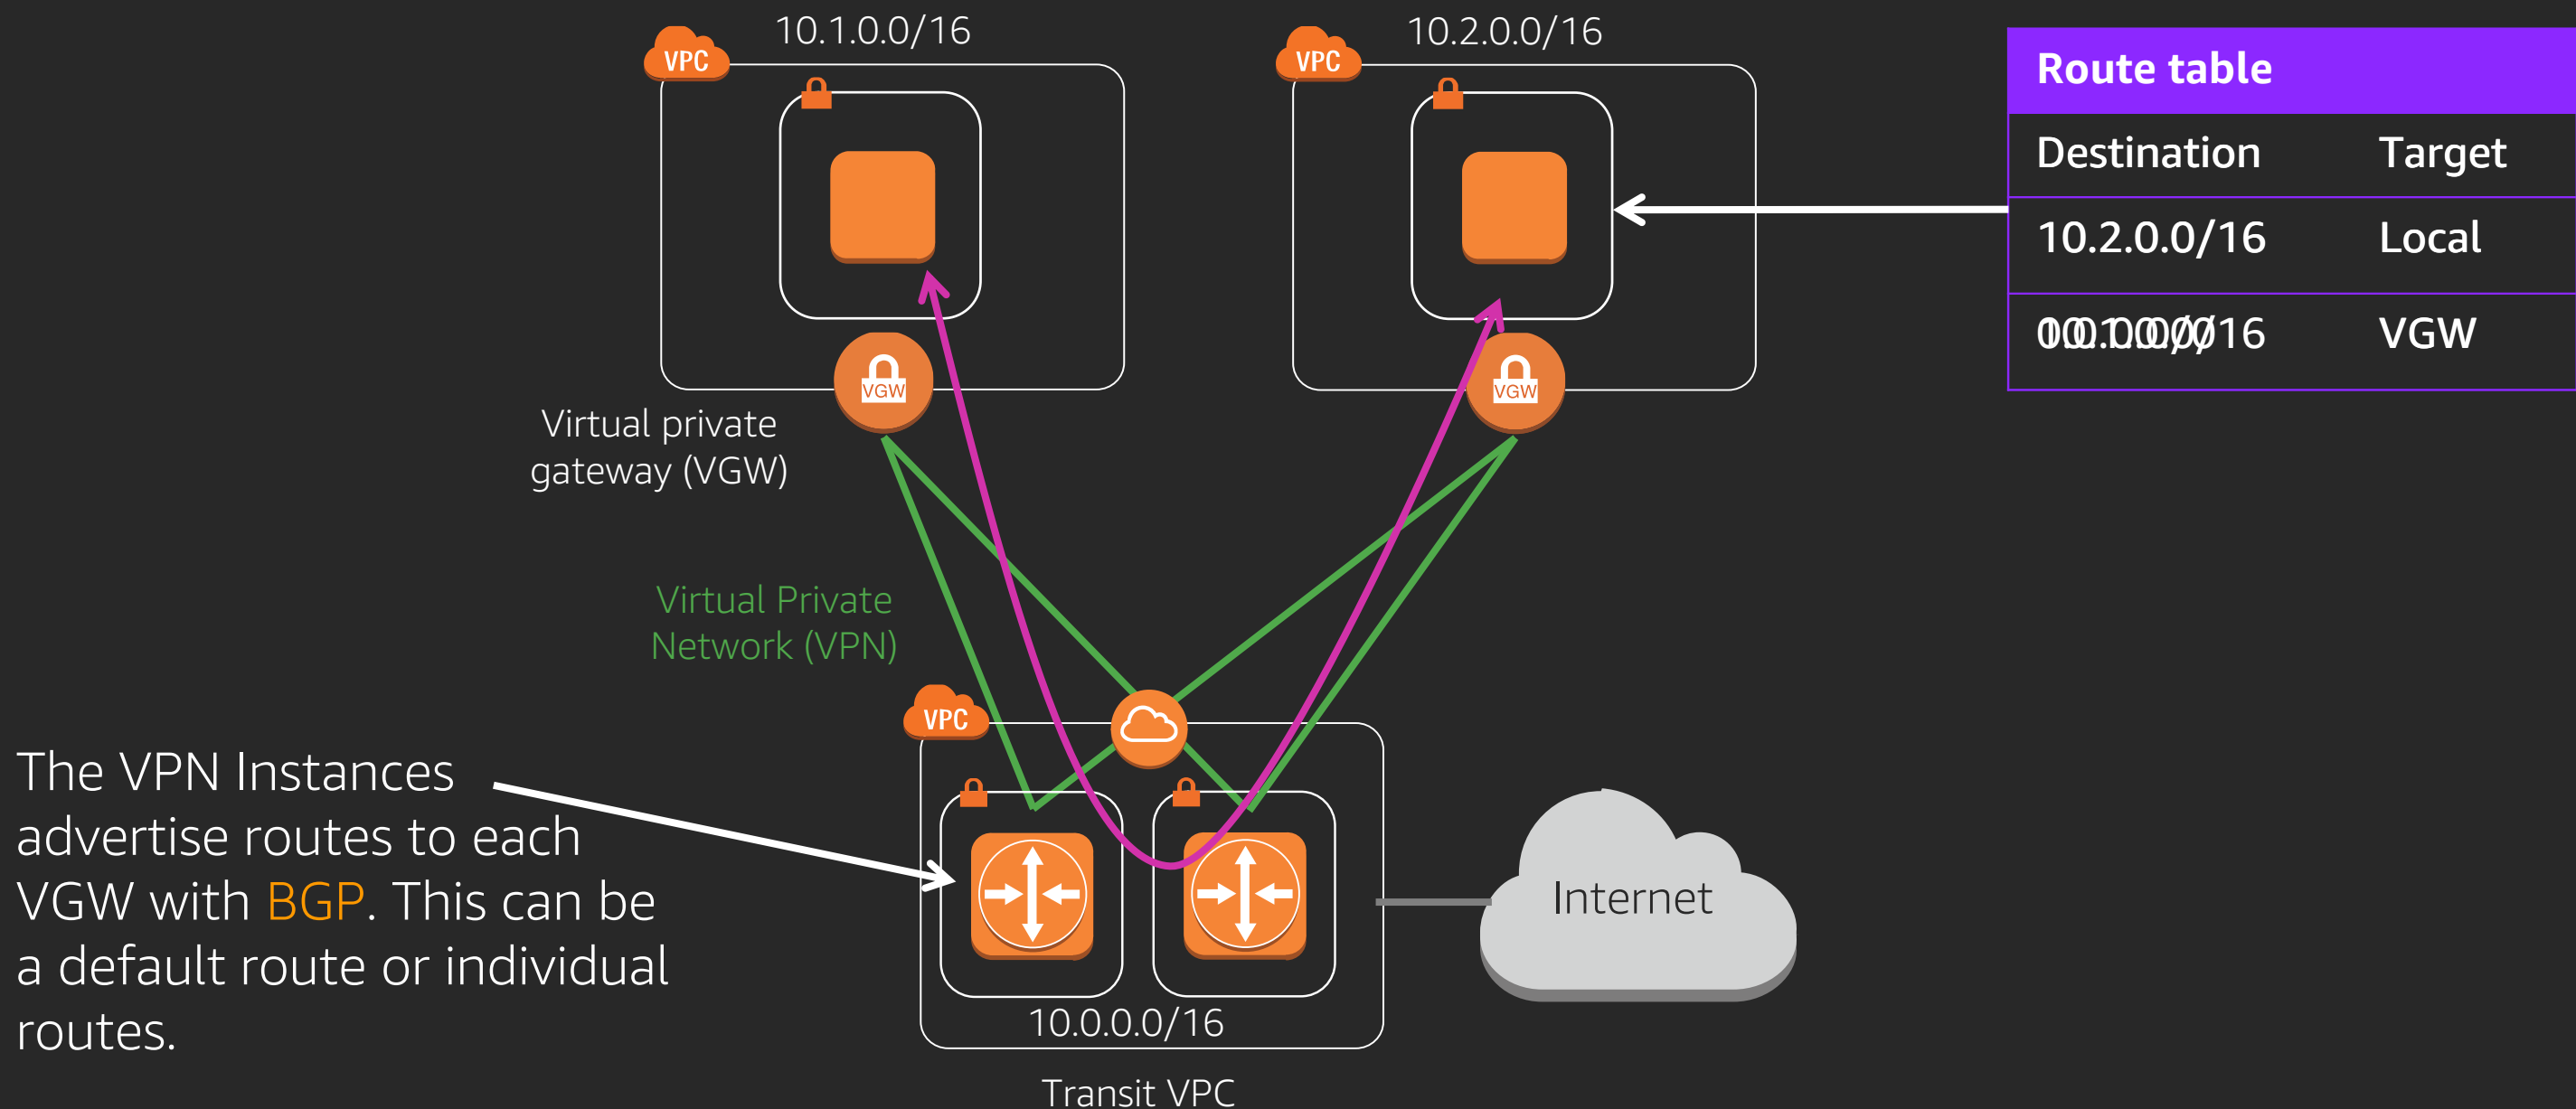
New



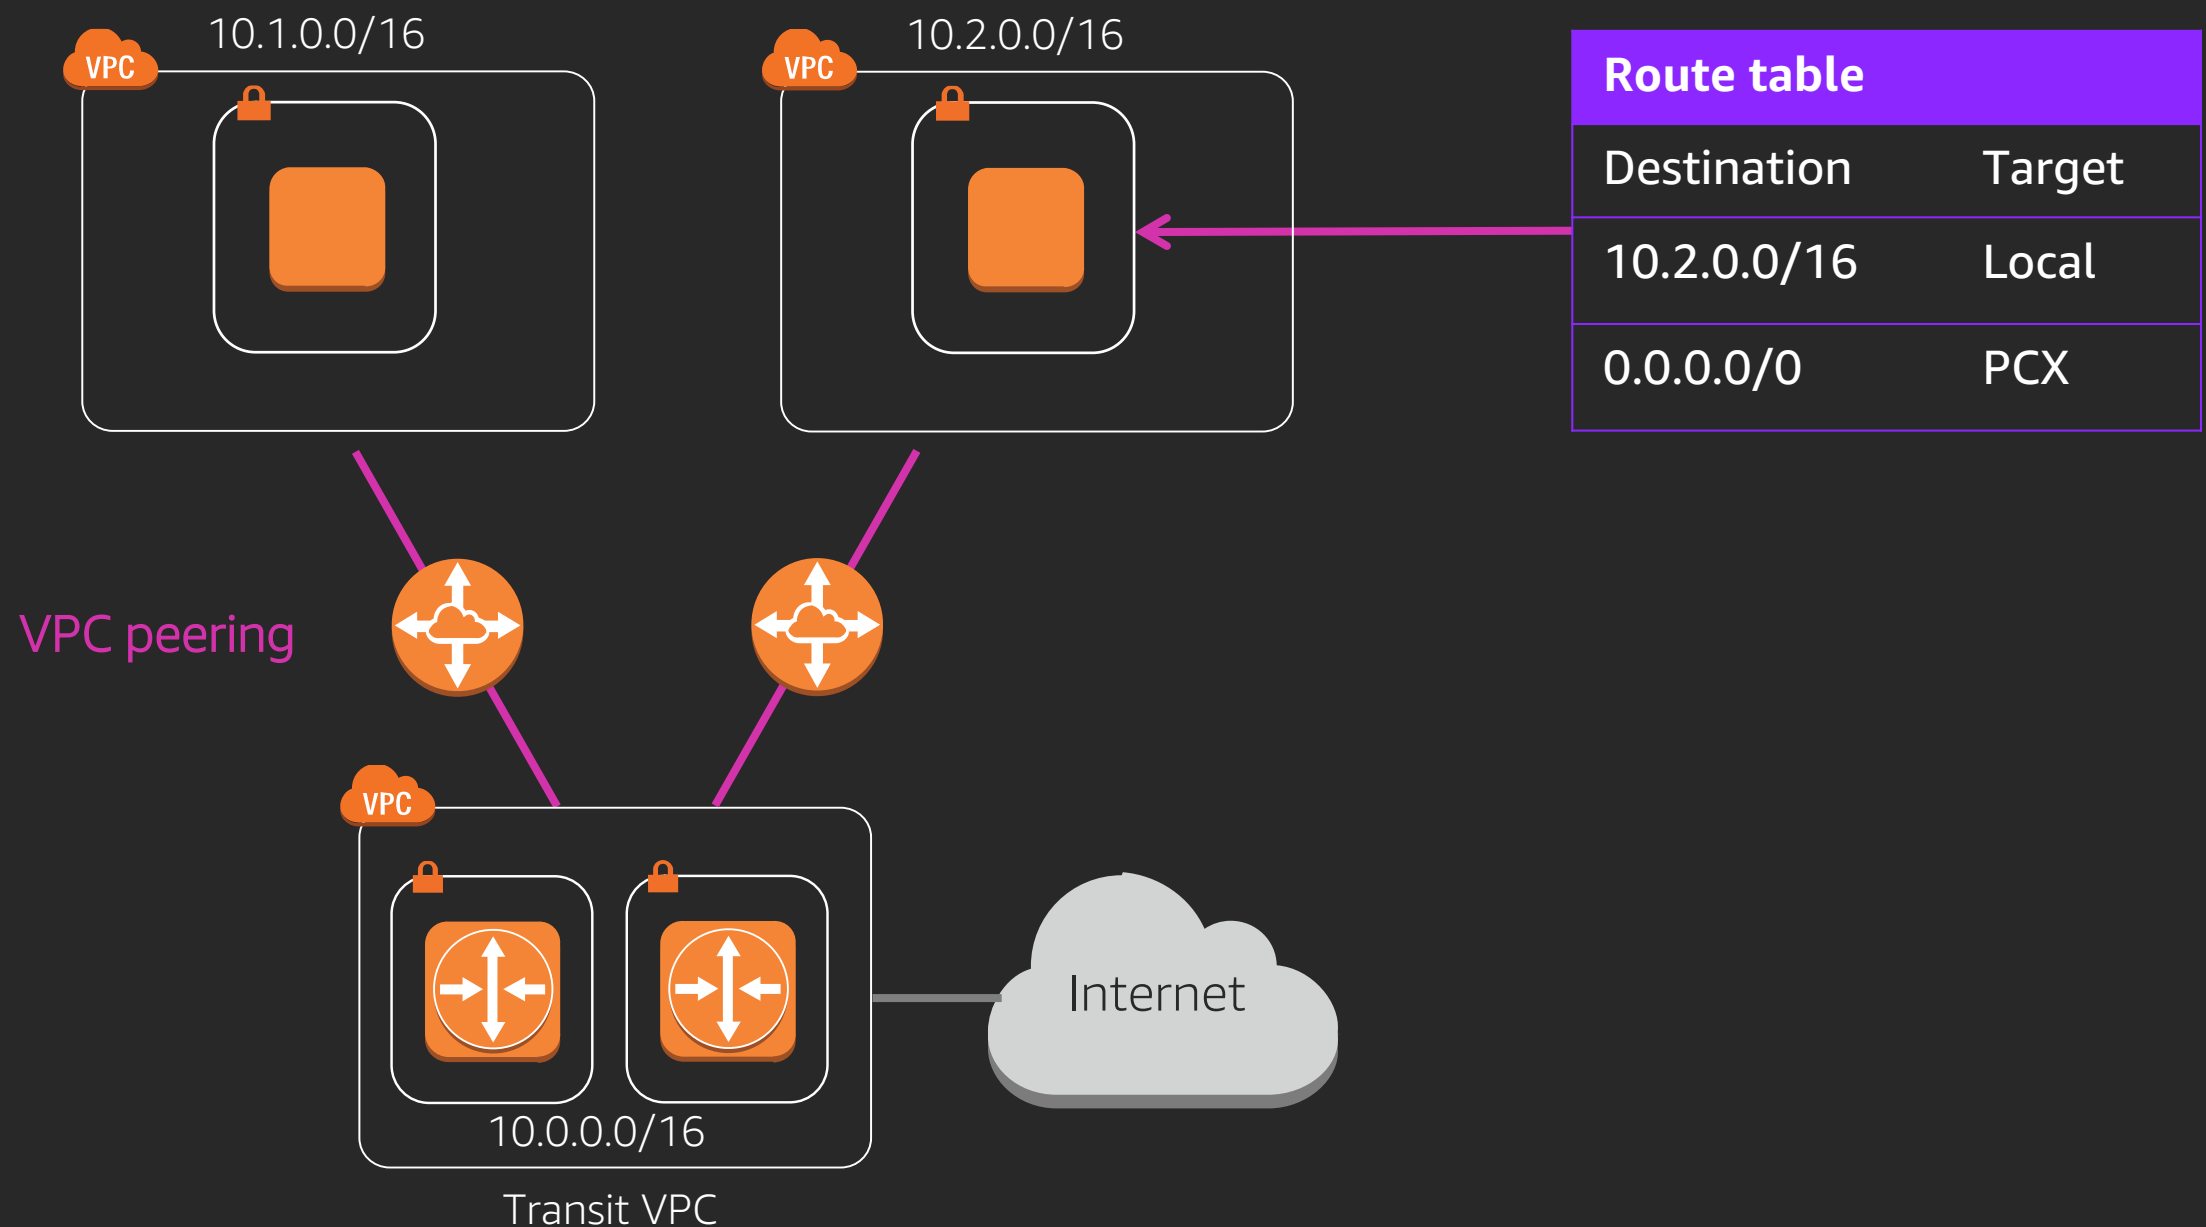


Transit VPC Mechanics


Transit VPC: Routing



Why doesn't peering work?



Why doesn't peering work?



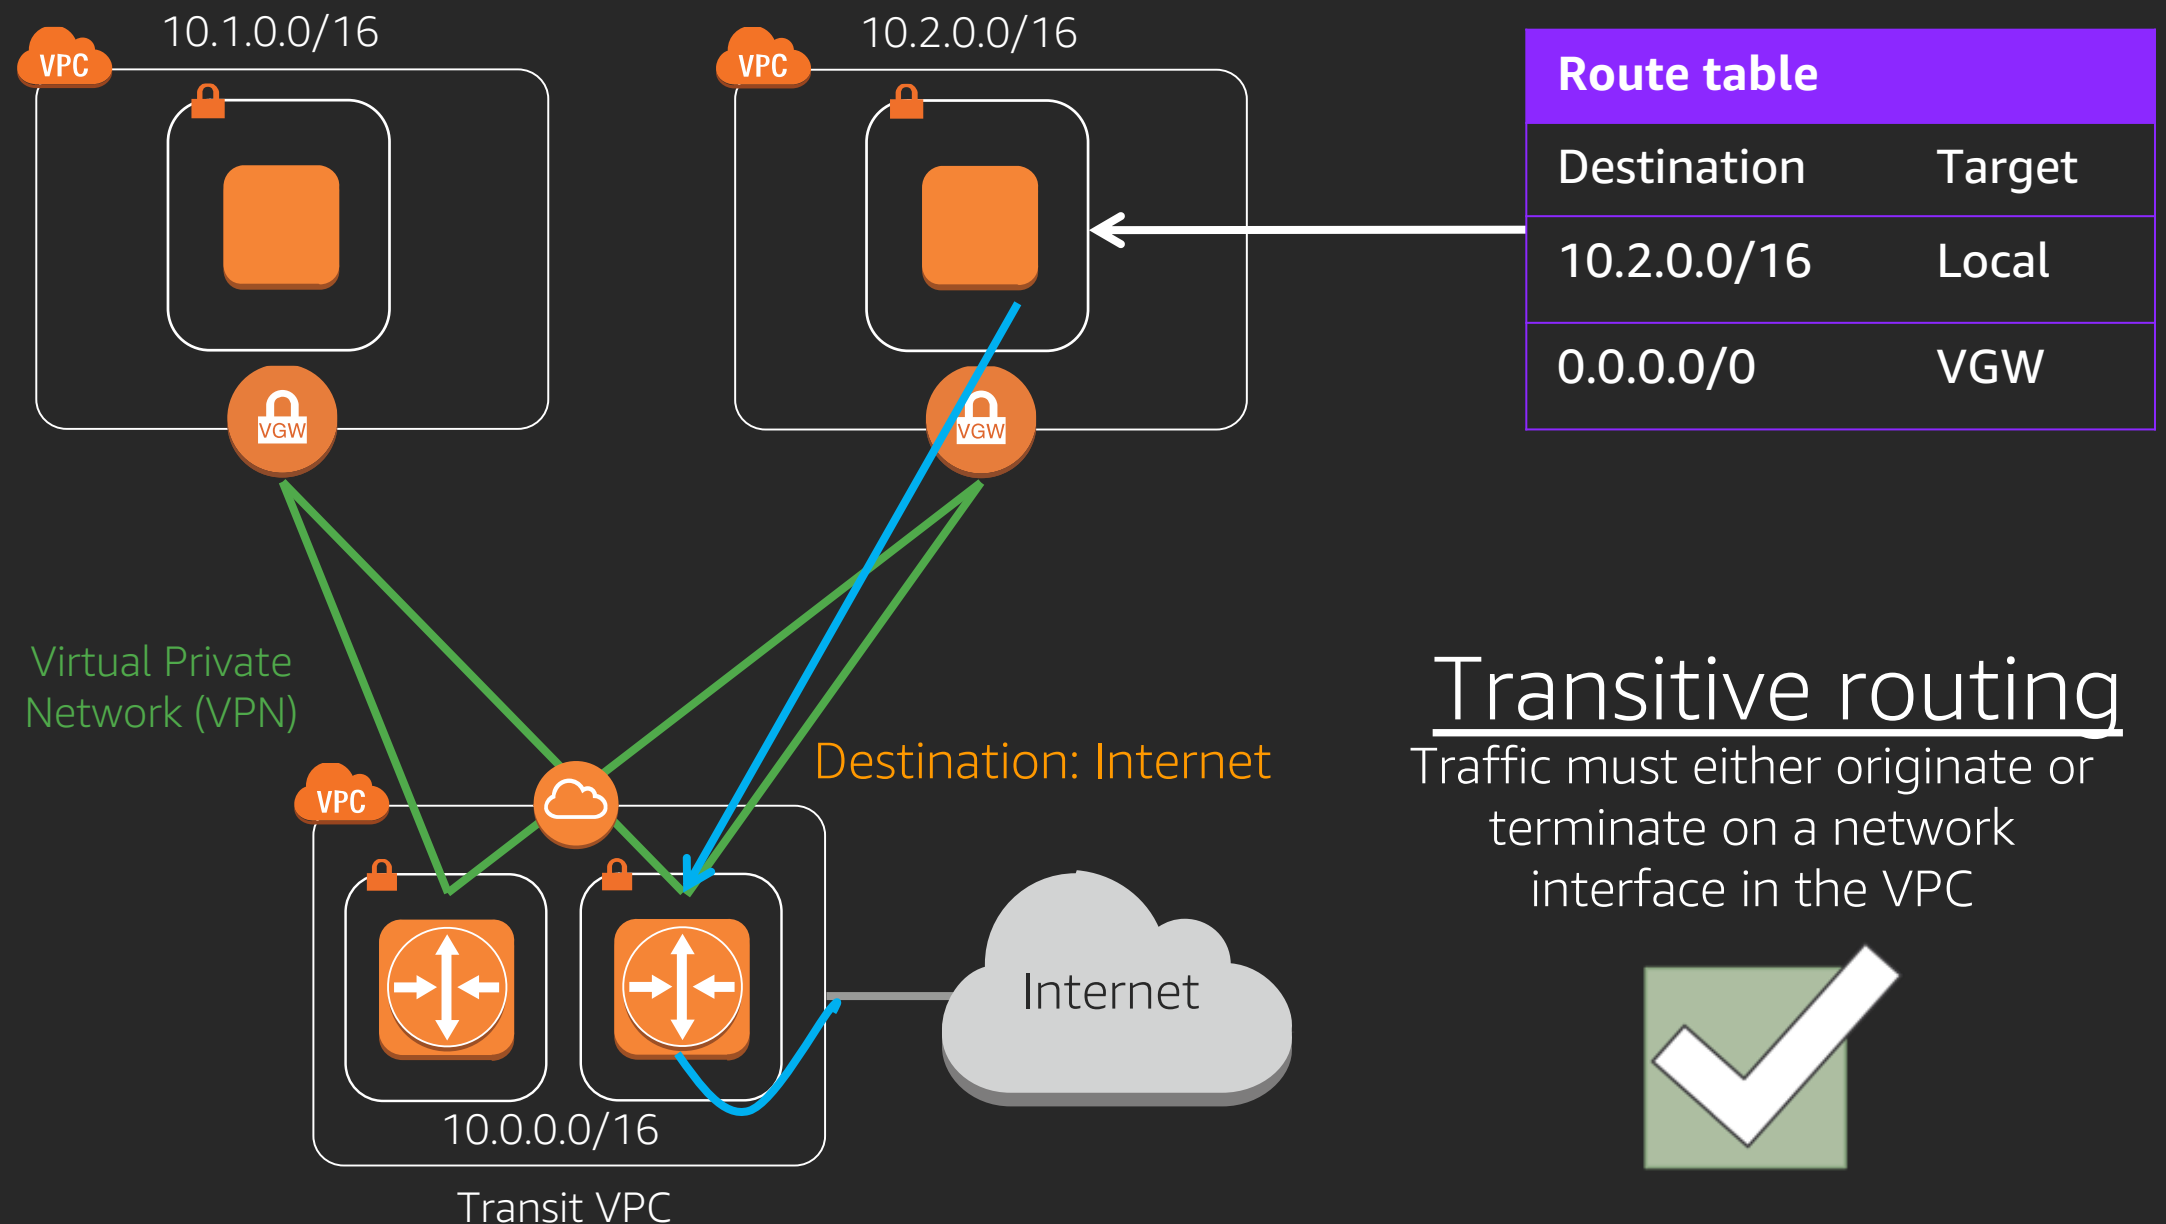
Route table

Destination	Target
10.2.0.0/16	Local
0.0.0.0/0	PCX

Transitive routing

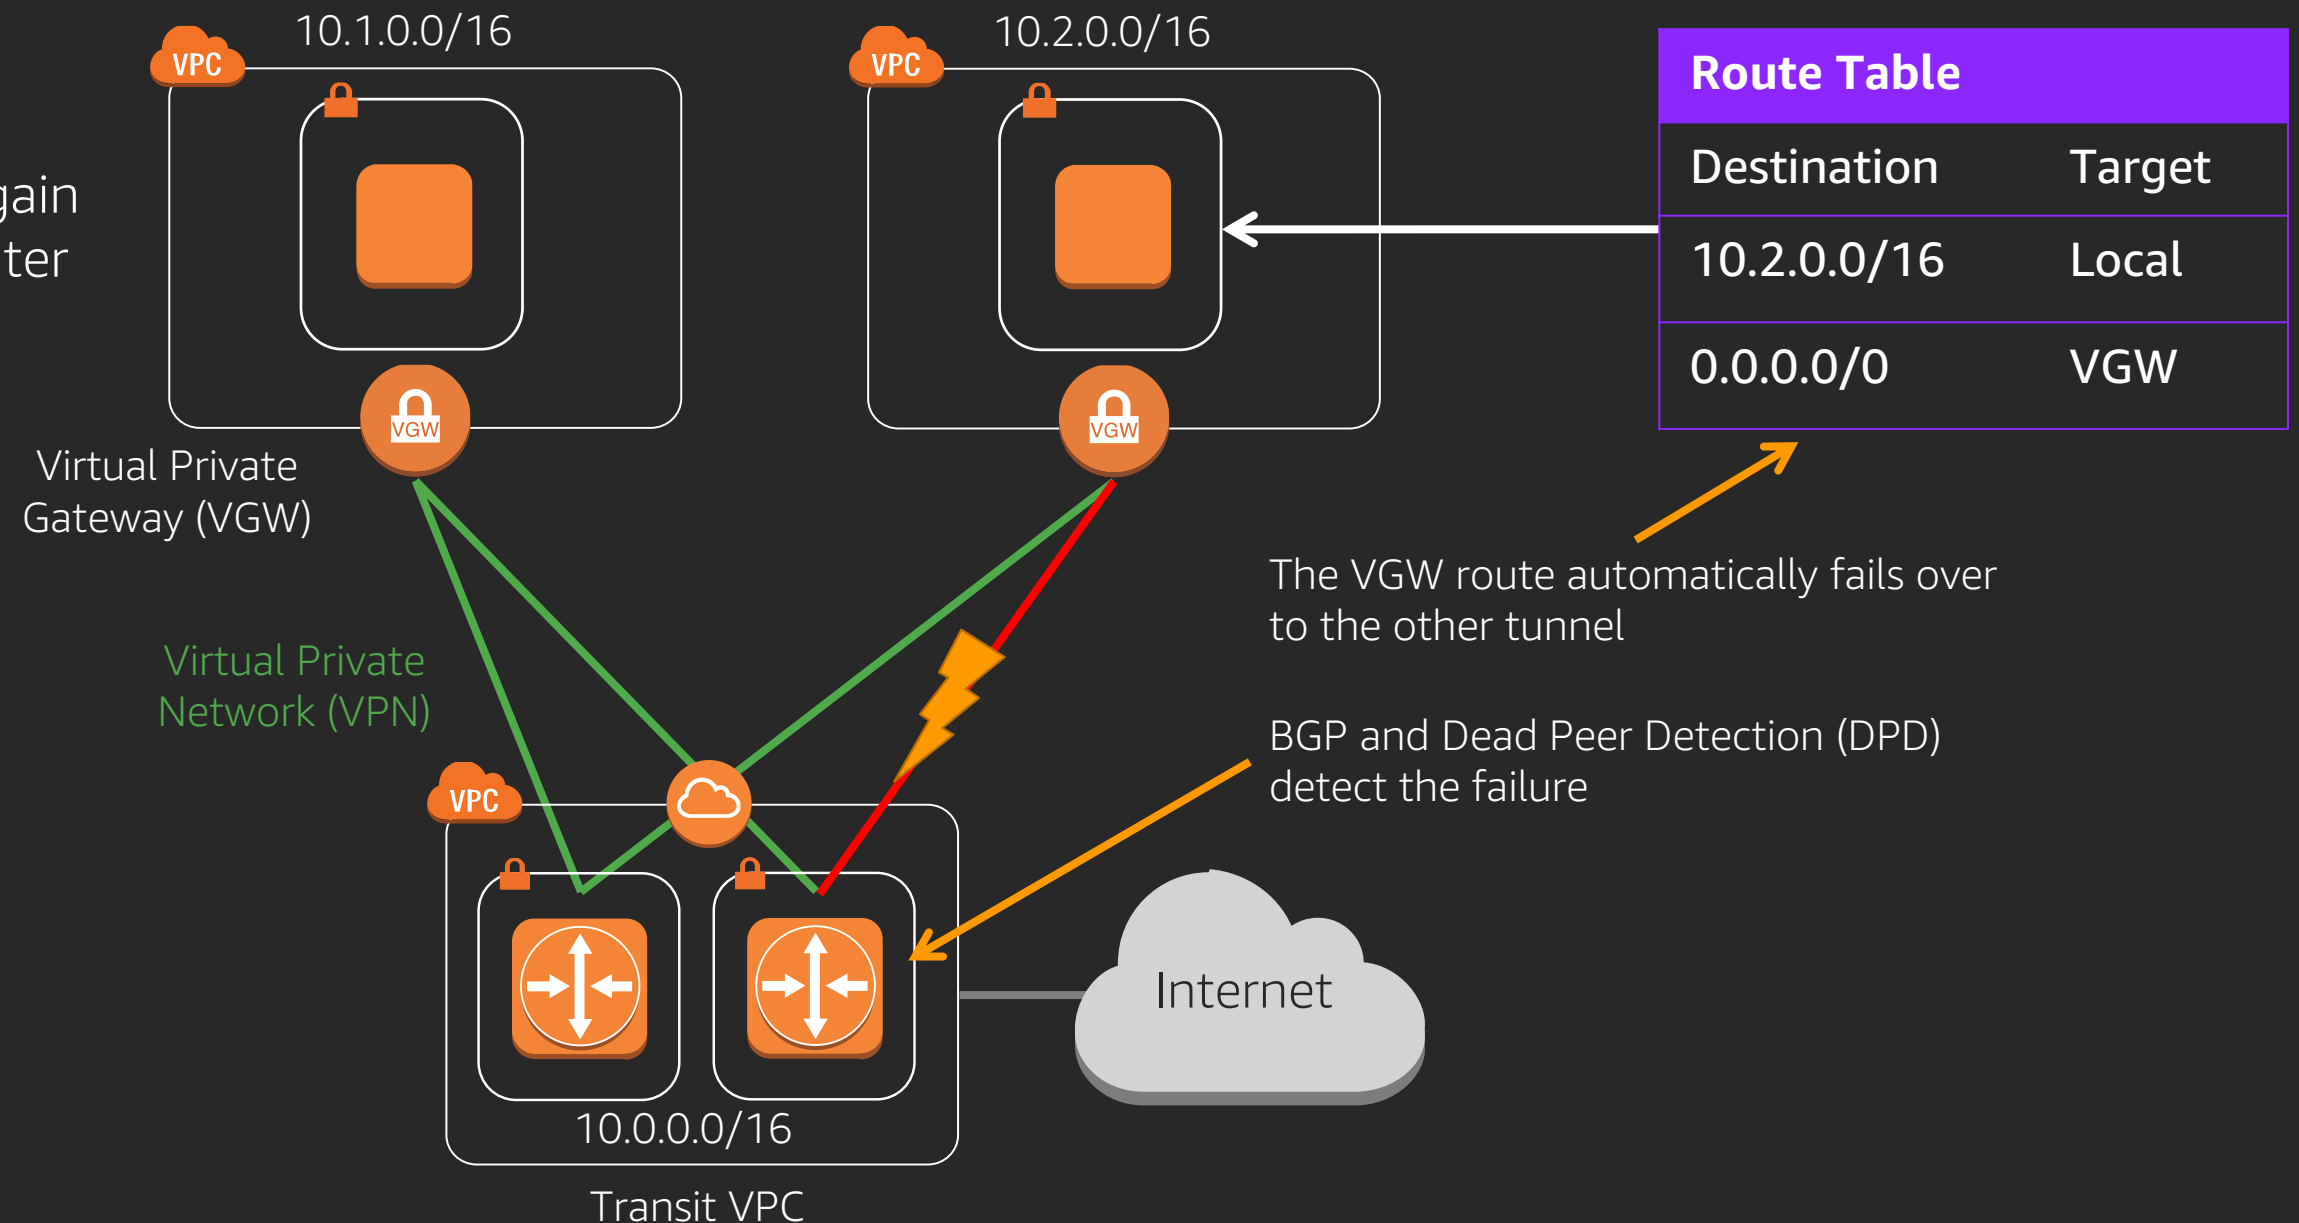
Traffic must either originate or terminate on a network interface in the VPC

Why does VPN work?



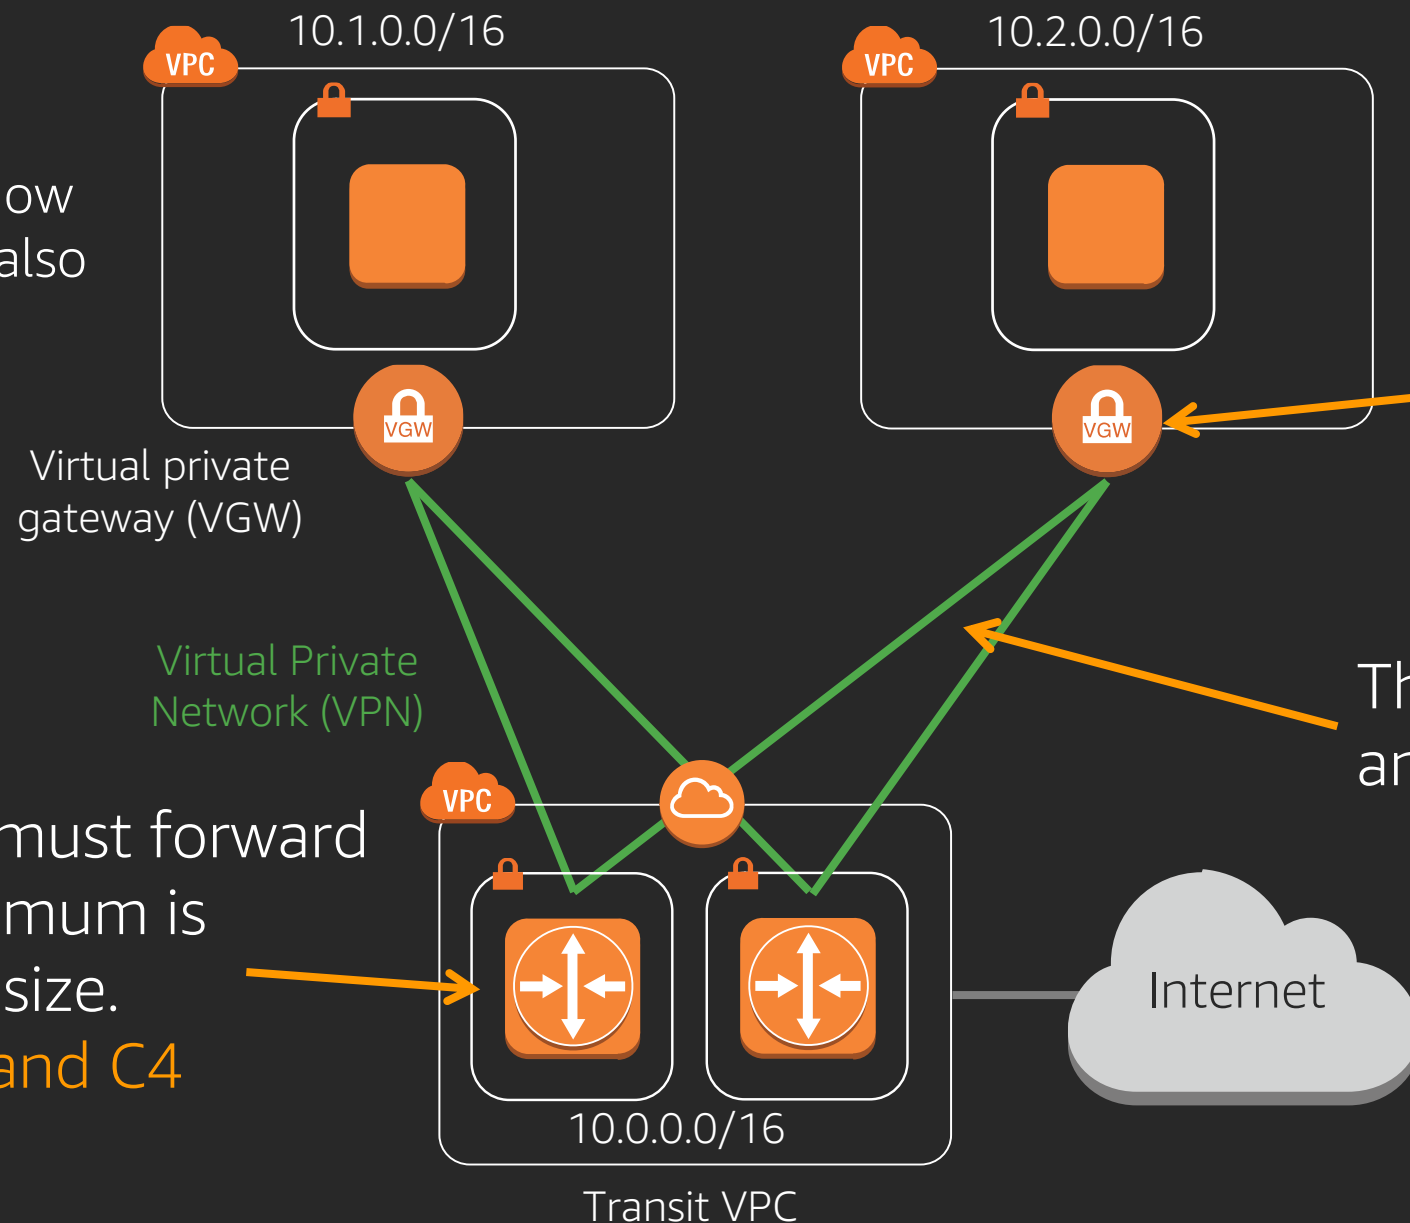
Transit VPC: Availability

Spoiler: We'll use this again with Transit Gateway later



Transit VPC: Performance

Spoiler: We'll need to know this for Transit Gateway also




The VGW will only choose a single tunnel for outbound traffic (1.25 gbps)


The VGW accepts packets on any tunnel or connection

The VPN instance must forward all traffic, the maximum is based on instance size.
~1-3 gbps on M4 and C4 families.

Transit VPC: Security Services



What is the AWS Transit Gateway?



Introducing: Transit Gateway New

Regional router

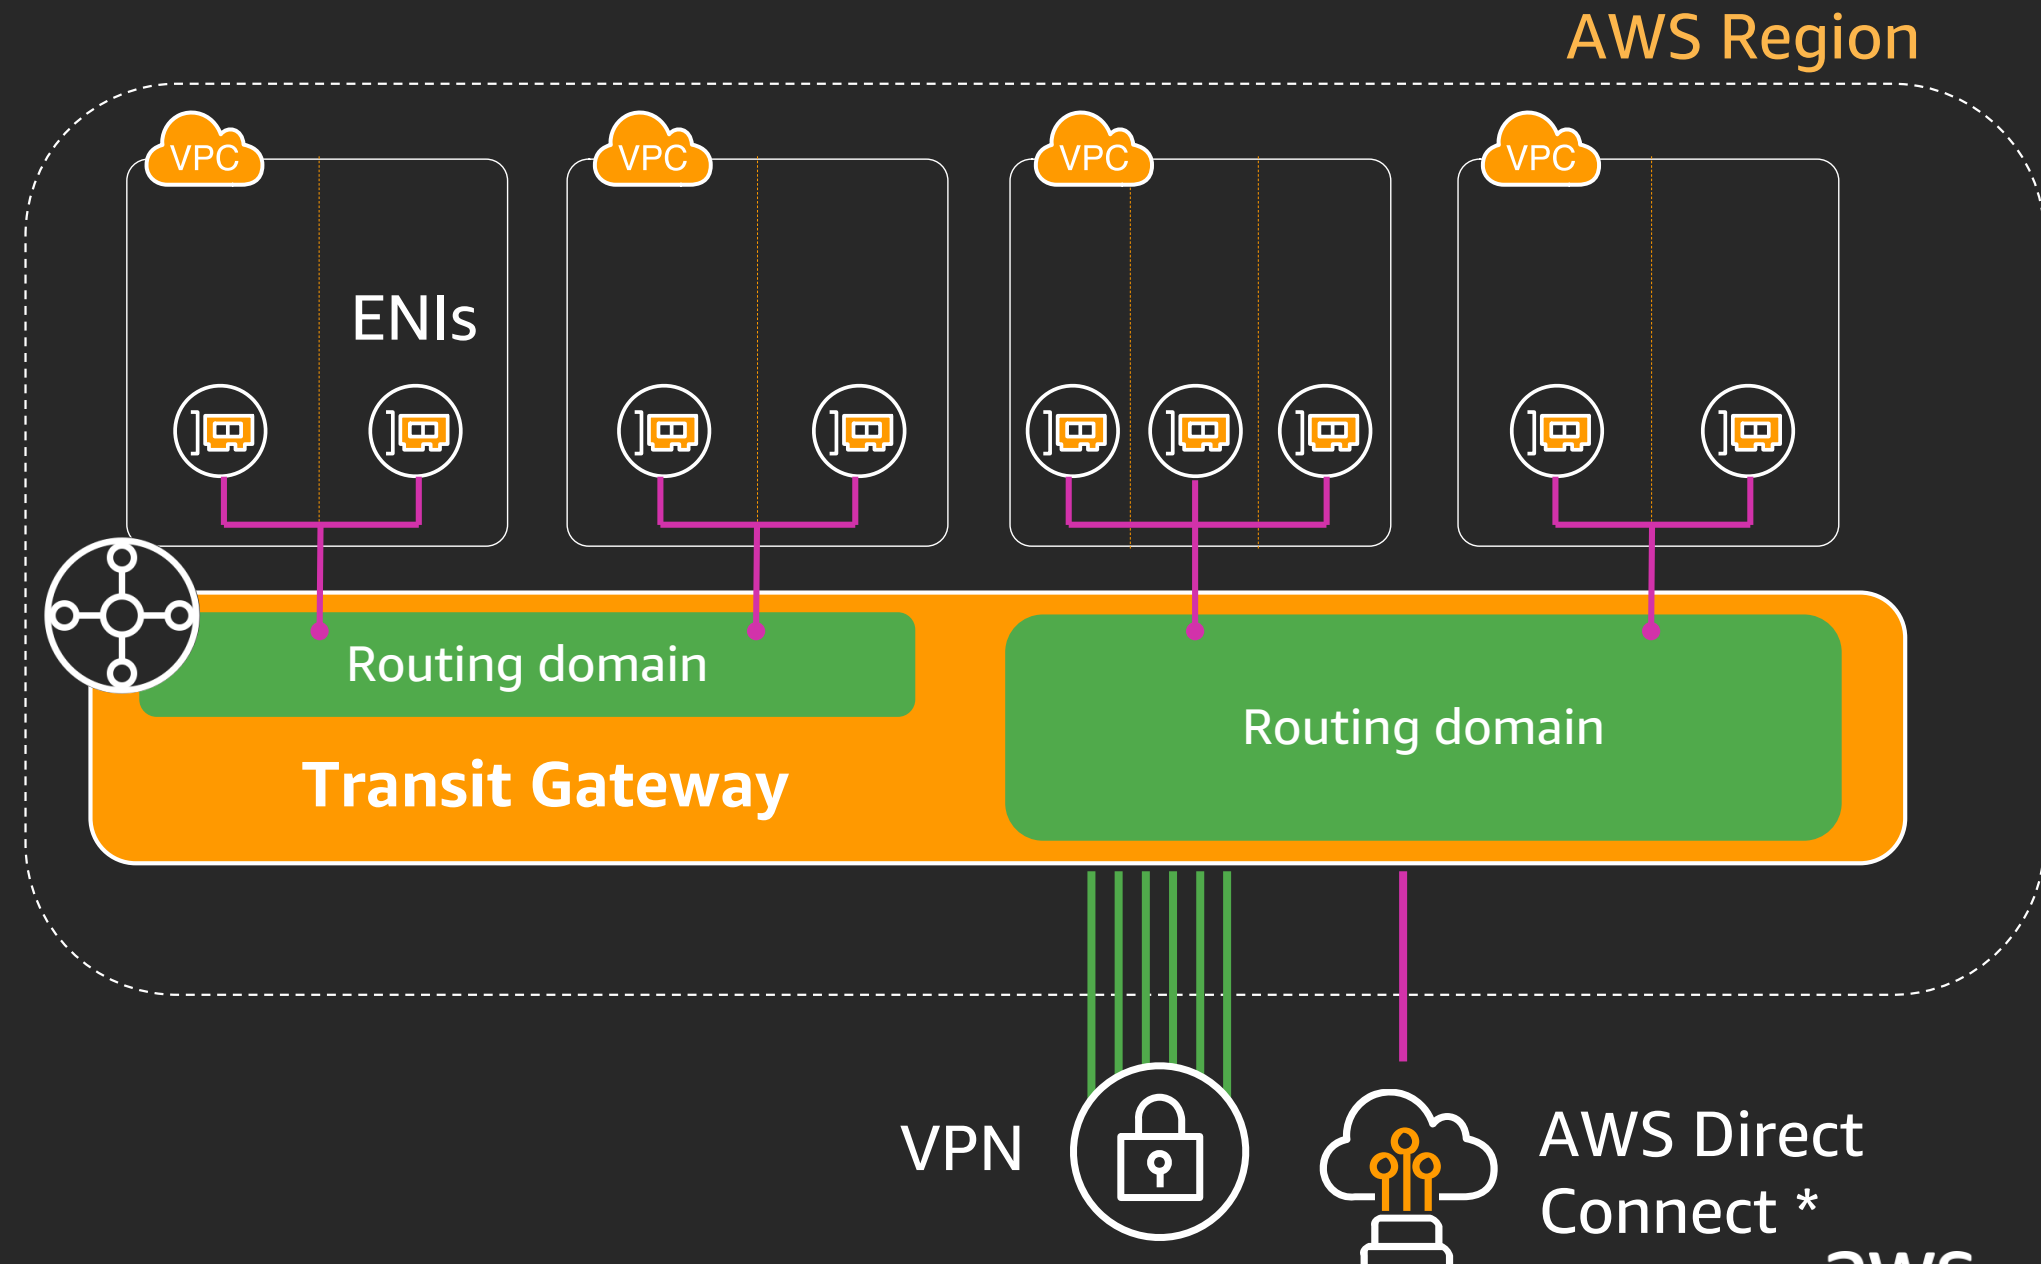
- Centralize VPN and AWS Direct Connect

Scalable

- Thousands of VPCs across accounts
- Spread traffic over many VPN connections

Flexible routing

- Network interfaces in subnets
- Control segmentation and sharing with routing

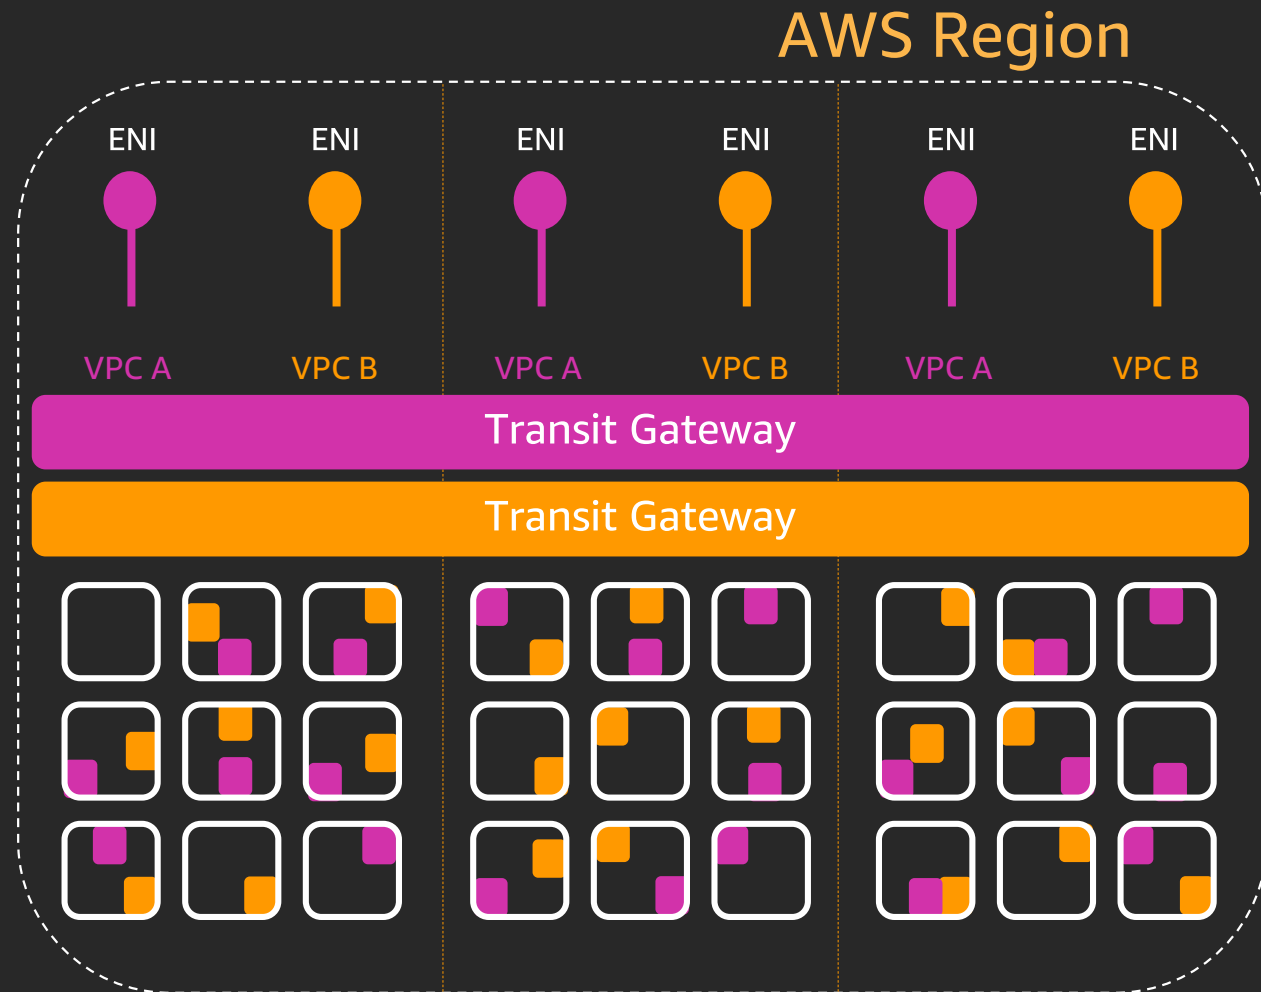


AWS
re:Invent

Available Q1 2019



AWS HyperPlane and AWS Transit Gateway



Attachments

- One network interface per Availability Zone
- Highly available per Availability Zone
- Network capacity shards
- Tens of microseconds of latency

AWS HyperPlane

- Horizontally-scalable state management
- Terabits of multi-tenant capacity
- Supports NLB, NAT Gateway, Amazon EFS and now Transit Gateway

Transit Gateway example time!

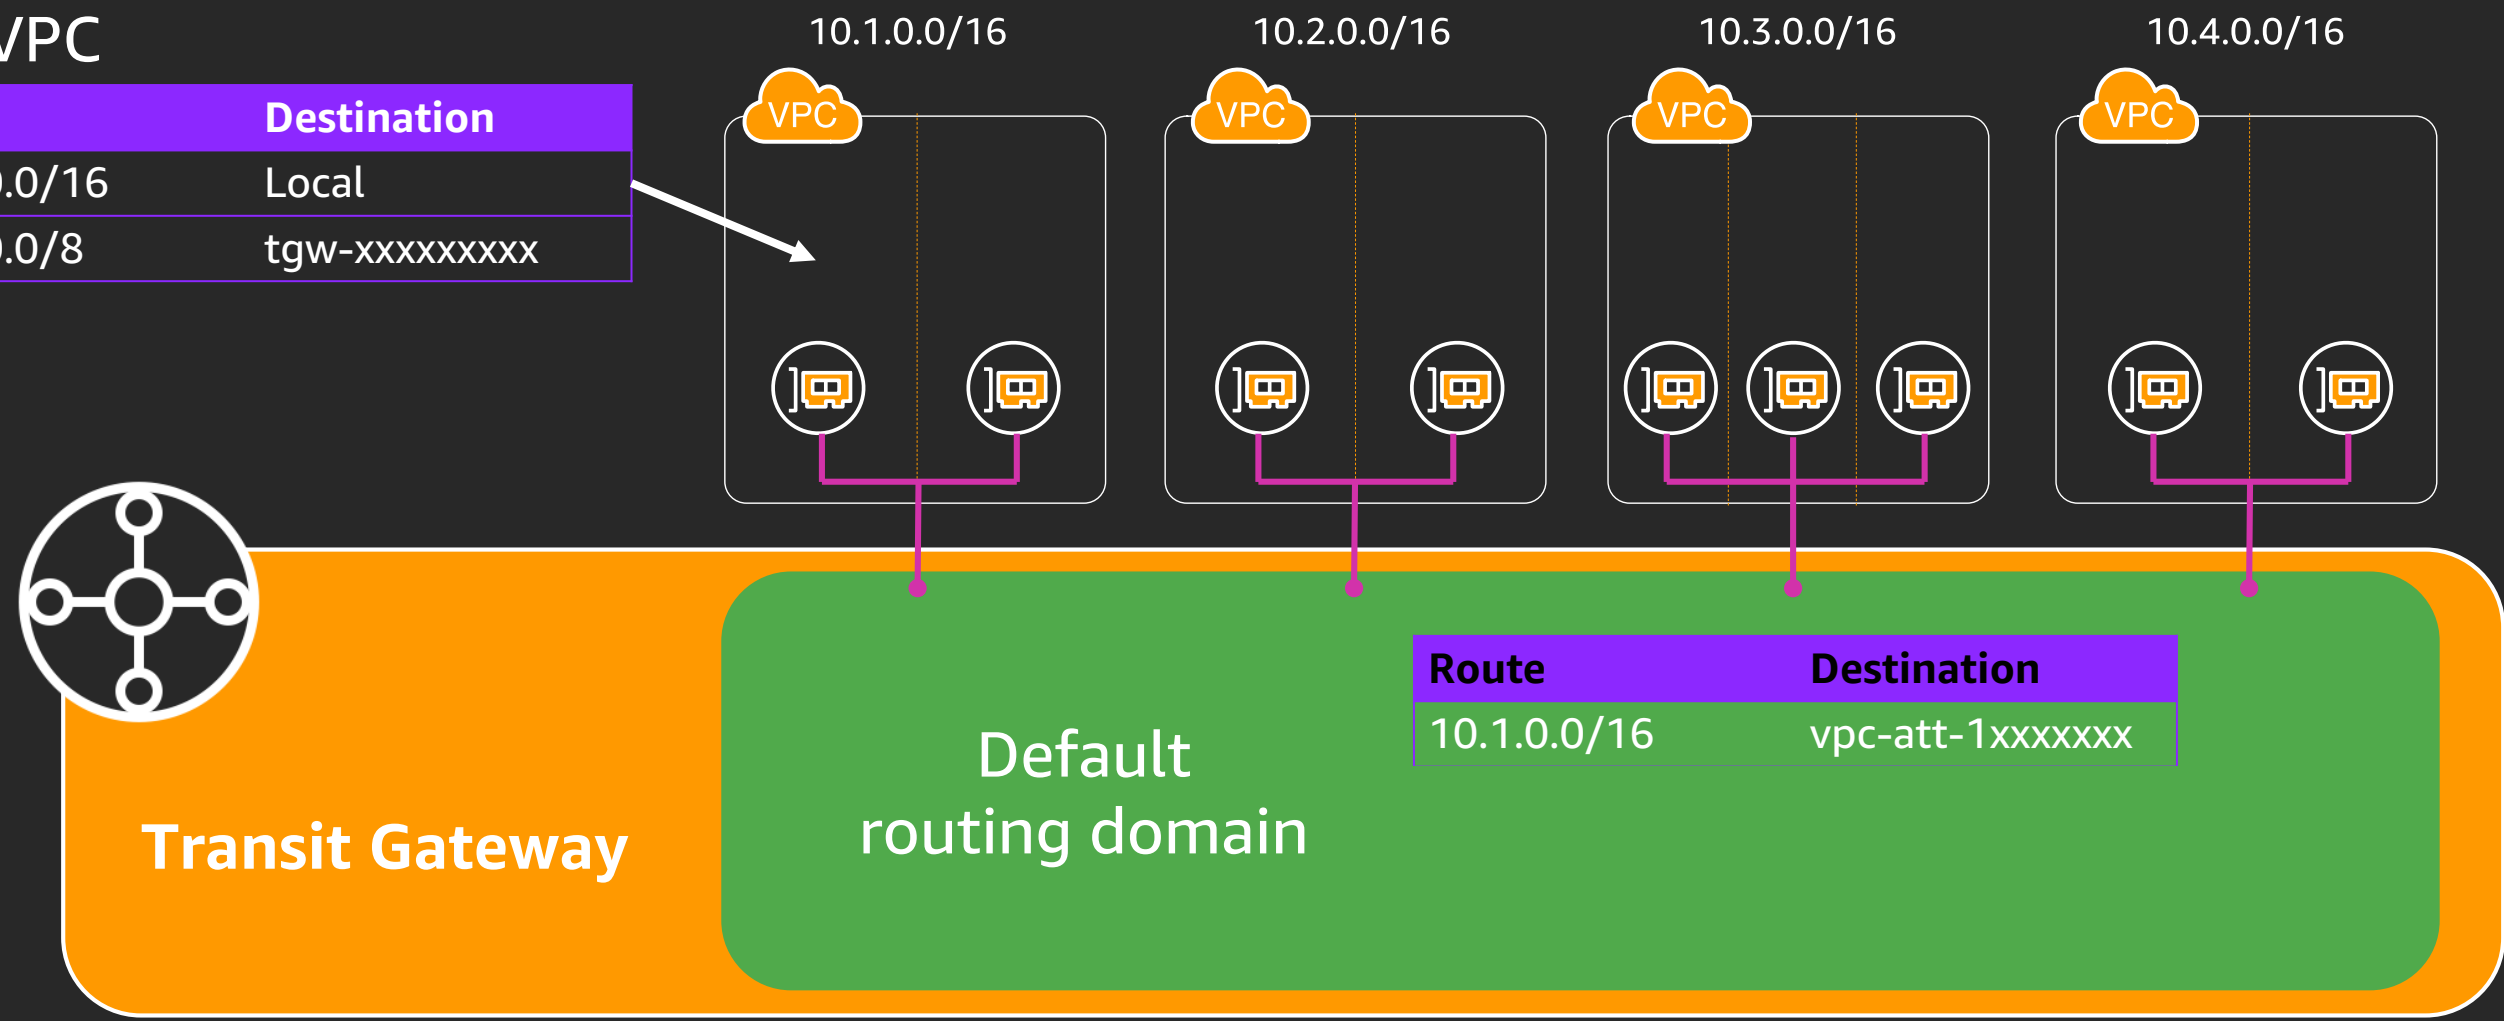
Flat: Every VPC should talk to every VPC!

Isolated: Don't let anything talk! Send everything back over VPN!

Flat: Transit Gateway route domains (route tables)

Per VPC

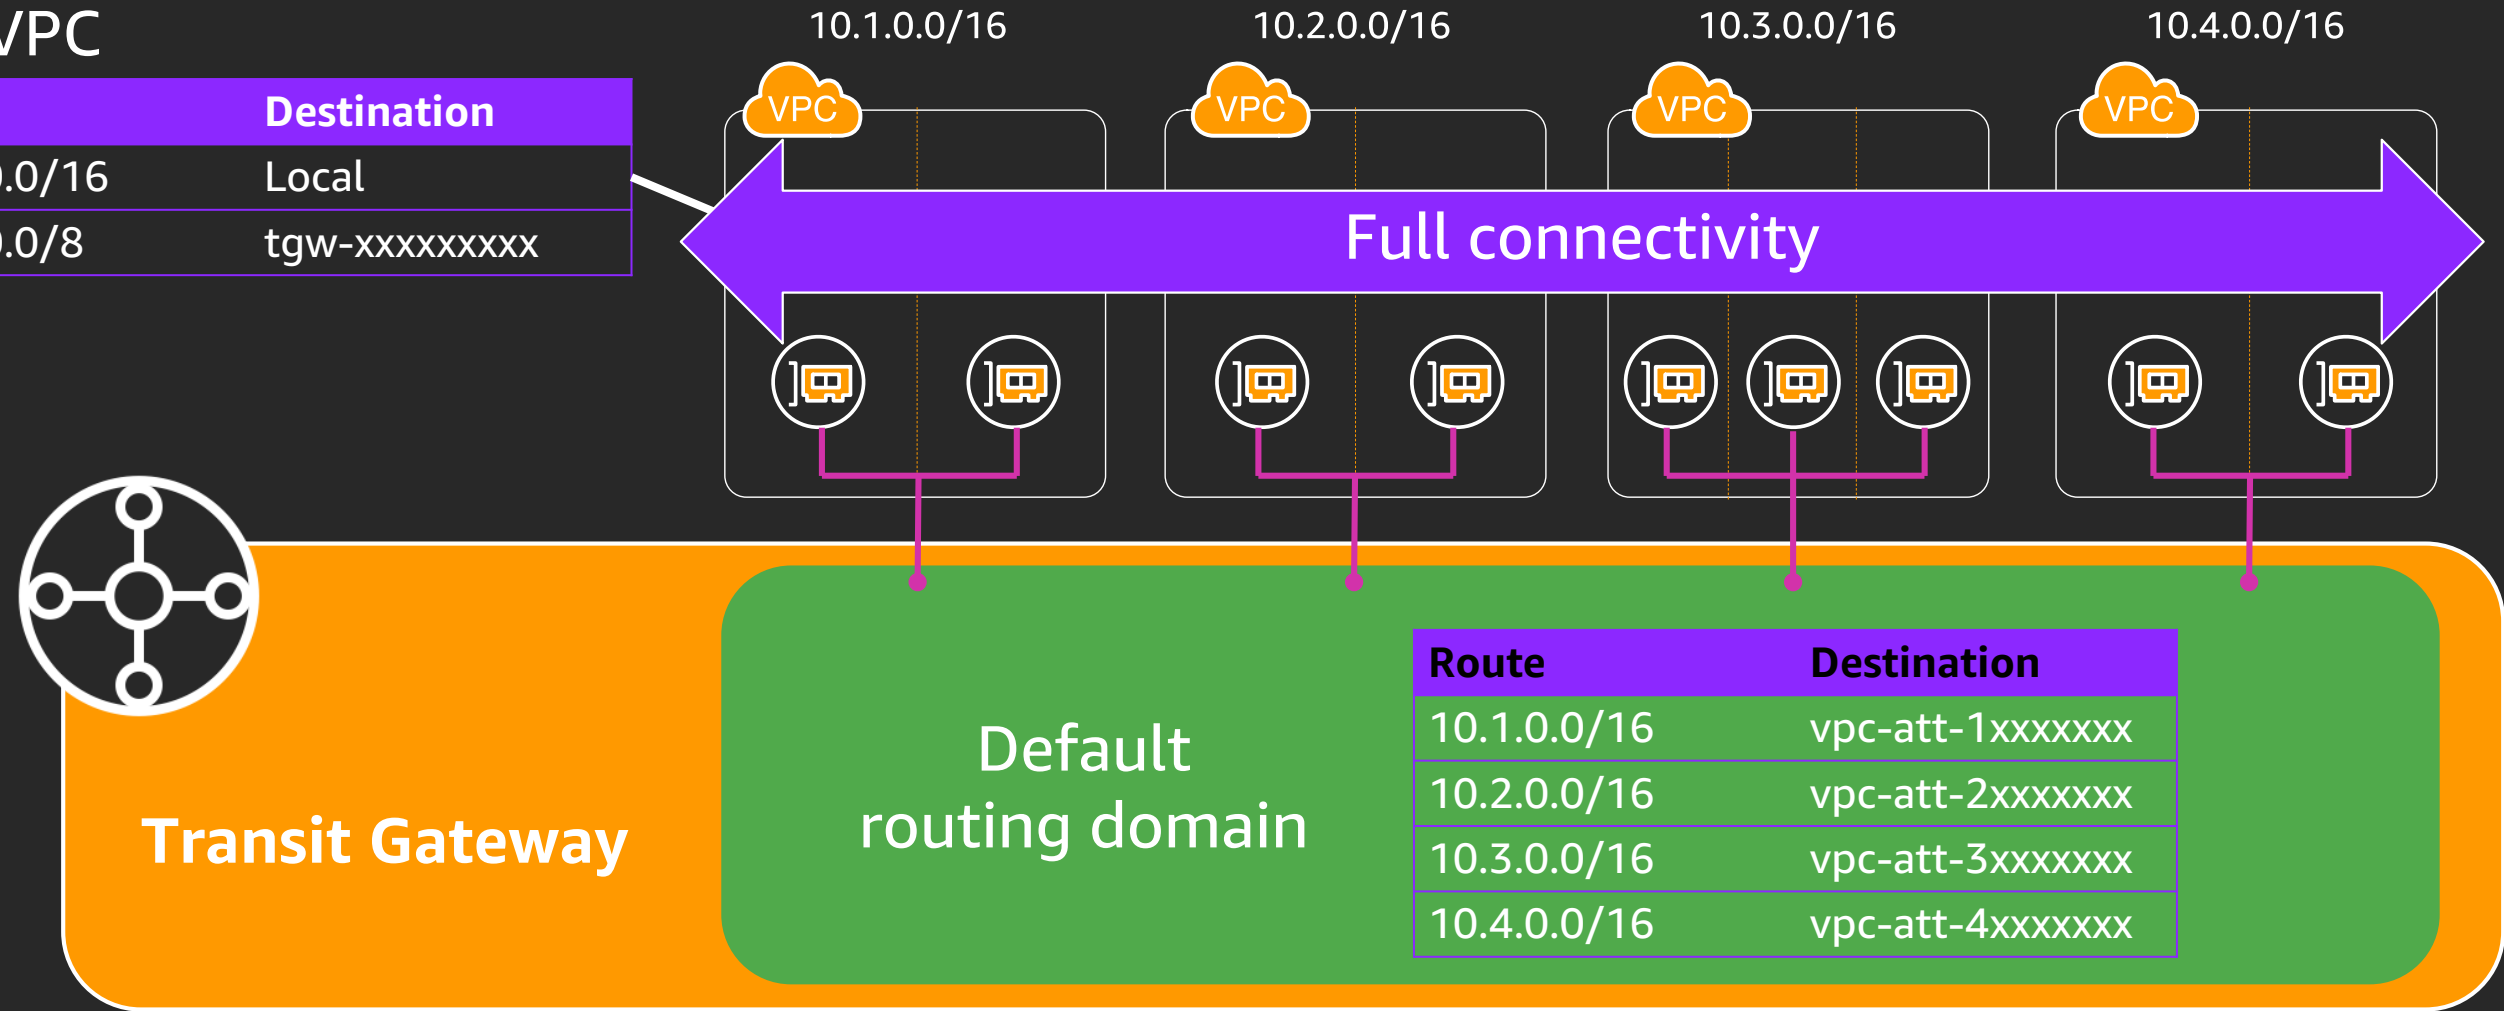
Route	Destination
10.1.0.0/16	Local
10.0.0.0/8	tgw-xxxxxxxxx



Flat: Transit Gateway route domains (route tables)

Per VPC

Route	Destination
10.1.0.0/16	Local
10.0.0.0/8	tgw-xxxxxxxxx

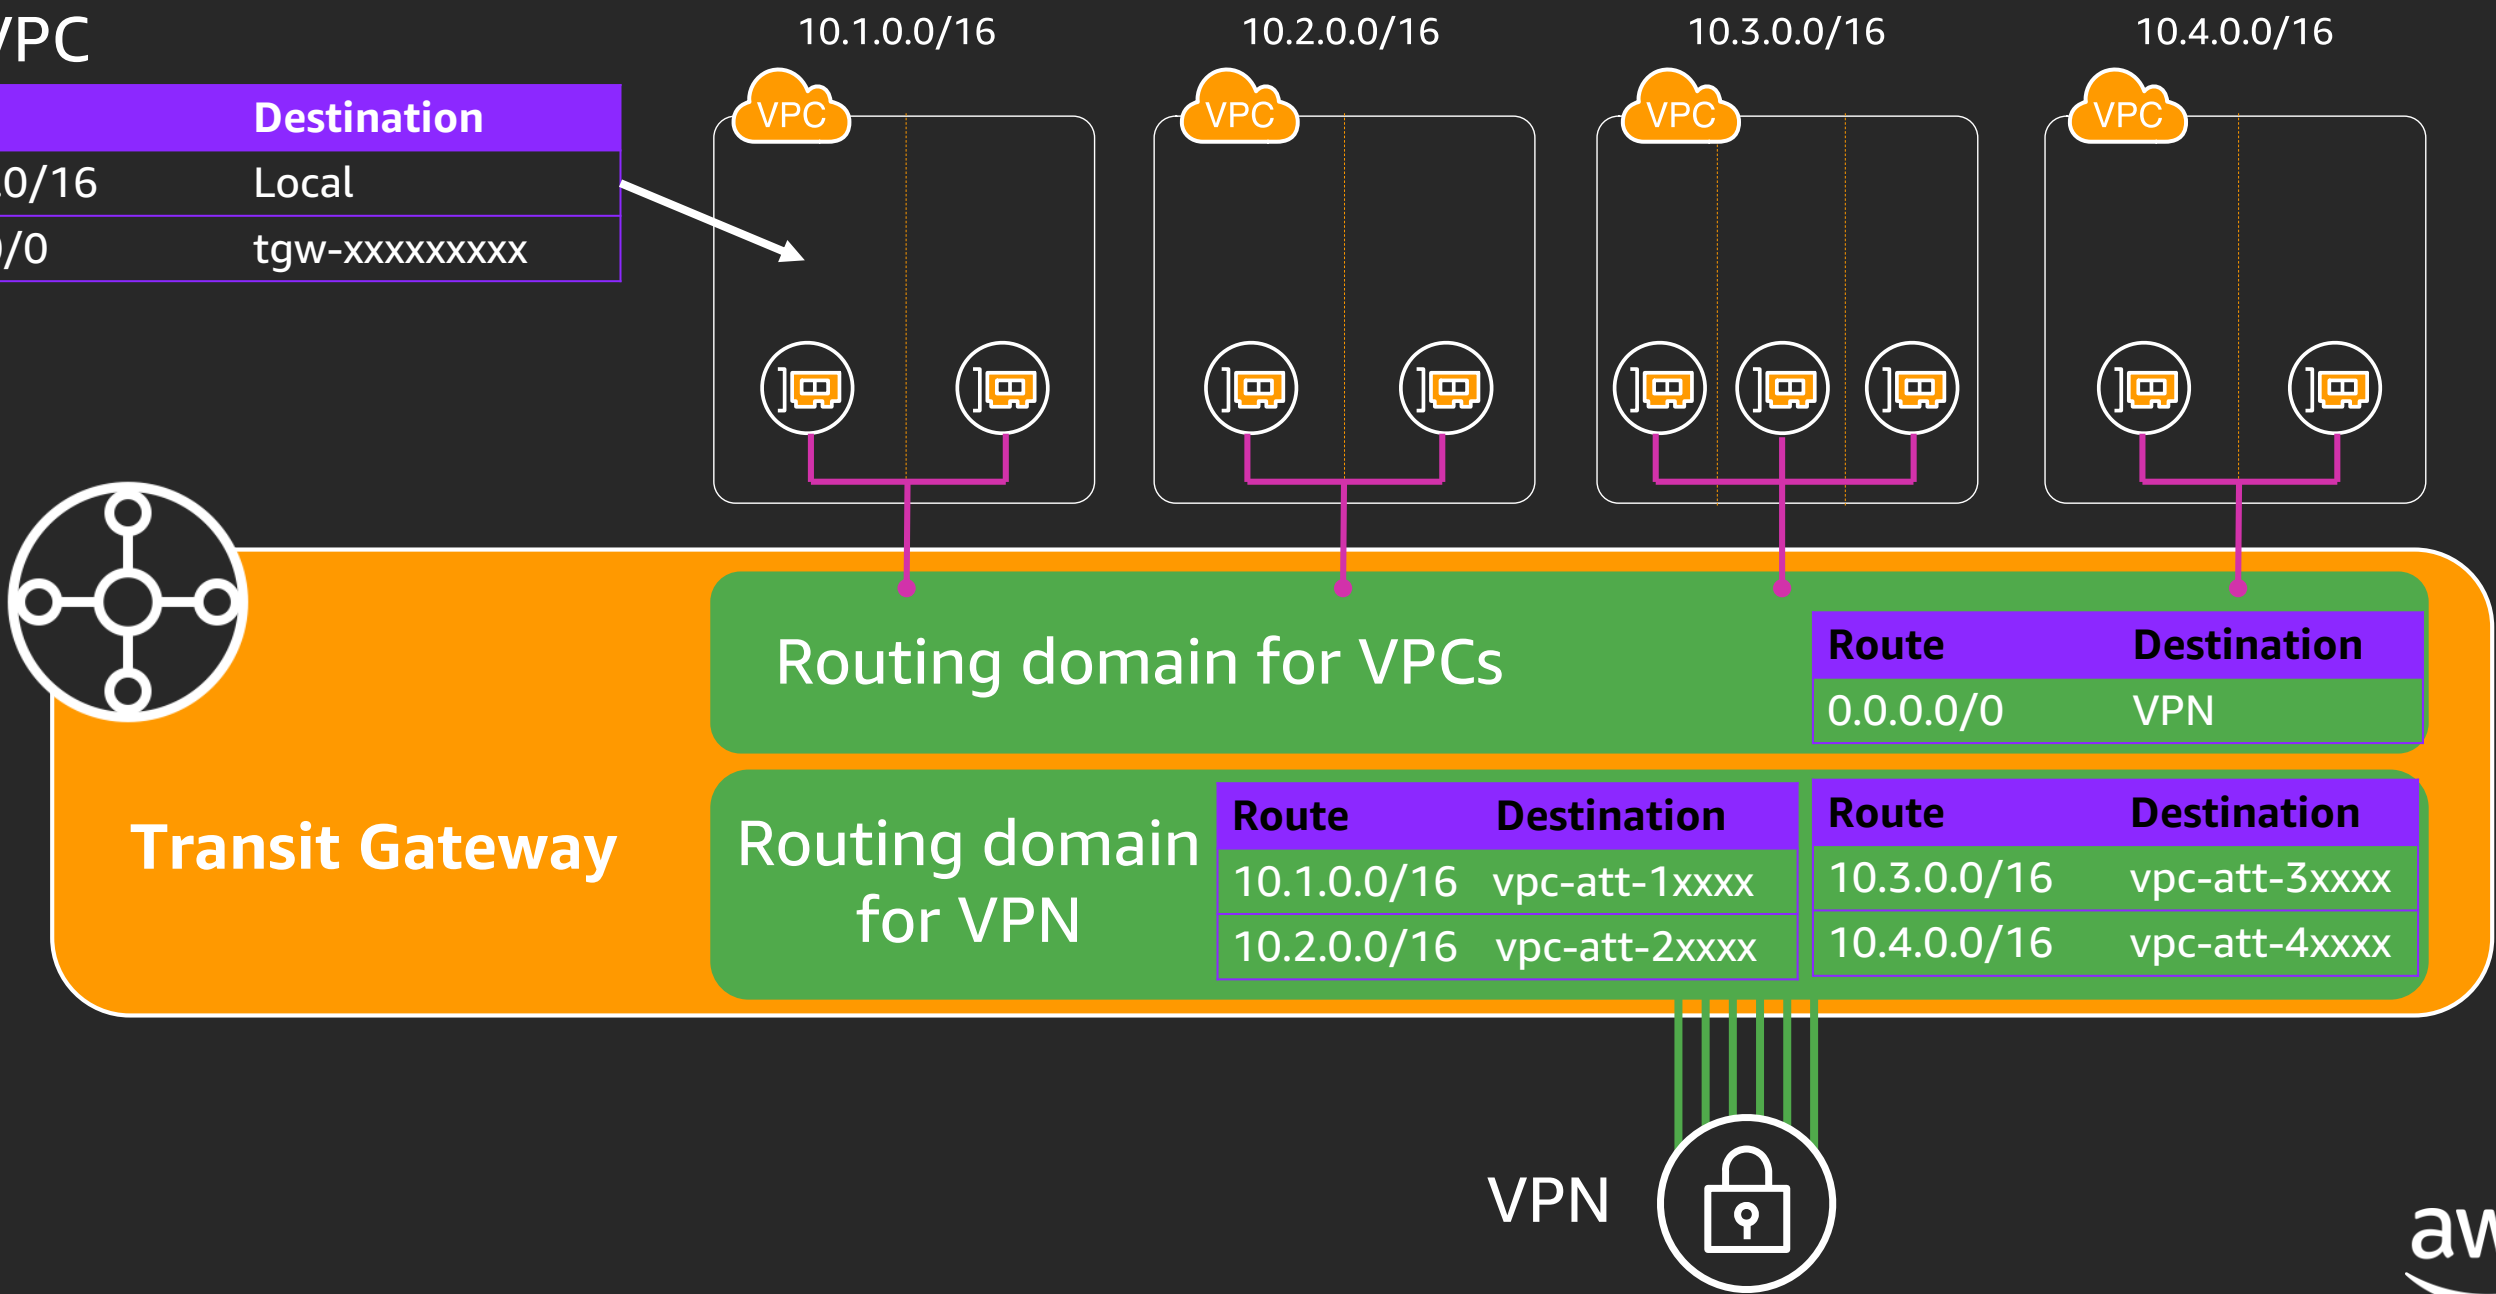


Wording warning: In this presentation a route domain is a route table of a Transit Gateway

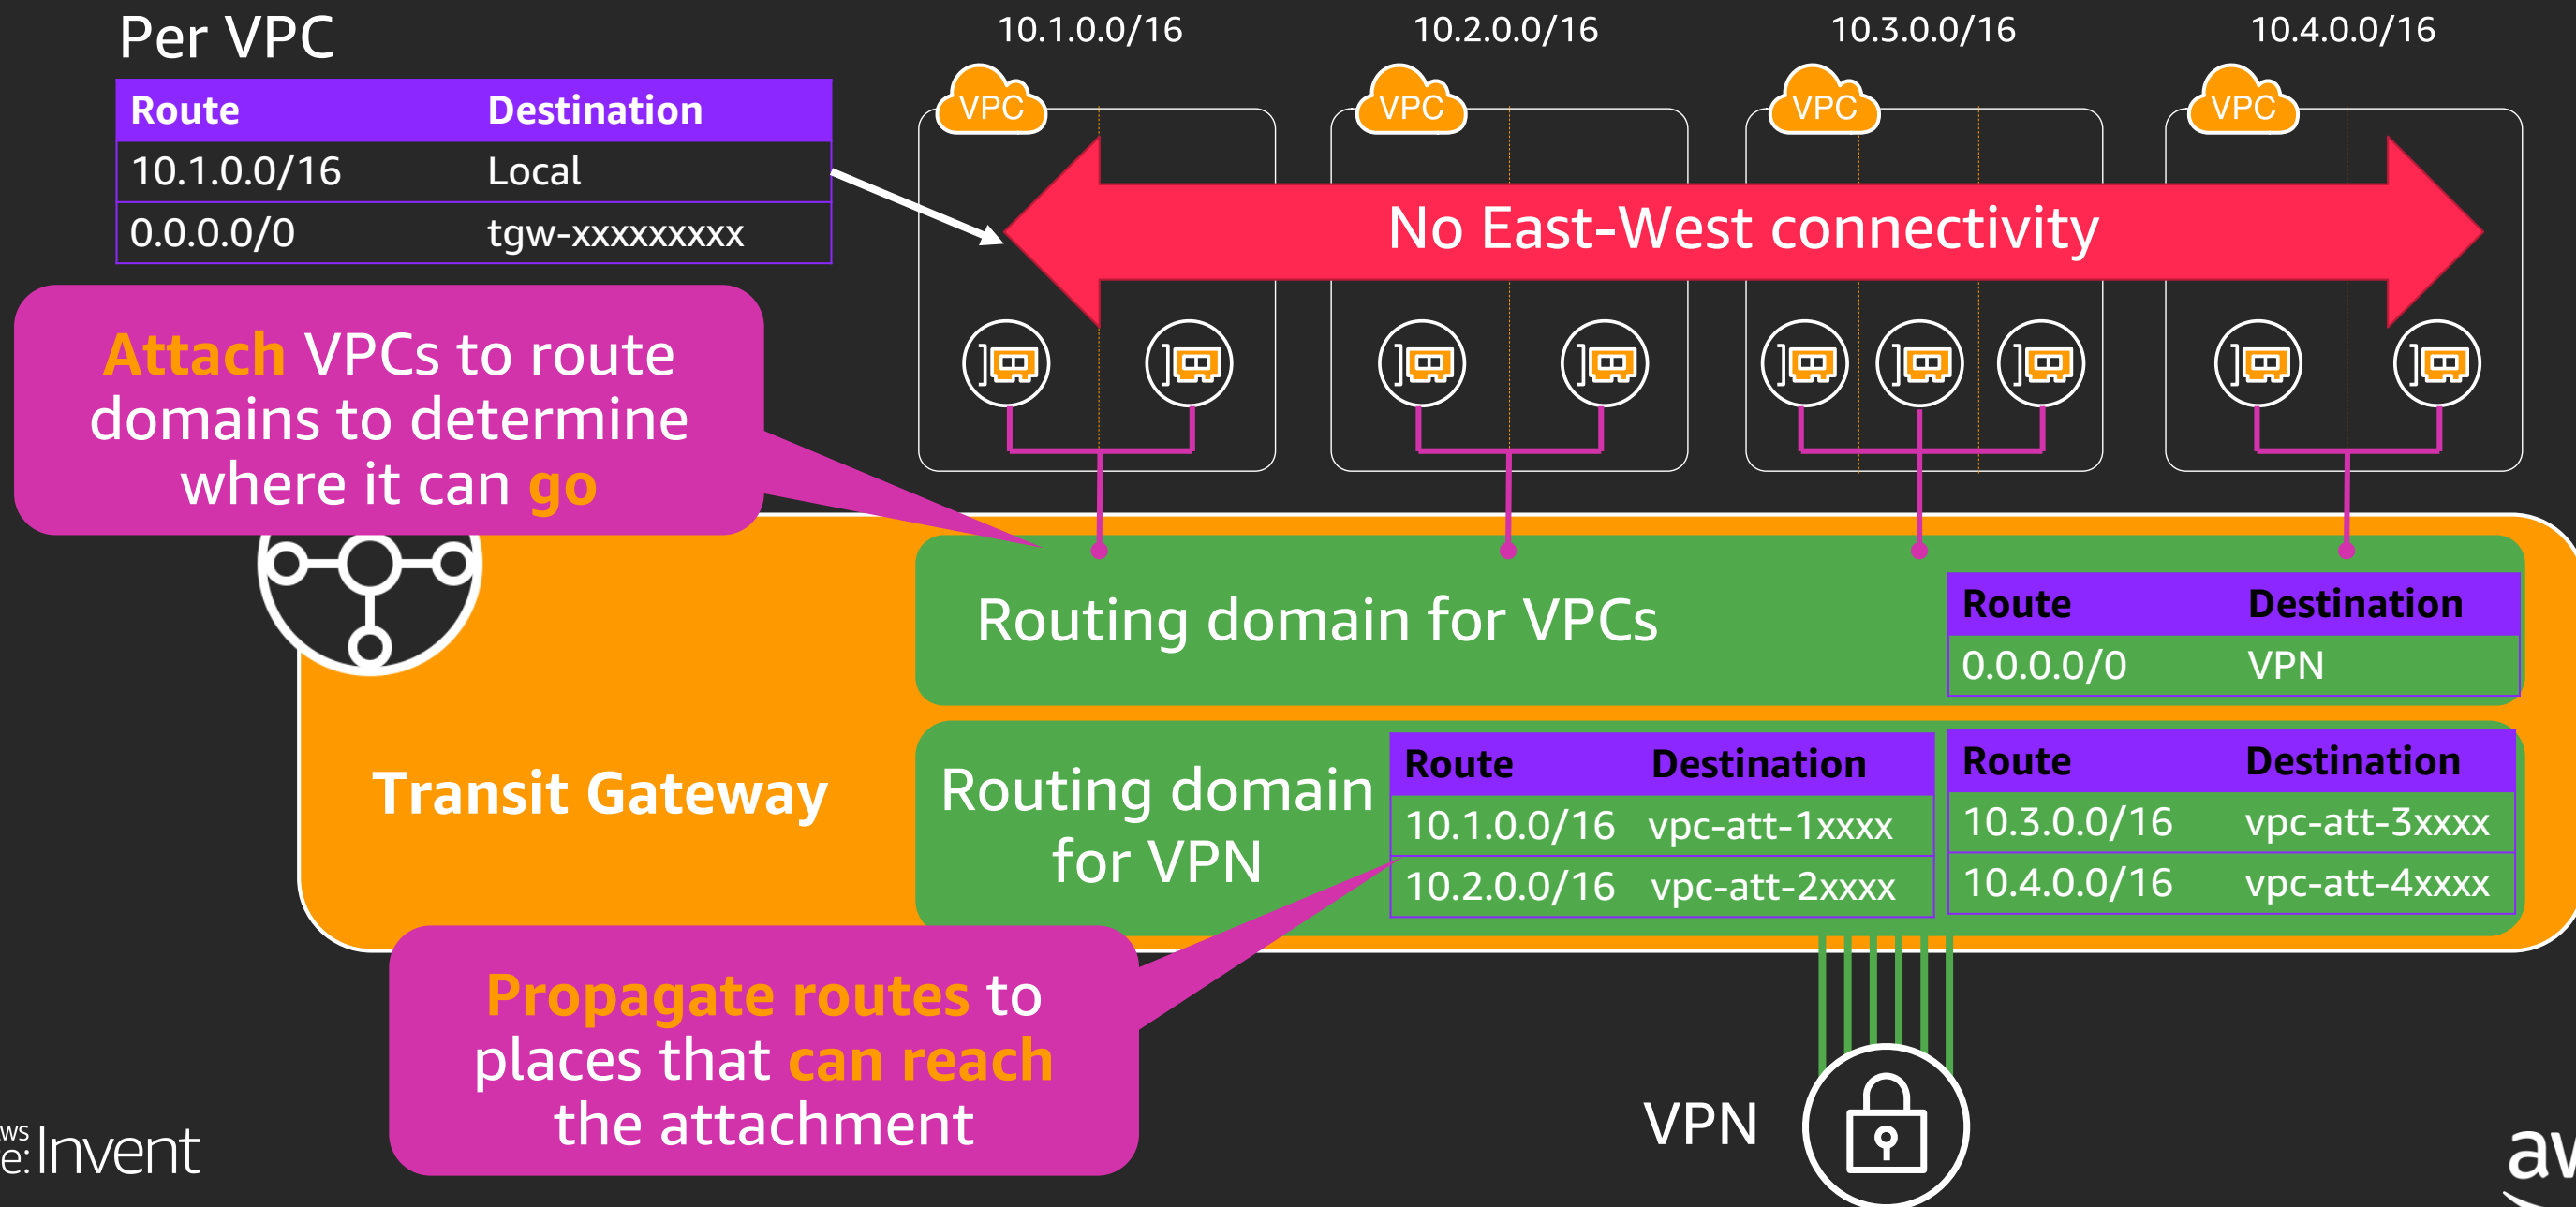
Isolated: Transit Gateway route domains

Per VPC

Route	Destination
10.1.0.0/16	Local
0.0.0.0/0	tgw-xxxxxxxxx



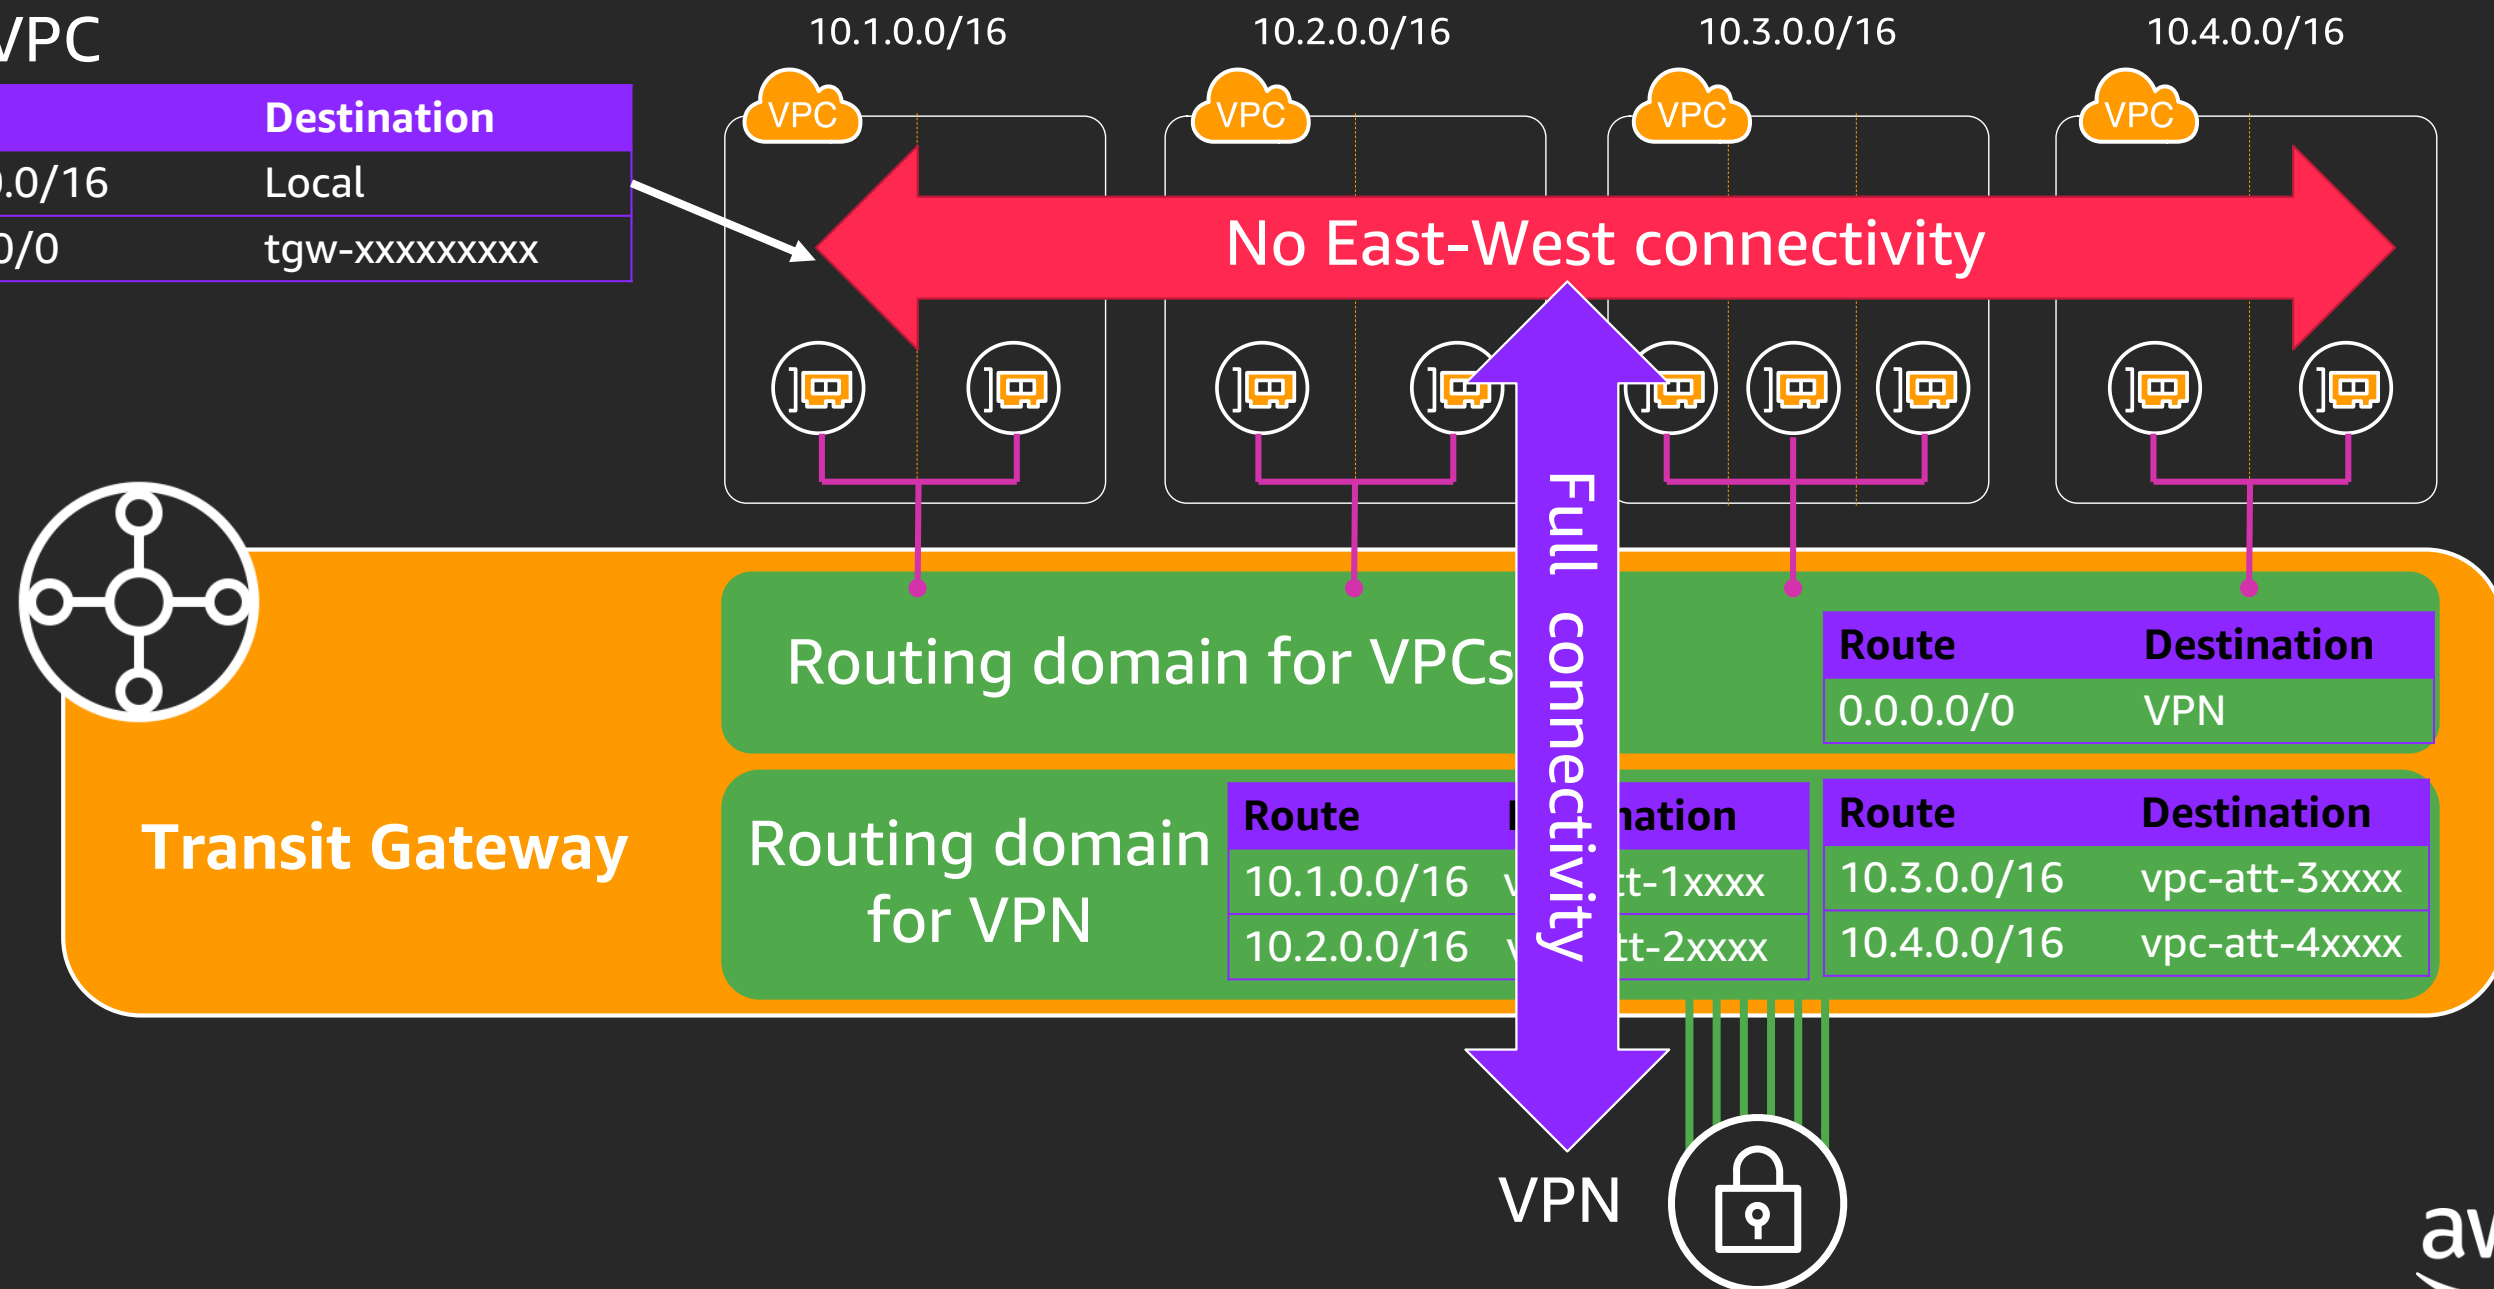
Isolated: Transit Gateway route domains




Isolated: Transit Gateway route domains

Per VPC

Route	Destination
10.1.0.0/16	Local
0.0.0.0/0	tgw-xxxxxxxxx



Quick comparison: Transit Gateway and Transit VPC



Transit VPC

- Customer managed instances
- Uses VPN and virtual private gateways
- Hard to scale and manage
- Difficult to segment

Transit Gateway

- AWS native service
- Uses elastic network interfaces
- Scales horizontally
- Flexible segmentation


Transit Gateway details

Find on YouTube

NET 331: NEW LAUNCH: Introduction to Transit Gateway

Are there any reasons to use a Transit VPC?

We will cover how adding Transit Gateway makes these easier



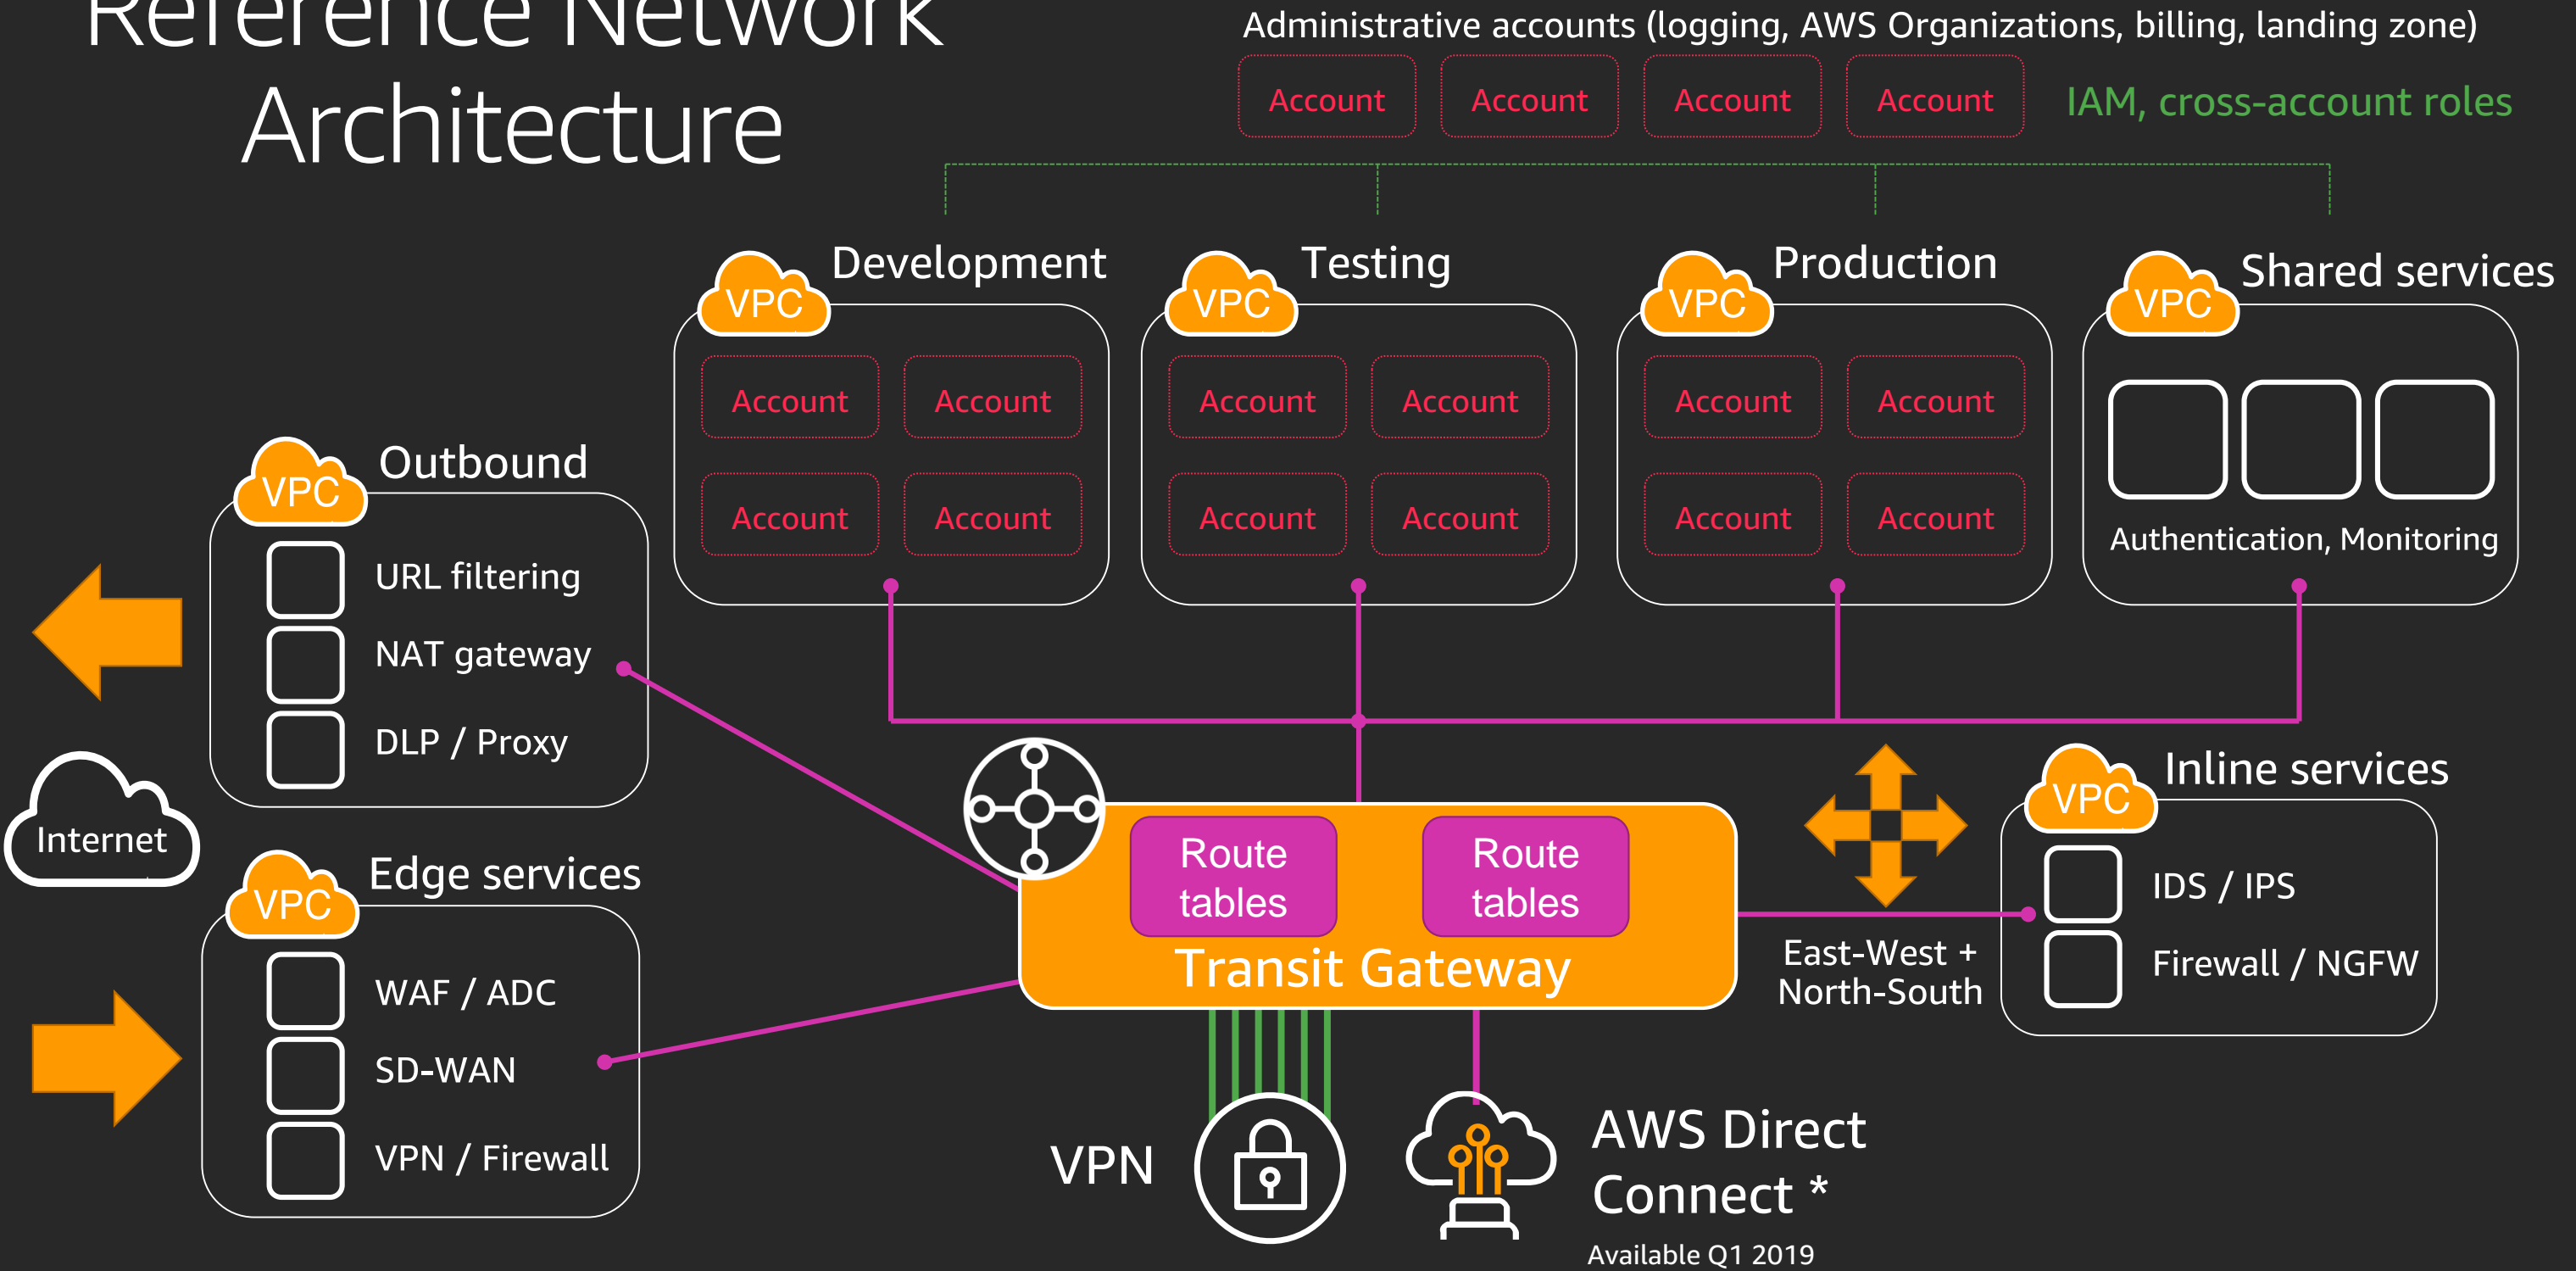
- You currently use one, and it works for you
 - Migration to Transit Gateway
- Additional visibility and monitoring
- Automated VPC networking using tagging
- You want to use additional services:
 - Security features
 - SD-WAN
 - NAT
 - Proprietary features

We're only adding things


You can use all existing options with Transit Gateway:

- VPC peering
- AWS Direct Connect
- Elastic Load Balancing
- AWS PrivateLink
- AWS CloudWatch metrics
- AWS CloudFormation
- Transit VPC

Reference Network Architecture




Architecture walk through




Account strategy



Segmentation model




Shared services




Connectivity



Network services




Multi-region options



Account
Strategy



Segmentation Model



Shared Services



Connectivity



Network
Services



Multi-Region
Options

Account strategy

Account and VPC segmentation

Larger VPCs or accounts

AWS Identity and Access Management
Strict security groups and routing
Identifying resources with tags

Policy and IAM

Smaller VPCs or accounts

Automation of infrastructure
AWS Direct Connect and VPN standards
Subnet and routing standards

Infrastructure and Networking

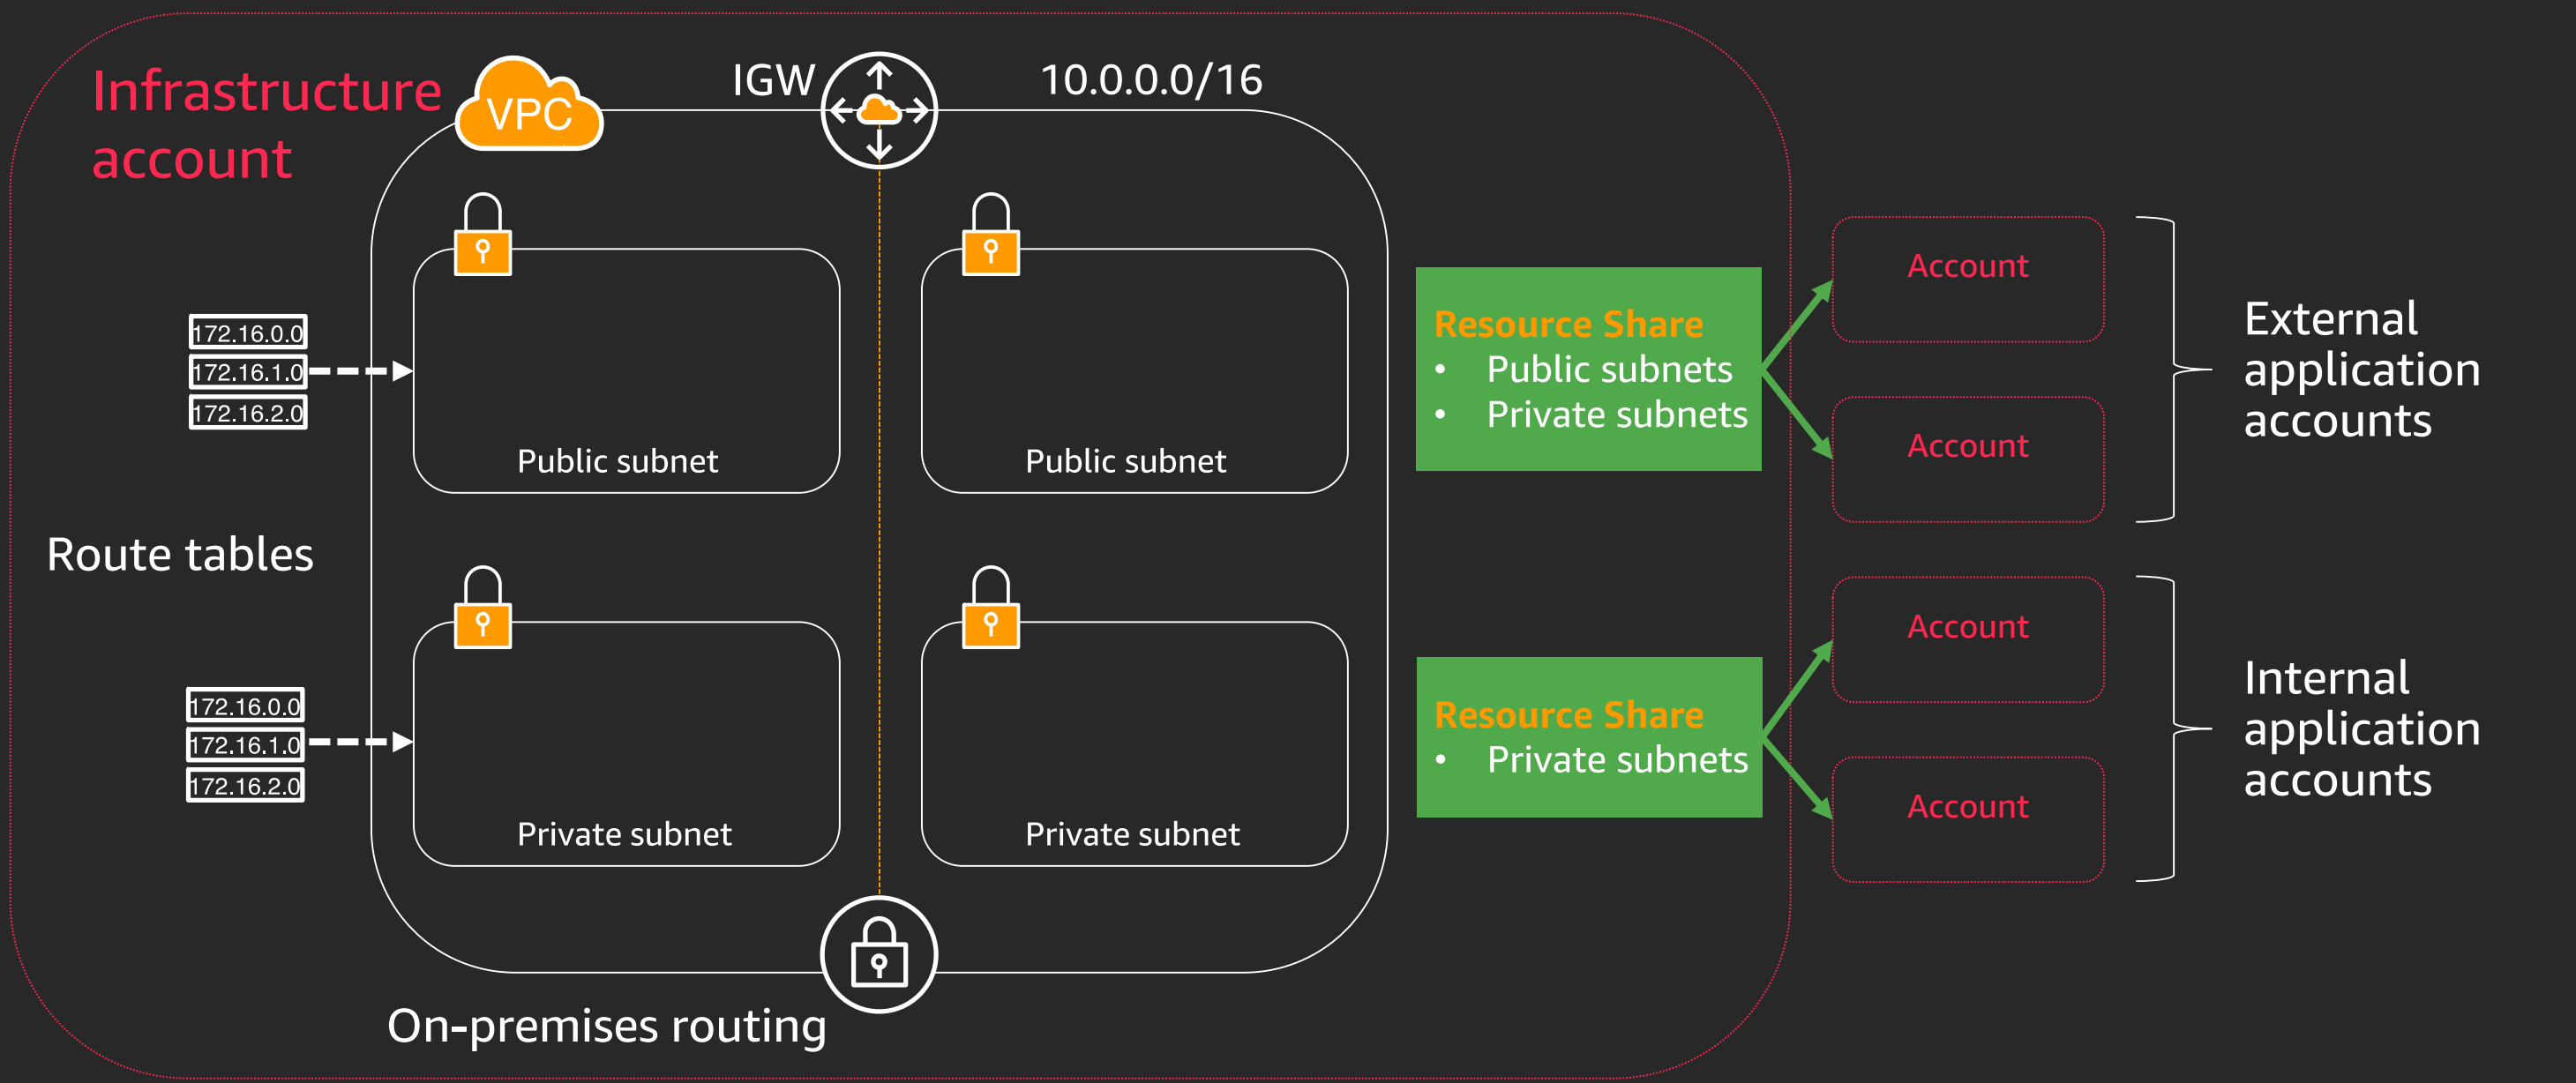
Why not both?

Provide granular account control
with centralized infrastructure

VPC Sharing and Resource Access Manager

Share subnets between accounts in an AWS Organization


New



VPC Sharing and Resource Access Manager

Account owners only see subnets and their resources

New



Account

External
accounts

Account

Internal
accounts

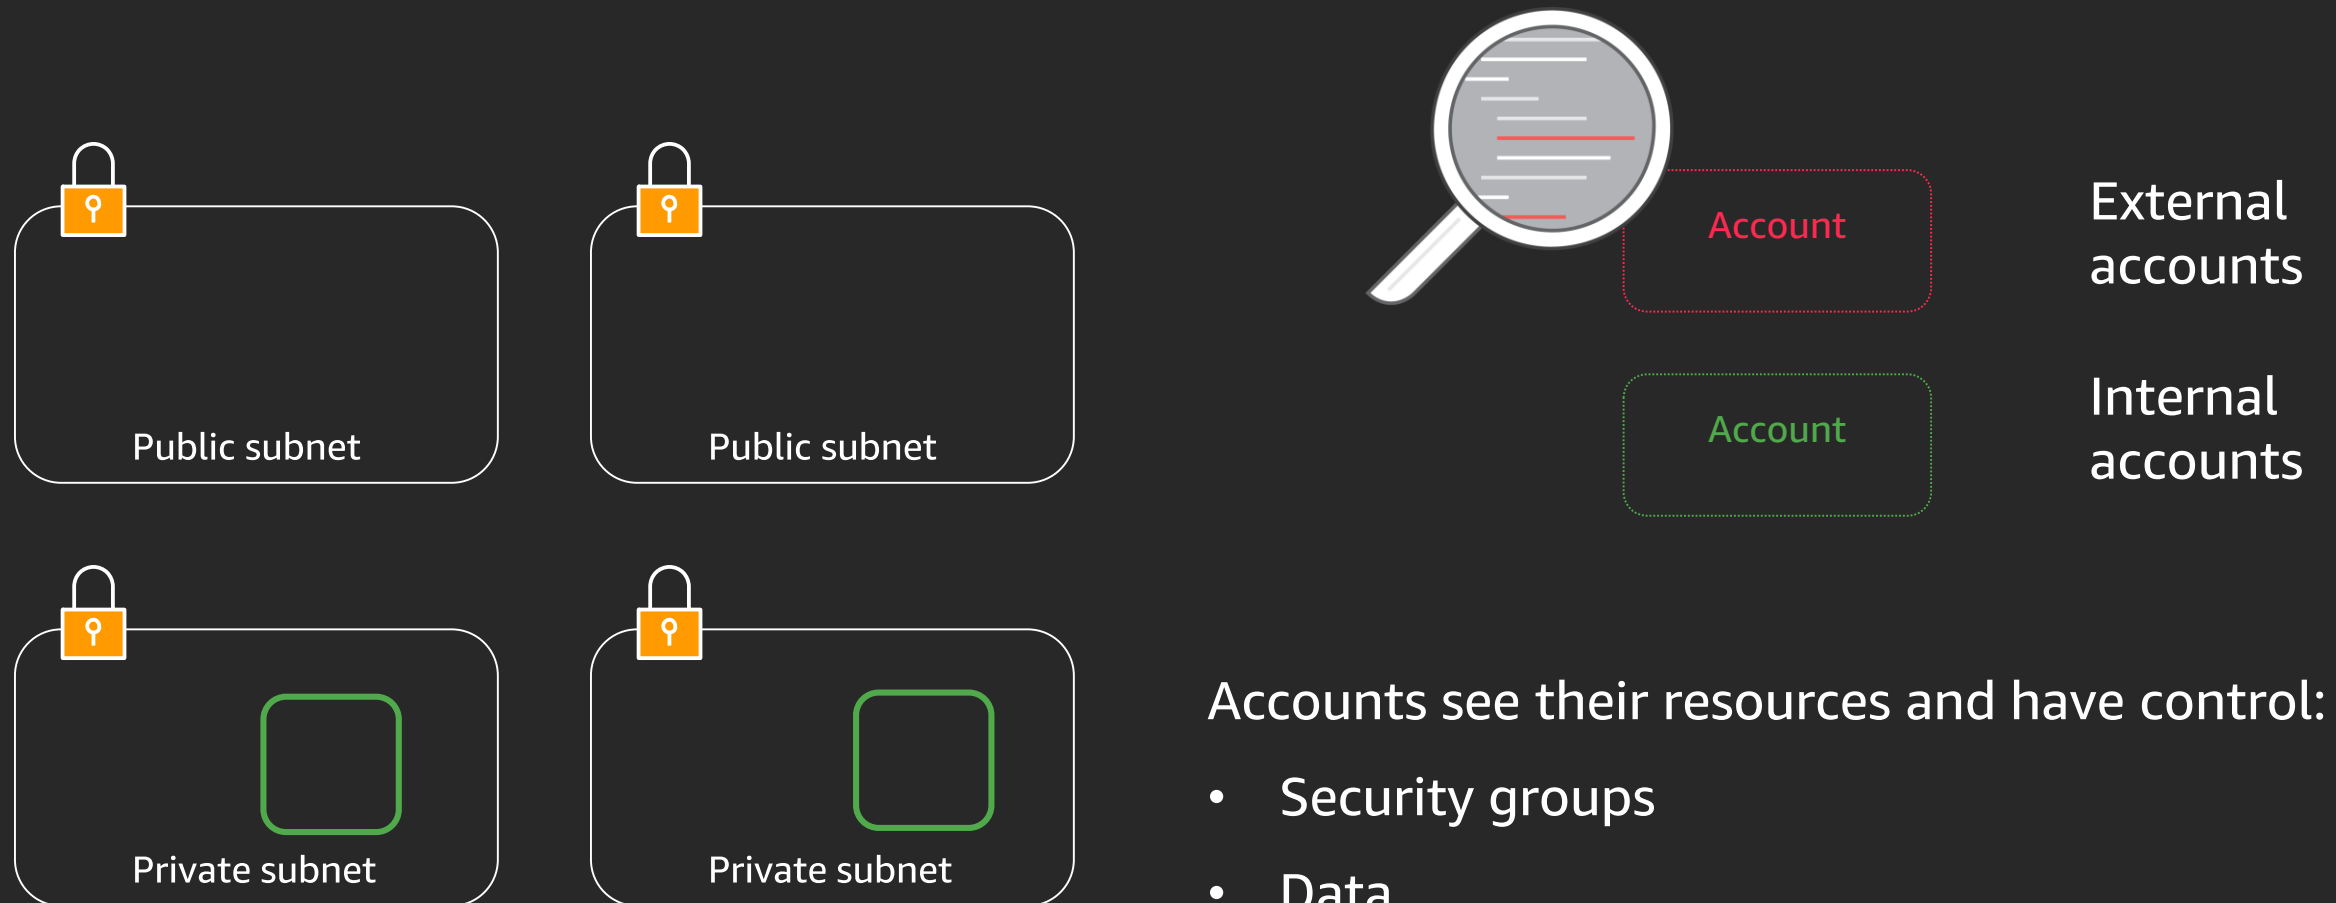
Accounts see their resources and have control:

- Security groups
- Data
- Instance details
- Account configuration

VPC Sharing and Resource Access Manager

Account owners only see subnets and their resources


New



Accounts see their resources and have control:


- Security groups
- Data
- Instance details
- Account configuration

VPC Sharing benefits




Separation of duties

- Infrastructure strictly controls routing, IP addresses, and VPC structure
- Developers own their resources, accounts, and security groups



Less unused resources

- Higher density subnets, add up to 5 additional CIDRs
- More efficient use of VPN and AWS Direct Connect



Decouple accounts and networks

- Account protection and billing without additional infrastructure
- Many accounts with fewer networks
- Avoid VPC peering charges

Other account considerations


One size does not need to fit all

- Example: production may use separate VPCs, development can use a shared VPC
- **AWS Transit Gateway can handle large amounts of VPCs if needed**

VPC Sharing works within an AWS Organization

VPC Sharing doesn't restrict resource utilization


- NAT gateways, VPN, subnet address space, and security groups have shared limits
- VPC Sharing doesn't change any VPC limits, only account limits
- Give highly scalable services like AWS Lambda dedicated IP space



Account
Strategy



Segmentation Model



Shared Services



Connectivity




Network
Services



Multi-Region
Options

Segmentation



Segmentation: Decision inputs

Relationship between accounts, VPCs, and tenants?

- Do accounts and tenants trust each other?
- Is the current network segmentation intentional or a side effect?

Who owns security and networking?

- Each team or a centralized team?

Compliance and governance requirements?

- Scope can be reduced at an account or a VPC level

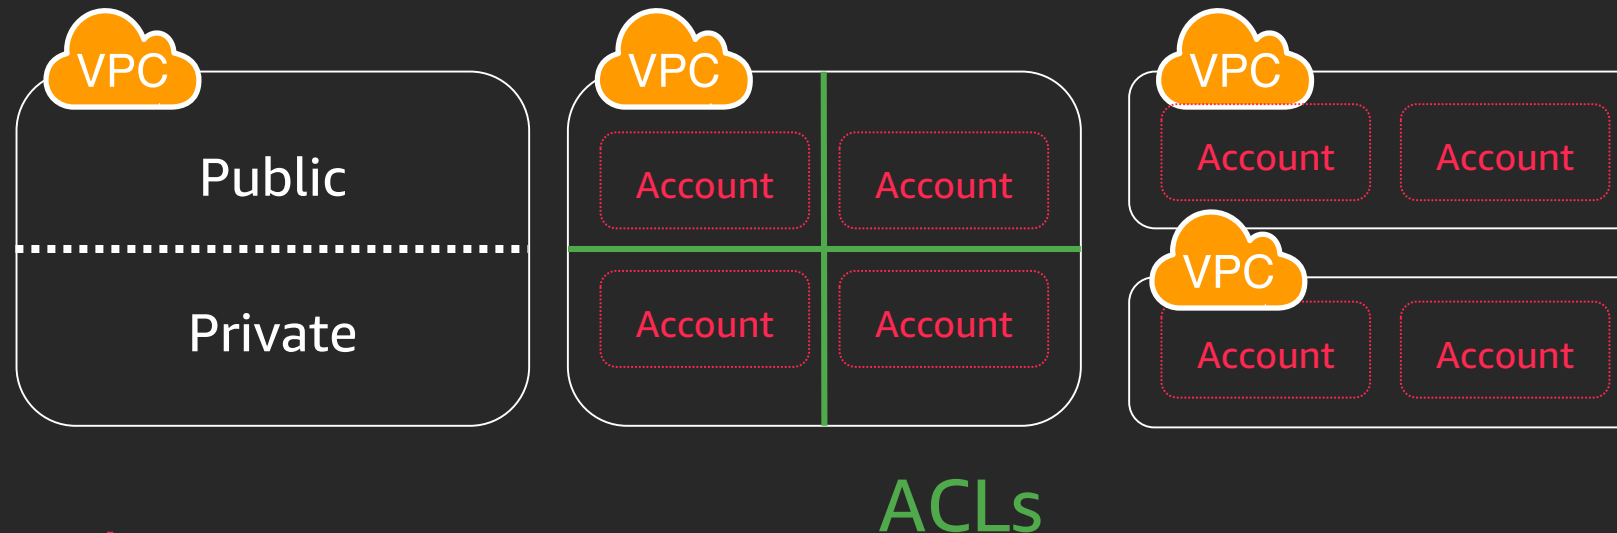
Segmentation options: Layers

Inside the account

- IAM users and roles
- Security groups

At the VPC

- Route tables
- Network ACLs
- Separate VPCs



Baseline security

Tenant
configuration

IAM: Control actions and privileges inside the account between users and role

Security groups: Whitelist ports, protocols, and other security groups for network access

Tenant and infrastructure

Shared Security line

Network security

Infrastructure
configuration

Route tables: Route table policy defines what VPC resources can access on the network

Network ACLs: Fence off access between specific subnets, ports, or destinations.

Separate VPCs: Full separation from other tenants.

Segmentation options: Layers

Inside the account

- IAM users and roles
- Security groups

At the VPC

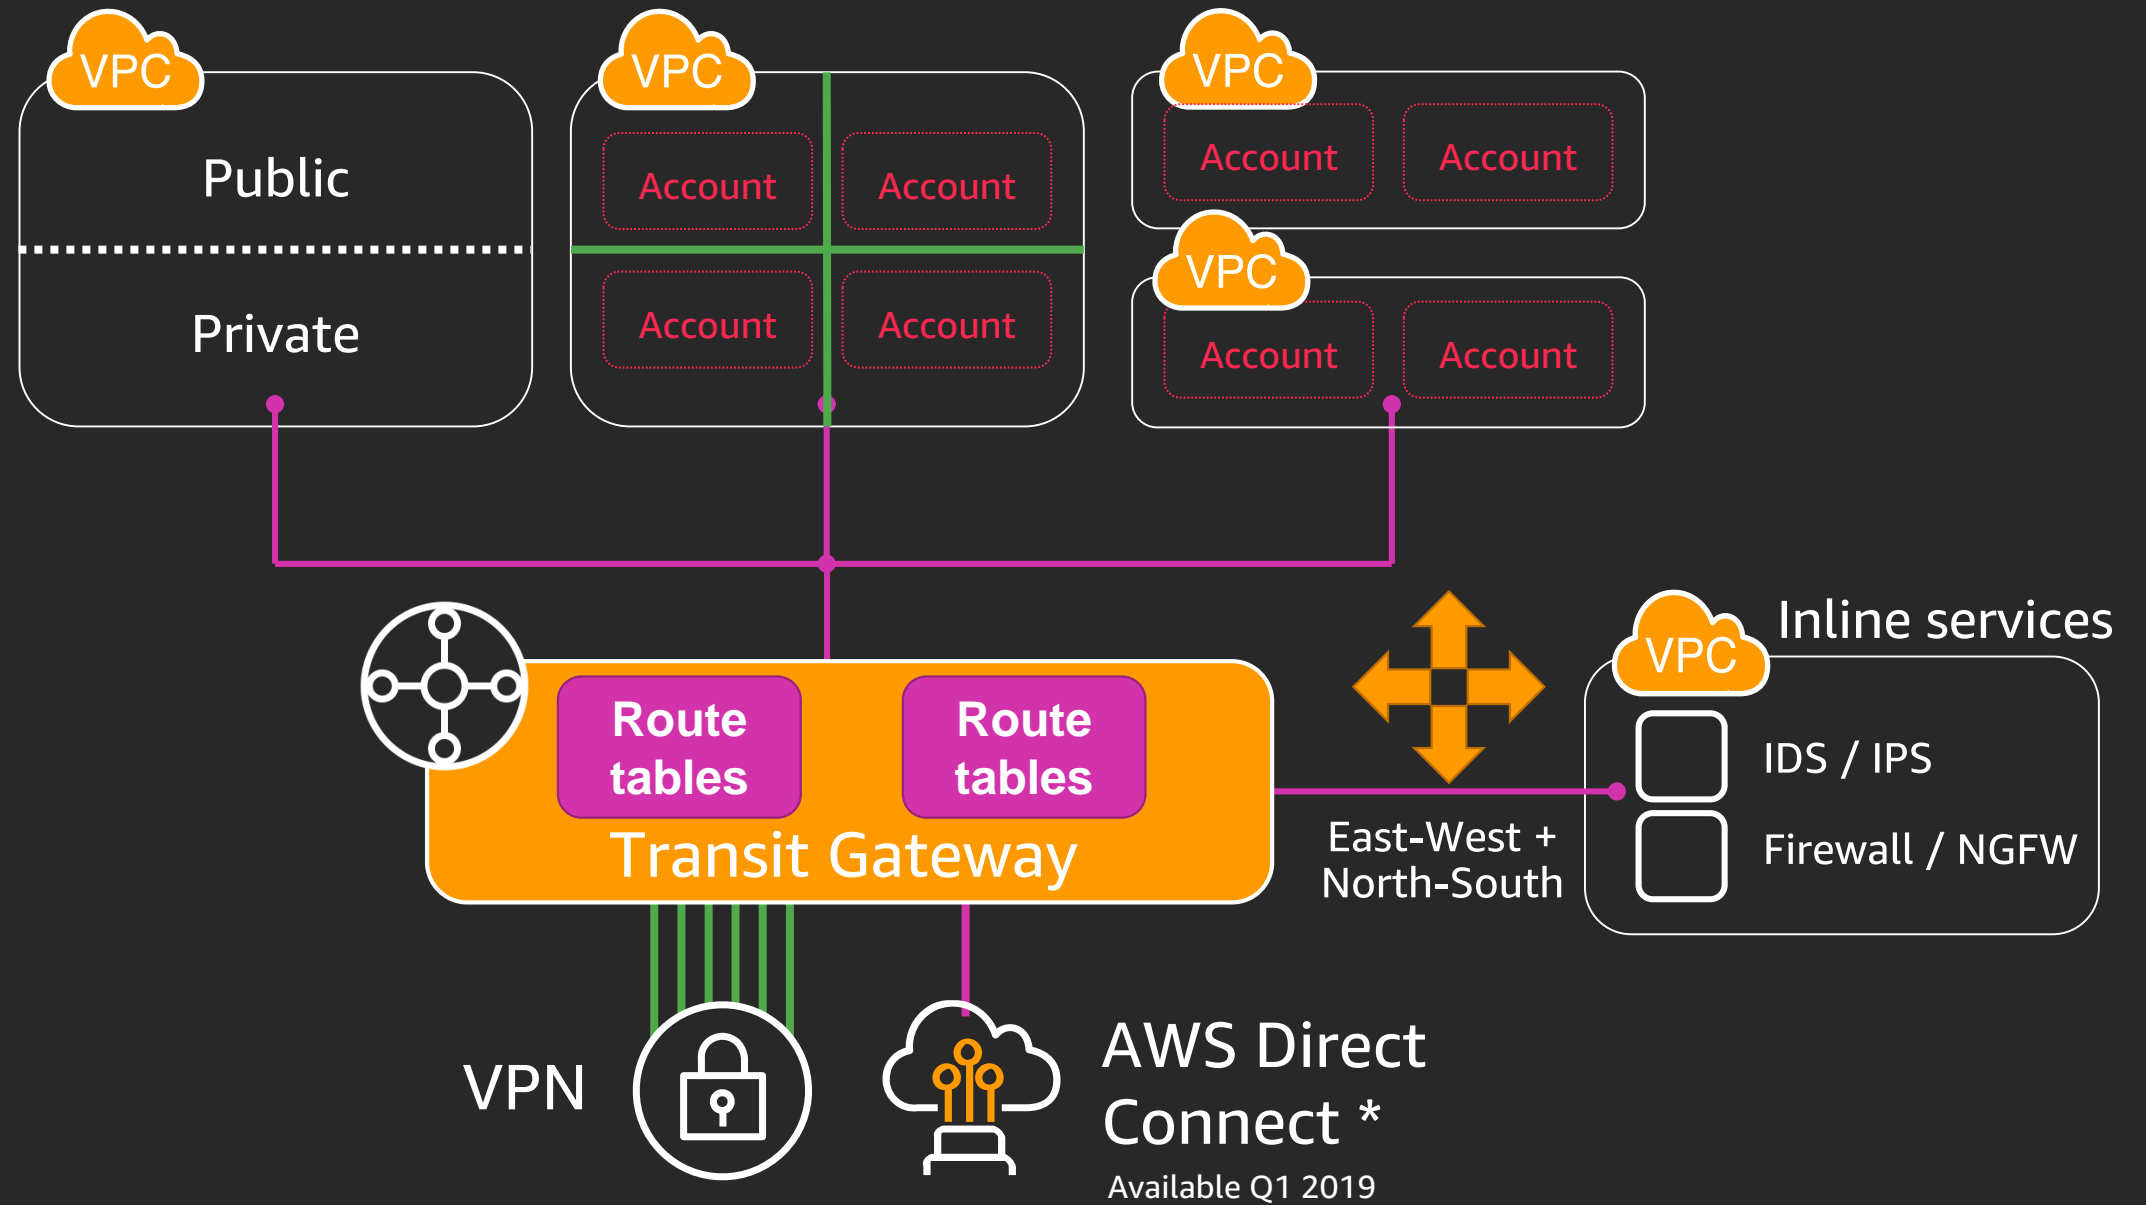
- Route tables
- Network ACLs
- Separate VPCs

Transit Gateway

- Route tables

Security services

- Firewalls
- Proxies
- Intrusion Detection / Prevention

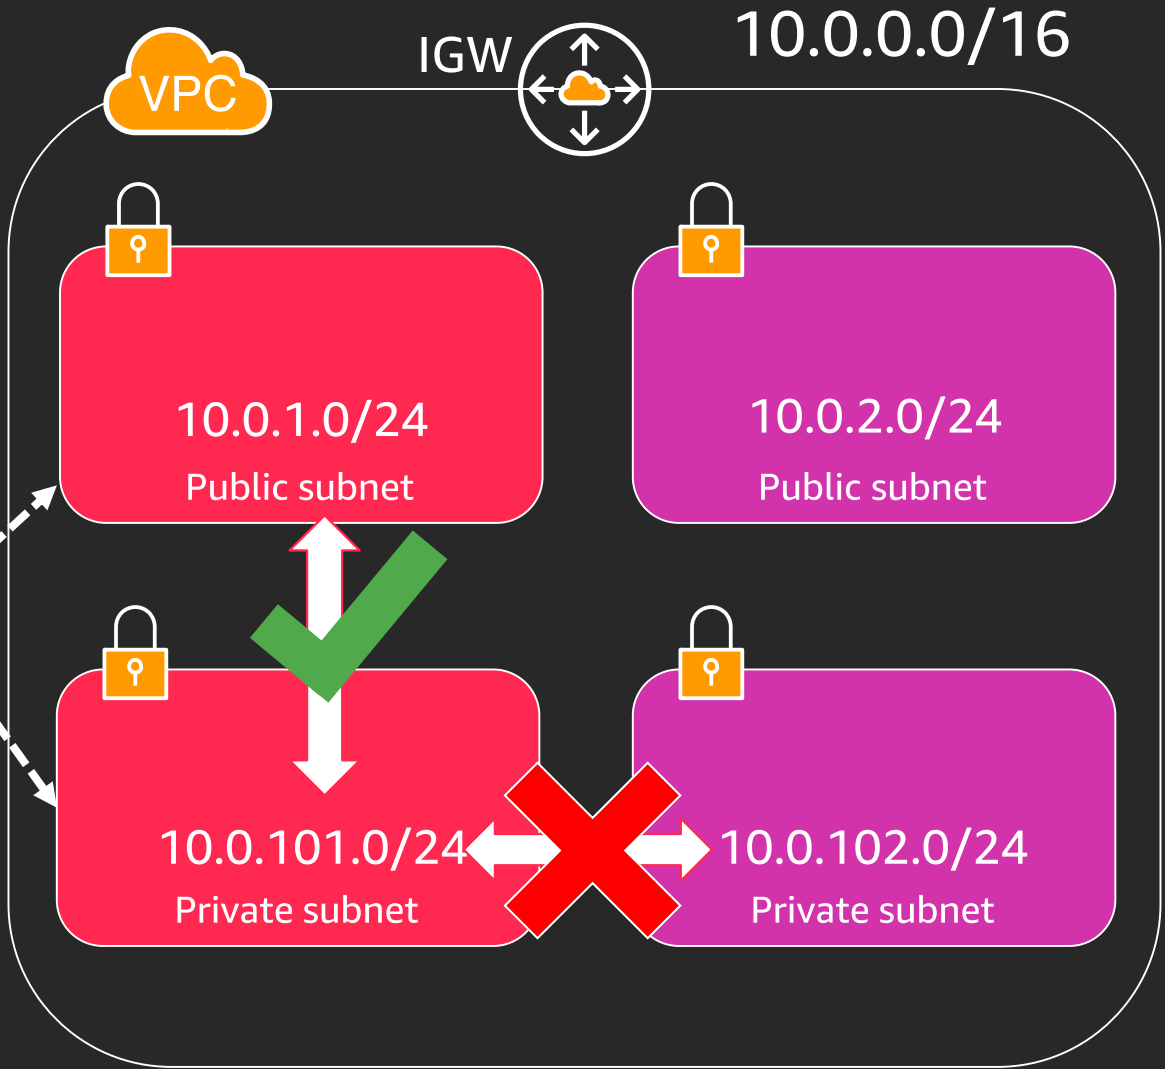


Segmentation in a Shared VPC with network ACLs

Mimic behavior of a single VPC:

- Allow traffic between subnets
- Deny inbound traffic from other tenants

Inbound network ACL		
#	Source	Action
100	10.0.1.0/24	ALLOW
101	10.0.101.0/24	ALLOW
200	10.0.0.0/16	DENY
300	0.0.0.0/0	ALLOW



Resource share

- Public subnets
- Private subnets

Account

Account

Resource share

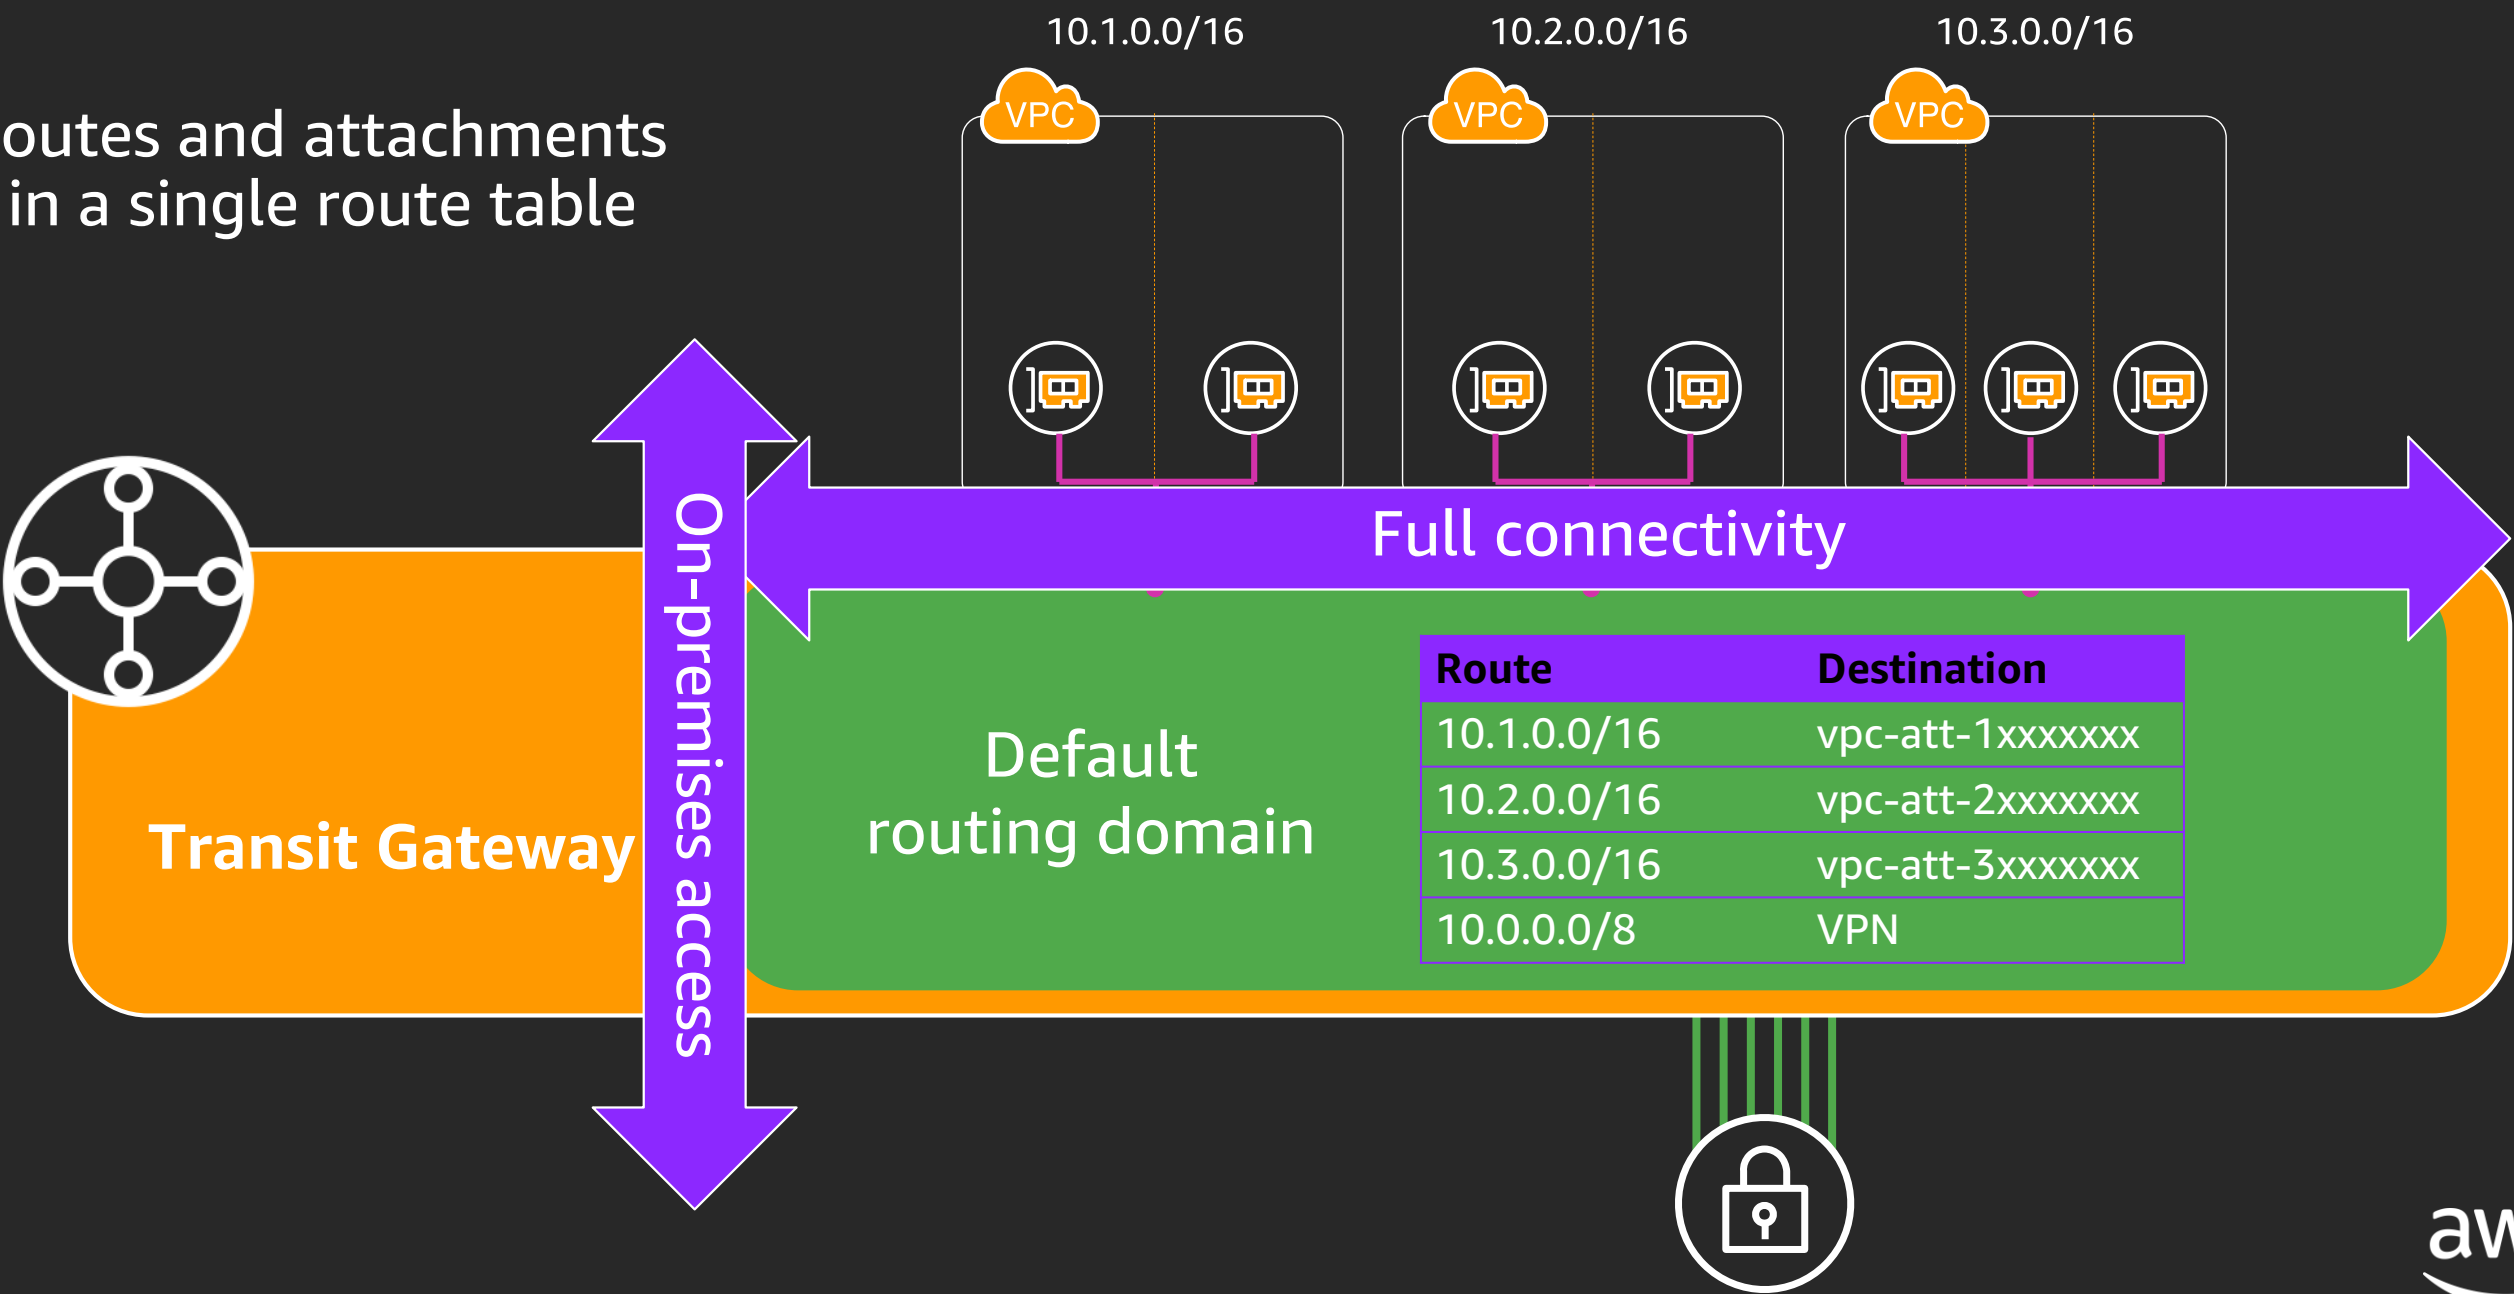
- Public subnets
- Private subnets

Account

Account

Flat: Transit Gateway route domains

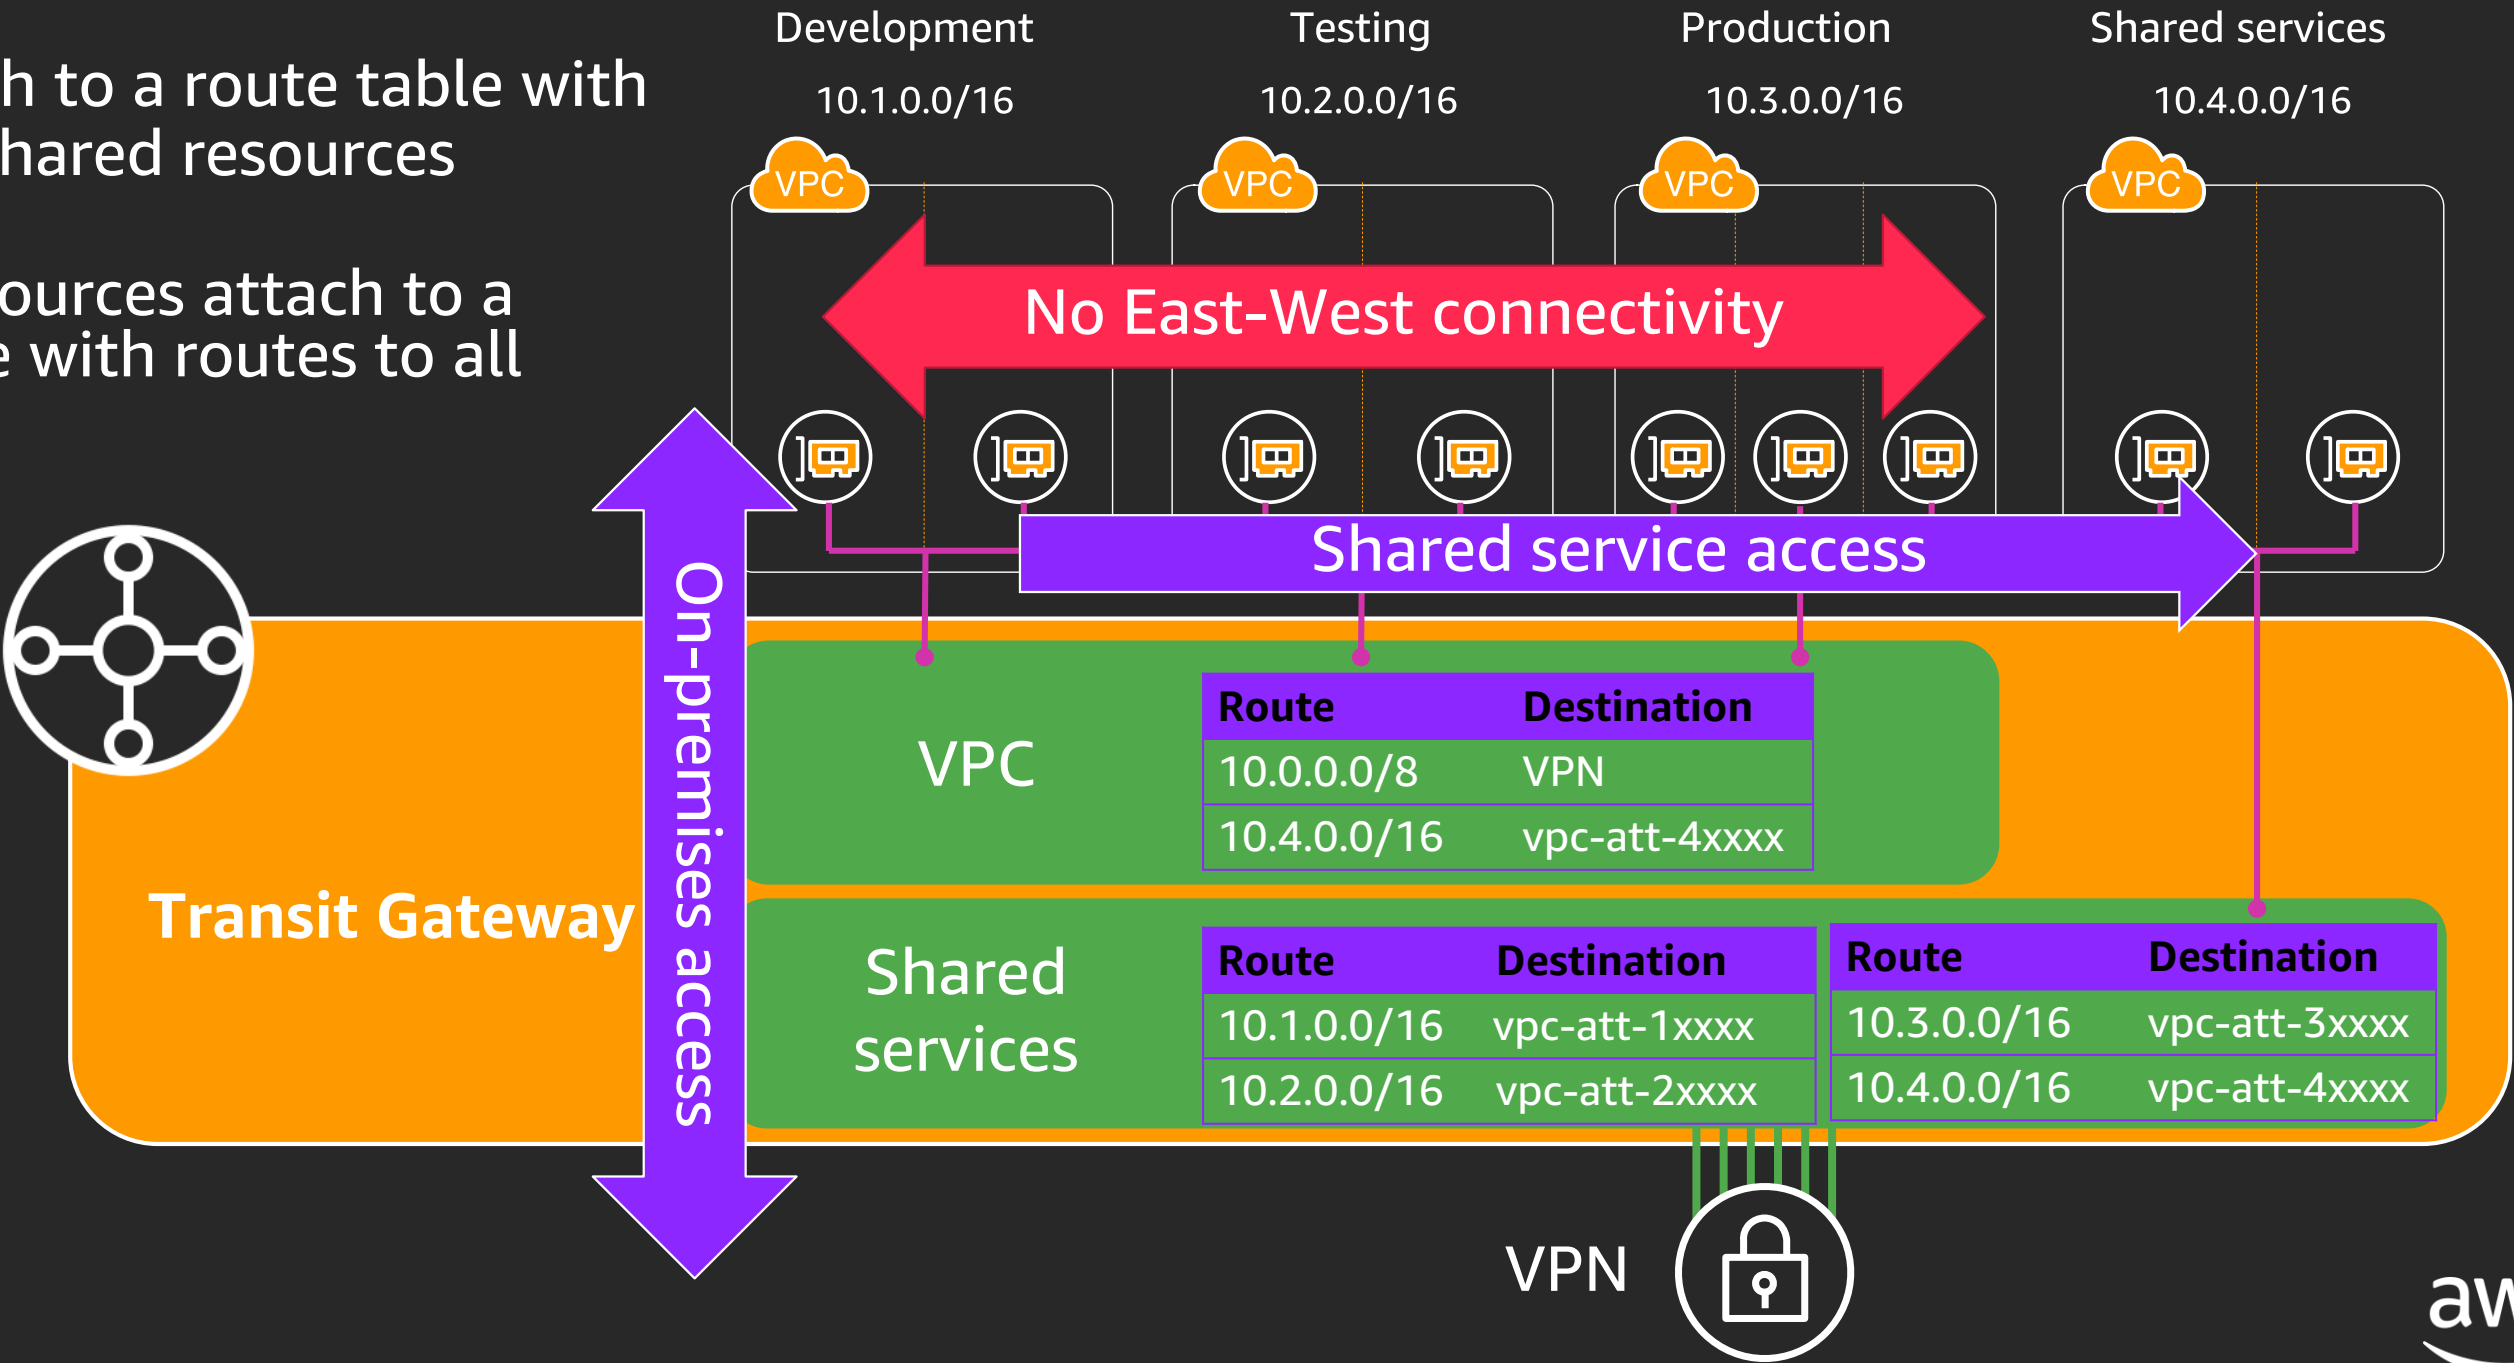
All routes and attachments are in a single route table



Isolated: Transit Gateway route domains

VPCs attach to a route table with routes to shared resources

Shared resources attach to a route table with routes to all resources



Segmentation considerations: Where to start

Security groups and IAM are effective and proven

- Encourage IAM and security group use and monitor security configuration

Shared VPCs


- Tenants should limit access from the internet and other tenants
- VPCs using VPC peering are likely to benefit from Shared VPCs
- Design around resource and limit contention

Separate VPCs

- Often the best security decision is the simplest. Separate VPCs are simple.
- Use separate VPCs for strong network segmentation and resource isolation
- Transit Gateway removes the scaling issues with many VPCs (peering, VPN, routes)

Transit Gateway route tables define multi-VPC policy


- Consider isolating environments (dev and prod) and allow access to shared resources



Account
Strategy



Segmentation Model



Shared Services



Connectivity




Network
Services




Multi-Region
Options

Shared services




Shared services connectivity options

VPC peering




- One-to-one connectivity
- Scales to 100 VPCs
- Security groups across VPCs
- Inter-region peering

AWS PrivateLink




- One-to-many connectivity
- Highly scalable
- Supports overlapping CIDRs
- Uses Elastic Load Balancing
- Load balancing and hourly endpoint costs

Transit VPC



- Shared services as a spoke
- Bandwidth constrained
- Complex management
- Instance and licensing costs


AWS Transit Gateway



- Many-to-many or one-to-many with route tables
- Highly scalable
- Hourly per AZ endpoint costs


Shared services connectivity options at scale

VPC Peering




- 1-to-1 connectivity
- Scales to 100 VPCs
- Security groups across VPCs
- Inter-region peering

AWS PrivateLink




- One-to-many connectivity
- Highly scalable
- Supports overlapping CIDRs
- Uses Elastic Load Balancing
- Load balancing and hourly endpoint costs

Transit VPC



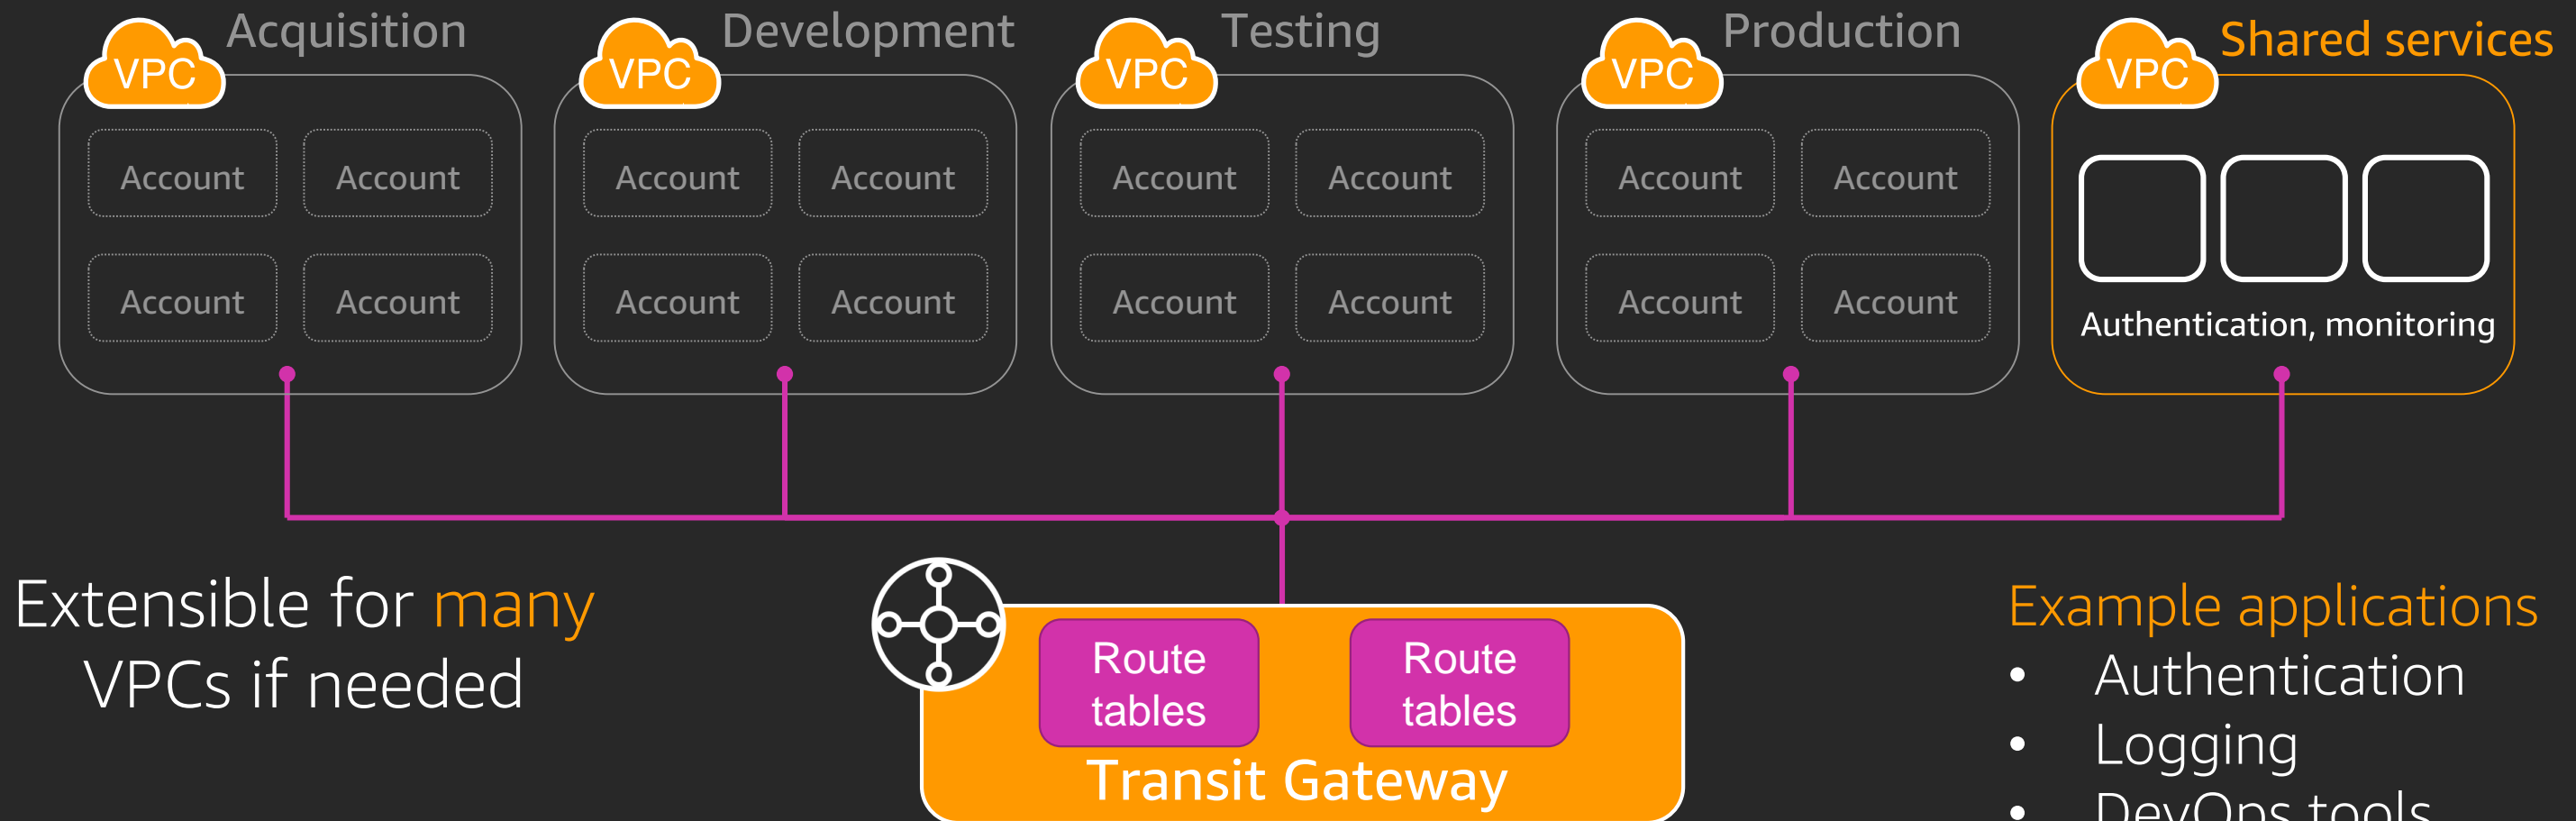
- Shared services as a spoke
- Bandwidth restricted
- Complex management
- Instance and licensing costs

AWS Transit Gateway



- Many-to-many or one-to-many with route tables
- Highly scalable
- Hourly per AZ endpoint costs


Shared services with Transit Gateway



Works with flat or isolated segmentation


Using Transit Gateway and PrivateLink

AWS PrivateLink



- One-to-many connectivity
- Highly scalable
- Supports overlapping CIDRs
- Uses Elastic Load Balancing
- Load balancing and hourly endpoint costs

AWS Transit Gateway



- Many-to-Many or one-to-many with route tables
- Highly scalable
- Hourly per AZ endpoint costs

Scope: Application shared services

Trust model: No mutual trust

Dependencies: Load balancing and application architecture

Scale: Thousands of spoke VPCs

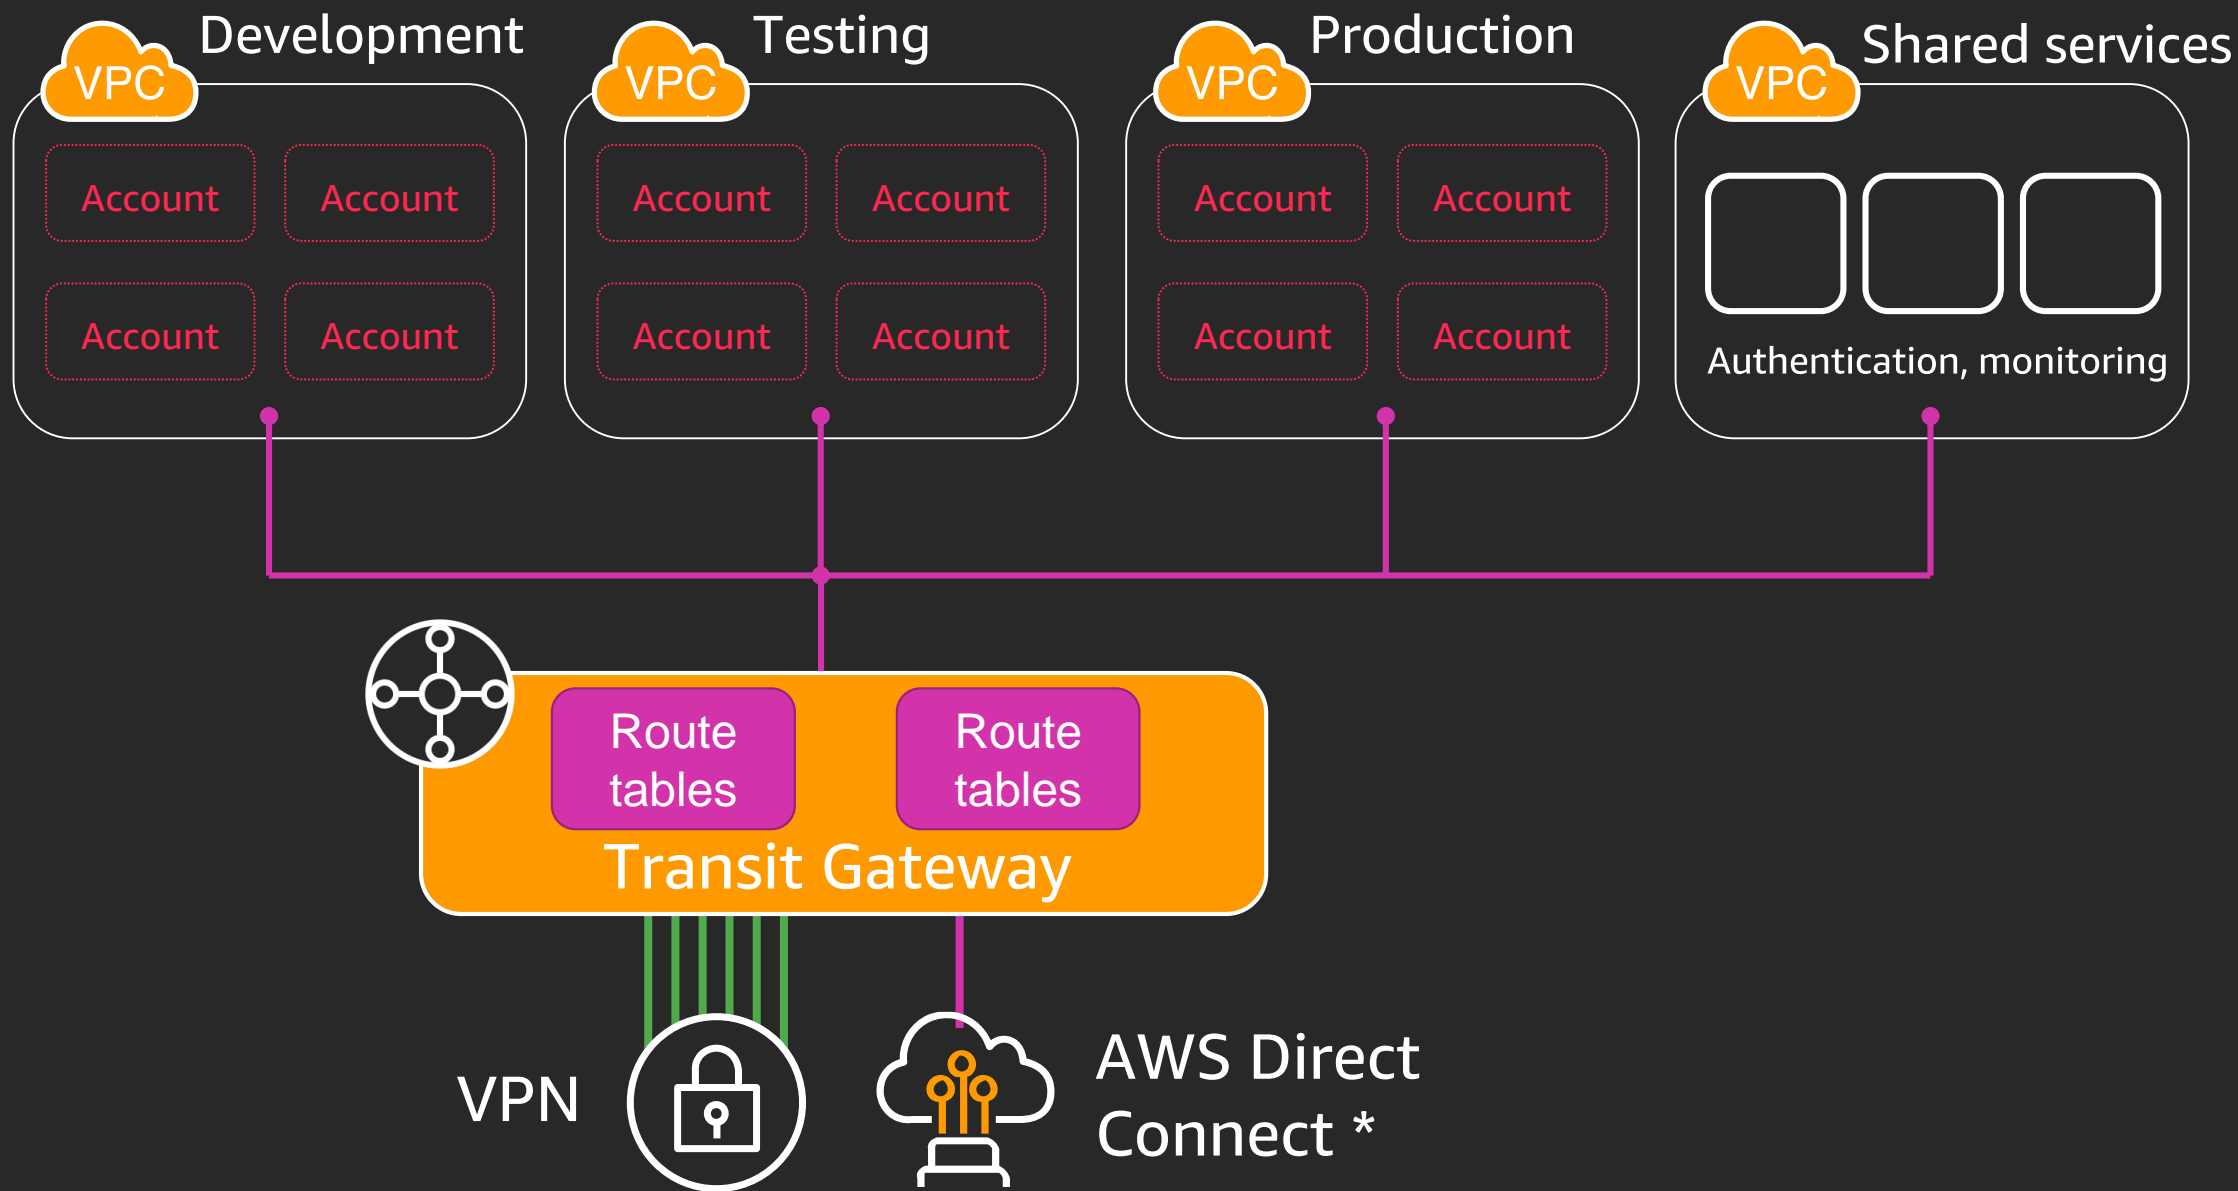




Scope: Network shared services to many VPCs

Trust model: Per VPC trust, centralized control

Dependencies: Centralized control of the Transit Gateway


Scale: Thousands of spoke VPCs

Connecting on-premises




Connecting to on-premises

Virtual Private Gateway VPN




- Per VPC
- 1.25 gbps per tunnel
- Encrypted in transit

AWS Direct Connect




- Per VPC (50 per port)
- Multiple VPCs with Direct Connect gateway
- No bandwidth restraint

Amazon EC2 customer VPN



- Per VPC or multiple (Transit VPC)
- Bandwidths vary by instance type
- AWS Marketplace options
- Scalability is generally limited by management complexity


AWS Transit Gateway VPN



- Multiple VPCs
- Add VPN connection as needed
- 1.25 gbps per tunnel
- Roadmap: AWS Direct Connect


Connecting to On-premises at Scale

Virtual Private Gateway VPN




- Per VPC
- 1.25 gbps per tunnel
- Encrypted in transit

AWS Direct Connect




- Per VPC (50 per port)
- Multiple VPCs with Direct Connect gateway
- No bandwidth restraint

Amazon EC2 Customer VPN



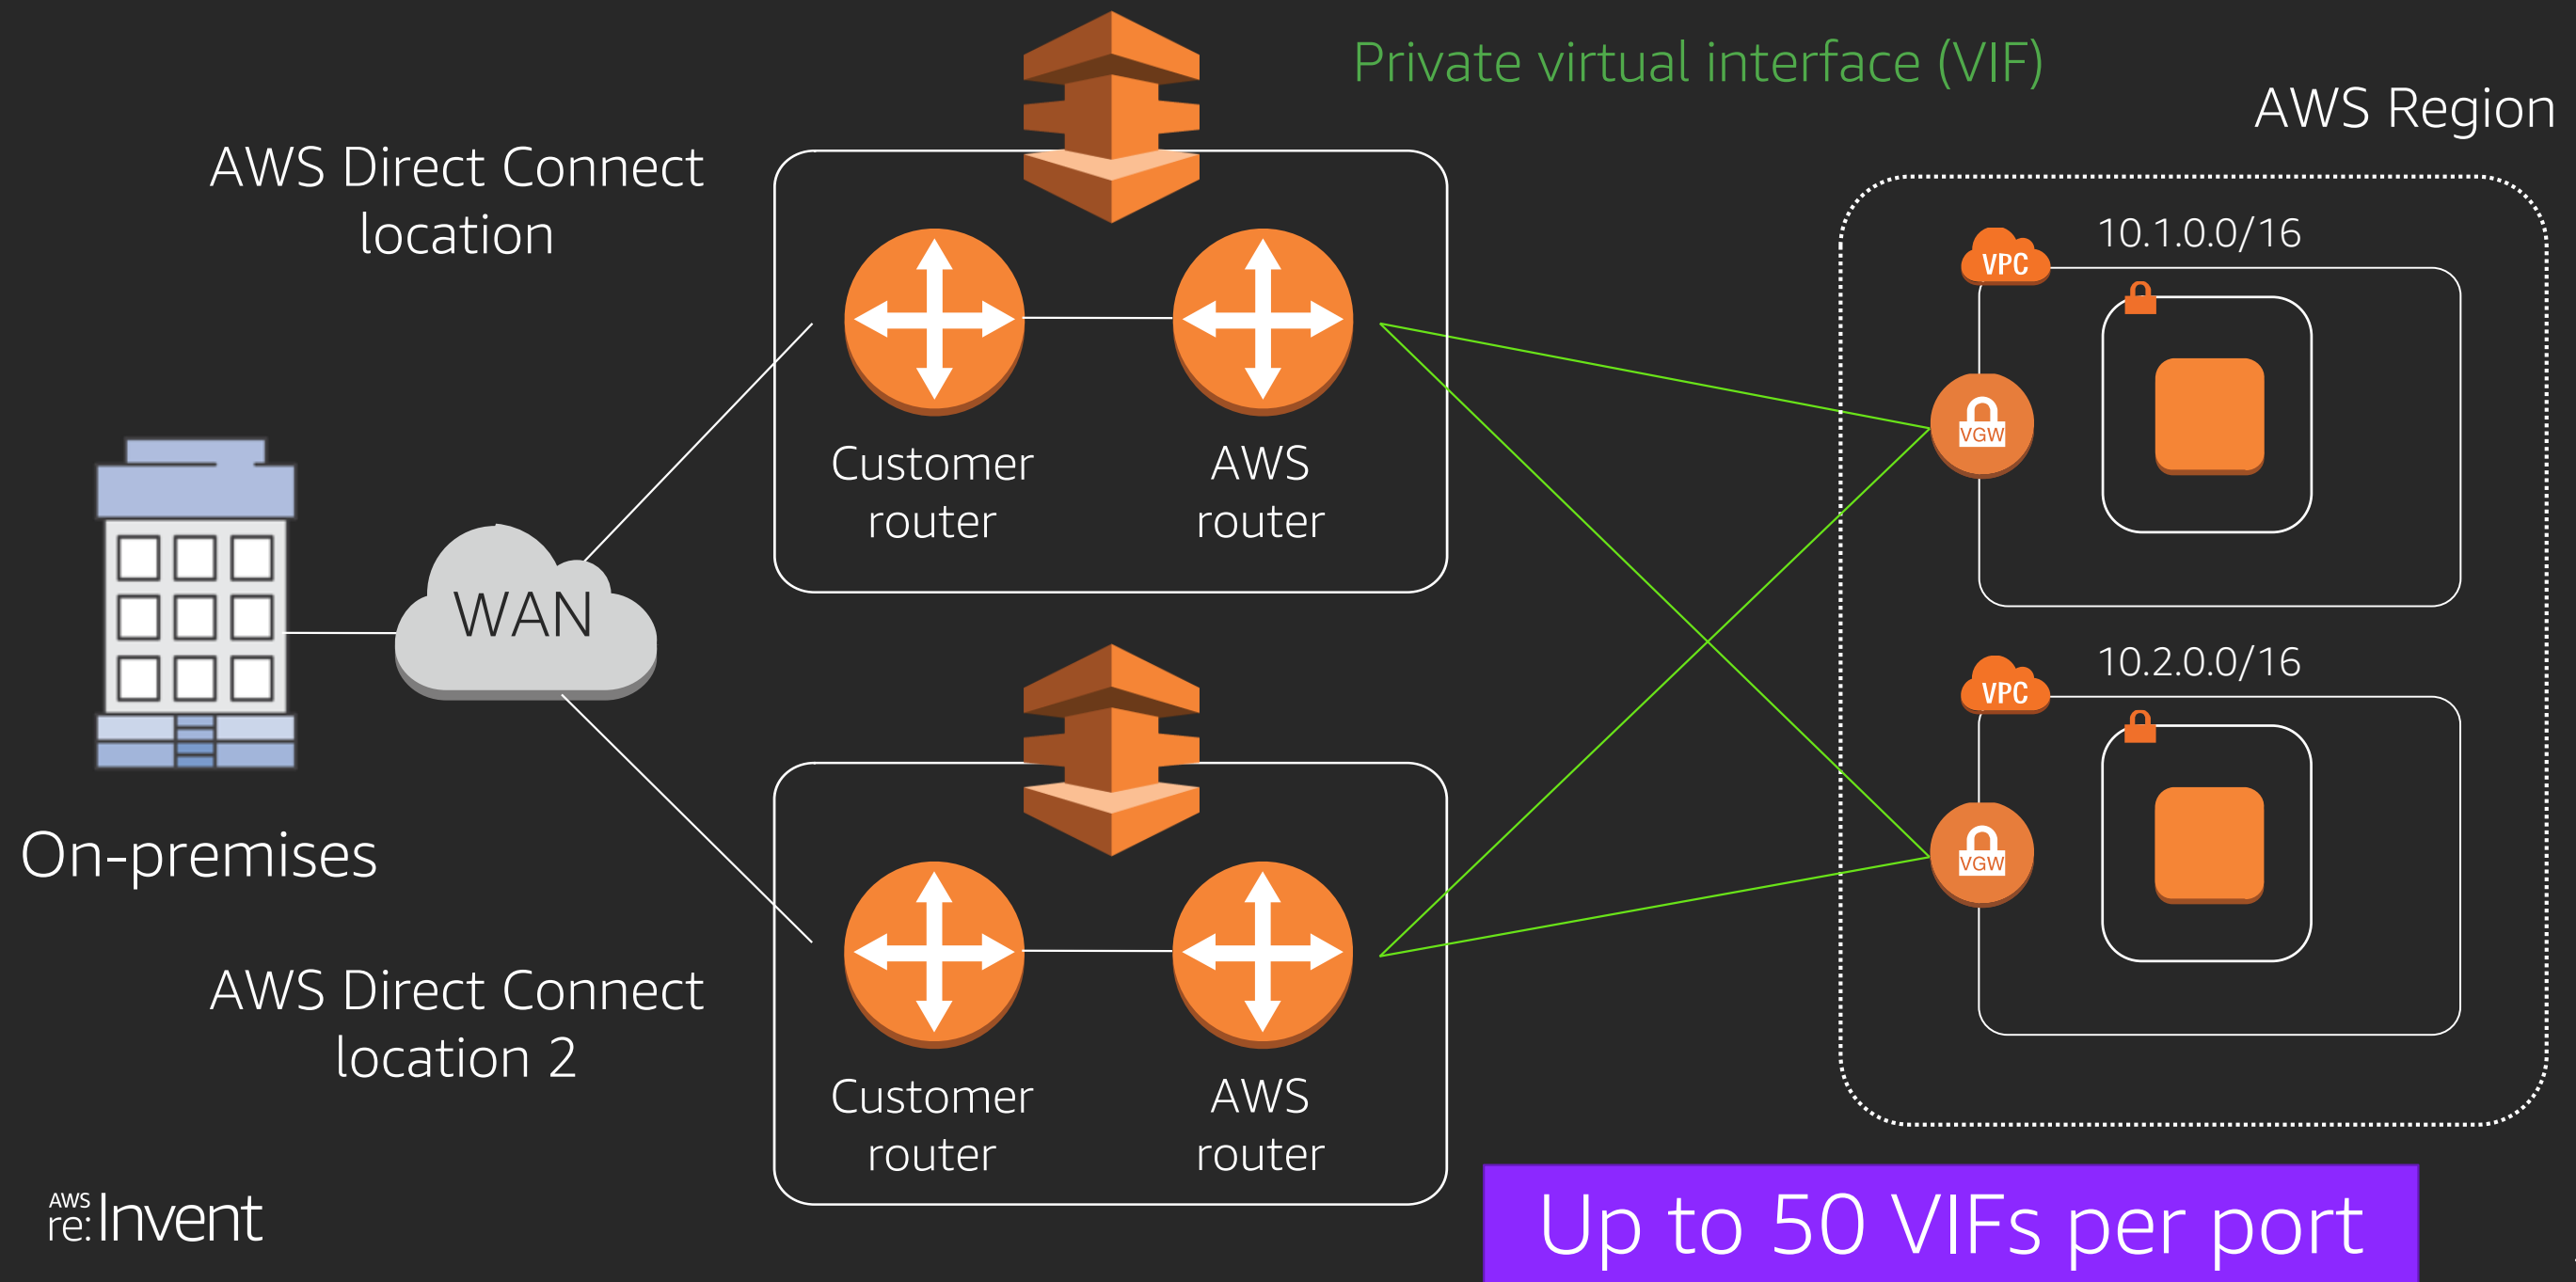
- Per VPC or multiple (Transit VPC)
- Bandwidths vary by instance type
- AWS Marketplace options
- Scalability is generally limited by management complexity

AWS Transit Gateway VPN

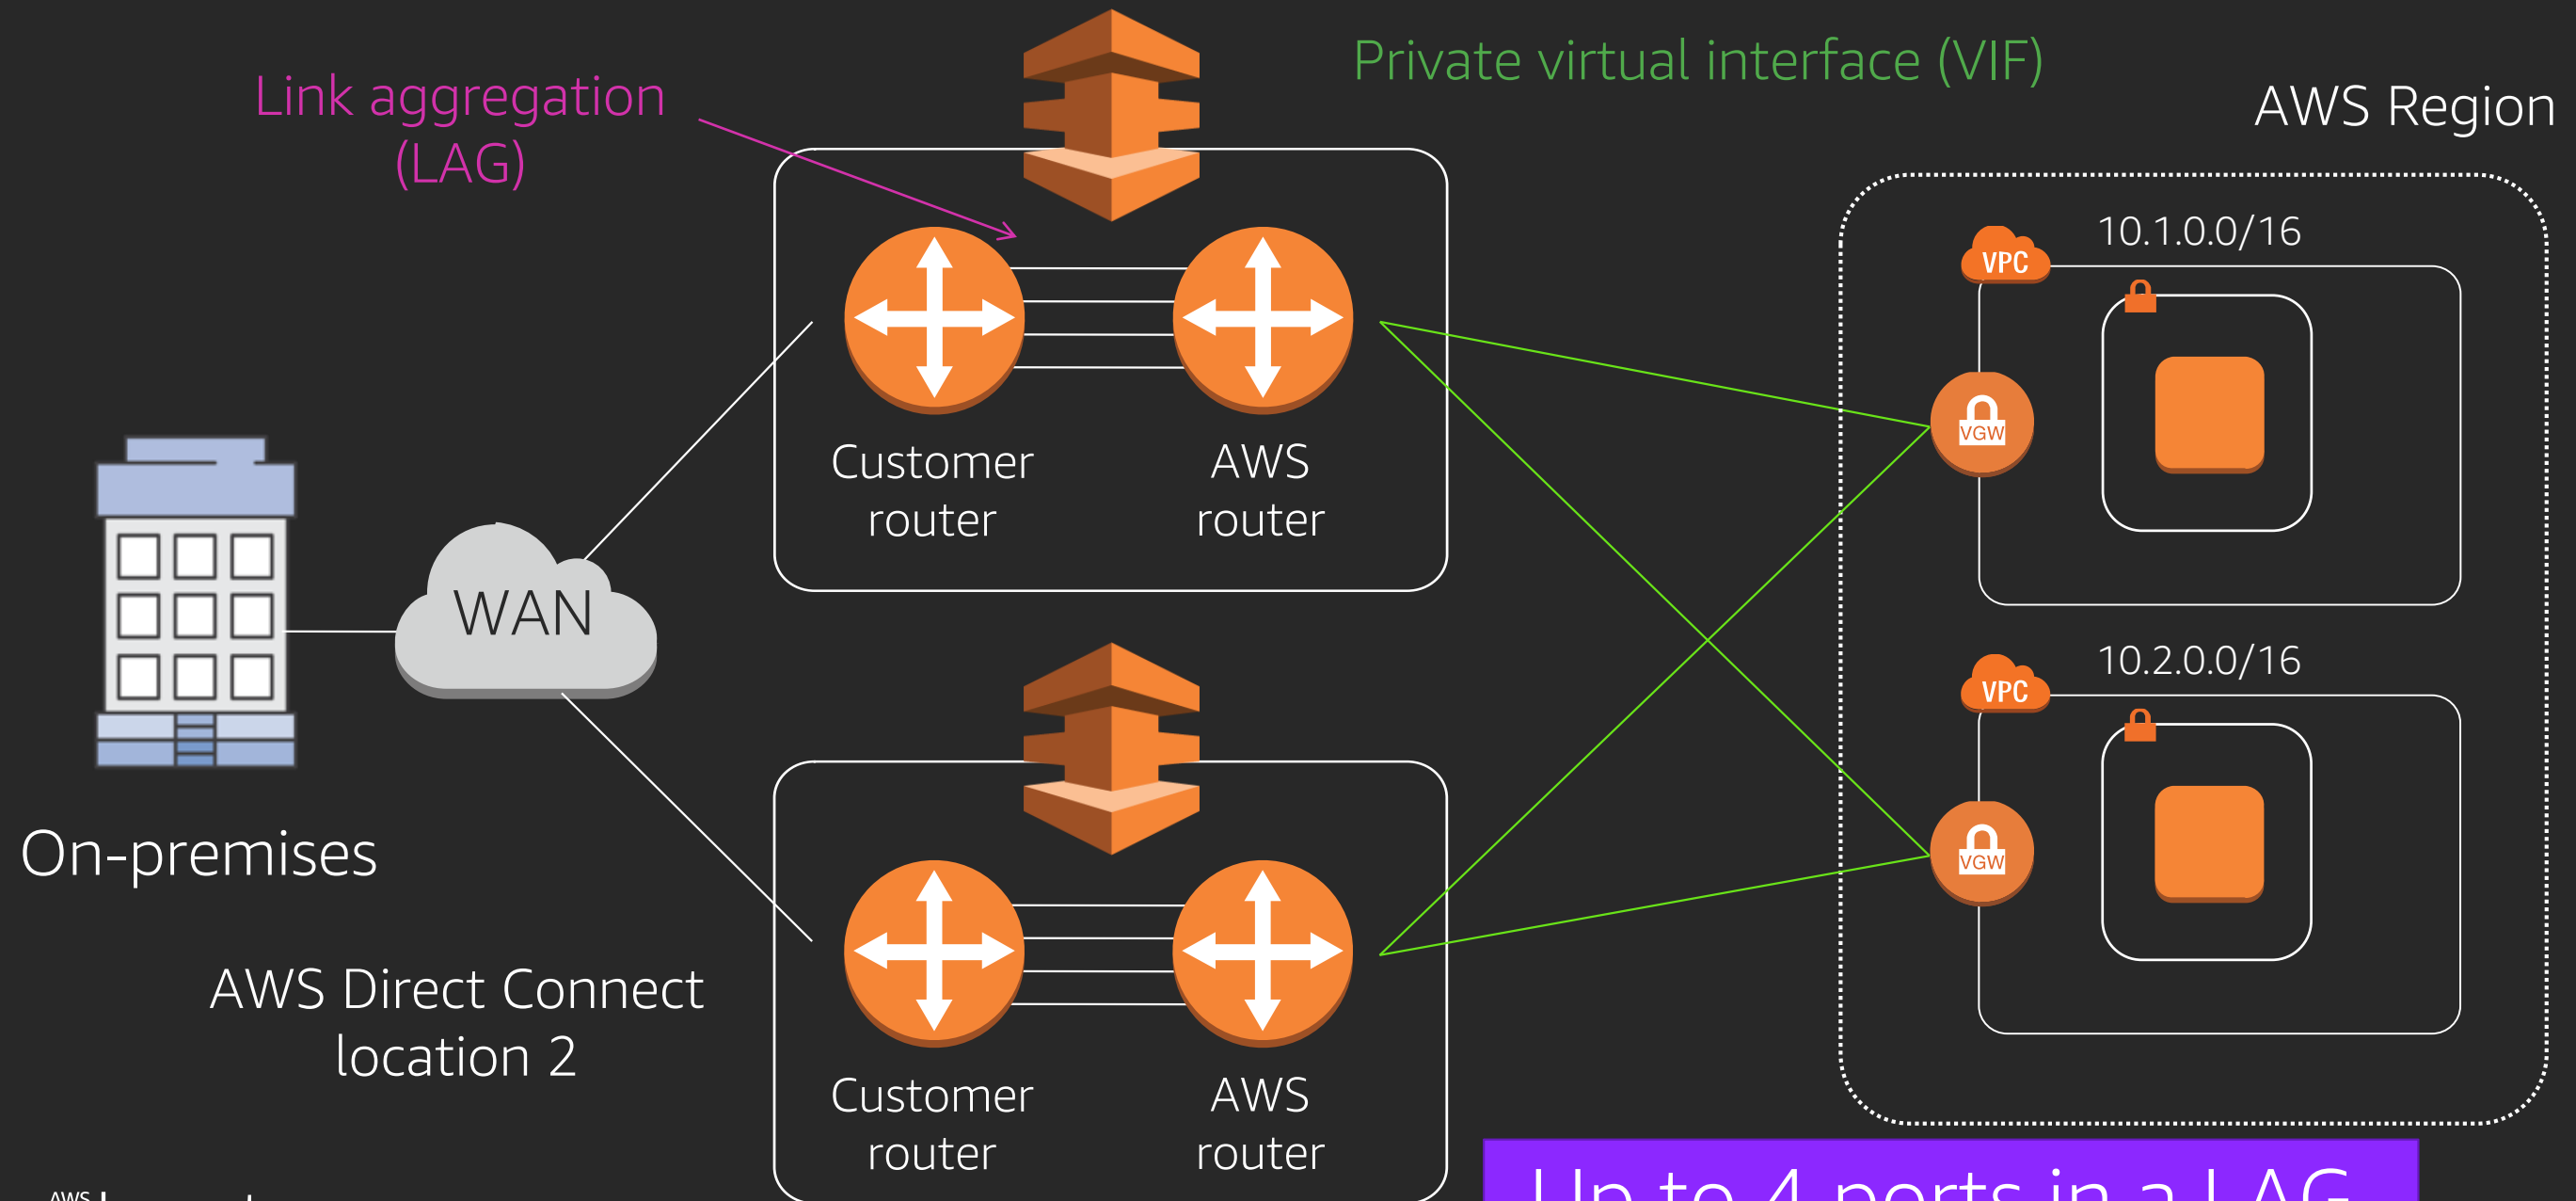


- Multiple VPCs
- Add VPN connection as needed
- 1.25 gbps per tunnel
- Roadmap: AWS Direct Connect

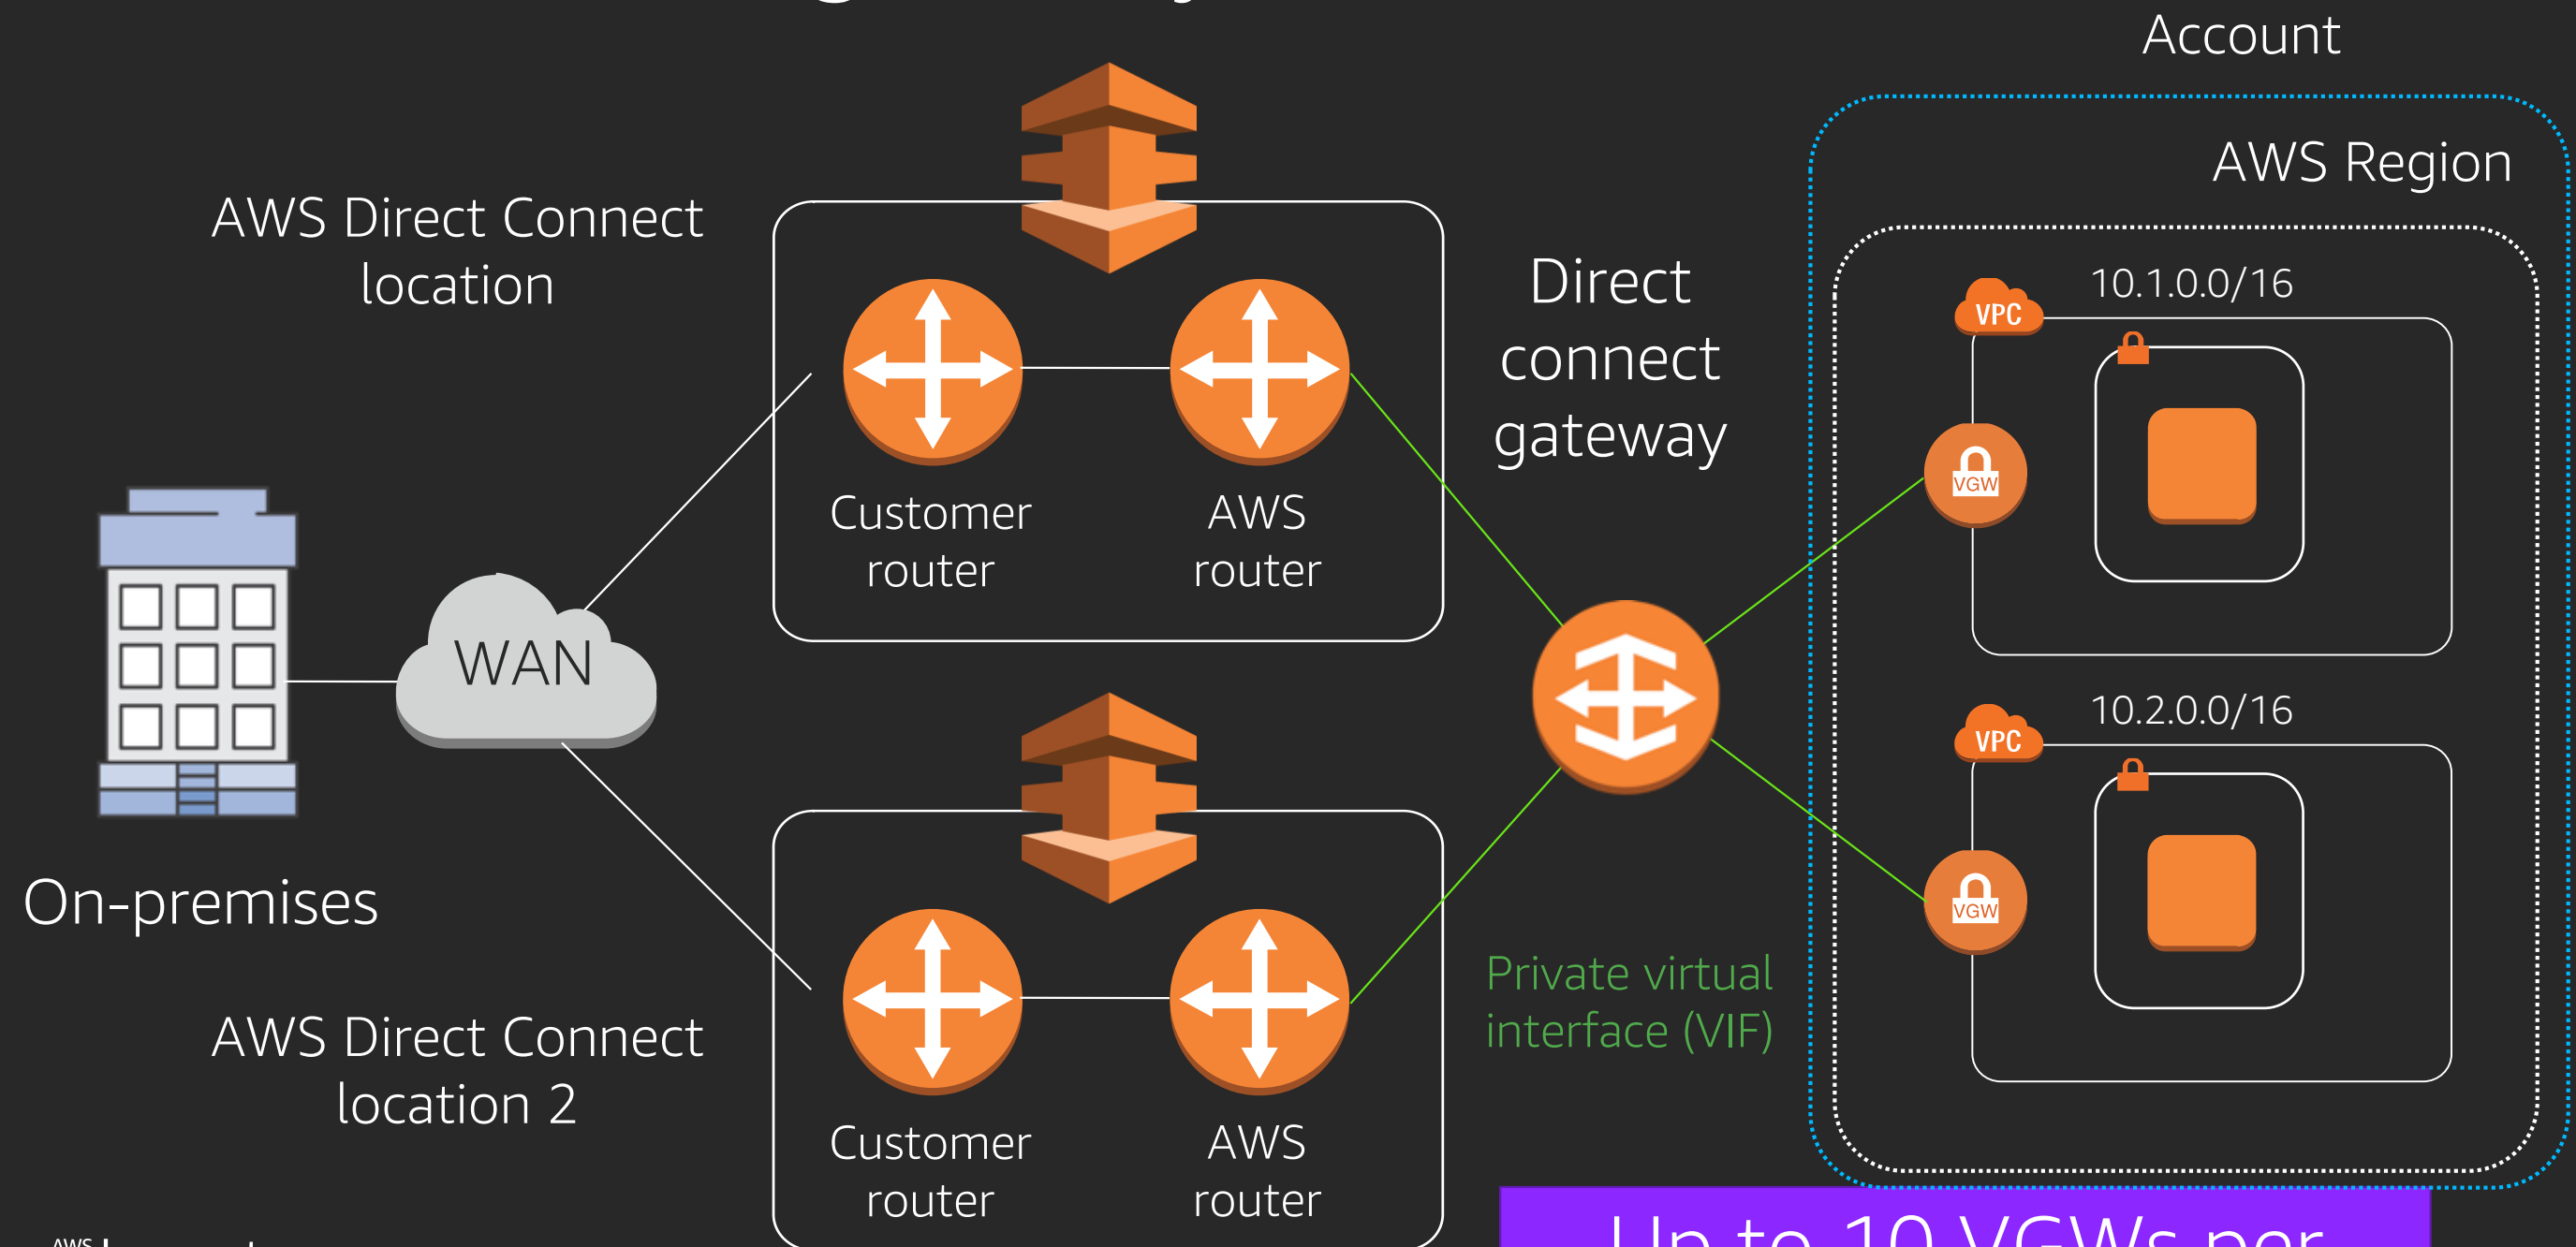
AWS Direct Connect to Many VPCs



AWS Direct Connect: Link Aggregation



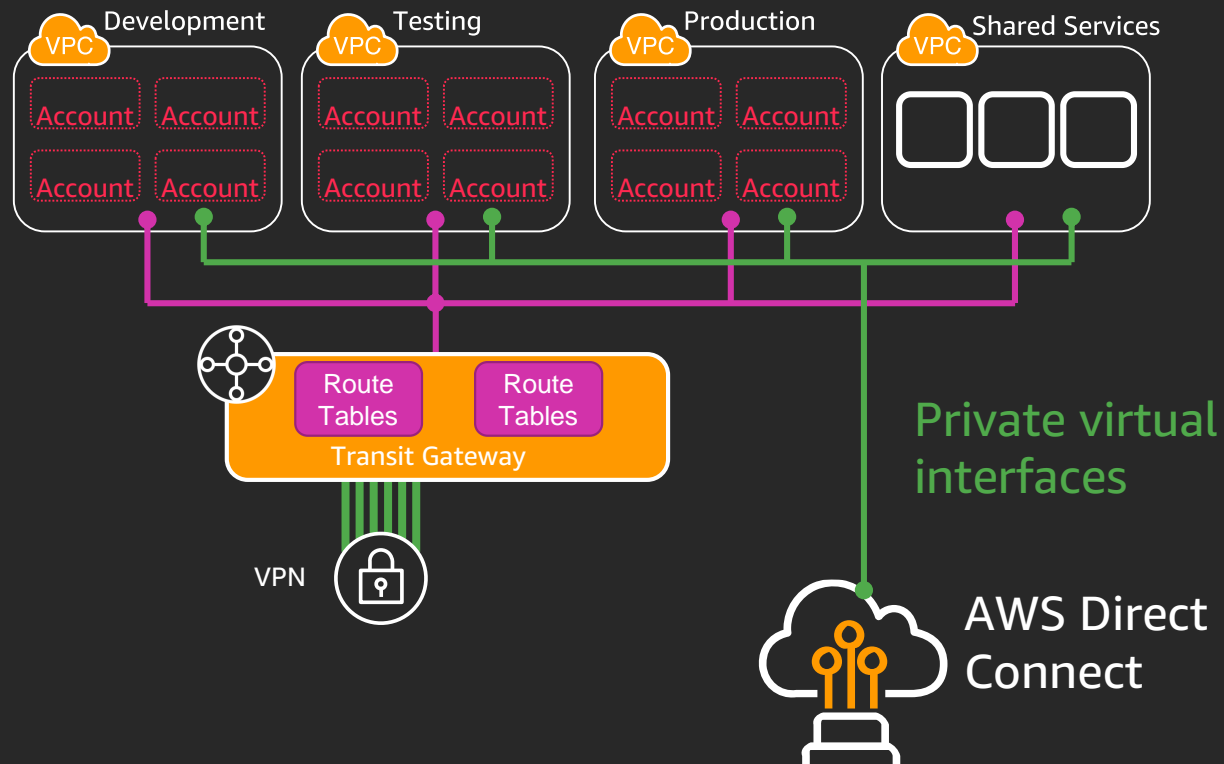
Direct Connect gateway



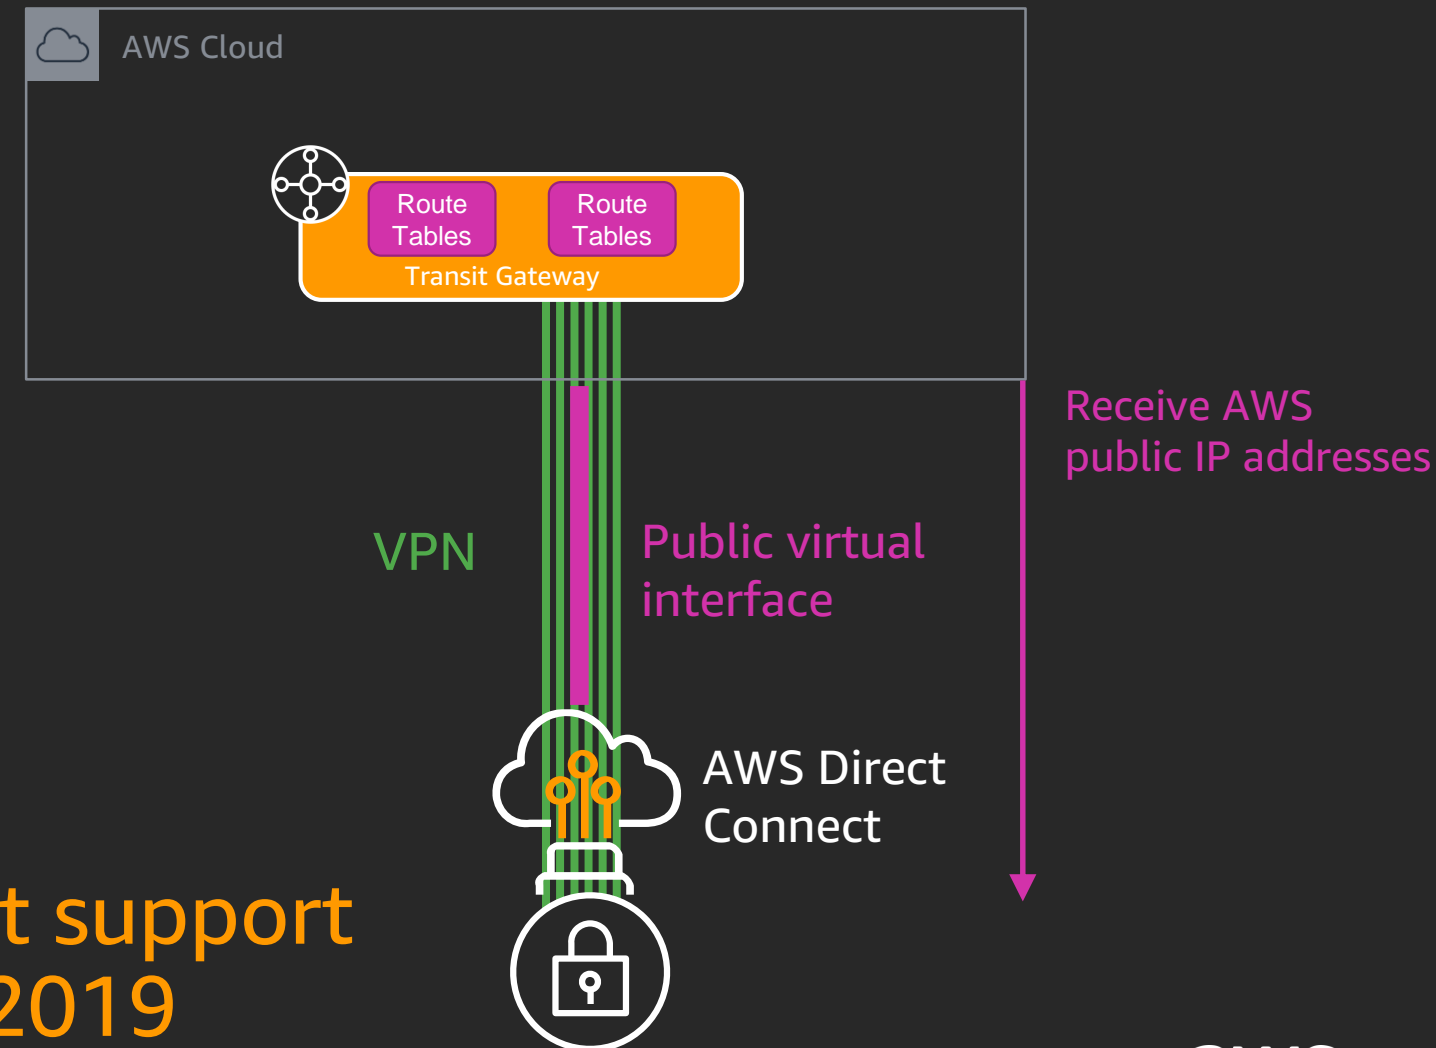
Up to 10 VGWs per direct connect gateway

AWS Direct Connect and Transit Gateway

Use Direct Connect in parallel



Use VPN over a Direct Connect public virtual interface (VIF)




Native Direct Connect support
planned for Q1 2019

AWS Direct Connect and Transit Gateway

Use an edge services VPC in front of a private virtual interface

- More detail in the network services section
- Also how used to **migrate or extend** existing Transit VPCs
- Helpful for single-VIF (<1 Gbps) Direct Connect
- Can be used for North-South inspection use-cases




VPN With Transit Gateway

Consolidate VPN at the Transit Gateway (TGW)

- VPN acts similar to the Virtual Private Gateway (VGW)
 - Bandwidth, configuration, APIs, cost, and experience
 - VPN is attached to a TGW instead of a VGW
 - Same 1.25 gbps bandwidth per tunnel applies

Encryption to the edge of many VPCs

- Traffic is encrypted until it's inside the VPC
- Does not natively encrypt traffic between VPCs
 - Inter-region VPC peering does



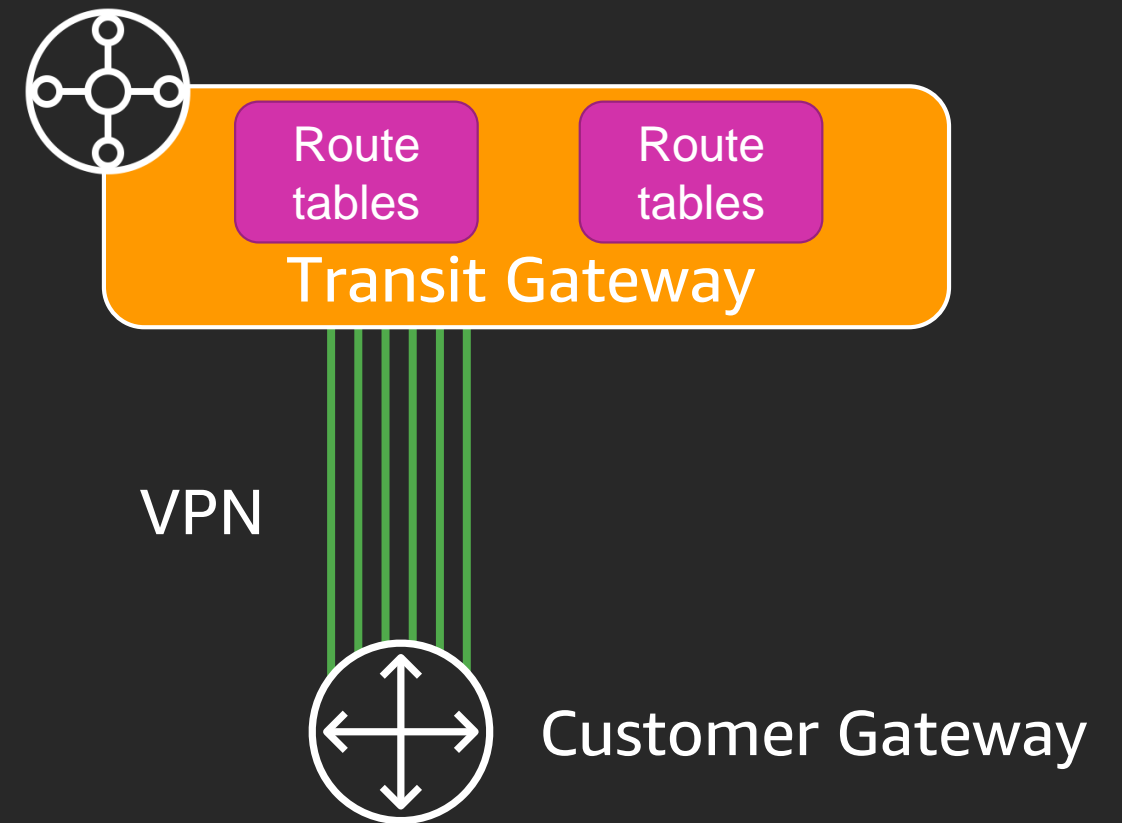
VPN with Transit Gateway: Add more bandwidth

Support for spreading traffic across many tunnels

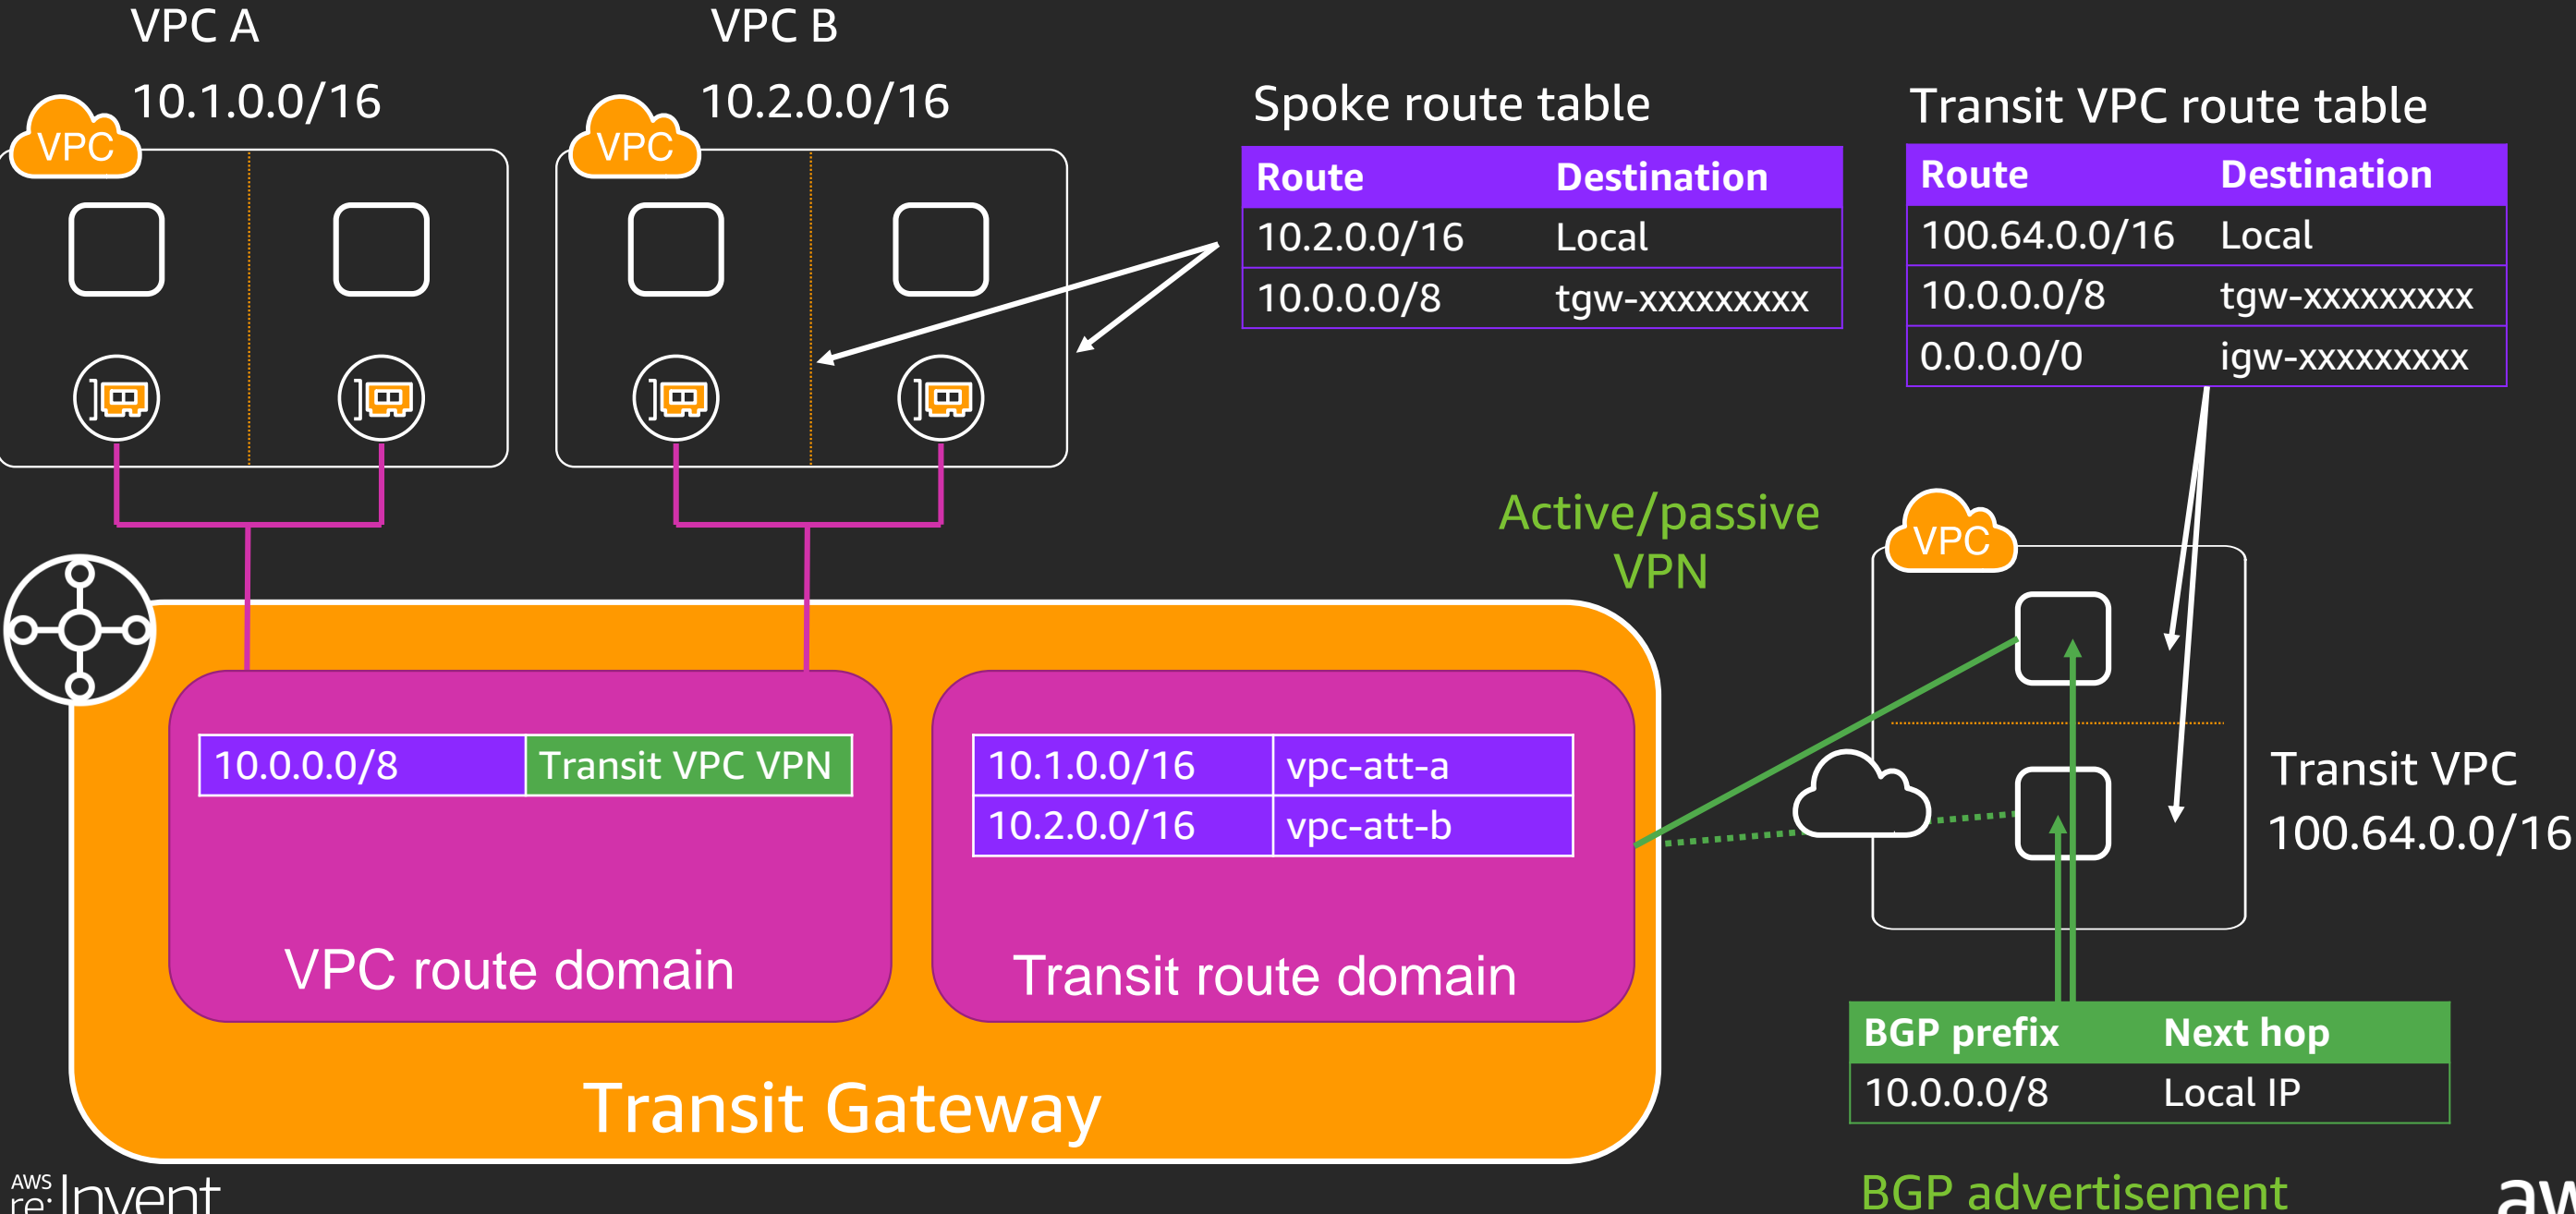
- Equal Cost Multi-Path (ECMP) support with BGP multi-path
- Tested up to 50 Gbps of traffic
- Split traffic into smaller flows, multi-part uploads, etc.

Check your on-premises configuration

- Multi-path BGP
- ECMP support, amount of equal paths, reverse-path forwarding/spoofing checks
- Only supported with BGP, not static routing




Transit VPC 1.1





Neat. But, why?






Account
Strategy



Segmentation Model




Shared Services



Connectivity



Network
Services

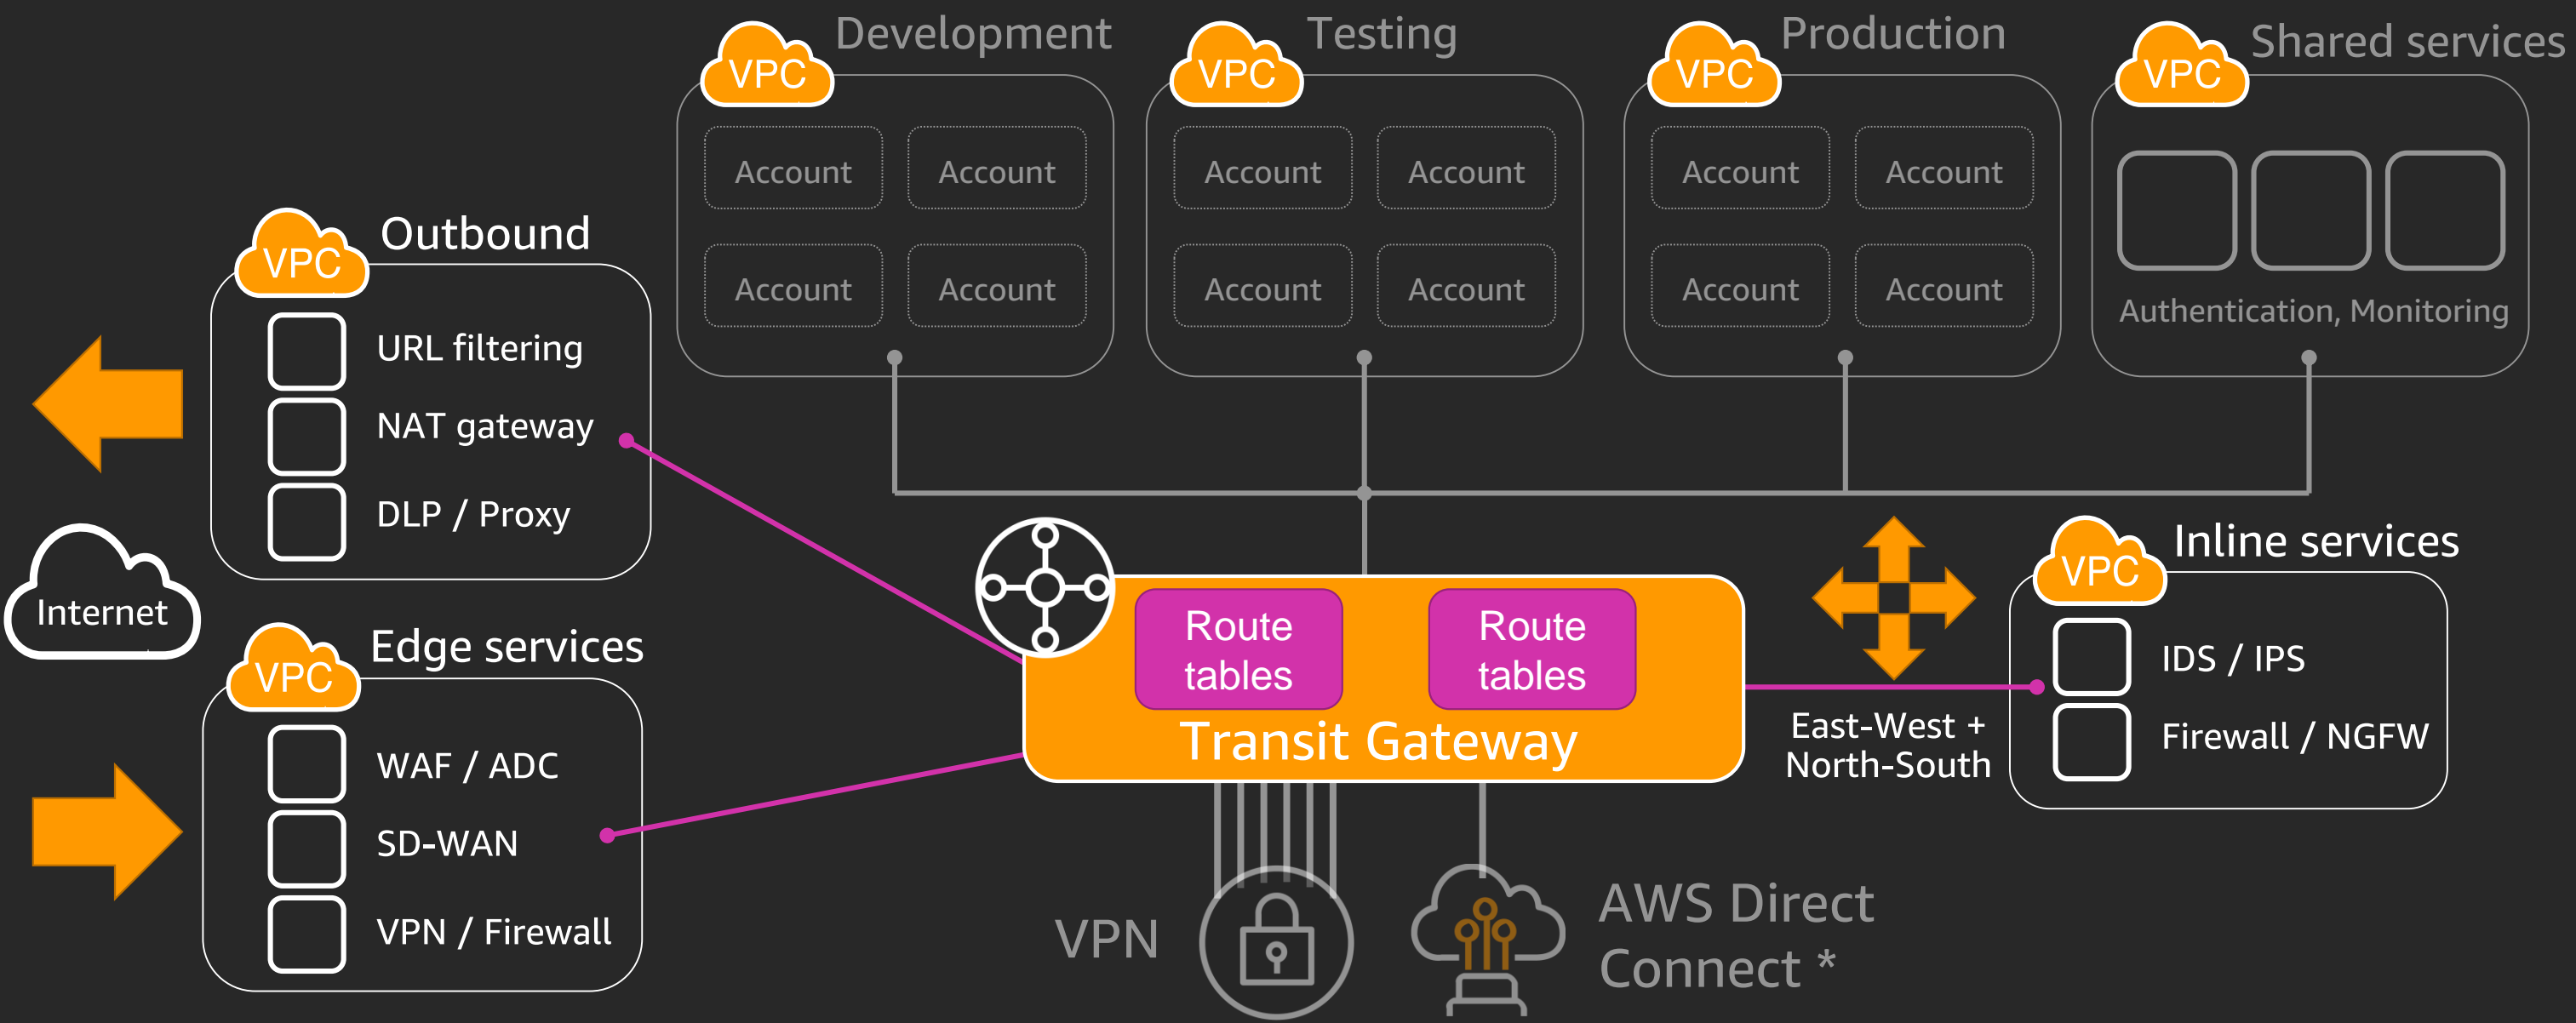


Multi-Region
Options

Network services

Reference Network Architecture

Optional network services



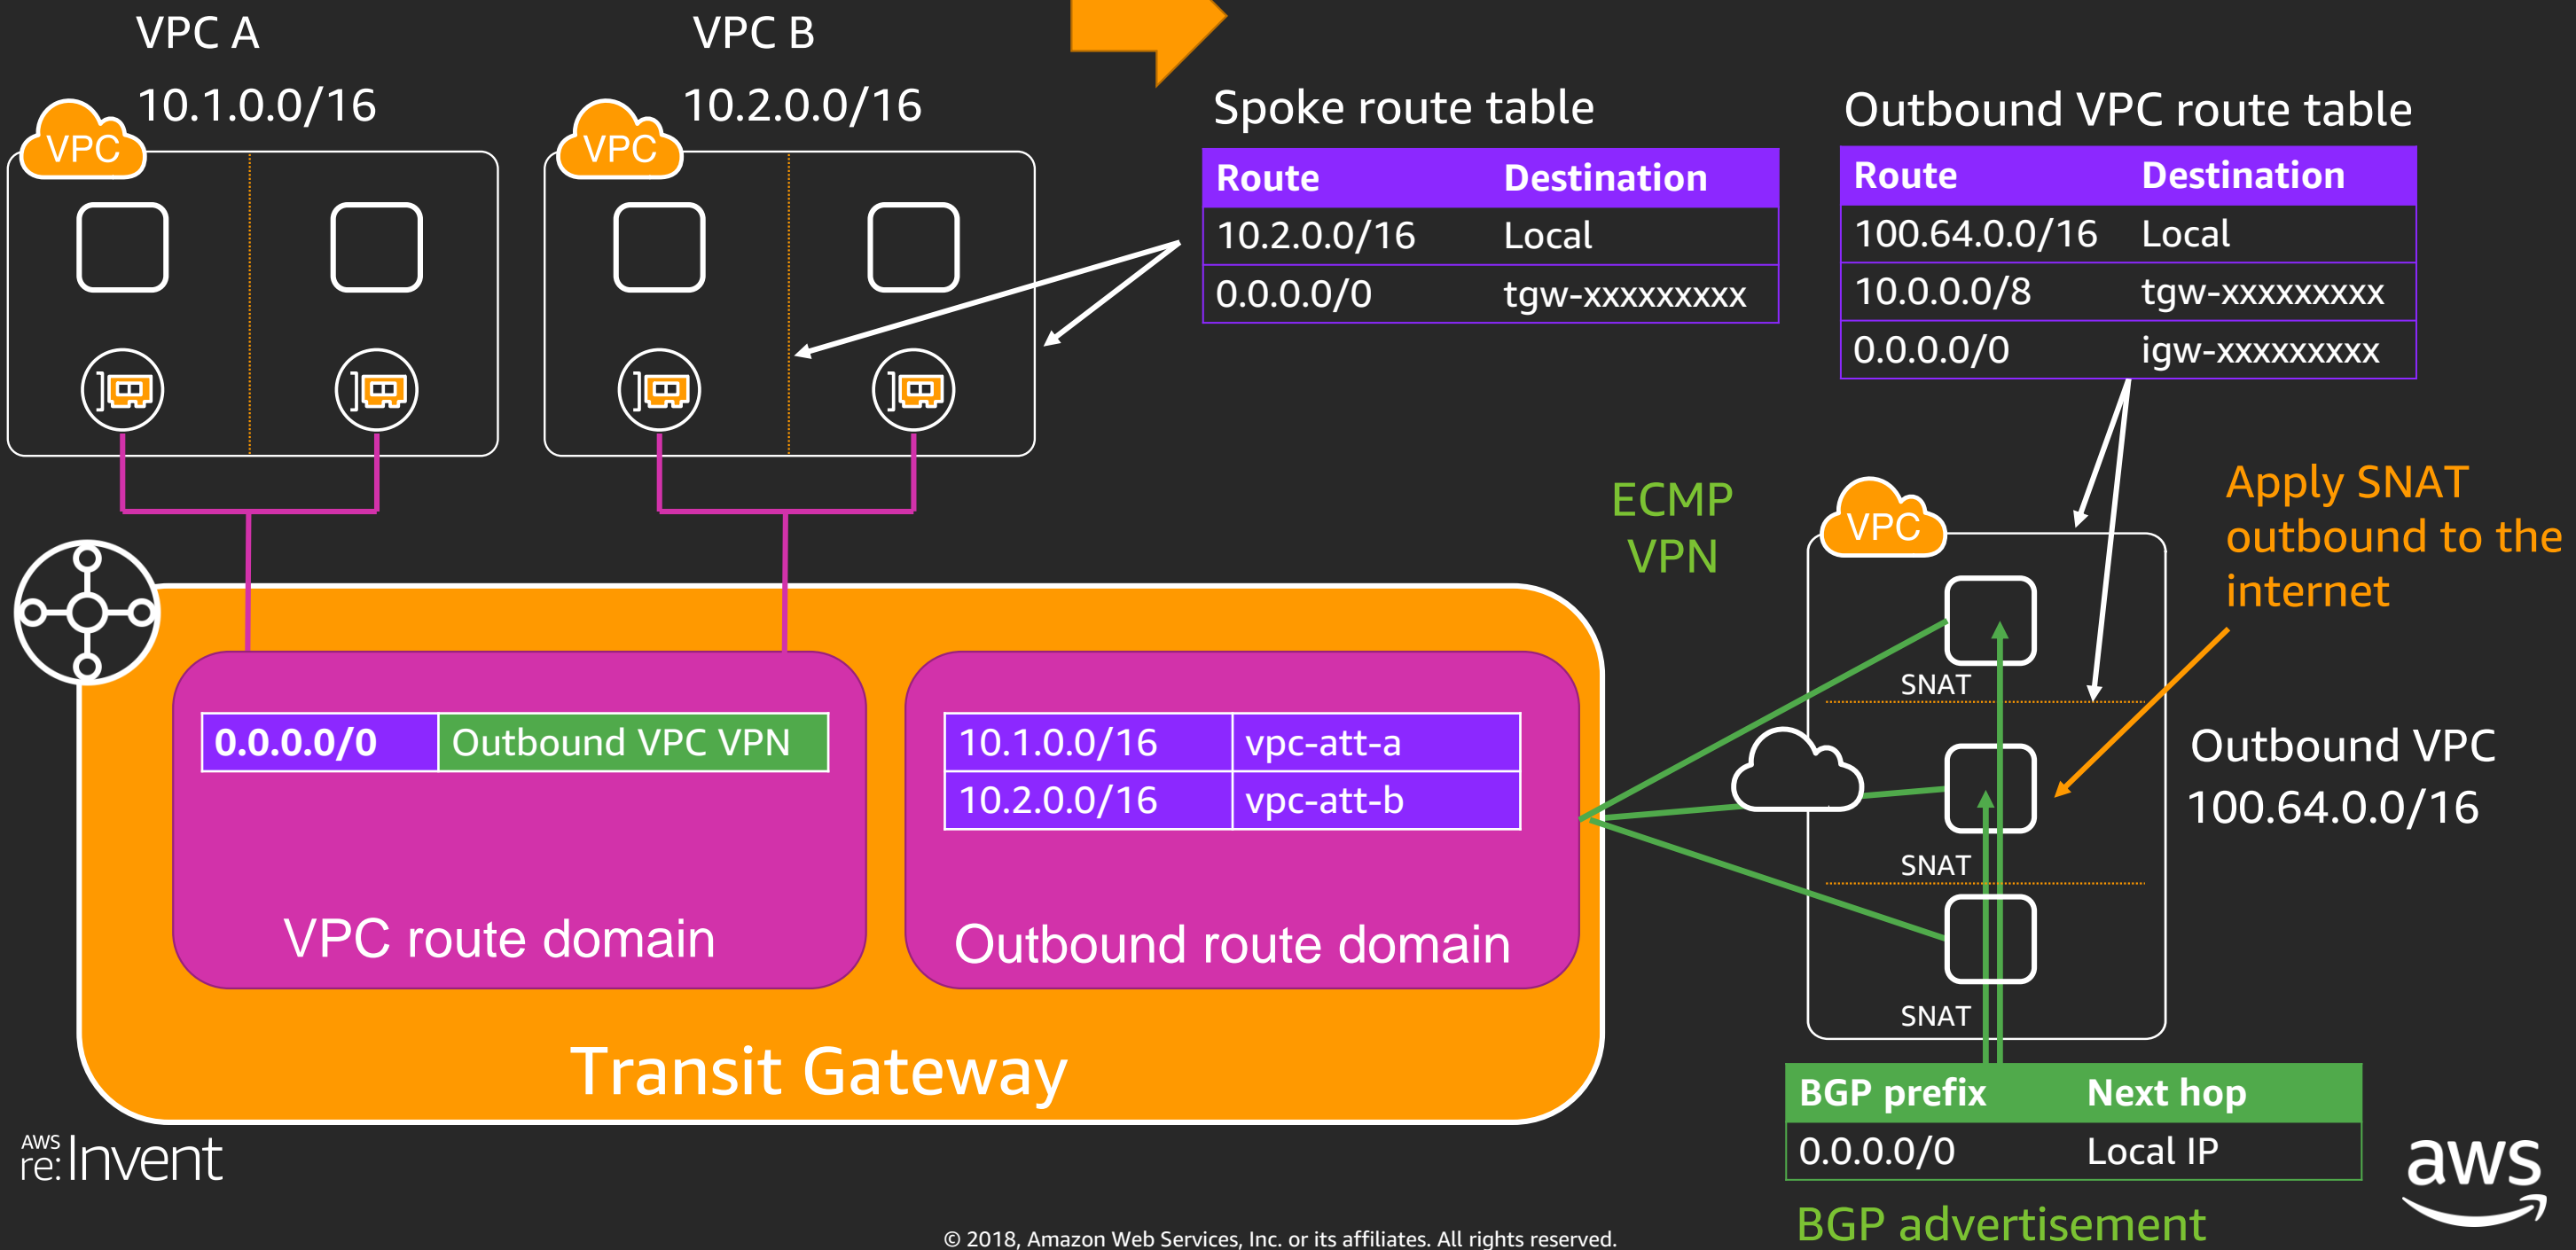
Do I need to put service each into their own VPC?

No, but let's understand the routing separately first.

Outbound services VPC

Use cases:

URL filtering, NAT gateway, data-loss prevention (DLP), web proxy services



VPN service insertion design notes

Instance must be able to support:

- **VPN** to the Transit Gateway
- **BGP** to the Transit Gateway (ECMP requirement)
- **Source NAT** to the internet

Performance

- IPsec overhead
- Compatible with auto-scaling architectures
- **No cumulative bandwidth limit**

High availability

- BGP and VPN Dead Peer Detection handle failover
- **No API calls required for fault tolerance**
- Optionally place instances in Amazon EC2 automatic recovery

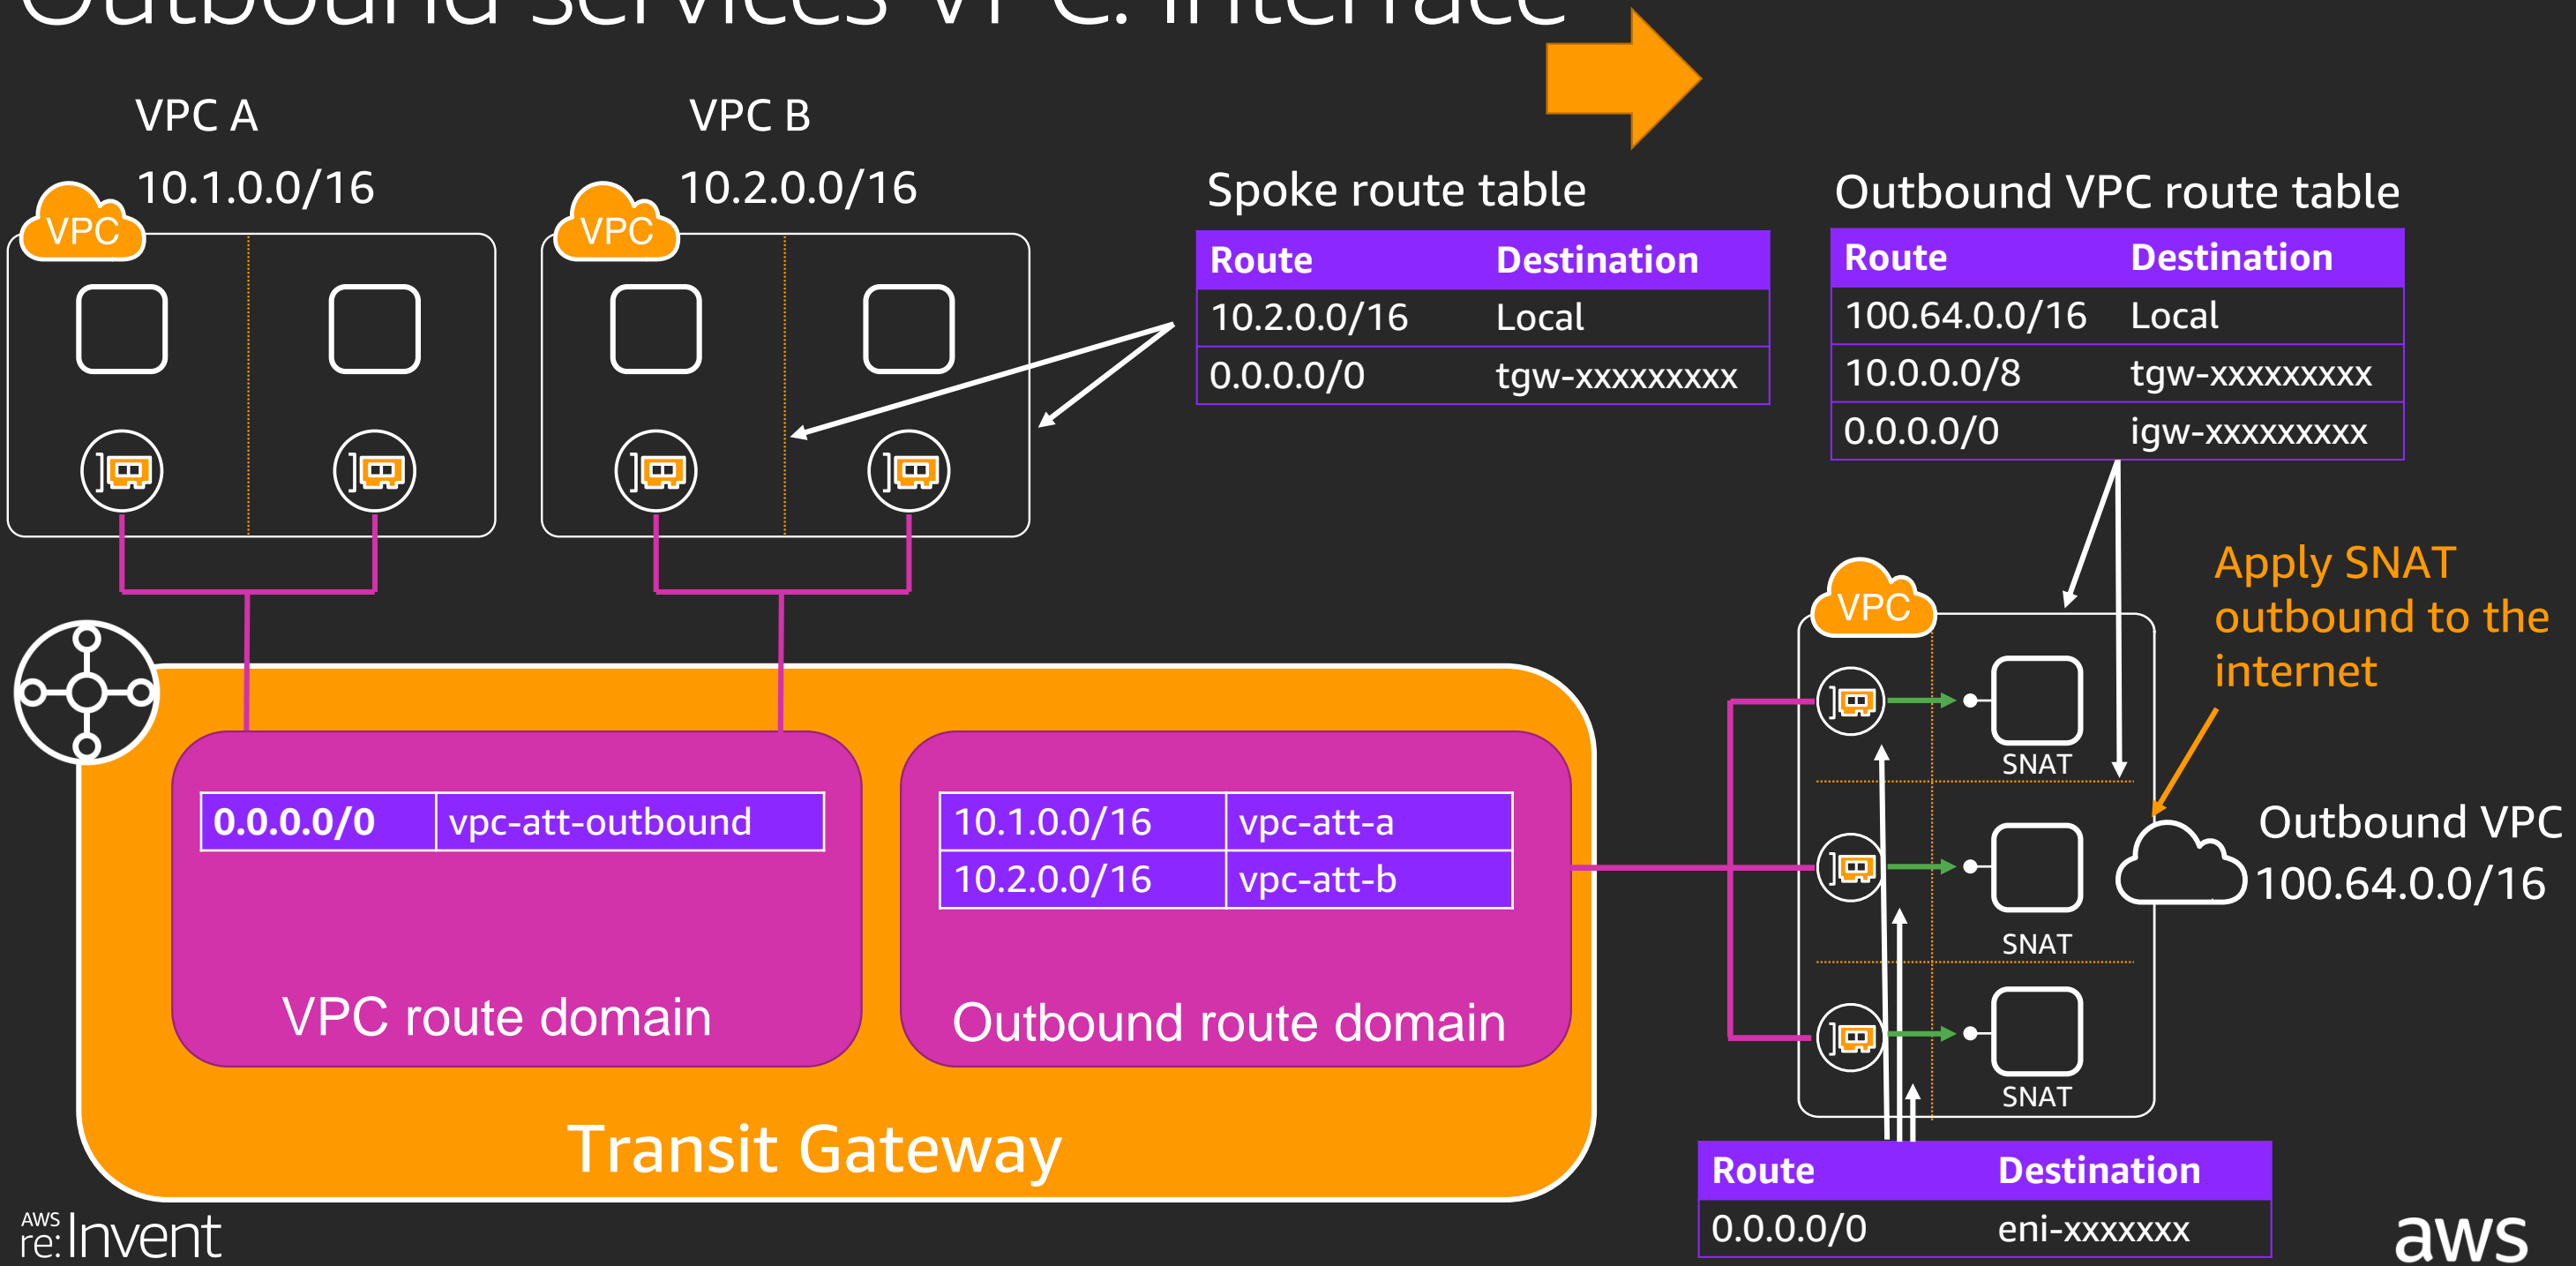
Stateful services

- Use Source NAT to guarantee the return flow to the same instance

Horizontally scalable service pattern

Preferred method if the service supports BGP, VPN and NAT.

Outbound services VPC: Interface



Interface service insertion design notes

Instance must be able to support:

- Source NAT to the internet

Performance

- No overhead (8500 MTU)
- Limited to one Transit Gateway attachment per Availability Zone, so one route table
- Traffic is forwarded within the same Availability Zone if possible
 - Likely that traffic isn't evenly distributed across instances

High availability

- There are no built-in health checks for the VPC routes, requires monitoring and management
- Optionally place instances in Amazon EC2 automatic recovery

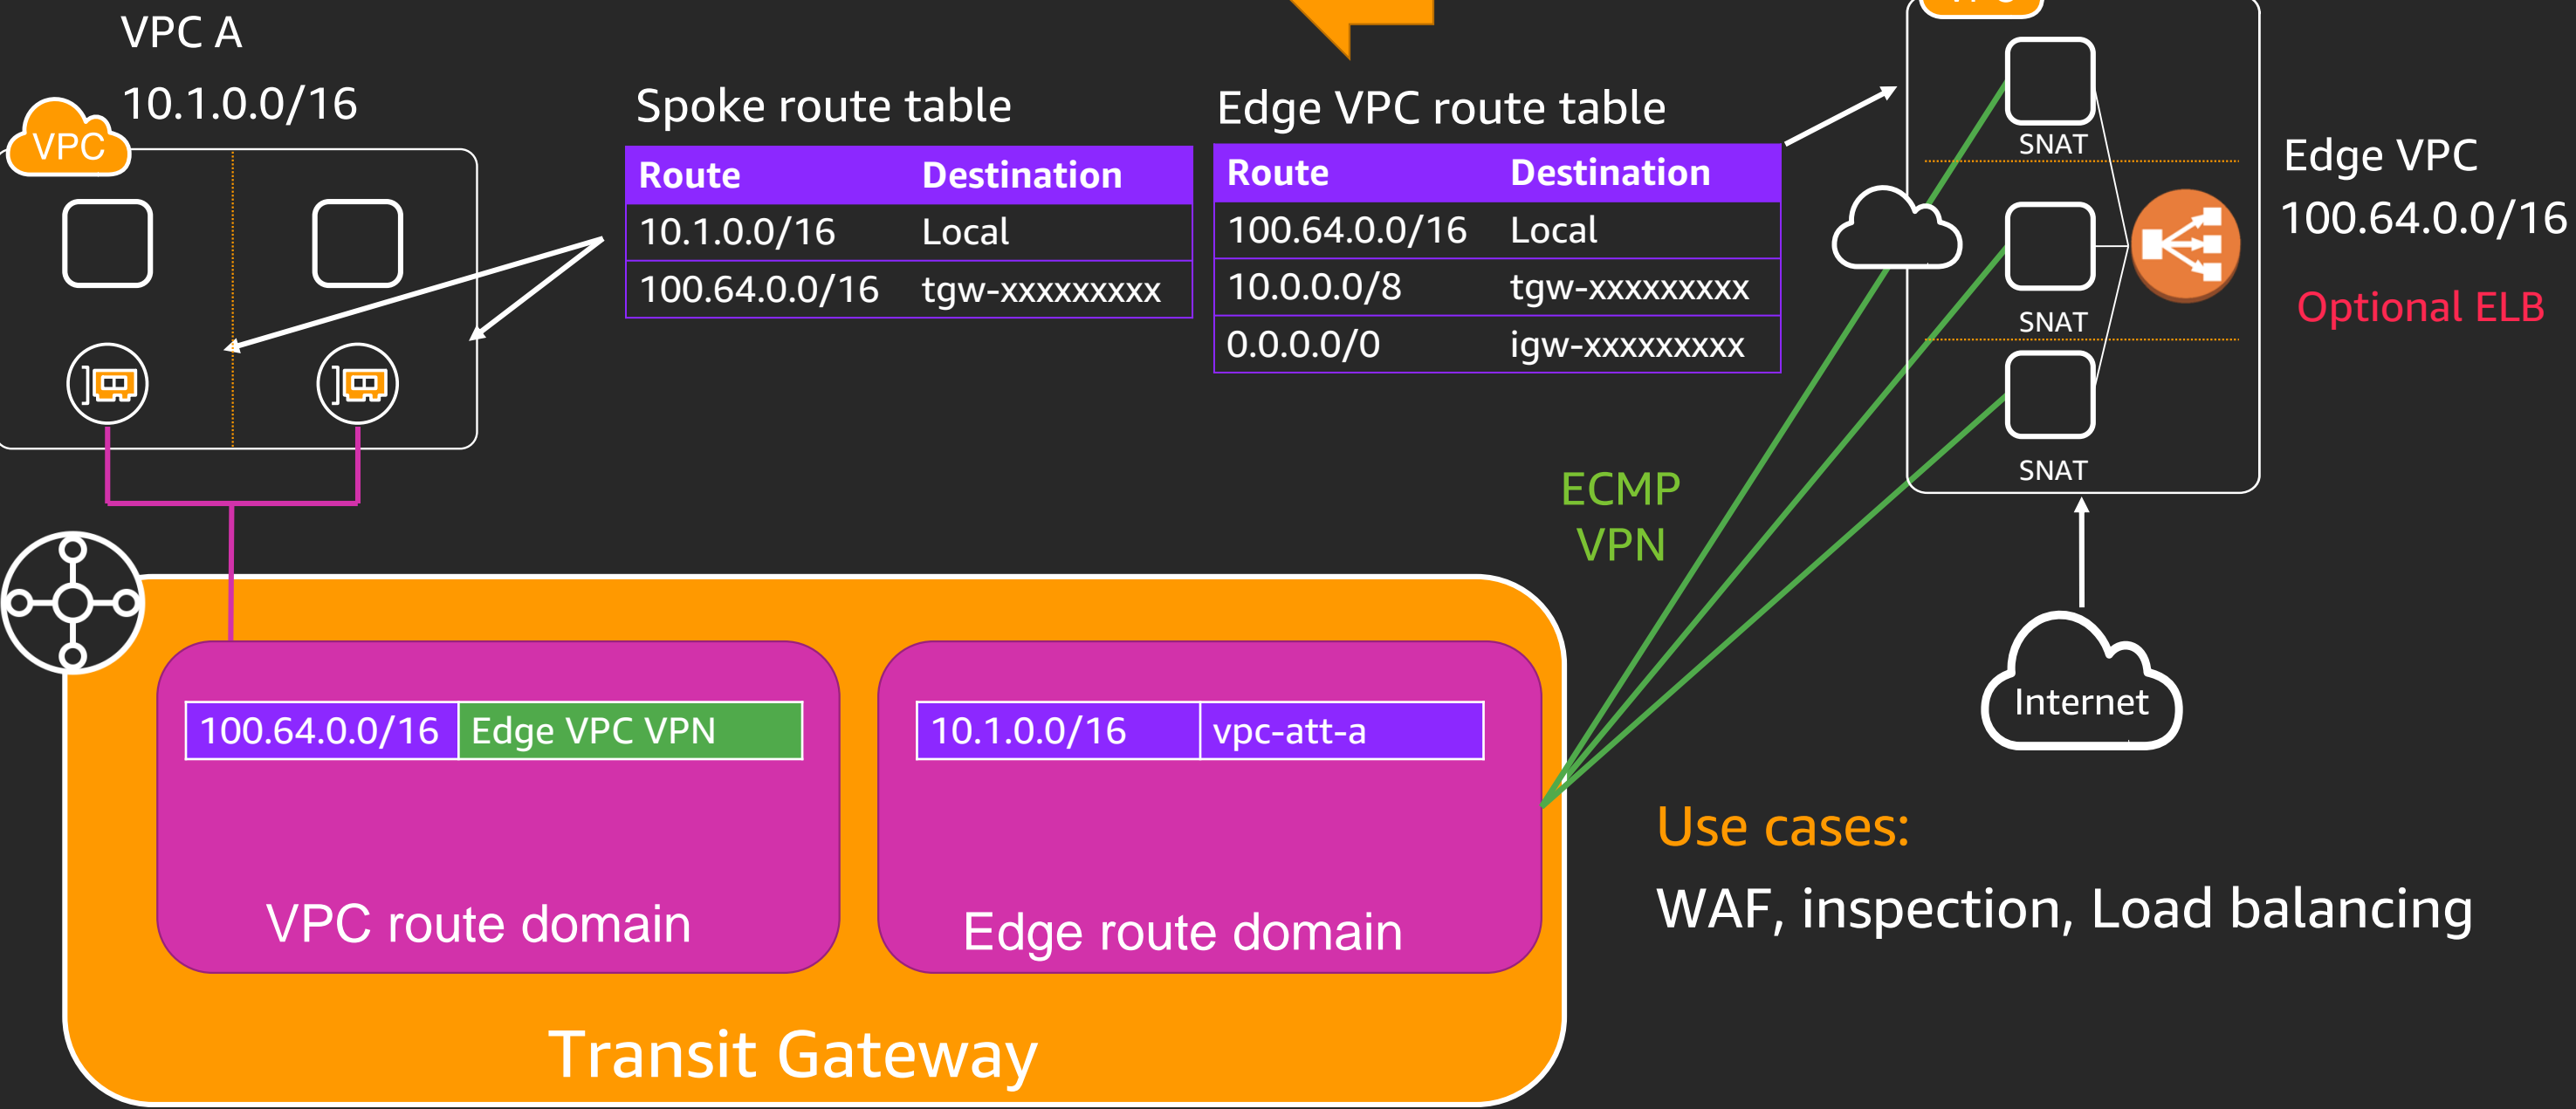
Stateful services

- Use Source NAT to guarantee the return flow to the same instance

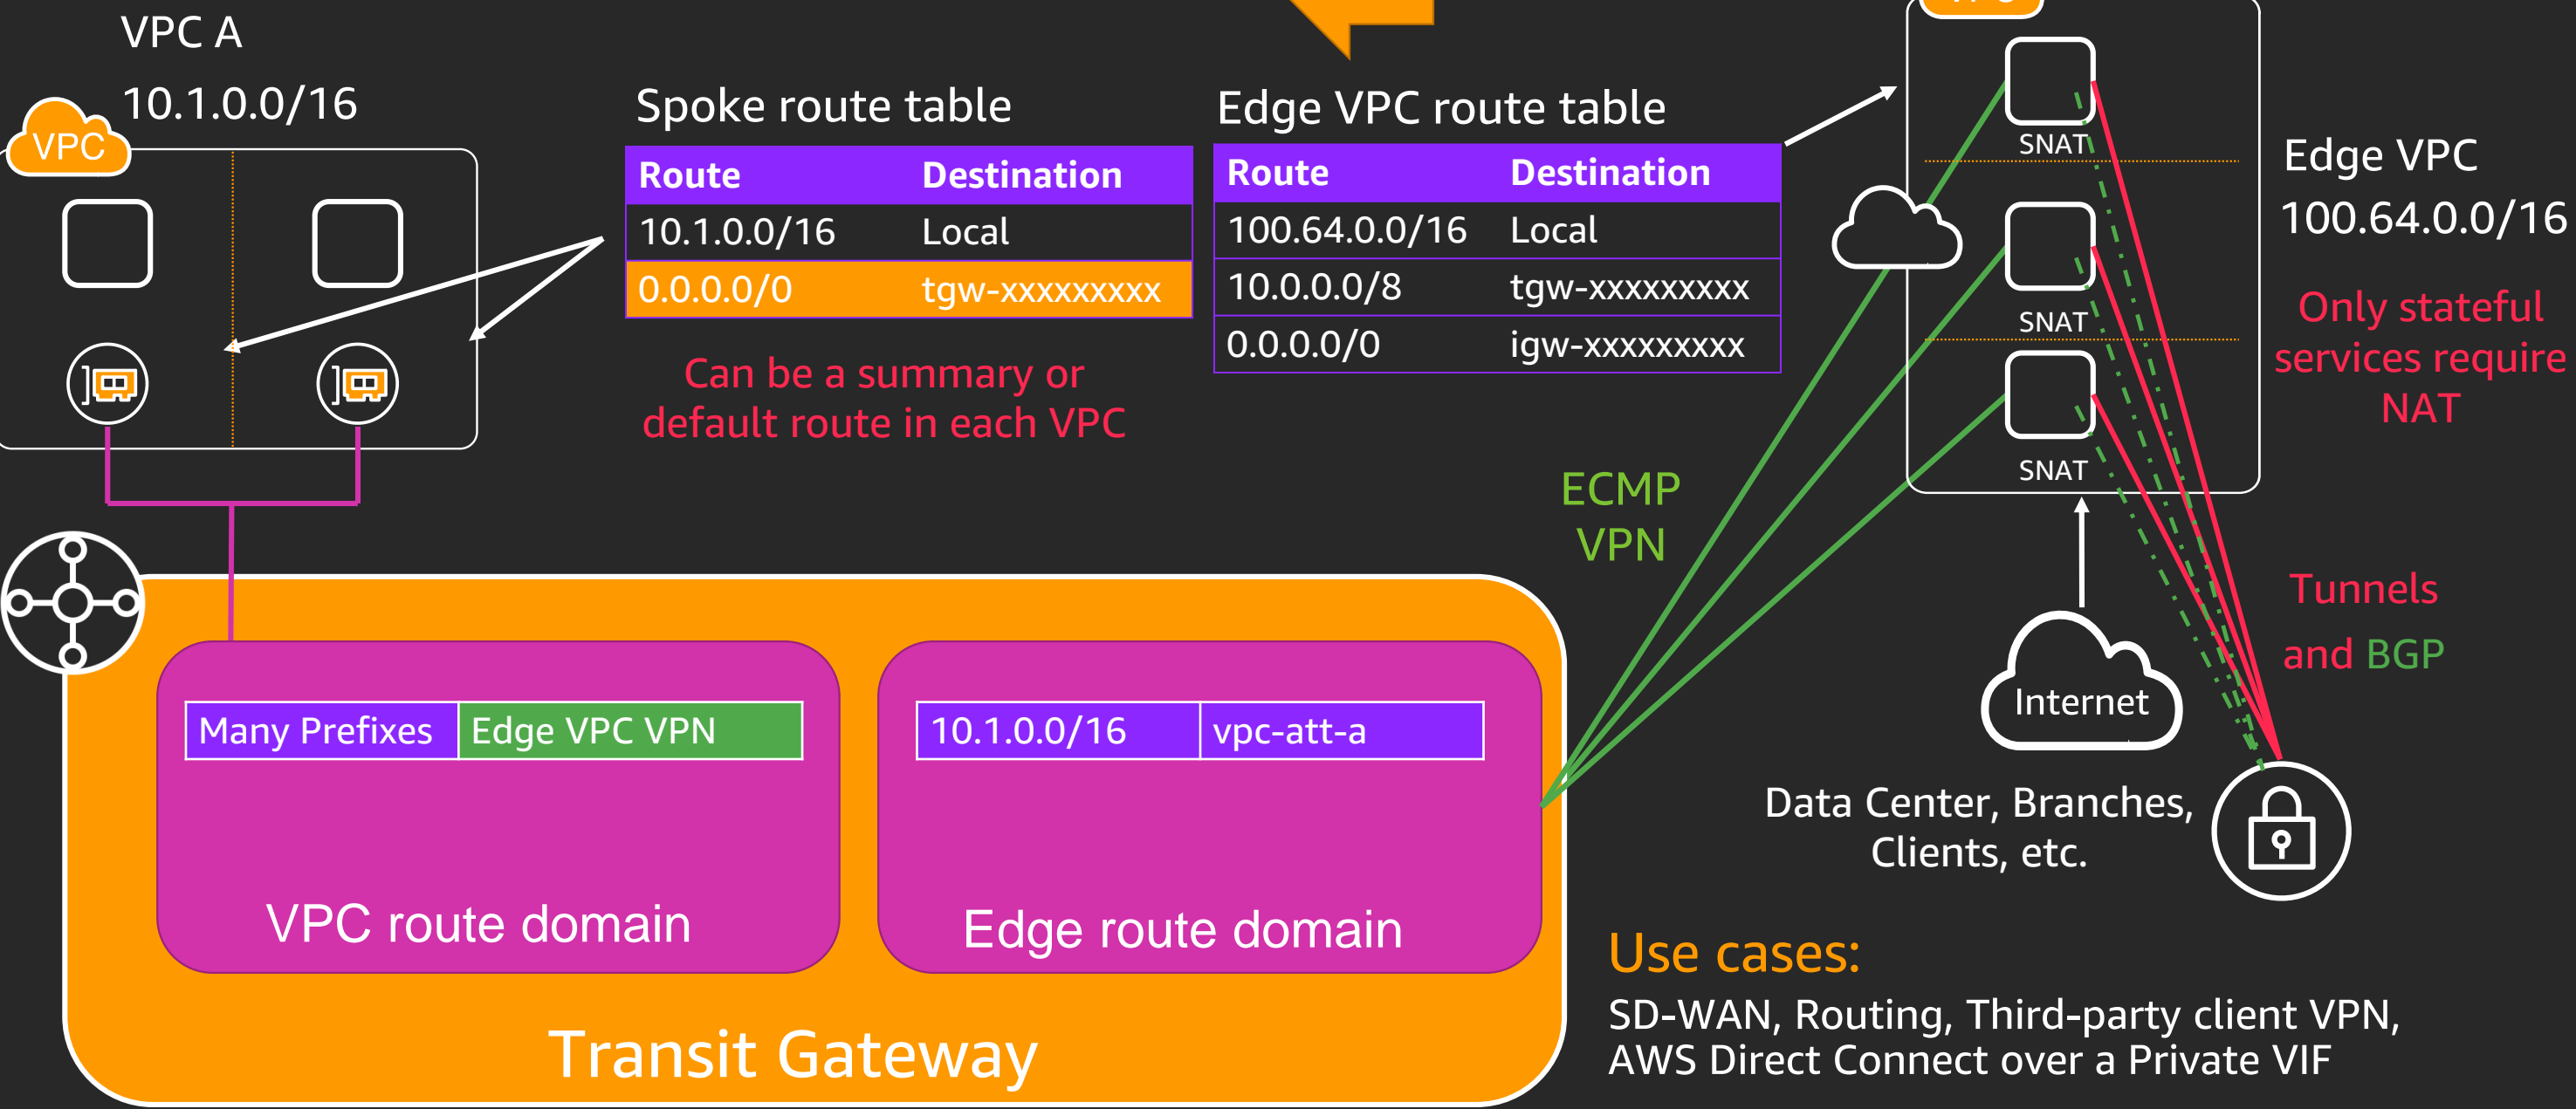
Simpler performance pattern

Stay within the performance of a single service instance (worst-case scenario) and configure your own high availability checks.


Edge services VPC: Ingress



Edge services VPC: SD-WAN



Reminder:



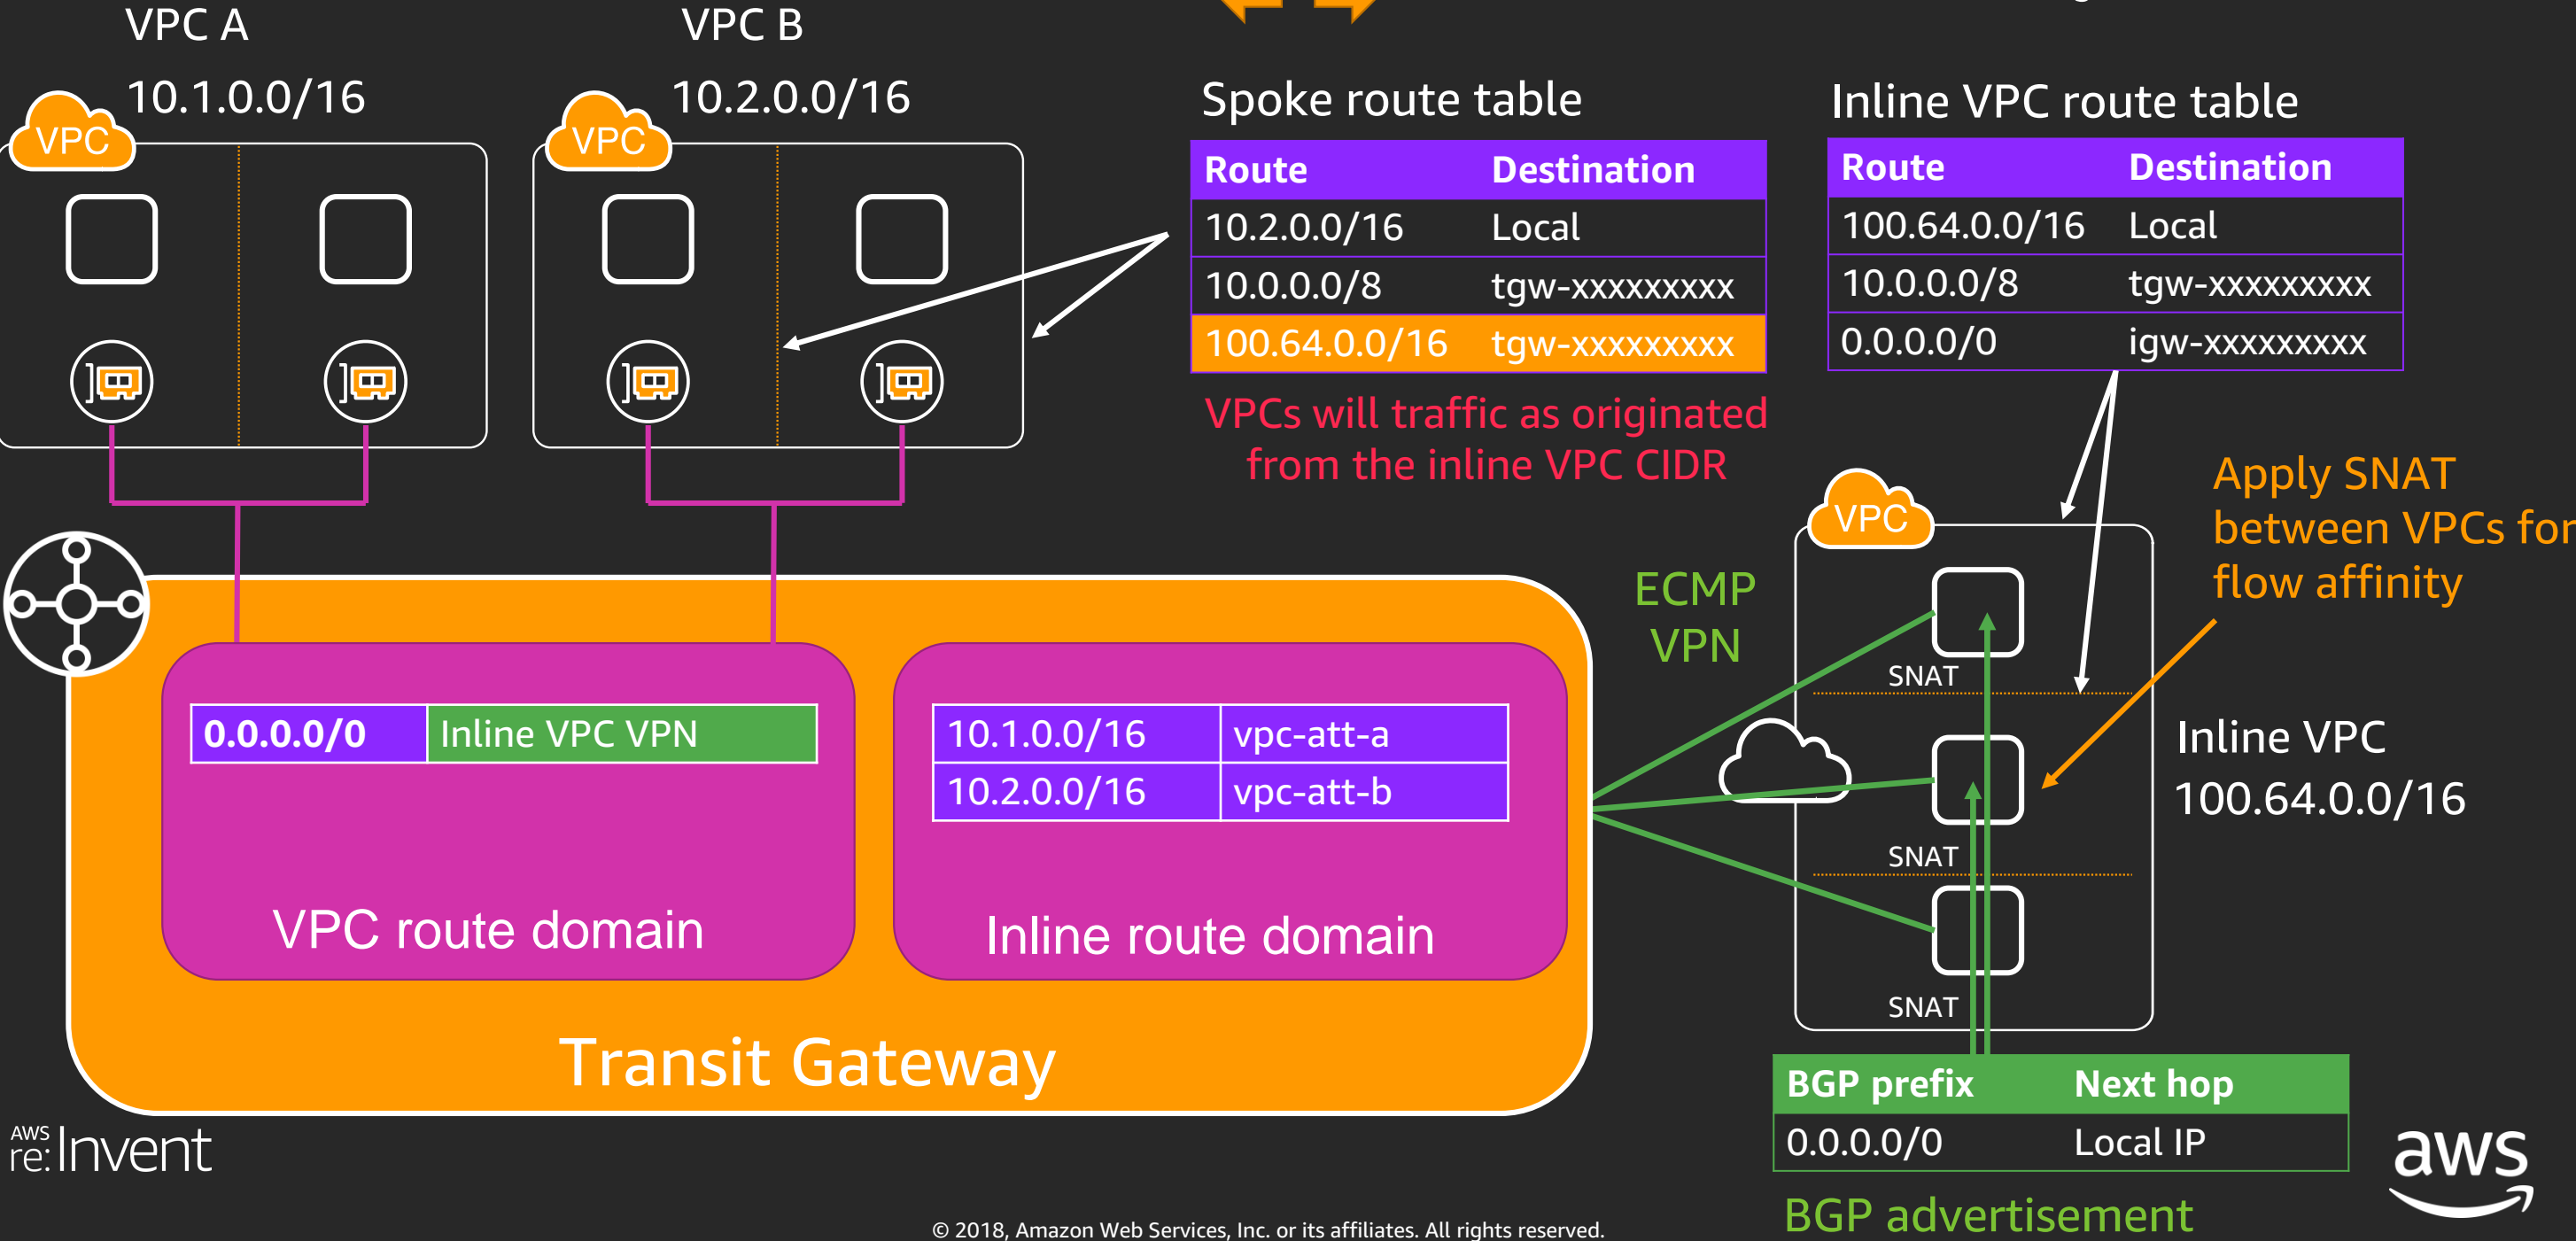
Existing network services or DMZs
may be **convenient**, but they may
also be **the problem**.

Remember to evaluate operational processes, alternatives, and automation

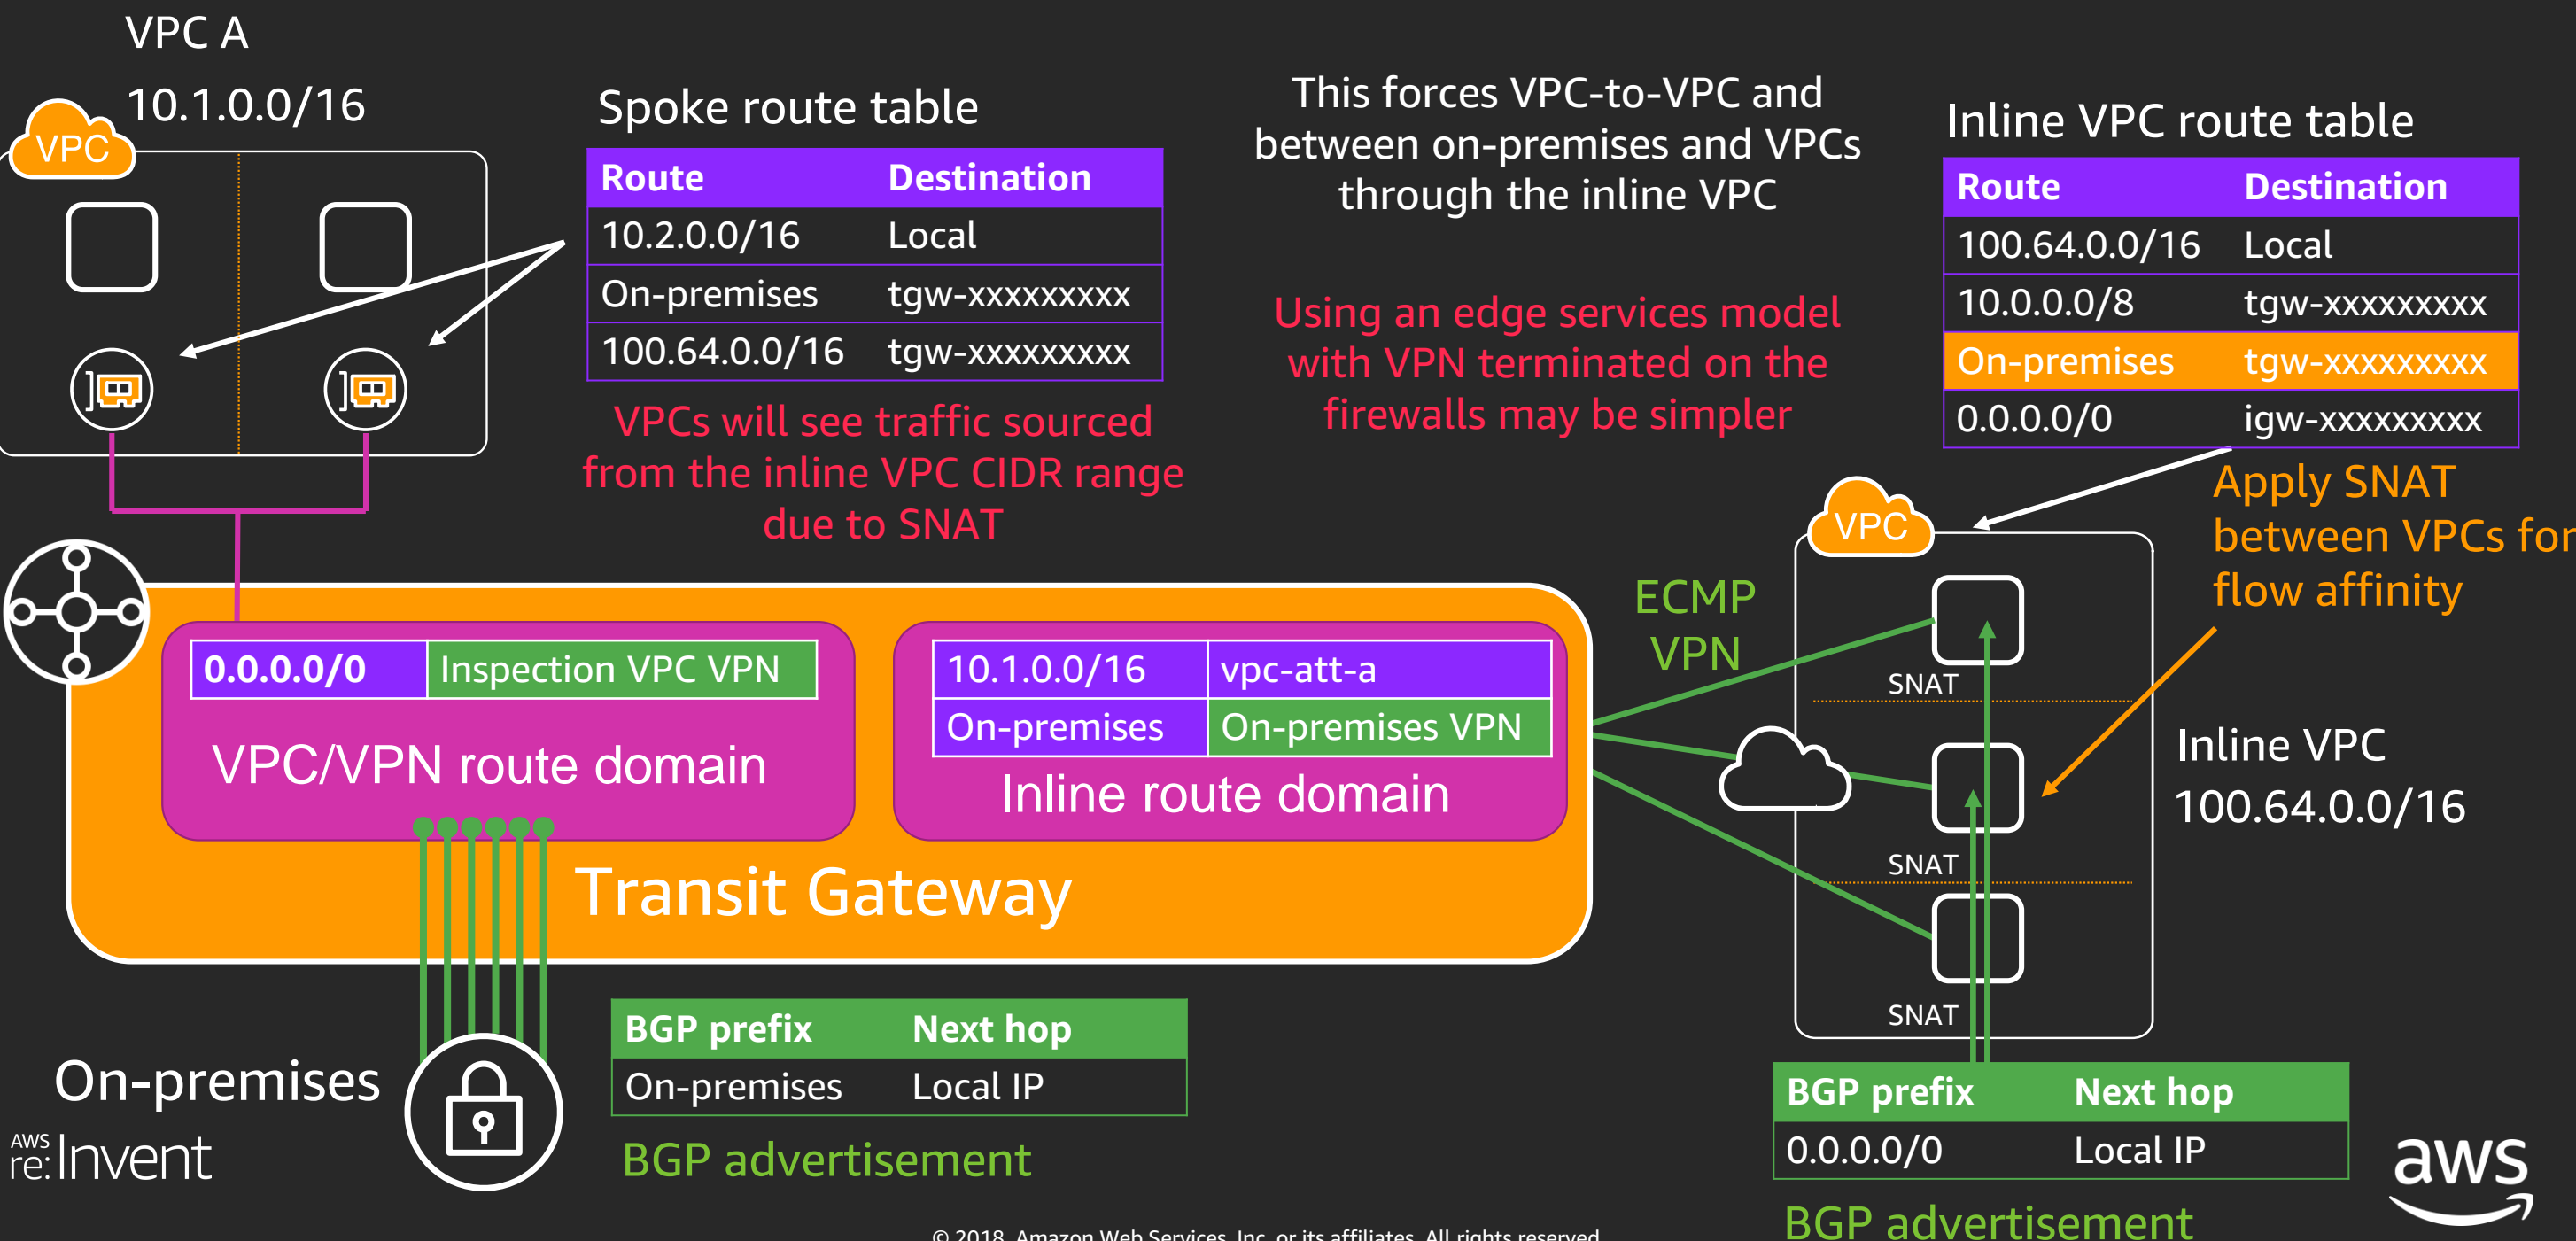
VPC to VPC service insertion

Use cases:

Intrusion detection/prevention (IDS/IPS), firewalls, NextGen Firewalls (NGFW), Unified Threat Management (UTM)



VPC to on-premises service insertion



Transit Gateway launch partners



O E I M



Check Point[®]
SOFTWARE TECHNOLOGIES LTD.

O E I M



CISCO
vEdge SD-WAN

O E I M



O E I M

FORTINET[®]

O E I M



VM-Series

O E I M



O E I M



HashiCorp
Terraform

O E I M

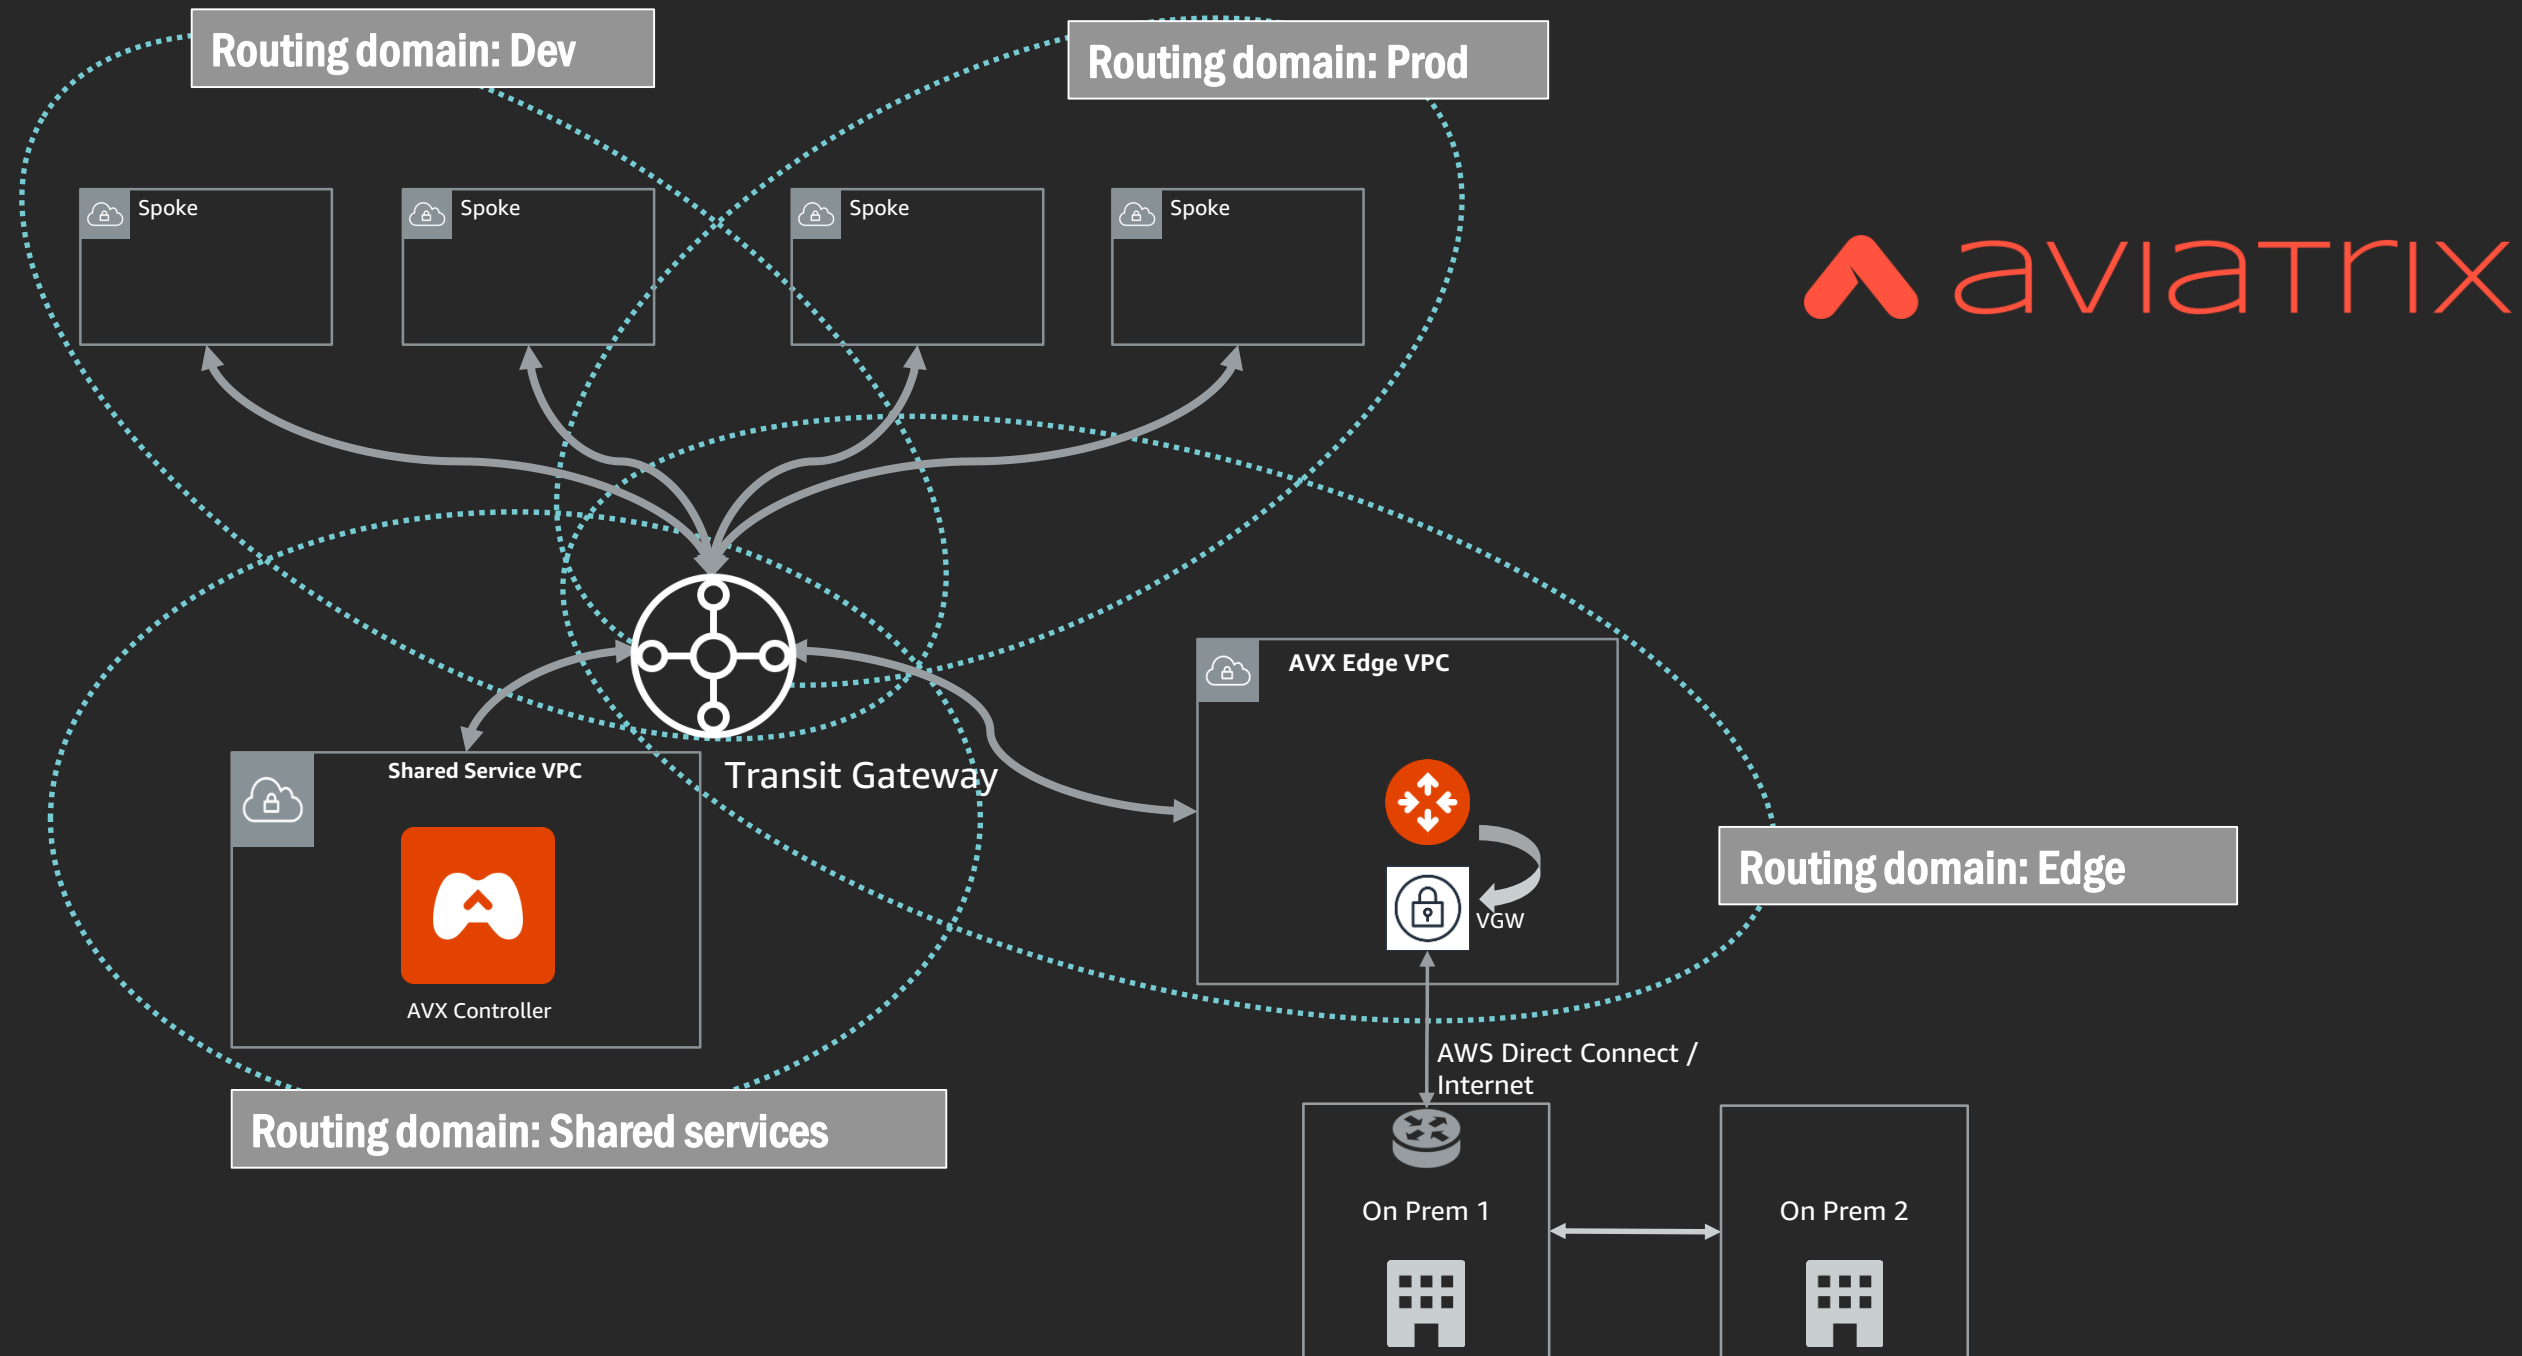
➡ Outbound services

⬅ Edge services

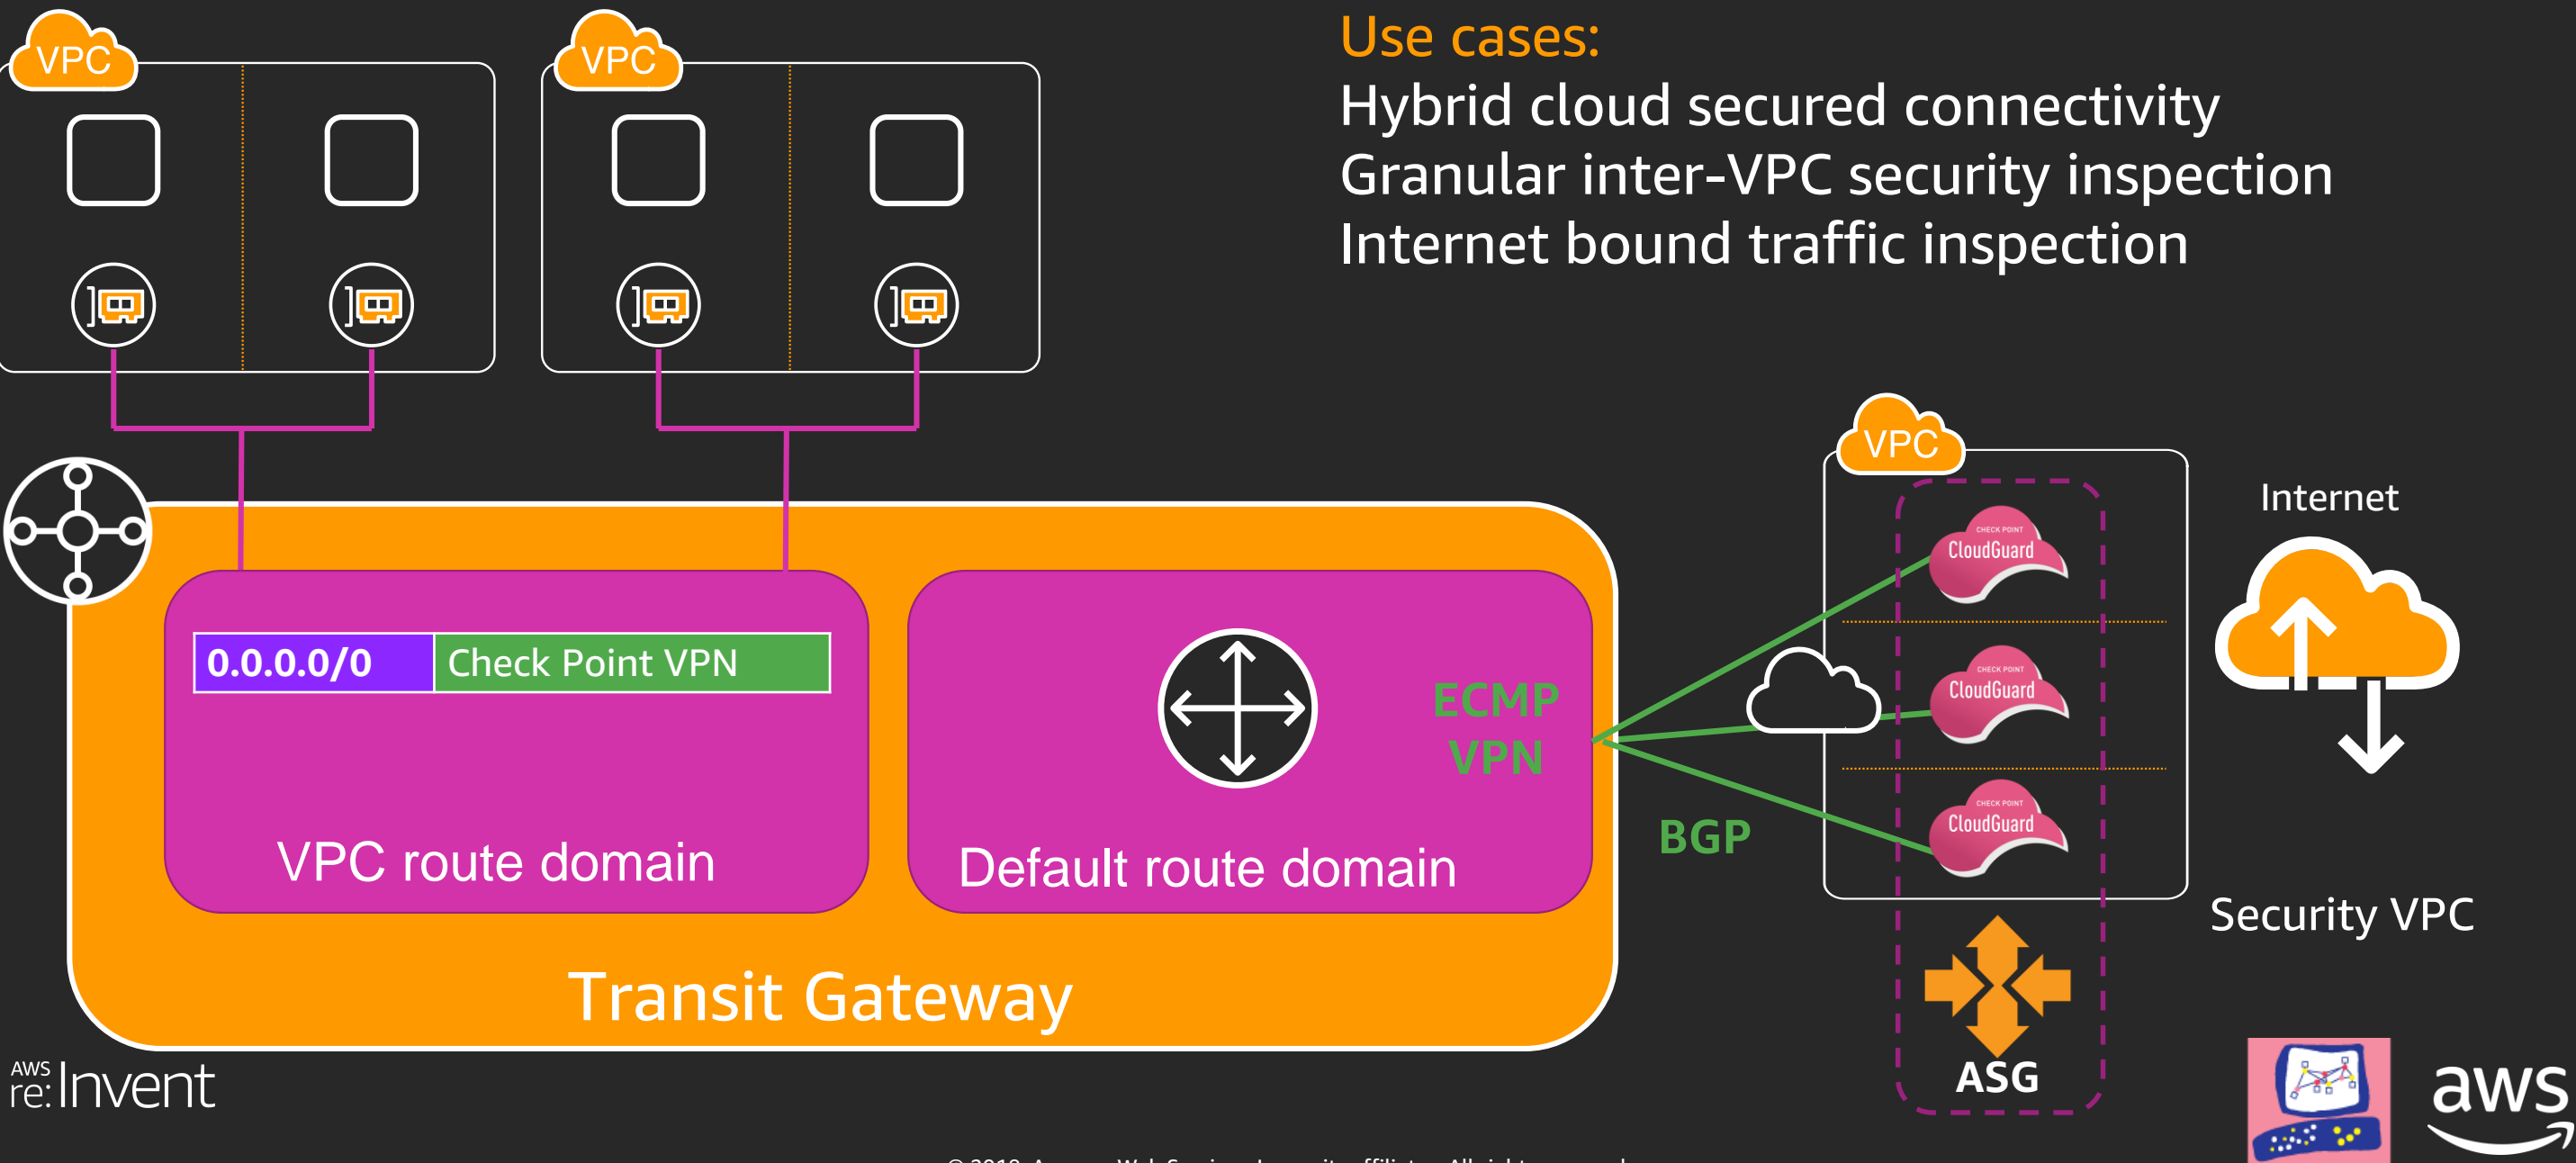
⬆⬆⬆⬆ Inline services

M Management

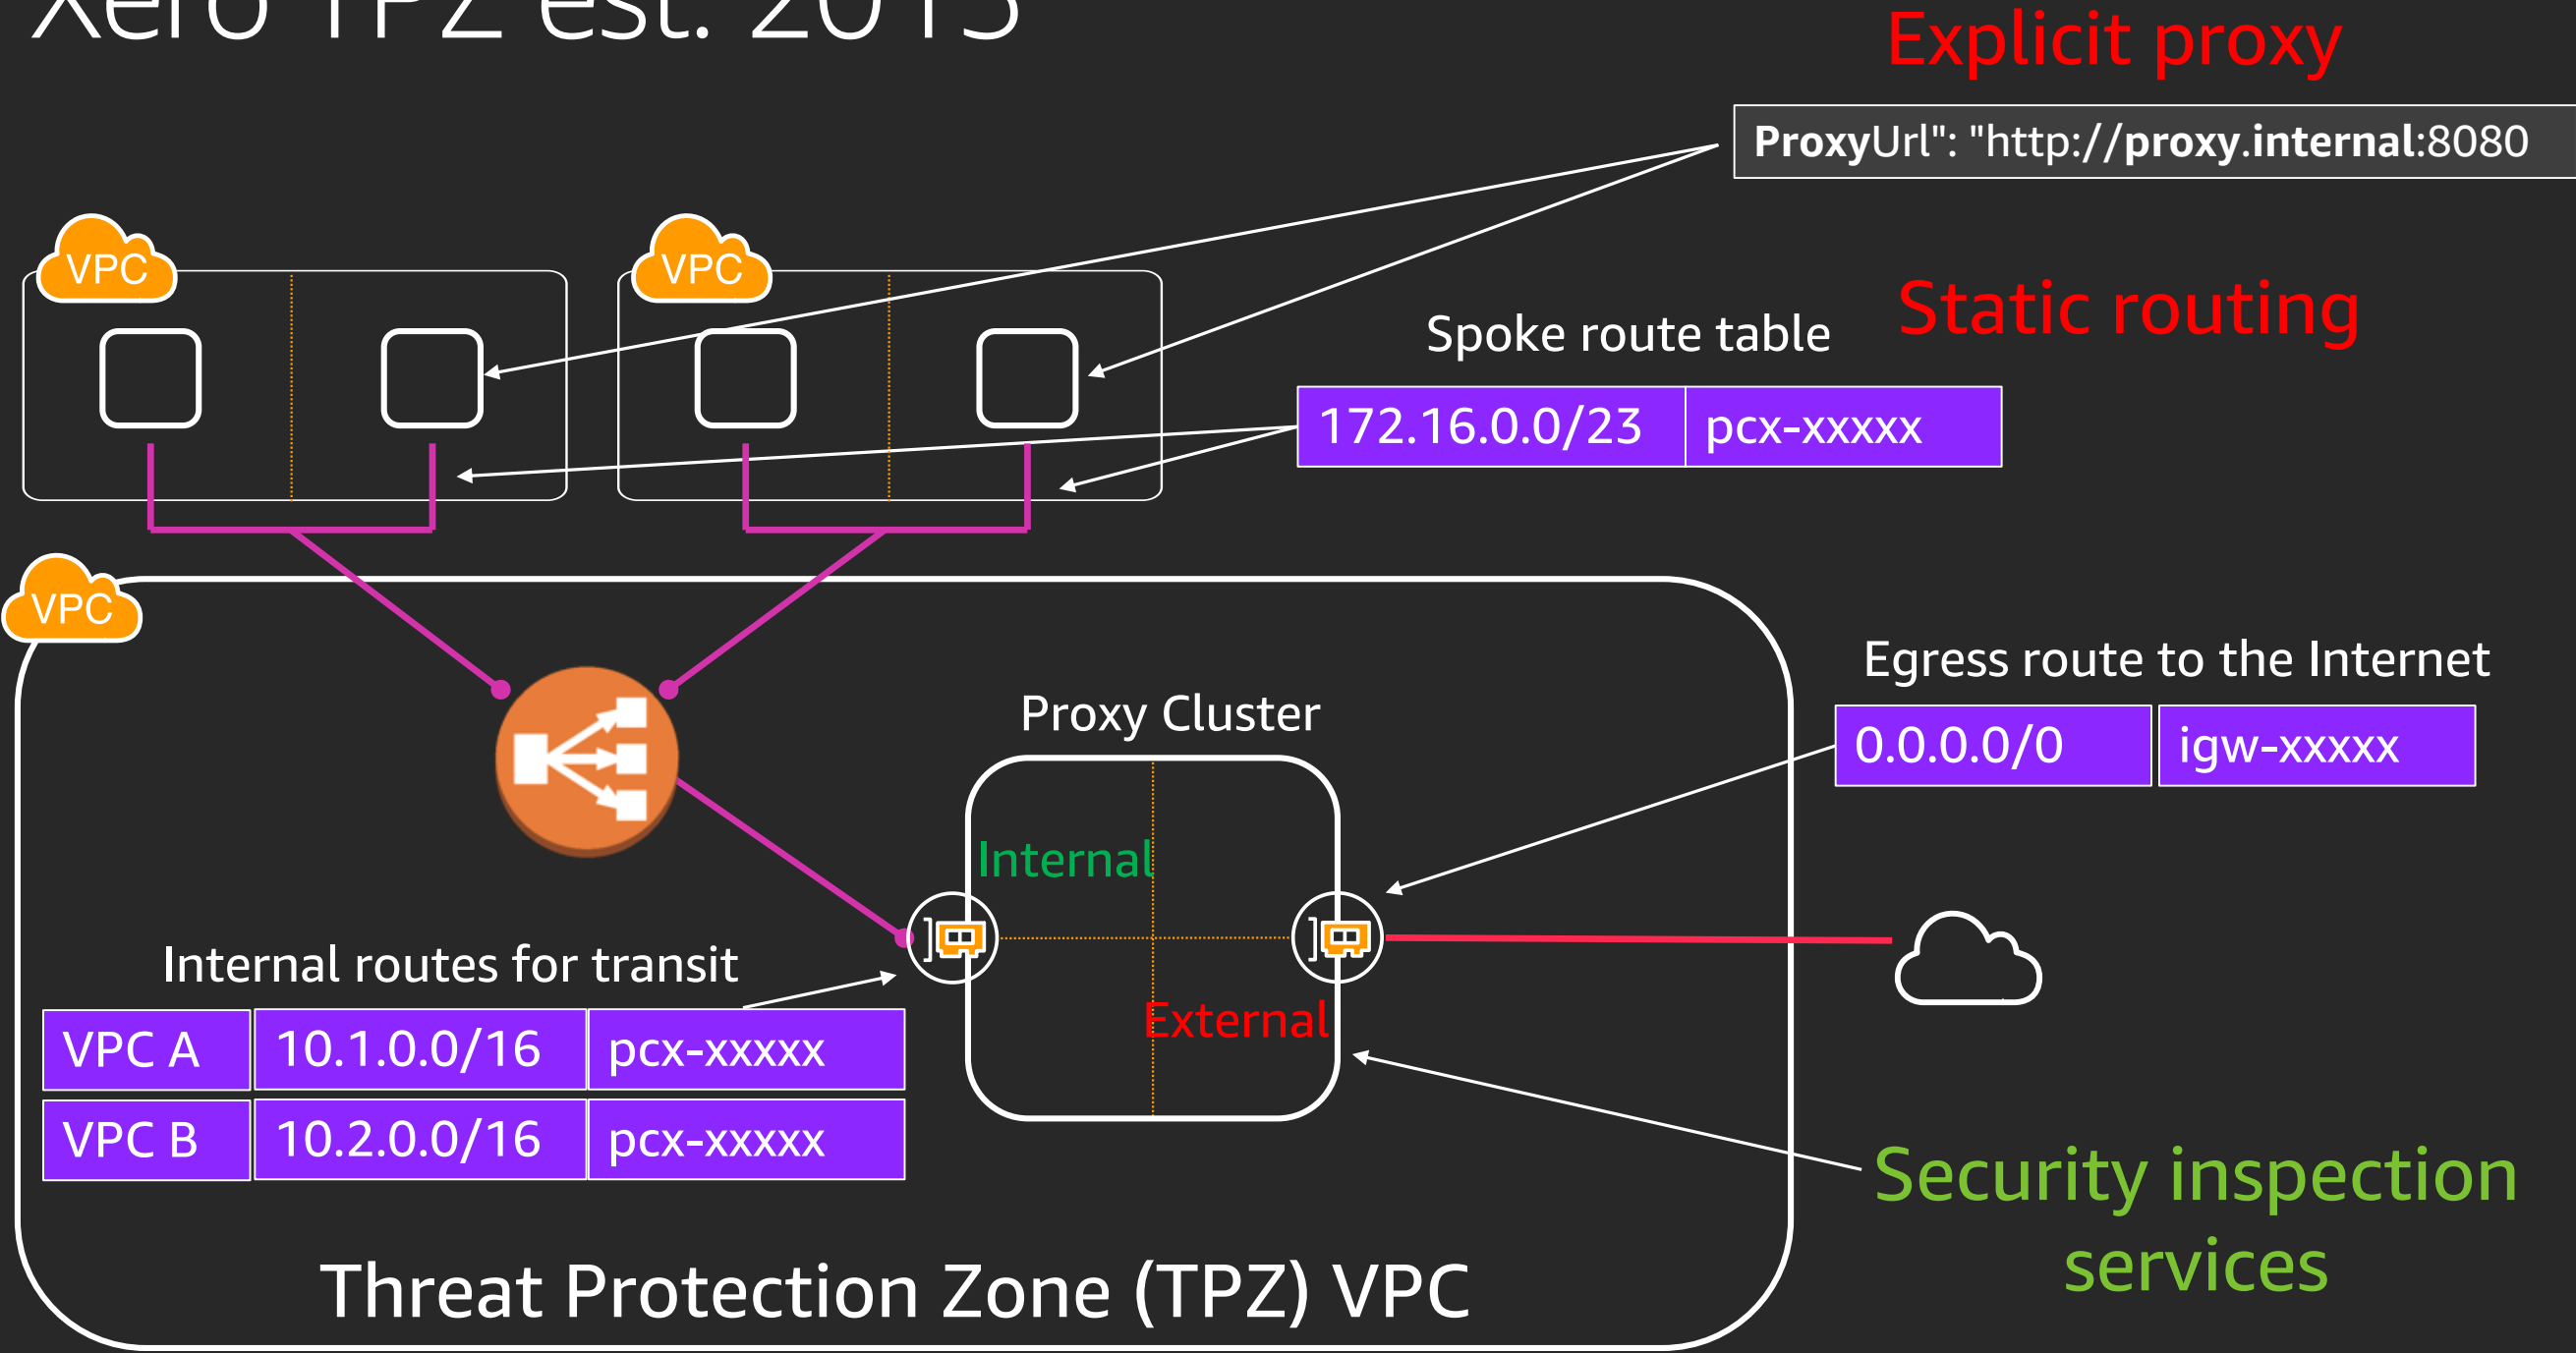
Orchestration: Dev & prod isolated transit network



Check Point Auto-Scaling integration




Xero TPZ est. 2015




Xero TPZ future state

Security inspection services

Security inspection services






Account
Strategy



Segmentation Model



Shared Services



Connectivity




Network
Services

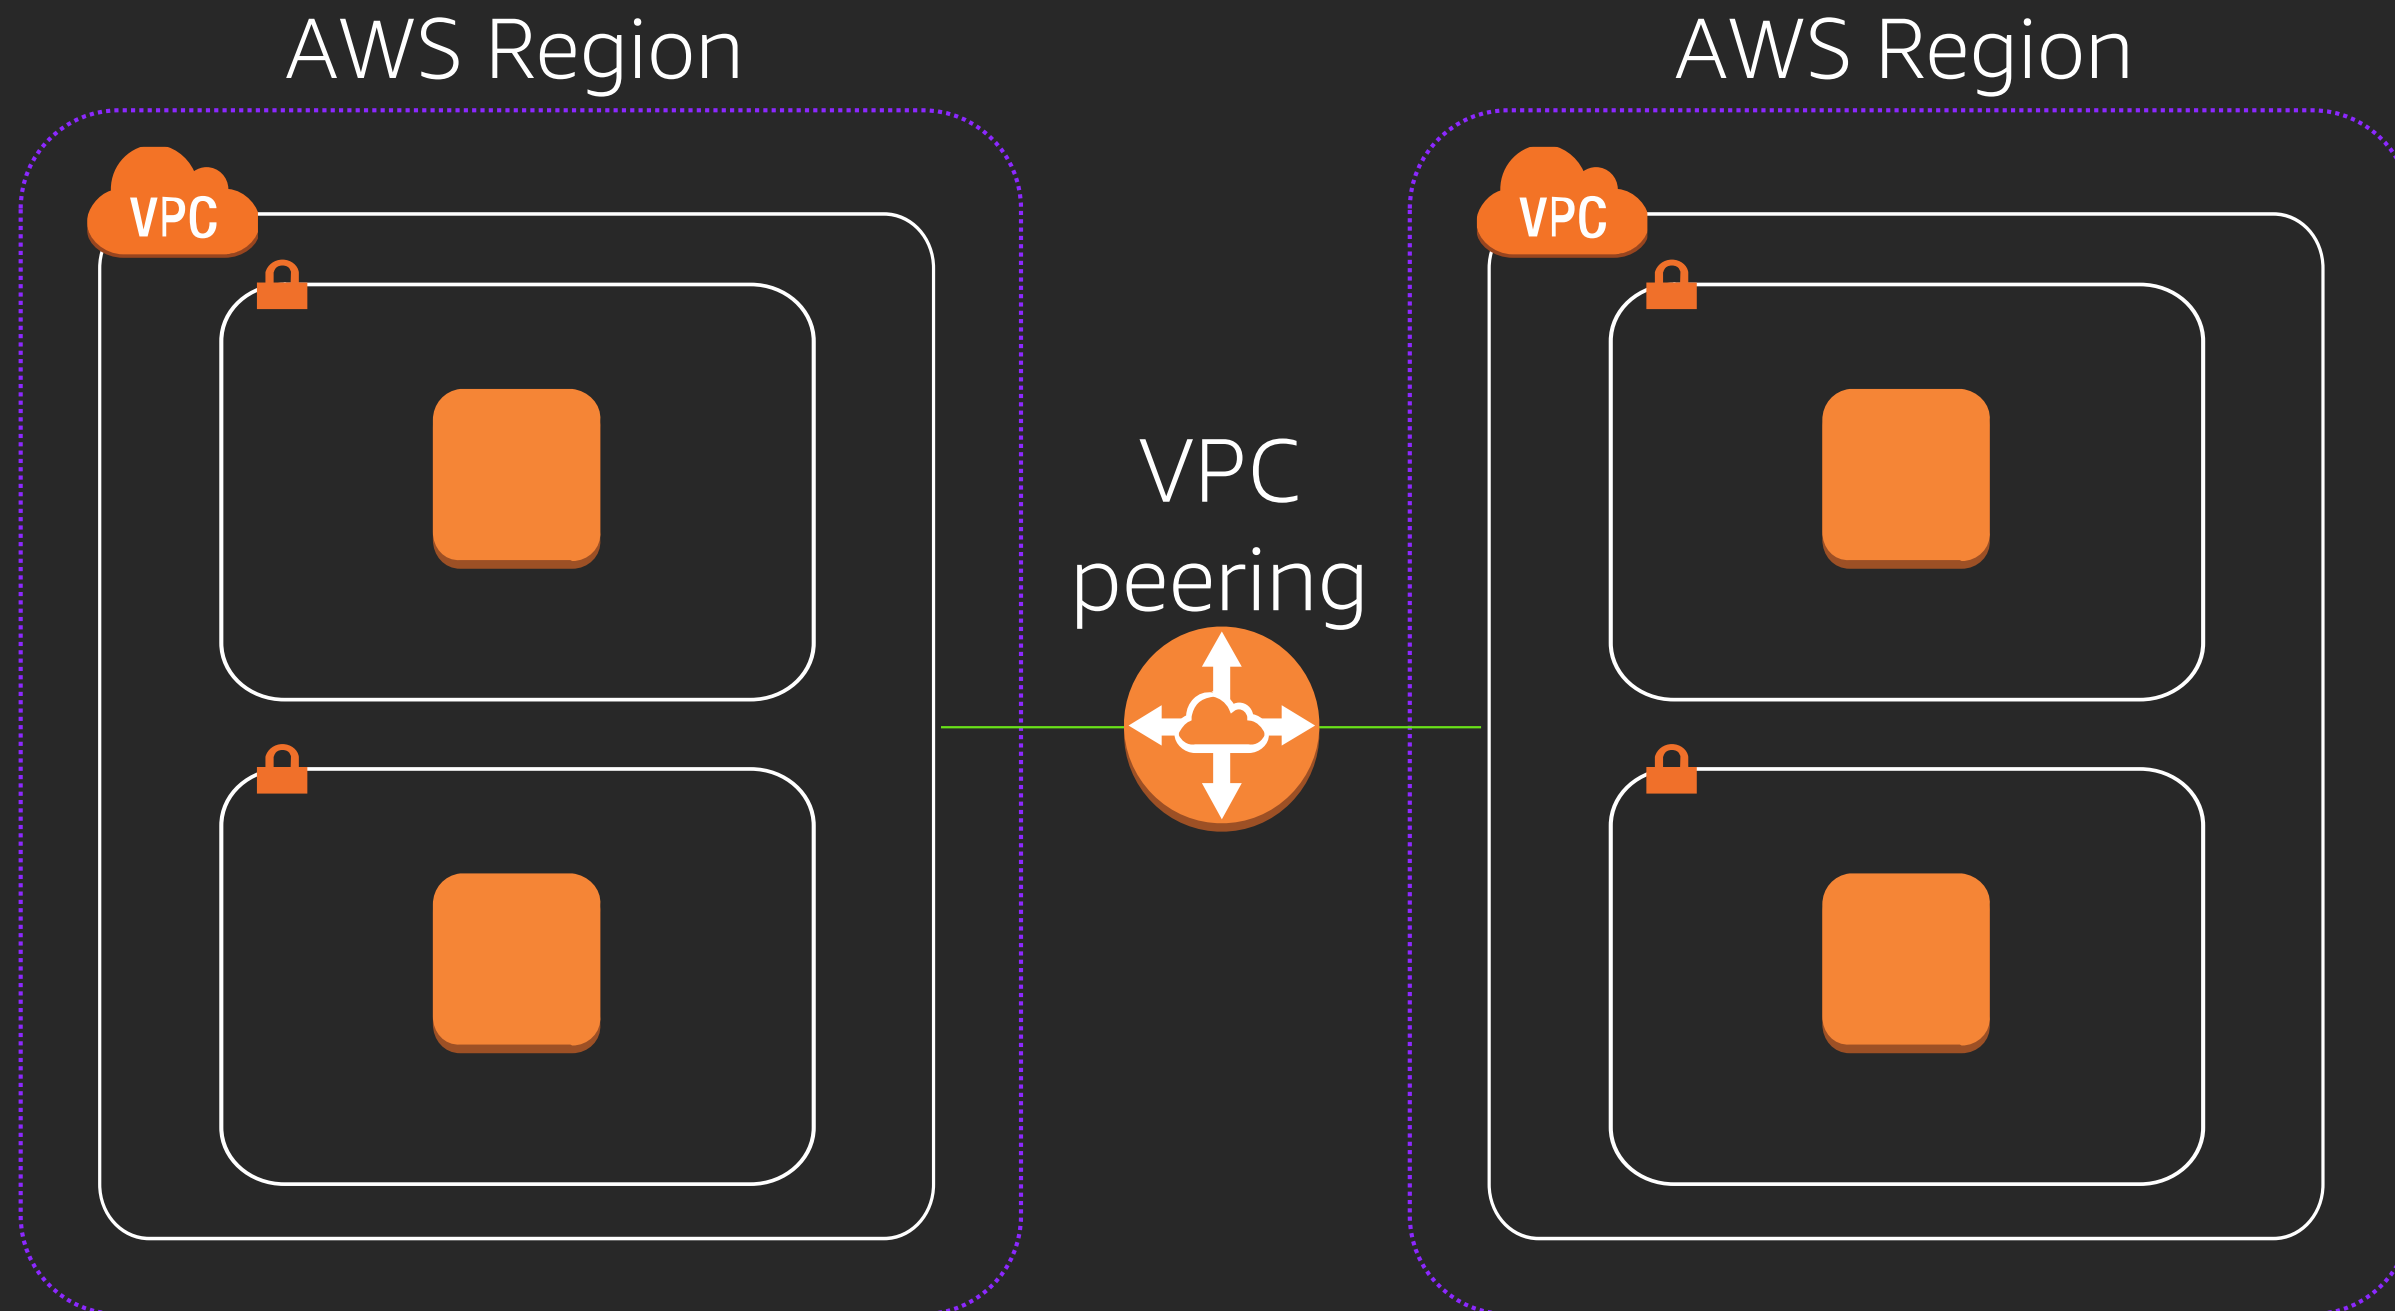


Multi-Region
Options

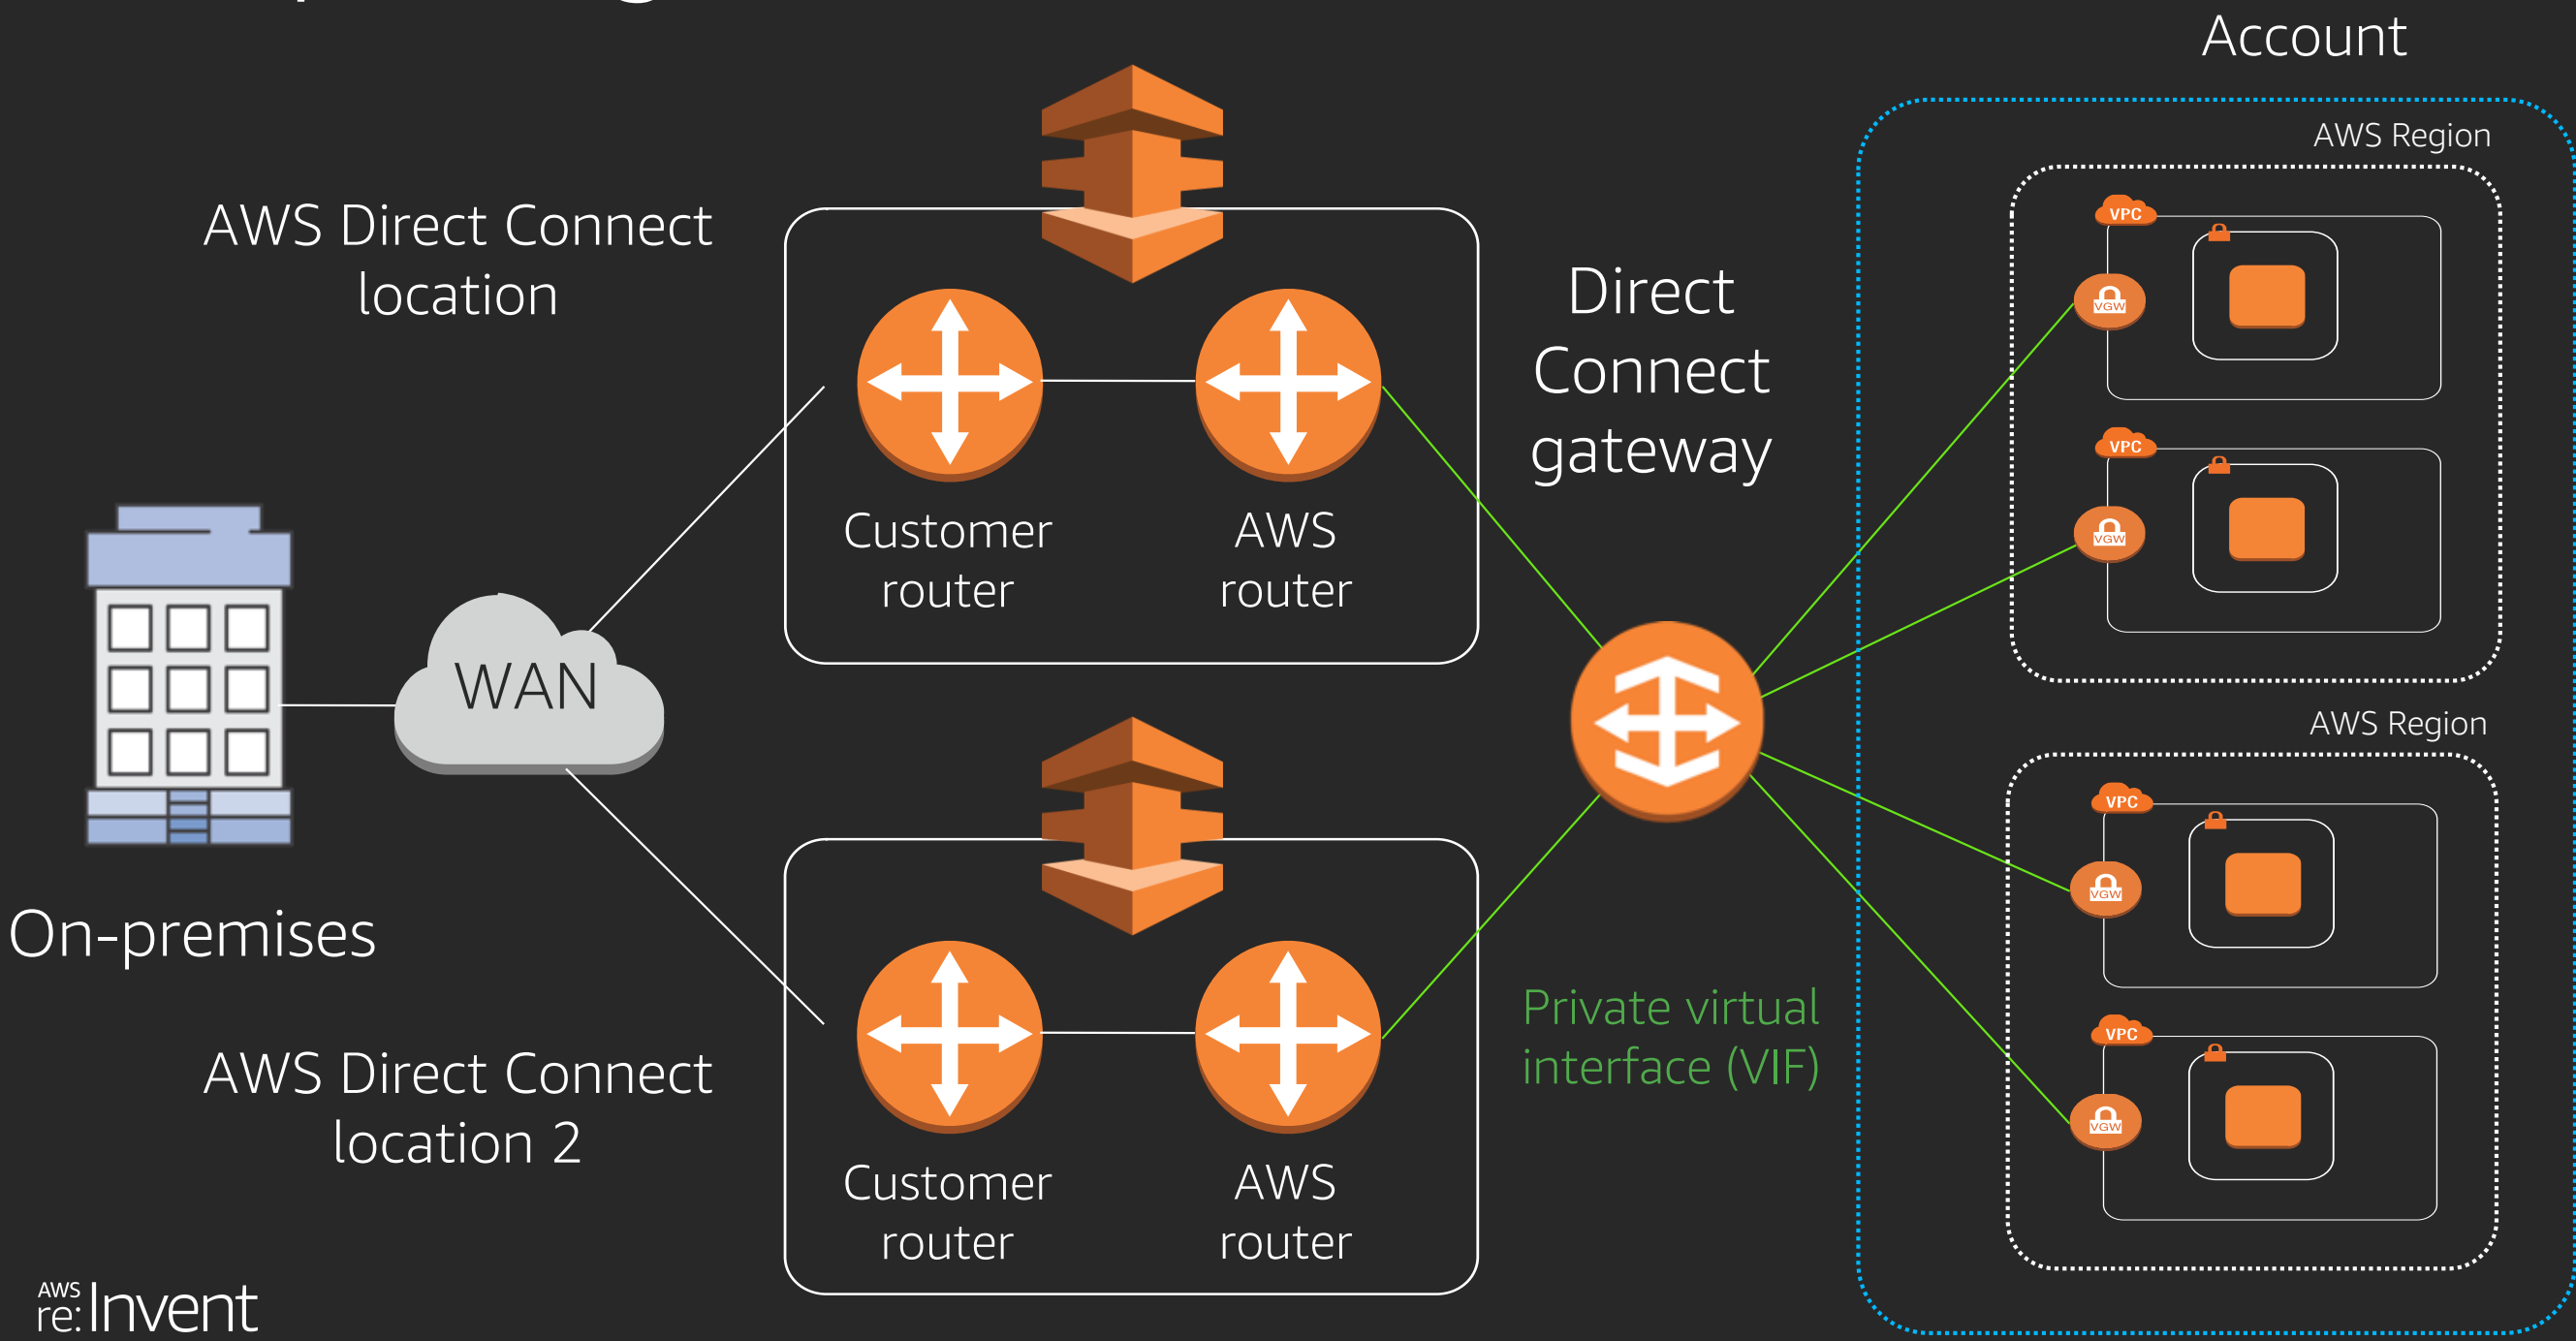
Multiple Regions



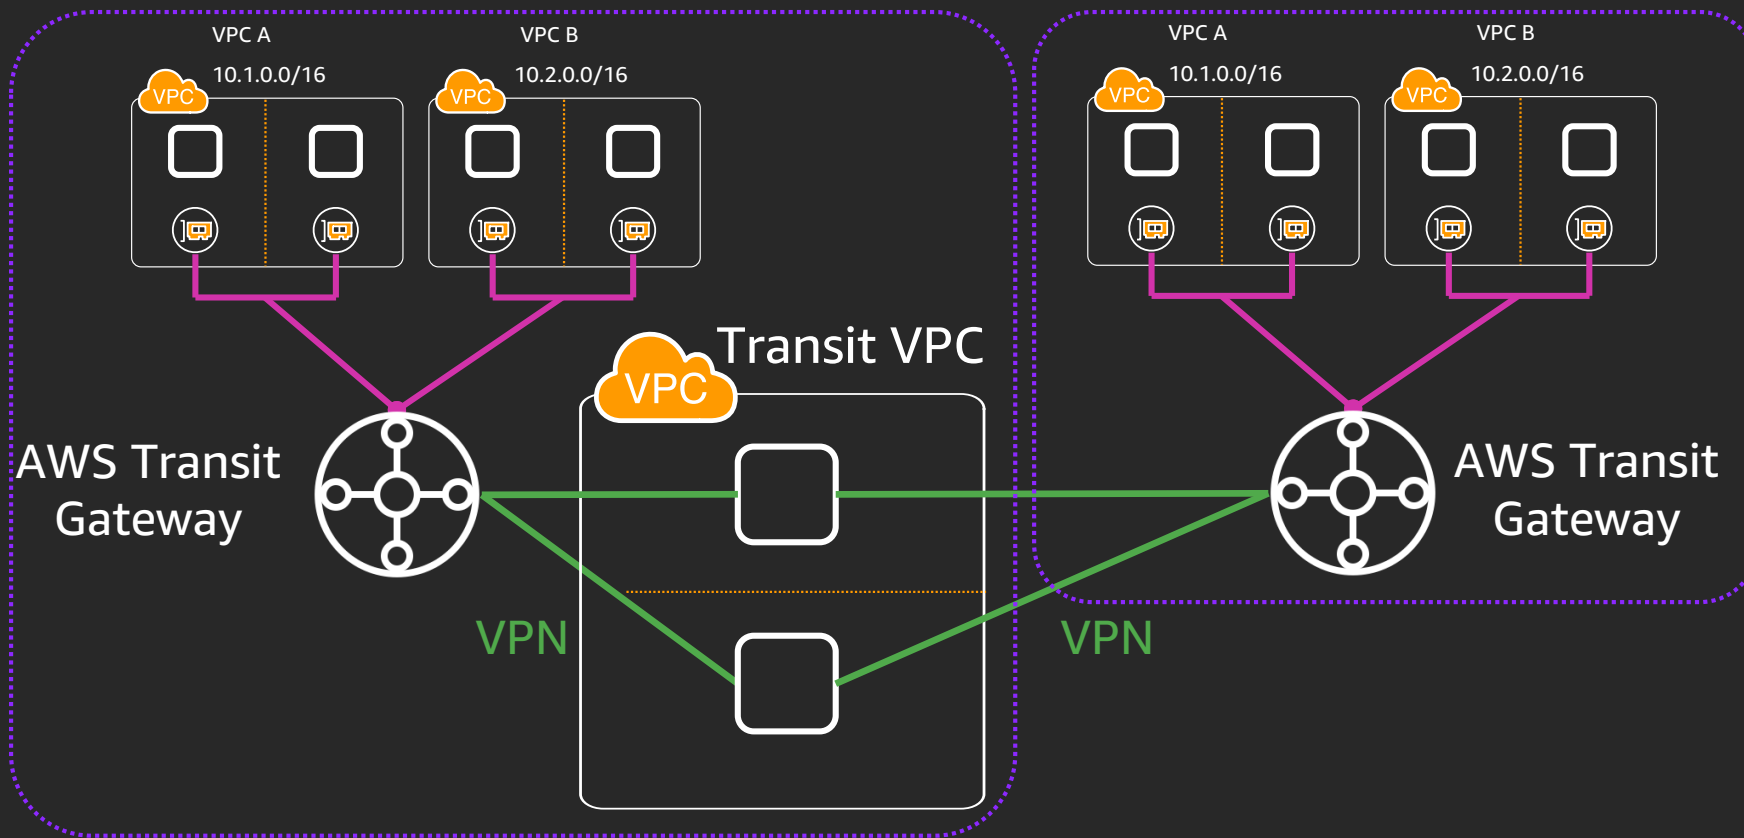
Inter-region VPC peering



Multiple Regions



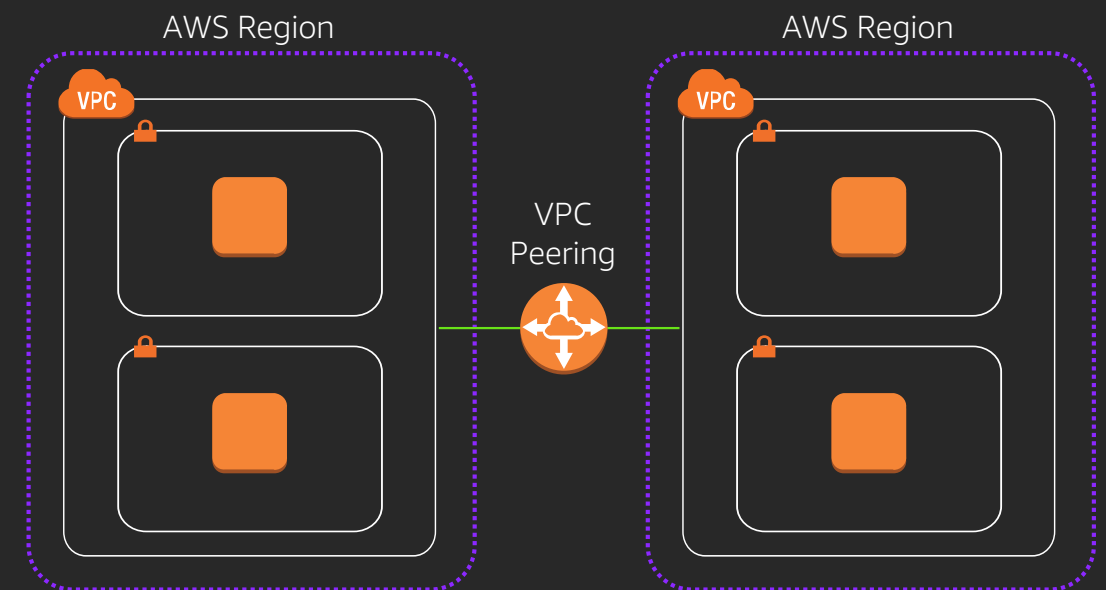
Transit Gateway in multiple Regions



Connecting Regions with VPN

Transit Gateway inter-region support coming soon!

Inter-region peering



Conclusions

Takeaways

We have tools and architectures that horizontally **scale to many VPCs**

There's **wiggle room** for your specific use cases

Use services in combination to **meet scale and security requirements**

Advice

- Networking changes fast, **no more crystal balls**
- **Start simple!** Stay simple. Reduce complexity to smaller scopes
- Segment and modify as needed
- Experiment and test



Thank you!

Nick Matthews
@nickpowpow



Please complete the session
survey in the mobile app.