

# PayPal Data Lake journey

Deepak Mohanakumar Chandramouli, Big Data Platform Eng



## Teradata

- Over 12 years of experience building Data Warehouse Solutions.
- 9+ years of experience working with Teradata based Warehouse Applications

## 2016

- Engineering - Foundational Big Data Applications for PayPal's Data Lake
- Apache Spark Based Ingestion Framework, Data Quality Framework, Dedup & Snapshot processes for Transactional Datasets.

## 2017

- Big Data Platform Engineering: to build PayPal's Gimel - Data Platform
- Gimel powers Data Lake vision & many other Big Data Engineering Use-cases in PayPal Inc.

- PayPal at Scale
  - Business
  - Technology
- PayPal - Datalake
- Compute
  - Challenges
  - Solution
- PayPal Gimel - Bigger Picture
- THE EDGE OF NEXT
- Data Access
  - Challenges
  - Solution
- PayPal Gimel – Next Steps



# PayPal at Scale

Business &  
Technology


# PayPal Scale | Business

- One of the world's largest internet payment companies
- PayPal platform includes **Braintree**, **Venmo**, **Paydiant**, **PP Credit** and **Xoom**

## Q3 2017 Results



## 2016 Full Year Results






# PayPal Scale | Technology



95+ PB  
Data



40,000+  
Yarn Jobs Per Day



6+ compute  
Supported - Hive, Pig, MR, Spark, Beam, Presto  
Others - Flink

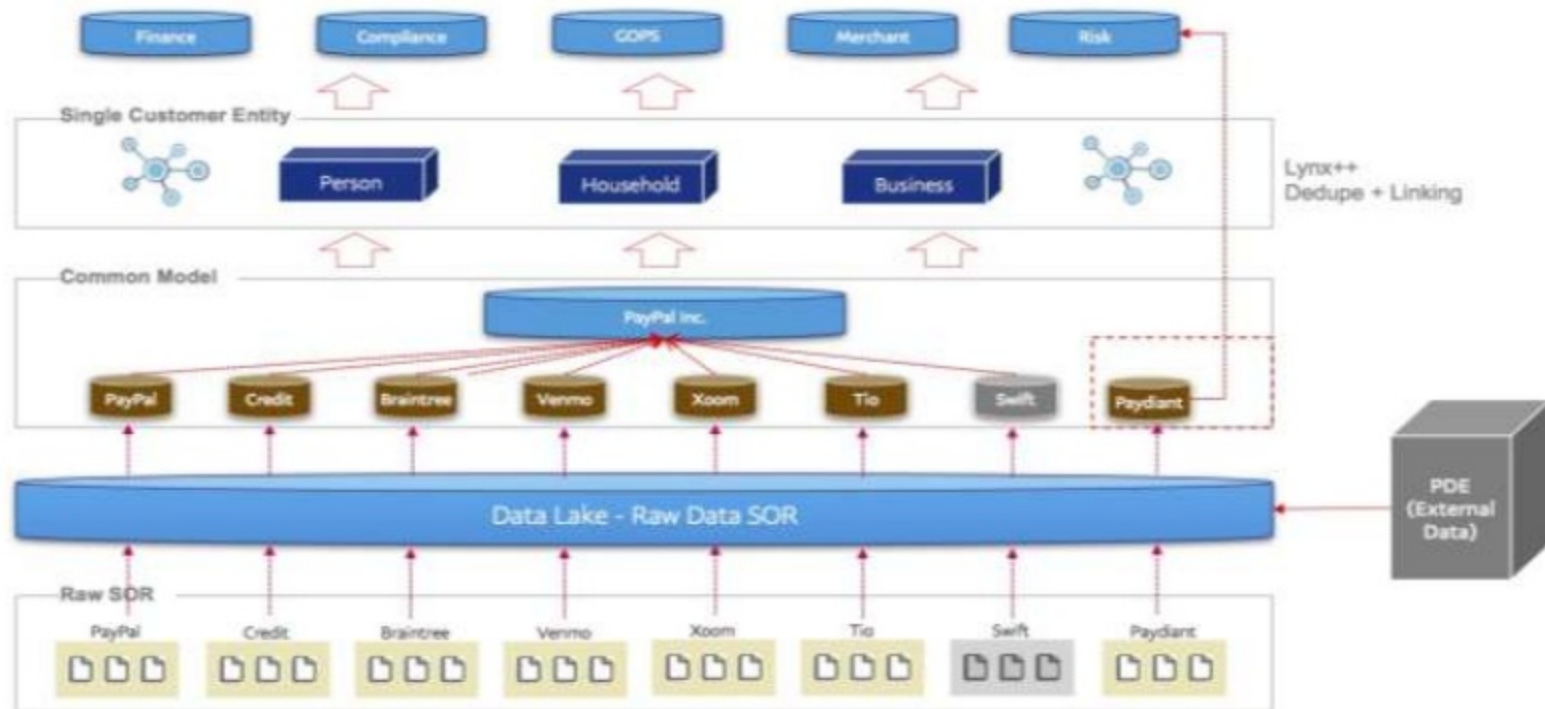
# PayPal - Datalake

Bigger Picture  
Complexities

# Data Lake | Bigger Picture

## Cross Property Data Lake – Conceptual View

Single source for all of PayPal and its subsidiaries data at the most granular level to serve the various needs across PayPal as well as the subsidiaries.





# Data Lake | Recap | Complexities

---

## Cross Property Integration

- **PayPal, Braintree, Venmo, Xoom, Credit, Tio, Paydient, More** to Come in future
- Each Property – Unique set of infrastructure & Storages
- Integration & Data Pipelines - involves **Multiple Storages** such as Oracle, MySQL, Kafka, Cloud Bases systems.

## Compute / Processing Capabilities

- Multiple Clusters – Each has a purpose
- Large Scale / **Complex Data Systems**
- **Variety of Use Cases** - Stream, Batch, Interactive, ML
- **Multiple Compute Options** – Hive, Spark, Flink, Beam, Presto, Map-Reduce
- Data & Operational **Security** !

## Data Engineering

- **Multiple Storages** – Cloud, No-SQL, RDBMS, Messaging, In-Memory
- **100's of APIs** to access different storages – SQL & Programmatic (Python, Scala, Java...)
- **Complexity** – Many-fold if one considers different compute engines – spark, hive, flink
- Time to develop and Time to Market – goes through "**learning curves**"

## Data Catalog

- **Puzzle**
- Awareness - **in what form & where data is available?**
- "Storage-specific" dataset creation
- **Dataset onboarding** – various ways
- Data assets – Lacking "**Centralized**" monitoring capabilities

2016

Beginning of the Data Lake Services

# Data Lake Services |2016



Simple, Governed, Unified.

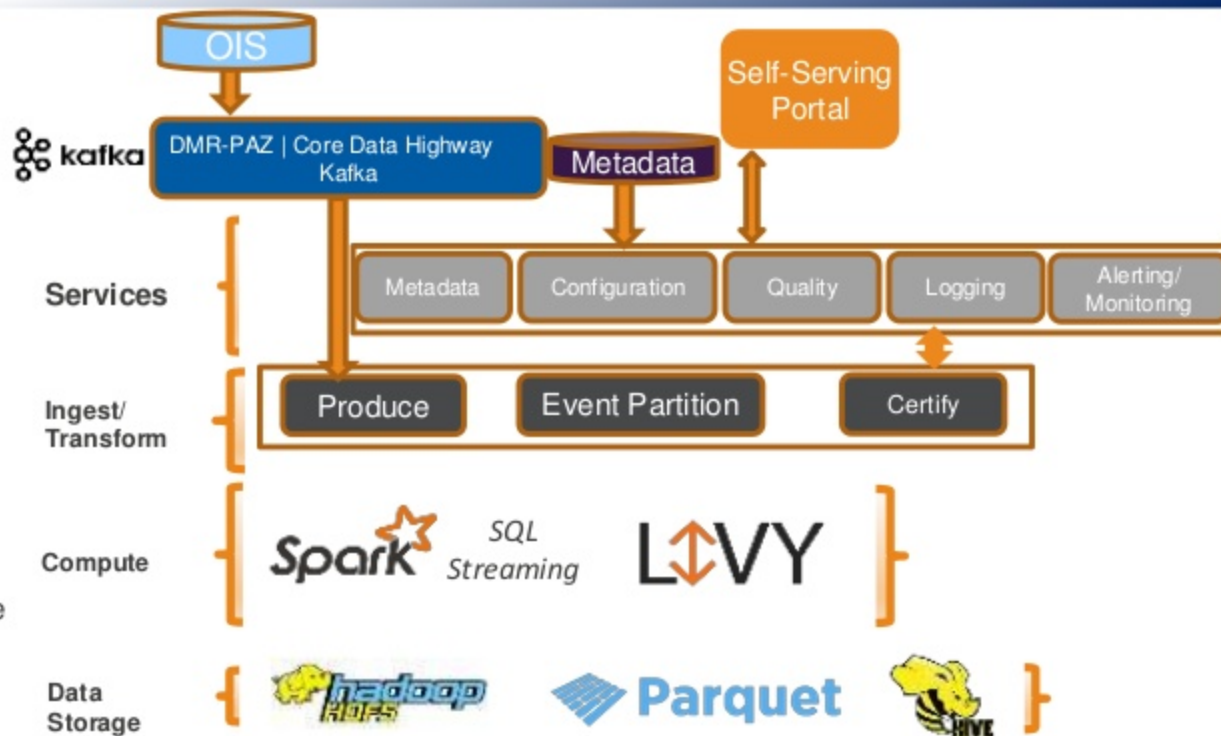
Quick Recap - Architecture

## Key Features

- Configuration/Metadata driven apps
- Spark based framework
- Self-Service Portal
- Data Quality Monitoring
- Metrics Collection
- Operational Dashboards

## Benefits

- Scalability
- User Experience
- High Speed Data Ingestion
- Converged Data Sources (PayPal, Braintree etc.)
- Secure and Governed Data



# Data Lake Services |2016 | Recap Challenges

## Data Governance

- Security
- Metadata
- Data Quality & Metrics

## Use cases

- Rapid growth rate of new business use cases
- Onboarding

## Data Lineage

- Most common Challenge !

## Data Models & Best Practices

- Data Layout
- Data Models

## Data across Property - Integration

- PayPal, Credit, Braintree, Venmo, Paydiant, TIO, Xoom

## Technology & Skillset

- Emerging/New technologies. Example - Spark, Flink, Beam, Kylo, ...
- Talent & Learning Curve

## Best Practices & Common Frameworks

- Logging
- Monitoring
- Alerting
- Auditing
- Operational Dashboards

## Administration

- Hosts Management
- Software Management

## Users Experience

- Data Engineers
- Analysts
- Scientists

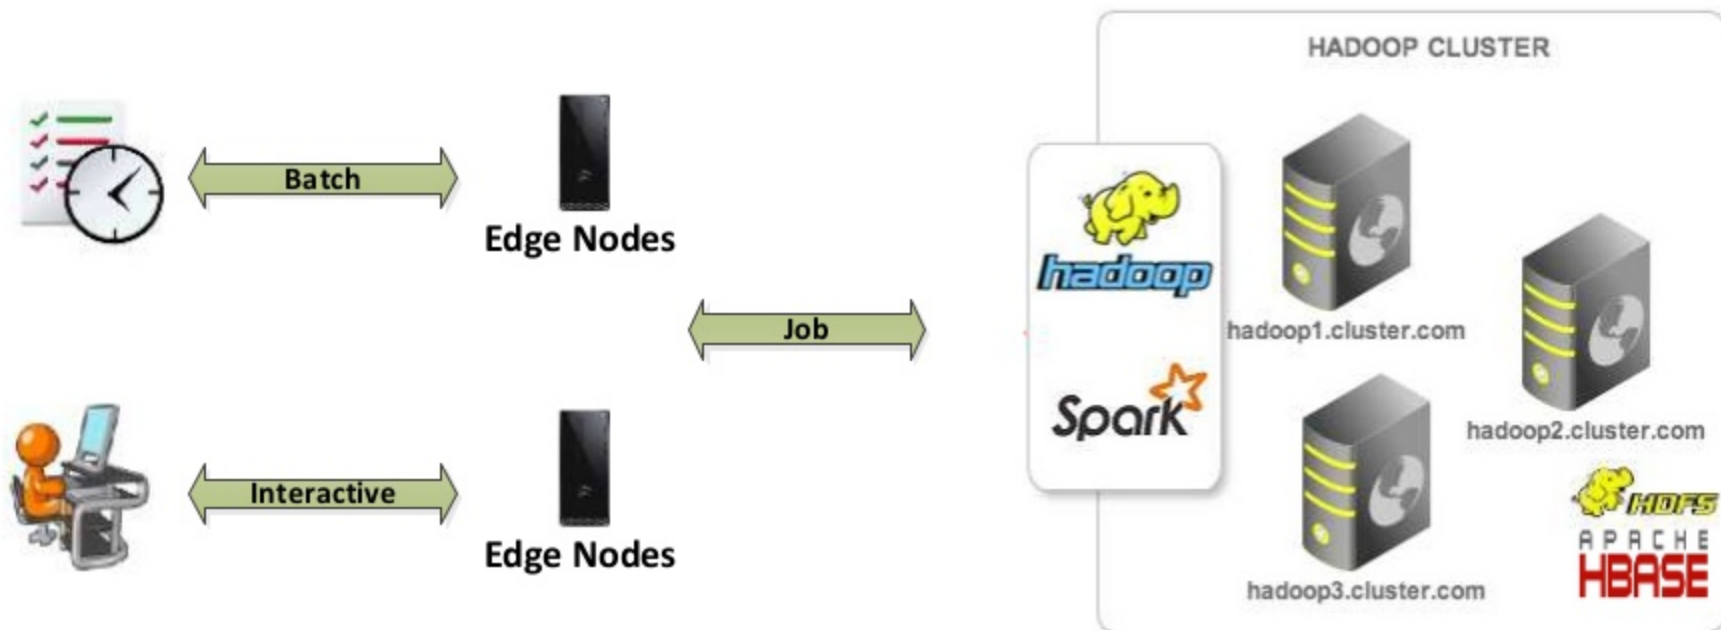
*Challenges Faced with...*

Compute




# Spark on Yarn| Typically

## Deployment




## Administrators




- Need extensive support and maintenance for CLI
- Need to deploy entire stack of software
- Need to sync configurations across systems
- Need extensive testing of jobs before any upgrade

## Developers




- No REST-friendly API
- No low-latency/sub-seconds execution
- No cache sharing across jobs
- No modularity and easy-restartability

## Analysts/Scientists



- No easy way of interactive applications
- No multi-tenancy support and private workspace
- No direct spark sql execution
- No Kerberos integration

## Operations/Security



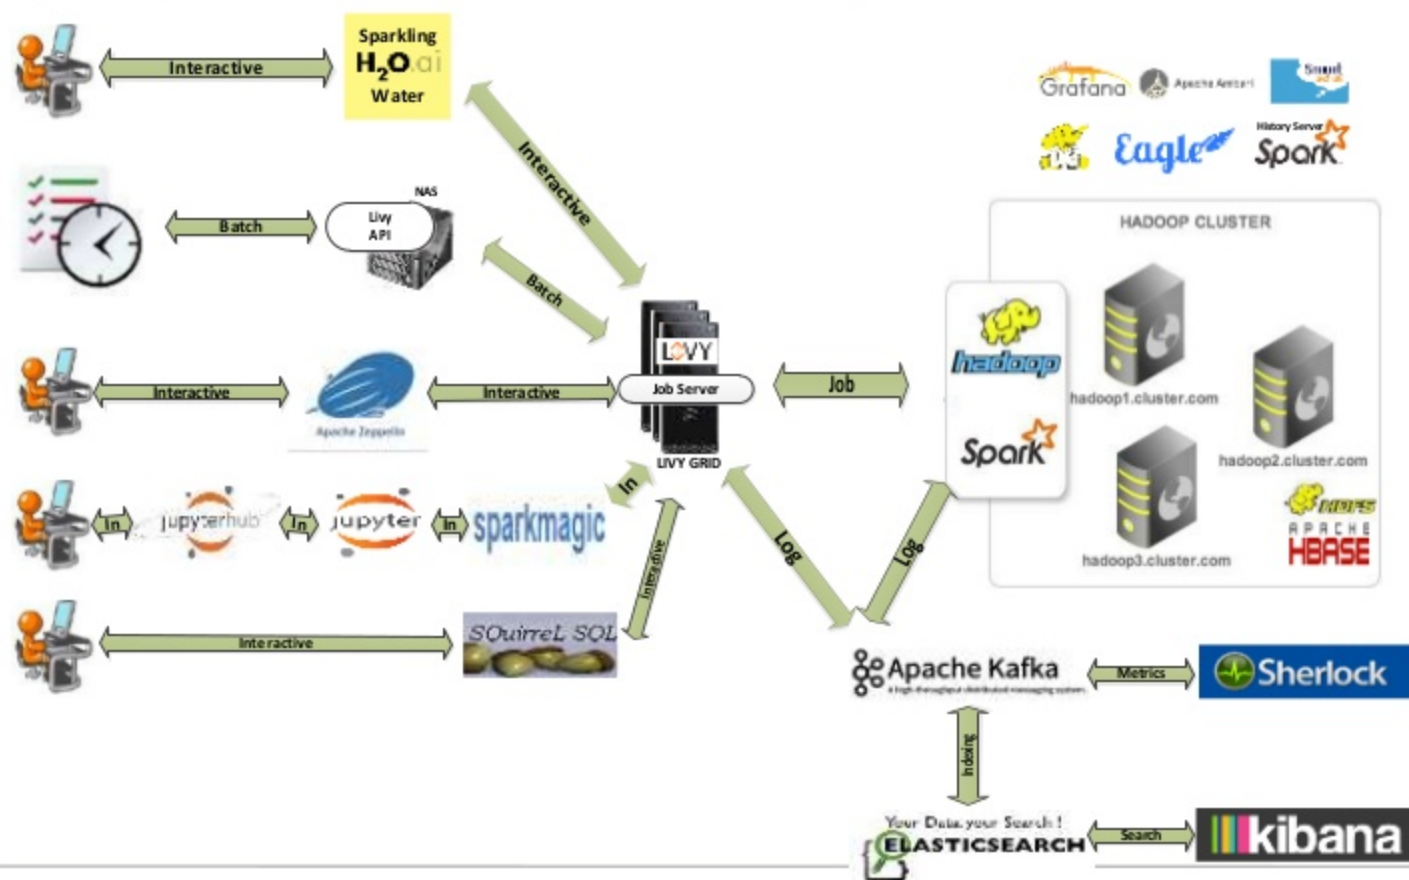
- Different ways of jobs execution and coding standards
- No uniform logging, monitoring and alerting
- Limited audit and control
- No statement level history or metrics

*Solution !*

# Gimel Compute Platform



# Data Lake | Compute Platform



## Administrators

- ✓ Less maintenance on CLI
- ✓ Deploy software stack only on Job Server
- ✓ Configurations at one place
- ✓ Easy platform/software upgrade

## Developers

- ✓ REST-friendly and Docker-friendly
- ✓ Low-latency/sub-seconds execution
- ✓ **Sharing** cache across jobs
- ✓ Modularity and easy restartability

## Analysts/Scientists

- ✓ User friendly interactive applications
- ✓ Multi-tenancy and Private workspace
- ✓ Direct spark sql execution
- ✓ Kerberos Support

## Operations/Security

- ✓ Standardized coding and unified execution
- ✓ Uniformed logging, monitoring and alerting
- ✓ Fine-grained audit
- ✓ Complete statement level history and metrics

*Challenges Faced with...*

Data Access

# Data Platform | Data Access | Current State

## Spark Read From Hbase

```
import org.apache.hadoop.hbase.client.{HBaseAdmin, Result}
import org.apache.hadoop.hbase.{ HBaseConfiguration, HTableDescriptor }
import org.apache.hadoop.hbase.mapreduce.TableInputFormat
import org.apache.hadoop.hbase.io.ImmutableBytesWritable

import org.apache.spark._

object HBaseRead {
  def main(args: Array[String]) {
    val sparkConf = new SparkConf().setAppName("HBaseRead").setMaster("local[2]")
    val sc = new SparkContext(sparkConf)
    val conf = HBaseConfiguration.create()
    val tableName = "table1"

    System.setProperty("user.name", "hdfs")
    System.setProperty("HADOOP_USER_NAME", "hdfs")
    conf.set("hbase.master", "localhost:60000")
    conf.setInt("timeout", 120000)
    conf.set("hbase.zookeeper.quorum", "localhost")
    conf.set("zookeeper.znode.parent", "/hbase-unsecure")
    conf.set(TableInputFormat.INPUT_TABLE, tableName)

    val admin = new HBaseAdmin(conf)
    if (!admin.isTableAvailable(tableName)) {
      val tableDesc = new HTableDescriptor(tableName)
      admin.createTable(tableDesc)
    }

    val hBaseRDD = sc.newAPIHadoopRDD(conf, classOf[TableInputFormat], classOf[ImmutableBytesWritable], classOf[Result])
    println("Number of Records Found : " + hBaseRDD.count())
    sc.stop()
  }
}
```

# Data Platform | Data Access | Current State

## Spark Read From Elastic Search

```
import org.apache.spark.SparkContext
import org.apache.spark.sql.SQLContext
import org.apache.spark.SparkConf
import org.elasticsearch.spark.sql._

object ElasticSearchRead {
  def main(args: Array (link is external)[String (link is external)]) {

    val conf = new SparkConf().setAppName("ReadFromES")

    conf.set("spark.es.nodes", "10.111.222.333")
    conf.set("es.index.auto.create", "true")
    conf.set("spark.es.nodes.client.only", "true")
    conf.set("es.query", "?q=*")
    conf.set("es.resource", "index/type")

    val sc = new SparkContext(conf)

    val sqlContext = new org.apache.spark.sql.SQLContext(sc)

    val es_df=sqlContext.read.format("org.elasticsearch.spark.sql").load("index/type")

    es_df.registerTempTable("things")

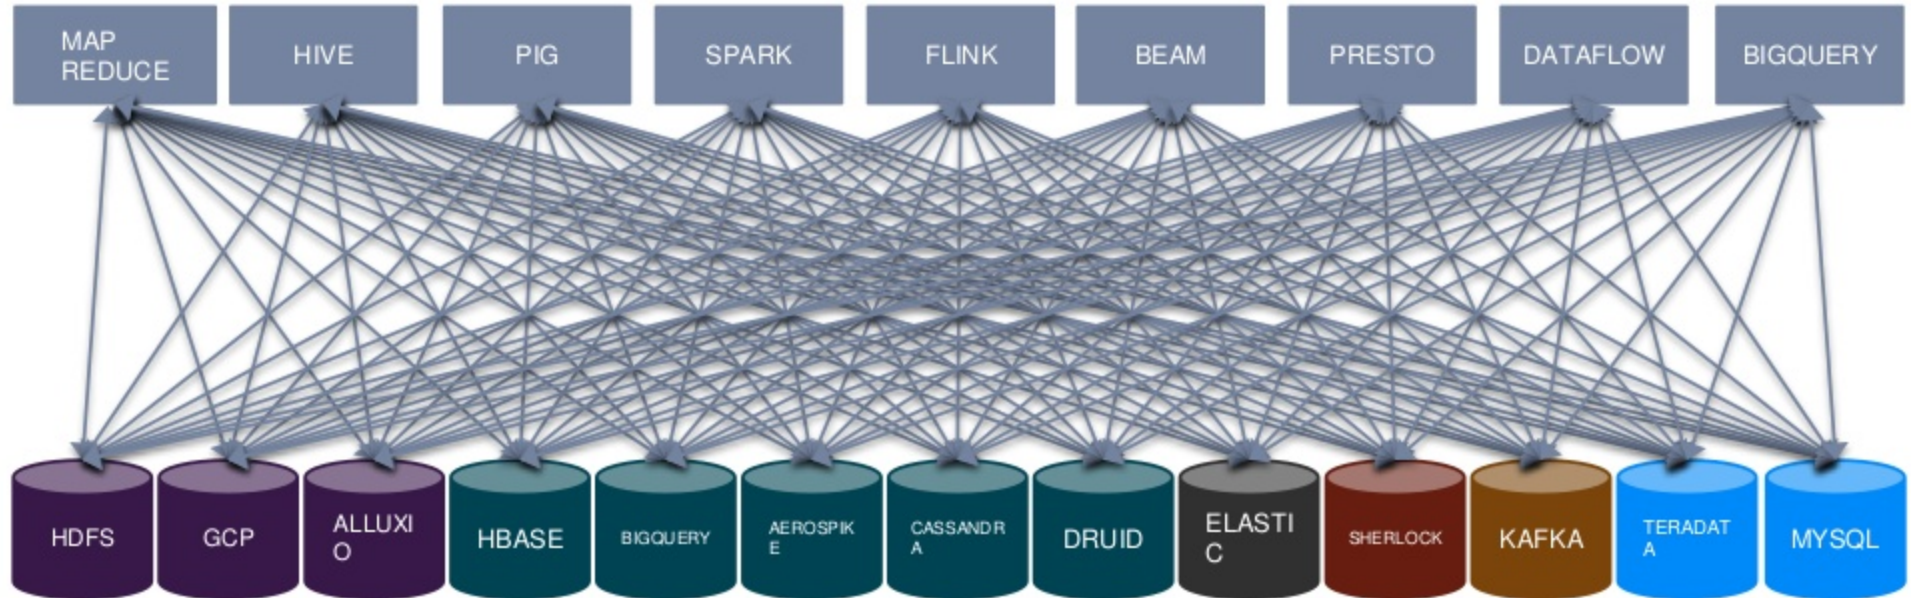
    val filtered_things = sqlContext.sql("select * from things where col = '1000'")

  }
}
```





# Data Access| Bigger Picture | Today





# Data Platform | Today

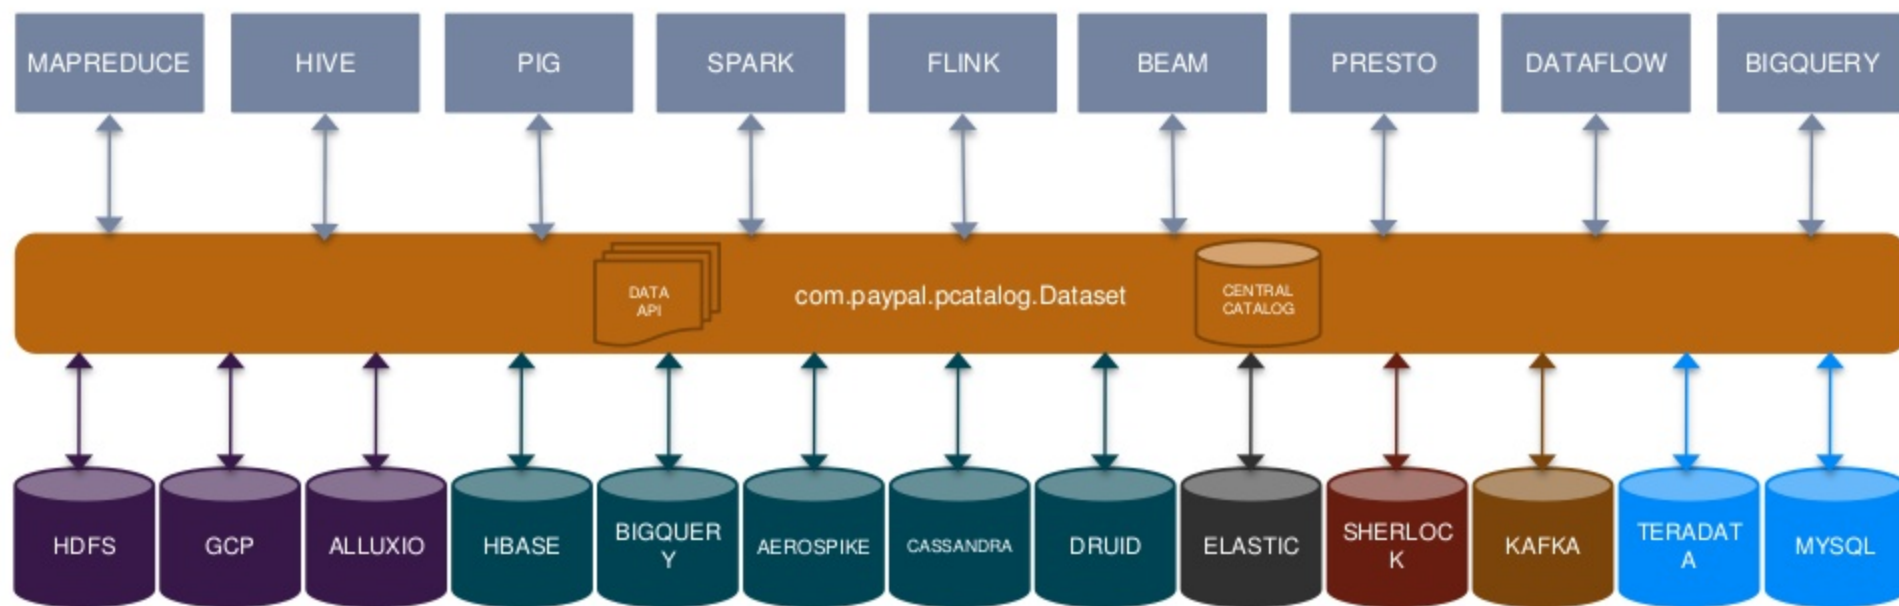
## Data Application Lifecycle - Current State



*Solution !*

# Gimel Data Platform

# Data Platform | Data Access | With Data API



- **Data API**

- Unified Data Access API
- Access from any Compute Engine
- Access any Storage

- **Data Catalog**

- Global Catalog - Explore at one place
- Onboard at one place
- Dataset Management- Track/Approve Onboard Request
- Dataset Stats

## Spark Read From Hbase

```
import org.apache.hadoop.hbase.client.{HBaseAdmin, Result}
import org.apache.hadoop.hbase.{HBaseConfiguration, HBaseDescriptor}
import org.apache.hadoop.hbase.mapreduce.TableInputFormat
import org.apache.hadoop.hbase.io.ImmutableBytesWritable

import org.apache.spark._

object HBaseRead {
  def main(args: Array[String]) {
    val sparkConf = new SparkConf().setAppName("HBaseRead").setMaster("local[*]")
    val sc = new SparkContext(sparkConf)
    val conf = HBaseConfiguration.create()
    val tableInfo = "table1"

    System.setProperty("hadoop.home.dir", "null")
    System.setProperty("hadoop.job.jar", "null")
    conf.set("hbase.master", "localhost:9000")
    conf.set("hbase.zookeeper.quorum", "localhost")
    conf.set("hbase.zookeeper.property.zoo3.version", "3.4.6")
    conf.set("hbase.zookeeper.quorum", "localhost")
    conf.set("hbase.zookeeper.property.zoo3.version", "3.4.6")
    conf.set("hbase.zookeeper.property.zoo3.version", "3.4.6")

    val table = new HBaseTable(conf)
    if (table.existsTable(tableInfo)) {
      val tableDesc = new HBaseDescriptor(tableInfo)
      table.createTable(tableInfo)
    }

    val HBaseRDD = sc.newAPIHadoopRDD(conf, classOf[HBaseTableInputFormat], classOf[ImmutableBytesWritable], classOf[Result])
    sc.stop()
  }
}
```

## Spark Read From AeroSpike

```
import org.apache.spark._
import com.aerospike.spark.sql._
import com.aerospike.client.AerospikeClient
import com.aerospike.client.Bin
import com.aerospike.client.Key
import com.aerospike.client.Value

import org.apache.spark._

object AerospikeRead {
  def main(args: Array[String]) {
    val sparkConf = new SparkConf().setAppName("AerospikeRead").setMaster("local[*]")
    val sc = new SparkContext(sparkConf)
    val sqlContext = new SQLContext(sc)

    val thingsRDD = sqlContext.read
      .format("com.aerospike.spark.sql")
      .option("aerospike.namespace", "ns")
      .option("aerospike.partition", "1000")
      .option("aerospike.namespace", "namespace")
      .option("aerospike.key", "key")
      .load()

    thingsRDD.registerTempTable("things")

    val filteredThings = sqlContext.sql("select * from things where val = 100")
    println("Number of Records found : " + filteredThings.count())
    sc.stop()
  }
}
```

## Spark Read From Elastic Search

```
import org.apache.spark.SparkContext
import org.apache.spark.sql.SQLContext
import org.apache.spark.SparkConf
import org.elasticsearch.spark.sql._

object ElasticSearchRead {
  def main(args: Array[String]) {
    val conf = new SparkConf().setAppName("ElasticSearchRead")

    conf.set("spark.es.nodes", "10.111.222.333")
    conf.set("es.index.auto.create", "true")
    conf.set("spark.es.nodes.client.only", "true")
    conf.set("es.query", "type")
    conf.set("es.resource", "index/type")

    val sc = new SparkContext(conf)
    val sqlContext = new org.apache.spark.sql.SQLContext(sc)

    val es_df = sqlContext.read.format("org.elasticsearch.spark.sql").load("index/type")
    es_df.registerTempTable("things")

    val filteredThings = sqlContext.sql("select * from things where val = 1000")
  }
}
```

## Spark Read From Druid

```
import org.apache.spark._
import com.druid.spark._

import org.apache.spark._

object DruidRead {
  def main(args: Array[String]) {
    val sparkConf = new SparkConf().setAppName("DruidRead").setMaster("local[*]")
    val sc = new SparkContext(sparkConf)
    val sqlContext = new SQLContext(sc)

    val druidRDD = sqlContext.read
      .format("com.druid.spark")
      .option("druid.url", "http://localhost:8888")
      .option("druid.query", "select * from druid")
      .load()

    druidRDD.registerTempTable("things")

    val filteredThings = sqlContext.sql("select * from things where val = 100")
    println("Number of Records found : " + filteredThings.count())
    sc.stop()
  }
}
```



## With Data API

```
val df = Dataset.Read("dataset_from_hbase")

val df = Dataset.Read("dataset_from_elastic")

val df = Dataset.Read("dataset_from_aerospike")

val df = Dataset.Read("dataset_from_druid")

val df = Dataset.Read("dataset_from_hdfs")

val df = Dataset.Read("dataset_from_kafka")

val df = Dataset.Read("dataset_from_sherlock")

val df = Dataset.Read("dataset_from_cassandra")

val df = Dataset.Read("dataset_from_solr")

val df = Dataset.Read("dataset_from_mysql")

val df = Dataset.Read("dataset_from_teradata")

val df = Dataset.Read("dataset_from_oracle")

val df = Dataset.Read("dataset_from_bigquery")

val df = Dataset.Read("dataset_from_bigtable")

val df = Dataset.Read("dataset_from_alluxio")
```



## Spark Read From Hbase

```
import org.apache.hadoop.hbase.client.{HBaseAdmin, Result}
import org.apache.hadoop.hbase.{HBaseConfiguration, HBaseDescriptor}
import org.apache.hadoop.hbase.mapreduce.TableInputFormat
import org.apache.hadoop.hbase.io.ImmutableBytesWritable

import org.apache.spark._

object HBaseRead {
  def main(args: Array[String]) {
    val sparkConf = new SparkConf().setAppName("HBaseRead").setMaster("local[*]")
    val sc = new SparkContext(sparkConf)
    val conf = HBaseConfiguration.create()
    val tableDesc = "table1"

    System.setProperty("hadoop.job.name", "hdfs")
    System.setProperty("hadoop.job.jar", "hdfs")
    conf.set("hbase.master", "localhost:9090")
    conf.set("hbase.zookeeper.quorum", "localhost")
    conf.set("hbase.zookeeper.property", "localhost")
    conf.set("hbase.zookeeper.quorum", "localhost")
    conf.set("hbase.zookeeper.property", "localhost")
    conf.set("hbase.zookeeper.property", "localhost")

    val table = new HBaseTable(conf)
    if (table.exists(tableDesc)) {
      val tableDesc = new HBaseTableDesc(tableDesc)
      table.create(tableDesc)
    }

    val HBaseRead = sc.newHBaseTable(conf, classOf[HBaseTableFormat], classOf[ImmutableBytesWritable], classOf[Result])
    sc.stop()
  }
}
```

## Spark Read From AeroSpike

```
import org.apache.spark._
import com.aerospike.spark.sql._
import com.aerospike.client.AerospikeClient
import com.aerospike.client.Bin
import com.aerospike.client.Key
import com.aerospike.client.Value

import org.apache.spark._

object AerospikeRead {
  def main(args: Array[String]) {
    val sparkConf = new SparkConf().setAppName("AerospikeRead").setMaster("local[*]")
    val sc = new SparkContext(sparkConf)
    val sqlContext = new SQLContext(sc)

    val thingsRDD = sqlContext.read
      .format("com.aerospike.spark.sql")
      .option("aerospike.url", "127.0.0.1")
      .option("aerospike.port", "3000")
      .option("aerospike.namespace", "namespace")
      .option("aerospike.key", "key")
      .load()

    thingsRDD.registerTempTable("things")

    val filteredThings = sqlContext.sql("select * from things where key = 'key'")
    println("Number of Records found : " + filteredThings.count())
    sc.stop()
  }
}
```

## Spark Read From Elastic Search

```
import org.apache.spark.SparkContext
import org.apache.spark.sql.SQLContext
import org.apache.spark.SparkConf
import org.elasticsearch.spark.sql._

object ElasticSearchRead {
  def main(args: Array[String]) {
    val conf = new SparkConf().setAppName("ElasticSearchRead")

    conf.set("spark.es.nodes", "10.111.222.333")
    conf.set("es.index.auto.create", "true")
    conf.set("spark.es.nodes.client.only", "true")
    conf.set("es.query", "type")
    conf.set("es.resource", "index/type")

    val sc = new SparkContext(conf)

    val sqlContext = new org.apache.spark.sql.SQLContext(sc)

    val es_df = sqlContext.read.format("org.elasticsearch.spark.sql").load("index/type")
    es_df.registerTempTable("things")

    val filteredThings = sqlContext.sql("select * from things where col = 'value'")
  }
}
```



## With Gimet SQL

SELECT \* FROM dataset\_from\_hbase

SELECT \* FROM dataset\_from\_elastic

SELECT \* FROM dataset\_from\_aerospike

SELECT \* FROM dataset\_from\_druid

SELECT \* FROM dataset\_from\_hdfs

SELECT \* FROM dataset\_from\_kafka

SELECT \* FROM dataset\_from\_sherlock

SELECT \* FROM dataset\_from\_cassandra

SELECT \* FROM dataset\_from\_solr

SELECT \* FROM dataset\_from\_mysql

SELECT \* FROM dataset\_from\_teradata

SELECT \* FROM dataset\_from\_oracle

SELECT \* FROM dataset\_from\_bigquery

SELECT \* FROM dataset\_from\_bigtable

SELECT \* FROM dataset\_from\_alluxio



## All Kind of Hive SQLs Supported !

### Batch SQL



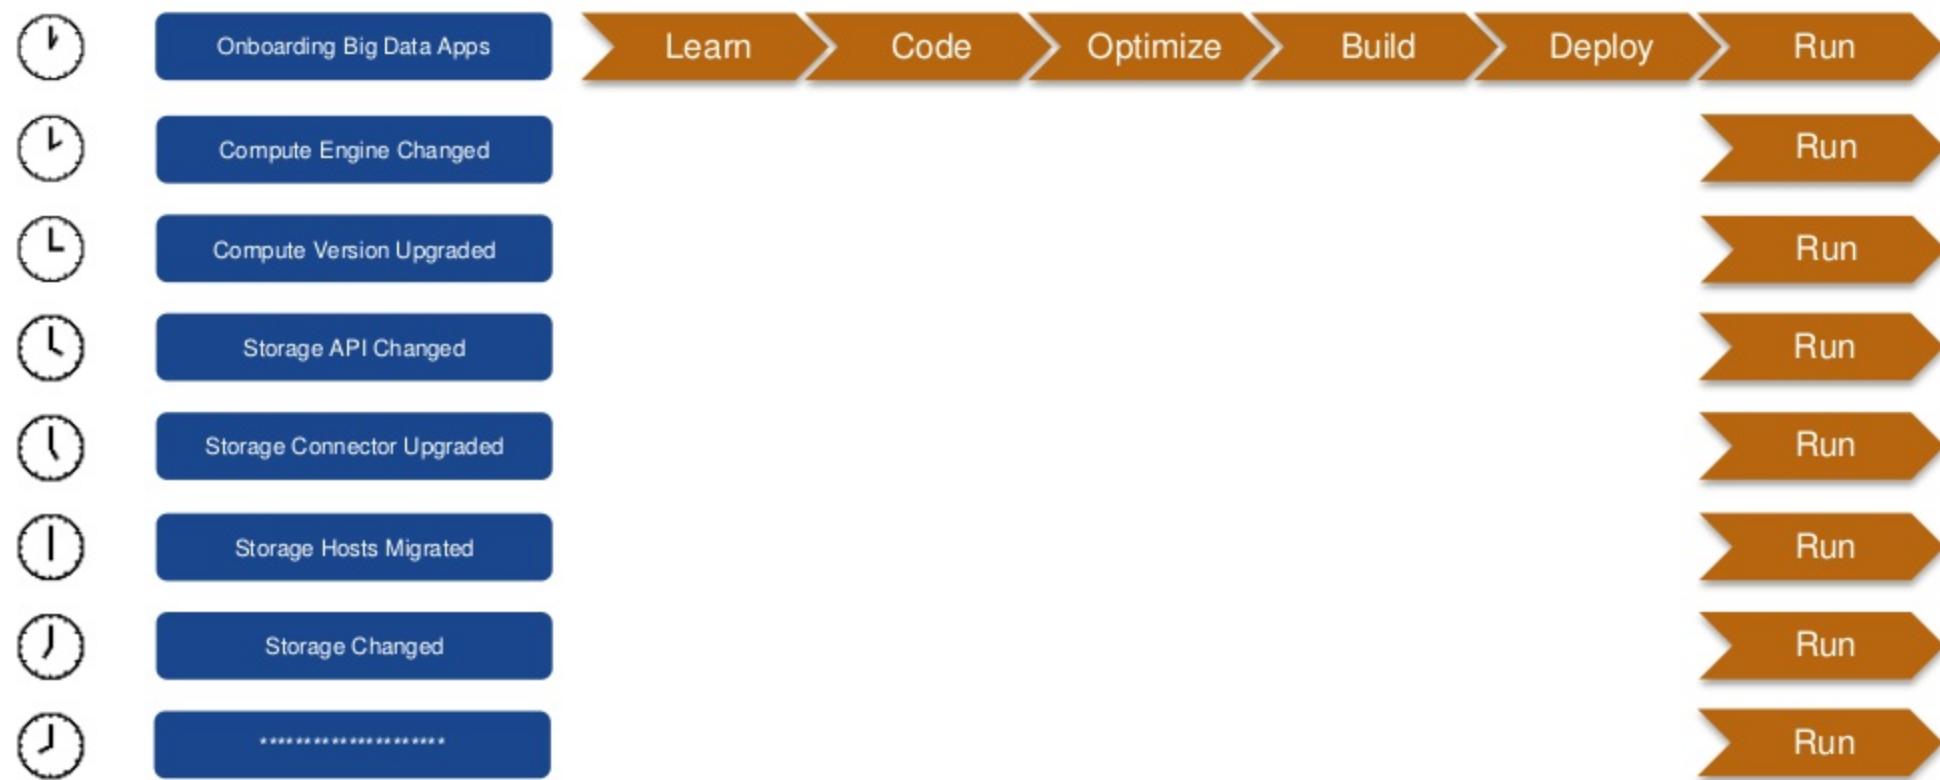
```
%%gimet-batch
insert into pcatalog.hdfs_dataset
partition(yyyy,mm,dd,hh,mi)
select kafka_ds.*,gimet_load_id
,substr(commit_timestamp,1,4) as yyyy
,substr(commit_timestamp,6,2) as mm
,substr(commit_timestamp,9,2) as dd
,substr(commit_timestamp,12,2) as hh
,case when cast(substr(commit_timestamp,15,2) as INT) <= 30 then "00" else
"30" end as mi
from pcatalog.kafka_dataset kafka_ds;
```

### Streaming SQL



```
%%gimet-stream
insert into pcatalog.hdfs_dataset
partition(yyyy,mm,dd,hh,mi)
select kafka_ds.*,gimet_load_id
,substr(commit_timestamp,1,4) as yyyy
,substr(commit_timestamp,6,2) as mm
,substr(commit_timestamp,9,2) as dd
,substr(commit_timestamp,12,2) as hh
,case when cast(substr(commit_timestamp,15,2) as INT) <= 30 then "00" else
"30" end as mi
from pcatalog.kafka_dataset kafka_ds;
```

# Data Platform | Data Application Lifecycle | With Data API



## Data Platform | Data API | Unlocking Additional Benefits

- **Explorer**

- View sample data
- Add capacity reservations
- Approval and Tracking

- **Discovery**

- Auto-discovery of datasets and importing into PCatalog

- **Dashboard**

- Operation dashboard: Metrics, statistics, trends, refresh status
- Admin dashboard: Capacity/growth/space statistics and action

- **Alert**

- User alerts to notify user datasets refresh delay and quality issue
- Admin alerts to notify datasets capacity and access violation

- **Query/BI Integration**

- Integrate catalog with Jupyter, Zeppelin and query any dataset

- **Out-of-the-box Benefits**

- Connector dependencies maintained by BDPE
- Change component version at one place. E.g. Hbase version
- Easy cluster upgrade
- Abstract storage from user. Just Prod, QA and Dev
- Add new compute engine to stack easily
- Add new storage system to stack easily
- Time to Develop
- Time to Market

*Presenting ....*

PayPal Gimmel


Data Platform  
+  
Compute Platform

NEXTGEN ANALYTICS DATA PLATFORM

Unified User Experience for any Compute Engine



Unified Access API for any Data Storage

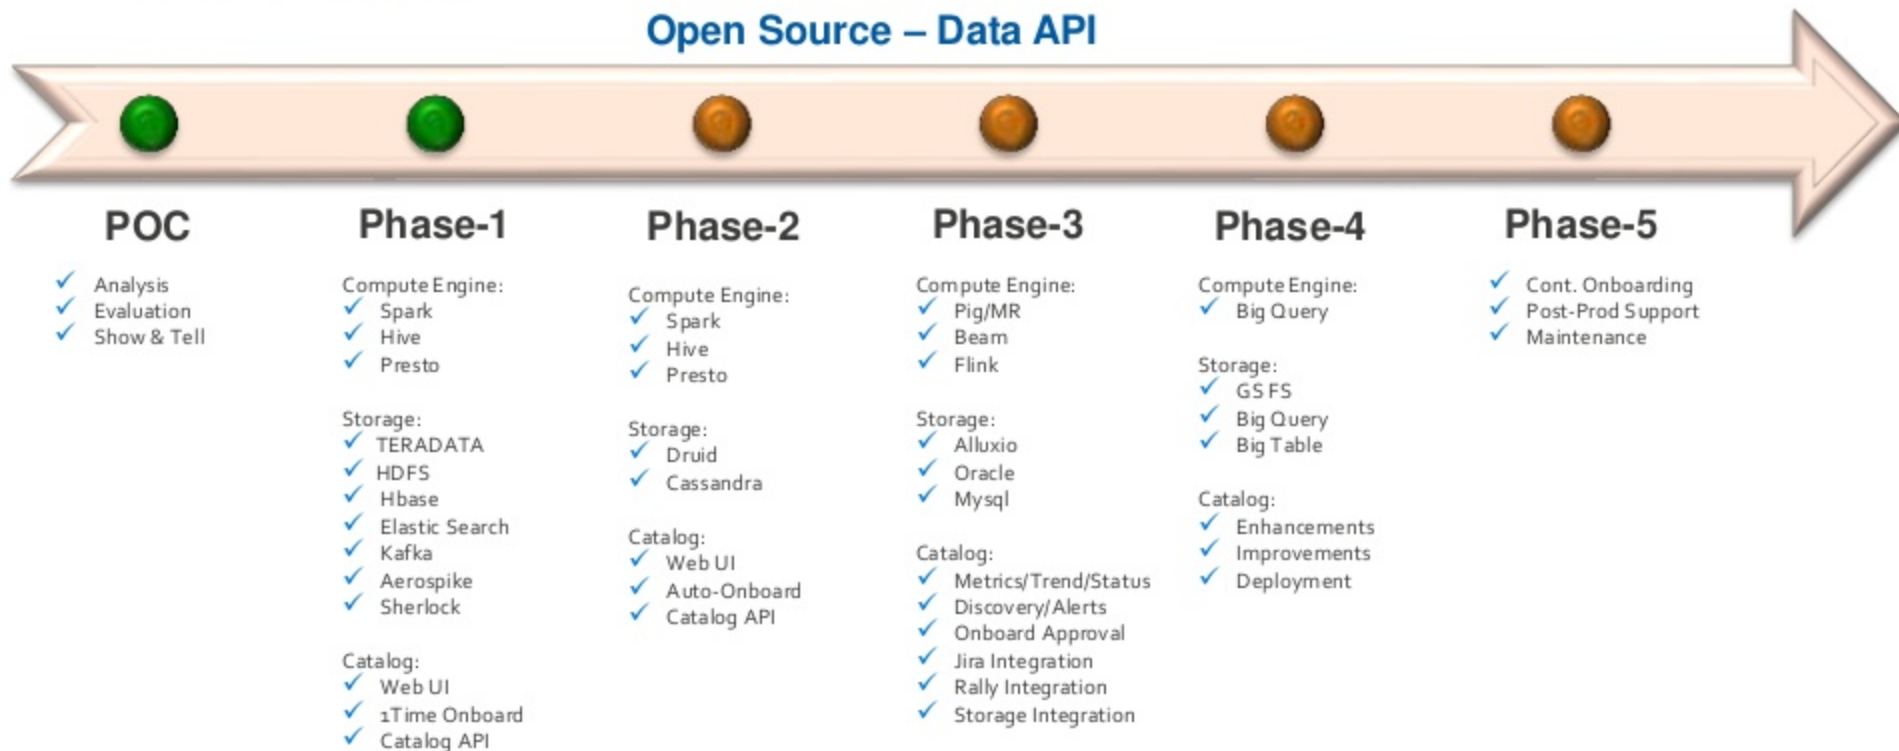


# Data Platform: Plan

2017 =====>

2018 =====>

Open Source – Data API







# Thank You!

**Rate This Session # 0348**

with the PARTNERS Mobile App

## **Questions/Comments**

Email: [dmohanakumarchan@paypal.com](mailto:dmohanakumarchan@paypal.com)

LinkedIn: <https://www.linkedin.com/in/deepakmc/>



Google Search Links Used for Logos of open source technologies :

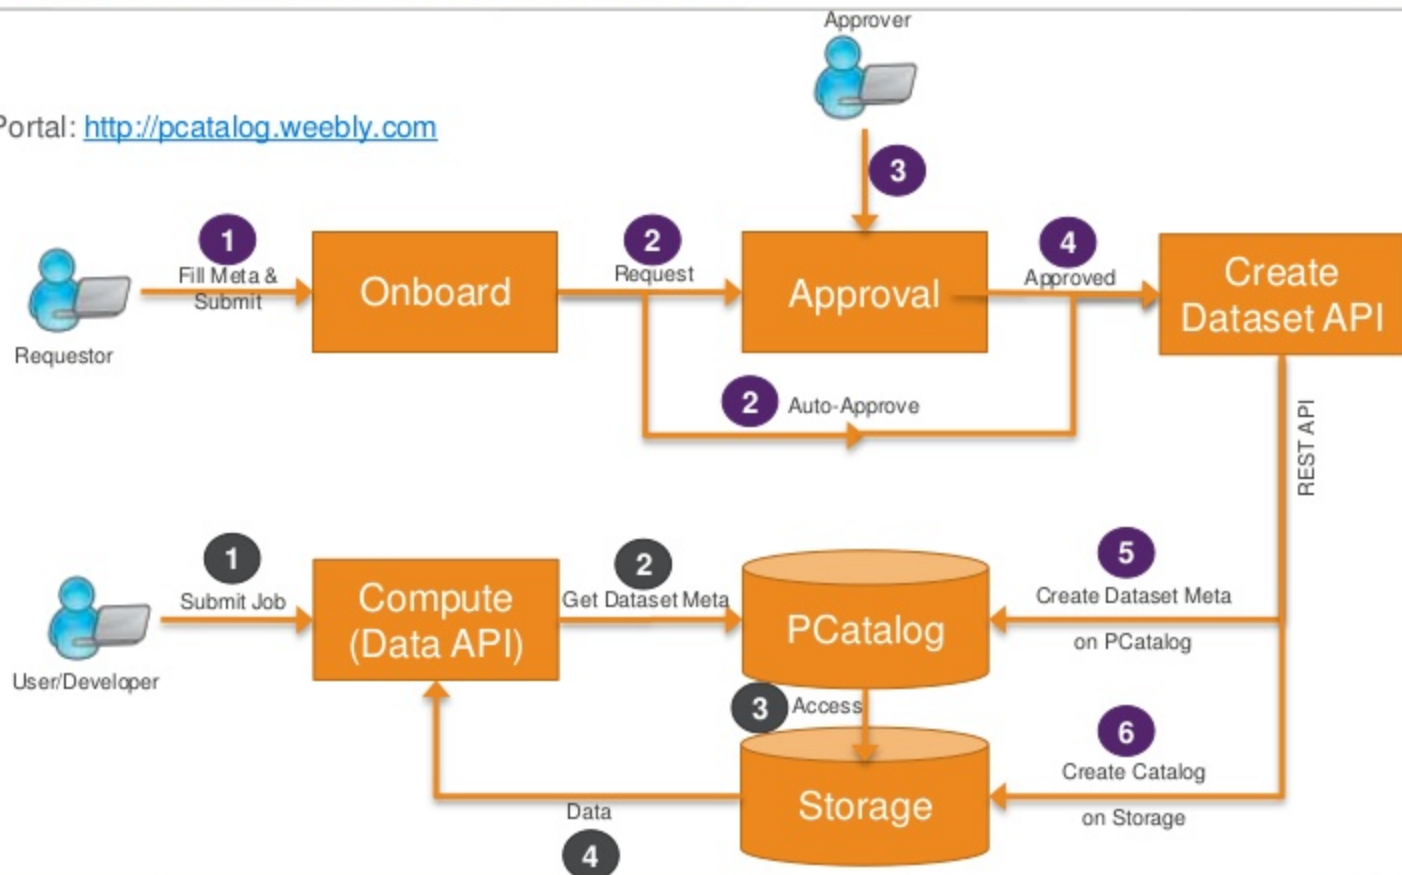
<https://www.google.com/search?q=big+data+stack+icons&source=lnms&tbm=isch&sa=X&ved=0ahUKEwiwnevLztzWAhVHsFQKHZp7CukQAUICigB&biw=1369&bih=739>

PayPal Q3 2017 & 2016 Full Year Results :

<https://www.paypal.com/us/webapps/mpp/about>

# Data Platform: Process Flow

- Sample Portal: <http://pcatalog.weebly.com>



- Data Access – Other challenges





# Data Platform: Plan Proposal

