

AWS

S U M M I T

# Deep Dive on AWS IoT

Olawale Oladehin, Sr. Solutions Architect, IoT

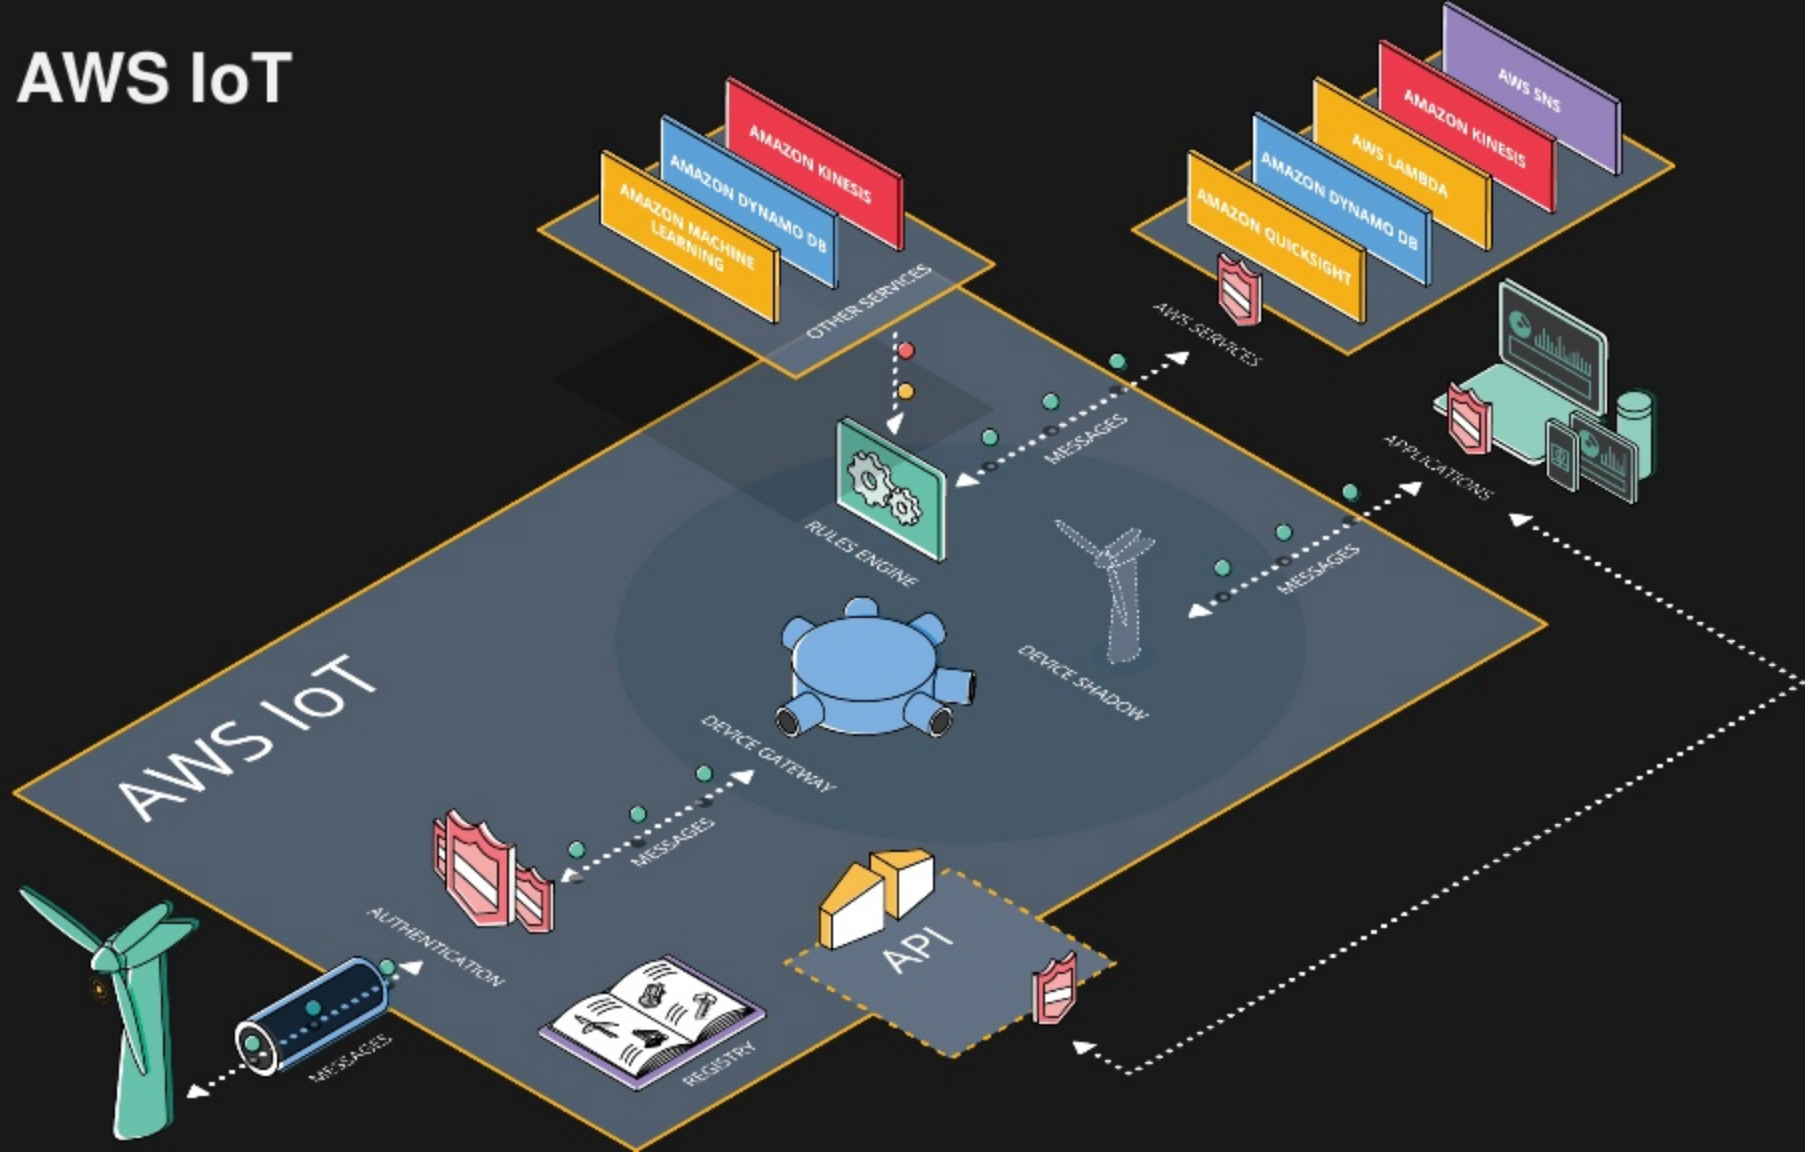
August 14, 2017



# Agenda

- Overview of AWS IoT
- Telemetry & analytics
- Cloud control
- Mobile control
- Lifecycle management
- Wrap-up

# AWS IoT

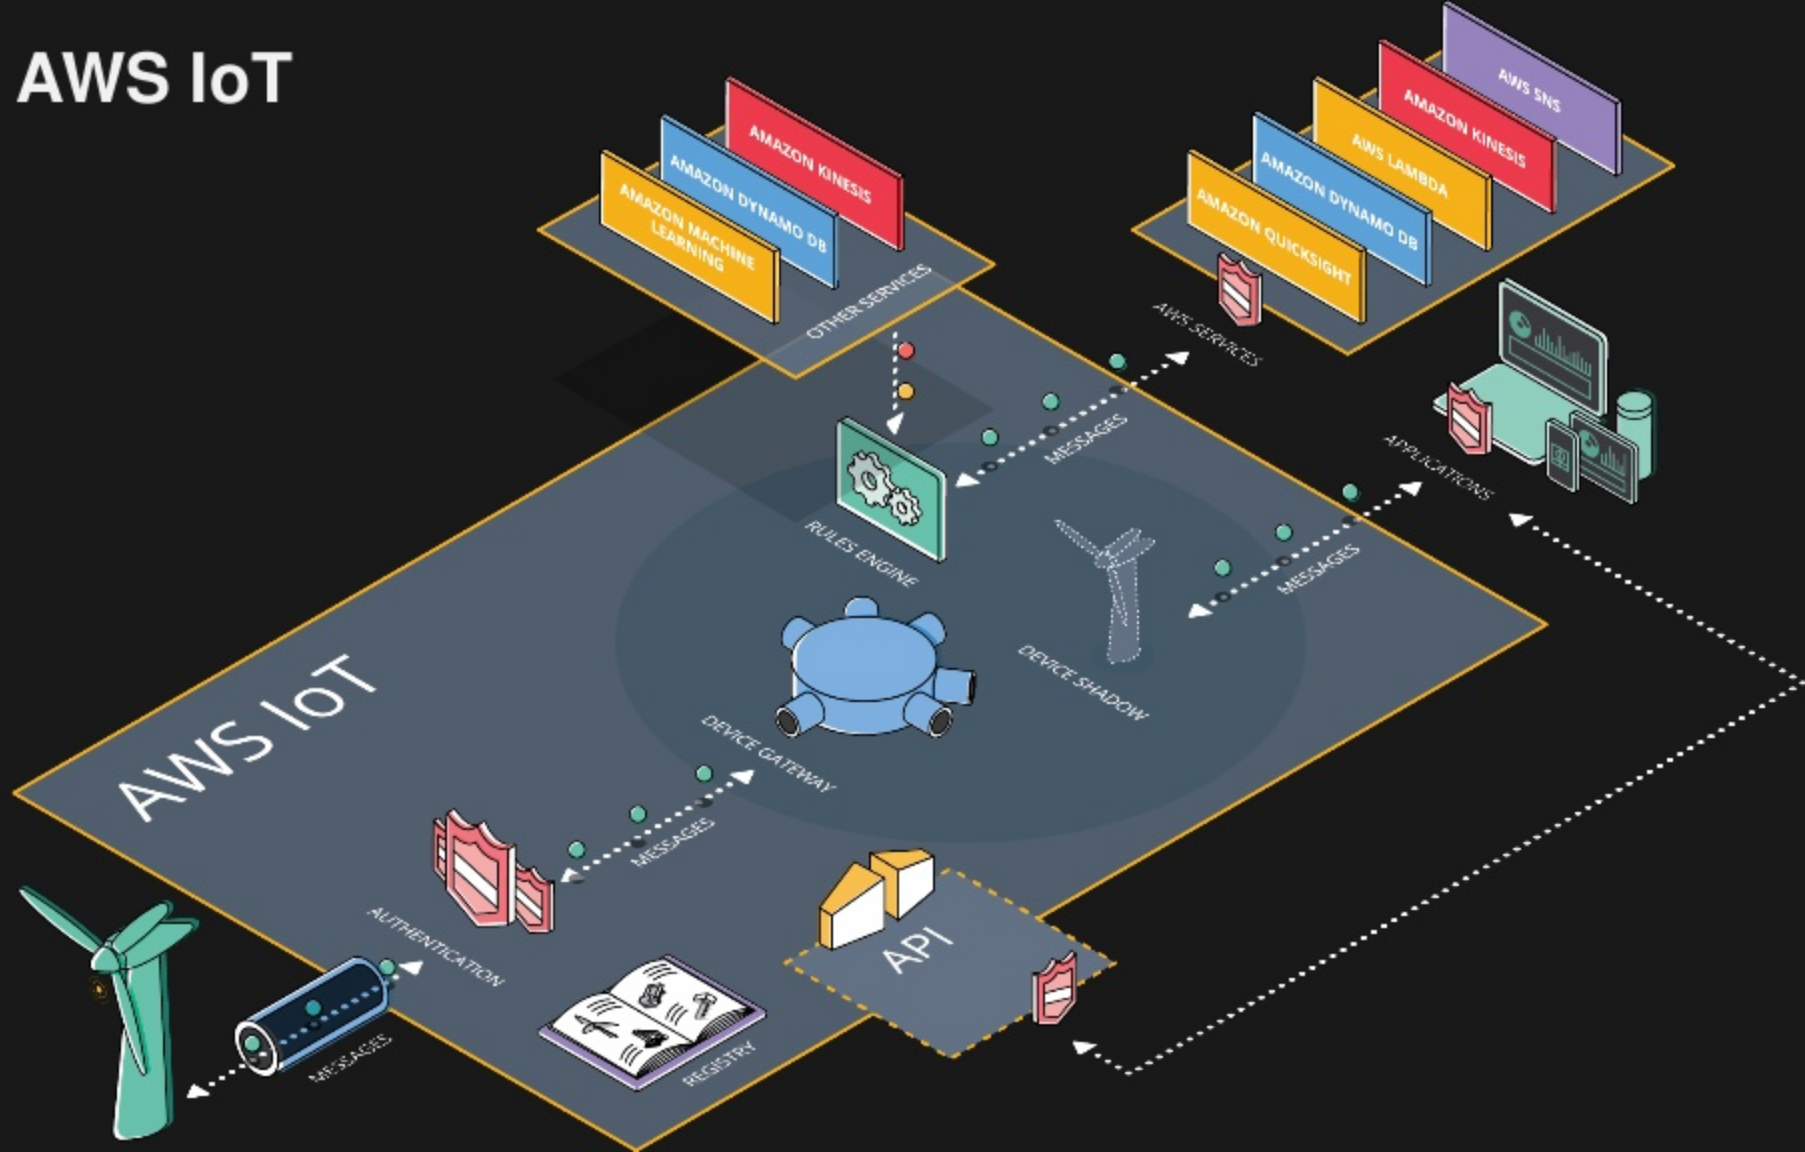


# Telemetry & Analytics




# AWS IoT Telemetry & Analytics

1. Connect devices
2. Send data
3. Collect and store the data
4. Do something with the data

# AWS IoT



# 1) Connect the devices

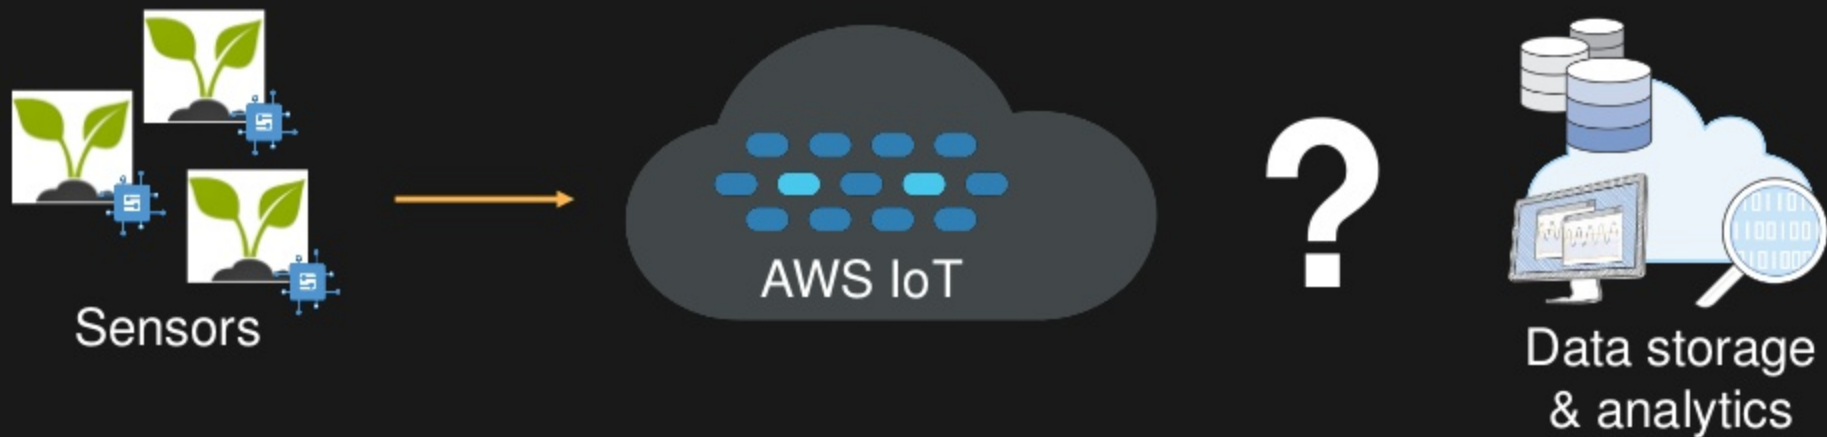
1. Provision a certificate 
2. Attach policy 
3. Connect over MQTT 

## 2) Send data

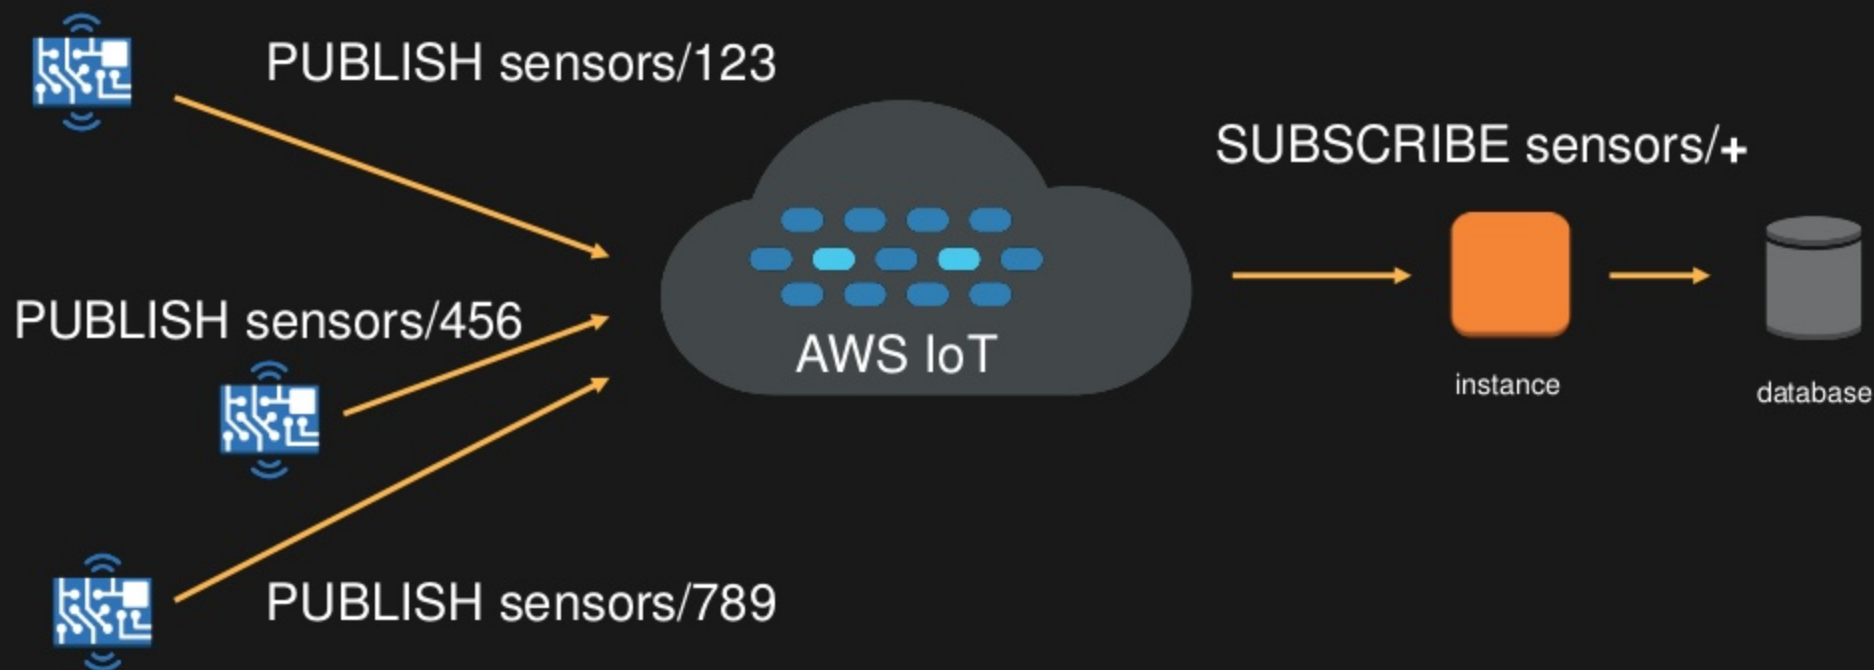
```
PUBLISH macdonald/sensors/123 (qos: 0)
{
  "timestamp": "2016-01-29T10:00:00",
  "temperature": 55
  "humidity": 39,
  "ph": 6.7
}
```



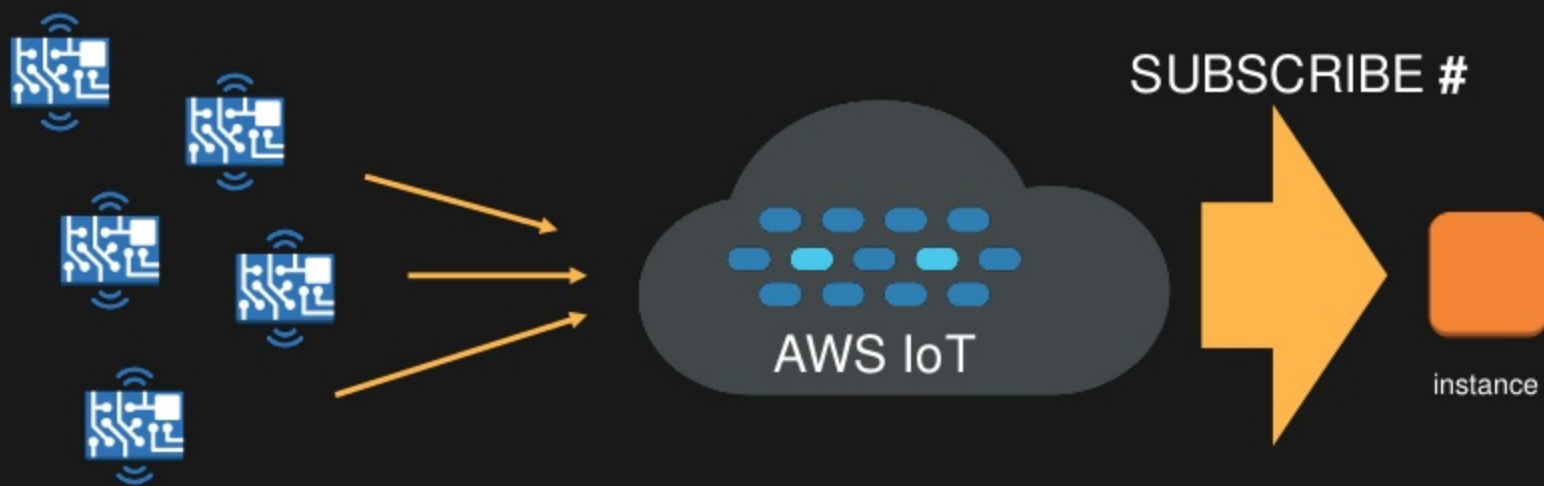
### 3) Collect the data



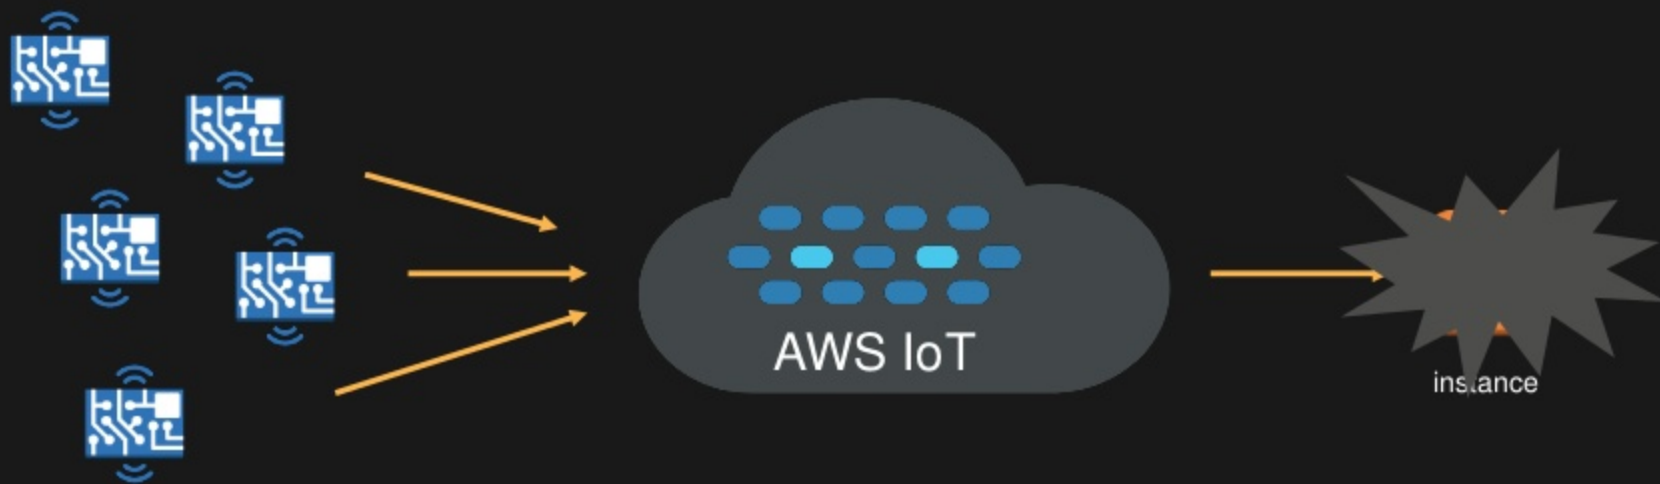
# Single consumer (don't do this)



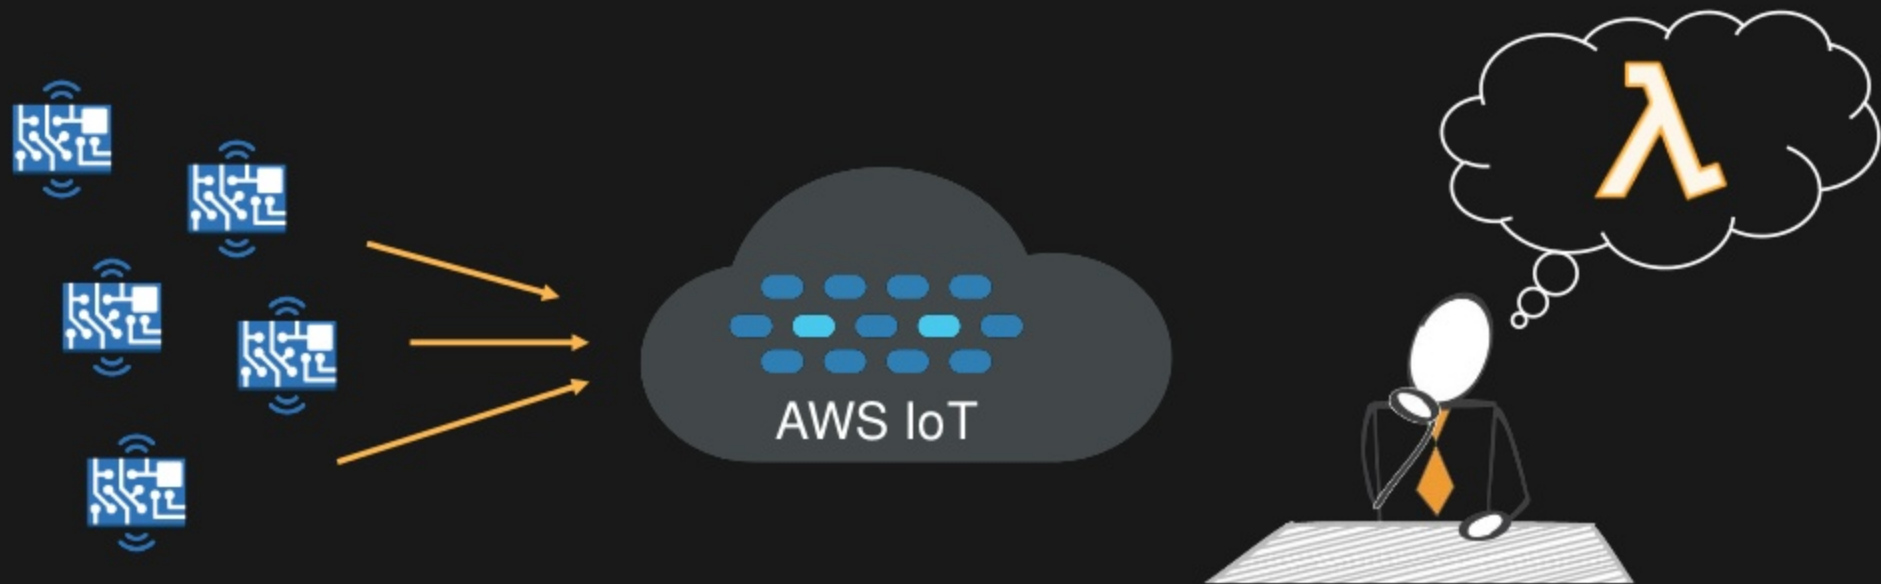
# Don't do this: scalability



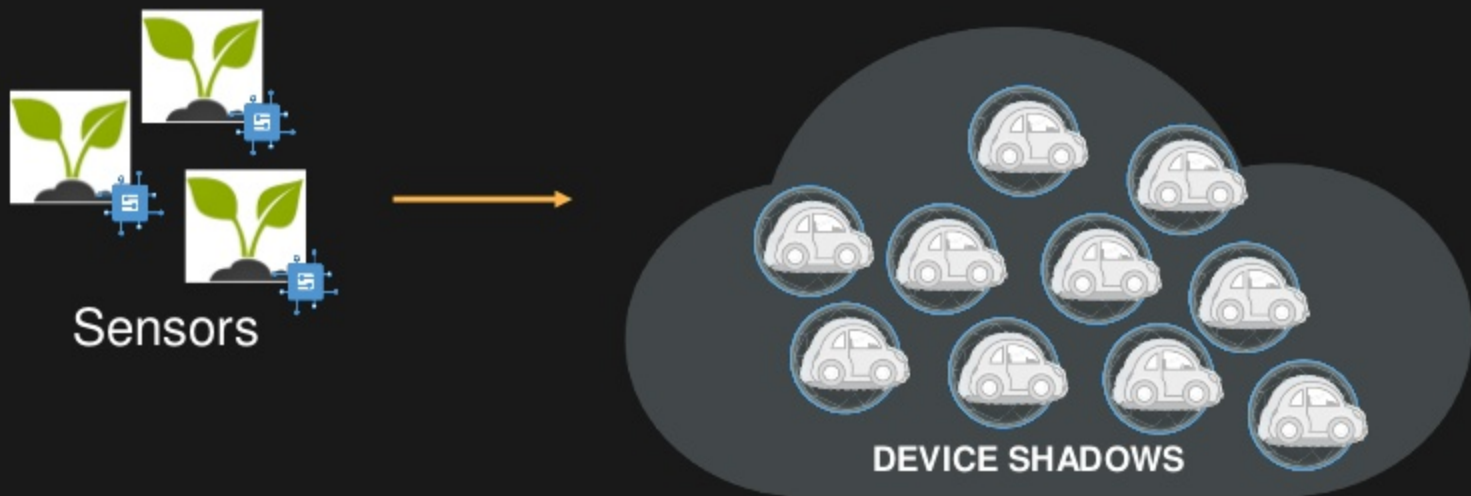
# Don't do this: availability



# Don't do this: maintainability



# Store it in the device shadow (don't do this)



# AWS IoT Rules Engine

Rules Engine connects AWS IoT to external endpoints and AWS services.



# Example rule

macdonald/sprinkler-456/temperature

```
{
  "rule": {
    "sql": "SELECT * AS message FROM 'macdonald/#'",
    "description": "Store all sensor data into dynamodb and firehose",
    "actions": [{
      "dynamoDB": {
        "tableName": "sensor_data",
        "roleArn": "arn:aws:iam::123456789012:role/aws_iot_dynamoDB",
        "hashKeyField": "sensor_id",
        "hashKeyValue": "${topic(2)}",
        "rangeKeyField": "timestamp",
        "rangeKeyValue": "${timestamp()}",
      }, {
        "firehose": {
          "roleArn": "arn:aws:iam::123456789012:role/aws_iot_firehose",
          "deliveryStreamName": "my_firehose_stream"
        }
      }
    ]
  }
}
```



# Example rule

macdonald/sprinkler-456/temperature

```
{
  "rule": {
    "sql": "SELECT * AS message FROM 'macdonald/#'",
    "description": "Store all sensor data into dynamodb and firehose",
    "actions": [{
      "dynamoDB": {
        "tableName": "sensor_data",
        "roleArn": "arn:aws:iam::123456789012:role/aws_iot_dynamoDB",
        "hashKeyField": "sensor_id",
        "hashKeyValue": "${topic(2)}",
        "rangeKeyField": "timestamp",
        "rangeKeyValue": "${timestamp()}",
      },
      {
        "firehose": {
          "roleArn": "arn:aws:iam::123456789012:role/aws_iot_firehose",
          "deliveryStreamName": "my_firehose_stream"
        }
      }
    ]
  }
}
```

# Example rule

macdonald/sprinkler-456/temperature

```
{
  "rule": {
    "sql": "SELECT * AS message FROM 'macdonald/#'",
    "description": "Store all sensor data into dynamodb and firehose",
    "actions": [{
      "dynamoDB": {
        "tableName": "sensor_data",
        "roleArn": "arn:aws:iam::123456789012:role/aws_iam_dynamoDB",
        "hashKeyField": "sensor_id",
        "hashKeyValue": "${topic(2)}",
        "rangeKeyField": "timestamp",
        "rangeKeyValue": "${timestamp()}",
      }, {
        "firehose": {
          "roleArn": "arn:aws:iam::123456789012:role/aws_iam_firehose",
          "deliveryStreamName": "my_firehose_stream"
        }
      }
    ]
  }
}
```

# Example rule

macdonald/sprinkler-456/temperature

```
{
  "rule": {
    "sql": "SELECT * AS message FROM 'macdonald/#'",
    "description": "Store all sensor data into dynamodb and firehose",
    "actions": [{
      "dynamoDB": {
        "tableName": "sensor_data",
        "roleArn": "arn:aws:iam::123456789012:role/aws_iot_dynamoDB",
        "hashKeyField": "sensor_id",
        "hashKeyValue": "${topic(2)}",
        "rangeKeyField": "timestamp",
        "rangeKeyValue": "${timestamp()}",
      },
      {
        "firehose": {
          "roleArn": "arn:aws:iam::123456789012:role/aws_iot_firehose",
          "deliveryStreamName": "my_firehose_stream"
        }
      }
    ]
  }
}
```

# IoT SQL functions

- Math
  - abs, acos, asin, atan, atan2, cos, cosh, etc...
  - round, ceil, floor, exp,
- Bitwise Ops
  - bitand, bitor, bitxor, and bitnot
- String and Arrays
  - chr, concat, encode, endswith, get, indexof, length, lower, upper, etc...
  - newuuid, parse\_time
- Hashing
  - md2, md5, sha1, sha224, sha256, sha384, sha512
- IoT
  - machinelearning\_predict(modelId, roleARN, record)
  - get\_thing\_shadow(thingName, roleARN)
  - principal, topic, timestamp, traceid, clientid

```
[...]  
"hashKeyValue": "${topic(2)}",  
"rangeKeyField": "timestamp"  
"rangeKeyValue": "${timestamp()}"  
[...]
```

# Different data scenarios

Want to run a lot of queries constantly?

- ➡ Use Amazon Kinesis Firehose to write into Amazon Redshift

Need fast lookups, e.g., in Rules or Lambda functions?

- ➡ Write into DynamoDB, add indexes if necessary

Have a need for heavy queries but not always-on?

- ➡ Use Firehose and S3, process with Amazon EMR or Athena.

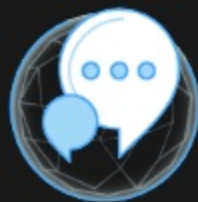
# Cloud Control

# Publish on/off to the sprinkler

SUBSCRIBE  
macdonald/sprinkler-456



Sprinkler



Device  
Gateway



Control  
logic

# Publish on/off to the sprinkler

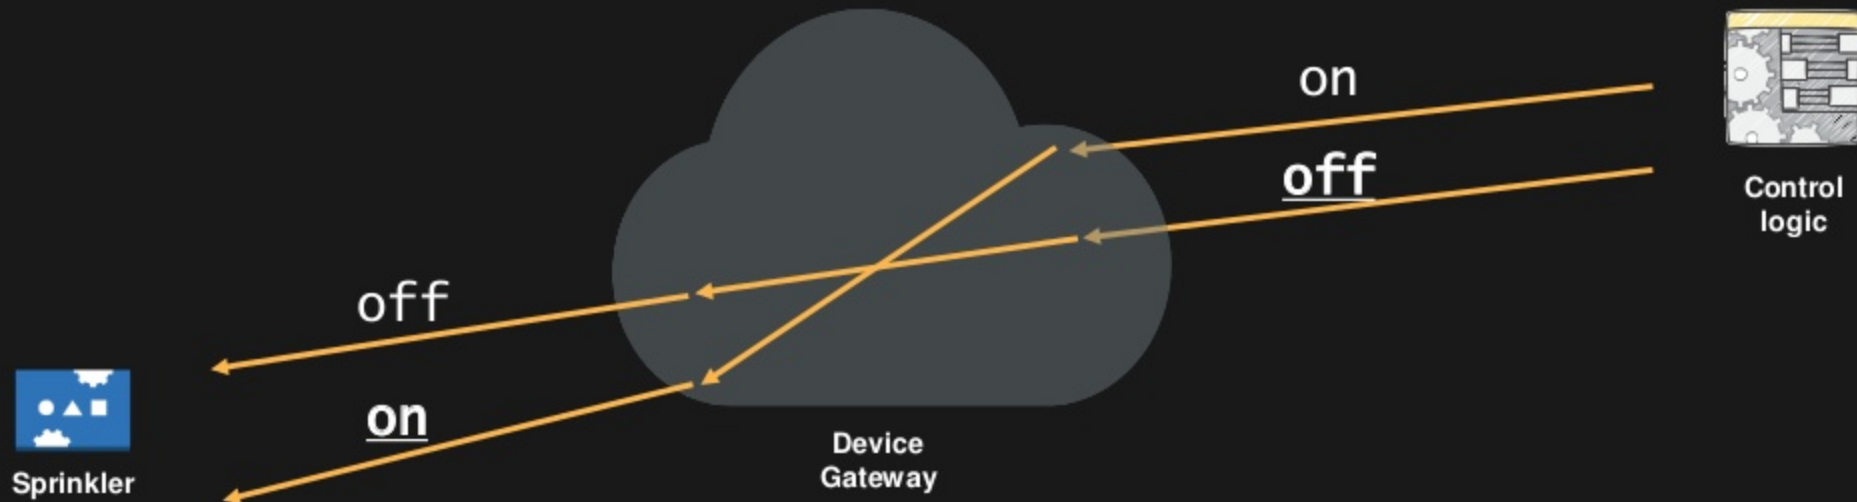




# Direct publishing: why not?



# Direct publishing: why not?



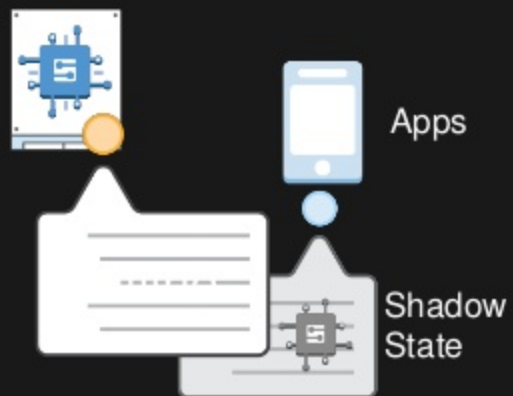
# Direct publishing: why not?

➡ Connection blips

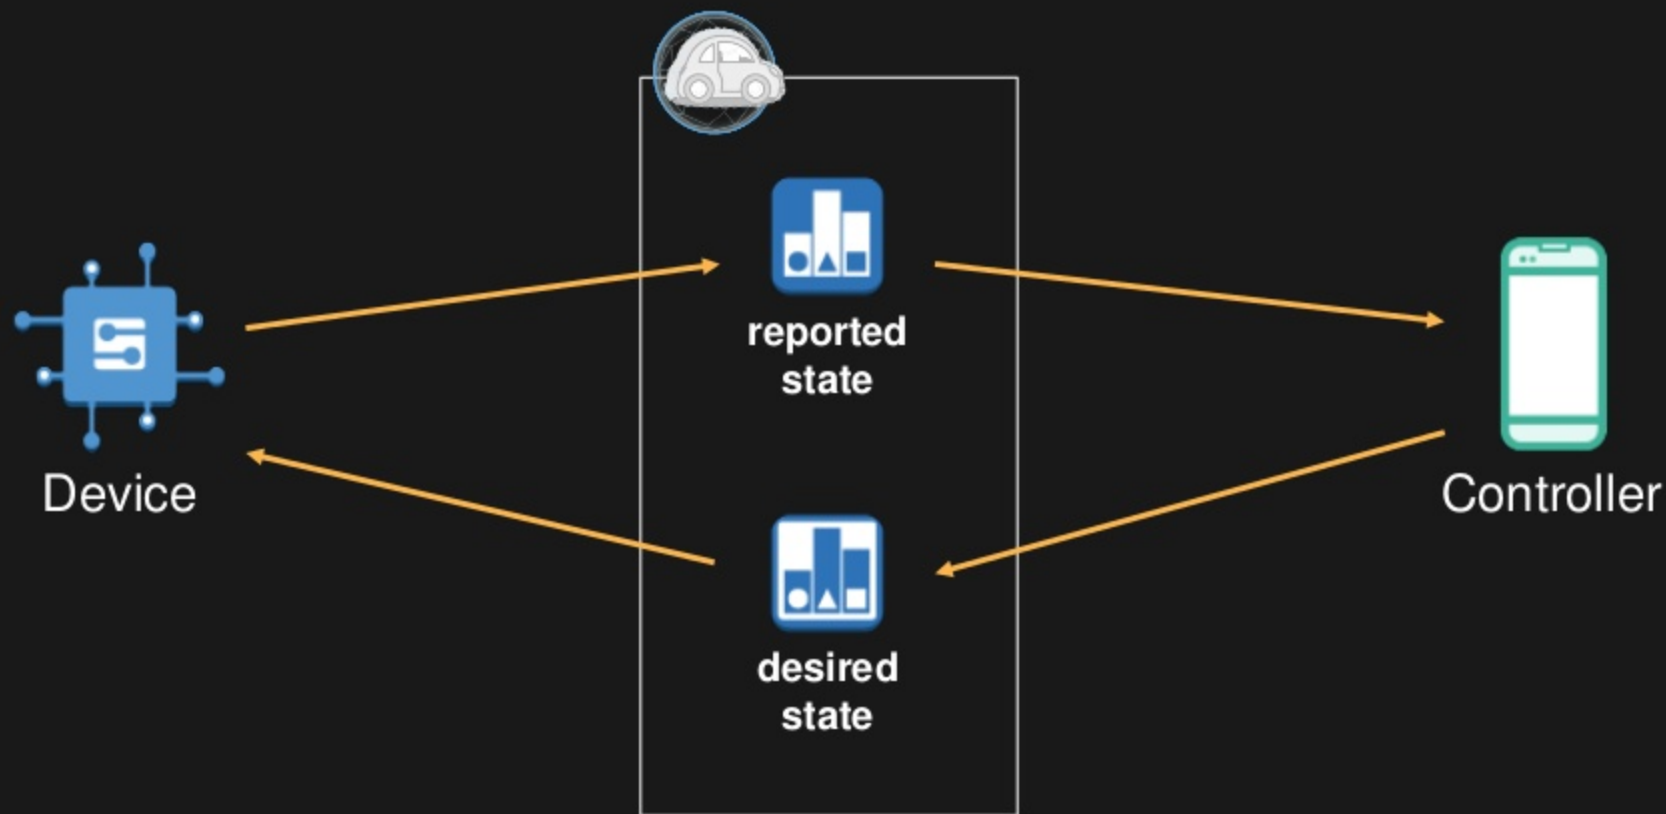
➡ Messages aren't ordered

So then what?

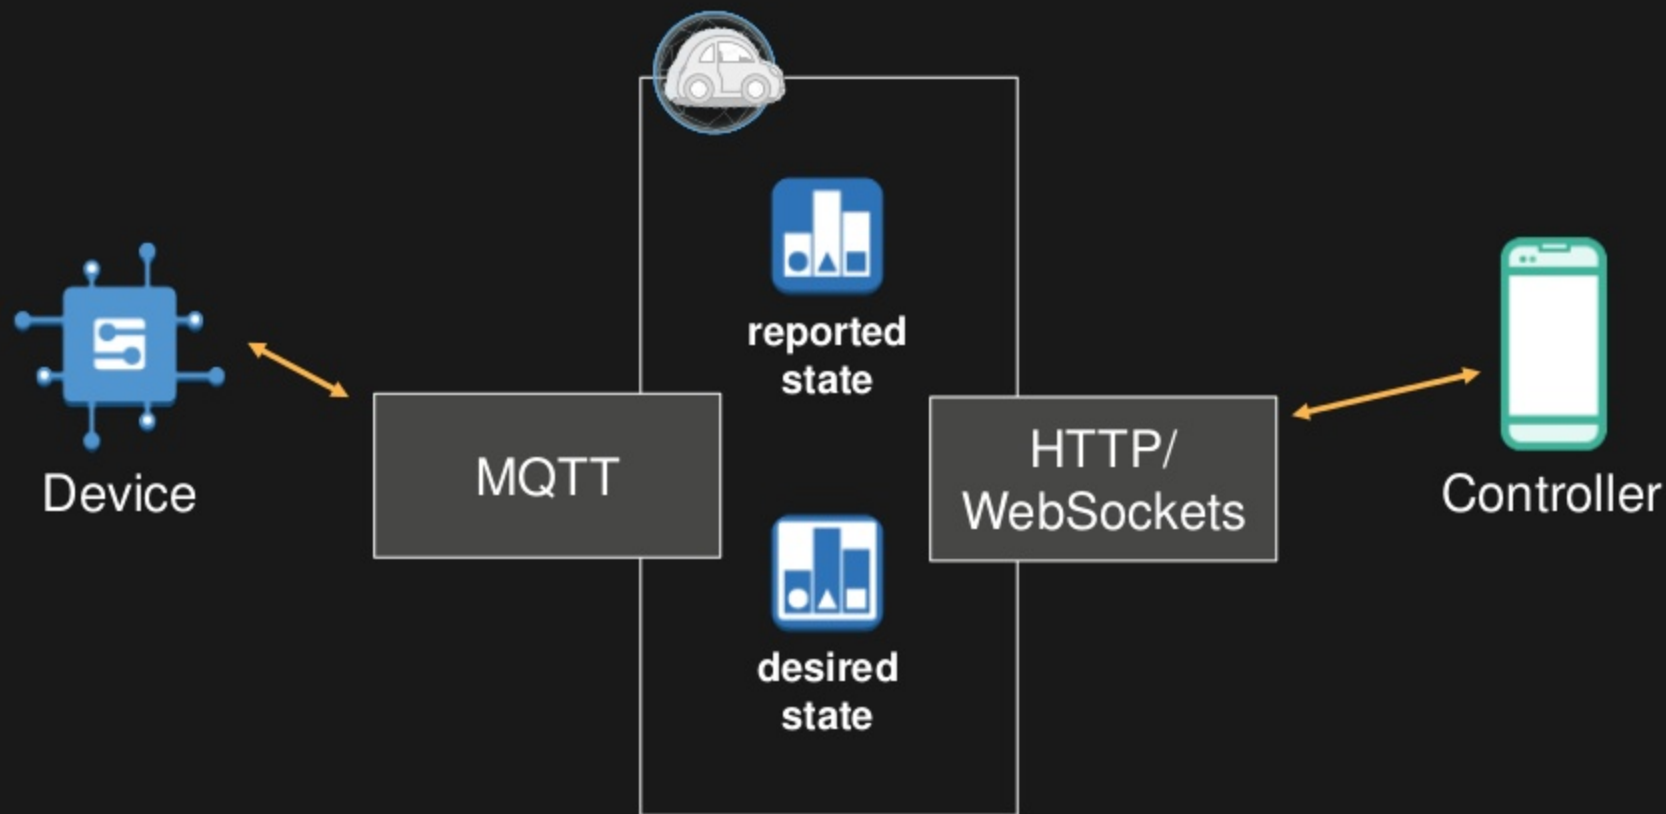
# Device Shadows



# Device Shadows



# Device Shadows



# AWS IoT Shadow - simple yet powerful



## Thing

Report its current state to one or multiple shadows  
Retrieve its desired state from shadow



## Shadow

Shadow reports delta, desired and reported states along with metadata and version

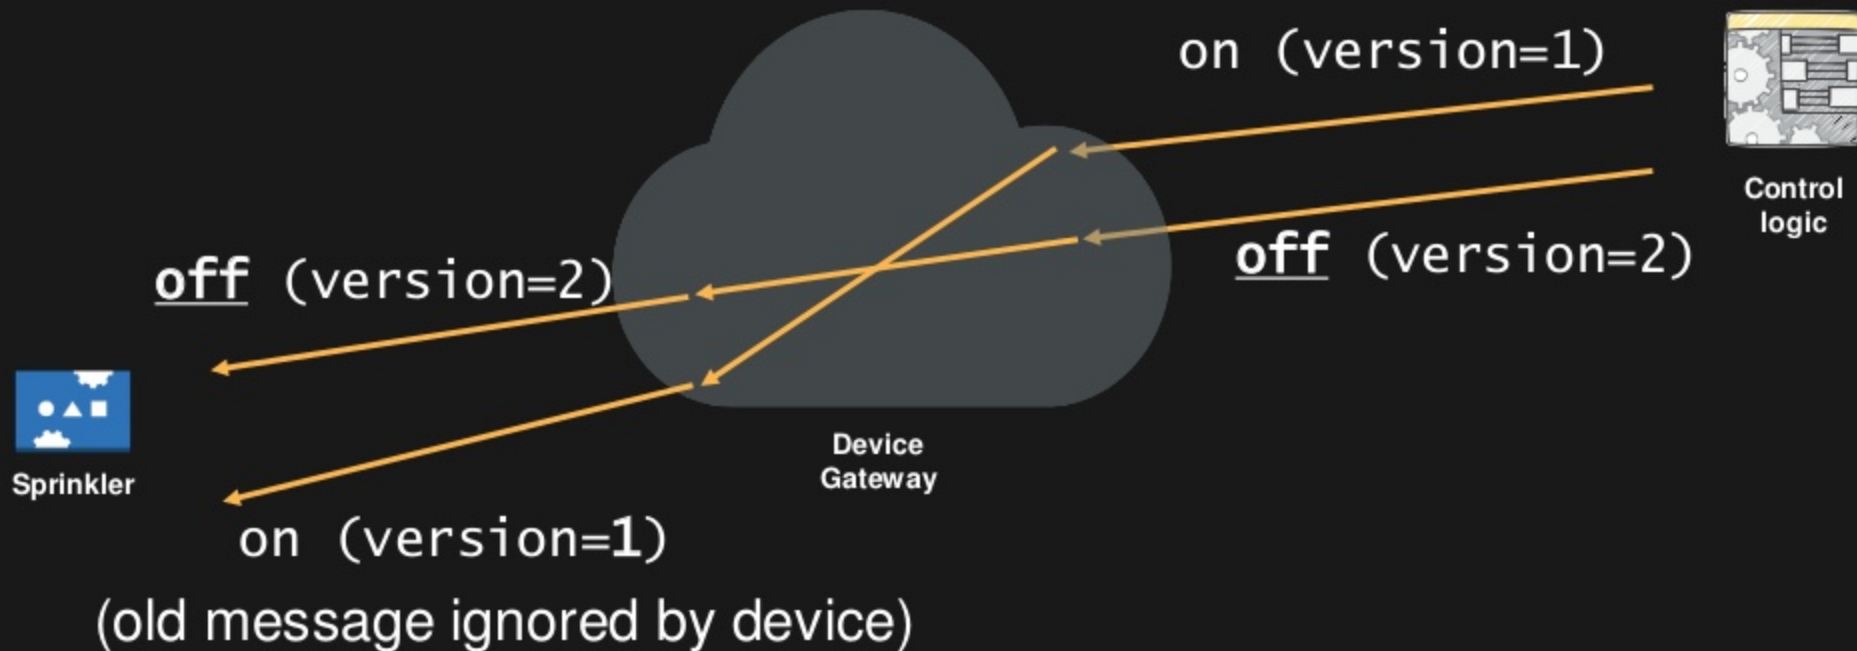


## Mobile App

Set the desired state of a device  
Get the last reported state of the device  
Delete the shadow

```
{  
  "state" : {  
    "desired" : {  
      "lights": { "color": "RED" },  
      "engine" : "ON"  
    },  
    "reported" : {  
      "lights" : { "color": "GREEN" },  
      "engine" : "ON"  
    },  
    "delta" : {  
      "lights" : { "color": "RED" }  
    }  
  },  
  "version" : 10  
}
```

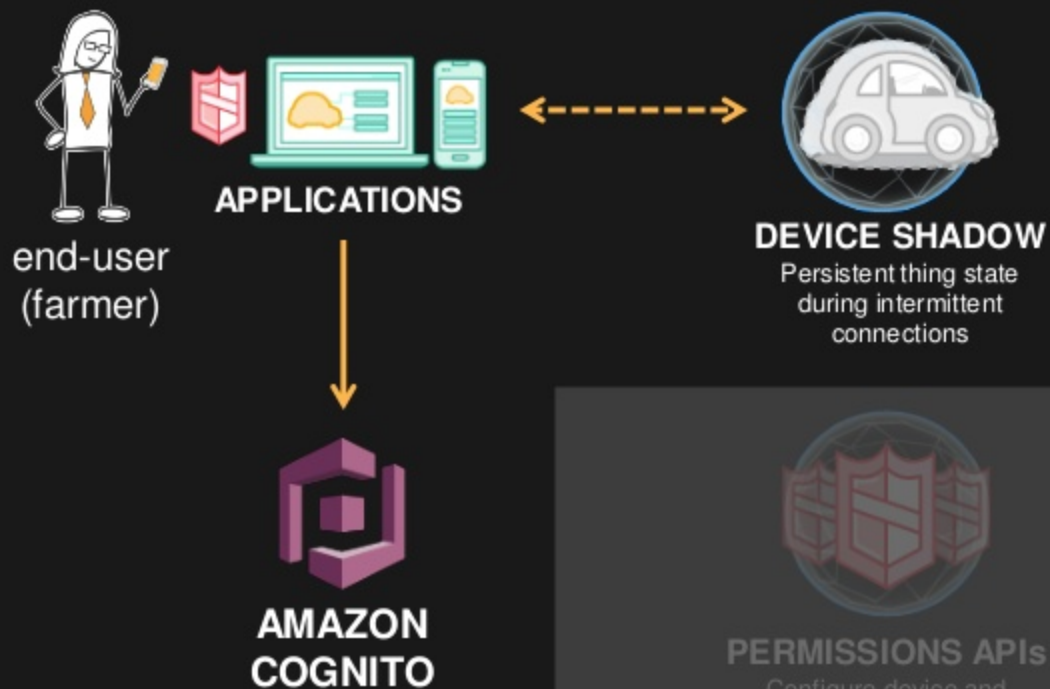
# Device Shadows and versioning





# Mobile Control

# Using Amazon Cognito with IoT



# Policy for Cognito with IoT

Cognito Identity = us-east-1:xxxx-yyyy-zzzz

```
aws iot attach-principal-policy --policy-name farm-sensors  
--principal us-east-1:xxxx-yyyy-zzzz
```

You will need a trusted entity to attach the Cognito principal to an IoT policy

- Only needed for iot-data plane calls such as DeleteThingShadow, UpdateThingShadow, GetThingShadow, Connect, Publish, and Subscribe
- Can use API Gateway, Cognito Sync Triggers, or other techniques for attaching the Cognito principal ID to the IoT policy

## Overall Cognito “pairing” workflow

1. Create a Cognito identity pool
2. Customer signs in using mobile app
3. Associate their user with their “farm”
4. Create a scope-down policy in IoT for their user
5. Attach that policy to their Cognito user in IoT

# Managing fine-grained permissions

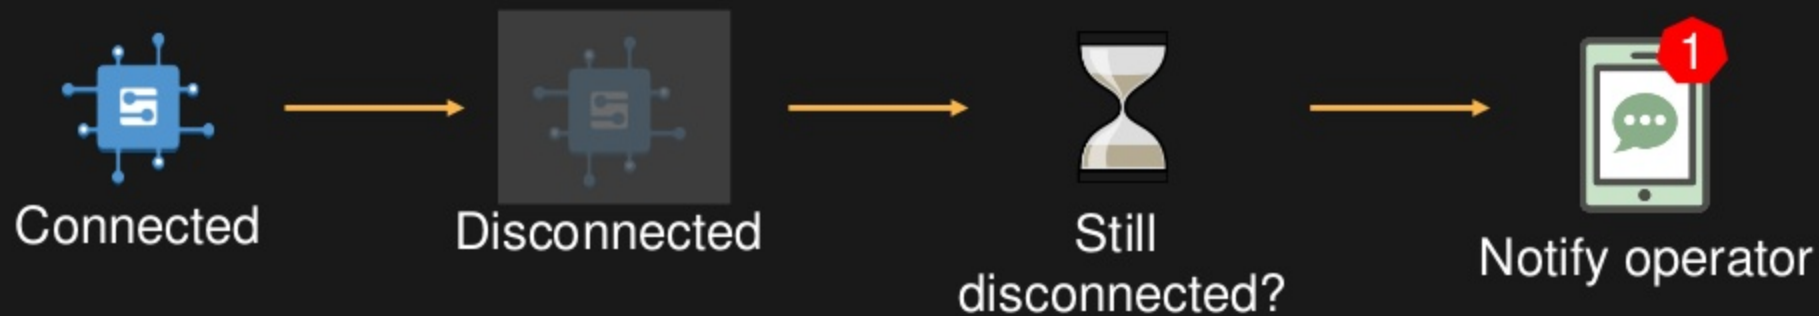
- One “farm owner” needs permissions to many shadows
  - "arn:aws:iot:...:thing/sprinkler123abc"
  - "arn:aws:iot:...:thing/sprinkler456def"
  - ...
- Listing each is tedious

# Best practice: thing name prefixing

- ➡ Prefix thing name with logical owner
  - sensor123abc -> macdona1d-sensor123abc
- ➡ IAM policies support wildcards
  - "arn:aws:iot:...:thing/sensor123abc"
  - "arn:aws:iot:...:thing/sensor123abc"
  - "arn:aws:iot:...:thing/sensor456def"
  - ...
  - "arn:aws:iot:...:thing/macdona1d-\*"

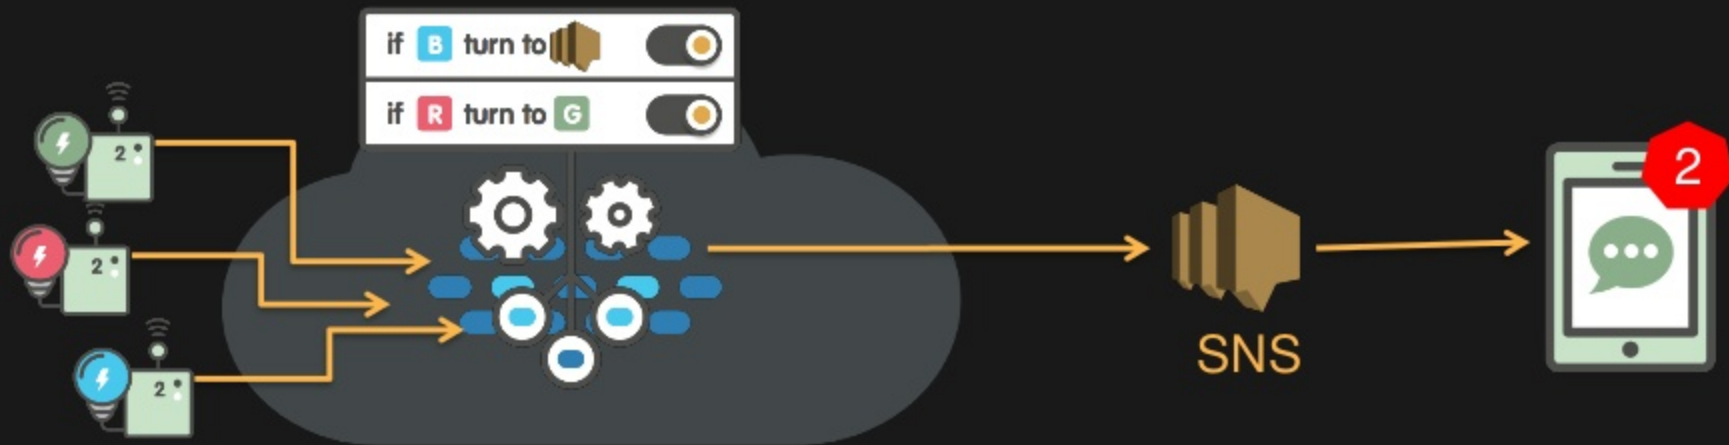
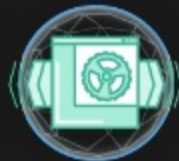
# Lifecycle Management

# Lifecycle workflow





# AWS IoT Rules Engine & Amazon SNS



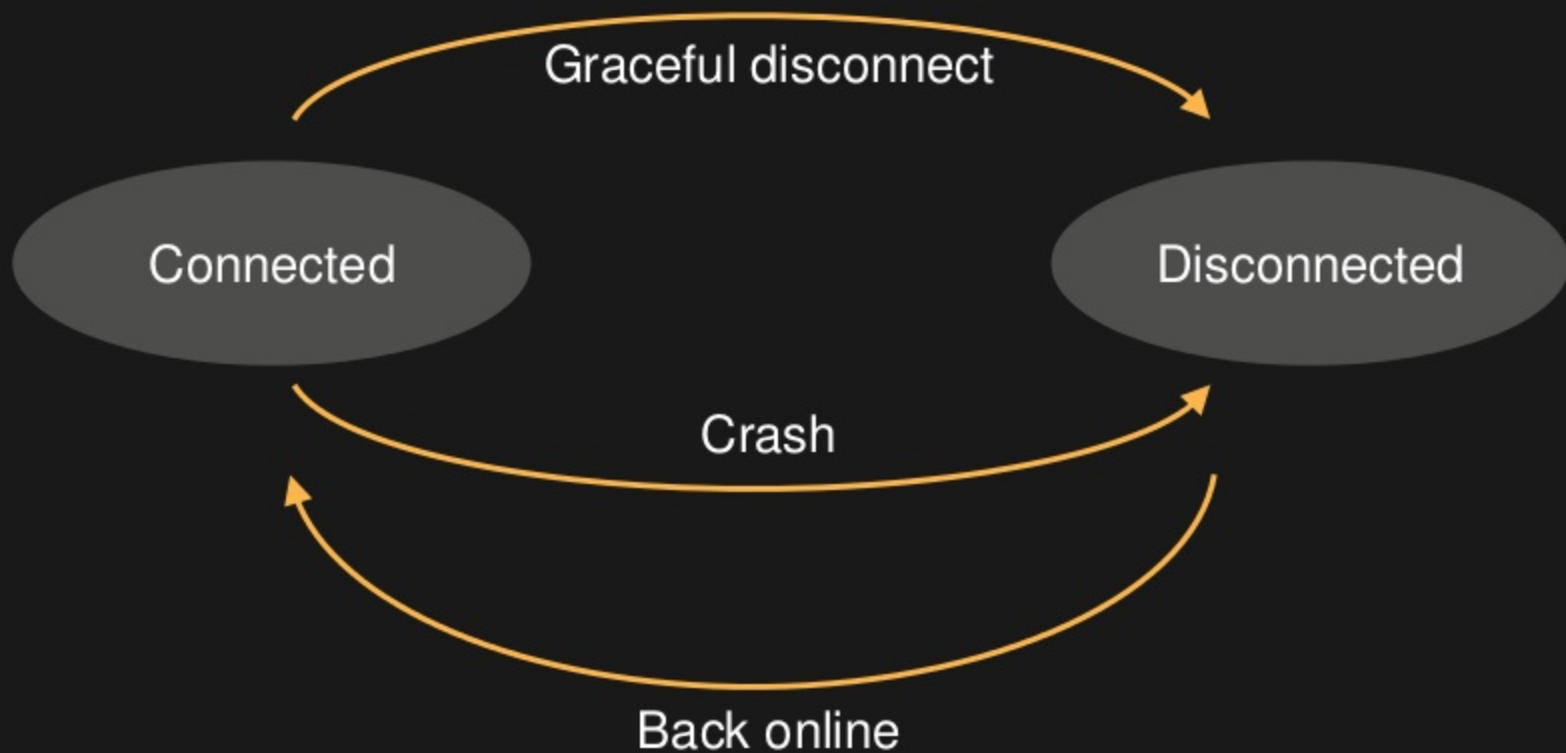
## Push Notifications

Apple APNS Endpoint, Google GCM Endpoint, Amazon ADM Endpoint, Windows WNS

## Amazon SNS -> HTTP Endpoint (or SMS or email)

Call HTTP-based third-party endpoints through SNS with subscription and retry support

# Detecting disconnects



# Lifecycle events

- Connect
  - PUBLISH lifecycle/sensor-123  
{"status": "online"}
- Disconnect (graceful)
  - PUBLISH lifecycle/sensor-123  
{"status": "offline"}
- Disconnect (crash)
  - PUBLISH lifecycle/sensor-123  
{"status": "offline", "isCrash": true}

# Handling lifecycle events

```
SELECT
    status,
    topic(2) as deviceId,
    timestamp() as time,
    isCrash
FROM lifecycle/#
WHERE status='offline'
```



## AWS Lambda function

- Look up mobile push ID for device owner
- Send SNS mobile push

# Delayed lifecycle events

SELECT

status,  
topic(2) as deviceId,  
timestamp() as time,  
isCrash

FROM lifecycle/#



## AWS Lambda function

- Store update device status in DynamoDB
- If offline: enqueue an SQS message with DelaySeconds



## SQS Message (15 minutes later)

- Double-check the status in DynamoDB
- Send SNS push notification if still offline



Amazon  
DynamoDB

Device	Status	Time
sensor-123	connected	11:30
...		

# Generating lifecycle events

- Connect
  - `PUBLISH lifecycle/sensor-123`  
`{"status": "online"}`
- Disconnect (graceful)
  - `PUBLISH lifecycle/sensor-123`  
`{"status": "offline"}`
- Disconnect (crash)
  - `PUBLISH lifecycle/sensor-123`  
`{"status": "offline", "isCrash": true}`

# Lifecycle events: connecting

## Automatic lifecycle message

`$aws/events/presence/connected/clientId`

```
{  
  "clientId": "a1b2c3d4e5f6a7b8c9d0e1f2a3b4c5d6",  
  "timestamp": 1460065214626,  
  "eventType": "connected",  
  "sessionId": "00000000-0000-0000-0000-000000000000",  
  "principalIdentifier": "000000000000/XXX:some-user/XXX:some-user"  
}
```

# Lifecycle events: disconnecting

## Automatic lifecycle message

`$aws/events/presence/disconnected/clientId`

```
{  
  "clientId": "a1b2c3d4e5f6a7b8c9d0e1f2a3b4c5d6",  
  "timestamp": 1460065214626,  
  "eventType": "disconnected",  
  "sessionId": "00000000-0000-0000-0000-000000000000",  
  "principalIdentifier": "000000000000/XXX:some-user/XXX:some-user"  
}
```



# Last Will and Testament (LWT)

CONNECT message parts:

Protocol: MQTT 3.1.1

ClientId: abc

KeepAlive: 60 seconds

LastWill PUBLISH message:

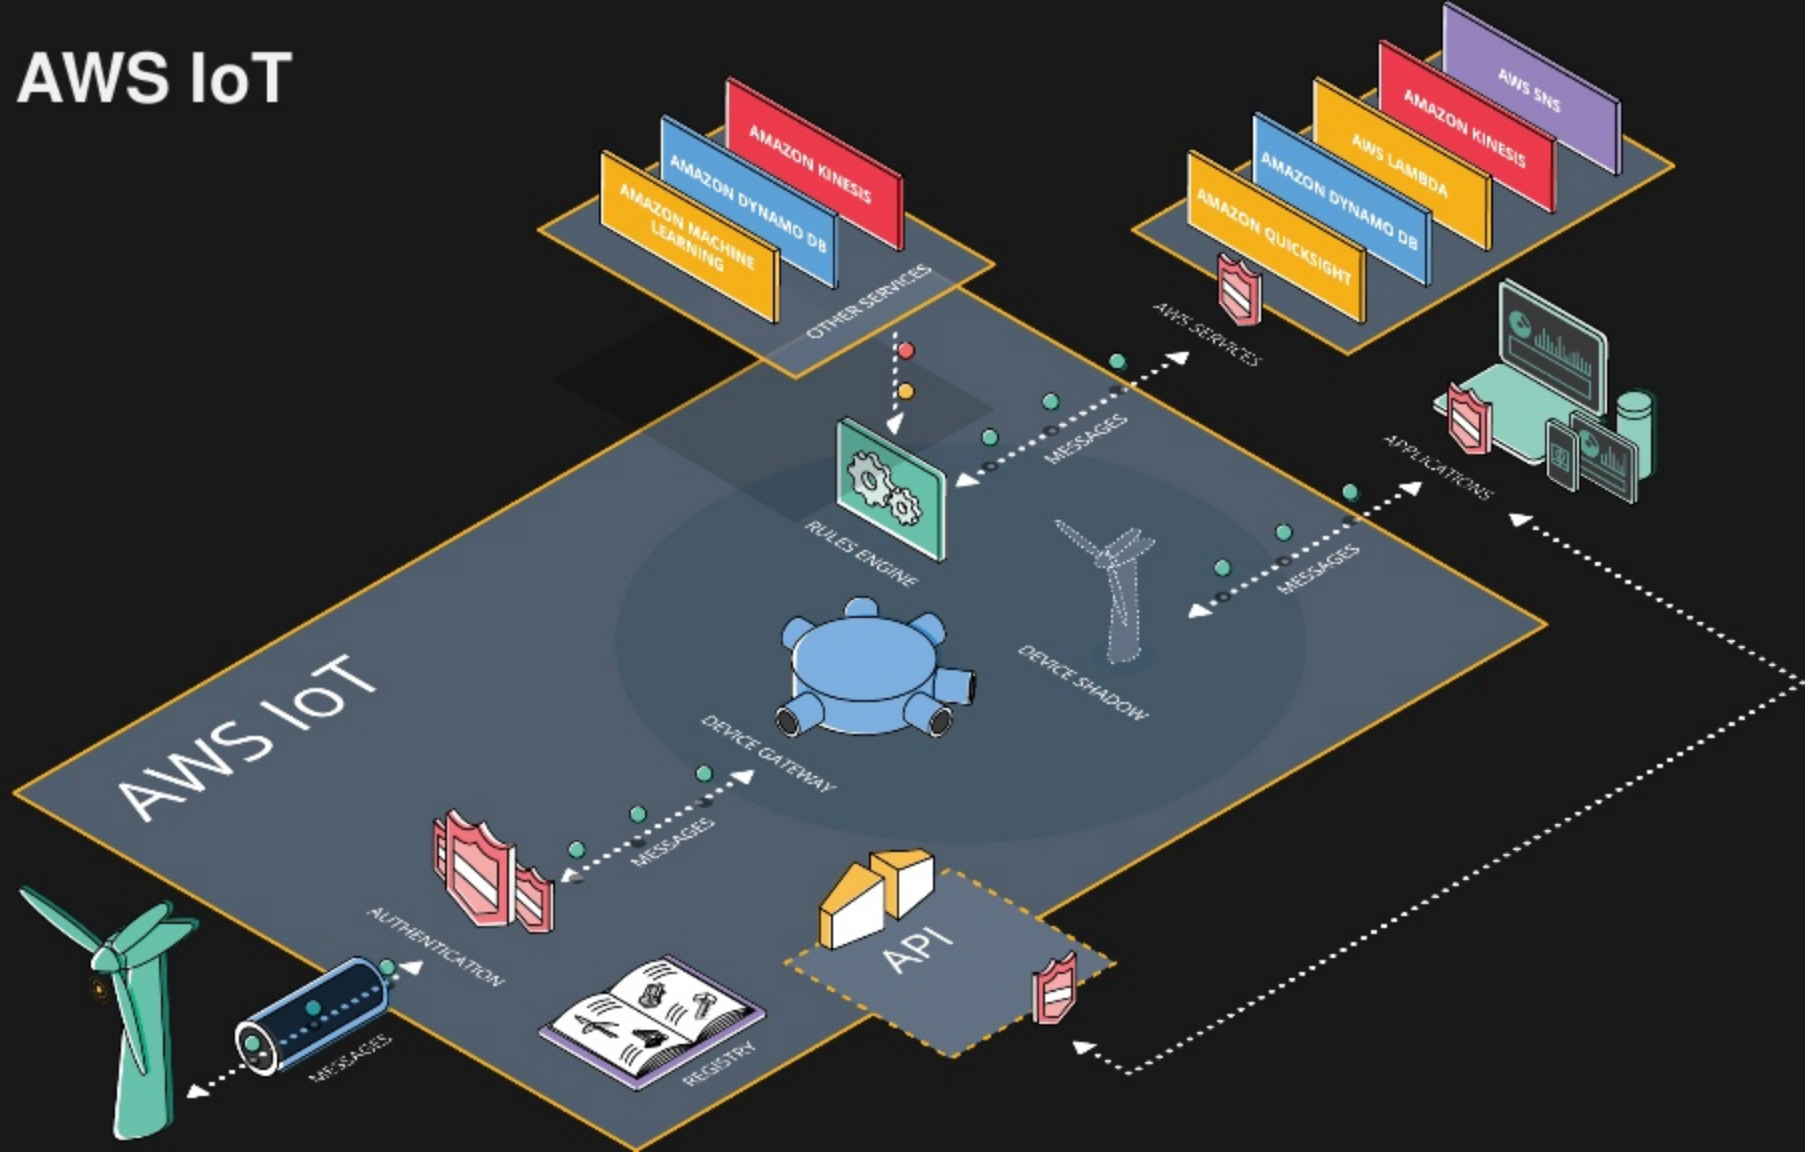
Topic: foo/bar

QoS: 1

Payload: {"foo": "bar"}

**Wrap-up**

# AWS IoT



# Key takeaways

- Messaging
  - Be careful with wide fan out
  - No message ordering guarantees
  - Avoid large fan in
  - WebSockets for Amazon Cognito authentication
- Rules
  - Send data to multiple data stores at the same time
  - Manage device lifecycle events
- Shadows
  - Designed for the real world: poor connectivity, out of order messages
  - Fine-grained control over software rollouts
  - Not ideal for storing time-series analytics data
- Security
  - One cert per device
  - Set fine-grained permissions for devices and Amazon Cognito users
  - Naming conventions can simplify policy management

AWS

S U M M I T

Thank you!

